



AWS  
re:Invent

**A R C 2 0 4 - R**

# Cost optimizing a workload

**Nathan Besh**

Cost Lead, Well-Architected  
Amazon Web Services

# Agenda

Introduction

Setup

Baseline

Cost optimization cycles (analyze & optimize)

Bonus: Well-Architected

# Introduction

- Workshop goal
- Simulate the improvements
- Present, then implement
  - A few slides on the topic
  - Hands-on labs
- Labs are available online
  - Stay together (don't worry if you don't finish)
  - Do them anytime

# Related Breakouts

ARC219 – AWS Cost Management tools for cost & usage optimization (200-300)

ENT206 – Optimize AWS costs and utilization with AWS management tools (100-200)

CMP410 – Save up to 90 percent on CI/CD and test workloads (spot)

ARC209 – Running lean architectures: How to be cost-effective on AWS (Presentation)

CMP207 – Manage, control, and optimize costs with native AWS products, ft. Intuit (Chalk talk)

CMP323 – Optimize performance and cost for your AWS compute (Presentation)

ENT204 – Managing your cloud financials as you scale on AWS (Presentation)

# Environment Setup

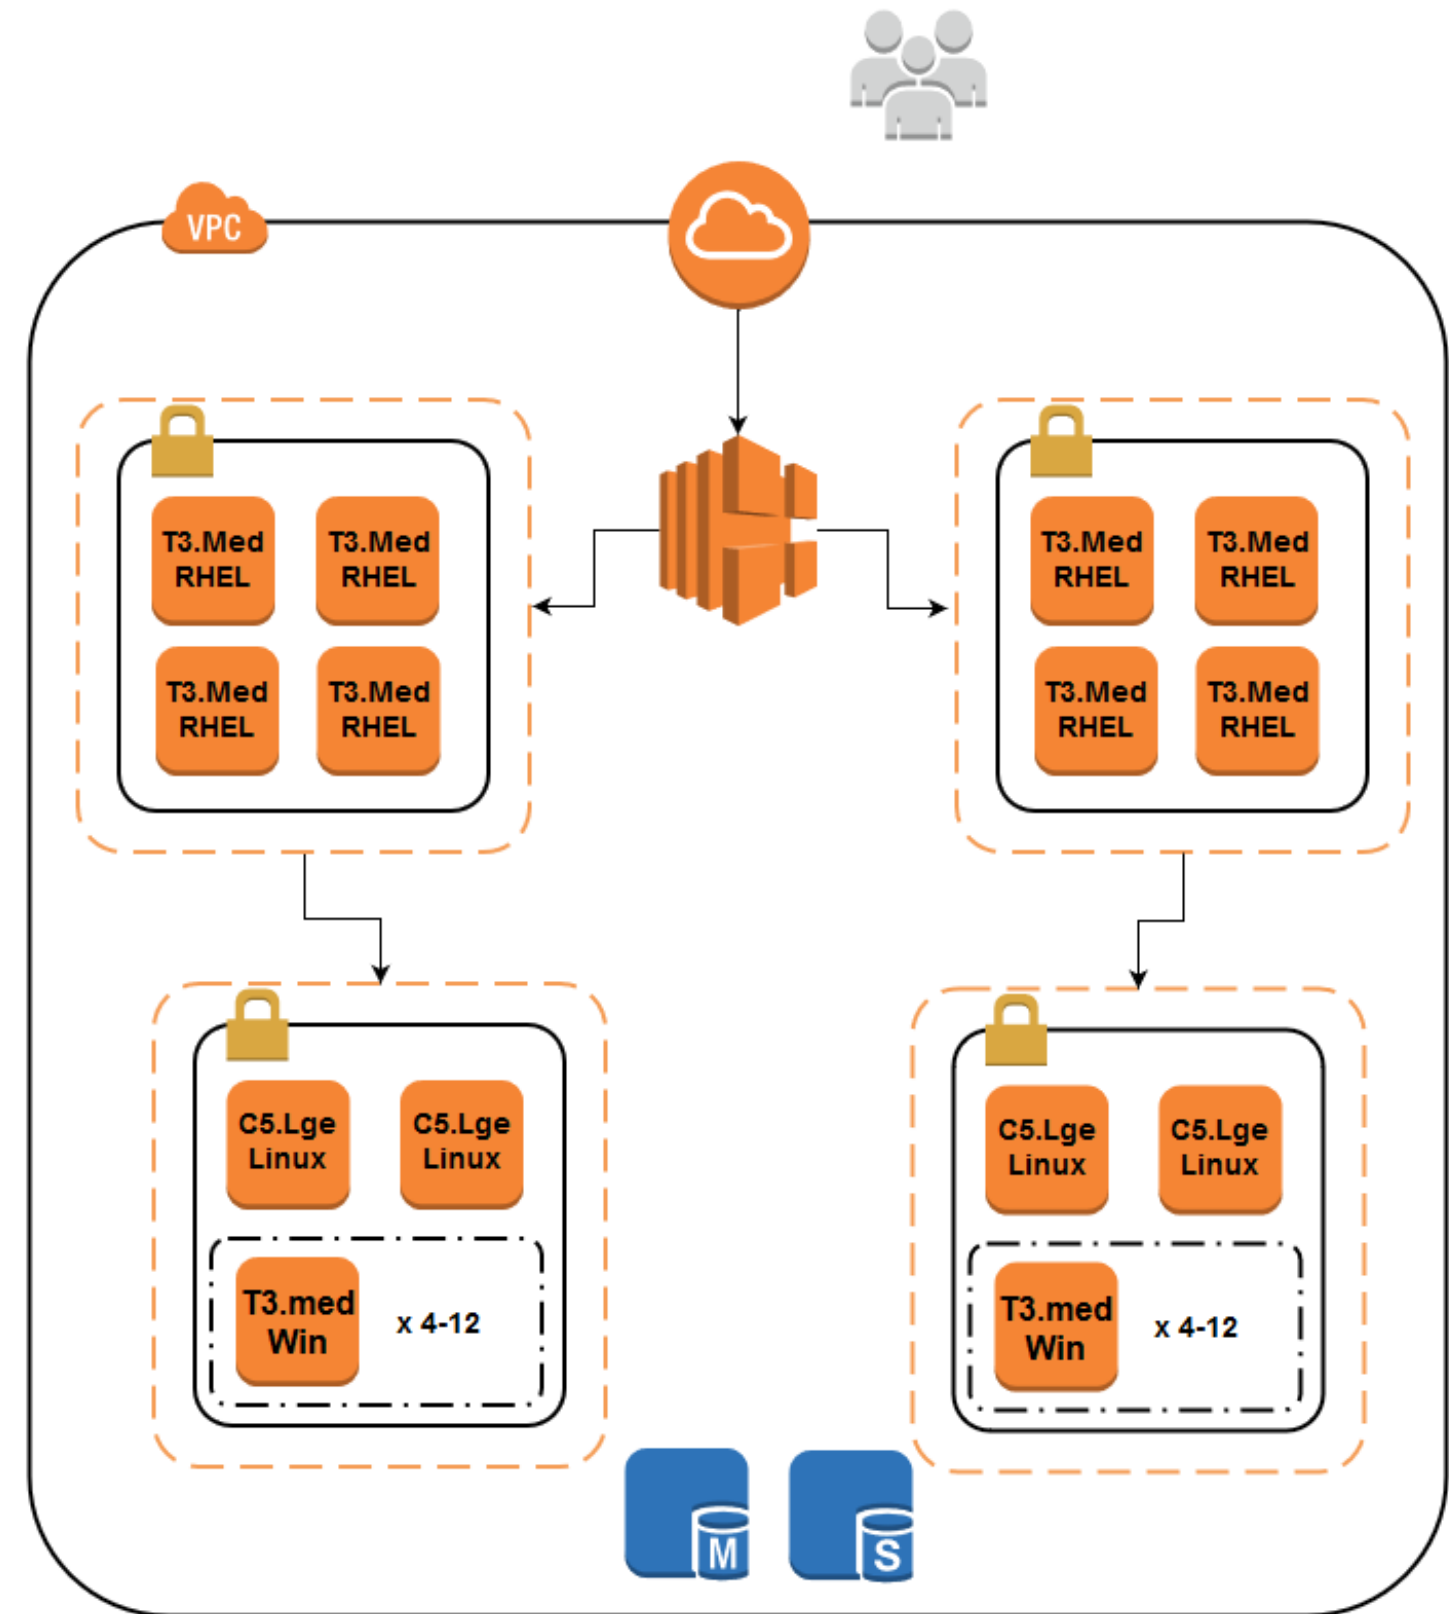
# Setup

## 3-Tier Web App

Front:  
8 x t3.medium RHEL  
Load Balancer

Application:  
4 x c5.large Linux  
4-12 (ASG) x t3.medium Windows

Database: Amazon Aurora



# Tools

- Application log files
- Cost and Usage Report (billing files)
- Amazon Simple Storage Service (Amazon S3)
- Amazon Athena
- Amazon QuickSight



# Process

- Copy application & billing files
- Perform analysis
- Decide action
- **Simulate implementation**
- Confirm optimization

# Lab

<https://wellarchitectedlabs.com/Cost/arc204.html>

Step 1: Setup



# Baseline

# Baseline

- How does my workload demand behave?
- How do my workload resourcing and costs behave?
- What is my workload efficiency?
- **Business focus & high level**
- What does your workload do? (business outcomes)
- What shows these outcomes? (log file)
- Workload resourcing/costs = Cost & Usage Report (CUR)

# Baseline

- Demand profile
- Workload profile
- Workload efficiency (outcomes/\$)

# Lab

<https://wellarchitectedlabs.com/Cost/arc204.html>

Step 2: Baseline



# Cycle #1: Licensing

# Licensing

- What is the cost of a software license?
- What is the cost of running licensed software?
- If software mandates additional resources, those resources are the cost of running the software & should be associated



# Lab

<https://wellarchitectedlabs.com/Cost/arc204.html>

Step 3: Licensing



# M5.xlarge – US-East-1 (N. Virginia) – Shared tenancy

	Linux	RHEL	Windows
Cost/hr	\$0.192	\$0.252	\$0.376

# T3.medium – US-East-1 (N. Virginia) – Shared tenancy

	Linux	RHEL	Windows
Cost/hr	\$0.0416	\$0.1016	\$0.06

# M5.xlarge – US-East-1 (N. Virginia) – Shared tenancy

- What's different when you launch?
  - Storage
  - Network traffic (updates, chatty OS)
- Console defaults, all free-tier eligible

	Linux + 8Gb GP2	RHEL + 10Gb GP2	Windows + 30Gb GP2
Cost/hr	\$0.193096 (\$0.192)	\$0.253370 (\$0.252)	\$0.380110 (\$0.376)

# M5.xlarge – US-East-1 (N. Virginia) – Shared tenancy

- Discounted Pricing Model
- 1-year, no upfront cost, regional/standard

	Linux + 8Gb GP2	RHEL + 10Gb GP2	Windows + 30Gb GP2
Cost/hr	\$0.124096 (\$0.192)	\$0.184370 (\$0.252)	\$0.311110 (\$0.376)

# T3.medium – US-East-1 (N. Virginia) – Shared tenancy

- Reserved Instances & Storage

	Linux	RHEL	Windows
Cost/hr	\$0.027096	\$0.087370	\$0.049110
License + Storage Cost/hr		\$0.060274	\$0.022014
% extra cost		3.2247 x	1.8124 x

# Cycle #2: Storage

# Storage

- Plan for scale
- Find unused storage efficiently
- Find what storage is used for, efficiently



# Storage

- Unused storage = remove the waste
- What storage is used for = another service?

# Lab

<https://wellarchitectedlabs.com/Cost/arc204.html>

## Step 4: Storage



# Cycle #3: Data Transfer

# Data Transfer

- AWS = managed network provider by default
- Networking is a resource (like compute)
- Accurate & transparent allocation drives efficient workloads
- Efficient compute, efficient storage, efficient network
- Ensure data transfer delivers value

# Data Transfer

- Pricing tips
  - Start with the source service
  - Look at VPC pricing (services inside a VPC)
  - **Gateway** VPC endpoints are free (Amazon S3 & Amazon DynamoDB)
  - **Interface** VPC endpoint / AWS PrivateLink is a private connectivity service
- How far from **your** resource is data traveling?
  - LAN = free
  - Remote/Router = paid

# Lab

<https://wellarchitectedlabs.com/Cost/arc204.html>

Step 5: Data Transfer



# Cycle #4: Pricing Models

# Pricing Models

- On Demand (OD)
- Spot: Instances, defined duration, fleet
- Reserved Instances (RI)
- Savings Plan (SP)




# Pricing Models: Flexibility & Discounts

Flexibility →

Standard RI	Regional RI (AZ)	Size Flex (AWS Linux) (AZ, Size)	Convertible RI (AZ, Size, Family, OS, tenancy)
Instance Saving Plan (AZ, Size, OS, tenancy)			Compute Savings Plan (AZ, Size, Family, OS, tenancy, Region, Service)
Highest Discount Up to 72%			High Discount Up to 66%

# Coverage vs. Utilization (RI)

- 100 instances of Linux, m5.large, Linux, 1-year commit, no upfront cost
  - On Demand: \$84,096
  - RI Discount = 36%
  - 50% coverage (50 RI/50 OD): \$68,766
  - 25% coverage (25 RI/75 OD): \$76,431
  - Purchase 128 RIs (need 100!): \$68,398
  - Purchase 143 RIs (need 100!): \$76,413
- 

# Coverage vs. Utilization (SP)

- 100 instances of Linux, m5.large, Linux, 1-year commit, no upfront cost
- On-Demand: \$84,096
- 100% coverage instance SP @ 36%: \$53,436
- 100% coverage compute SP @ 26%: \$62,196

74 Instances running: \$62,231 On-Demand
--

64 Instances running: \$53,821 On-Demand
--

Instance SP: \$385 (SAVING)	Compute SP: \$8,375 (LOSS)
-----------------------------	----------------------------

# Pricing Models = Commitment = Investment

- What is the risk?
  - Risk = investment over the full commitment
  - How quickly can I pay off the **full** commitment with savings?
  - Formula = (Monthly discounted amount x 12) / (Monthly On Demand)
  - Breakeven = wrong
- What is the return?
  - Savings from discount
  - Variable!
- What if it is unused?
  - DON'T Worry!

# Lab

<https://wellarchitectedlabs.com/Cost/arc204.html>

## Step 6: Pricing Models



# Bonus: Well-Architected Tool

# Cost Optimization: Where should I focus?

- Well-Architected Tool
- 9 Questions
  - Different focus areas of cost optimization
  - Best practices in each area
- Open discussion within your organization
  - No right or wrong
  - How you are achieving the outcomes
- Free learning tool for your organization
- Track your progress and improvement

# Lab

<https://wellarchitectedlabs.com/Cost/arc204.html>


Step 7: Well-Architected Tool






# Learn to architect with AWS Training and Certification

Resources created by the experts at AWS to propel your organization and career forward



Free foundational to advanced digital courses cover AWS services and teach architecting best practices



Classroom offerings, including Architecting on AWS, feature AWS expert instructors and hands-on labs



Validate expertise with the **AWS Certified Solutions Architect - Associate** or **AWS Certification Solutions Architect - Professional** exams

Visit [aws.amazon.com/training/path-architecting/](https://aws.amazon.com/training/path-architecting/)

# Thank you!

**Nathan Besh**

[natbesh@amazon.com](mailto:natbesh@amazon.com)



Please complete the session  
survey in the mobile app.