



AWS
re:Invent

CON413-R

Move your Machine Learning workloads on Amazon Elastic Kubernetes Service (EKS)

Jiaxin Shan

Software Engineer
AWS

Mike Stefaniak





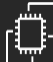






Product Manager
AWS

Arun Gupta

Open Source Technologist
AWS


Machine Learning Stack on EKS

ML Frameworks +
Infrastructure


FRAMEWORKS			INTERFACES		INFRASTRUCTURE						
 TensorFlow	 mxnet	PYTORCH	 GLUON	 Keras	 EC2 P3 & P3dn	 EC2 G4	 EC2 C5	 FPGAs	FSx  Lustre	 EFA	 Inferentia

Machine Learning Stack on EKS


Container Platform




Amazon EKS



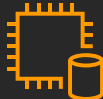
Optimized GPU AMI



Auto Scaling














Deep Learning Container



FSx CSI Plugin

ML Frameworks + Infrastructure

FRAMEWORKS			INTERFACES		INFRASTRUCTURE						
 TensorFlow	 mxnet	PYTORCH	 GLUON	 Keras	 EC2 P3 & P3dn	 EC2 G4	 EC2 C5	 FPGAs	 FSx Lustre	 EFA	 Inferentia

Machine Learning Stack on EKS


ML Platform



Container Platform



ML Frameworks + Infrastructure




Machine Learning Stack on EKS


ML Platform



Container Platform

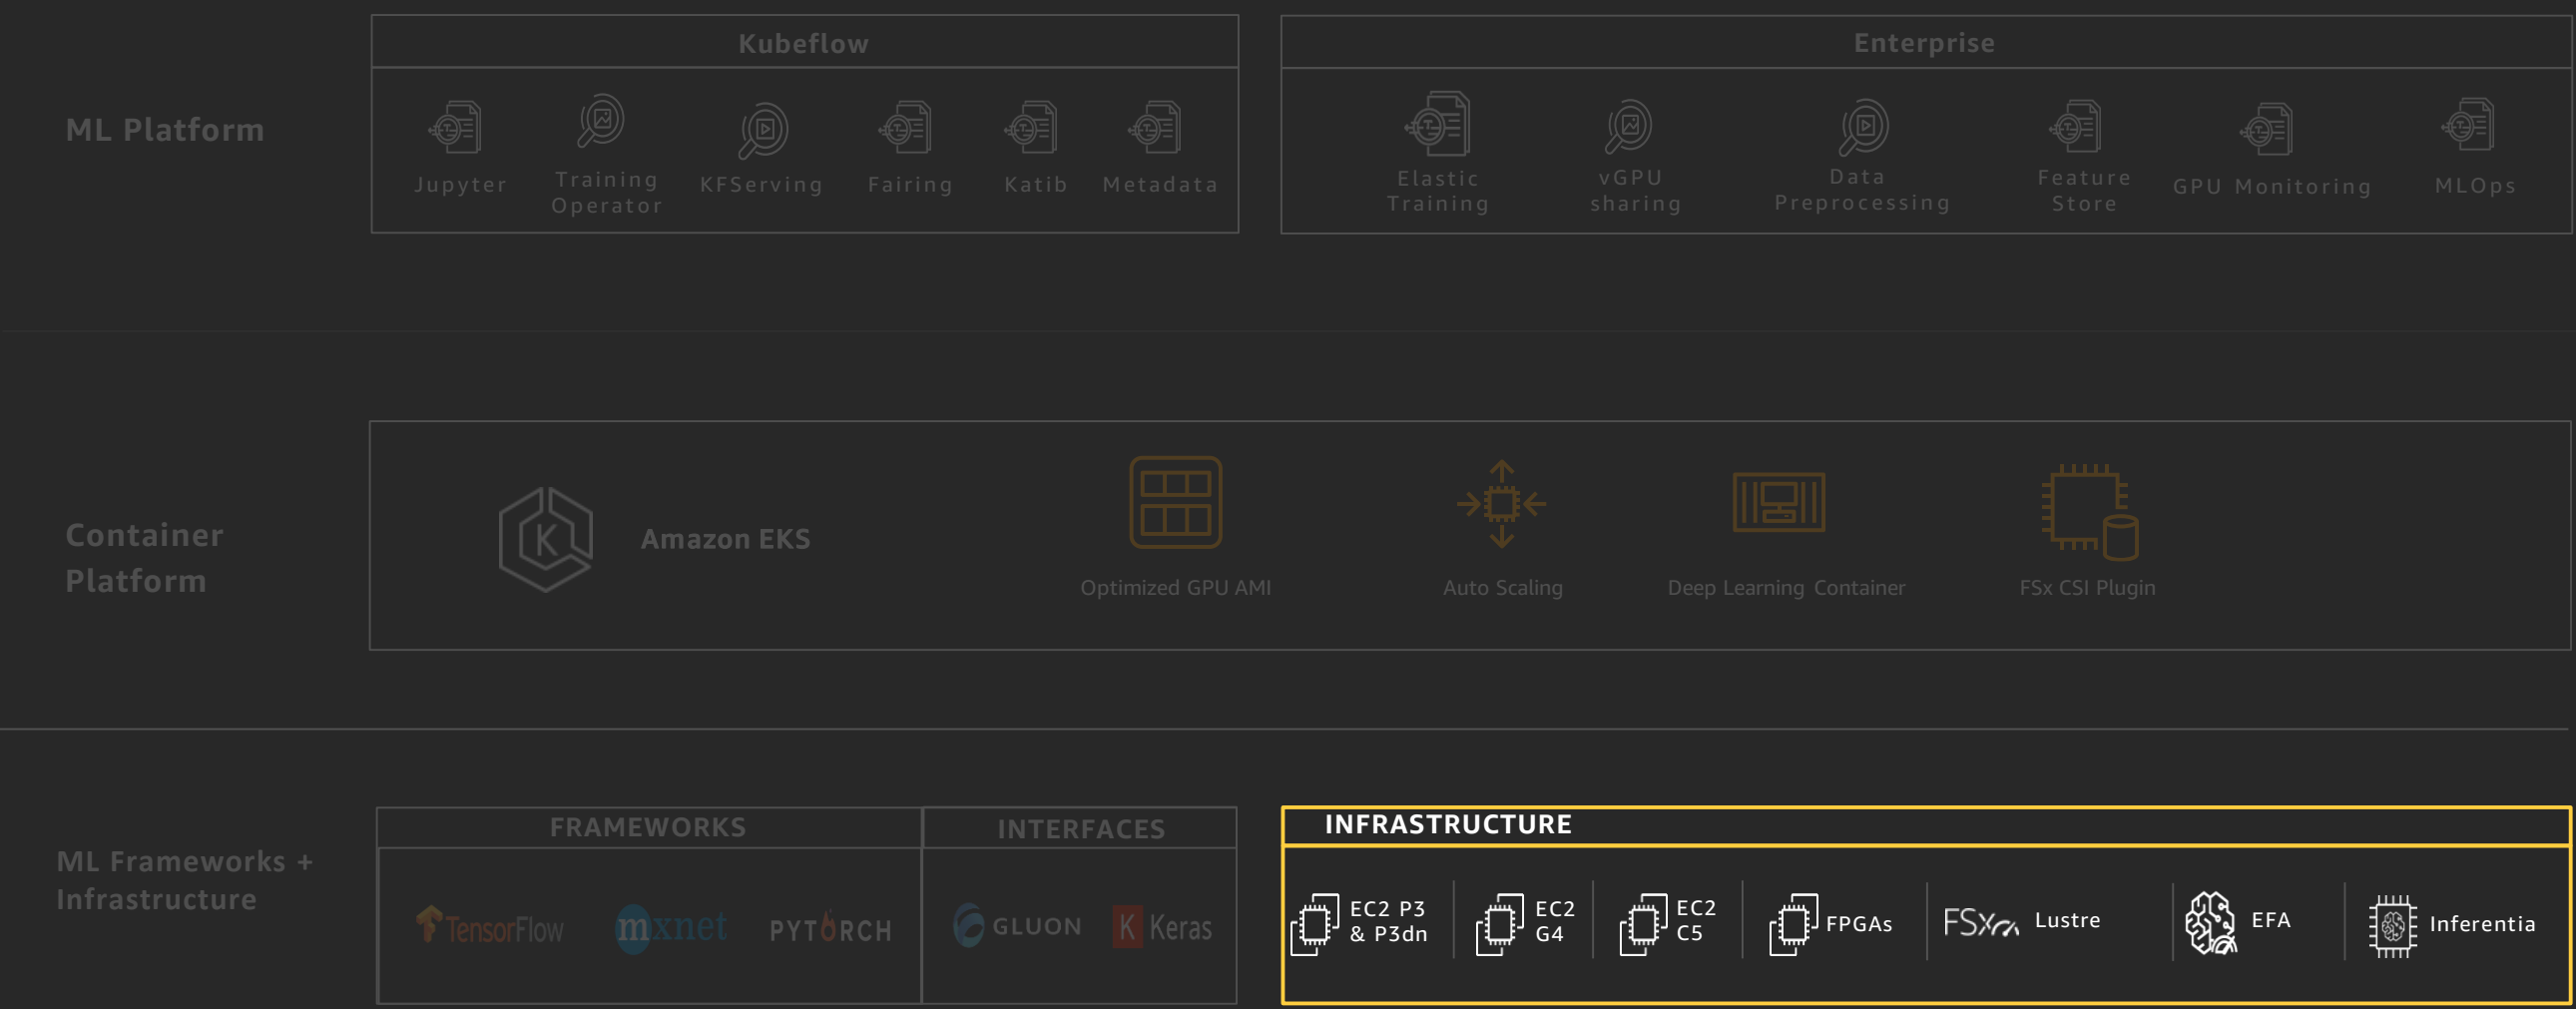


ML Frameworks + Infrastructure





Infrastructure

Machine Learning Stack on EKS




ML Infrastructure supported by EKS

Compute



Accelerated Instance Types

Networking



Elastic Fabric
Adapter

Storage




Amazon Elastic File
System



Amazon FSx
for Lustre

Inference

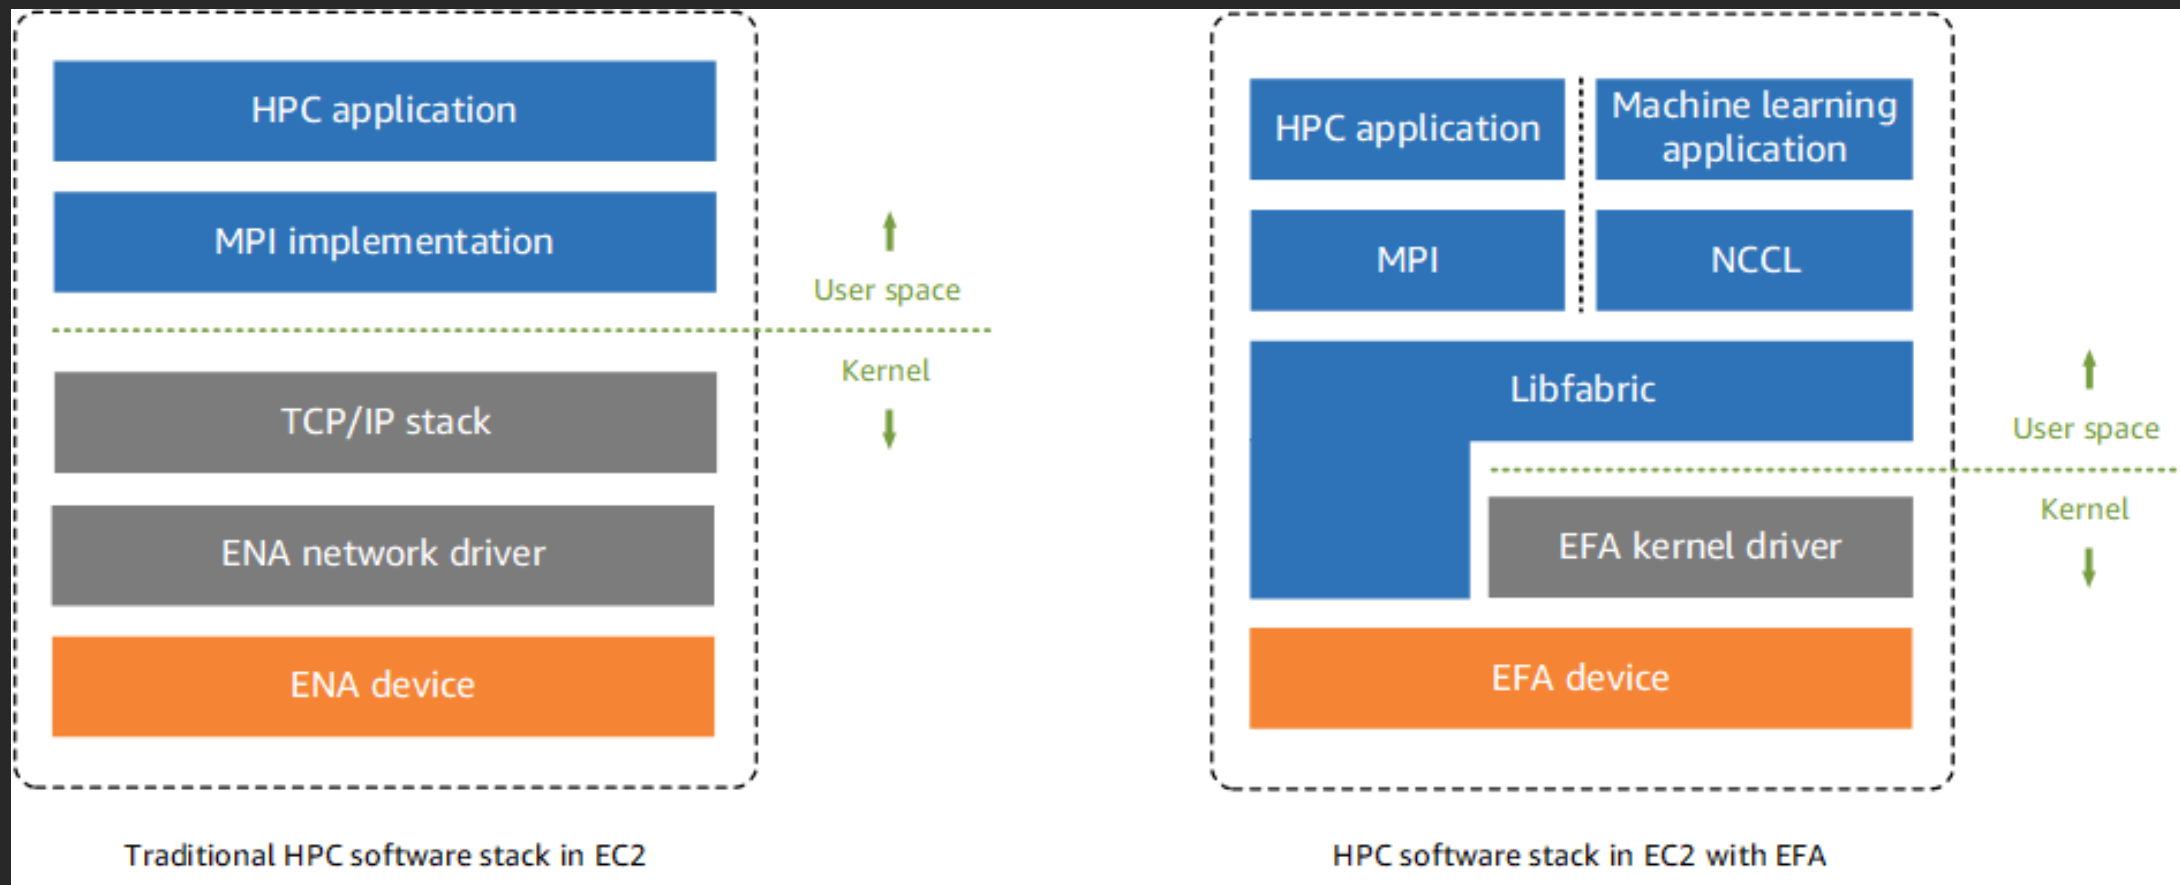


AWS Inferentia
(EKS support
coming soon)

Elastic Fabric Adapter

Use MPI or NCCL to interface with Libfabric API (bypass OS kernel)

Reduce overhead and enable HPC applications to run more efficiently



Distributed Storage



Amazon Elastic
File System

Cloud-native, shared NFS storage
solution

Mount shared filesystem into pods

Share datasets or models across teams




Amazon FSx
for Lustre

File system optimized for ML and HPC
workloads

Native integration with S3 for stored
training data

Read/write data up to hundreds of GB/sec
of throughput and millions of IOPS

AWS Inferentia (EKS support coming soon)



Built for high throughput,
low latency applications



High performance
alternative to a GPU



Order of magnitude
cost reduction

Each chip provides hundreds of TOPS (tera operations per second) of inference throughput to allow complex models to make fast predictions.


For even more performance, multiple Inferentia chips can be used together to drive thousands of TOPS of throughput

Supports all major frameworks used in the deep learning community including TensorFlow, Apache MXNet, and PyTorch, as well as models that use the ONNX format


Kubernetes

Machine Learning Stack on EKS


ML Platform




Container Platform




ML Frameworks + Infrastructure




Why Machine Learning on Kubernetes?



Composability



Portability



Scalability



Amazon EKS-Optimized GPU AMI



Built on top of the standard Amazon EKS-Optimized AMI


Includes packages to support Amazon P2/P3/G3/G4 instances

- NVIDIA drivers
- nvidia-docker2 package
- nvidia-container-runtime (as default runtime)


GPU Clock Optimization

AWS Deep Learning Containers

Optimized and customizable containers for deep learning environments



Pre-packaged
Docker container
images
fully configured
and validated



Best performance
and scalability
without tuning



Works with
Amazon EKS,
Amazon ECS,
and Amazon EC2

KEY FEATURES

Customizable
container images



Single and multi-node
training and inference


Cluster Autoscaler Improvements

- Add GPU Support
 - [autoscaler#1584](#) Move GPULabel and GPUPypes to cloud provider → GPU autoscaling supported for AWS
 - [autoscaler#1589](#) Consider GPU utilization in scaling down → GPU scale down performance optimization
- Prevent CA from removing a node with ML training job running
 - Annotate job "cluster-autoscaler.kubernetes.io/safe-to-evict": "false"
- Recommended to create GPU node group per AZ
 - Improve network communication performance
 - Prevent ASG rebalancing


Kubeflow

Machine Learning Stack on EKS


ML Platform



Container Platform



ML Frameworks + Infrastructure

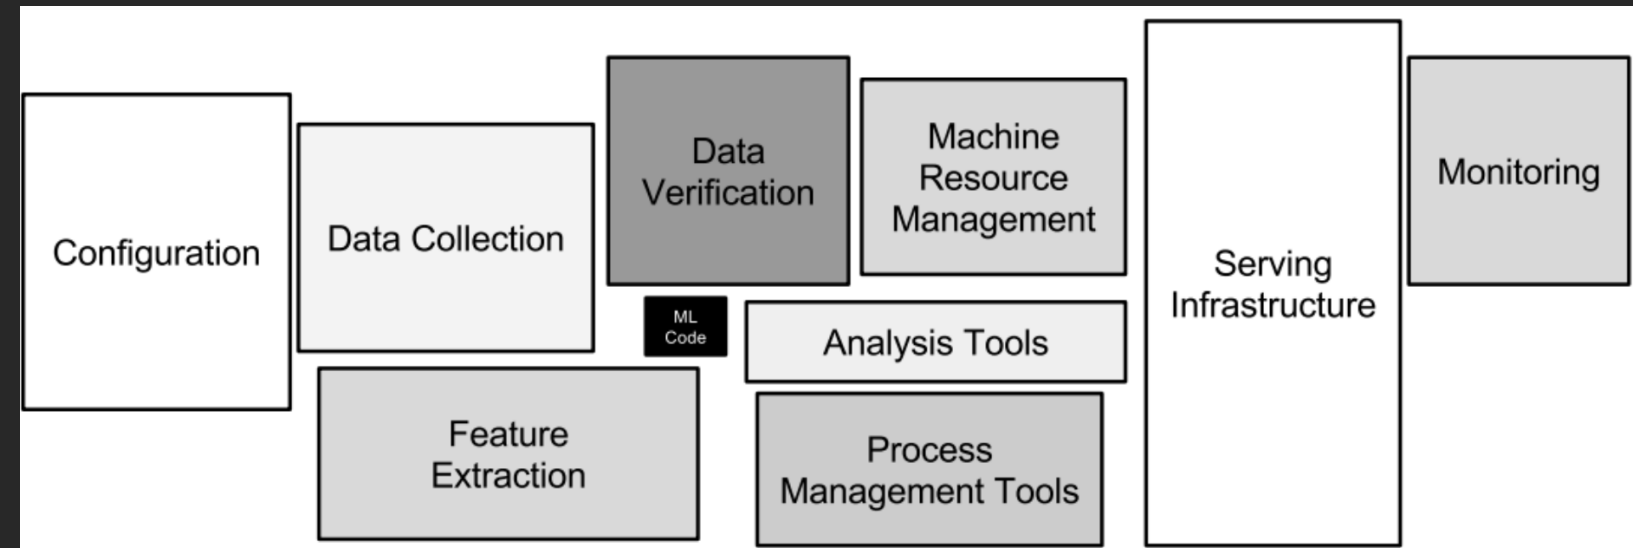


What is Kubeflow

Containerized machine learning platform

Makes it easy to develop, deploy, and manage portable, scalable end-to-end ML workflows on k8s

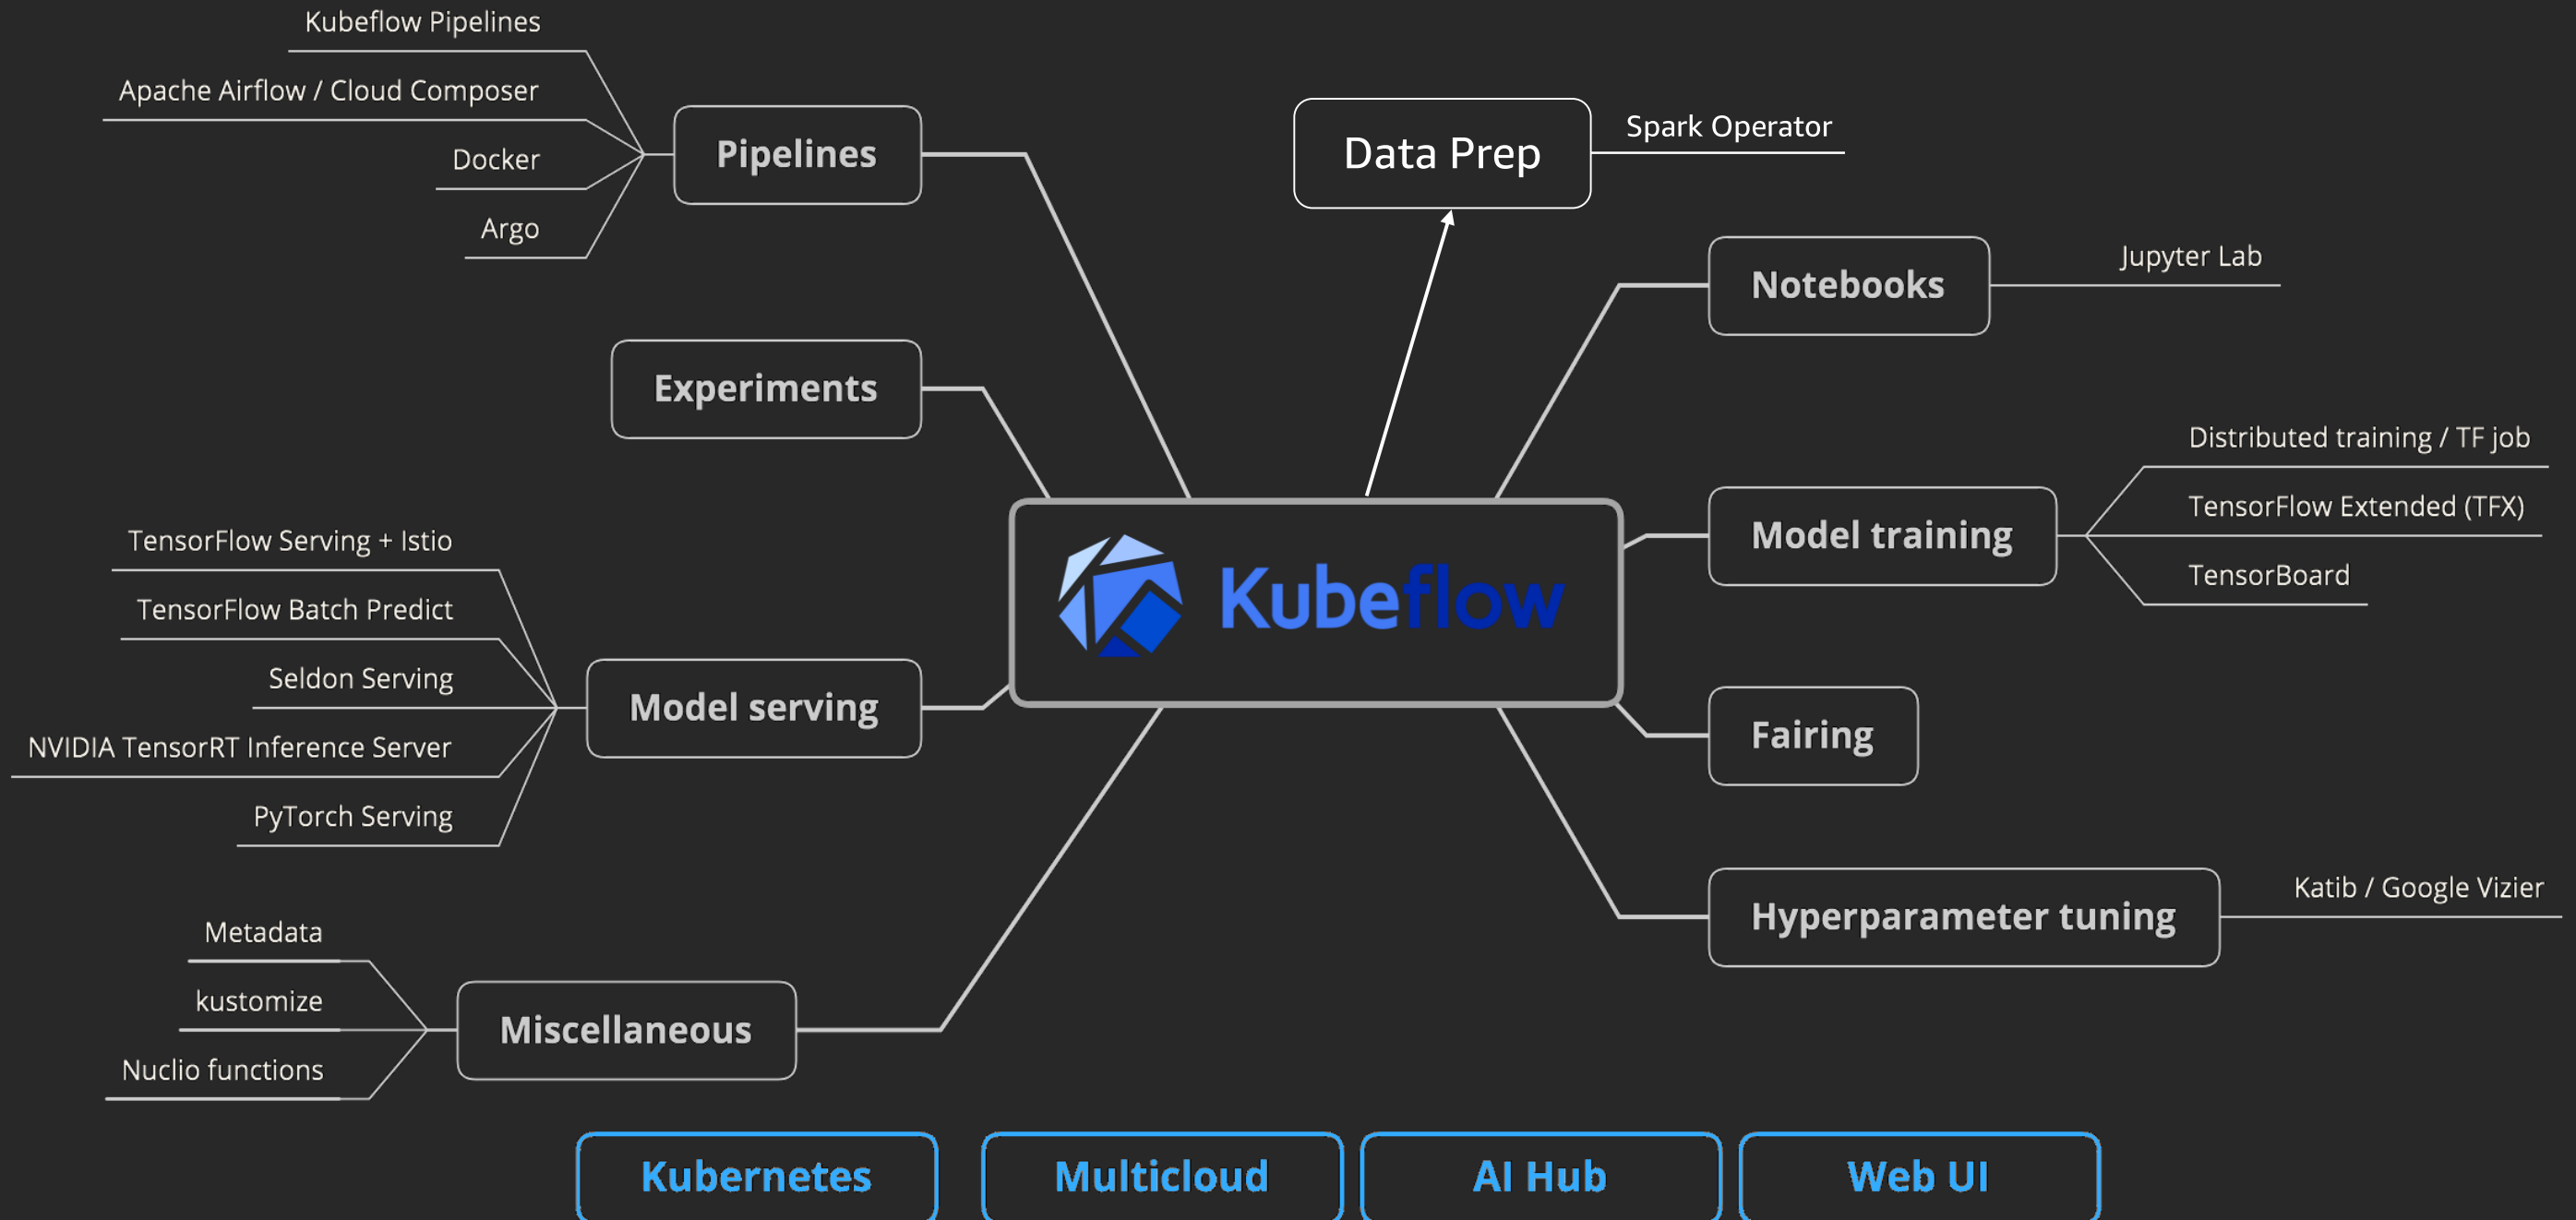
“Toolkit” – loosely coupled tools and blueprints for ML



End to End ML workflow – ML code is only a small component

<https://papers.nips.cc/paper/5656-hidden-technical-debt-in-machine-learning-systems.pdf>

Kubeflow Components

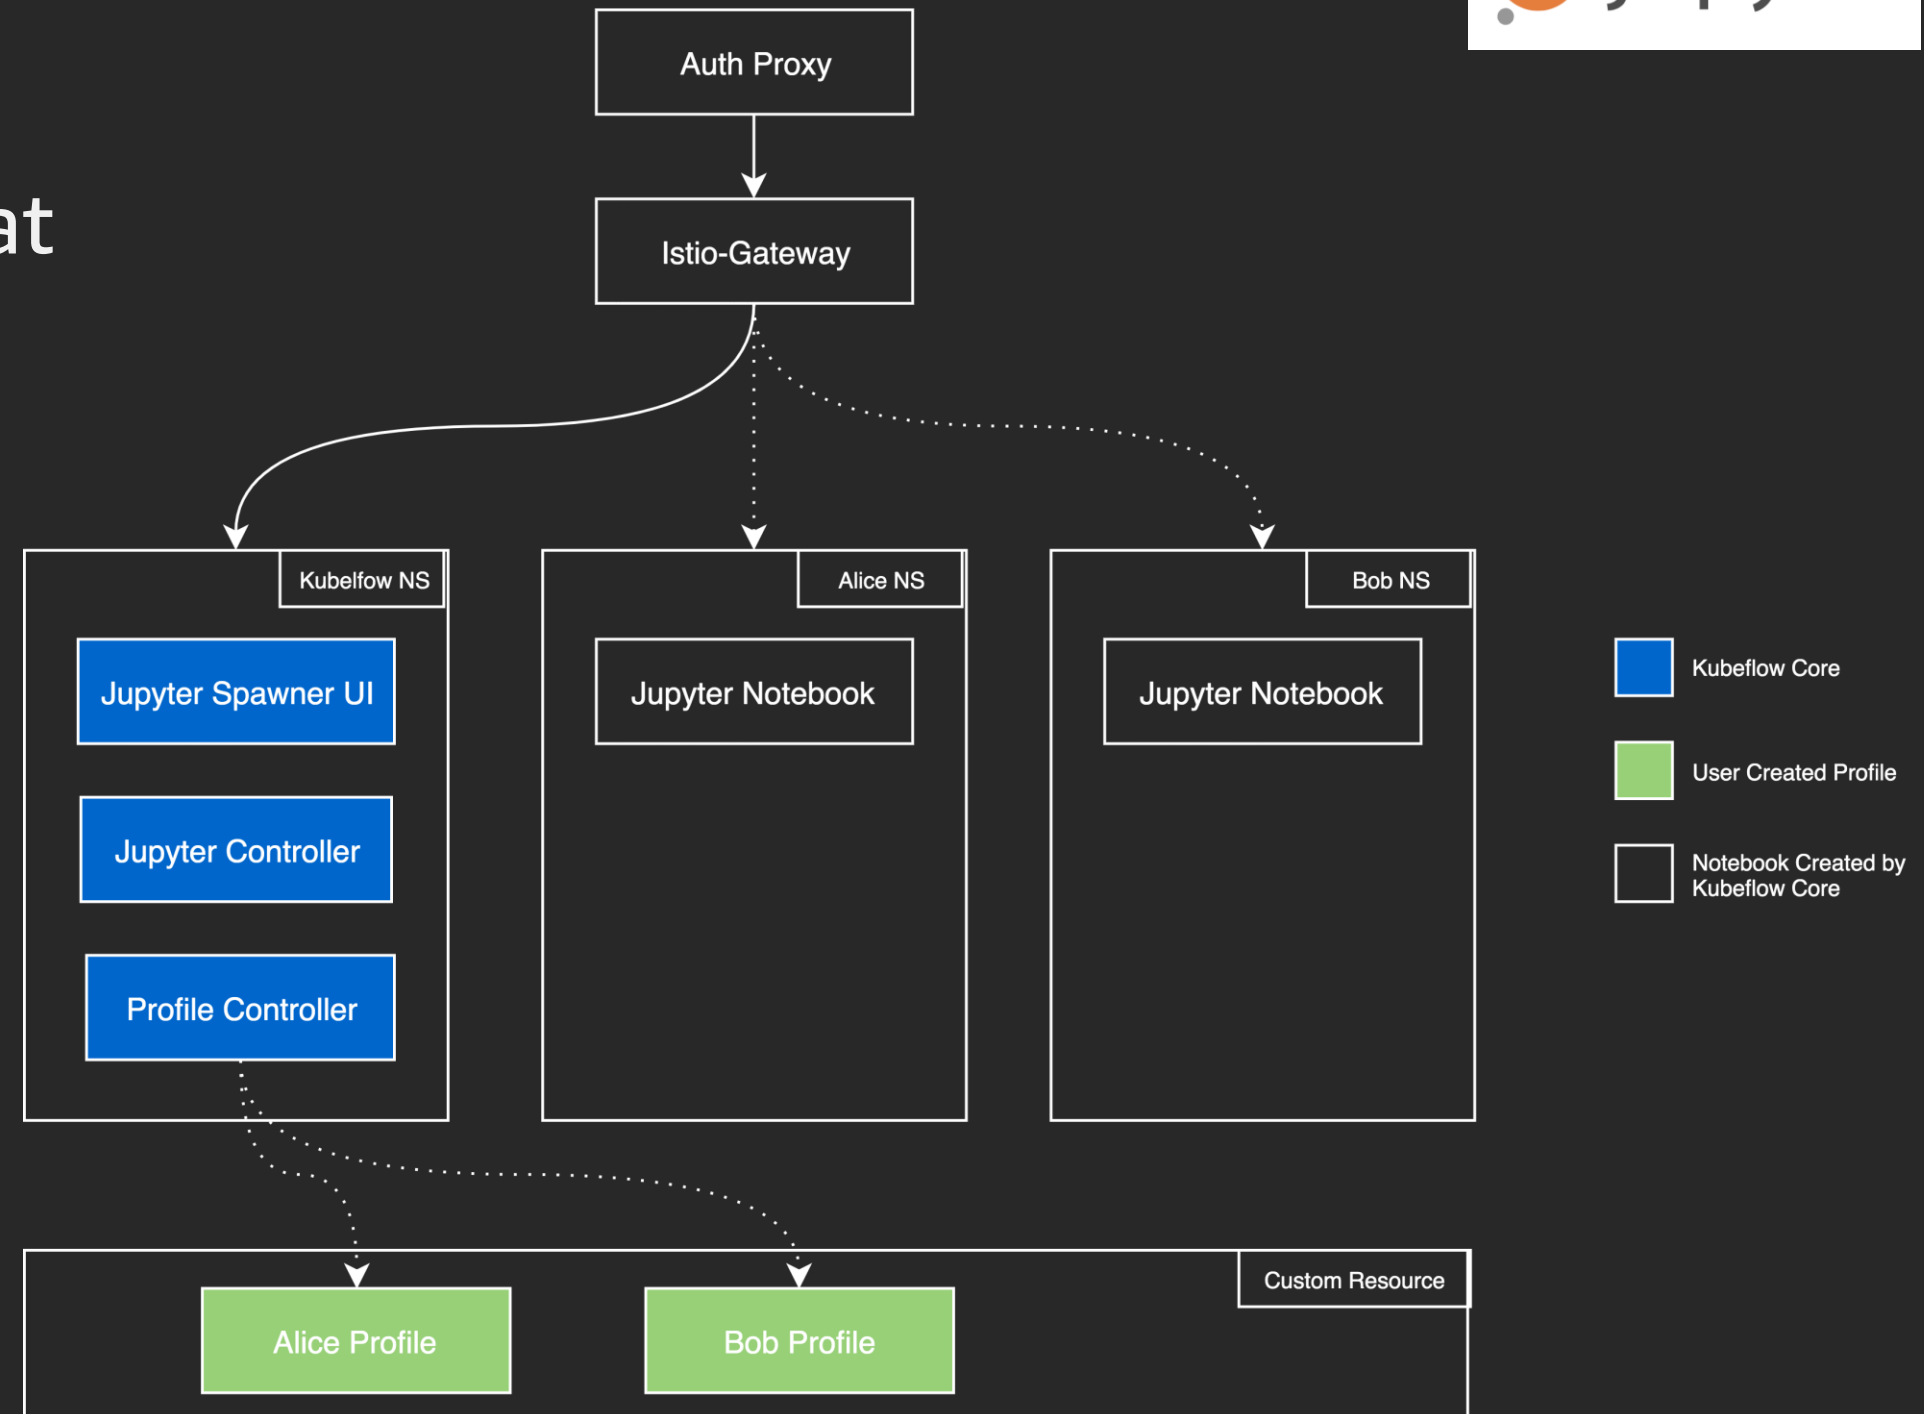


Jupyter Notebook



Create and share docs that contain live code, equations, visualizations, and narrative text

- UI to manage notebooks
- Integrate with RBAC/IAM
- Ingress/Service Mesh



Fairing

Python SDK to build, train and deploy ML models

- Easily package ML training jobs  

- Train ML models from notebook to k8s

- Streamline the model development process

Setup Kubeflow Fairing for training and prediction

```
from kubeflow import fairing
from kubeflow.fairing import TrainJob
from kubeflow.fairing.backends import KubeflowAWSBackend

from kubeflow import fairing

FAIRING_BACKEND = 'KubeflowAWSBackend'

AWS_ACCOUNT_ID = fairing.cloud.aws.guess_account_id()
AWS_REGION = 'us-west-2'
DOCKER_REGISTRY = '{}.dkr.ecr.{}.amazonaws.com'.format(AWS_ACCOUNT_ID, AWS_REGION)
S3_BUCKET = 'kubeflow-pipeline-data'

import importlib

if FAIRING_BACKEND == 'KubeflowAWSBackend':
    from kubeflow.fairing.builders.cluster.s3_context import S3ContextSource
    BuildContext = S3ContextSource(
        aws_account=AWS_ACCOUNT_ID, region=AWS_REGION,
        bucket_name=S3_BUCKET
    )

BackendClass = getattr(importlib.import_module('kubeflow.fairing.backends'), FAIRING_BACKEND)
```

Train an XGBoost model remotely on Kubeflow

```
from kubeflow.fairing import TrainJob
train_job = TrainJob(HousingServe, input_files=['ames_dataset/train.csv', "requirements.txt"],
                     docker_registry=DOCKER_REGISTRY,
                     backend=BackendClass(build_context_source=BuildContext))
train_job.submit()
```

Deploy the trained model to Kubeflow for prediction

```
from kubeflow.fairing import PredictionEndpoint
endpoint = PredictionEndpoint(HousingServe, input_files=['trained_ames_model.dat', "requirements.txt"],
                             docker_registry=DOCKER_REGISTRY,
                             service_type='ClusterIP',
                             backend=BackendClass(build_context_source=BuildContext))

endpoint.create()
```


Katib – Hyperparameter Tuning

Hyperparameter are parameters external to the model to control the training, e.g. learning rate, batch size, epochs

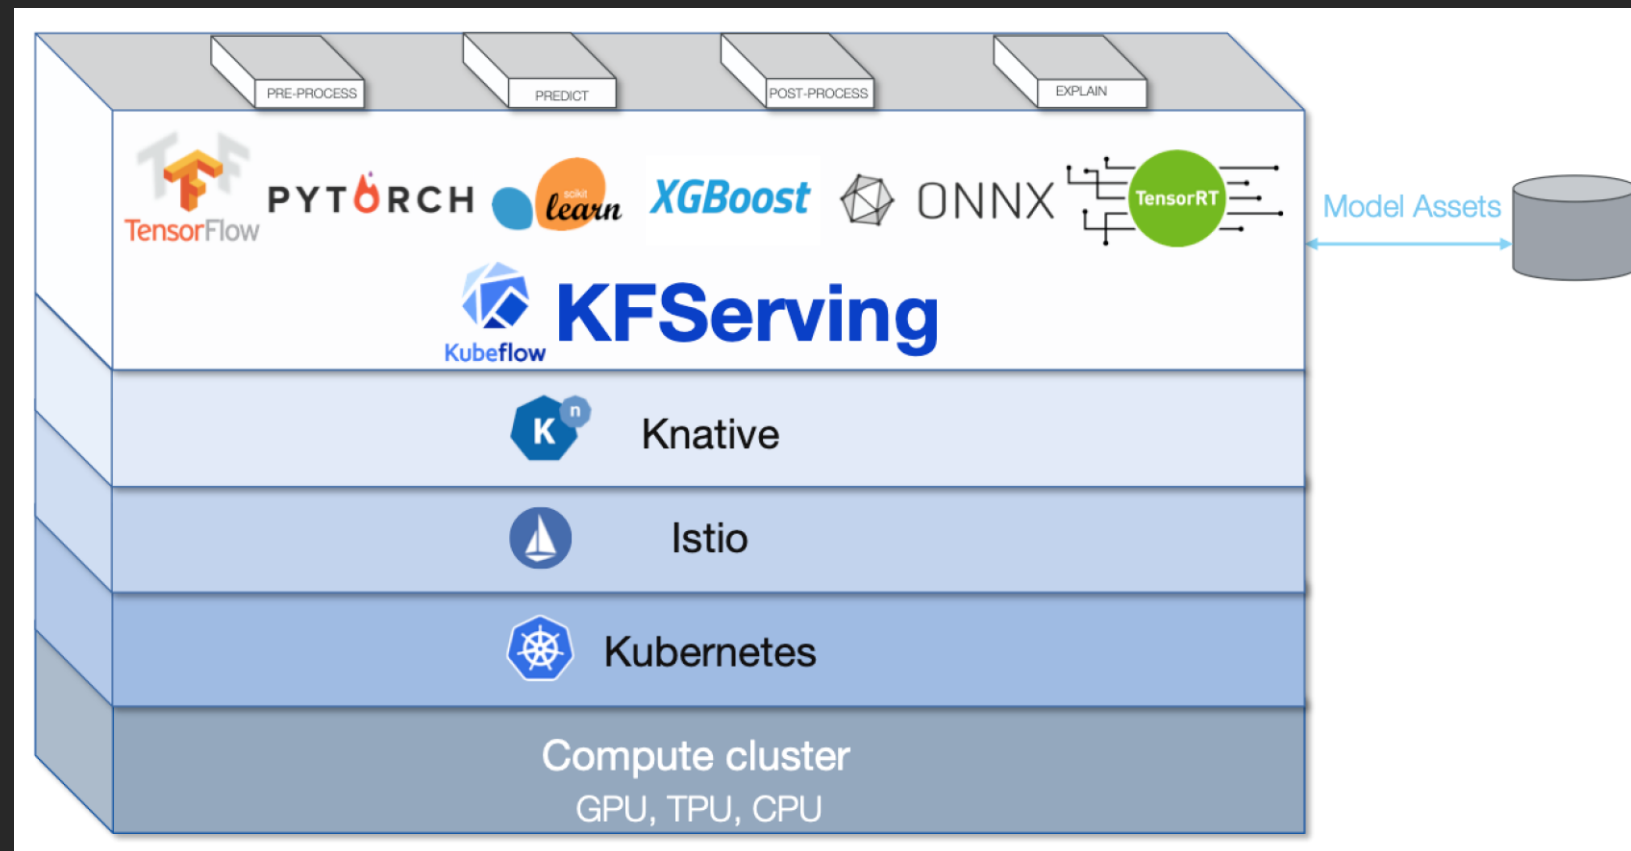
Tuning finds a set of hyperparameters that optimizes an objective function, e.g. find the optimal batch size and learning rate to maximize prediction accuracy

```
Hyperparameters
19  parameters:
20    - name: --lr
21      parameterType: double
22      feasibleSpace:
23        min: "0.01"
24        max: "0.03"
25    - name: --num-layers
26      parameterType: int
27      feasibleSpace:
28        min: "2"
29        max: "5"
30    - name: --optimizer
31      parameterType: categorical
32      feasibleSpace:
33        list:
34          - sgd
35          - adam
36          - ftrl
```

trialName	Validation-accuracy	accuracy	--lr	--num-layers	--optimizer
random-experiment-rfwwbnsd	0.974920	0.984844	0.013831565266960293	4	sgd
random-experiment-vxgwlgqq	0.113854	0.116646	0.024225789898529138	4	ftrl
random-experiment-wclrwlcq	0.979697	0.998437	0.021916171239020756	4	sgd
random-experiment-7lsc4pwb	0.113854	0.115312	0.024163810384272653	5	ftrl
random-experiment-86vv9vgv	0.963475	0.971562	0.02943228249244735	3	adam
random-experiment-jh884cxz	0.981091	0.999219	0.022372025623908262	2	sgd
random-experiment-sgtwhrgz	0.980693	0.997969	0.016641686851083654	4	sgd
random-experiment-c6vvz6dv	0.980792	0.998906	0.0264125850165842	3	sgd
random-experiment-vqs2xmfj	0.113854	0.105313	0.026629394628228185	4	ftrl
random-experiment-bv8lsh2m	0.980195	0.999375	0.021769570793012488	2	sgd
random-experiment-7vbnqc7z	0.113854	0.102188	0.025079750575740783	4	ftrl
random-experiment-kwj9drmg	0.979498	0.995469	0.014985919312945063	4	sgd

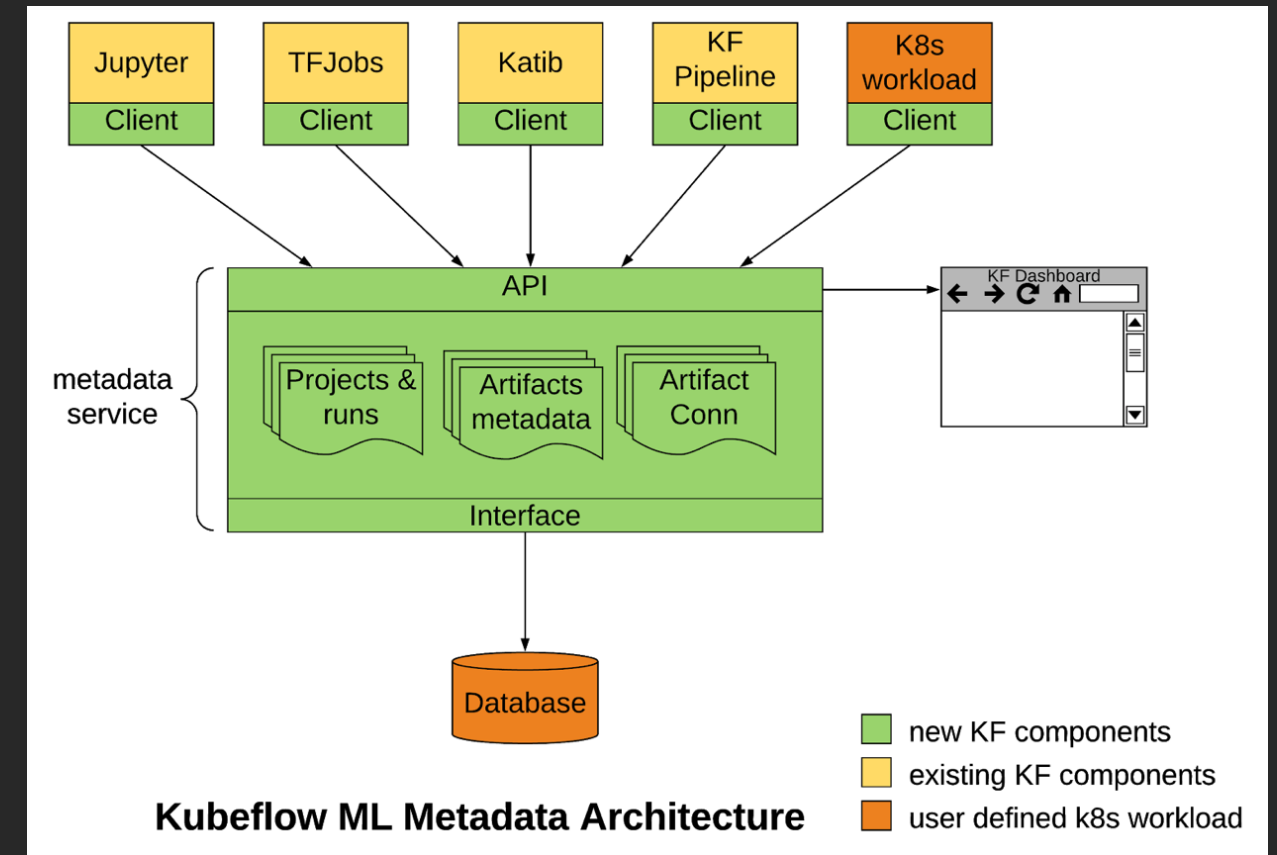
KFServing: Model serving and management

- Provides a Kubernetes CRD for serving ML models on arbitrary frameworks
- Encapsulates the complexity of autoscaling, networking and server configuration to bring features like scale to zero, transformations, and canary rollouts to your deployments
- Enables a simple, pluggable, and complete story for your production ML inference server by providing prediction, pre-processing, post-processing and explainability



Metadata – Model Tracking

- Metadata schema to track artifacts related to execution contexts
- Metadata API for storing and retrieving metadata
- Client libraries for end-users to interact with the Metadata service from their Notebooks or Pipelines code



```
def create_execution(self):
    workspace = metadata.Workspace(
        # Connect to metadata-service in namespace kubeflow in k8s cluster.
        backend_url_prefix="metadata-service.kubeflow:8080",
        name="xgboost-synthetic",
        description="workspace for xgboost-synthetic artifacts and executions")

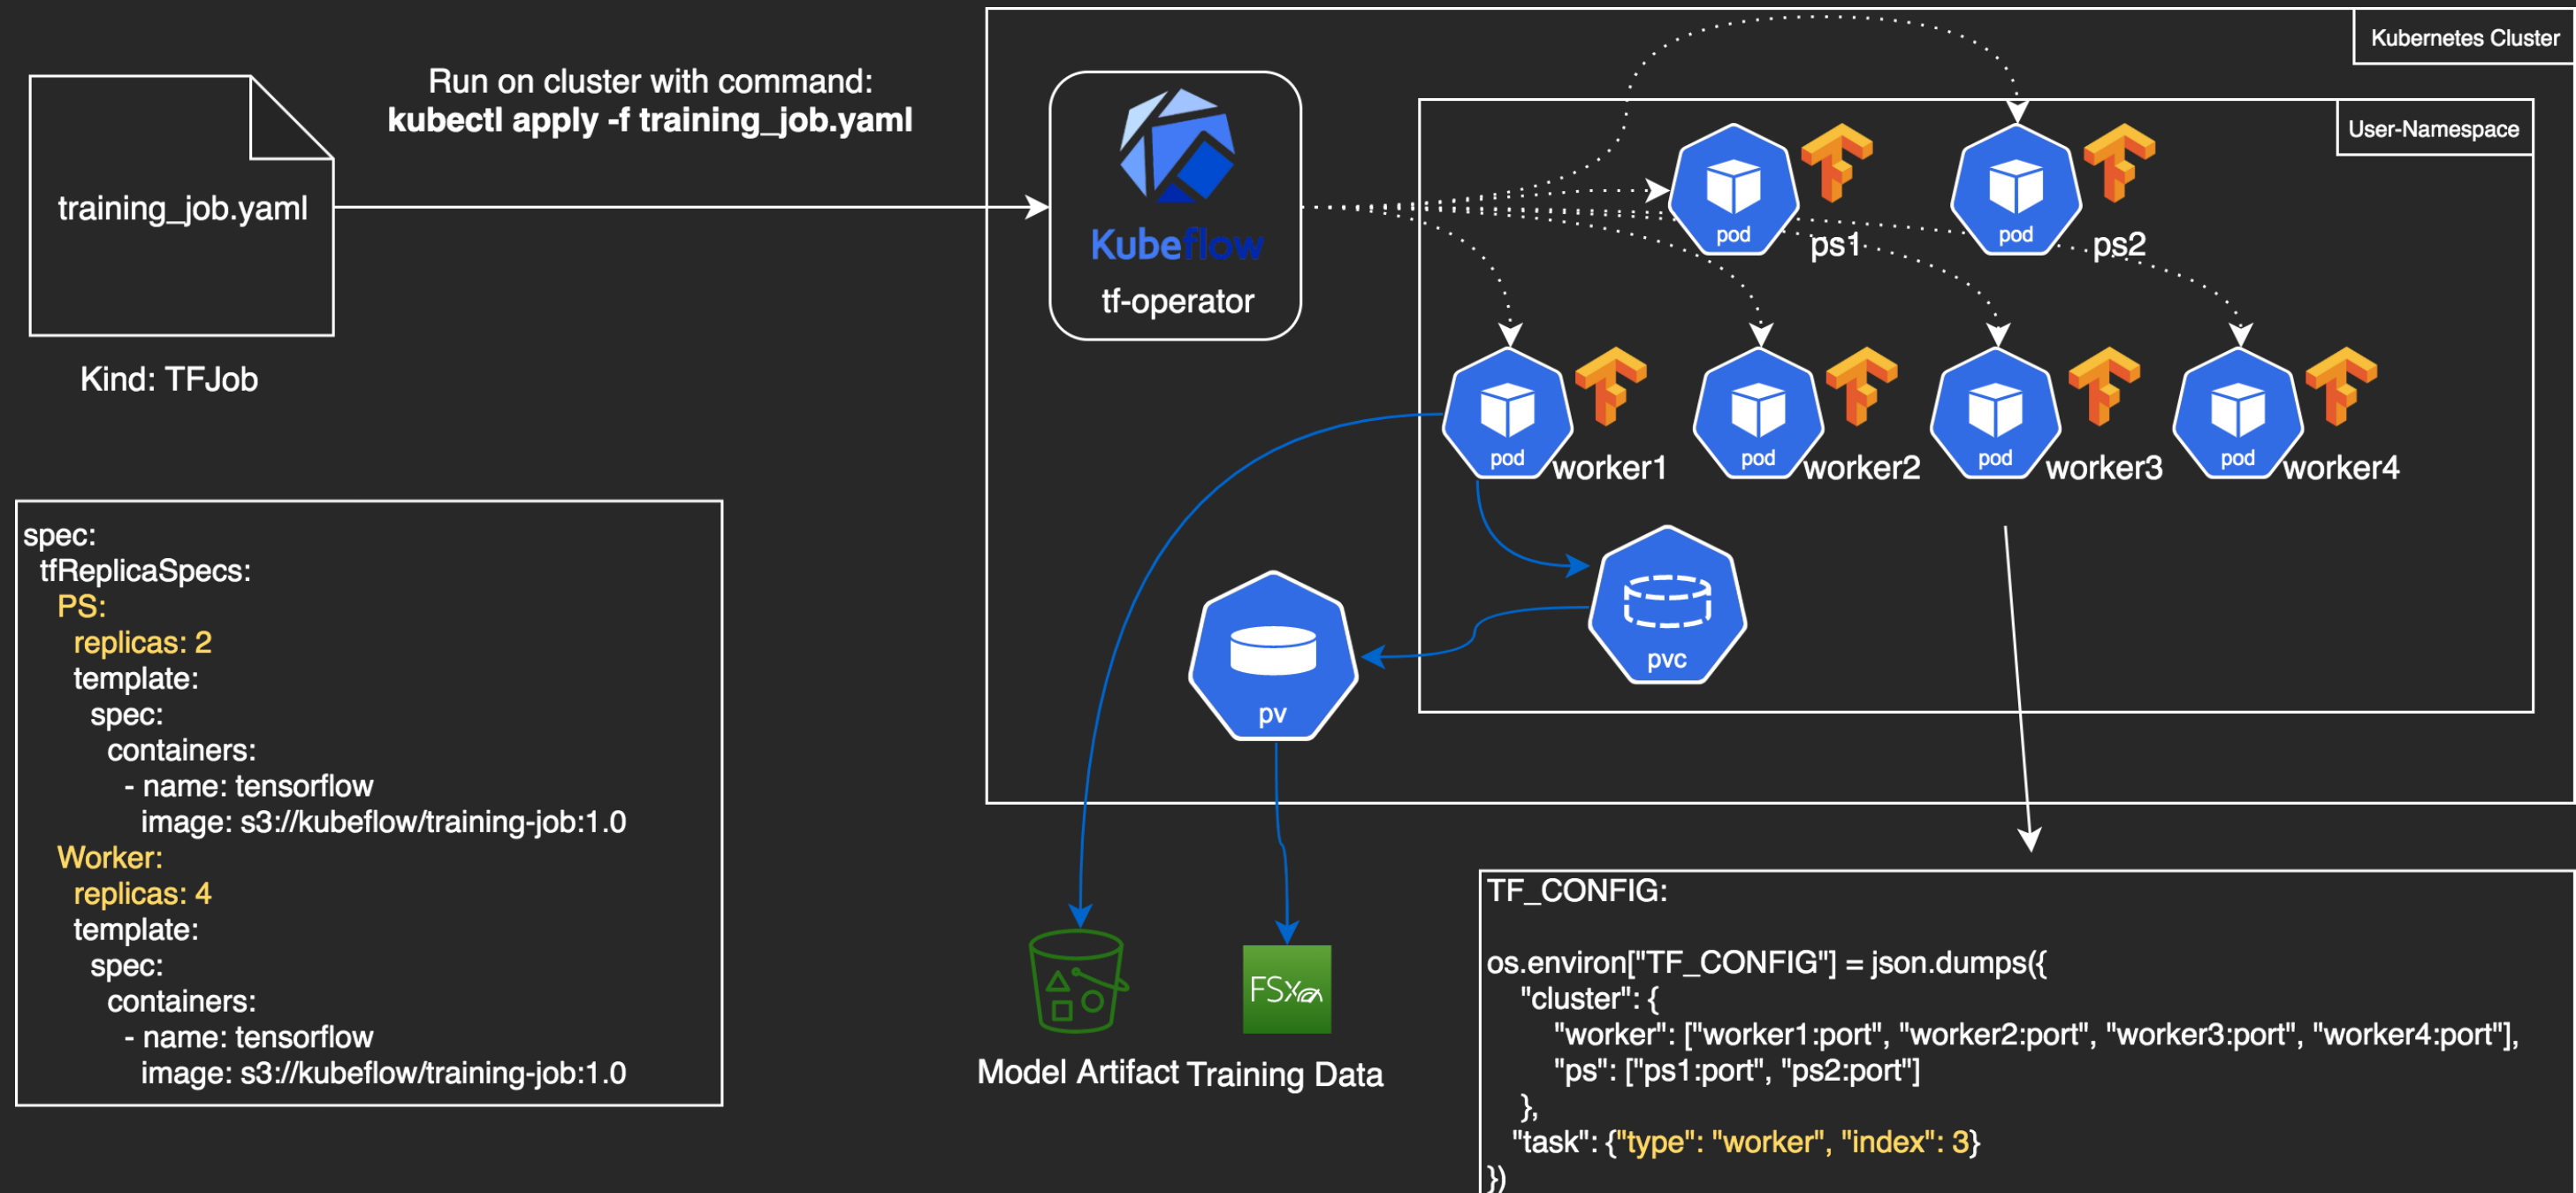
    r = metadata.Run(
        workspace=workspace,
        name="xgboost-synthetic-faring-run" + datetime.utcnow().isoformat("T"),
        description="a notebook run")

    return metadata.Execution(
        name = "execution" + datetime.utcnow().isoformat("T"),
        workspace=workspace,
        run=r,
        description="execution for training xgboost-synthetic")
```


The screenshot shows the Kubeflow ML Metadata Dashboard. The top navigation bar includes the Kubeflow logo and the user 'zhenghui'. The main section is titled 'Artifacts' and has a filter set to 'syn'. Below this is a table listing various artifacts.

Name	Version ↓	Type	URI	Workspace	Created at
synthetic-data	v1.0.0	kubeflow.org/alpha/data_set	file://path/to/dataset	xgboost-synthetic	7/19/2019, 10:06:06
housing-price-model	v0.0.3	kubeflow.org/alpha/model	mockup-model.dat	xgboost-synthetic	7/19/2019, 10:06:06
housing-price-model	v0.0.2	kubeflow.org/alpha/model	mockup-model.dat	xgboost-synthetic	7/19/2019, 10:04:43
housing-price-model	v0.0.1	kubeflow.org/alpha/model	mockup-model.dat	xgboost-synthetic	7/19/2019, 10:02:54
xgboost-synthetic-trai...		kubeflow.org/alpha/metrics	gcs://path/to/metrics	xgboost-synthetic	7/19/2019, 10:06:06

Distributed Training



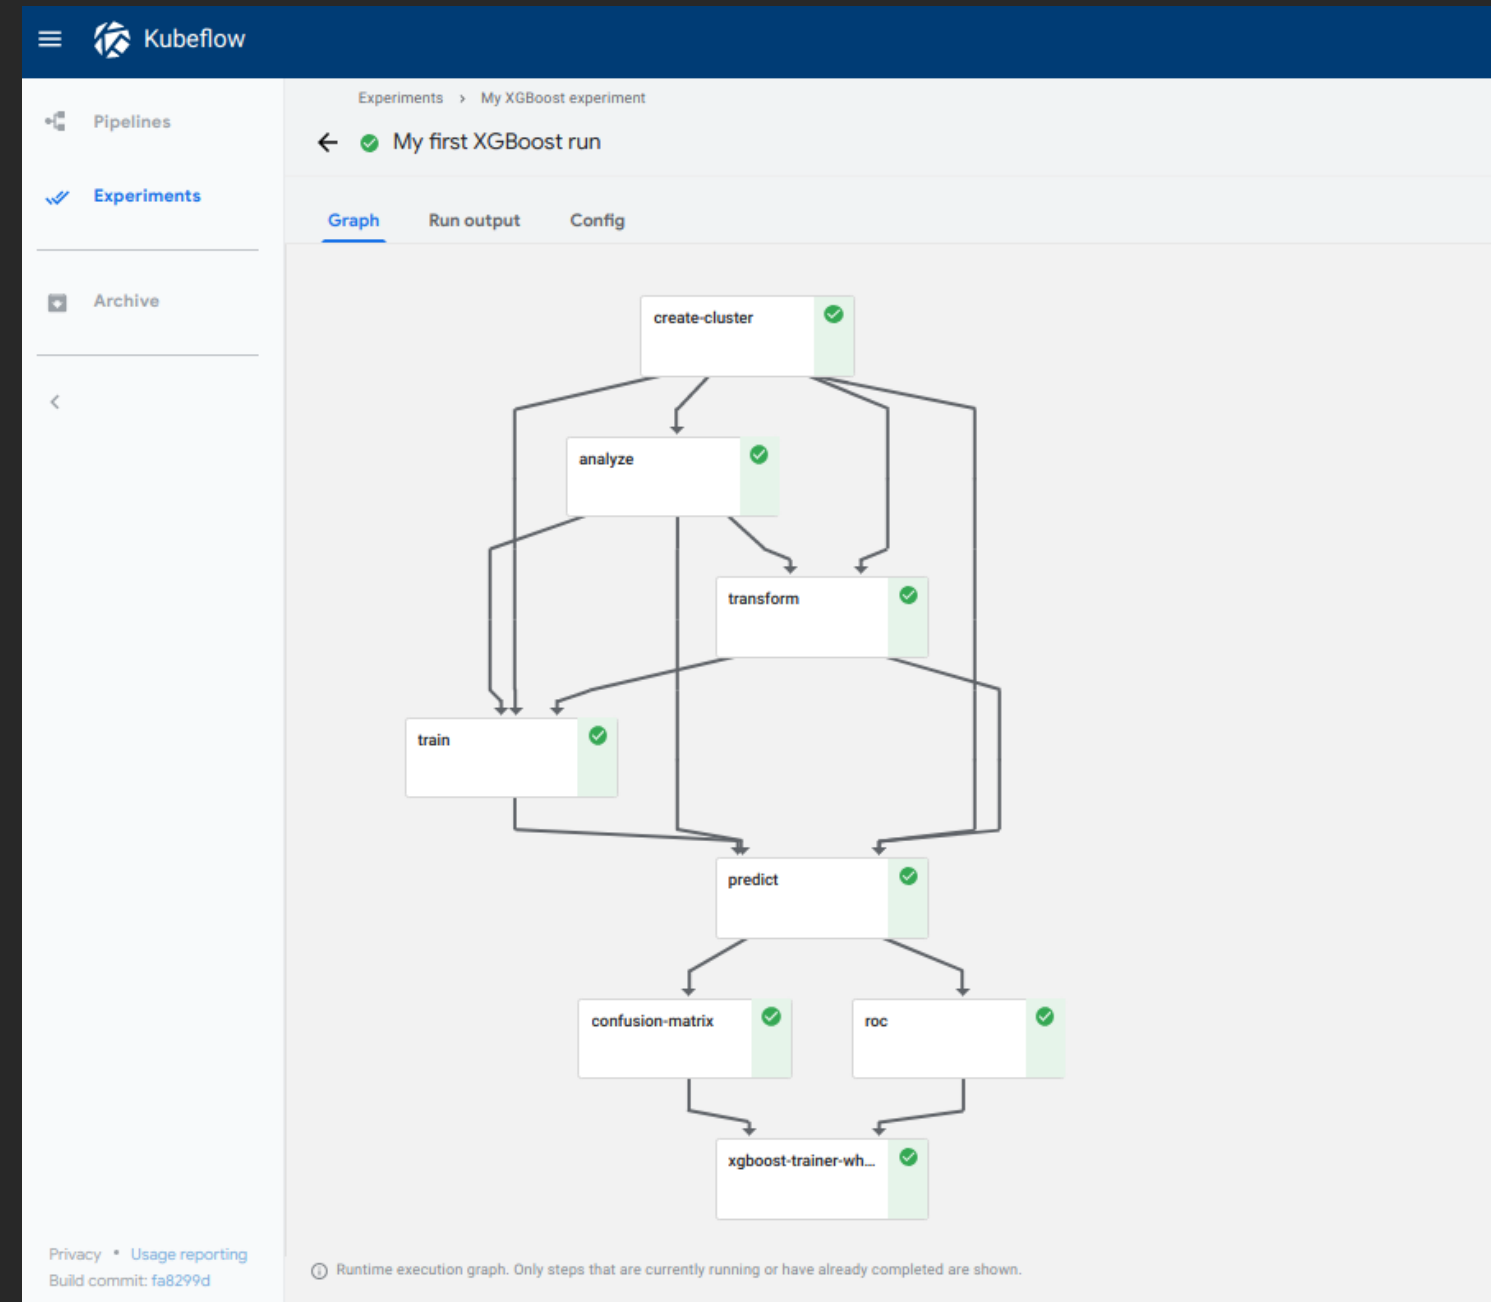
Best Practices for Optimizing Distributed Deep Learning Performance on Amazon EKS



<https://aws.amazon.com/blogs/opensource/optimizing-distributed-deep-learning-performance-amazon-eks/>

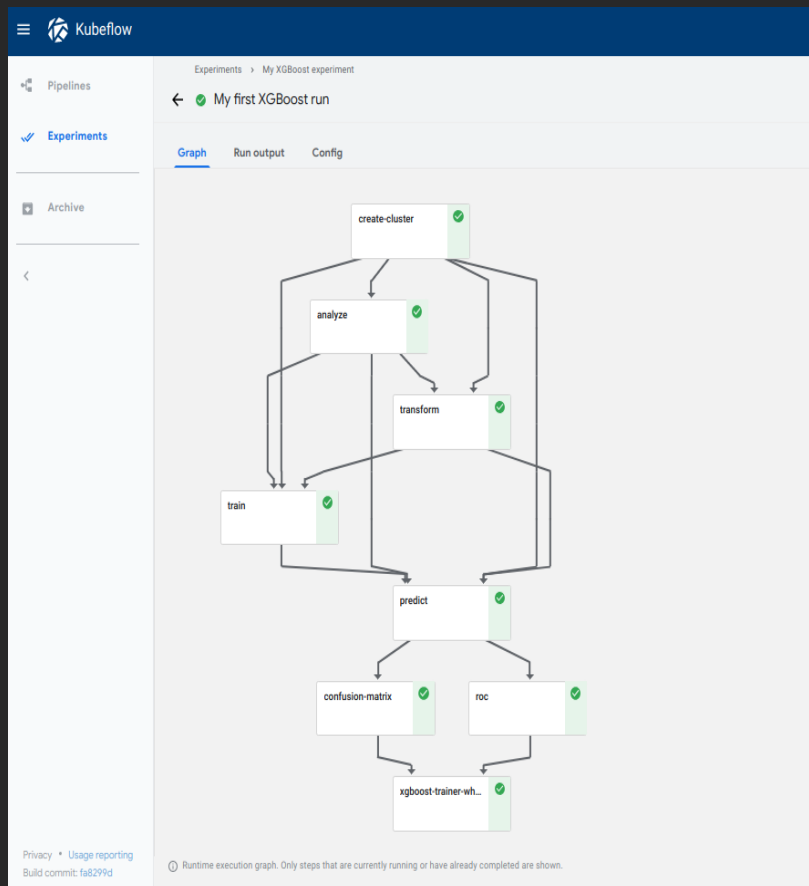
Pipelines – Machine Learning Job Orchestrator

- Compose, deploy, and manage end-to-end ML workflows
 - End-to-end orchestration
 - Easy, rapid, and reliable experimentation
 - Easy re-use
- Built using Pipelines SDK
 - `kfp.compiler`, `kfp.components`, `kfp.Client`
- Uses Argo under the hood to orchestrate resources



Creating Kubeflow Pipeline Components

Pipeline decorator



```
@dsl.pipeline(  
    name='Sample Trainer',  
    description=""  
)
```

Pipeline function

Pipeline component

```
def sample_train_pipeline(...):
```

```
    create_cluster_op = CreateClusterOp('create-cluster', ...)
```

```
    analyze_op = AnalyzeOp('analyze', ...)
```

```
    transform_op = TransformOp('transform', ...)
```

```
    train_op = TrainerOp('train', ...)
```

```
    predict_op = PredictOp('predict', ...)
```

```
    confusion_matrix_op = ConfusionMatrixOp('confusion-matrix', ...)
```

```
    roc_op = RocOp('roc', ...)
```


Compile pipeline

```
kfp.compiler.Compiler().compile(sample_train_pipeline, 'my-  
pipeline.zip')
```

Making Kubeflow a first class citizen on AWS

- Centralized and unified Kubernetes cluster logs in **Amazon CloudWatch**
- External traffic and authentication management with **ALB Ingress Controller**
- TLS and authentication with **AWS Certificate Manager** and **AWS Cognito**
- In-built **FSx CSI driver** w/S3 data repository integration to optimize training performance
- **Elastic File System** integration for common data sharing in JupyterHub
- Easier and customizable Kubeflow installation with kfctl and Kustomize support
- Kubeflow Pipeline integration with AWS Services – **Amazon EMR, Athena, SageMaker**
- Add **ECR integration** to Kubeflow Fairing
- Jupyter Notebook images with **AWS CLI installed and ECR support**
- Auto detect GPU worker nodes and install NVIDIA device plugin

Kubeflow 1.0 Arriving January 2020



http://bit.ly/kf_roadmap

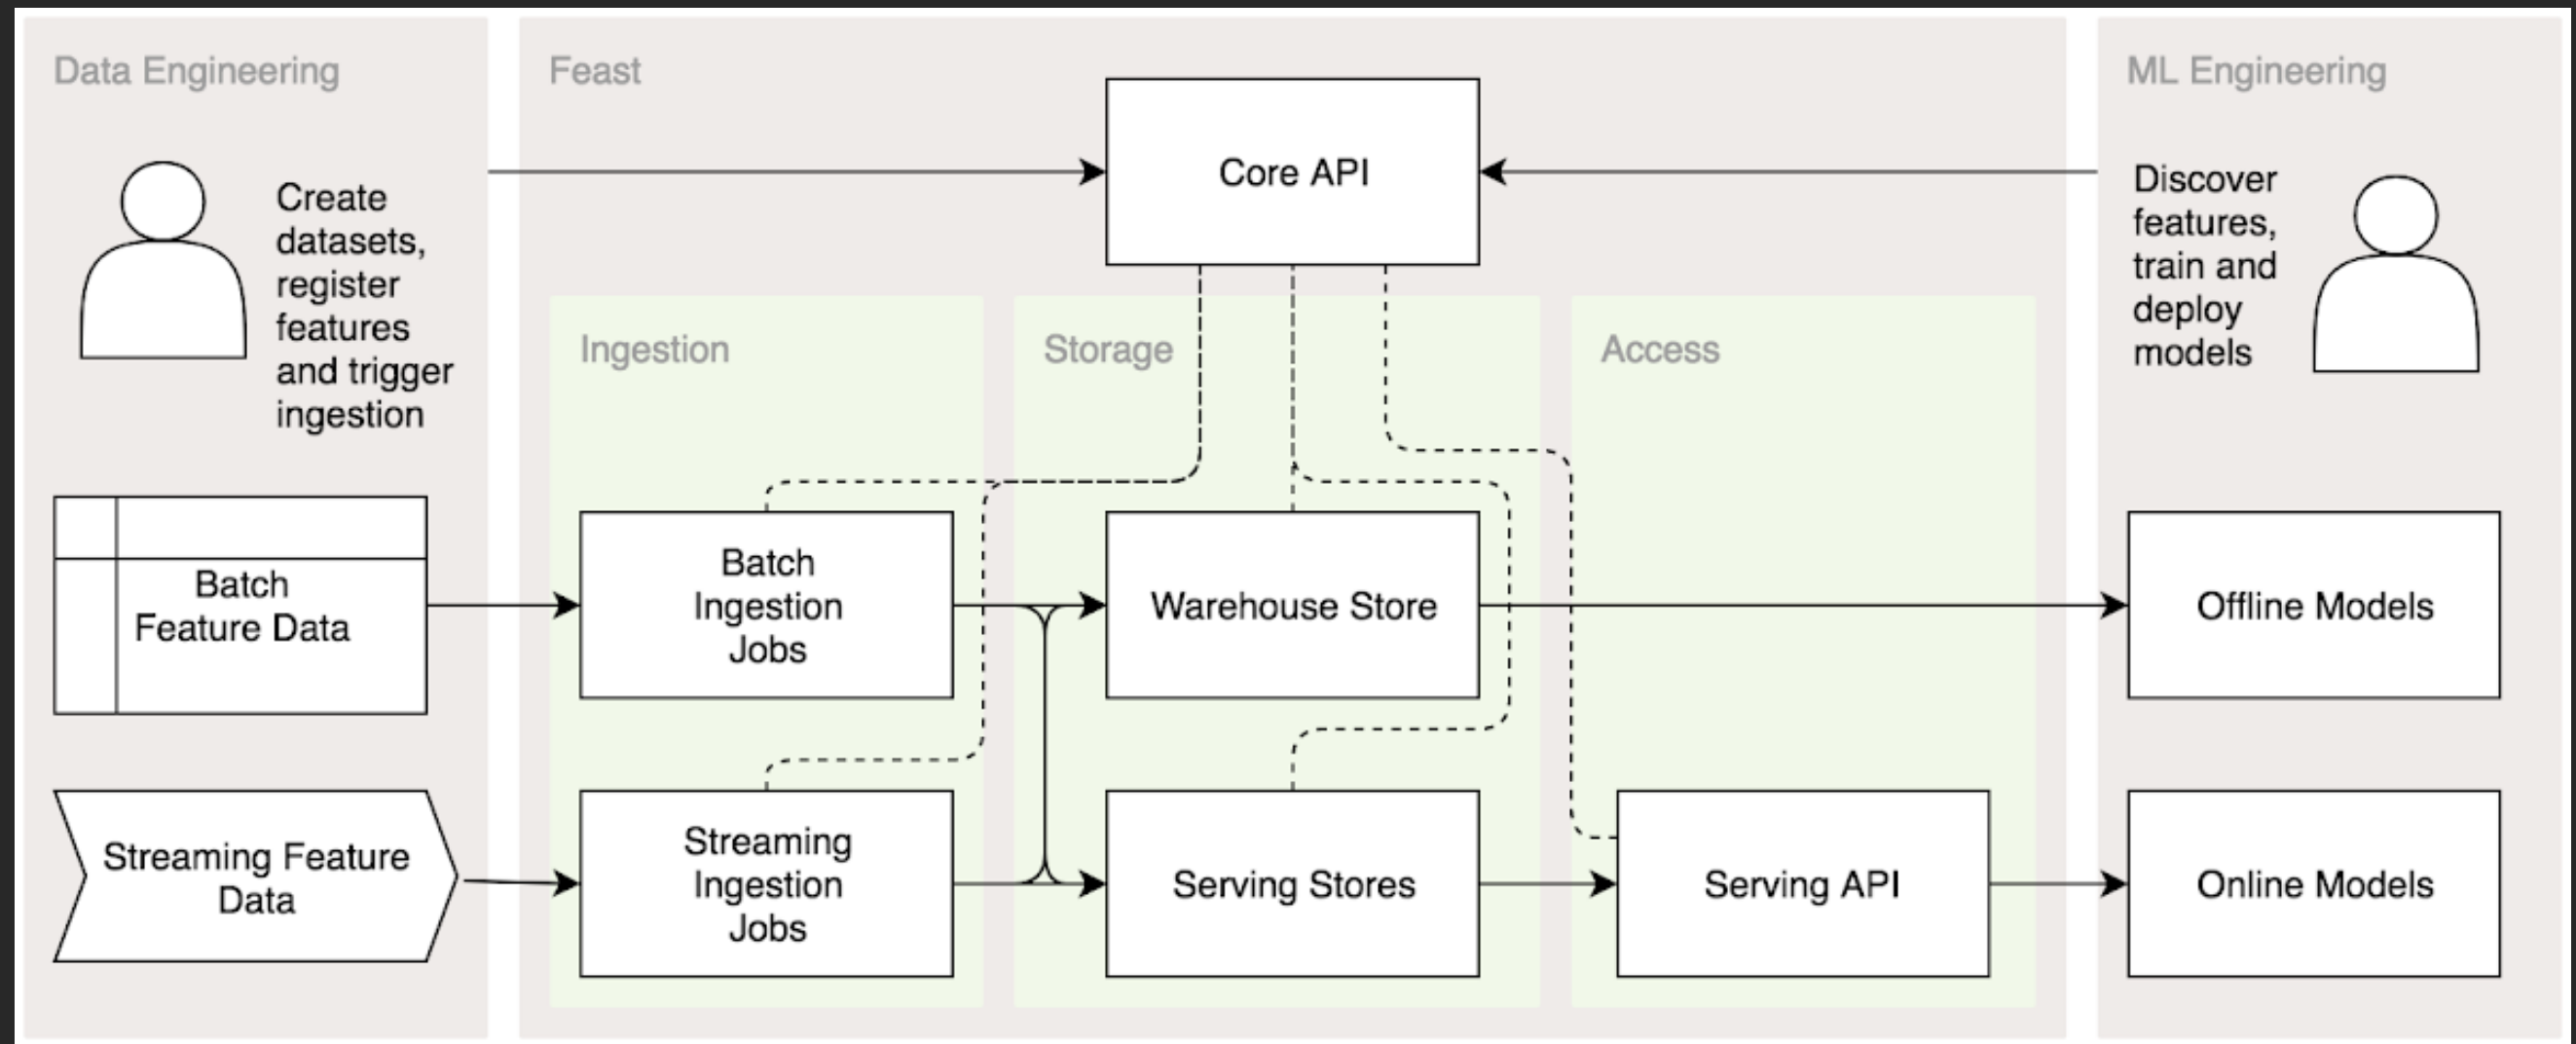


AWS Kubeflow Roadmap

Kubeflow v1.0 - Theme: Enterprise Readiness

- E2E examples and increased documentation on Kubeflow site
- Upstream testing for Kubeflow on AWS
- Support DIY K8S on AWS
- IAM Roles for Service Accounts integration with Jupyter notebooks
- Support for managed contributors

Feature store - Feast



- Discoverability and reuse of features
- Standardization of features
- Access to features for training and serving
- Consistency between training and serving

Enterprise

Machine Learning Stack on EKS


ML Platform



Container Platform



ML Frameworks + Infrastructure



Future Ideas

- Elastic Training
- Virtual GPU Device Plugin
- GPU Monitoring
- MLOps

Elastic Training

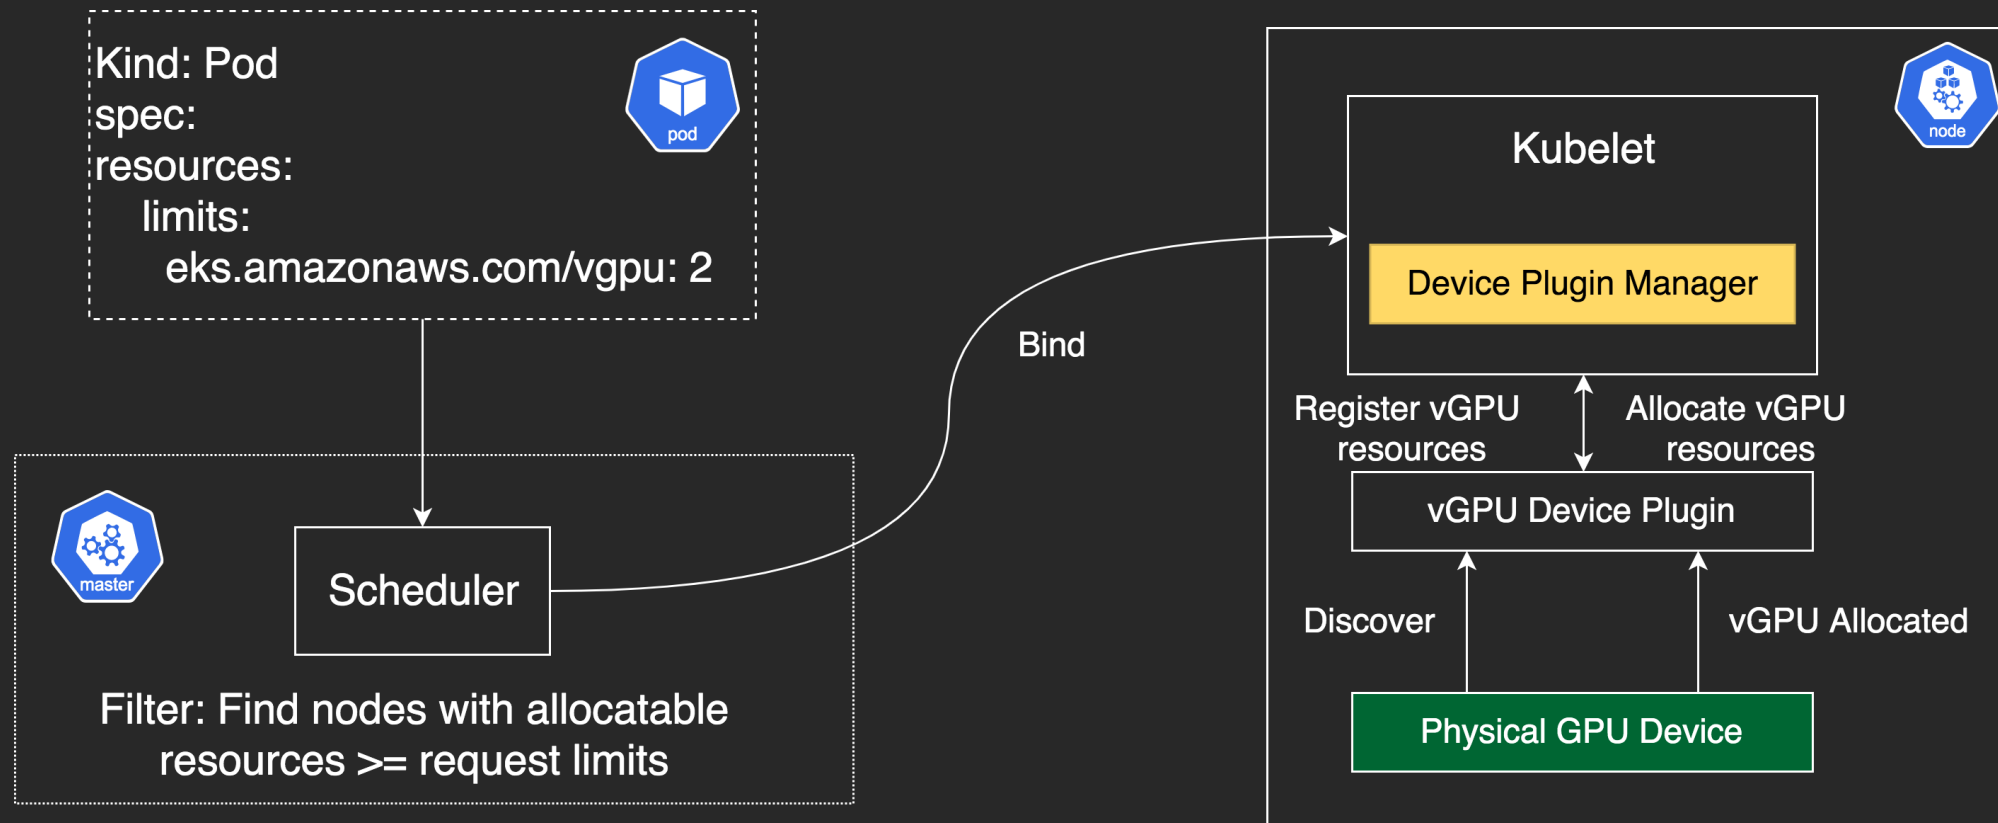
Fault Tolerant

- Enable Job Priority & Preemption
- Unlock SLA Critical Jobs to run on Spot instances

Elastic Scheduling

- Improve the GPU utilization rate

Virtual GPU for ML inference

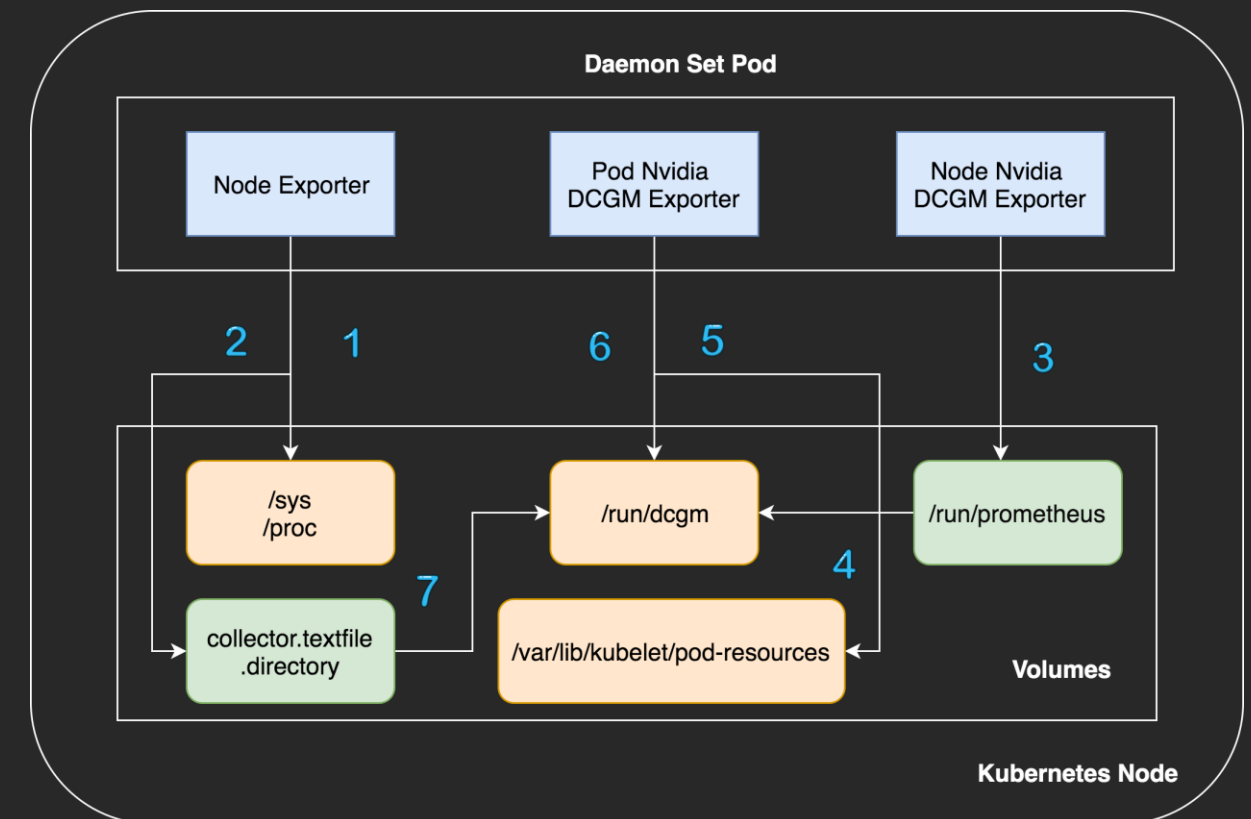
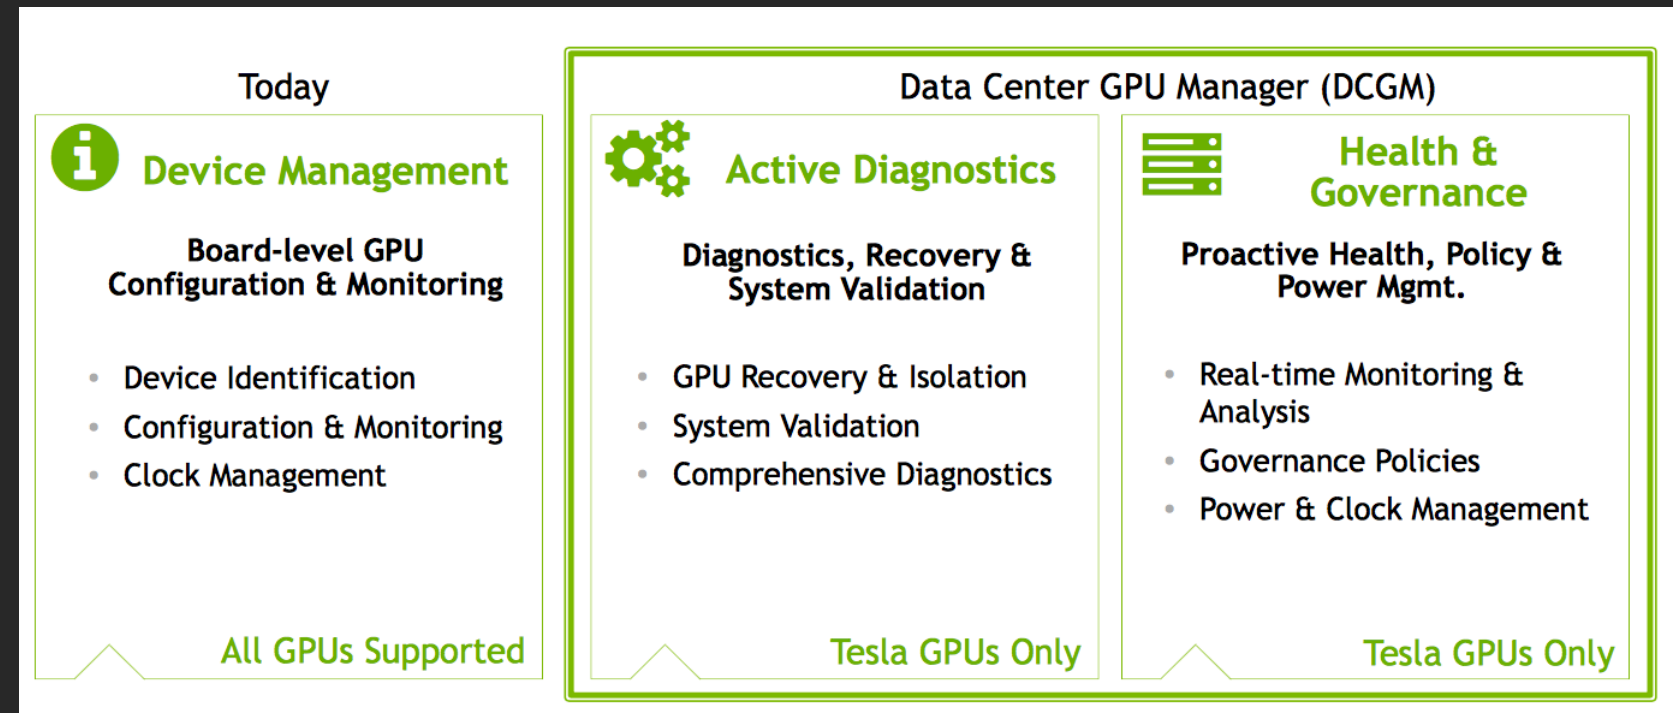


- **Map** physical GPU to virtual GPUs in vGPU Device Plugin
e.g. Physical GPU = 100 vGPUs
- vGPU Device Plugin **reports** discovered extended resources to API Server
- Kubernetes Scheduler **schedule** pods request Custom Resource e.g. eks.amazonaws.com/vgpu
- vGPU Device Plugin **allocate** Physical GPU

GPU Monitoring

Monitoring Stack

- Nvidia Management Library (NVML)
- Data Center GPU Manager (DCGM)
- CAdvisor accelerator metrics (CAdvisor)



MLOps

Data Scientists

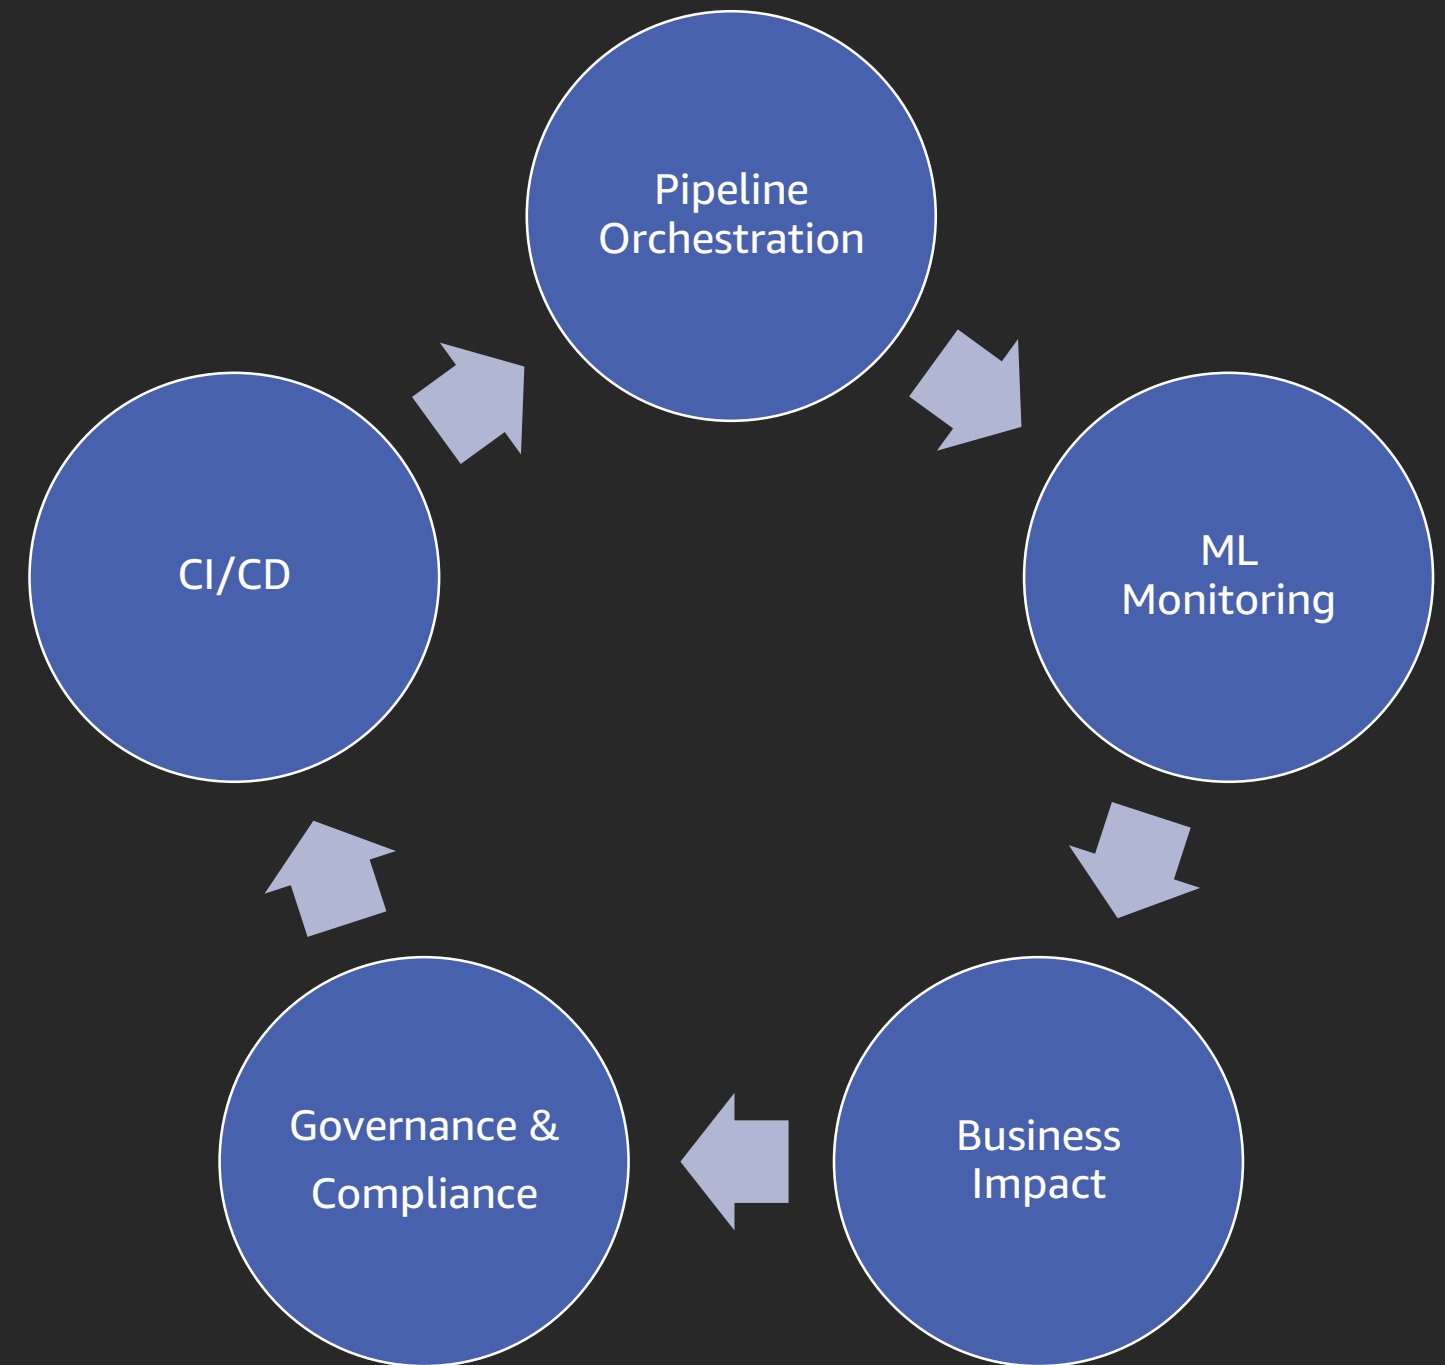
Hybrid, integrated, cloud-based dev env
Version control (scripts & artifacts)

DevOps

Seamless deployment of hybrid pipelines
Trigger-based scheduling & orchestration of runs
Monitoring & dashboard
Version control (runs & pipelines)

Business Analyst

Ensure Compliance
Explain-ability



Goal: Graduate ML into a first class citizen of software development

Thank you!

Jiaxin Shan

Software Engineer
AWS

Mike Stefaniak

Product Manager
AWS

Arun Gupta

Open Source Technologist
AWS



Please complete the session
survey in the mobile app.