

AWS
re:Invent

C O N 2 1 2 - R

Running Kubernetes at Amazon scale using Amazon EKS

Kevin Dillane

Sr. Software Engineer
Amazon.com

Cristian Munoz Urbancic

Sr. Software Engineer
Amazon.com

Bob Wise

General Manager
Elastic Kubernetes Service
(EKS)
Amazon Web Services

Agenda

About velocity

From monoliths to microservices: The rendering problem

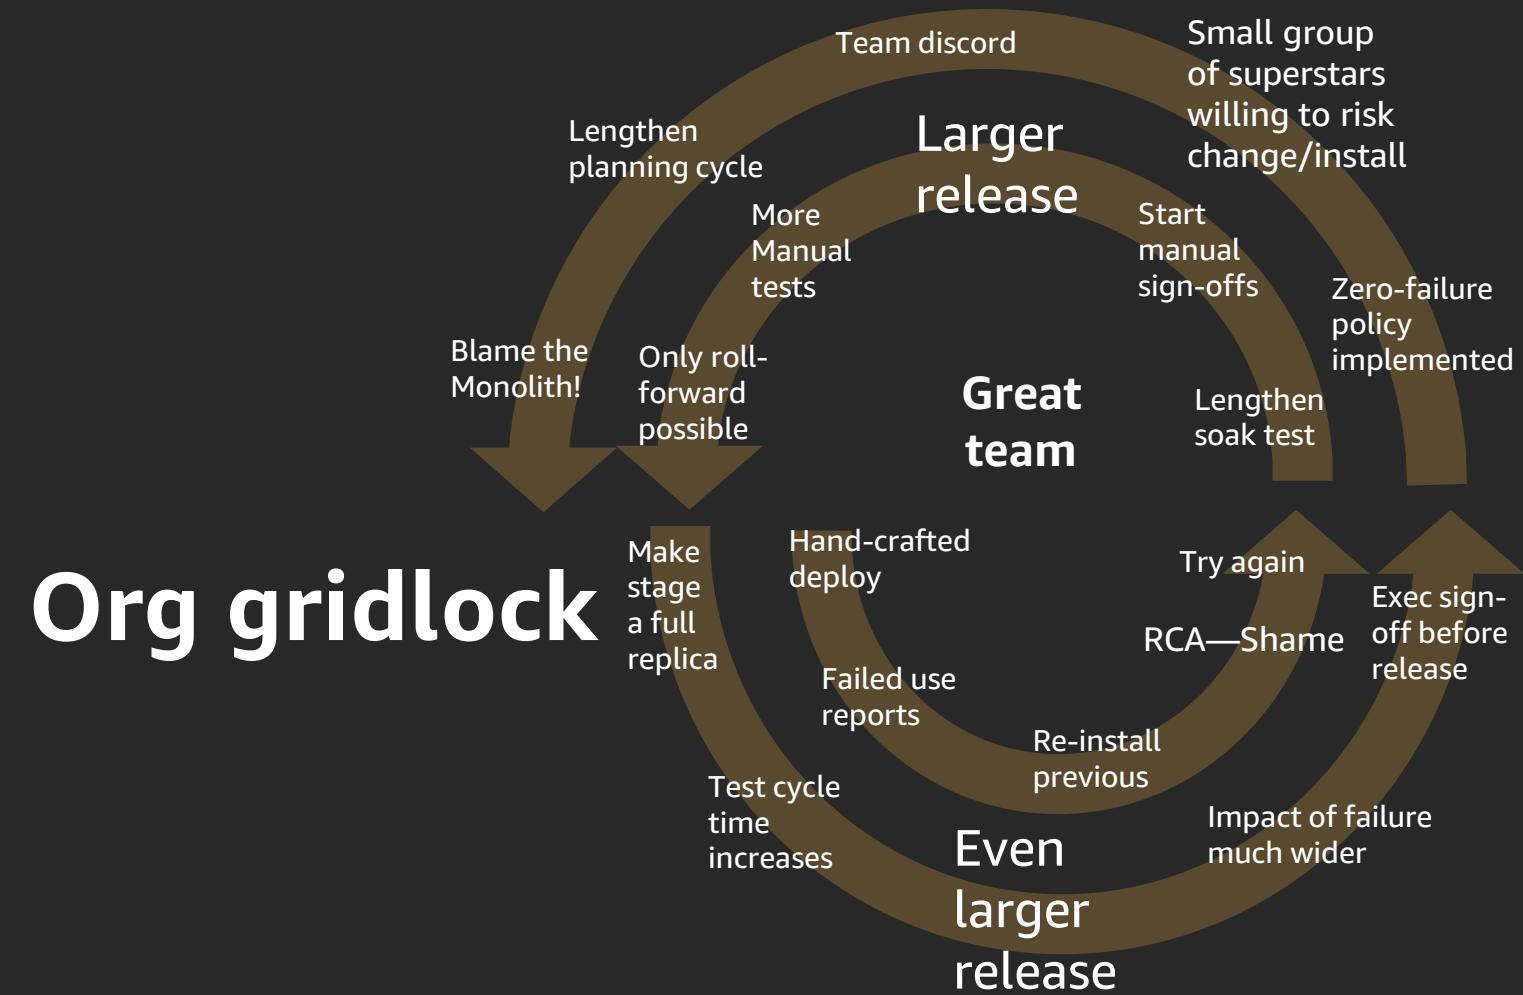
Improving ownership by owning less

CI/CD: Release features all the time!

How we learned to stop worrying and ❤️ Amazon scale

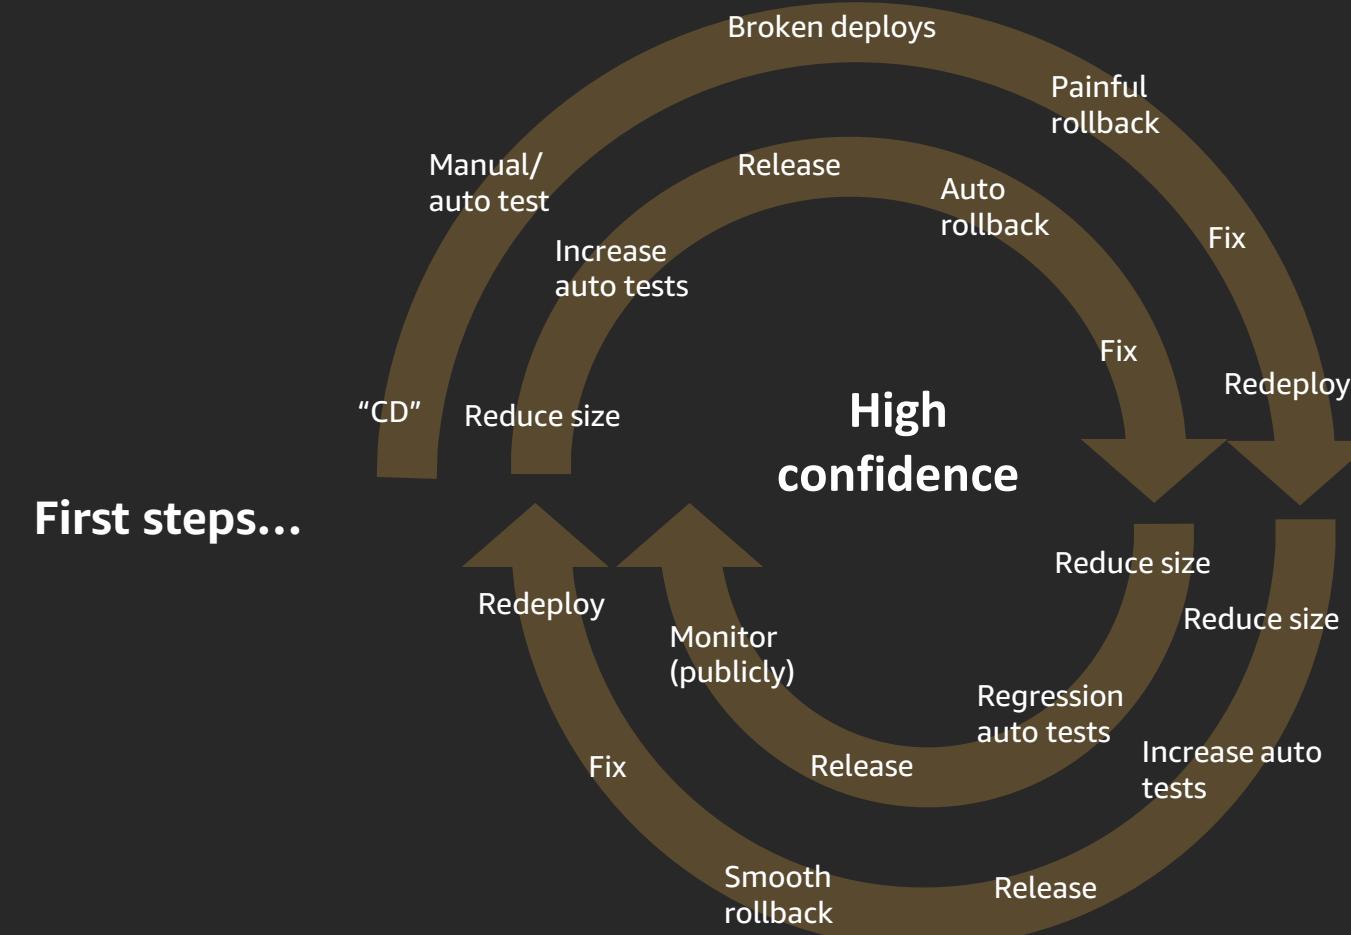
About velocity

Good intentions... Painful result



Spiral to slowness

Path to higher velocity...



Virtuous flywheel

Evolved teams

5x

Lower change
failure rate

440x

Faster from
commit to deploy

46x

More frequent
deployments

44%

More time spent on
new features & code

Containers on hosts

A host is a server, e.g., Amazon Elastic Compute Cloud (Amazon EC2) virtual machine. We run these hosts together as a cluster.

Our simple example web application is already containerized



Web
application

To start let's run 3 copies of our web application across our cluster of Amazon EC2 hosts

Host 1

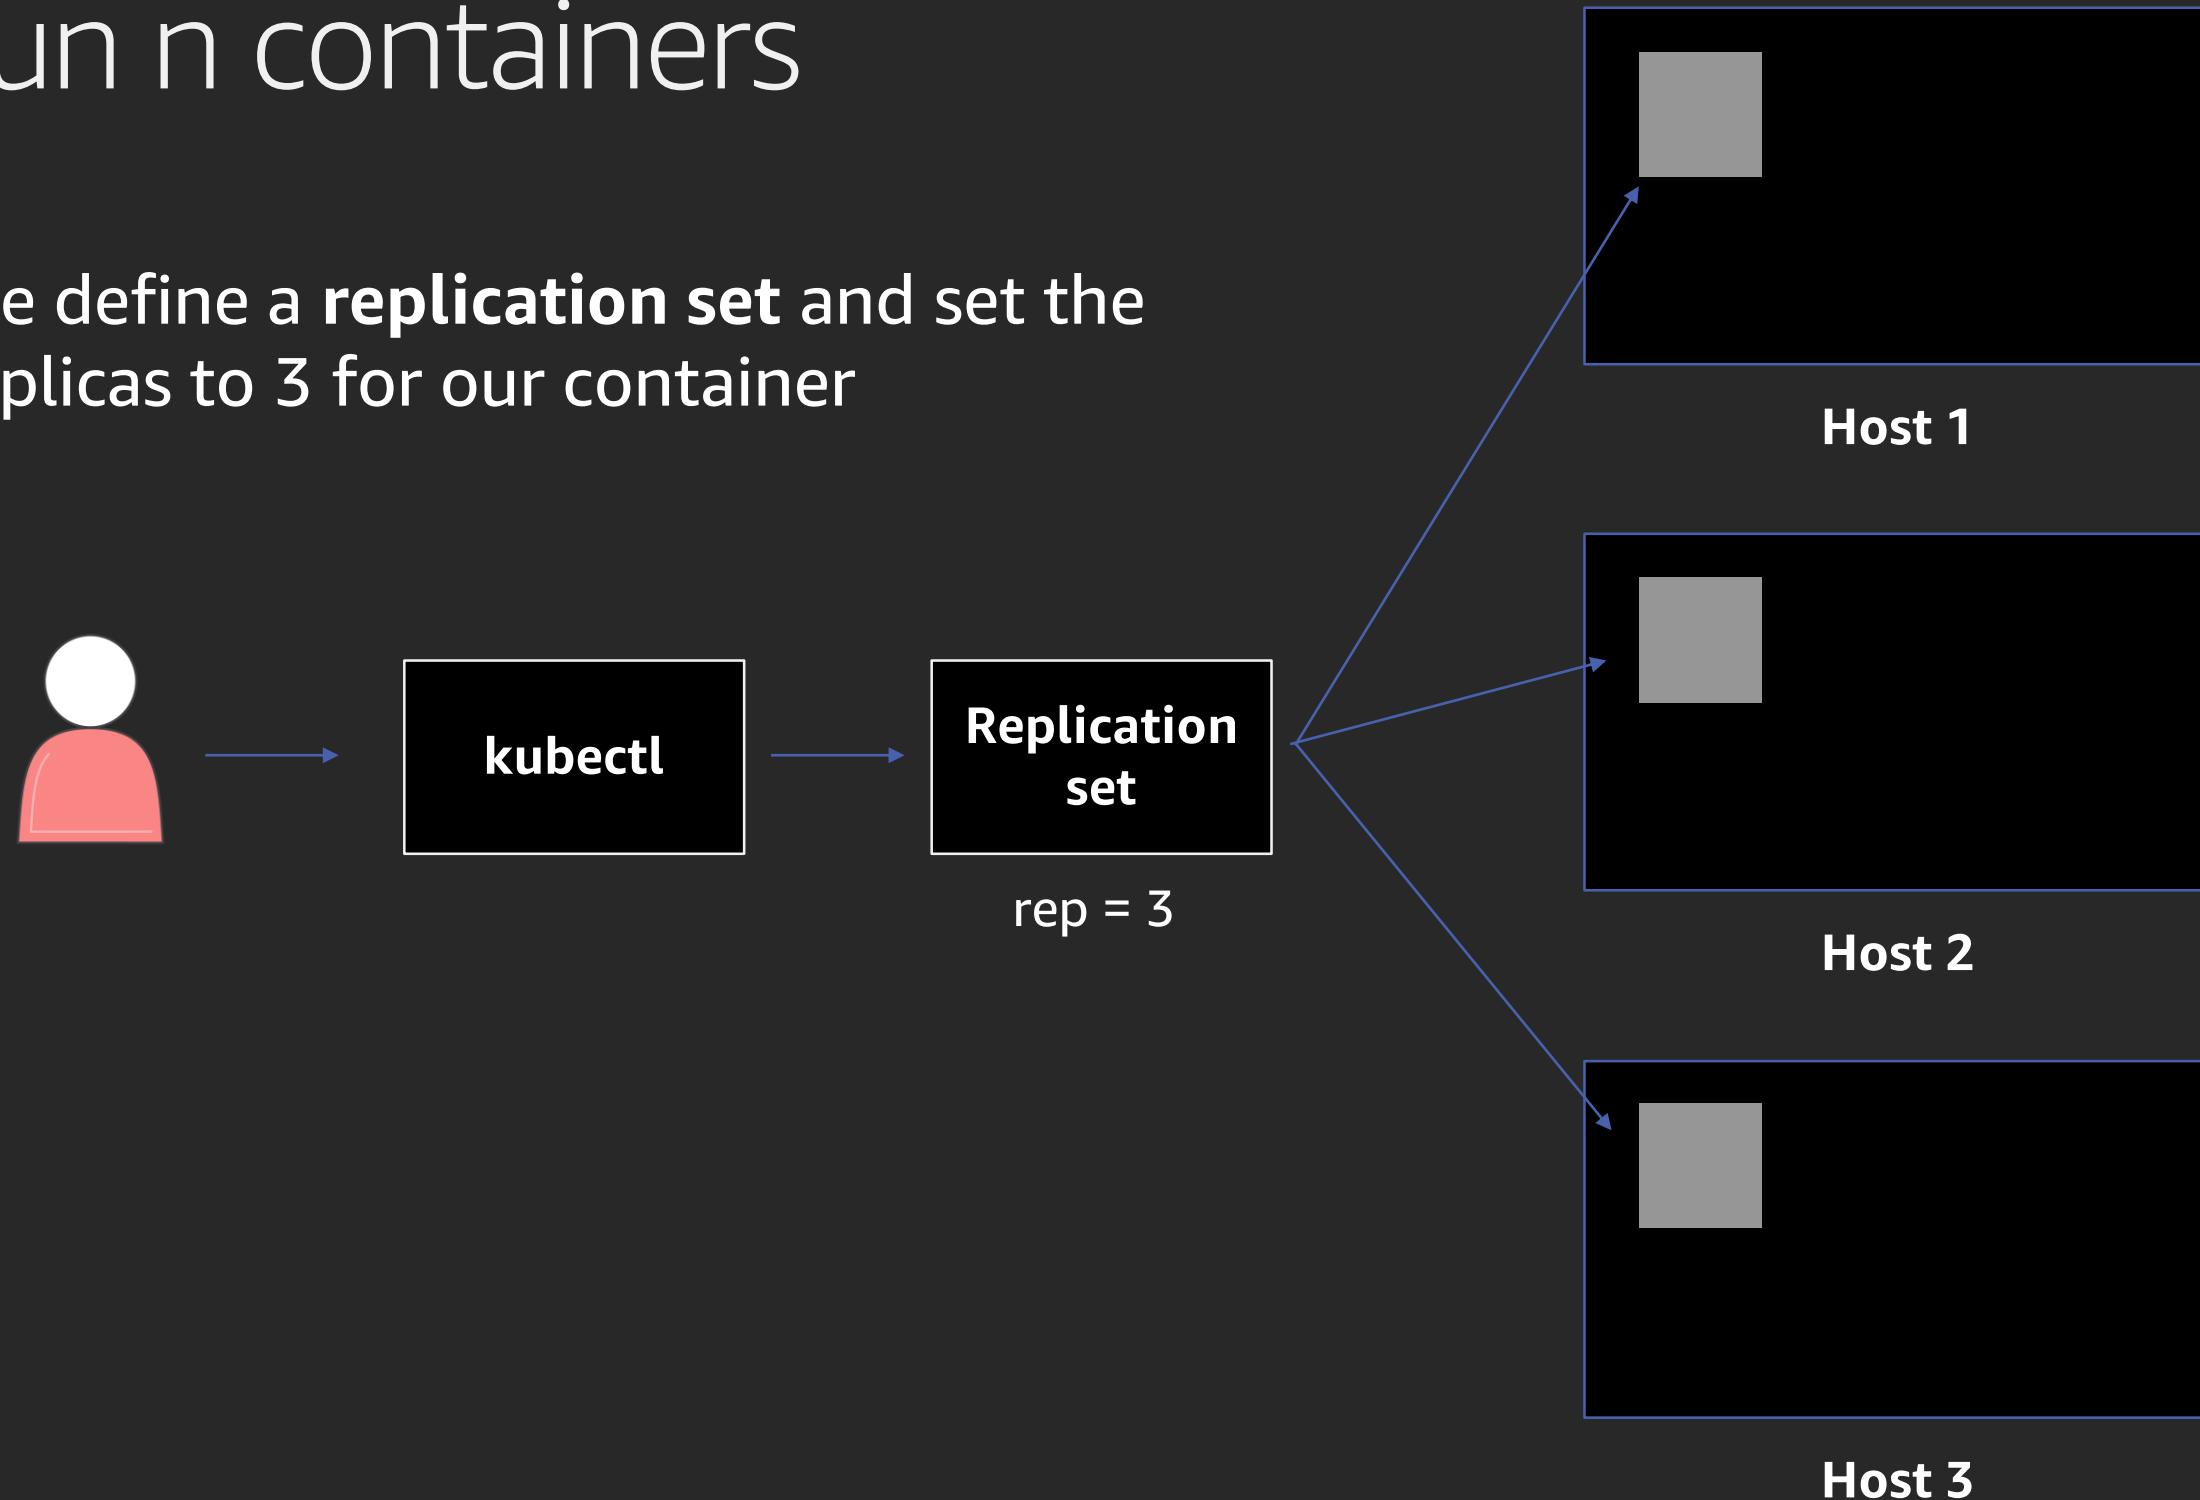
Cluster

Host 2

Host 3

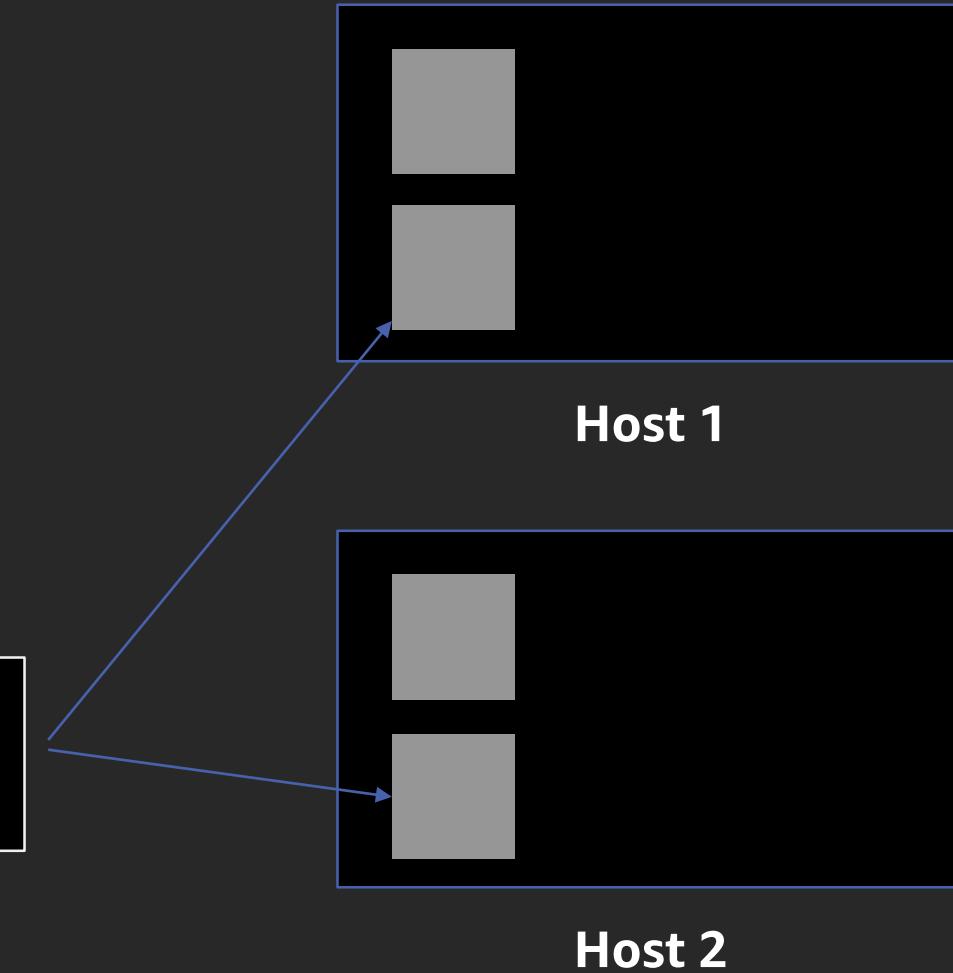
Run n containers

We define a **replication set** and set the replicas to 3 for our container



Scale up!

Need more containers?
Update the replication set!



The new containers are started on the cluster

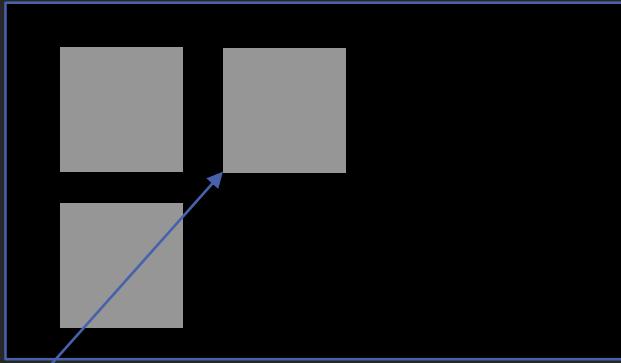
Untimely termination

Oh no! Our host has died!

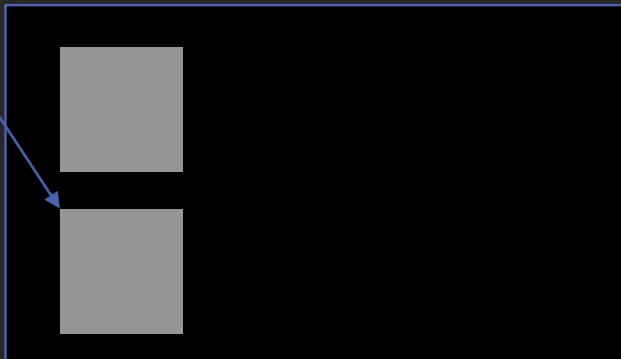
Replication set

rep = 5

Kubernetes notices only 3 of the 5 containers are running and starts 2 additional containers on the remaining hosts



Host 1



Host 3

Containers IRL

In production, we want to do more complex things like:

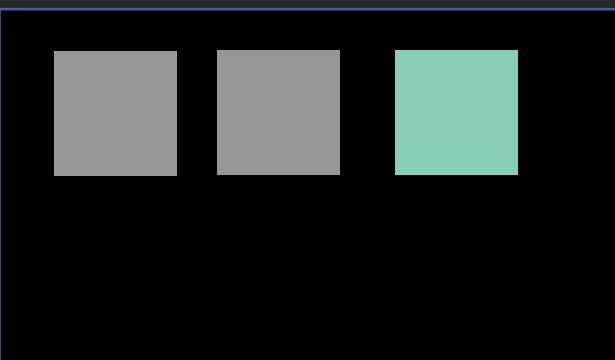
- Run a **service** to route traffic to a set of running containers
- Manage the **deployment** of containers to our cluster
- Run multiple containers **together** and specify how they run



Host 1



Host 2



Host 3

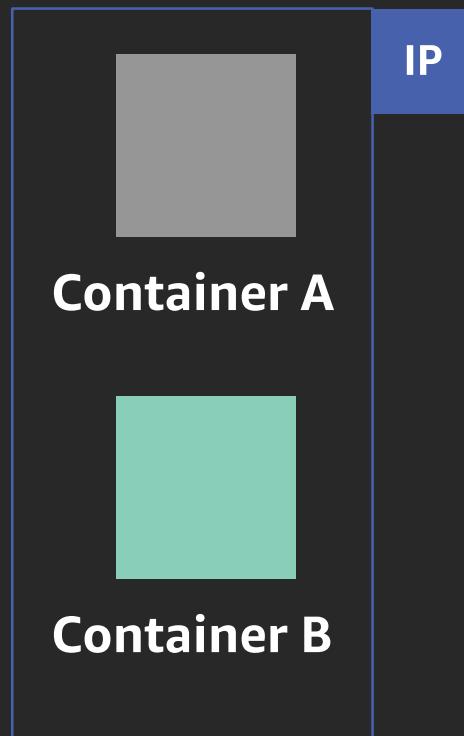
Pods

- Define how your containers should run
- Allow you to run 1 to n containers together

Containers in pods have

- Shared IP space
- Shared volumes
- Shared scaling (you scale pods, not individual containers)

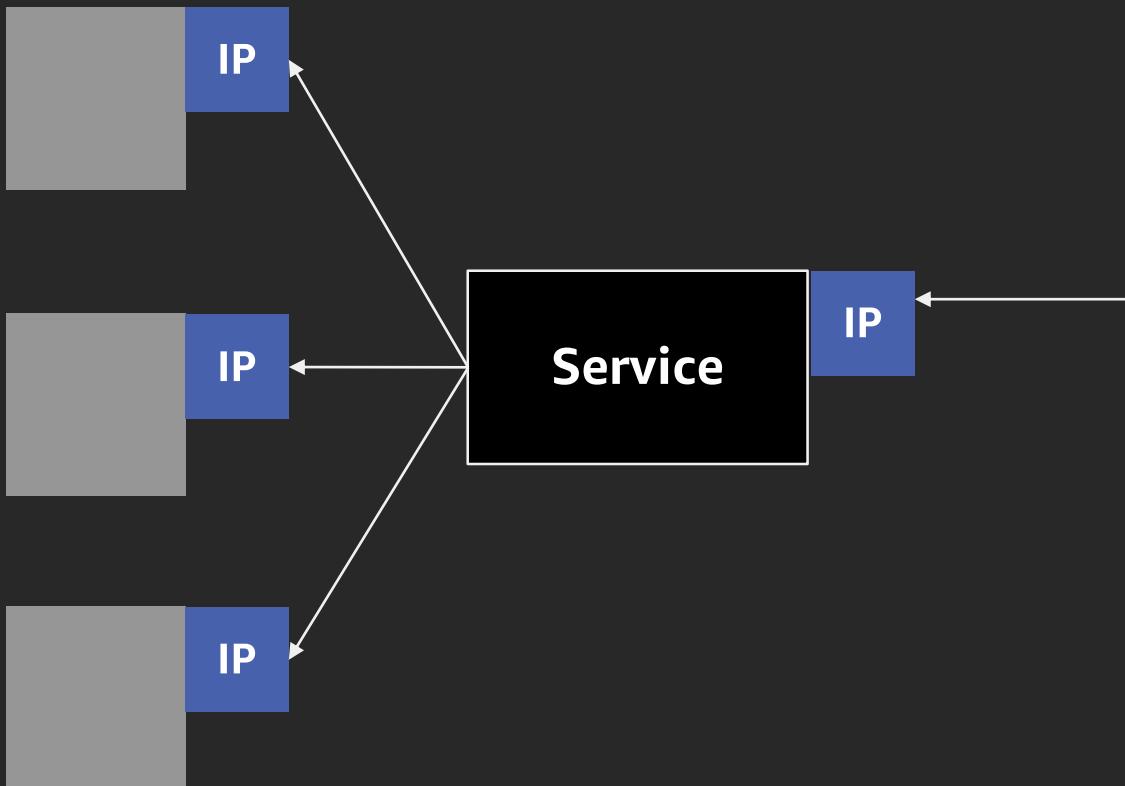
When containers are started on our cluster, they are always part of a pod
(even if it's a pod of 1)



Services

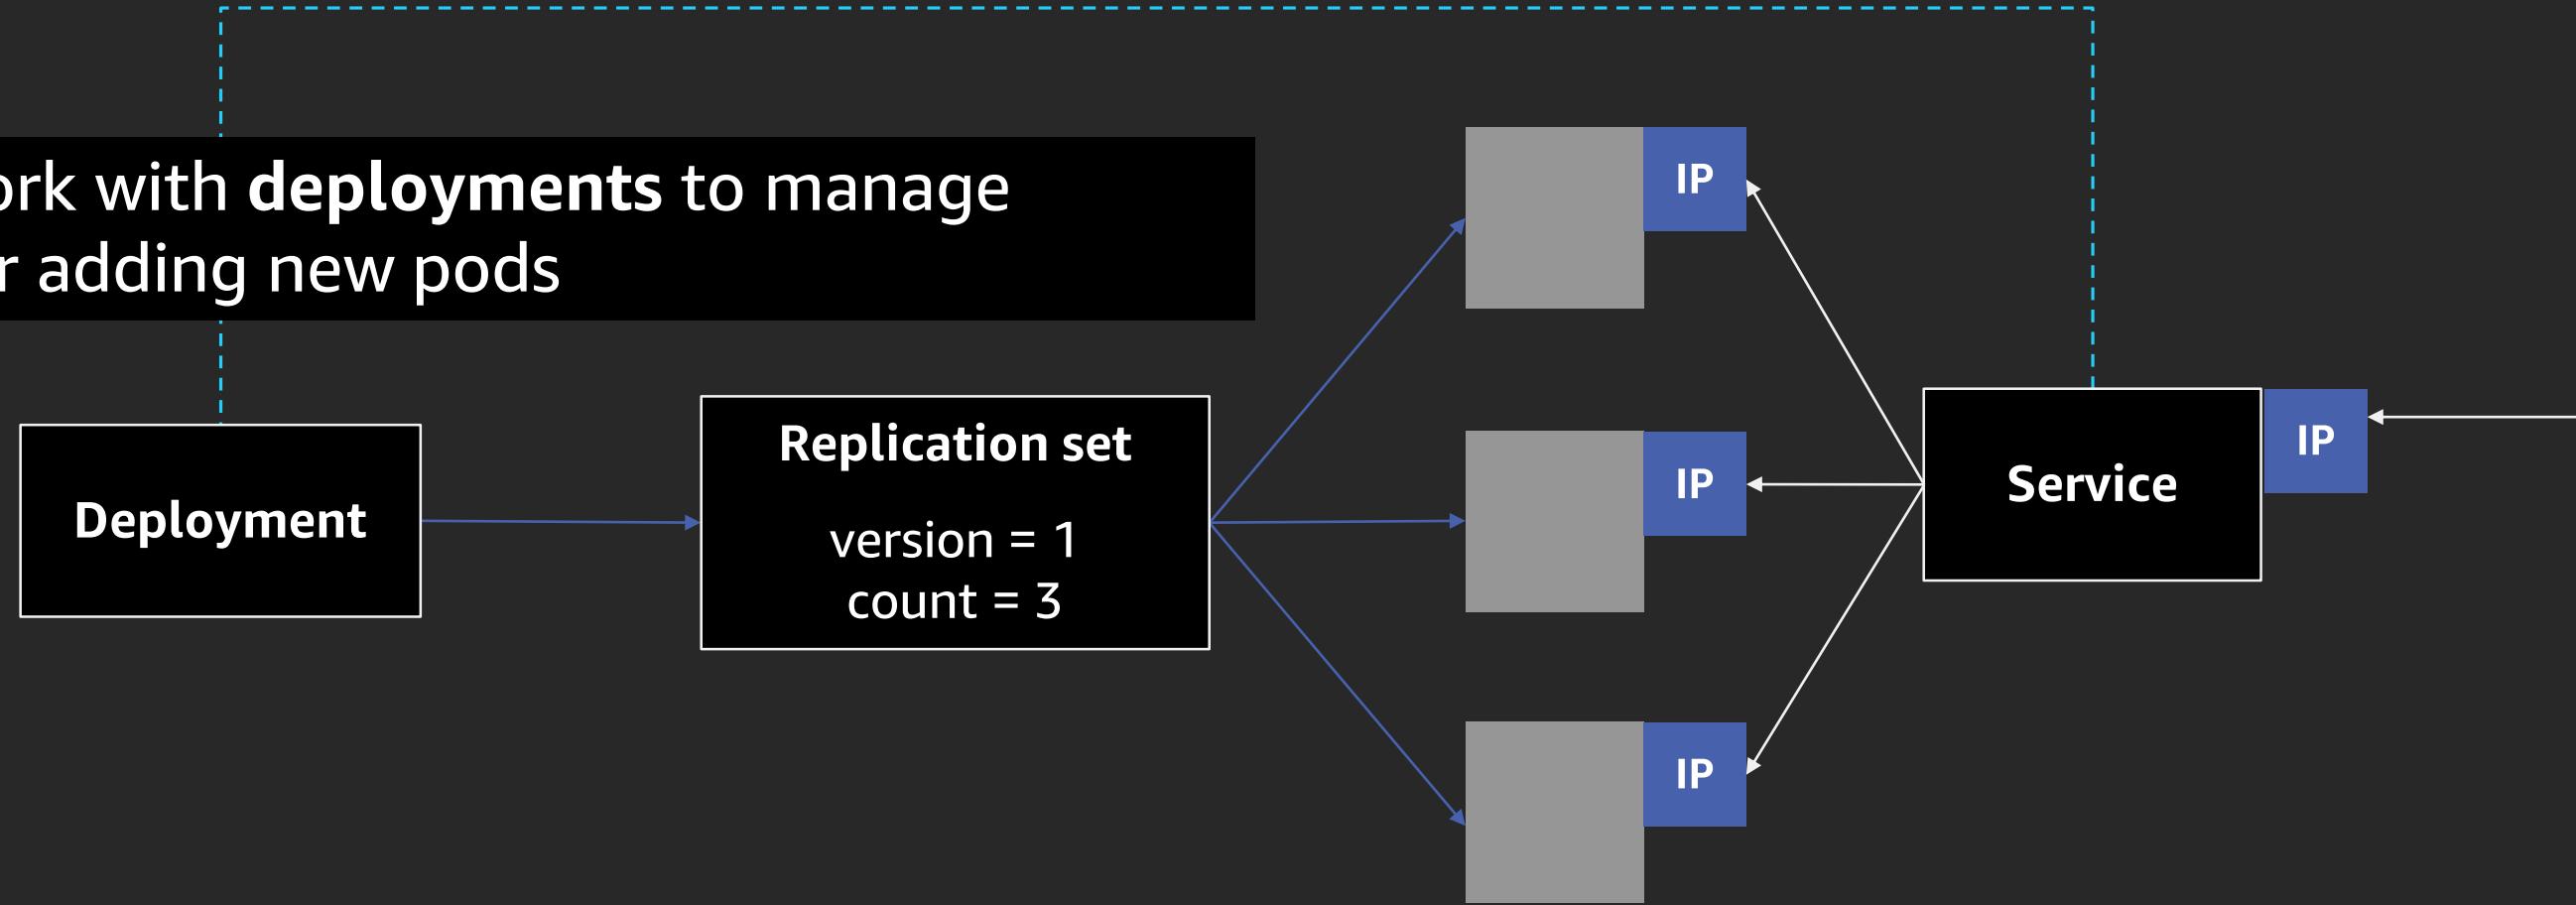
One of the ways traffic gets to your containers

- Internal IP addresses are assigned to each container
- Services are connected to containers and use labels to reference which containers to route requests to



Deployments

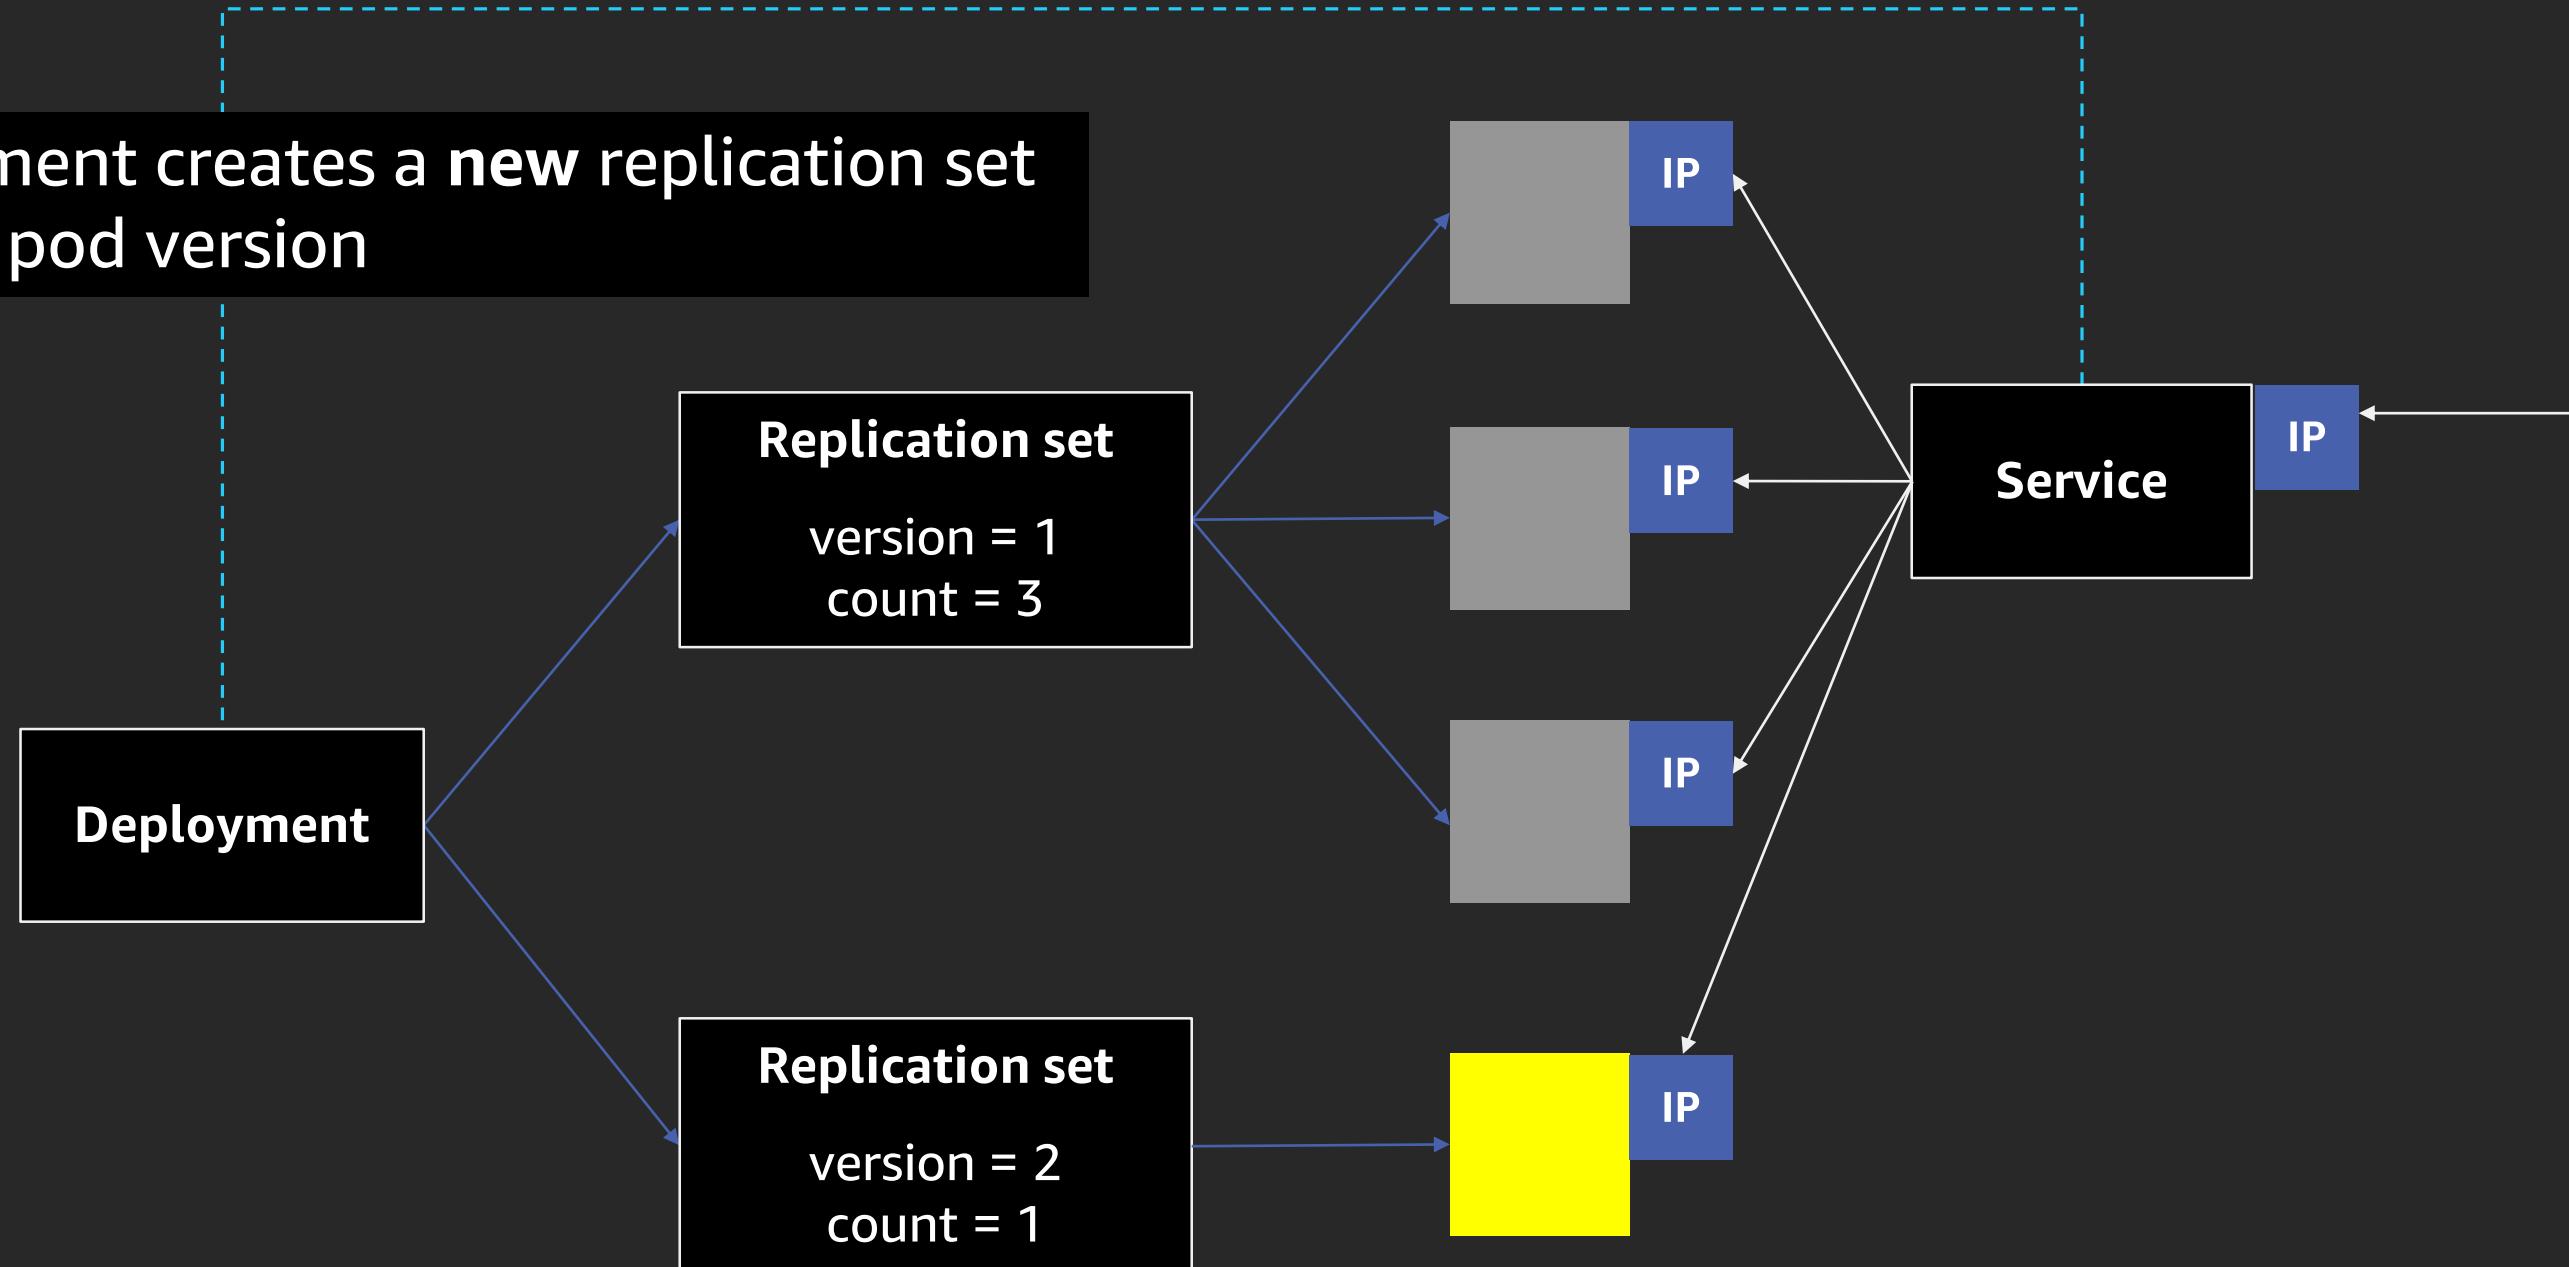
Services work with **deployments** to manage updating or adding new pods



Let's say we want to deploy a new version of our web application as a “canary” and see how it handles traffic

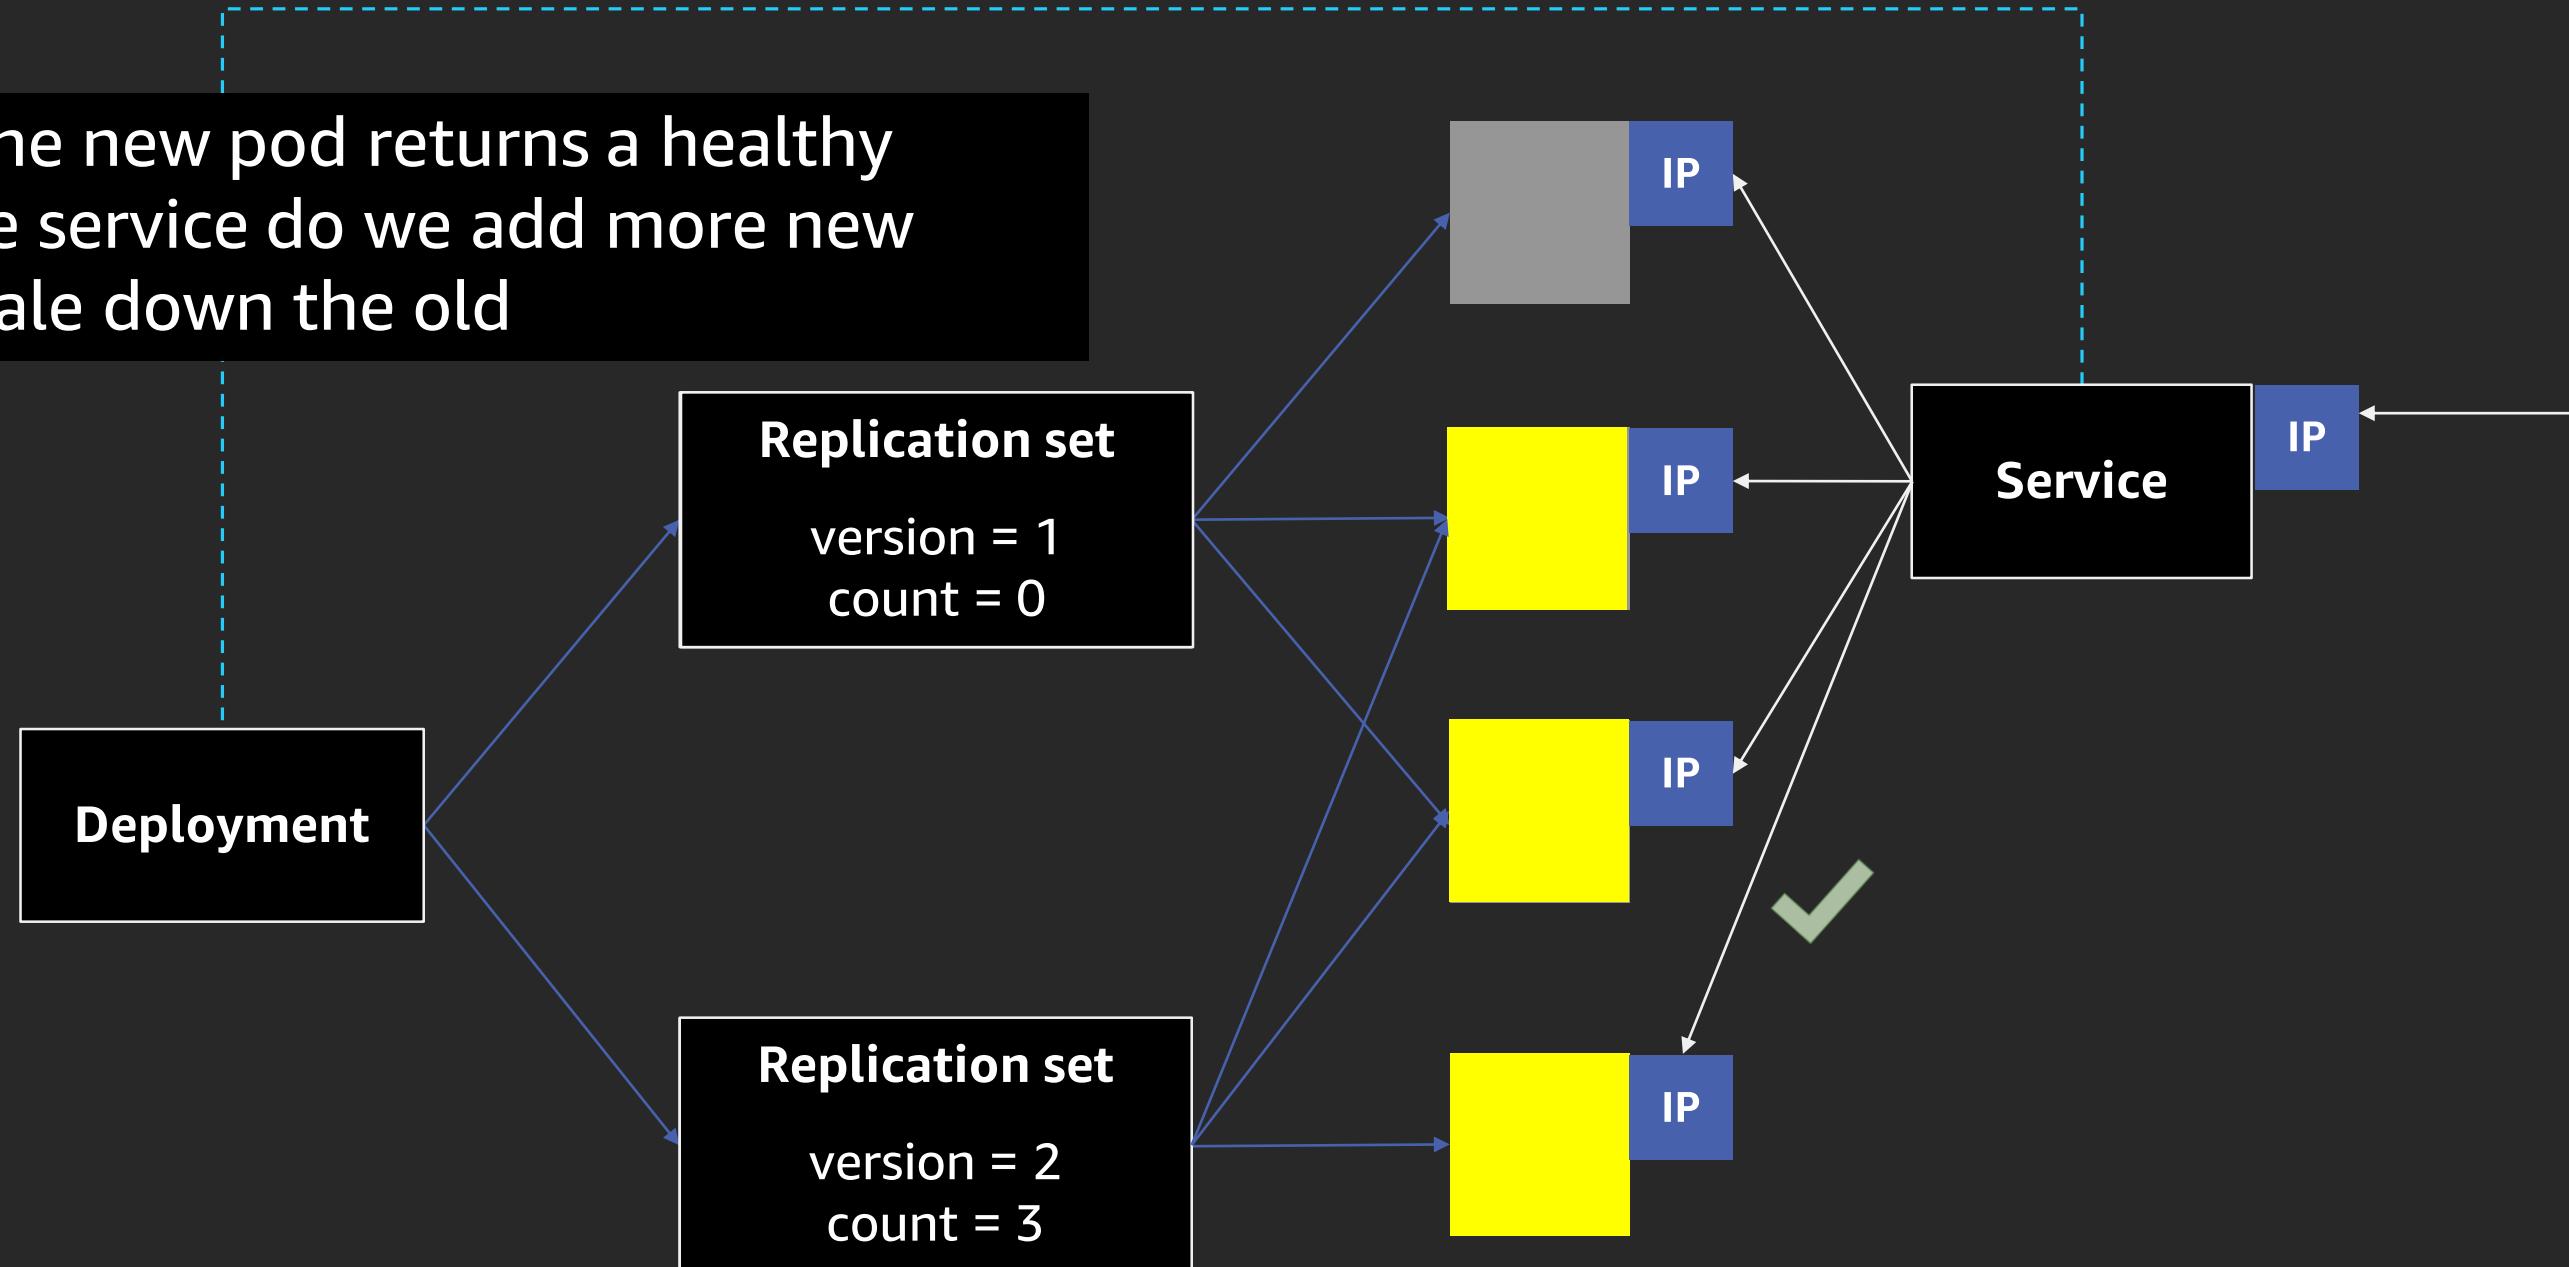
Deployments

The deployment creates a **new replication set** for our new pod version



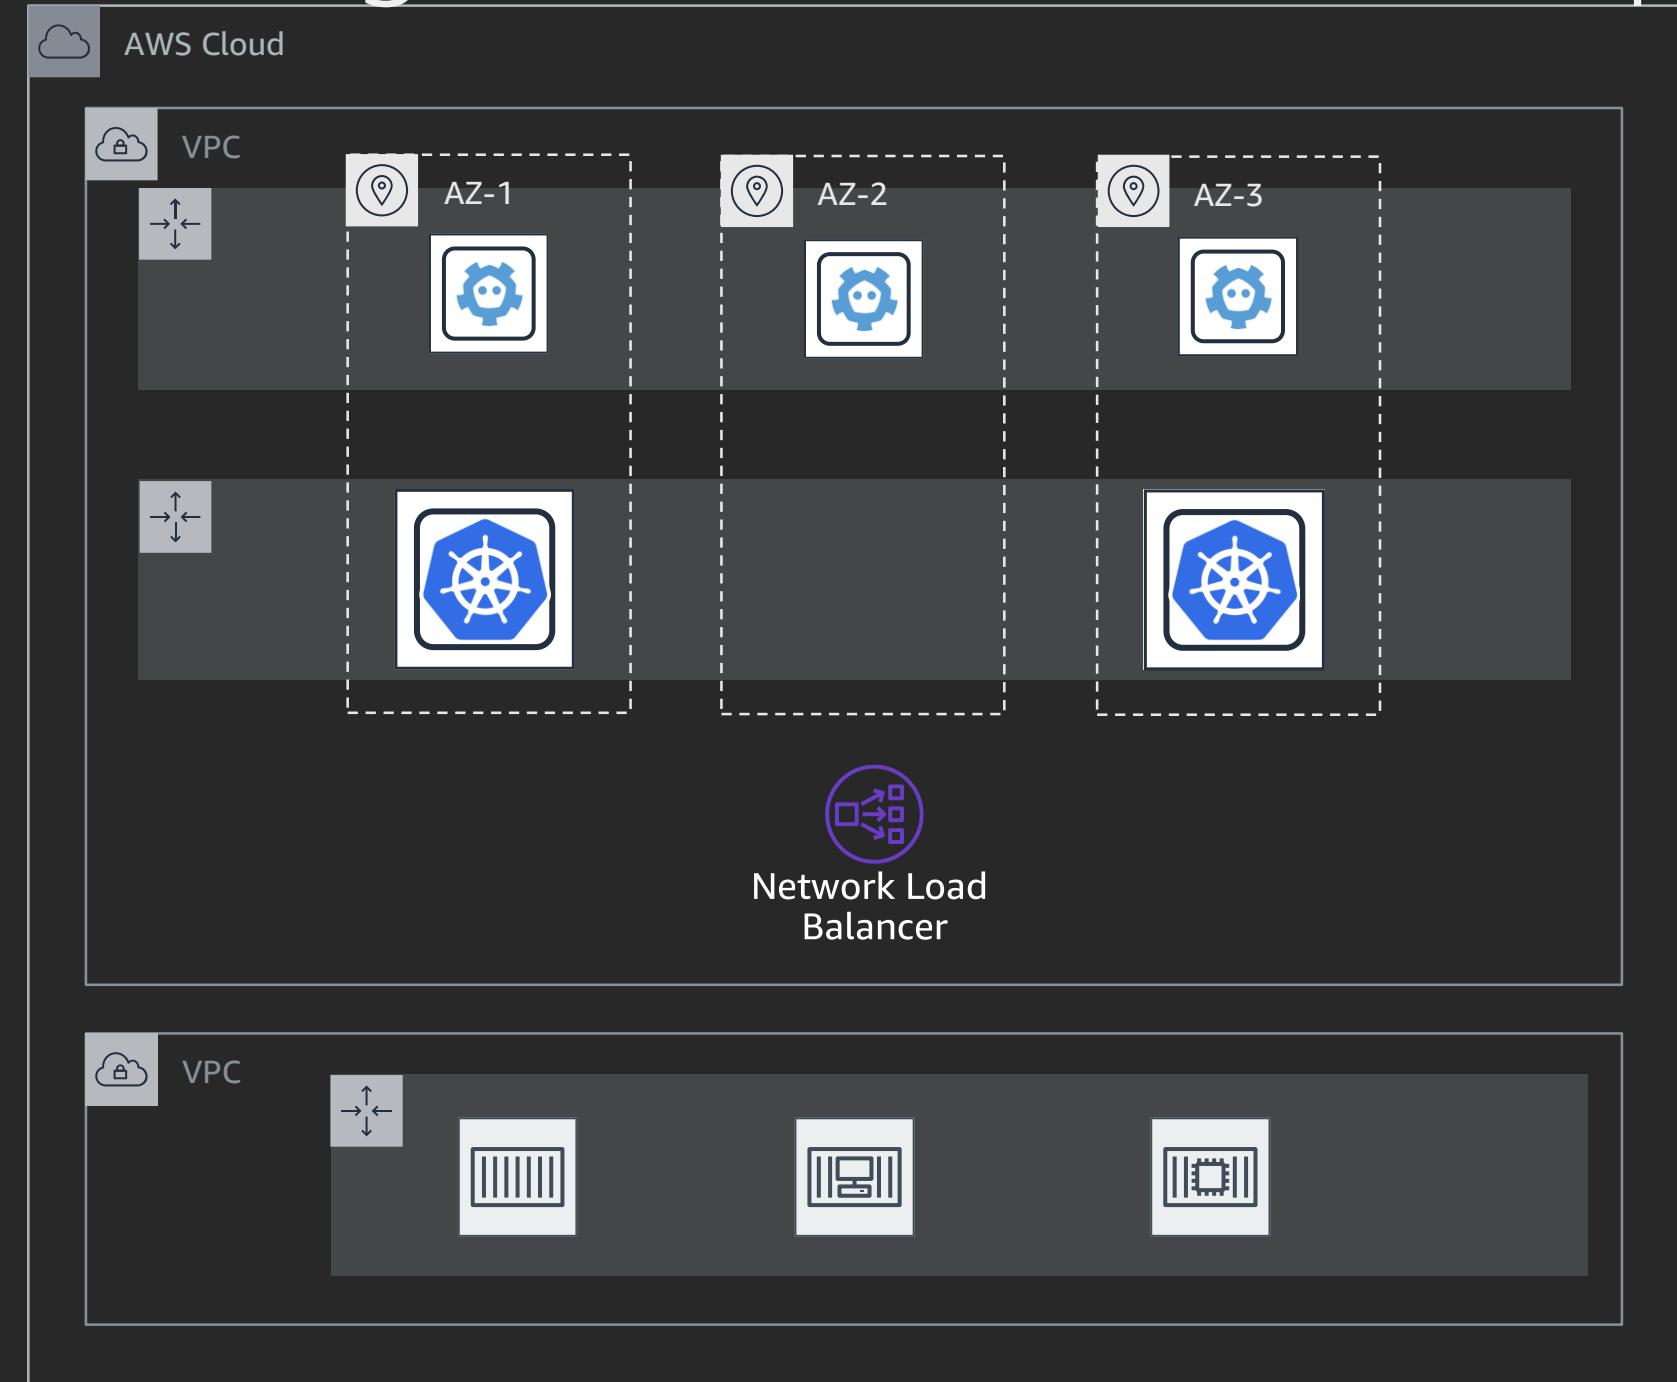
Deployments

Only after the new pod returns a healthy status to the service do we add more new pods and scale down the old



Amazon EKS: Managed Kubernetes control plane

Amazon Elastic
Kubernetes Service
(Amazon EKS)
managed control plane



Amazon EKS

- **Amazon EKS Noun.** \ 'a-mə-,zän, i:,keɪ,es\
 - *Production-ready, highly-available, managed Kubernetes service for AWS*
1. Platform for enterprises to run production-grade workloads
 2. Provides a native and upstream Kubernetes experience
 3. AWS integrations are seamless and eliminate undifferentiated heavy lifting
 4. EKS team actively contributes to the Kubernetes community

AWS: The best place to run Kubernetes

- AWS Managed Services option: Amazon EKS
 - Eksctl: Open source command line management using gitops
- Open source DIY options:
 - Kops
 - Kubeadm
 - Kube-Aws
 - Kubespray
- Partner supported options: MANY
- Mix and match!
 - Use Amazon Linux 2 Amazon EKS AMI with DIY
 - Use commercially supported OS with Amazon EKS
 - Use EKS CNI driver on DIY or partner or open source CNI with Amazon EKS

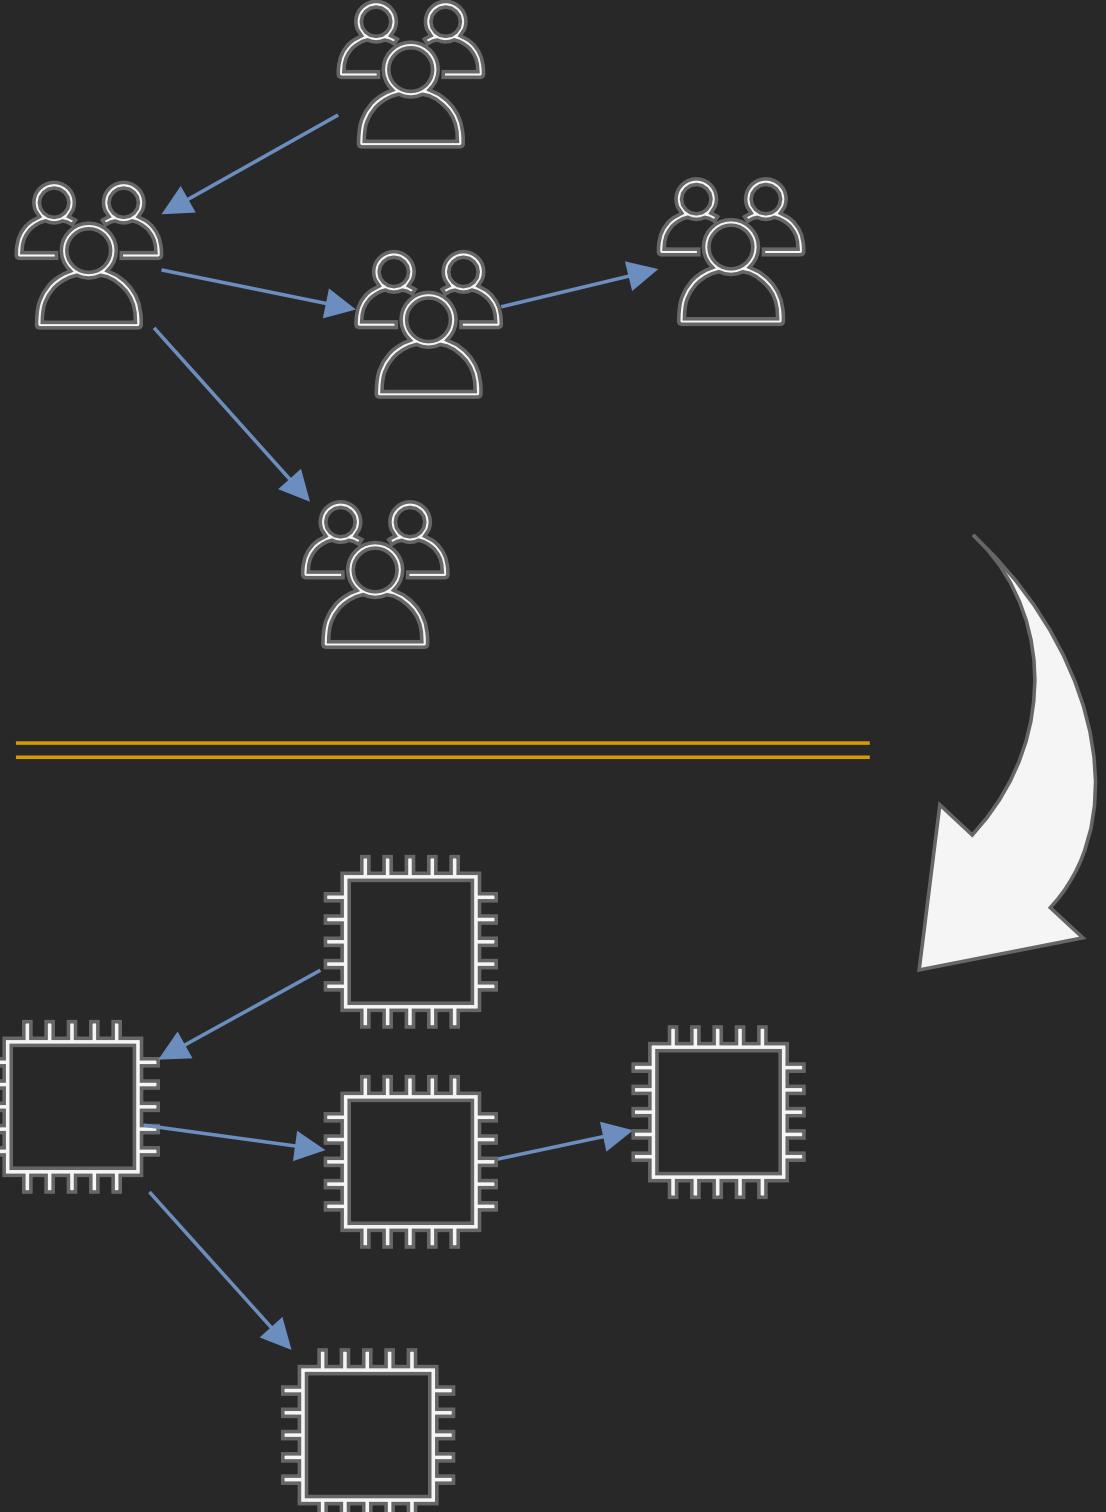
From monoliths to microservices: The rendering problem

From monoliths...

Two-pizza teams

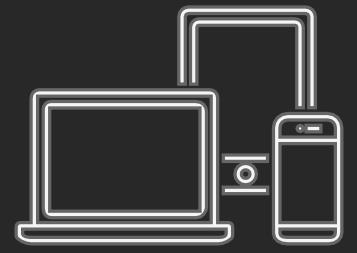
Service-oriented architecture

Traditionally data-oriented services

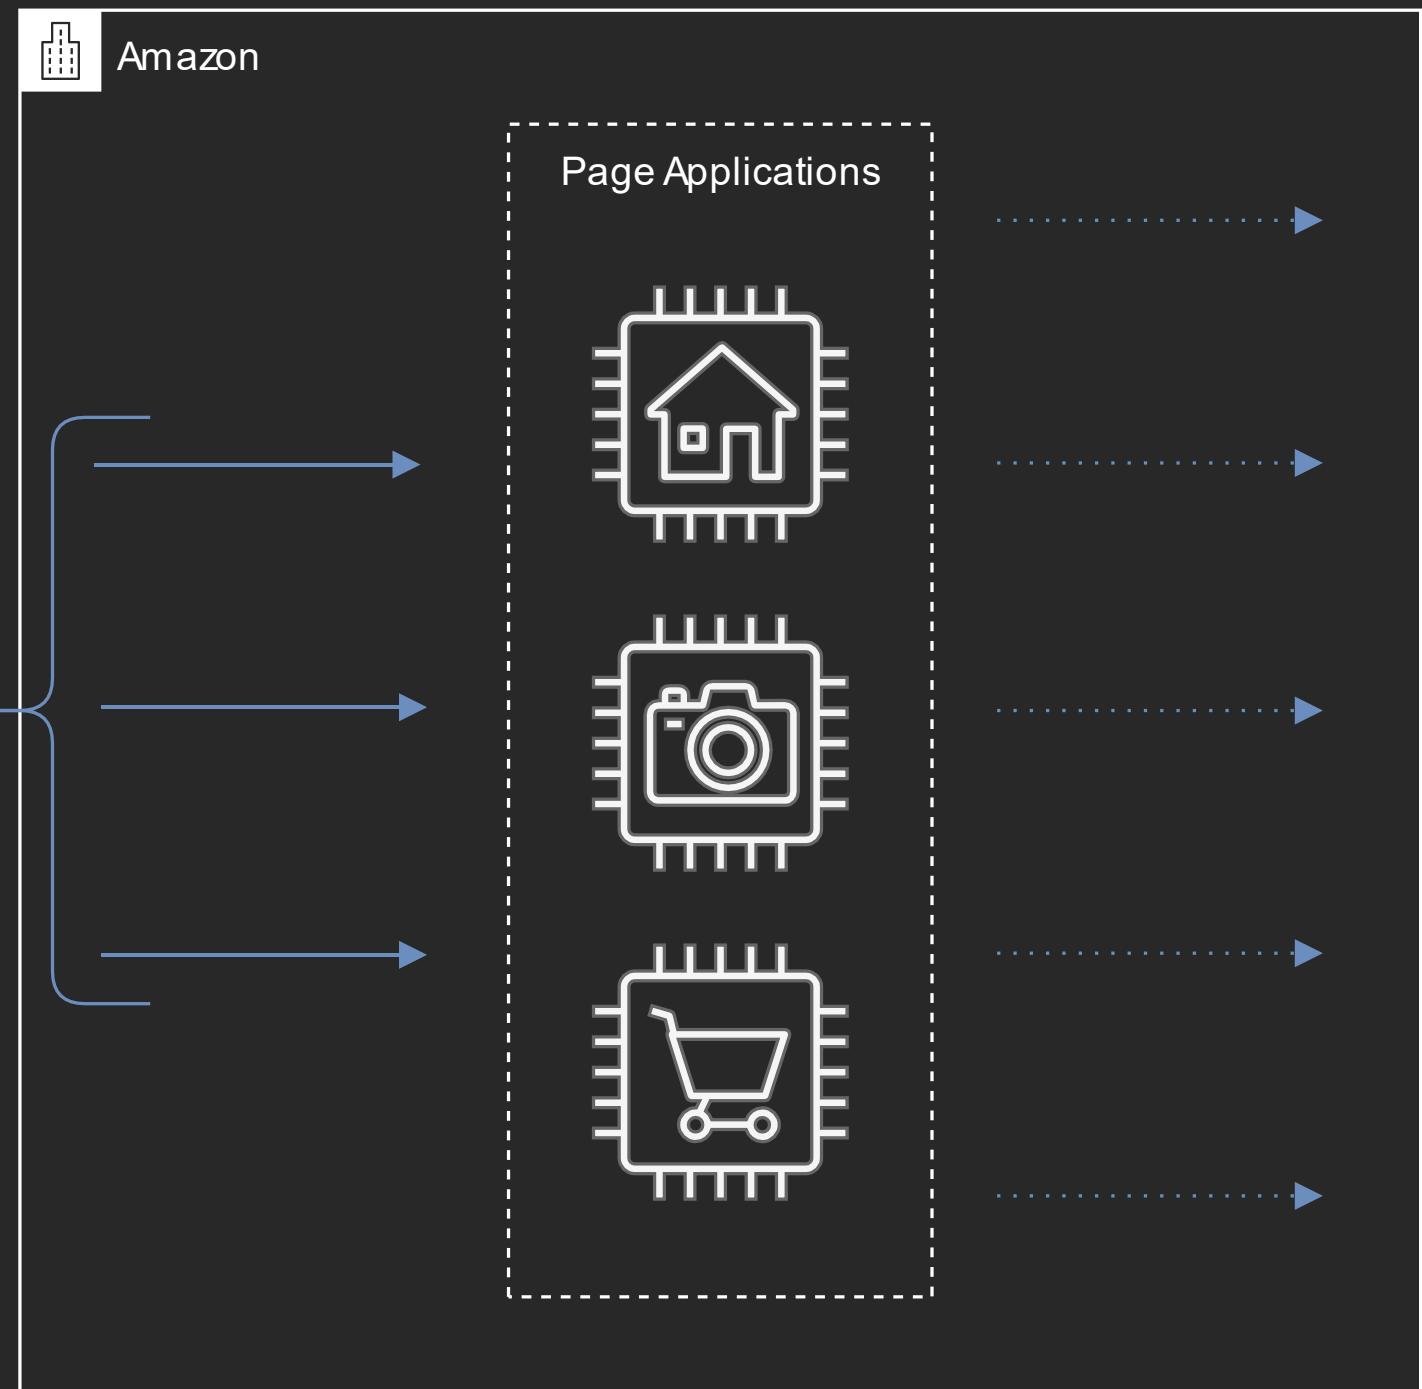


From monoliths...

Shopping experiences are organized around monolithic applications

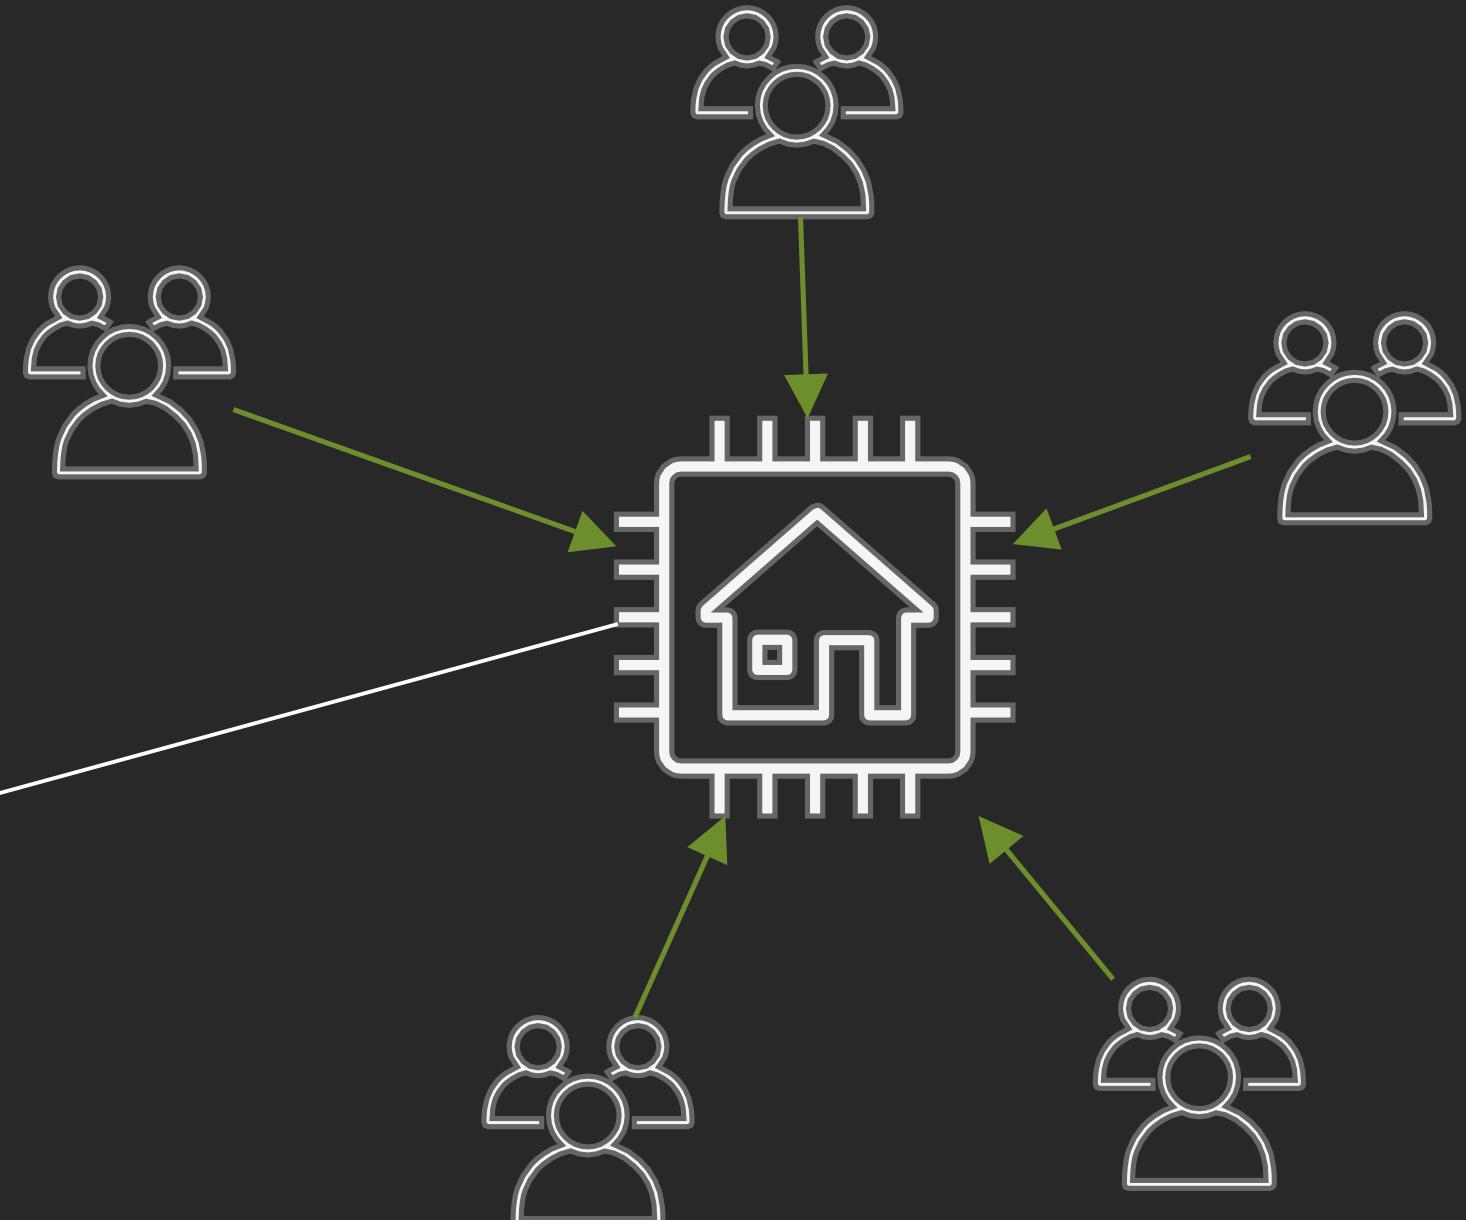
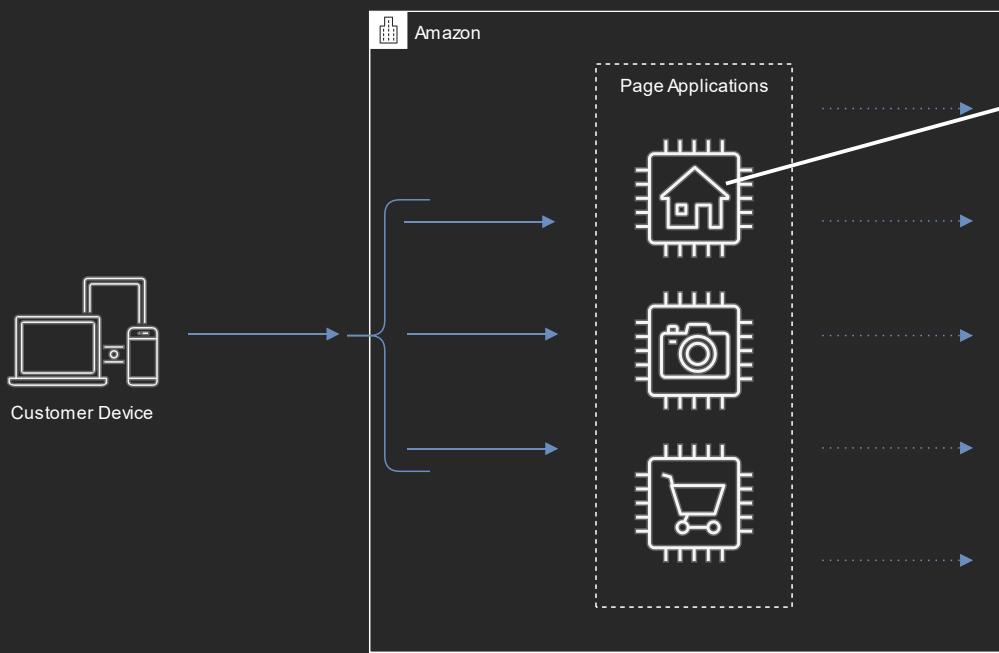


Customer Device



From monoliths...

Shopping experiences are organized around monolithic applications





Break it apart

BOOM

...To microservices

Independent release cycle

Rapid experimentation

Isolation across features

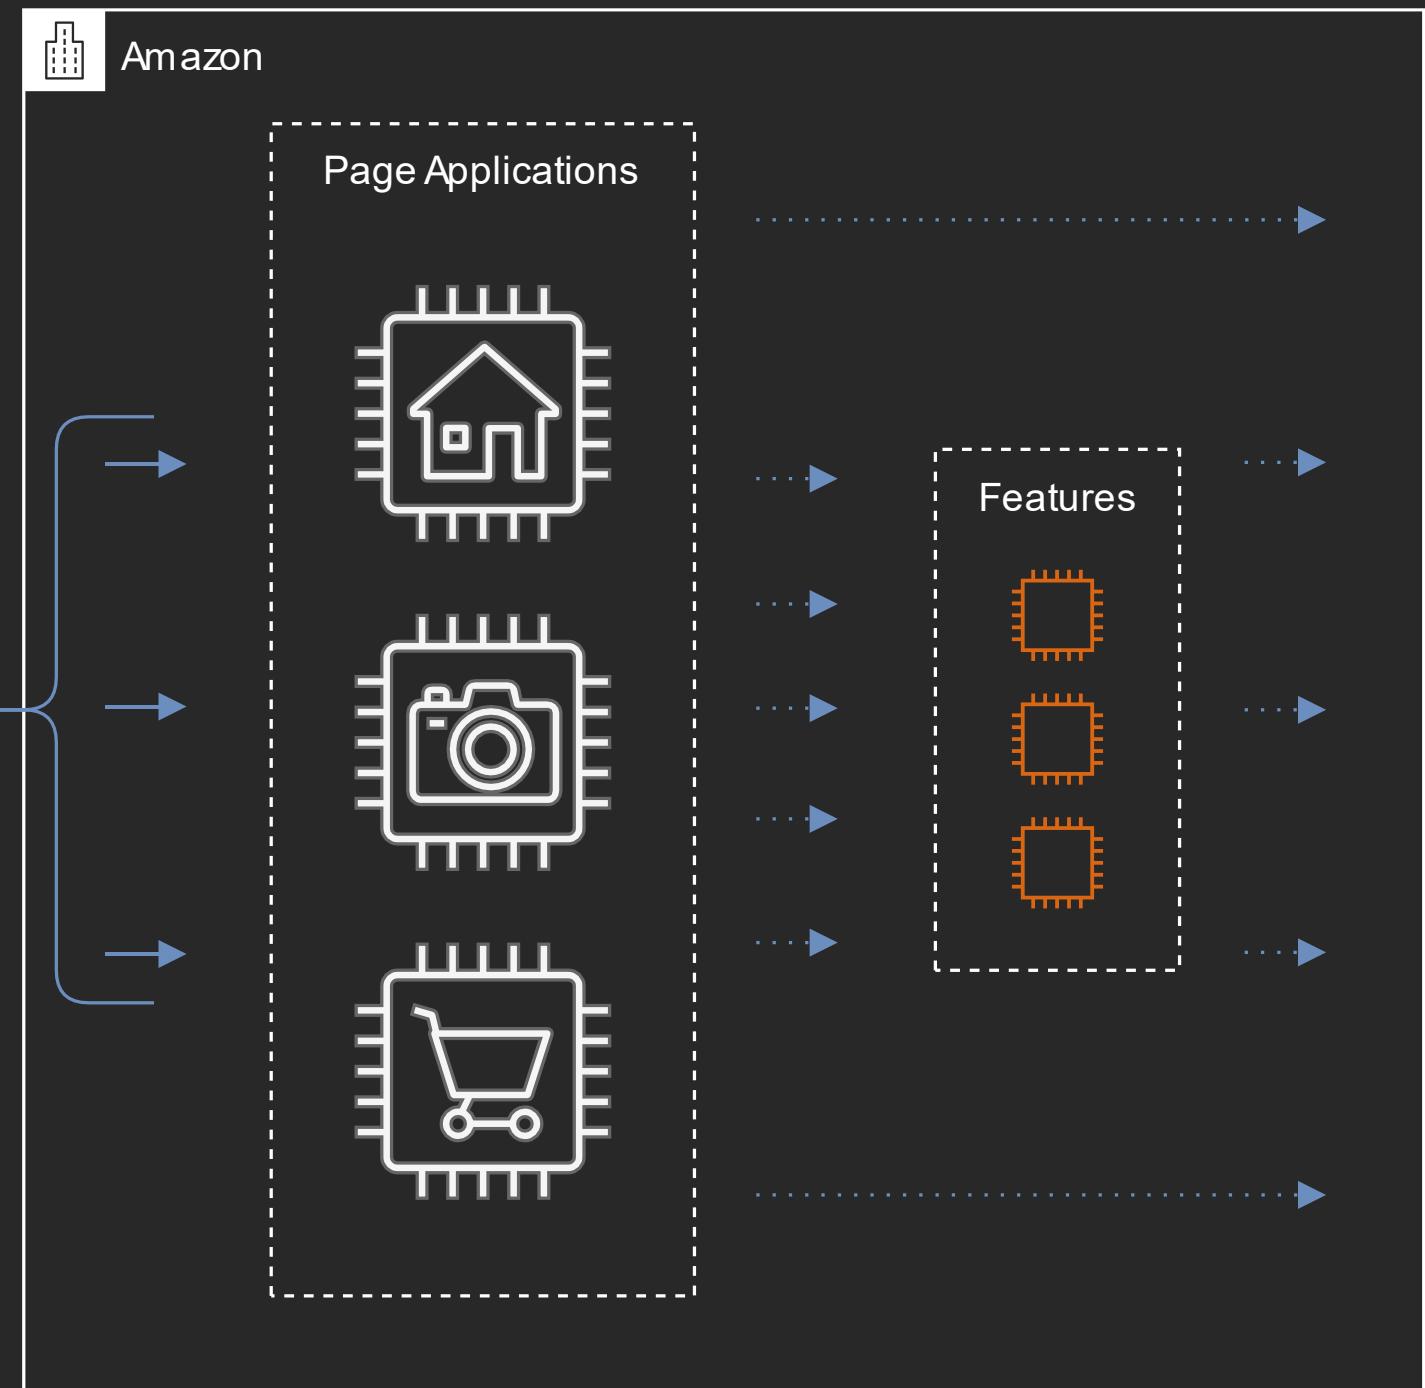


...To microservices

Decompose the shopping experience



Customer Device

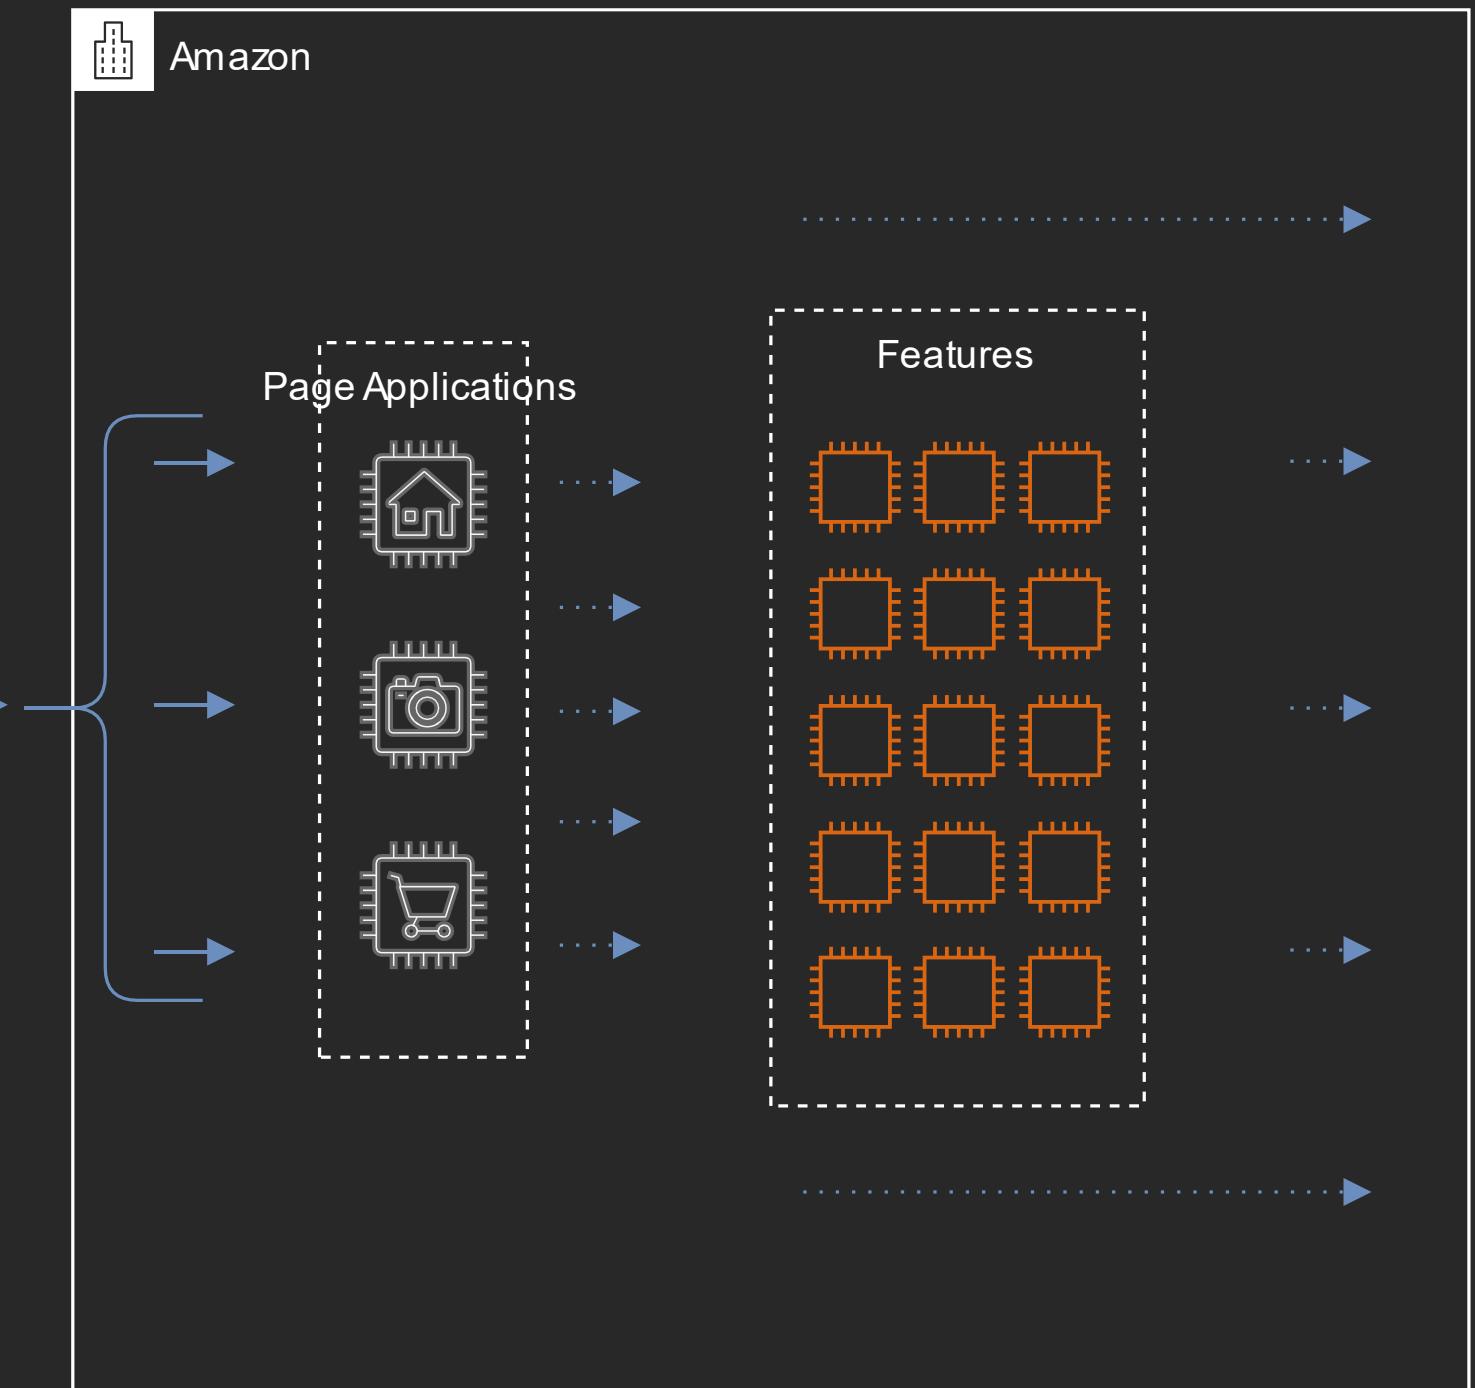


...To microservices

Decompose the shopping experience



Customer Device



Hello
Select your address

Best Sellers Gift Ideas New Releases Whole Foods Today's Deals

Shop Halloween décor

Deals, movies and music
all in one membership
prime

Deal of the Day



Women's shop by style



Cool



Casual



Classic



Boho

bissell

ICONpet
CORDLESS VACUUM

Save \$100
on your ICONpet purchase

Conditions apply.

Bissell ICONpet Cordless Stick Vacuum

★★★★★ 27

\$299.99

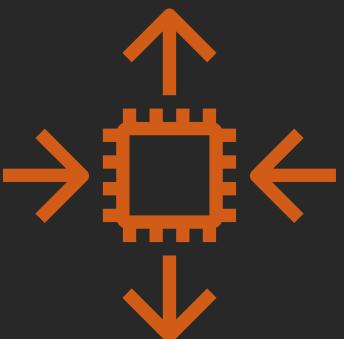
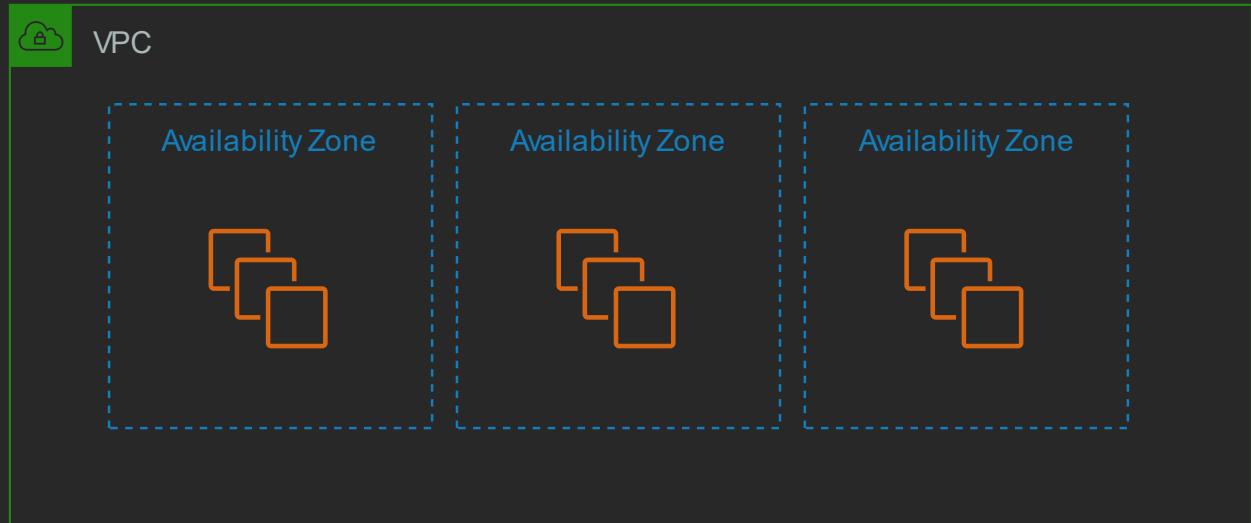
Add to Cart

Ad feedback

...To microservices

Hosted environment requirements

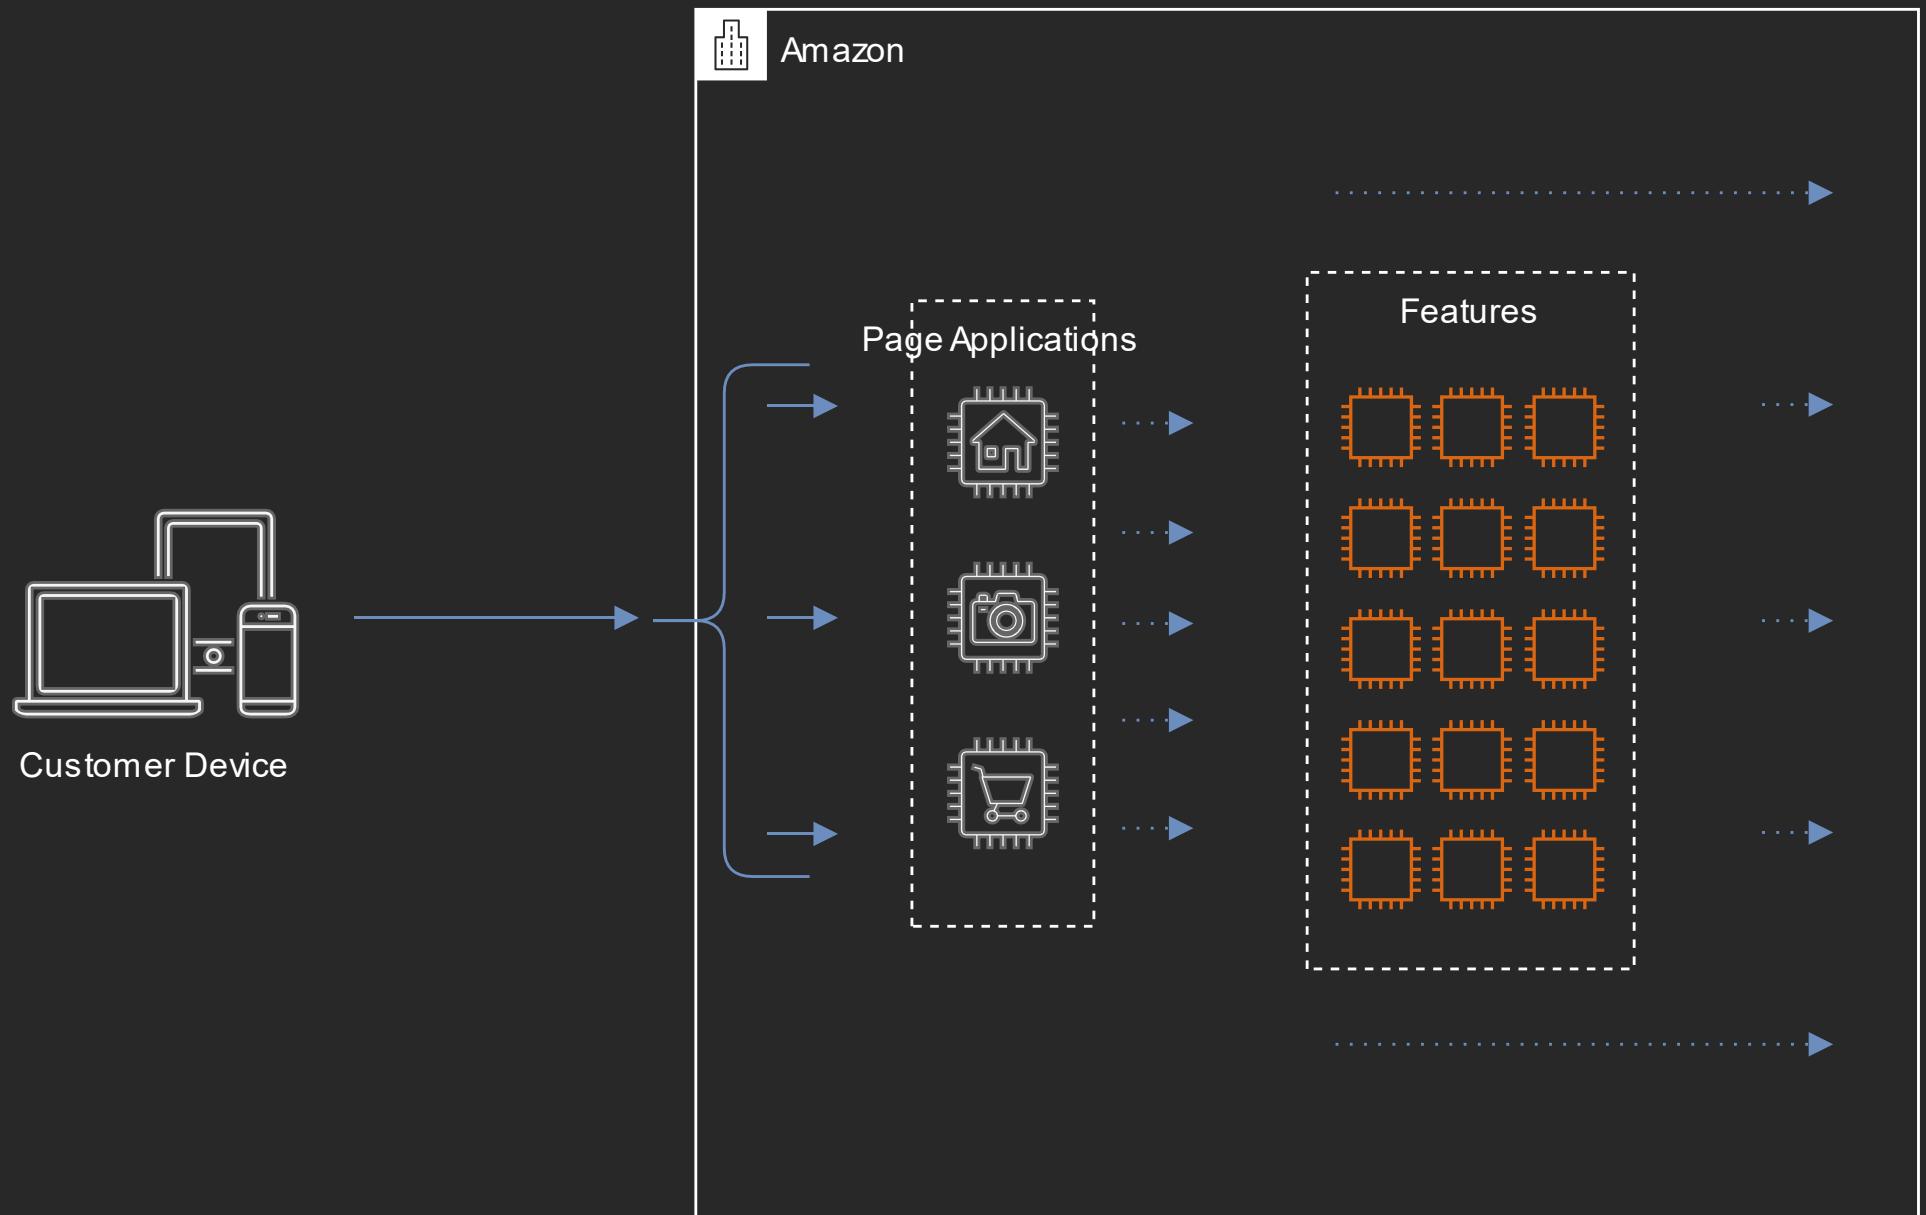
- Availability
- Scalability
- Reliability



...To microservices

Docker containers

Kubernetes clusters

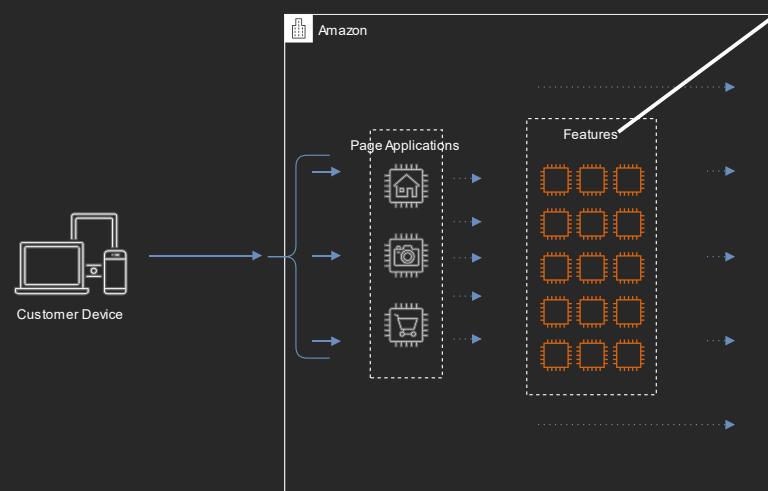


Improving ownership by owning less

Improving ownership by owning less

How did Amazon EKS help?

- Managed updates
- etcd management
- Masters scaling



Improving ownership by owning less

Managed updates

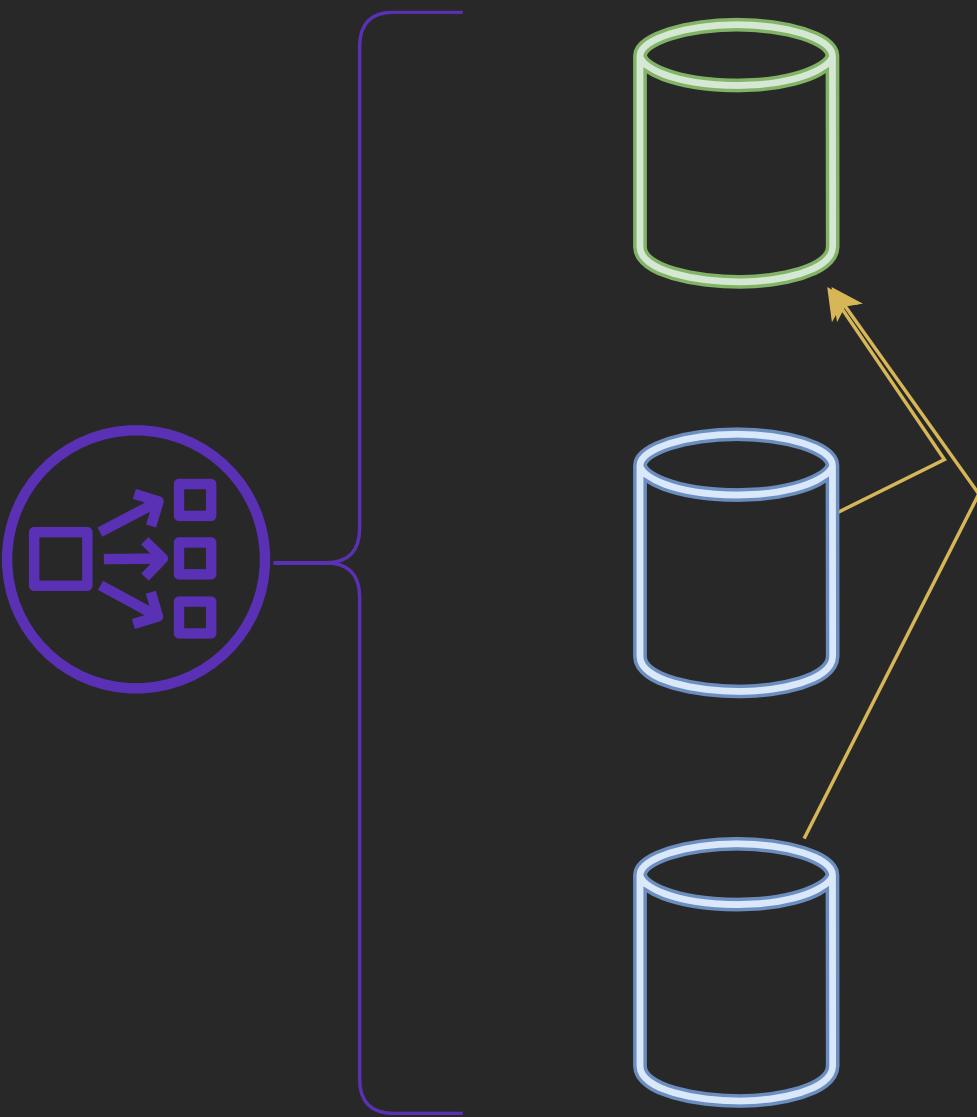
Common vulnerabilities and exposures (CVE)

Supported version updates



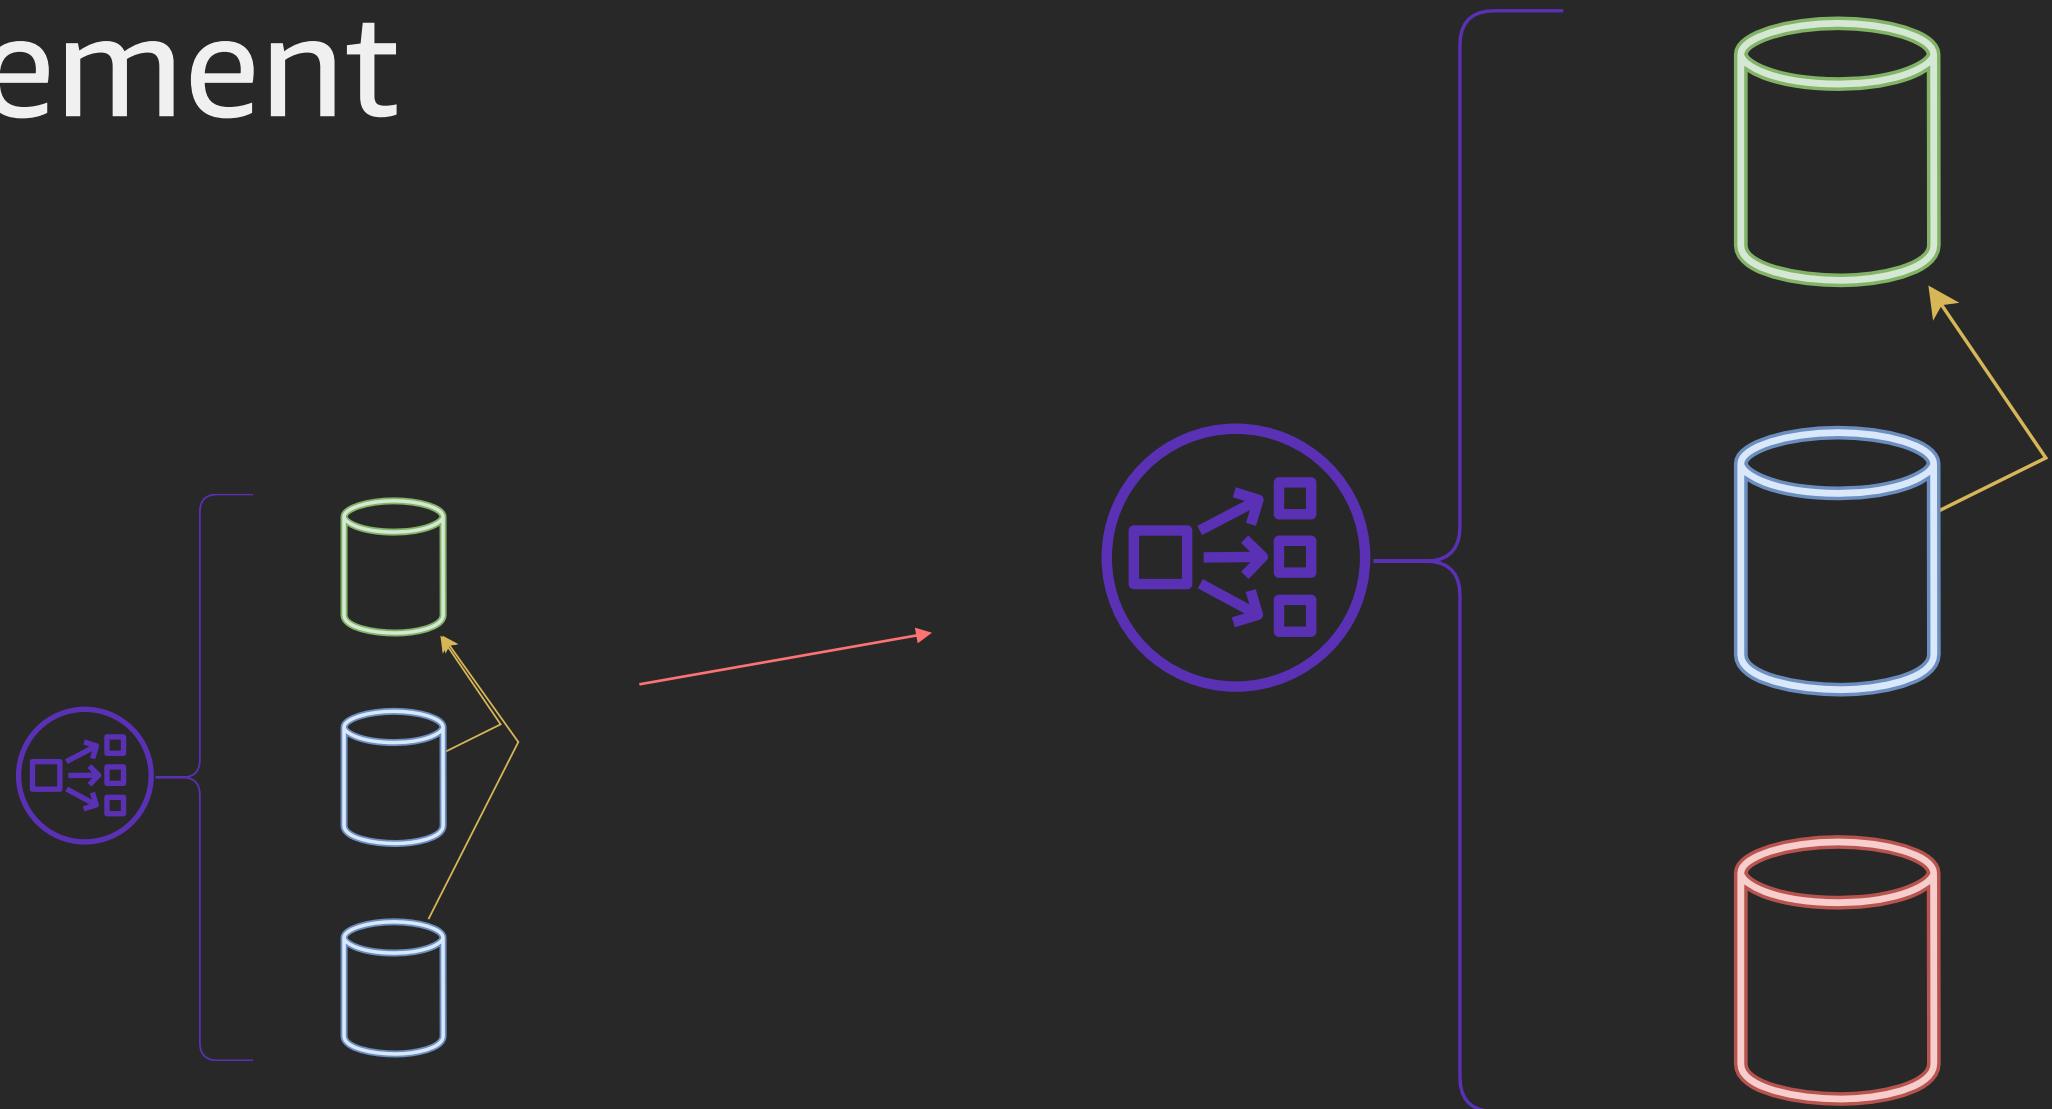
Improving ownership by owning less

etcd management



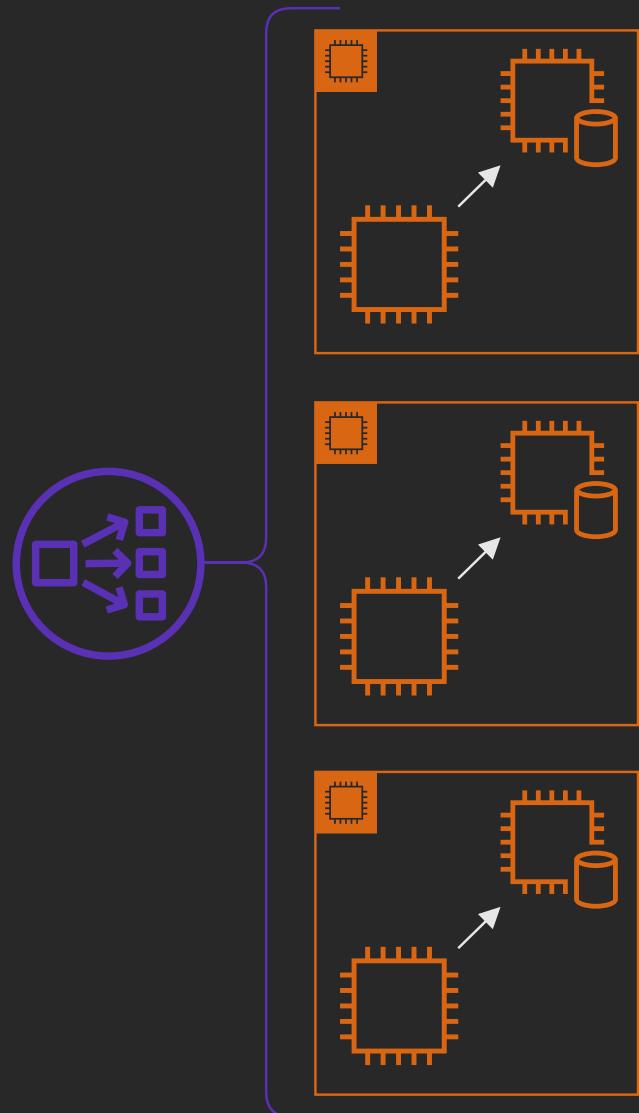
Improving ownership by owning less

etcd management



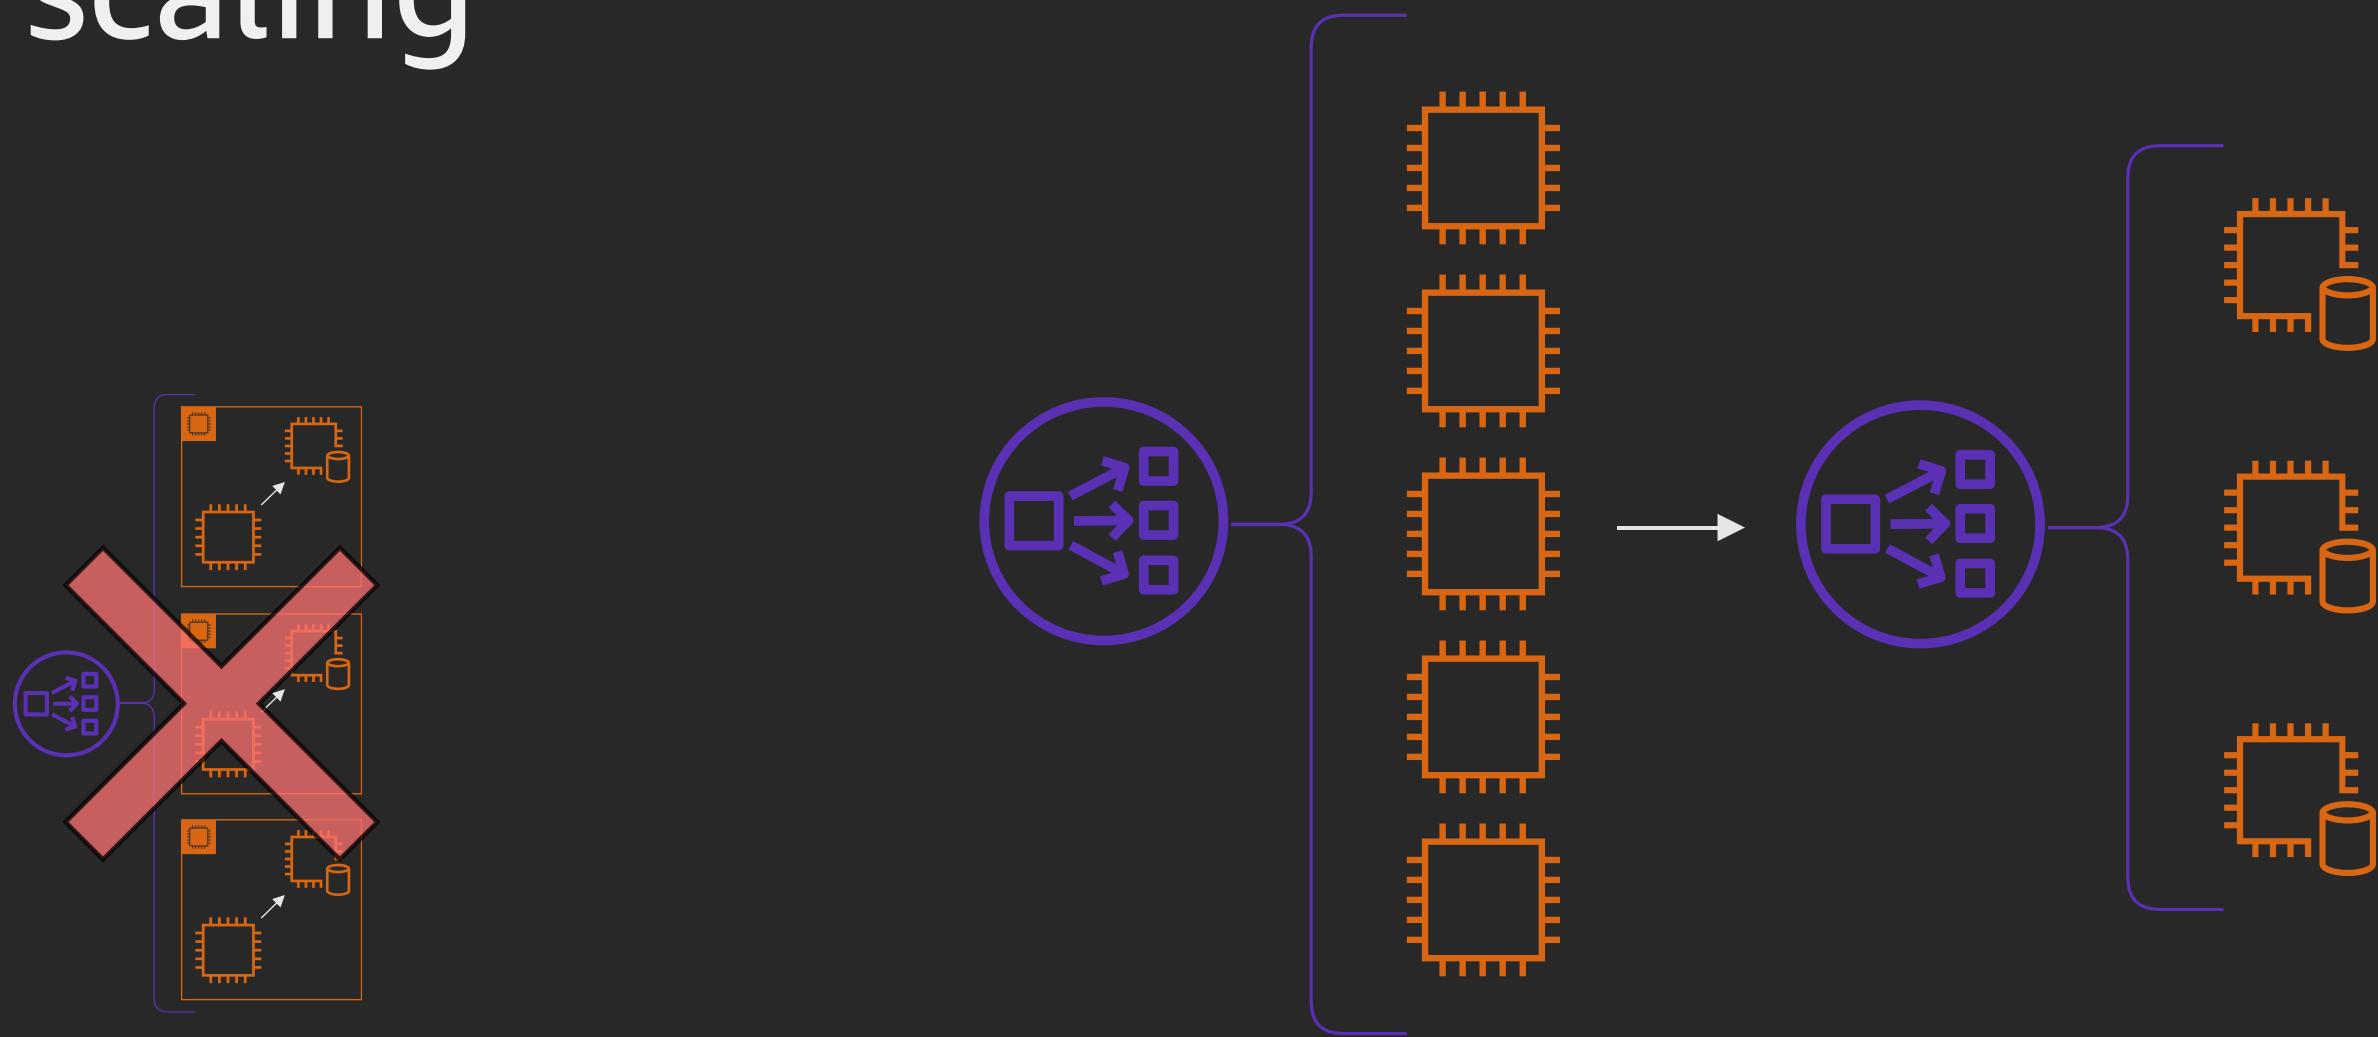
Improving ownership by owning less

Masters scaling



Improving ownership by owning less

Masters scaling

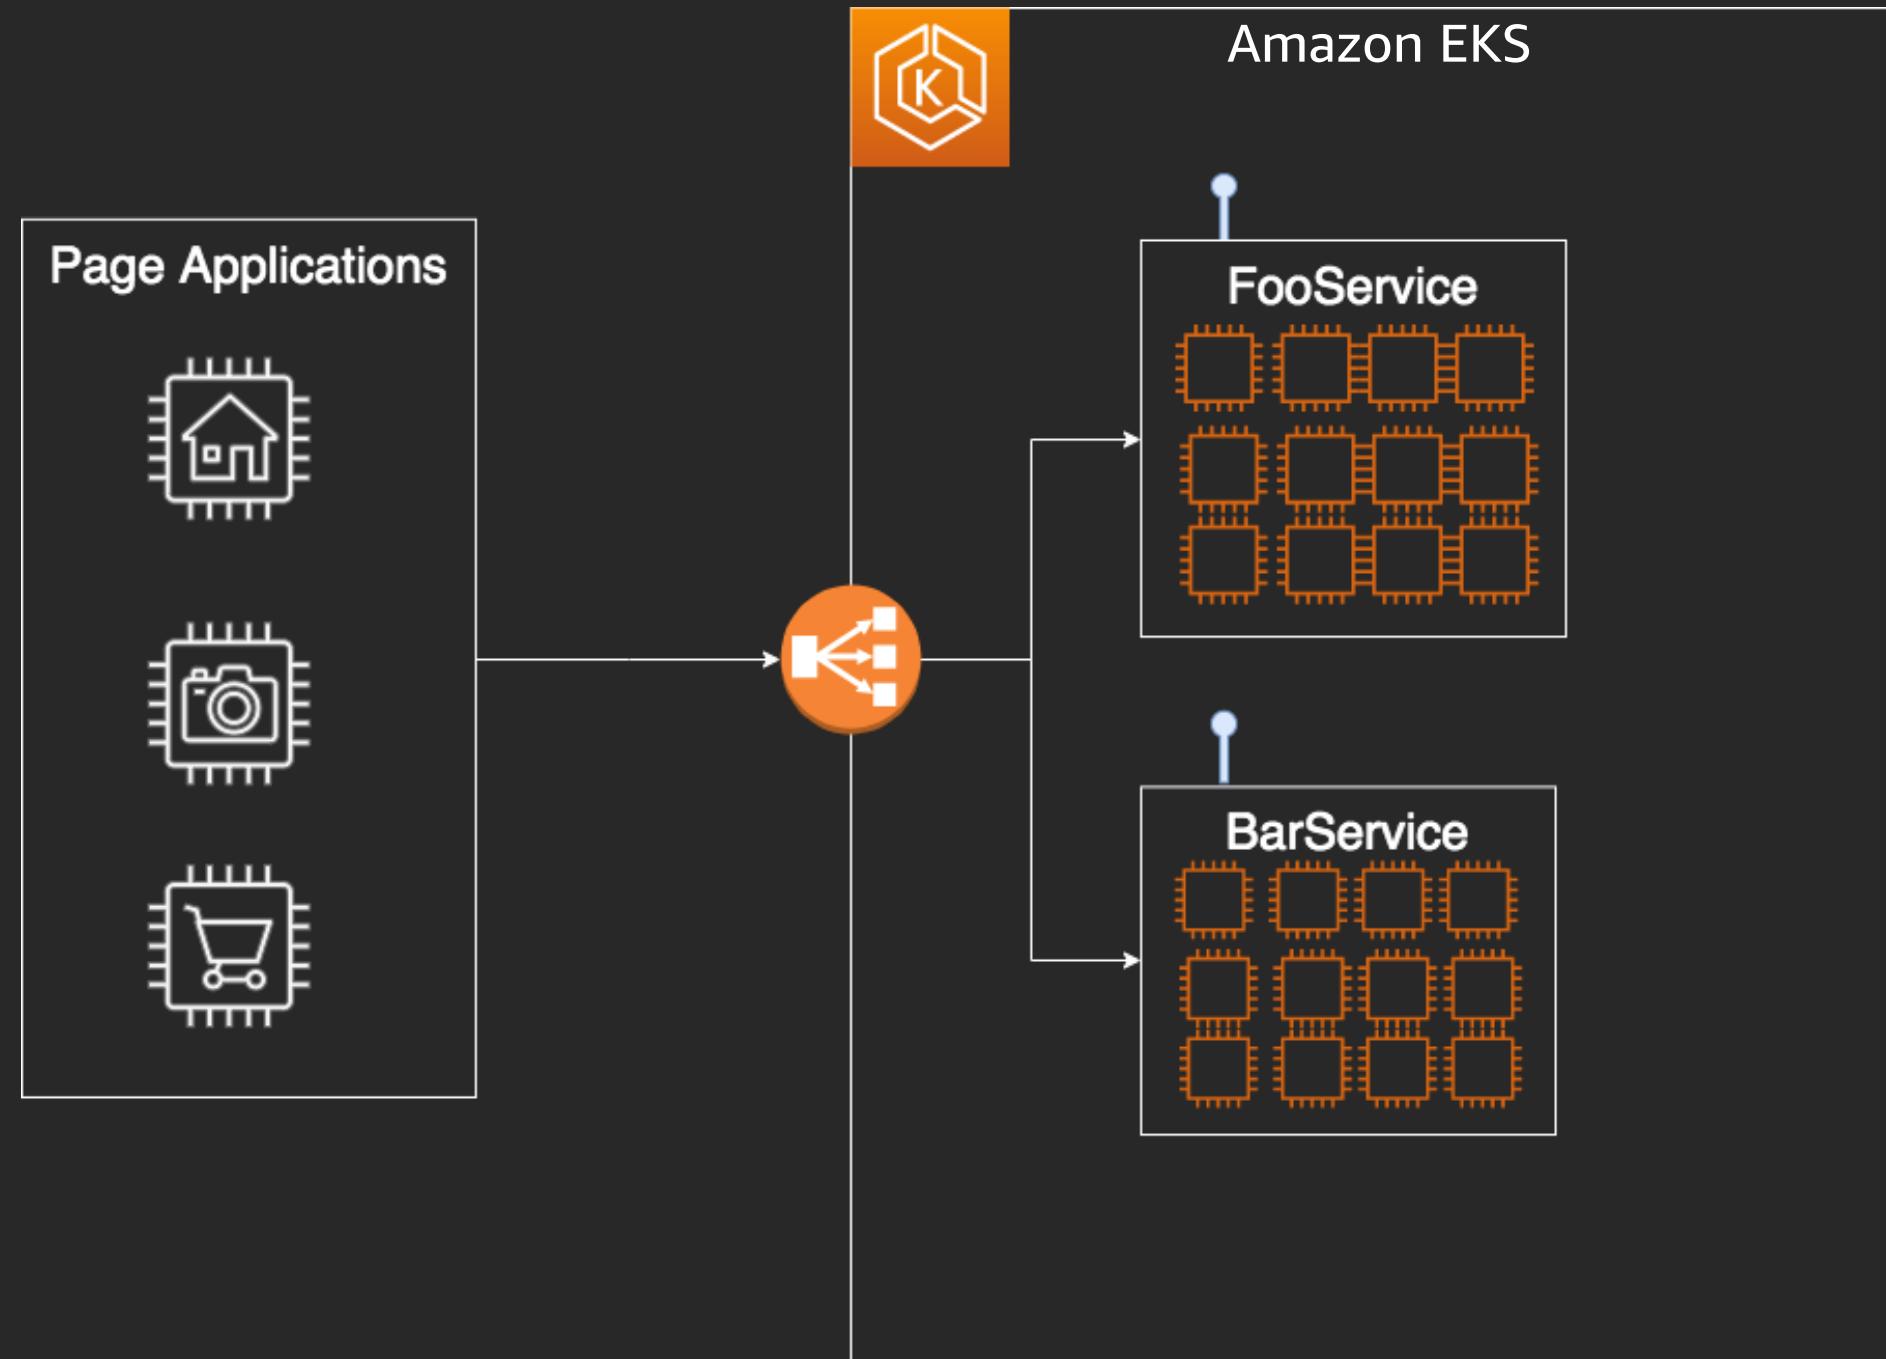


A photograph of a person's hands working on a green and red LEGO robot. The robot has a black motor and a blue gear. In the background, a computer monitor displays a video editing interface with multiple windows and tools. A keyboard and a mouse are also visible on the desk.

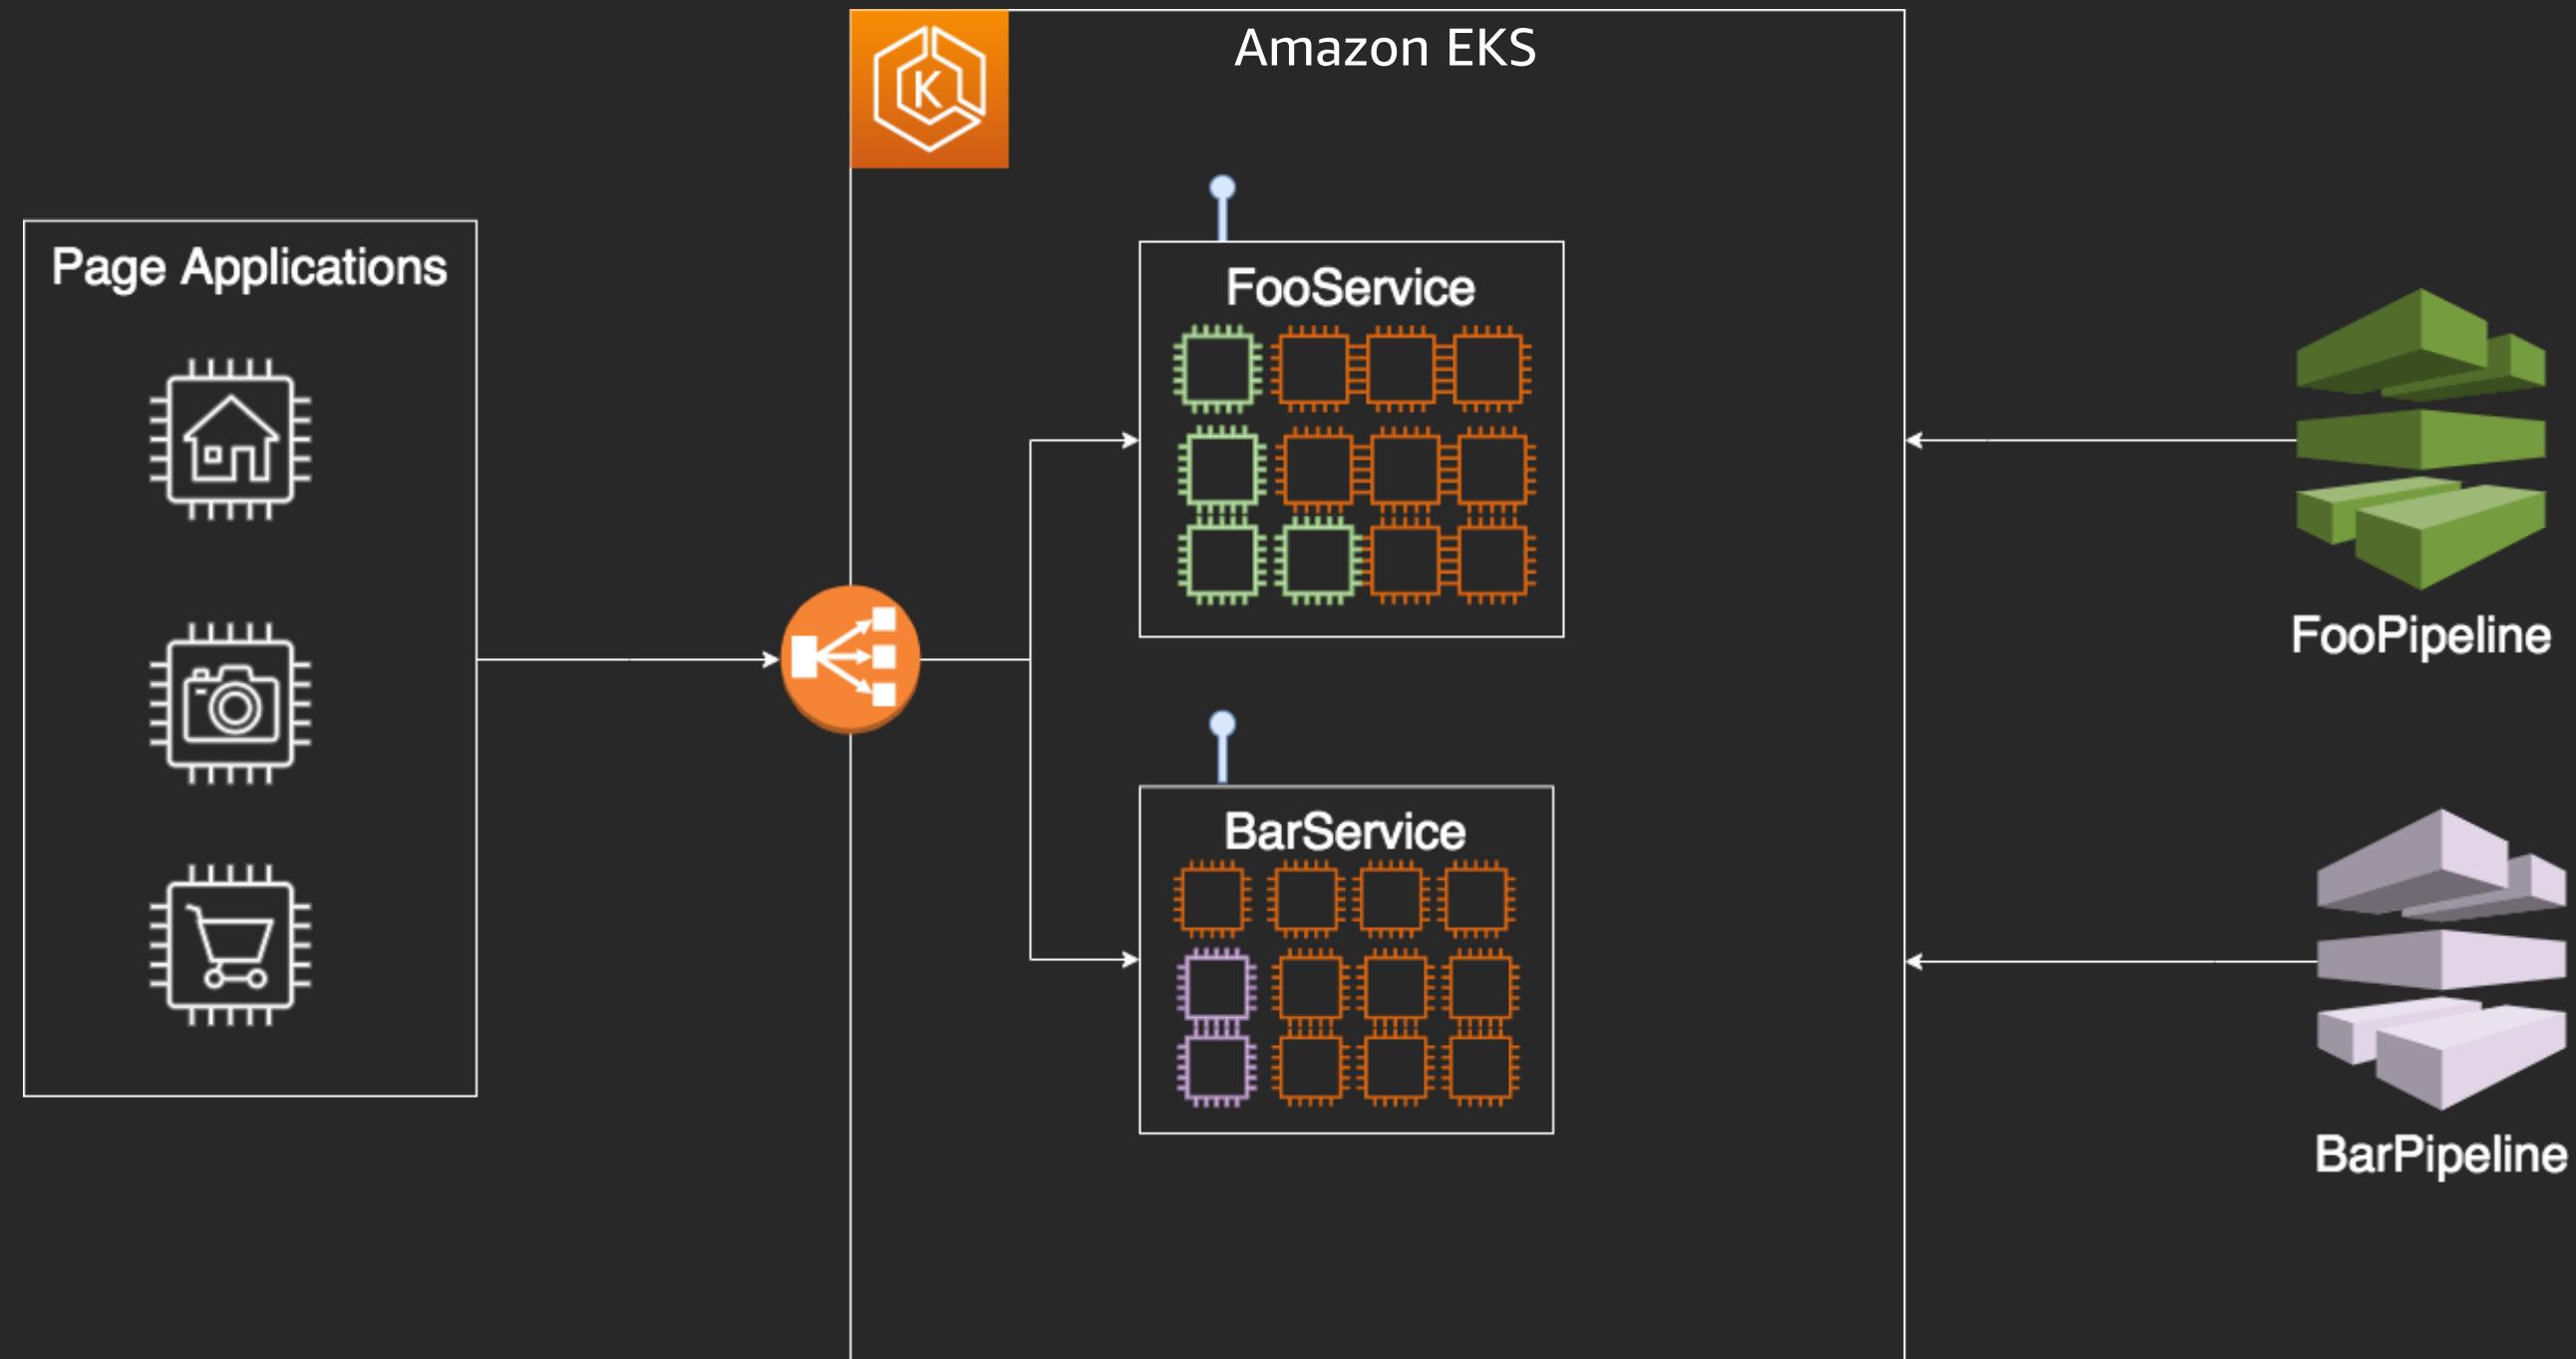
Focus on the problem

CI/CD: Release features all the time!

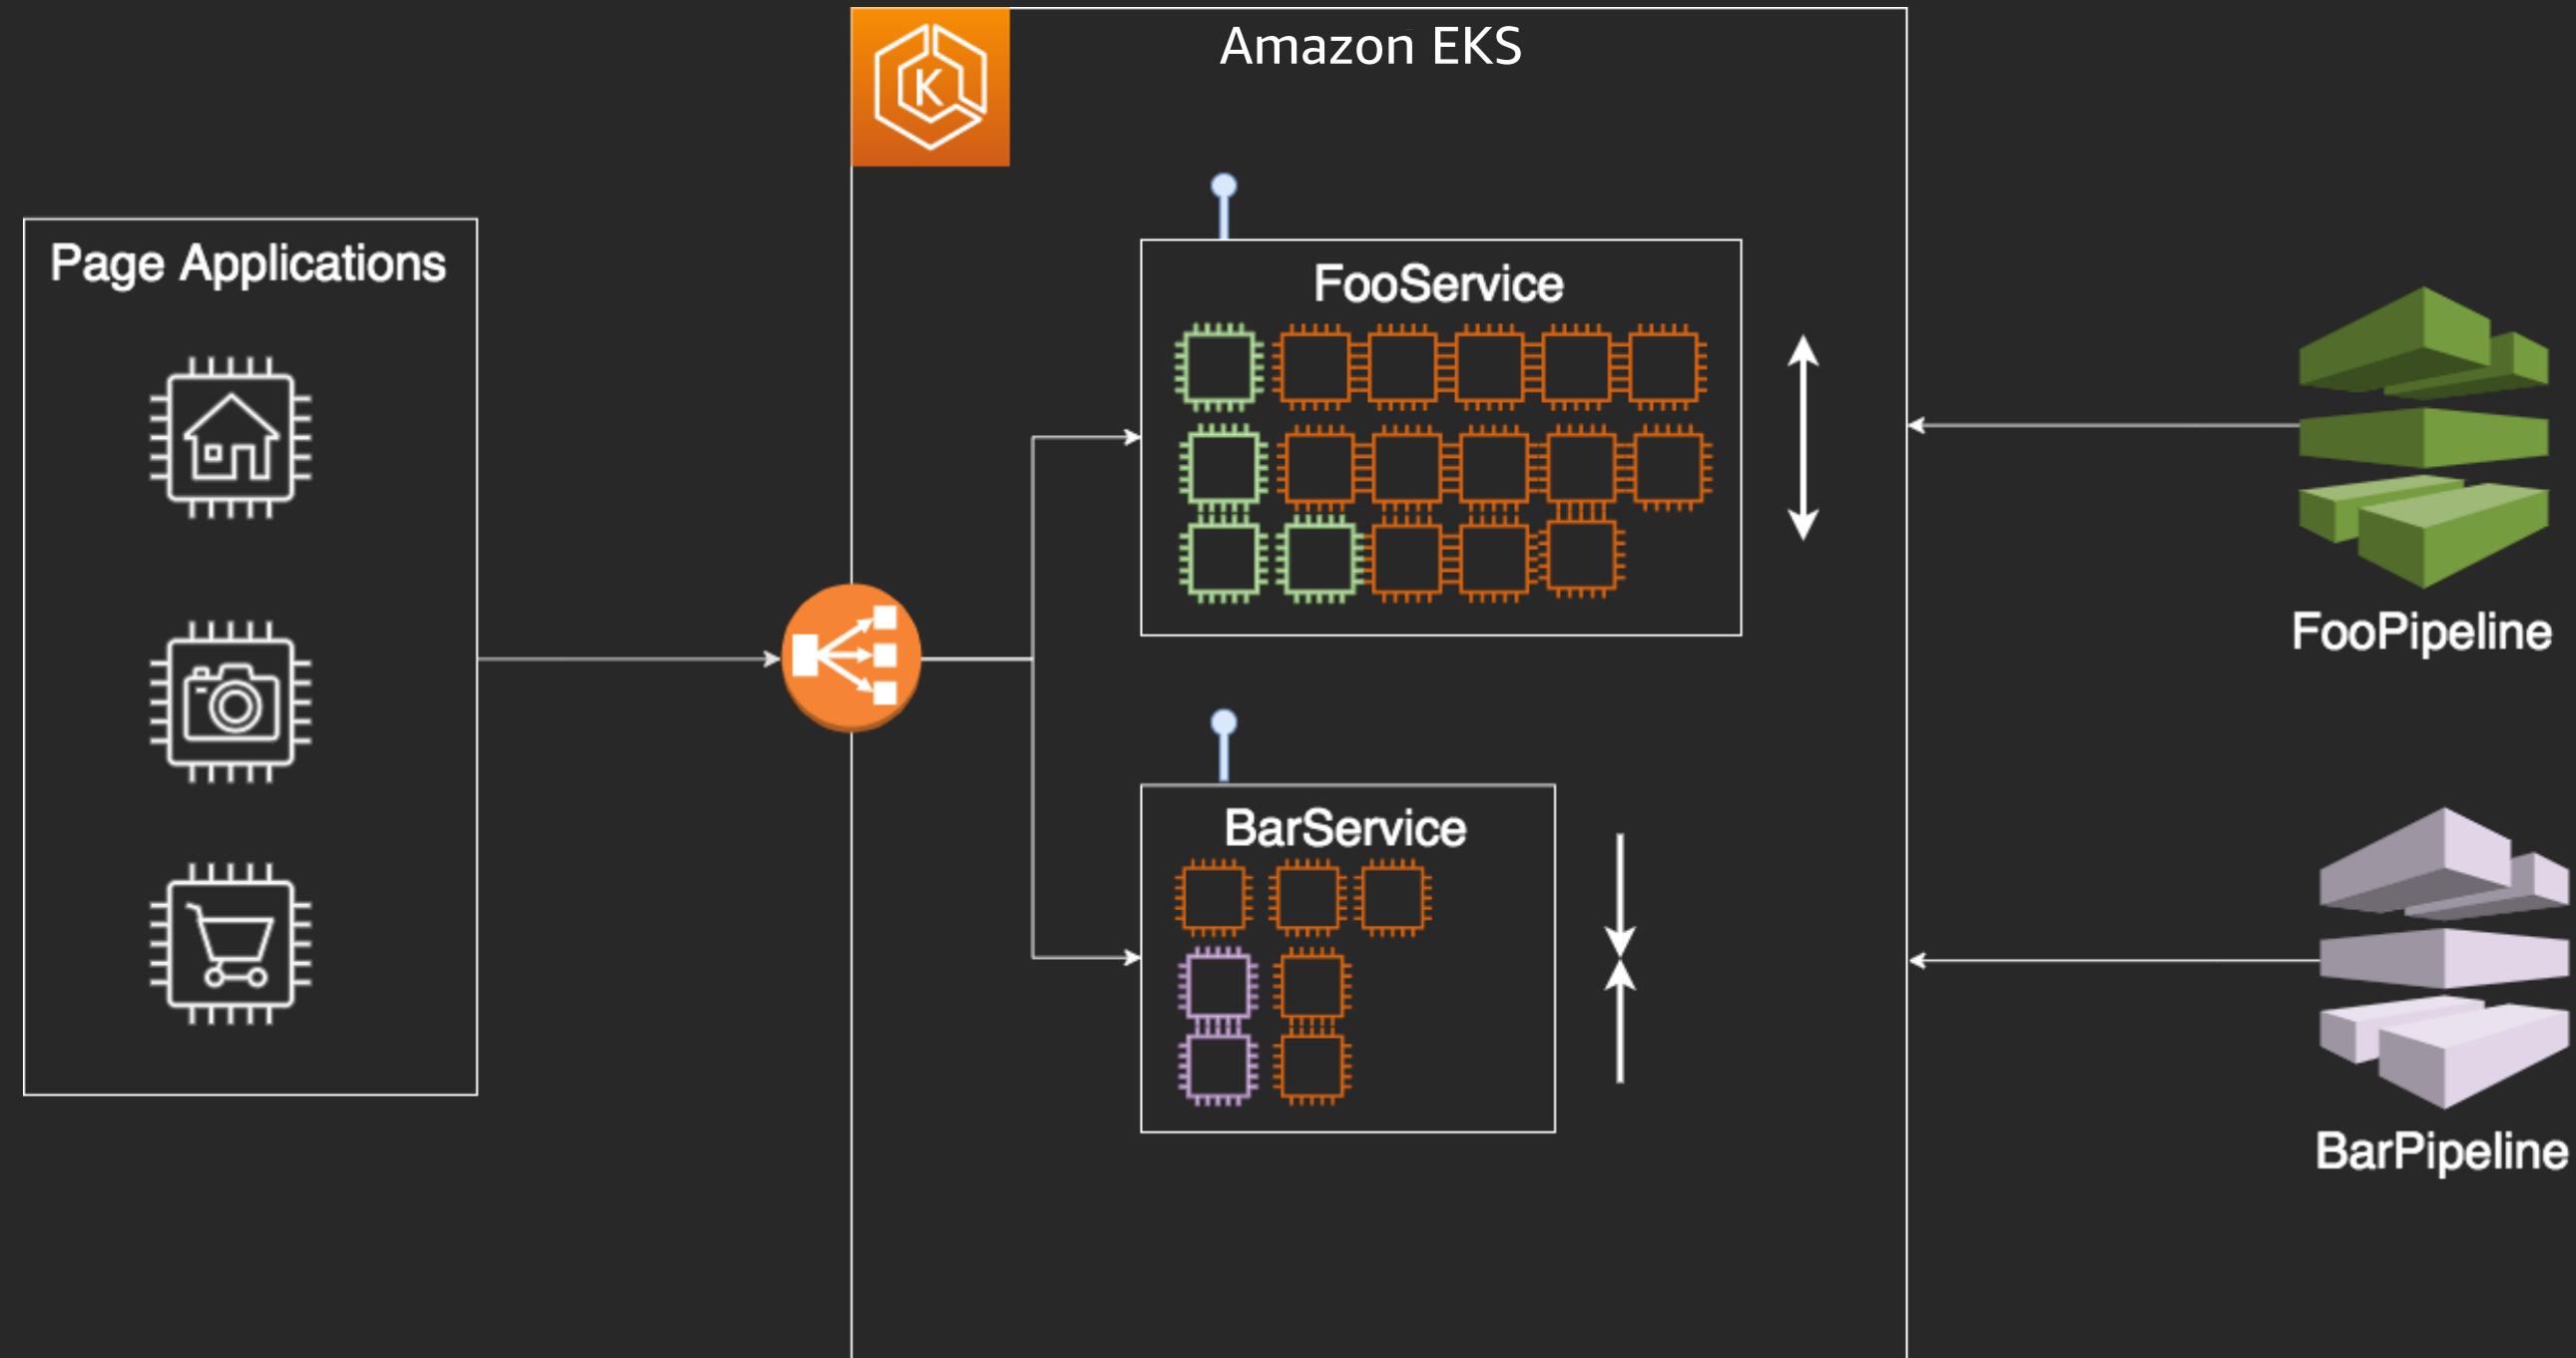
Containerized applications



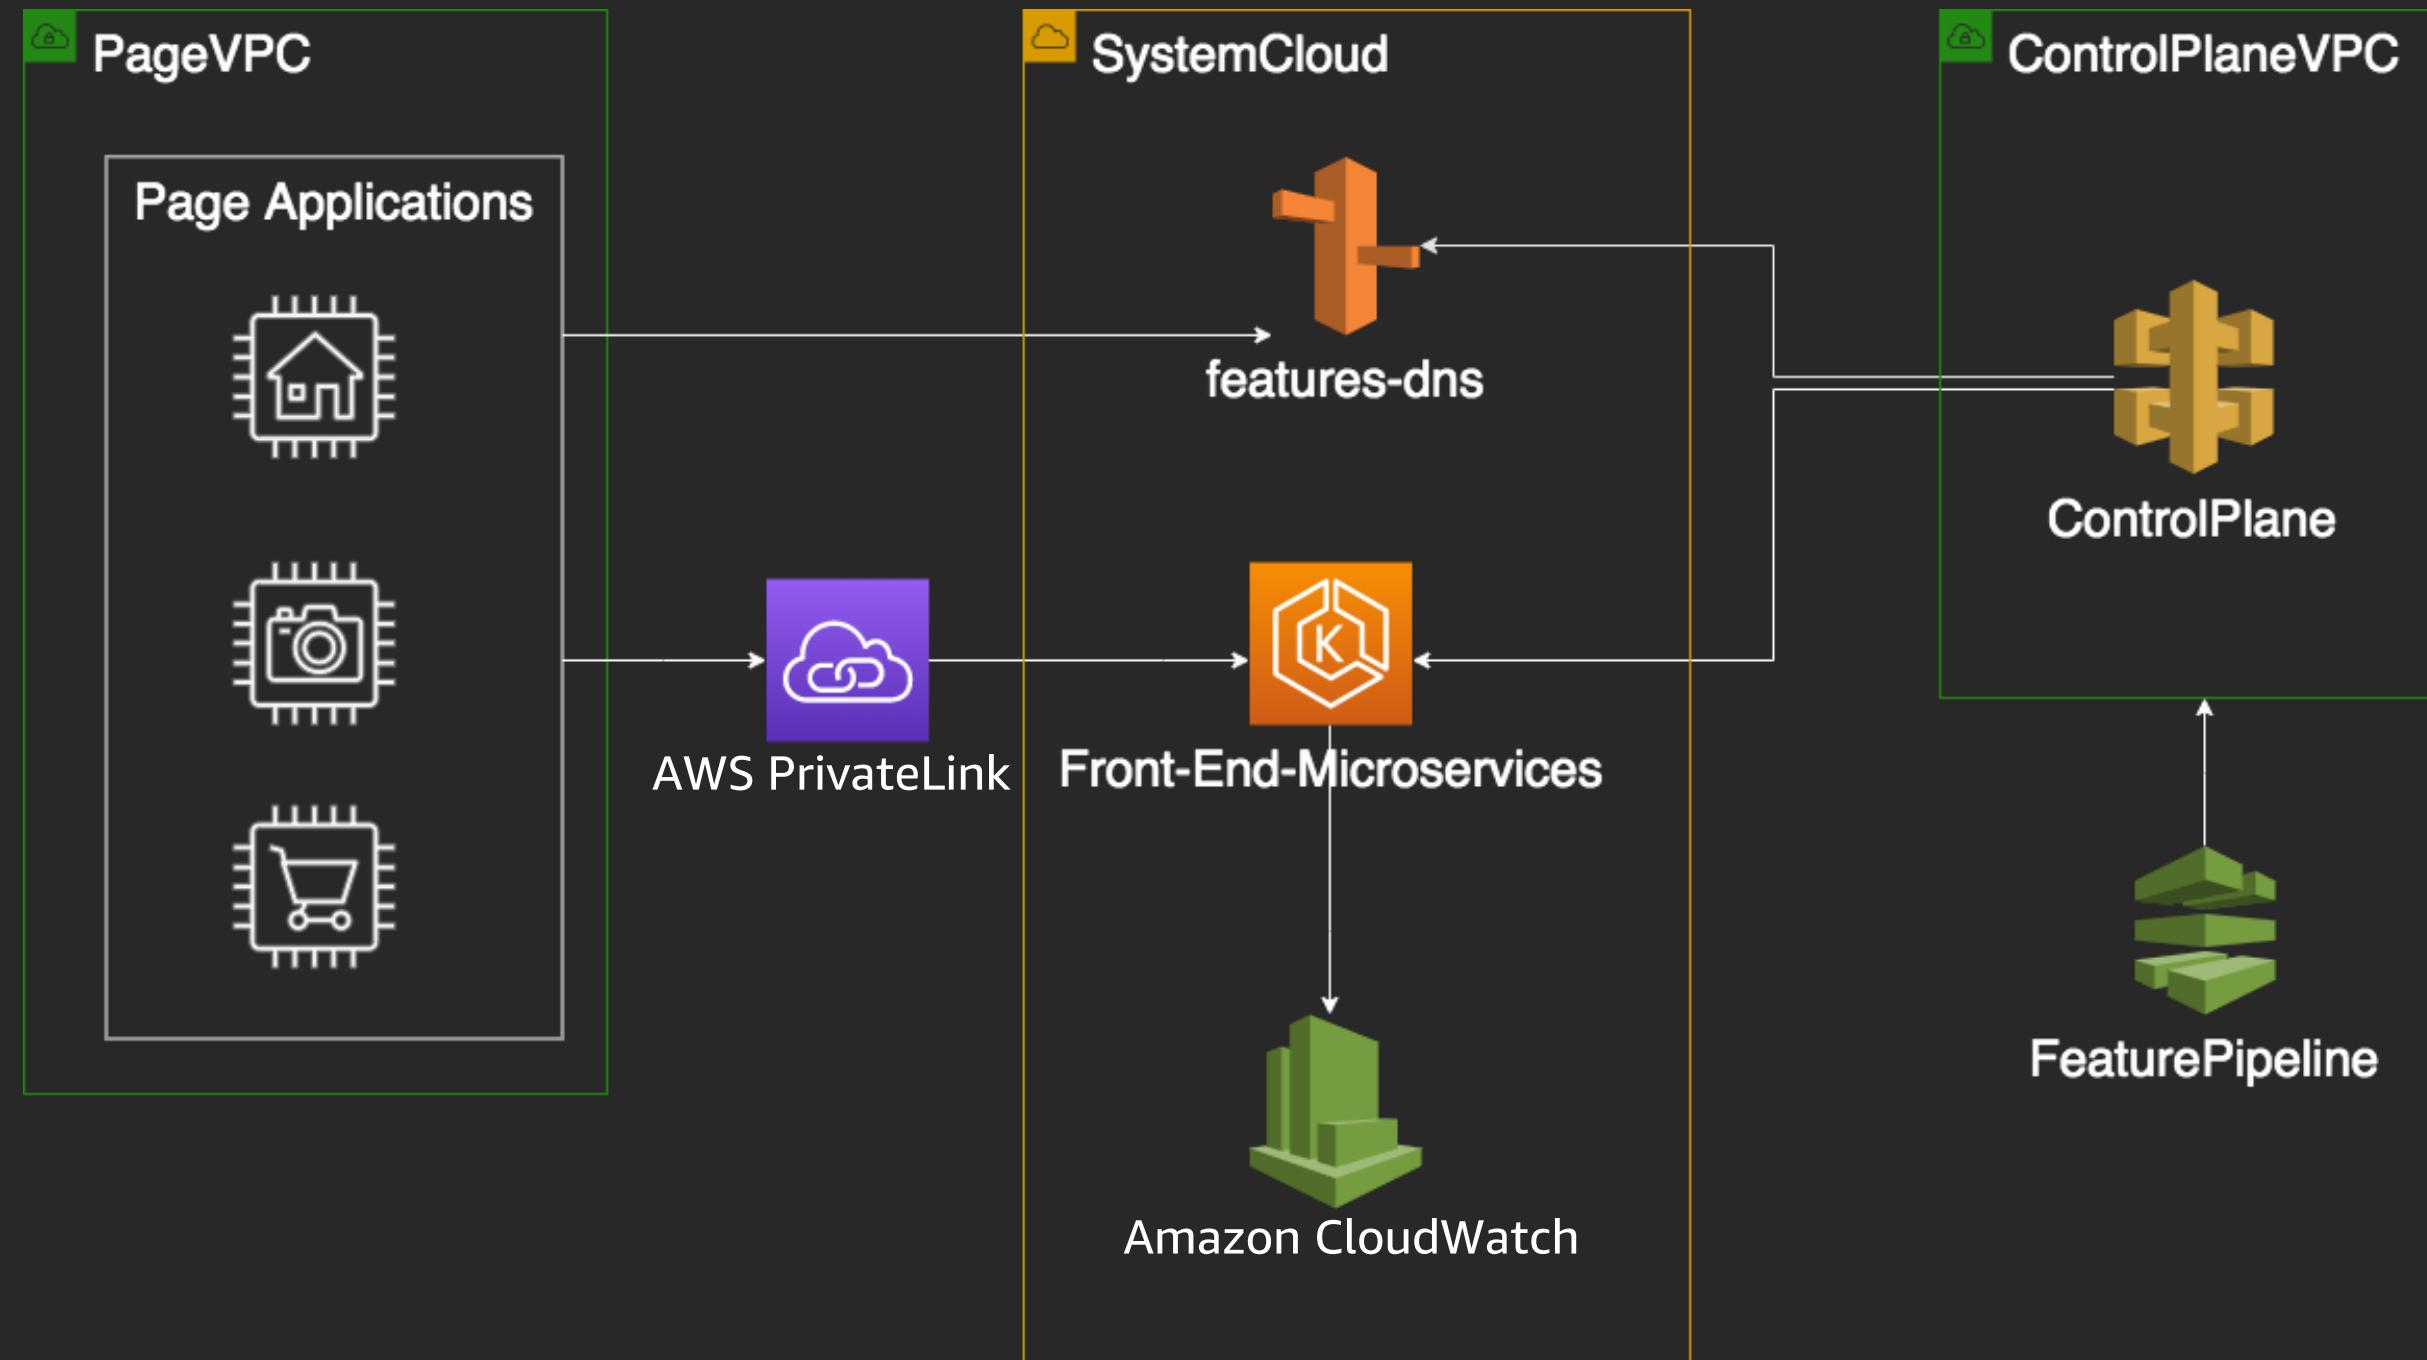
Gradual deployments



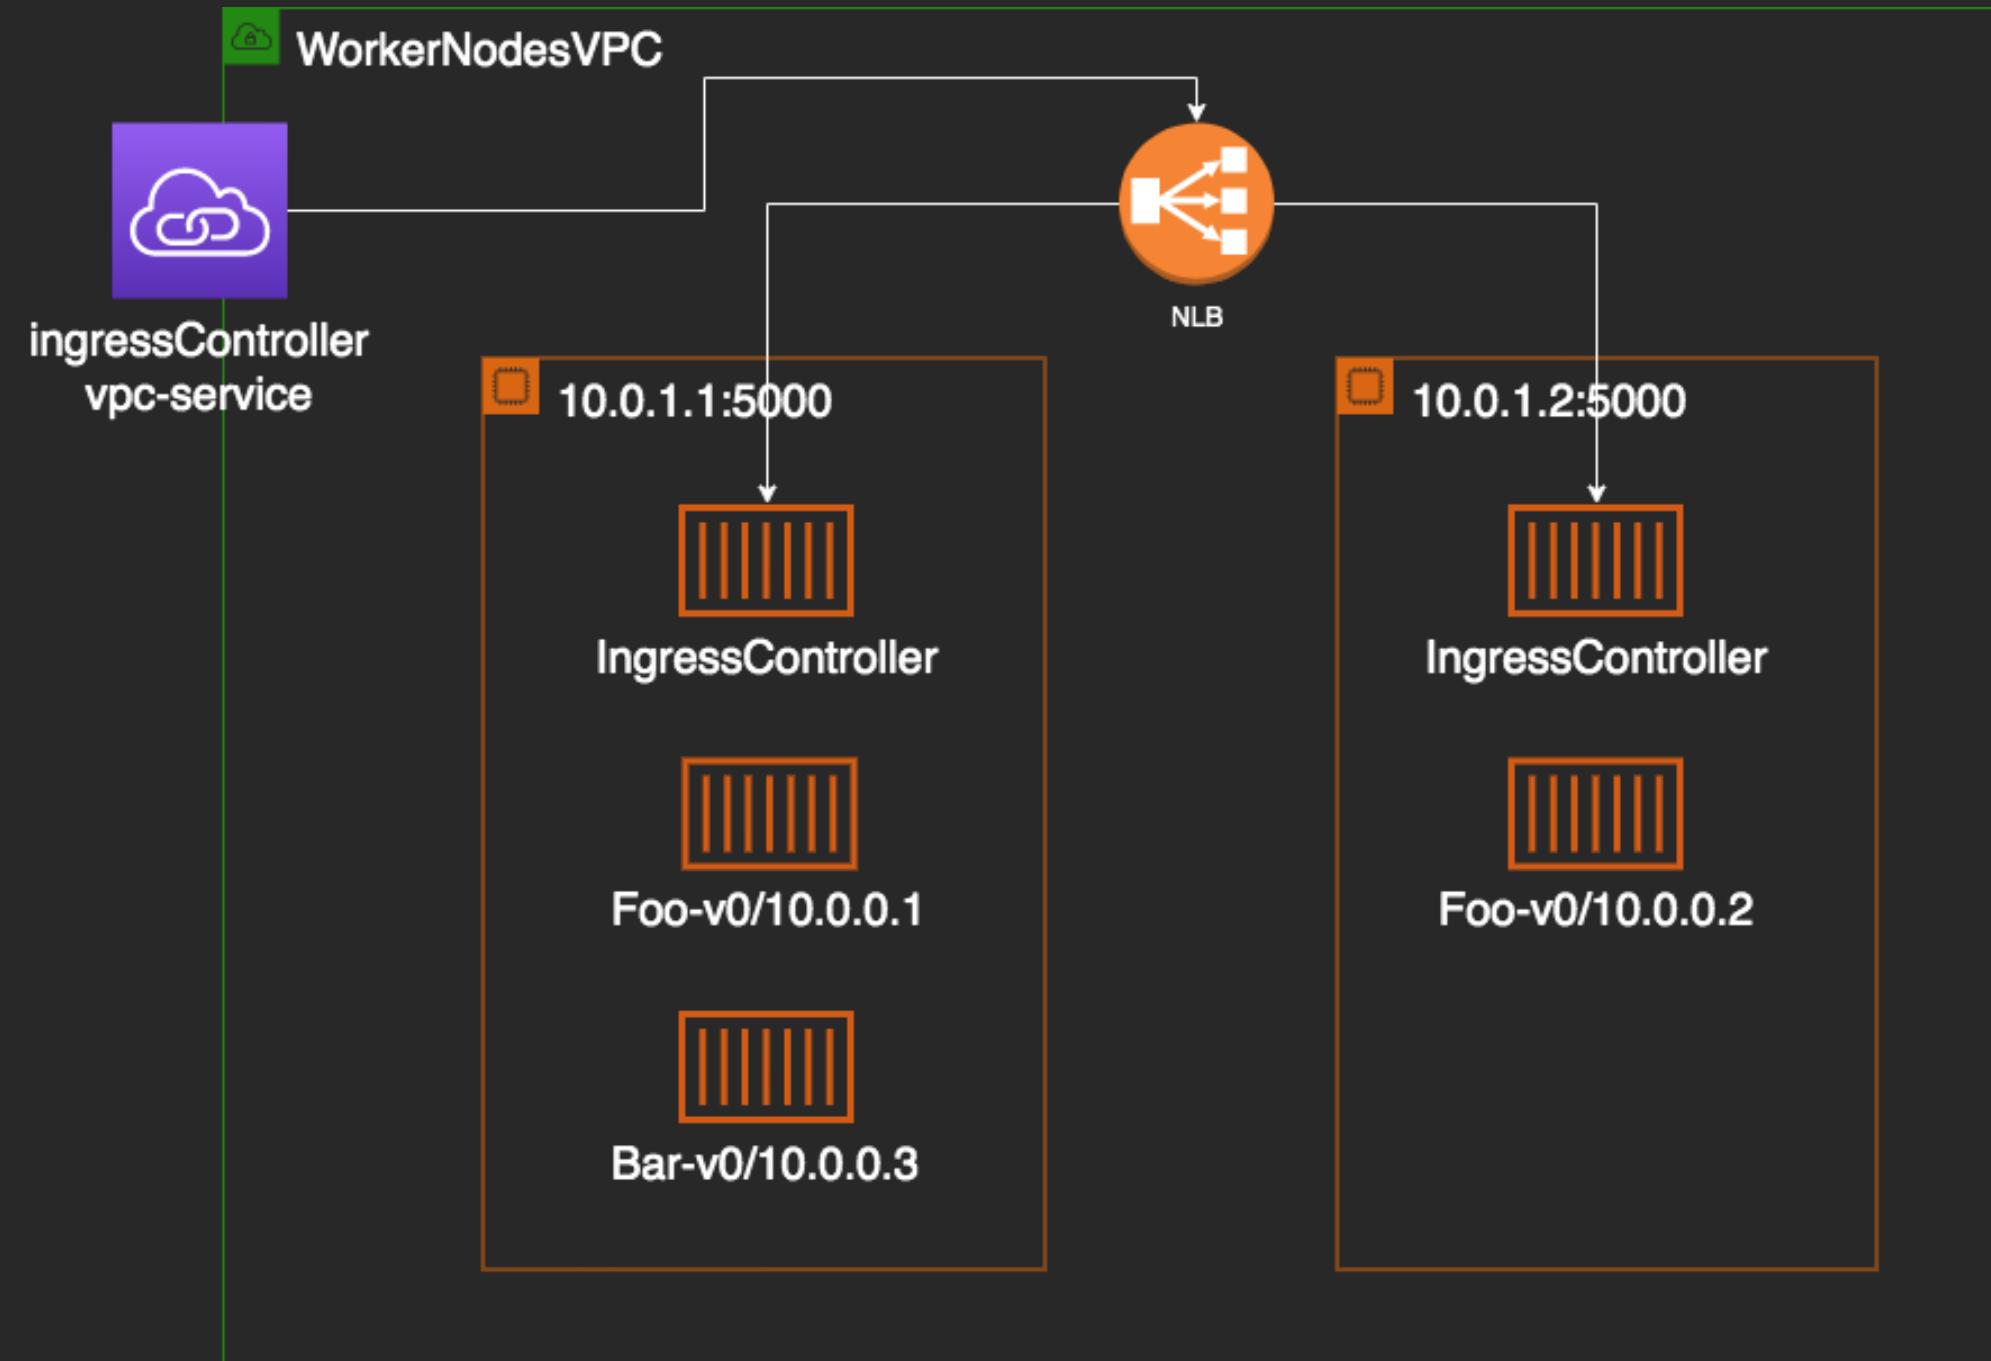
Highly available elasticity



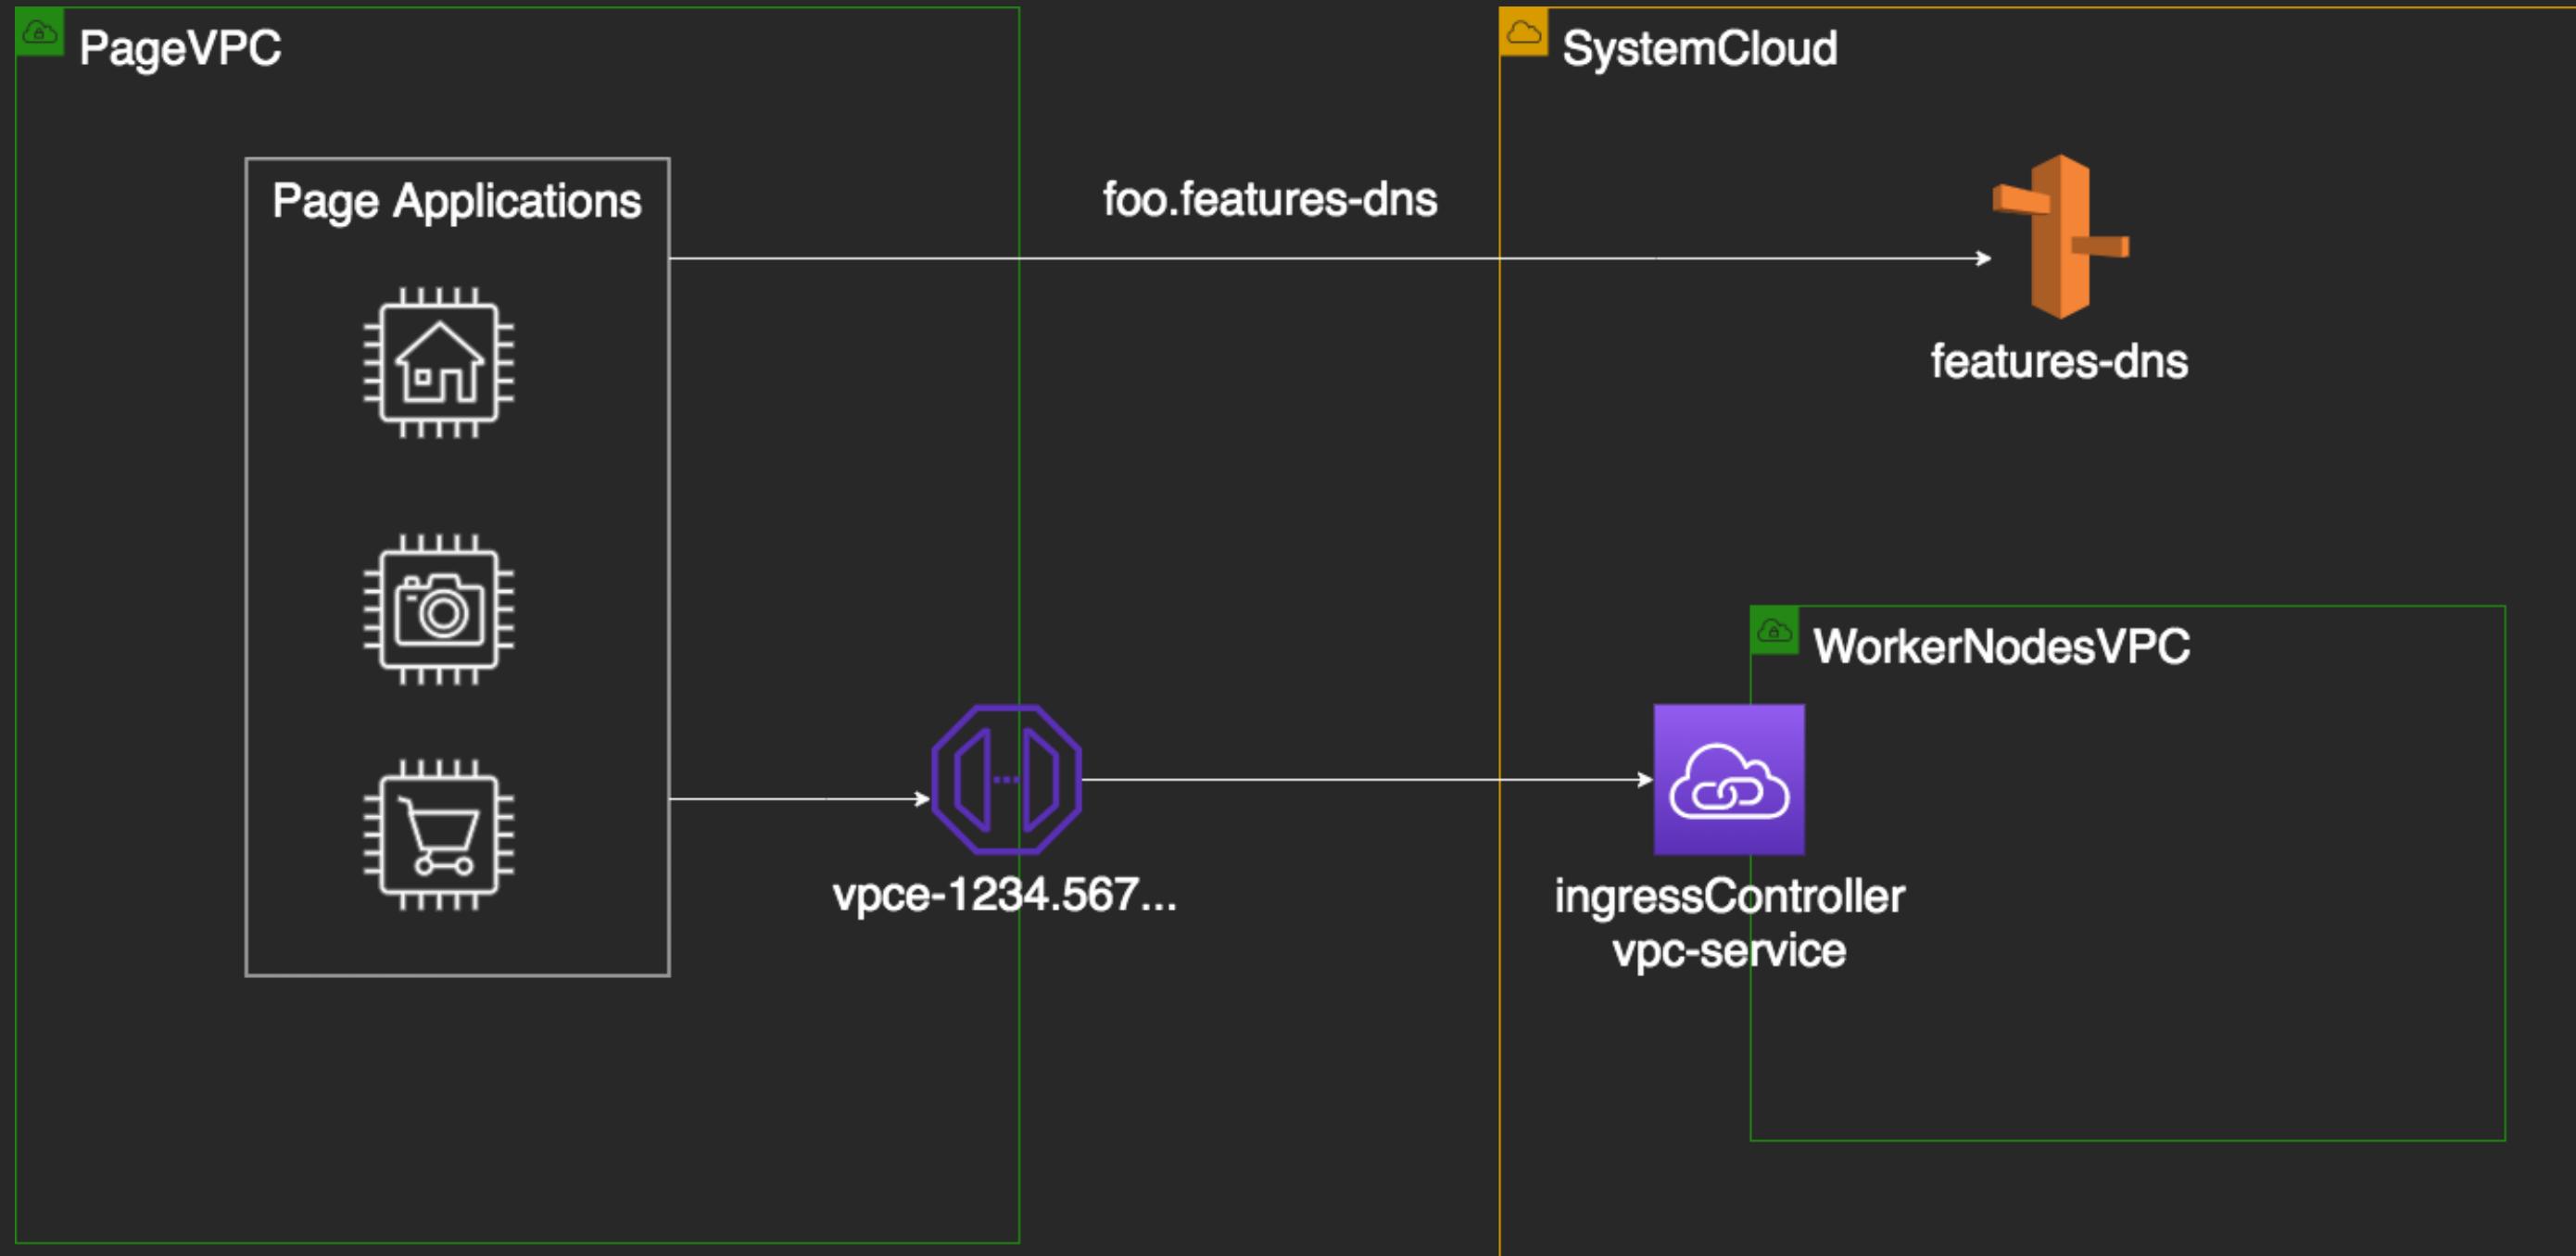
High-level architecture



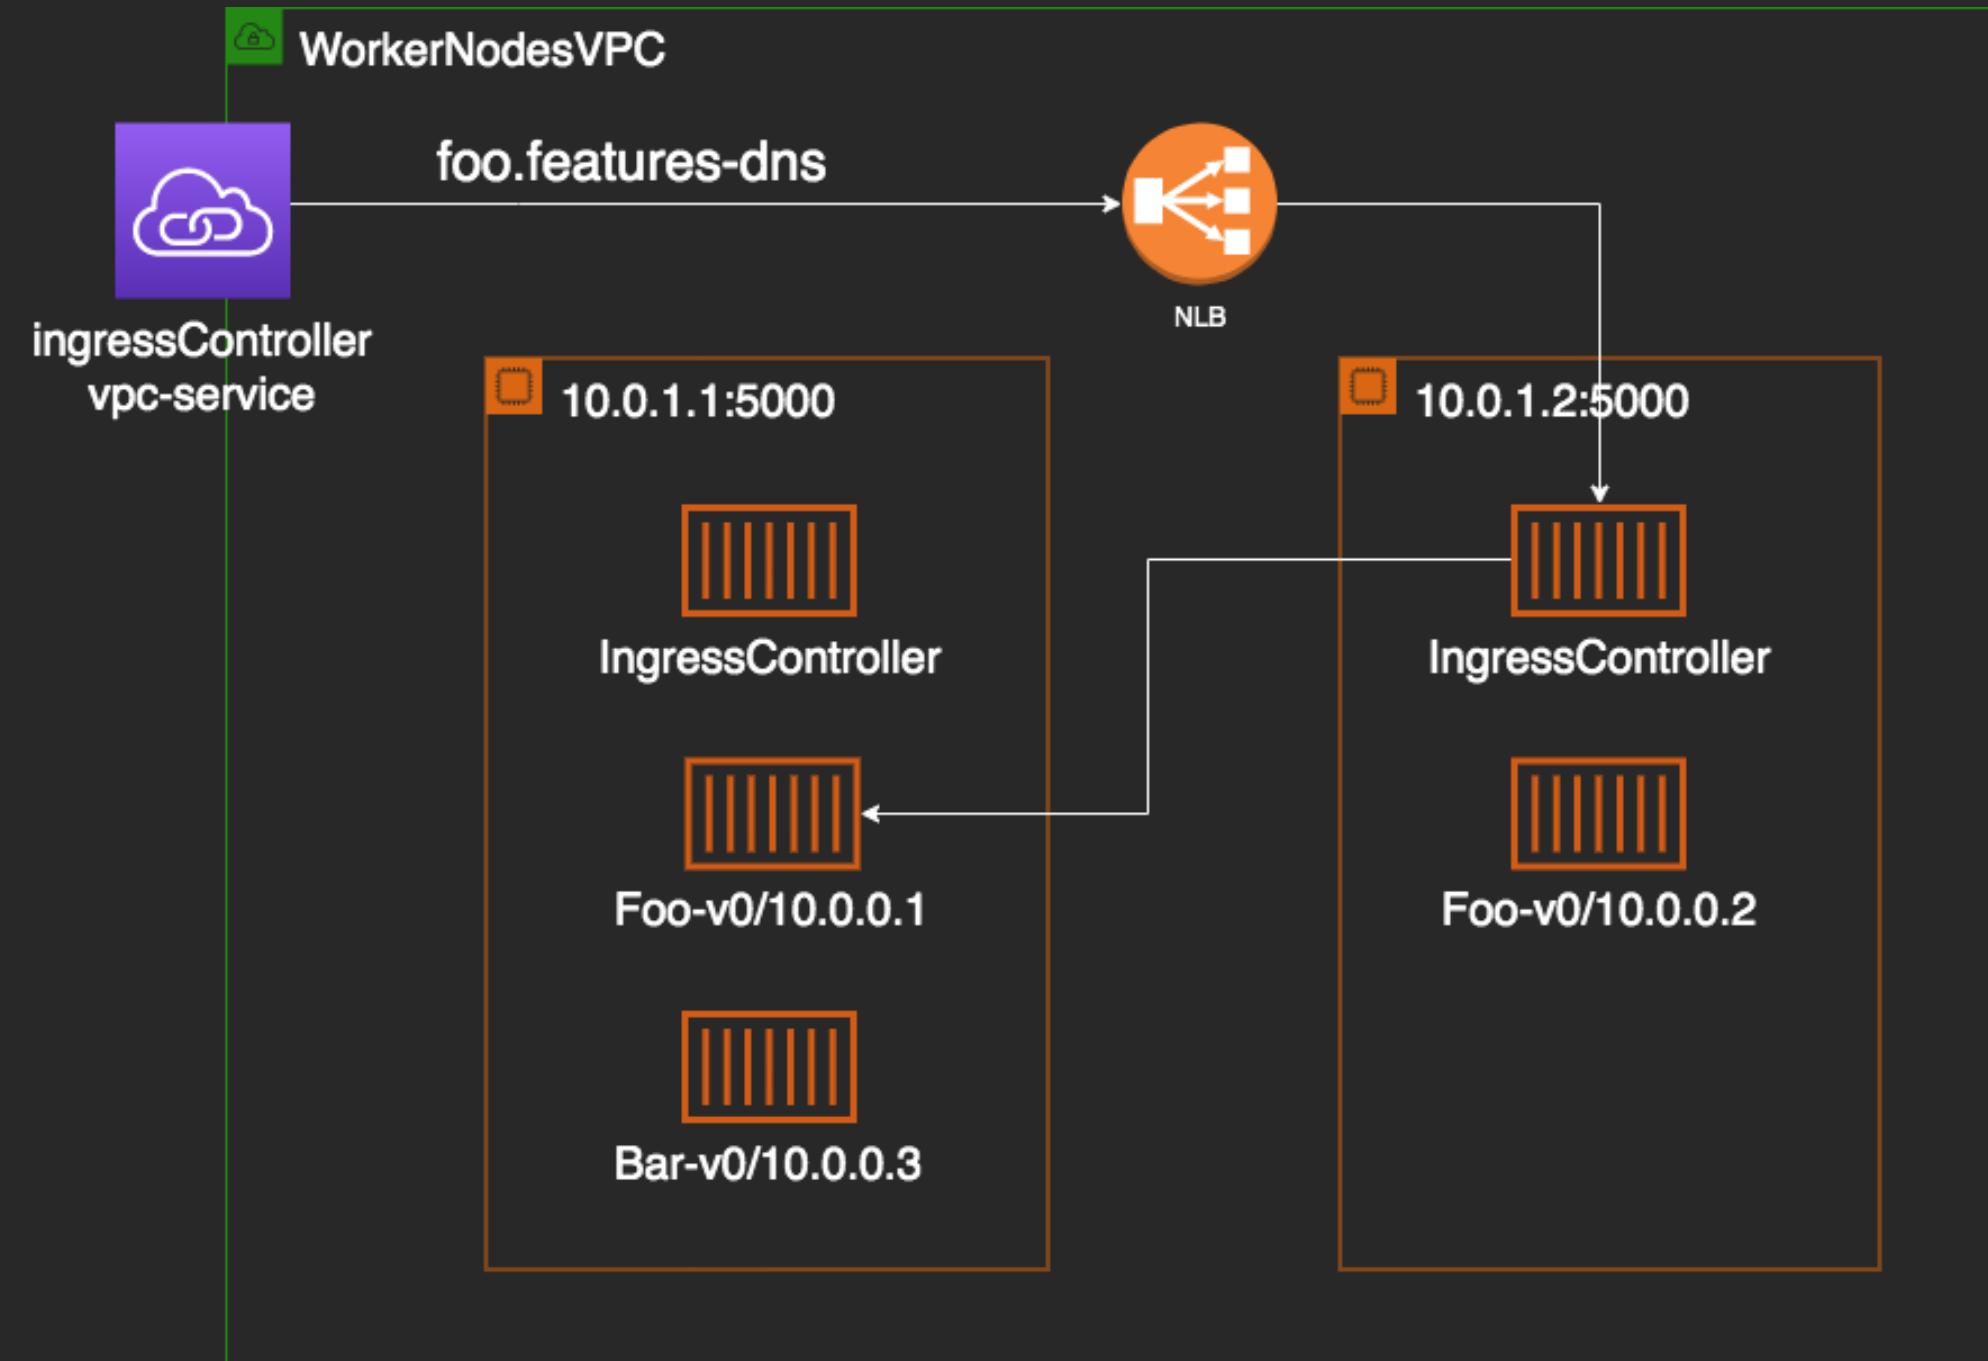
Running Foo/Bar in Amazon EKS



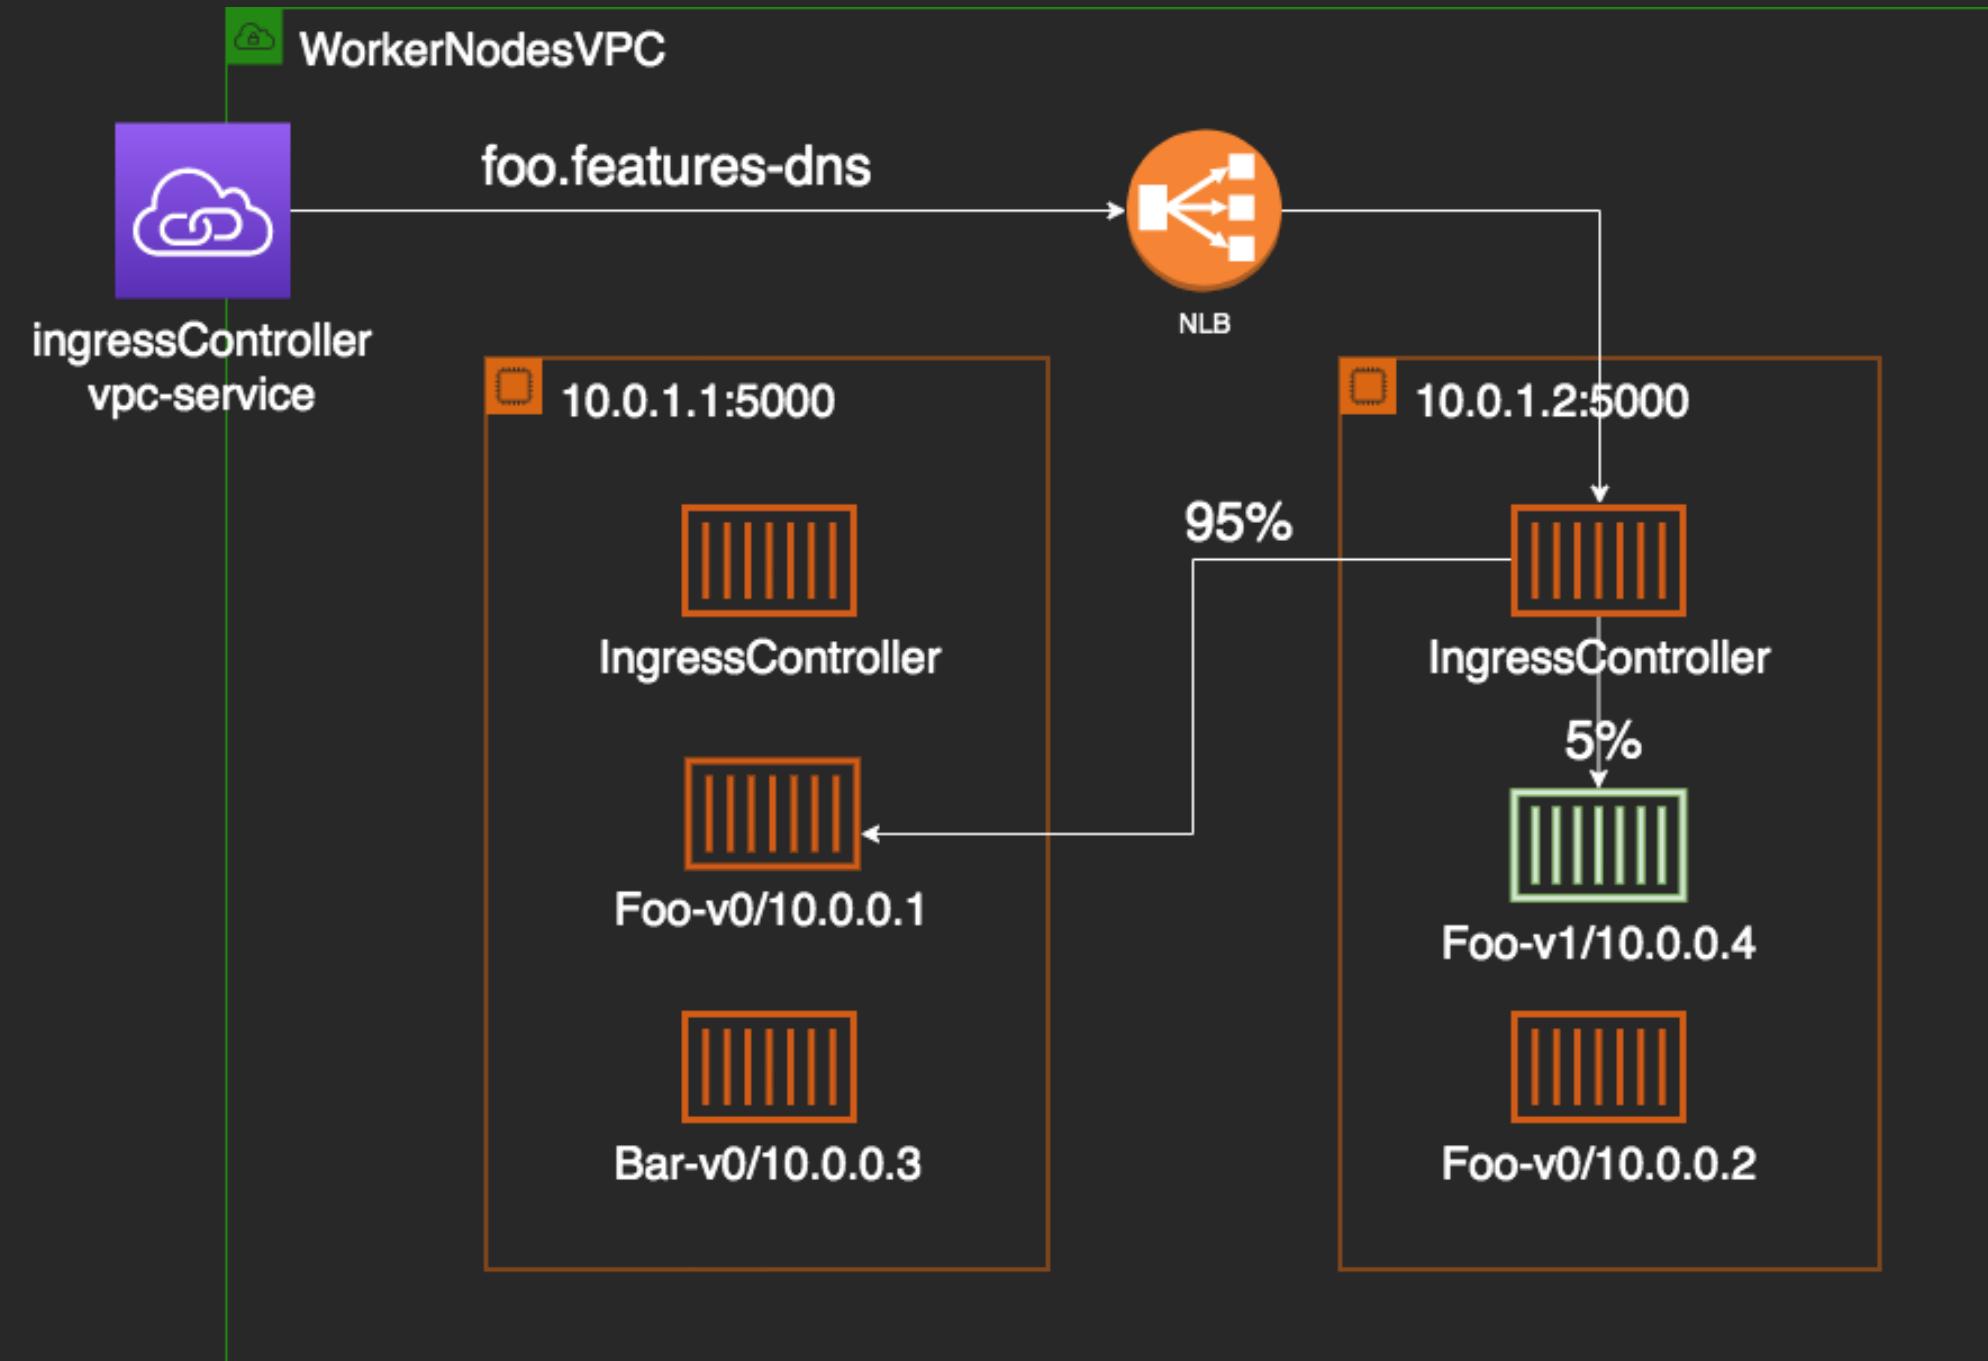
Routing requests to Foo



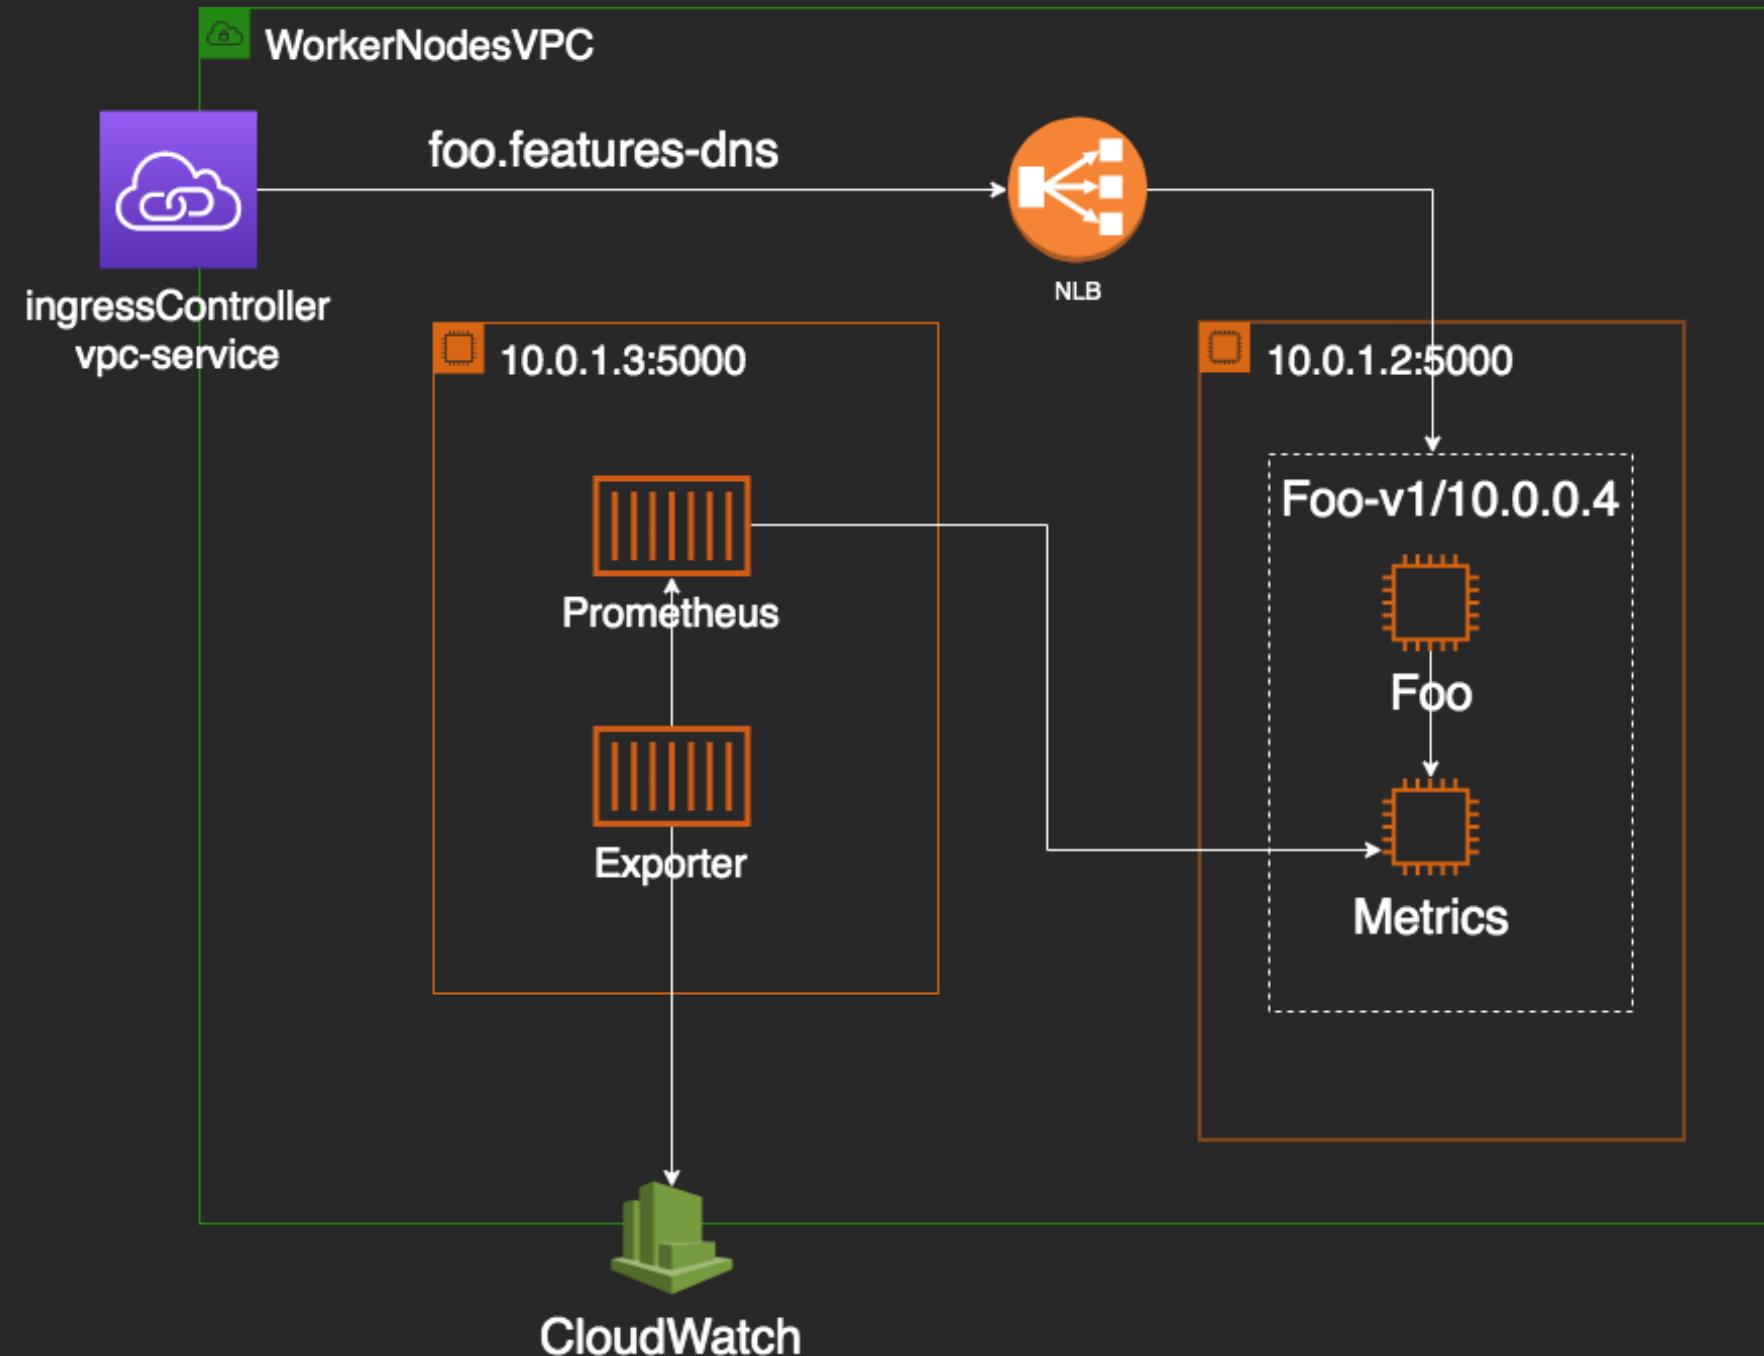
Routing requests to Foo



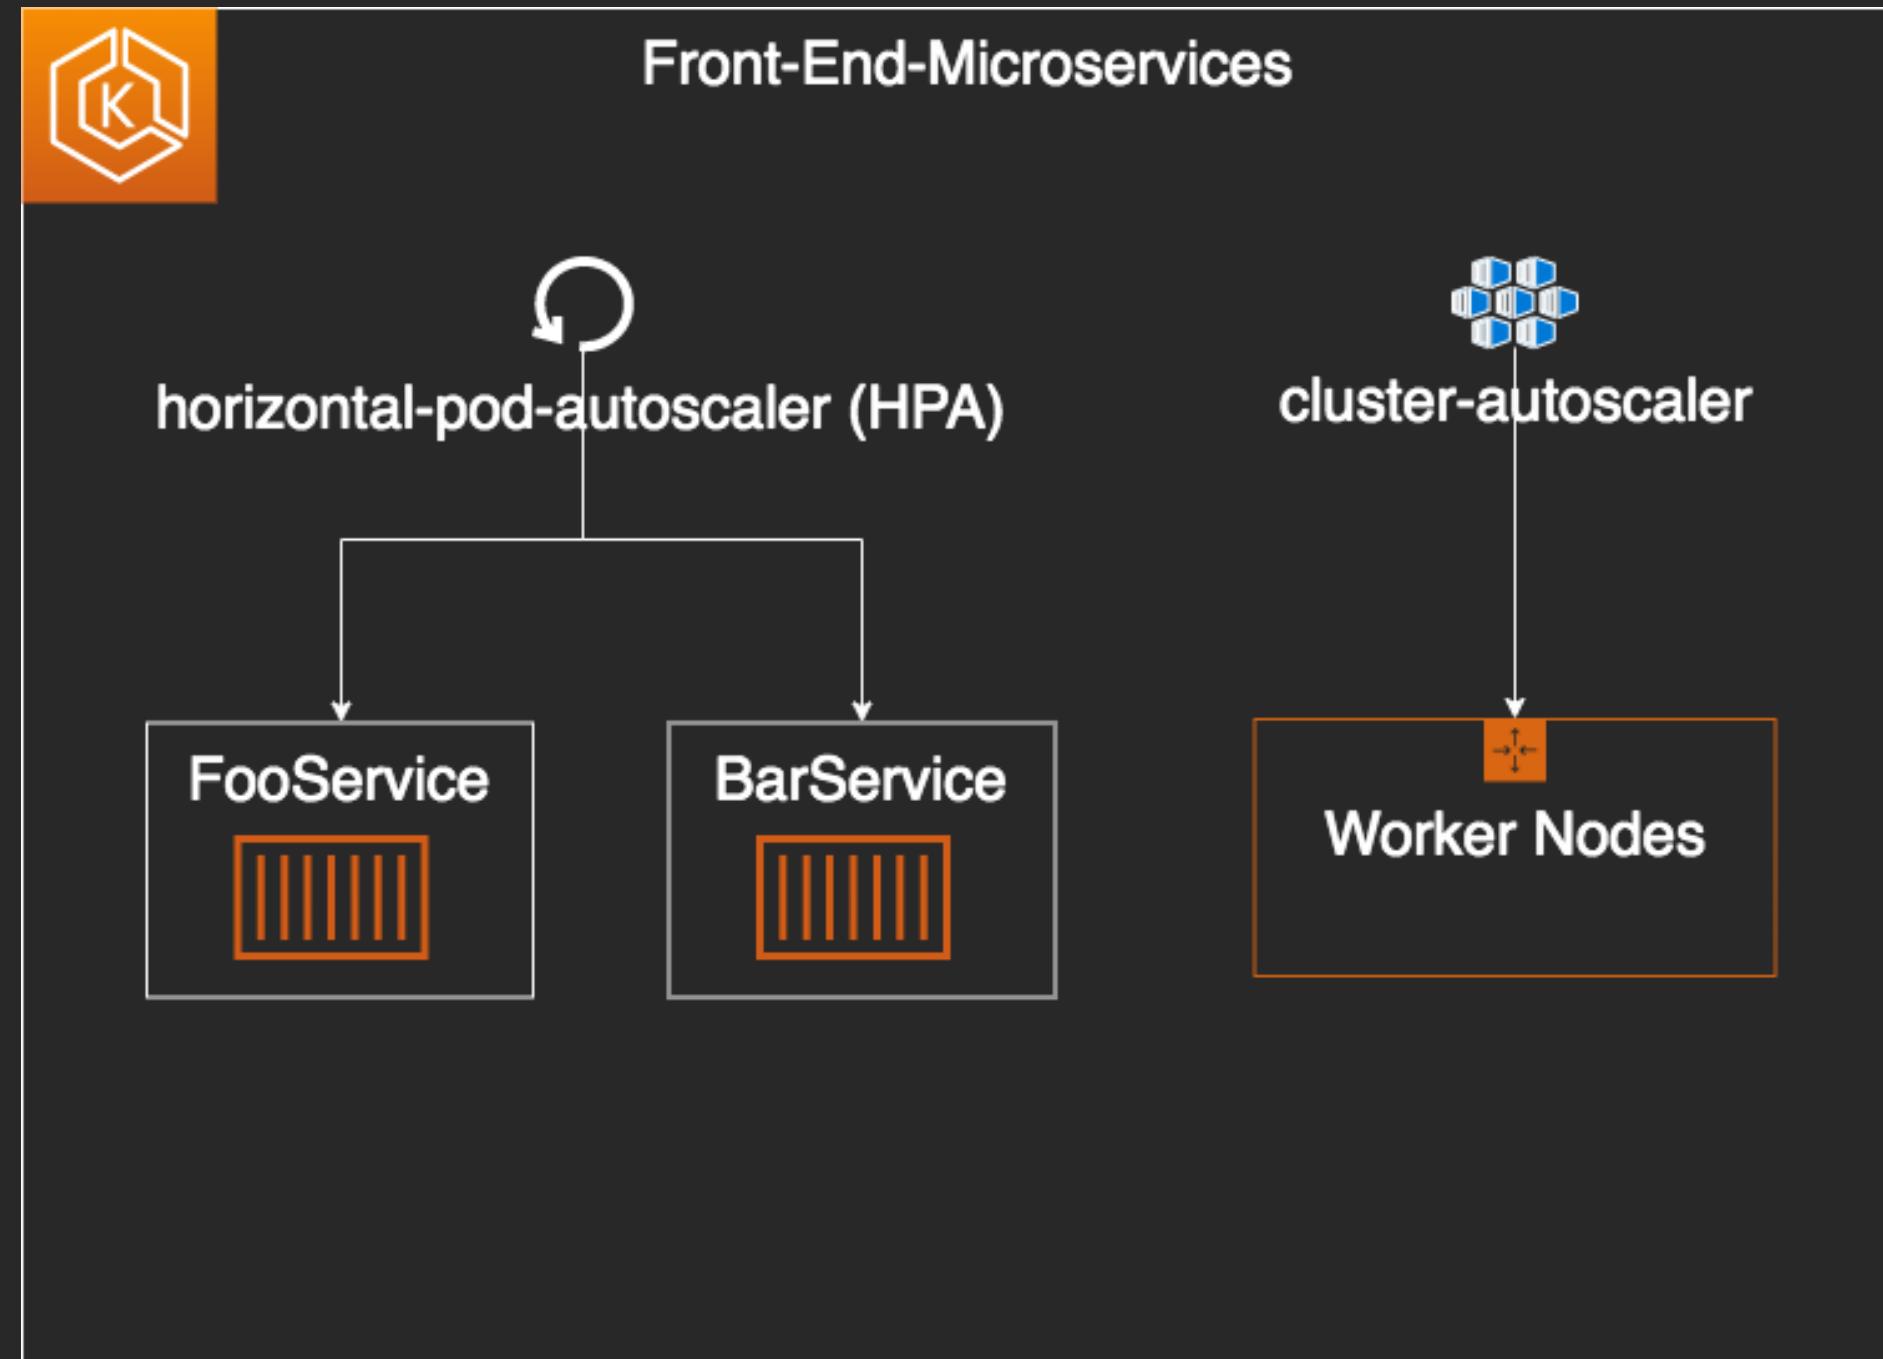
Gradually deploy Foo



Gradually deploy Foo



Highly available elasticity



Baby steps

- ✓ Running services
- ✓ Exposed endpoints
- ✓ Gradual deployments
- ❖ Elastic scaling



Running Foo without elasticity

- Success rate within SLAs
 - Small size payload (10kb)
 - Fast response (5ms)
- Drops on availability
 - Deployments
 - Scaling



How we enabled elasticity?

- Cold start time
- Success rate
- 60% CPU



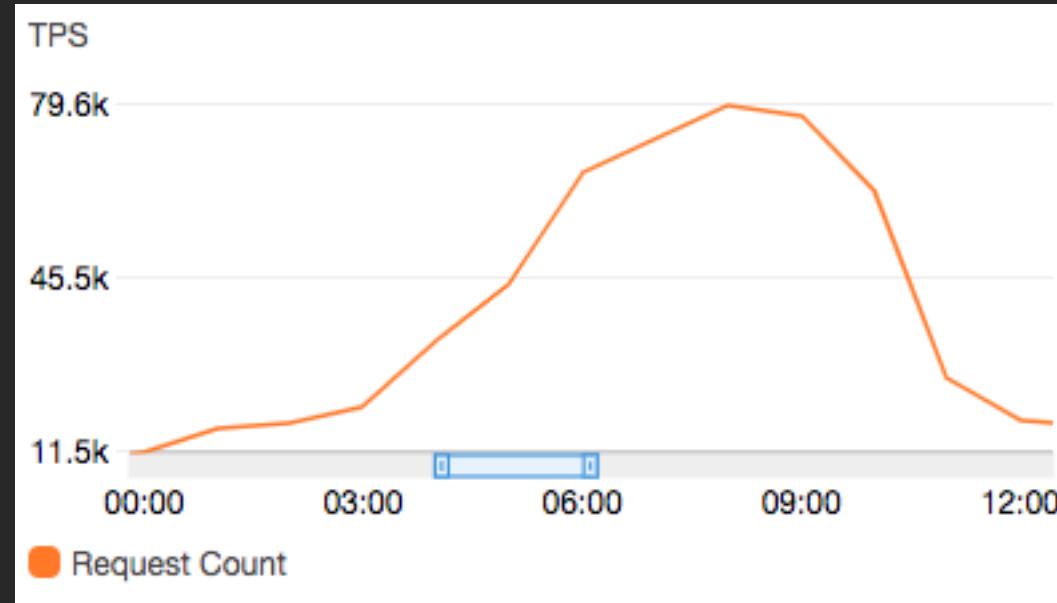
We did it!



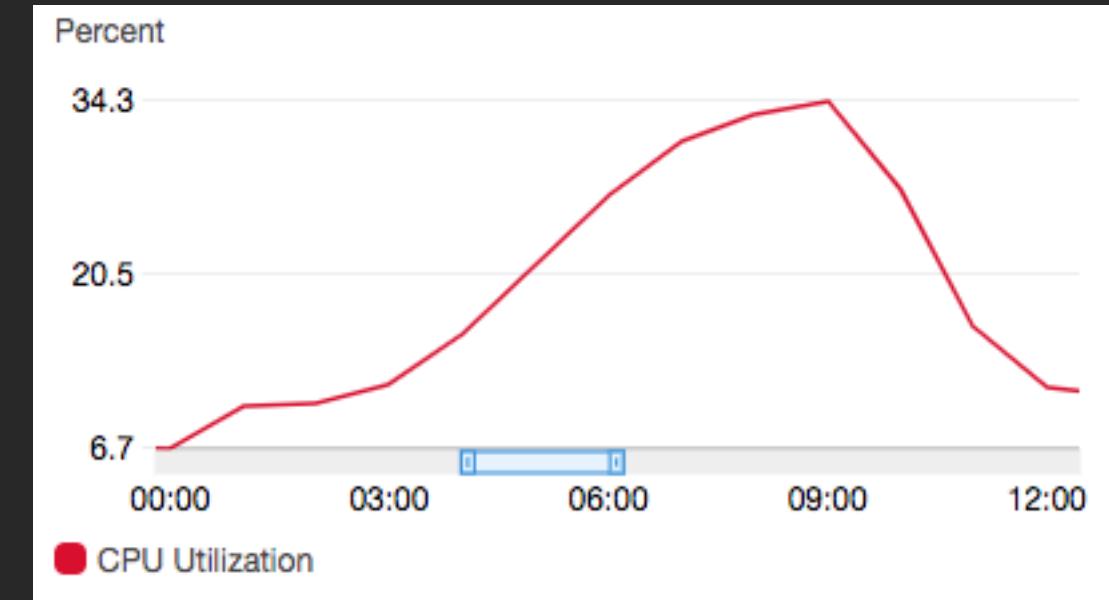
How we learned to stop worrying and ❤️ Amazon scale

Highly available elasticity

Request count

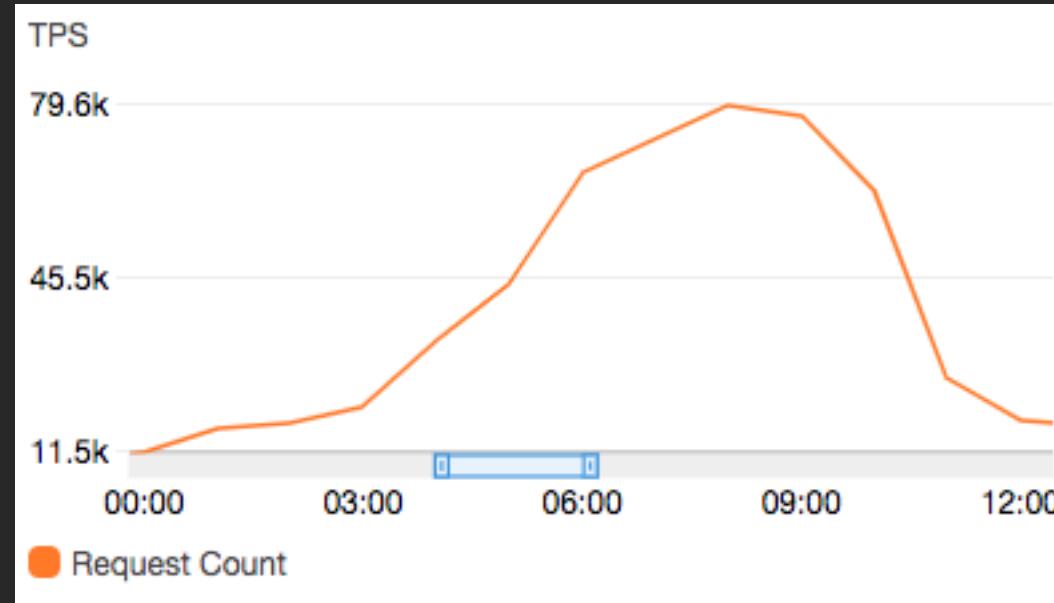


CPU utilization

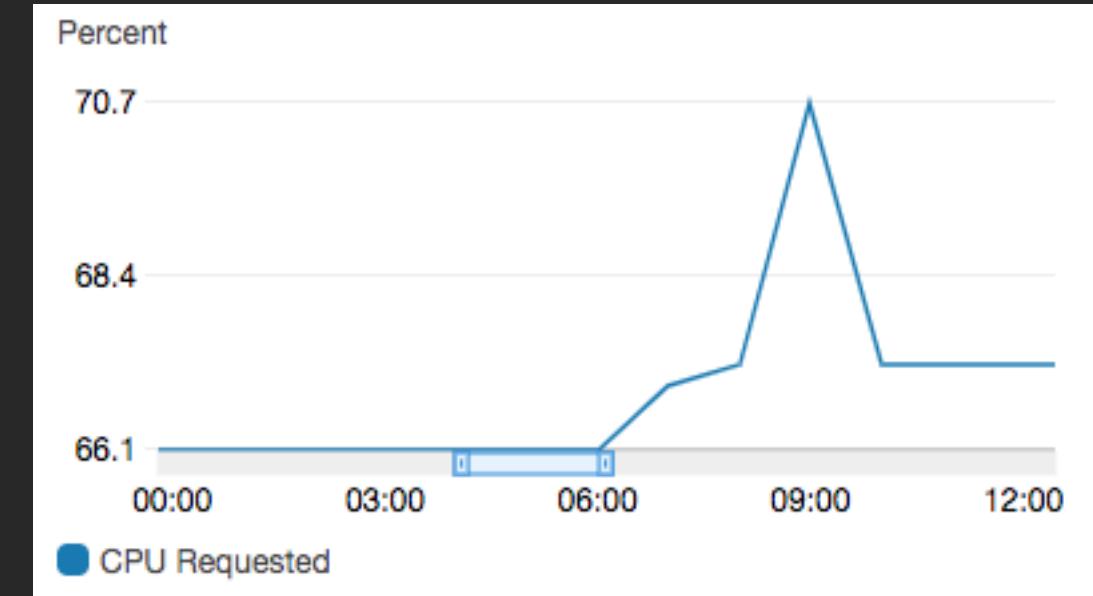


Highly available elasticity

Request count

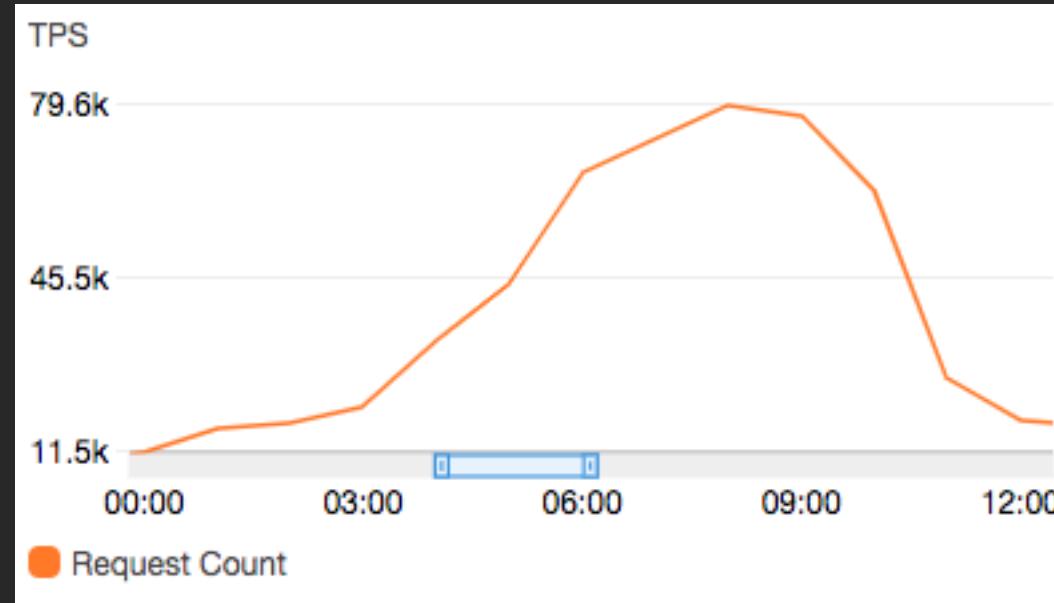


CPU requested

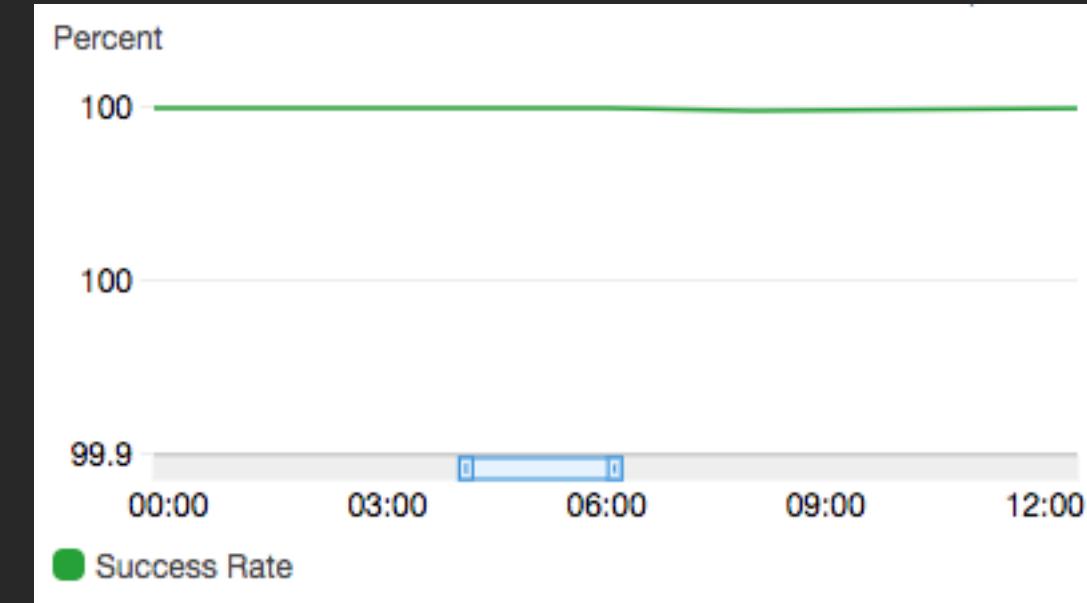


Highly available elasticity

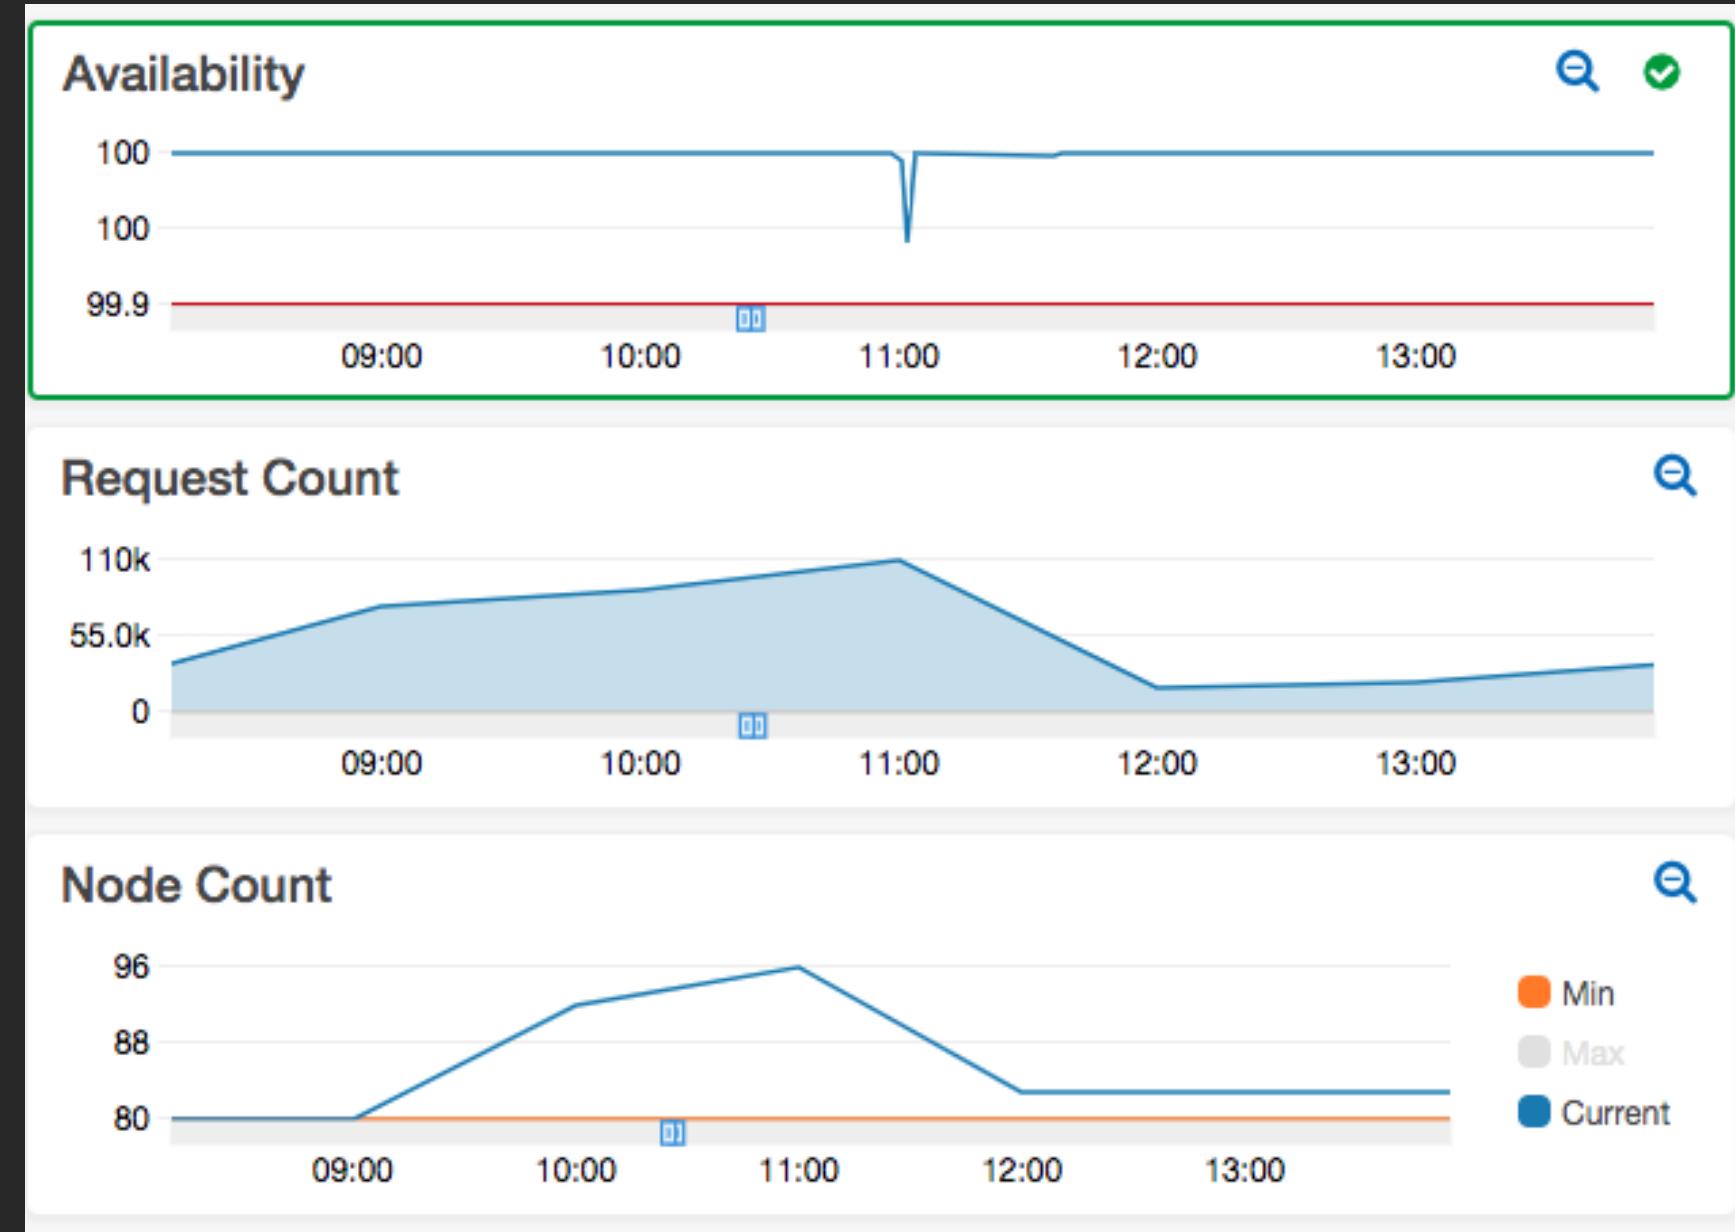
Request count



Success rate



Highly available elasticity



And we updated our cluster

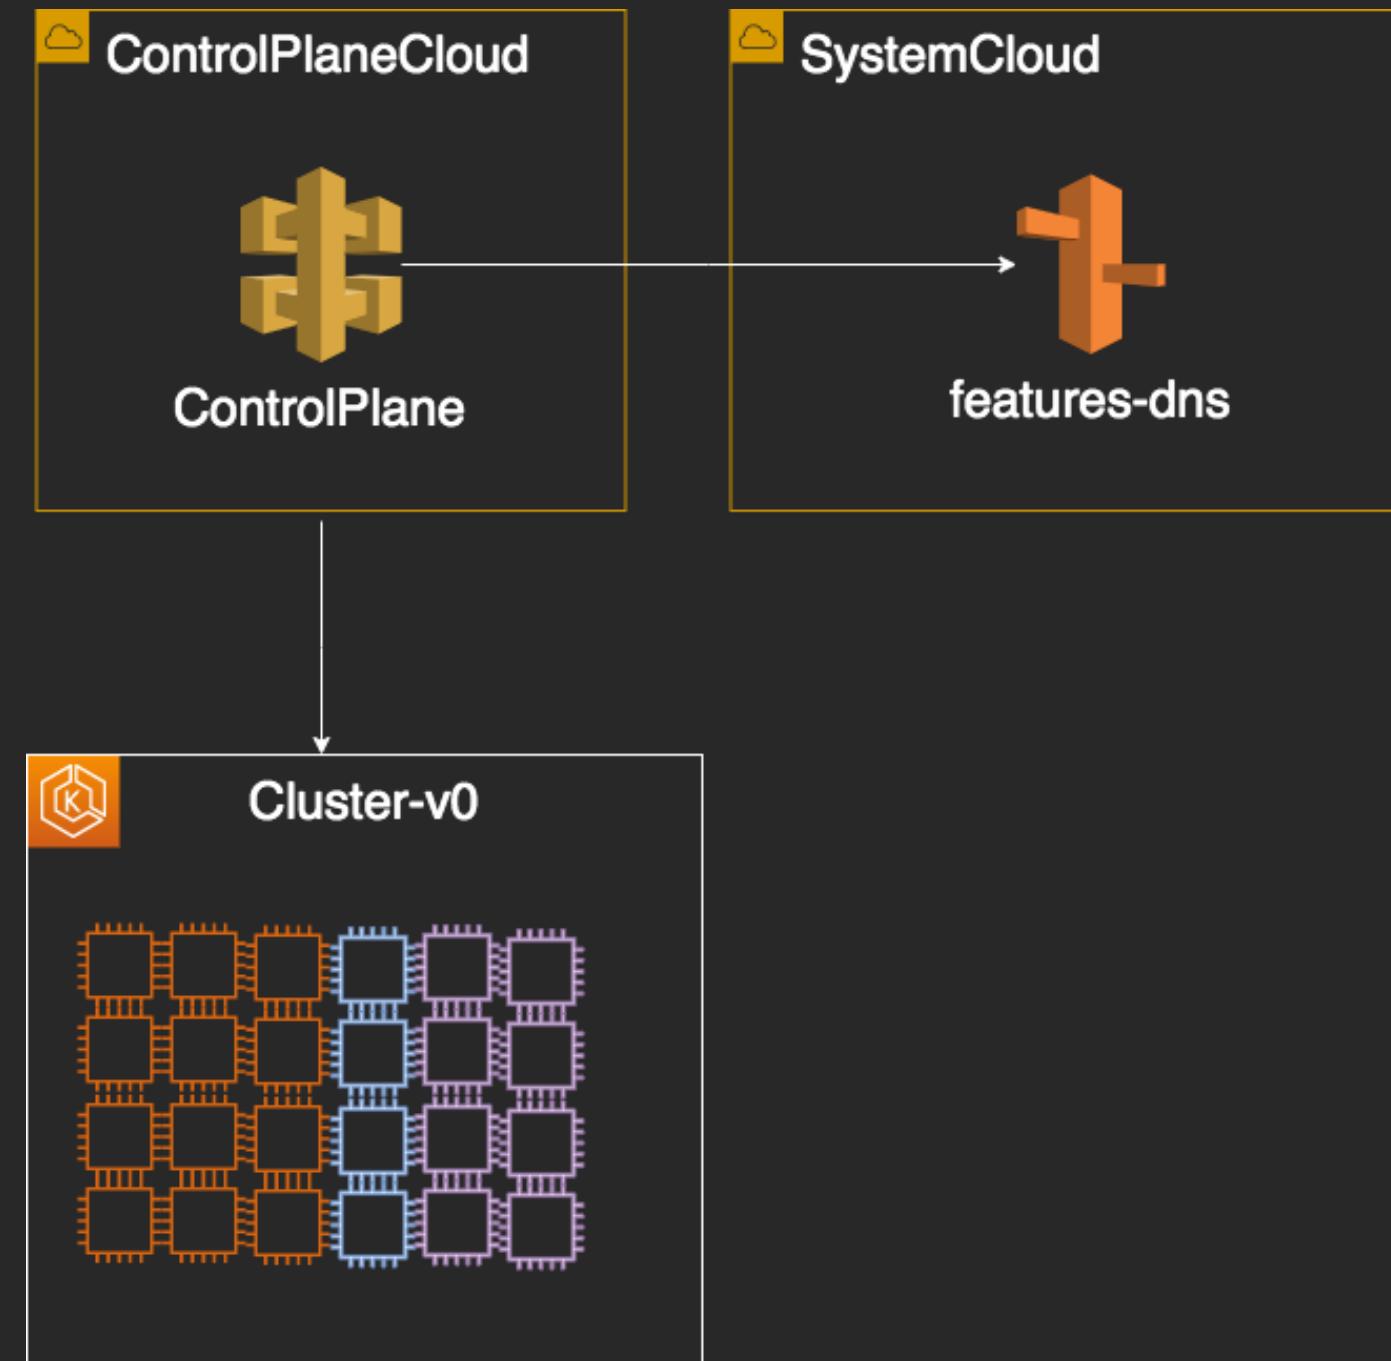


Ephemeral clusters

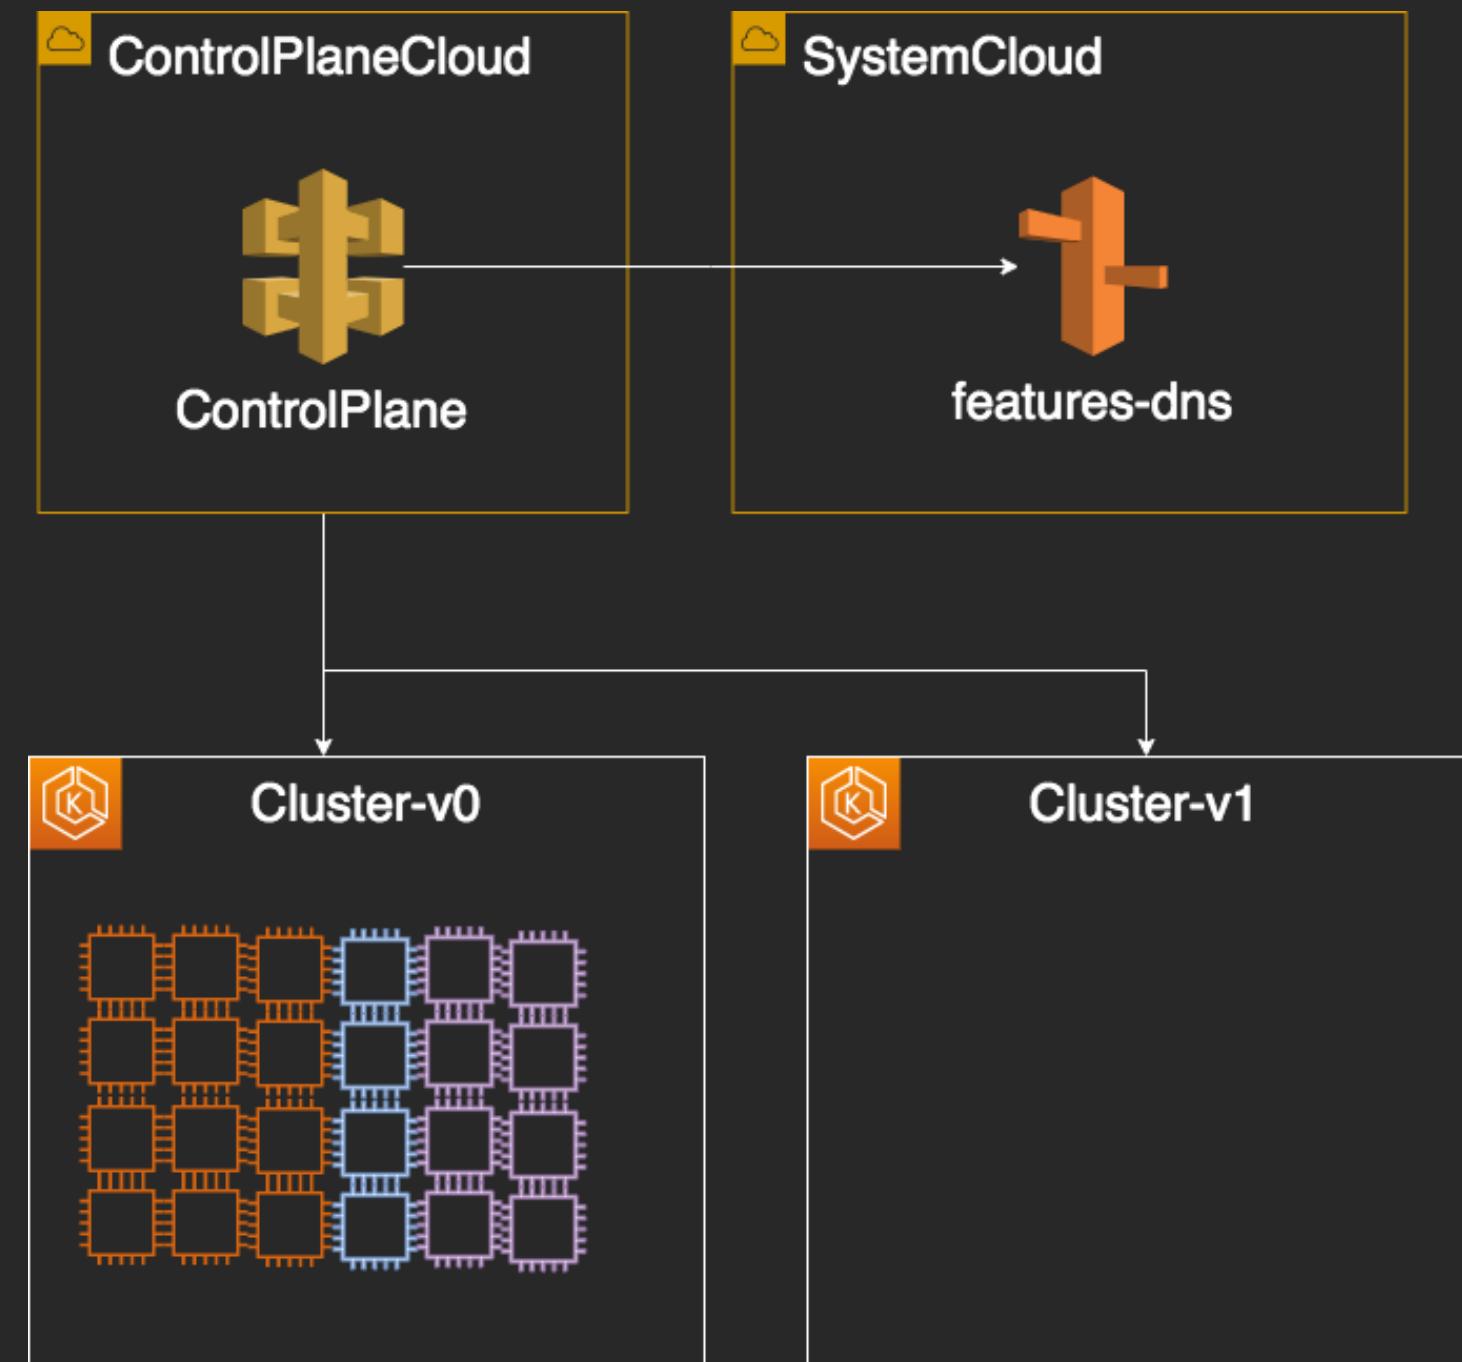
- Infrastructure automation
- Cluster rollouts



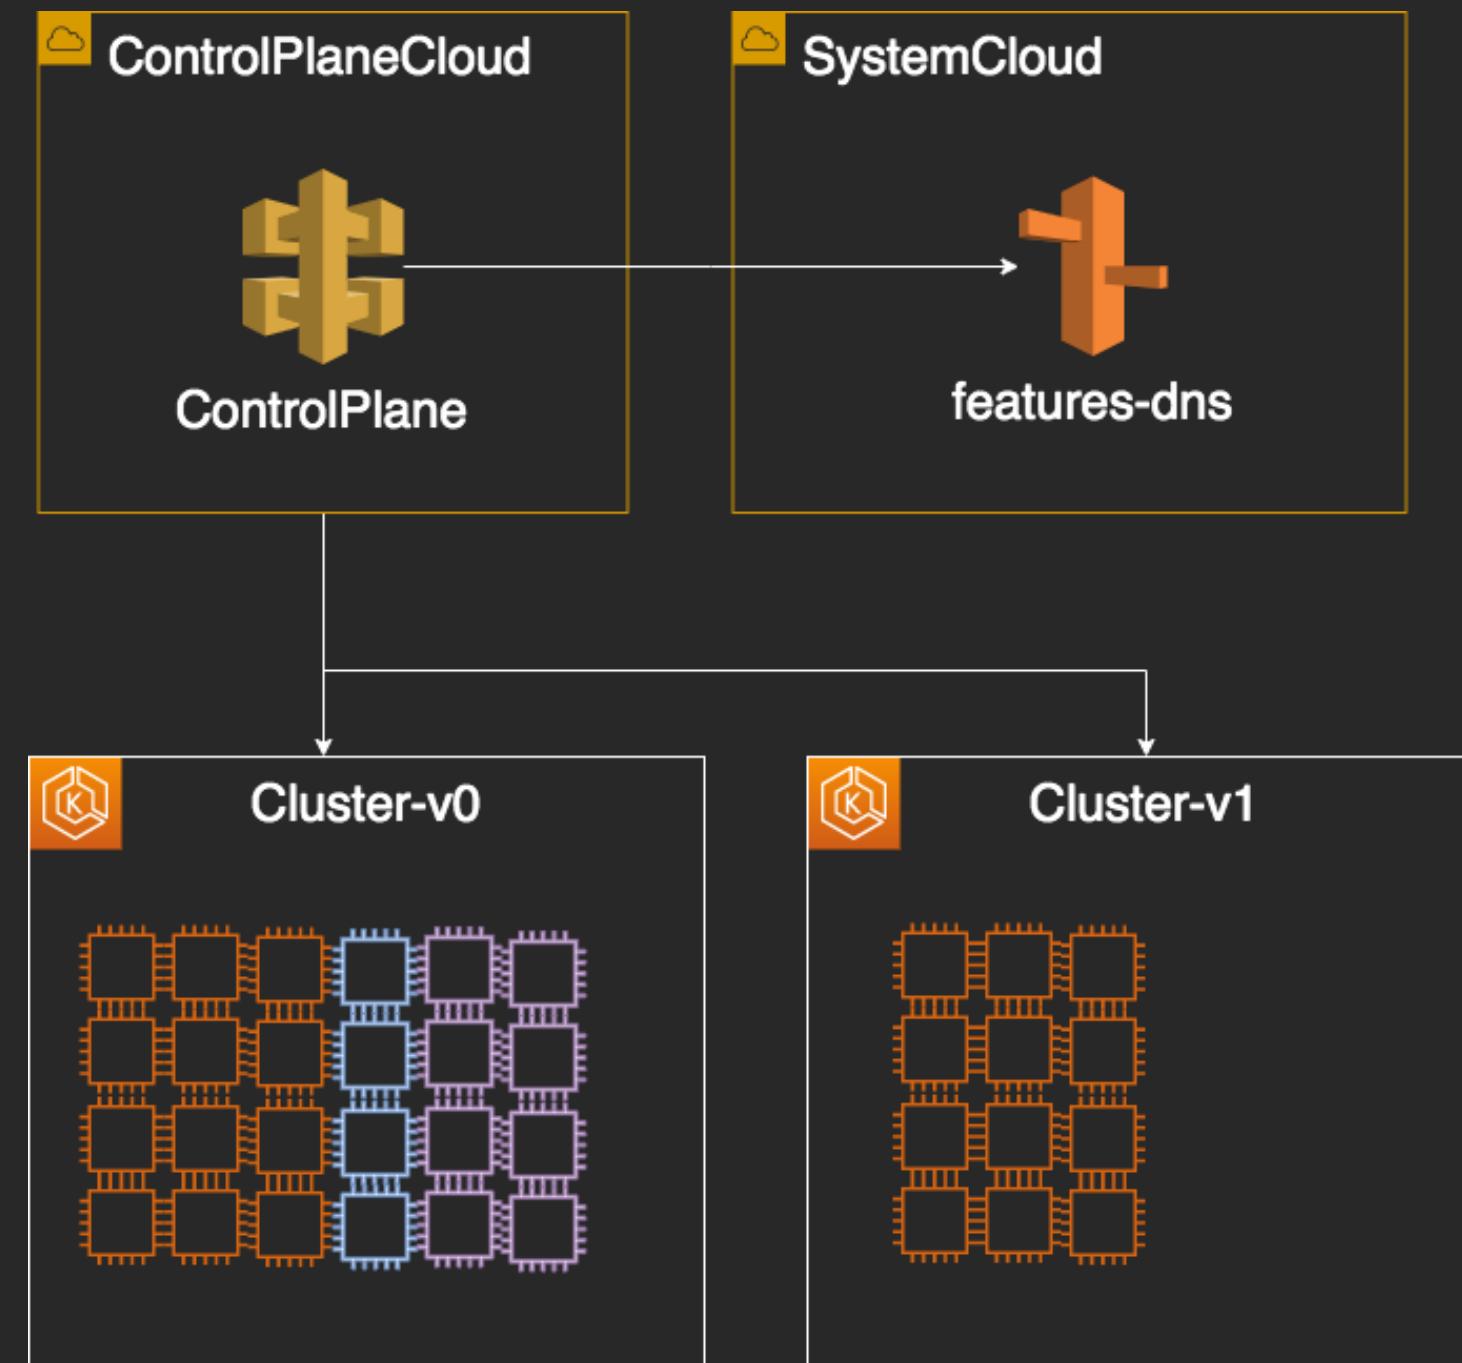
Cluster rollouts



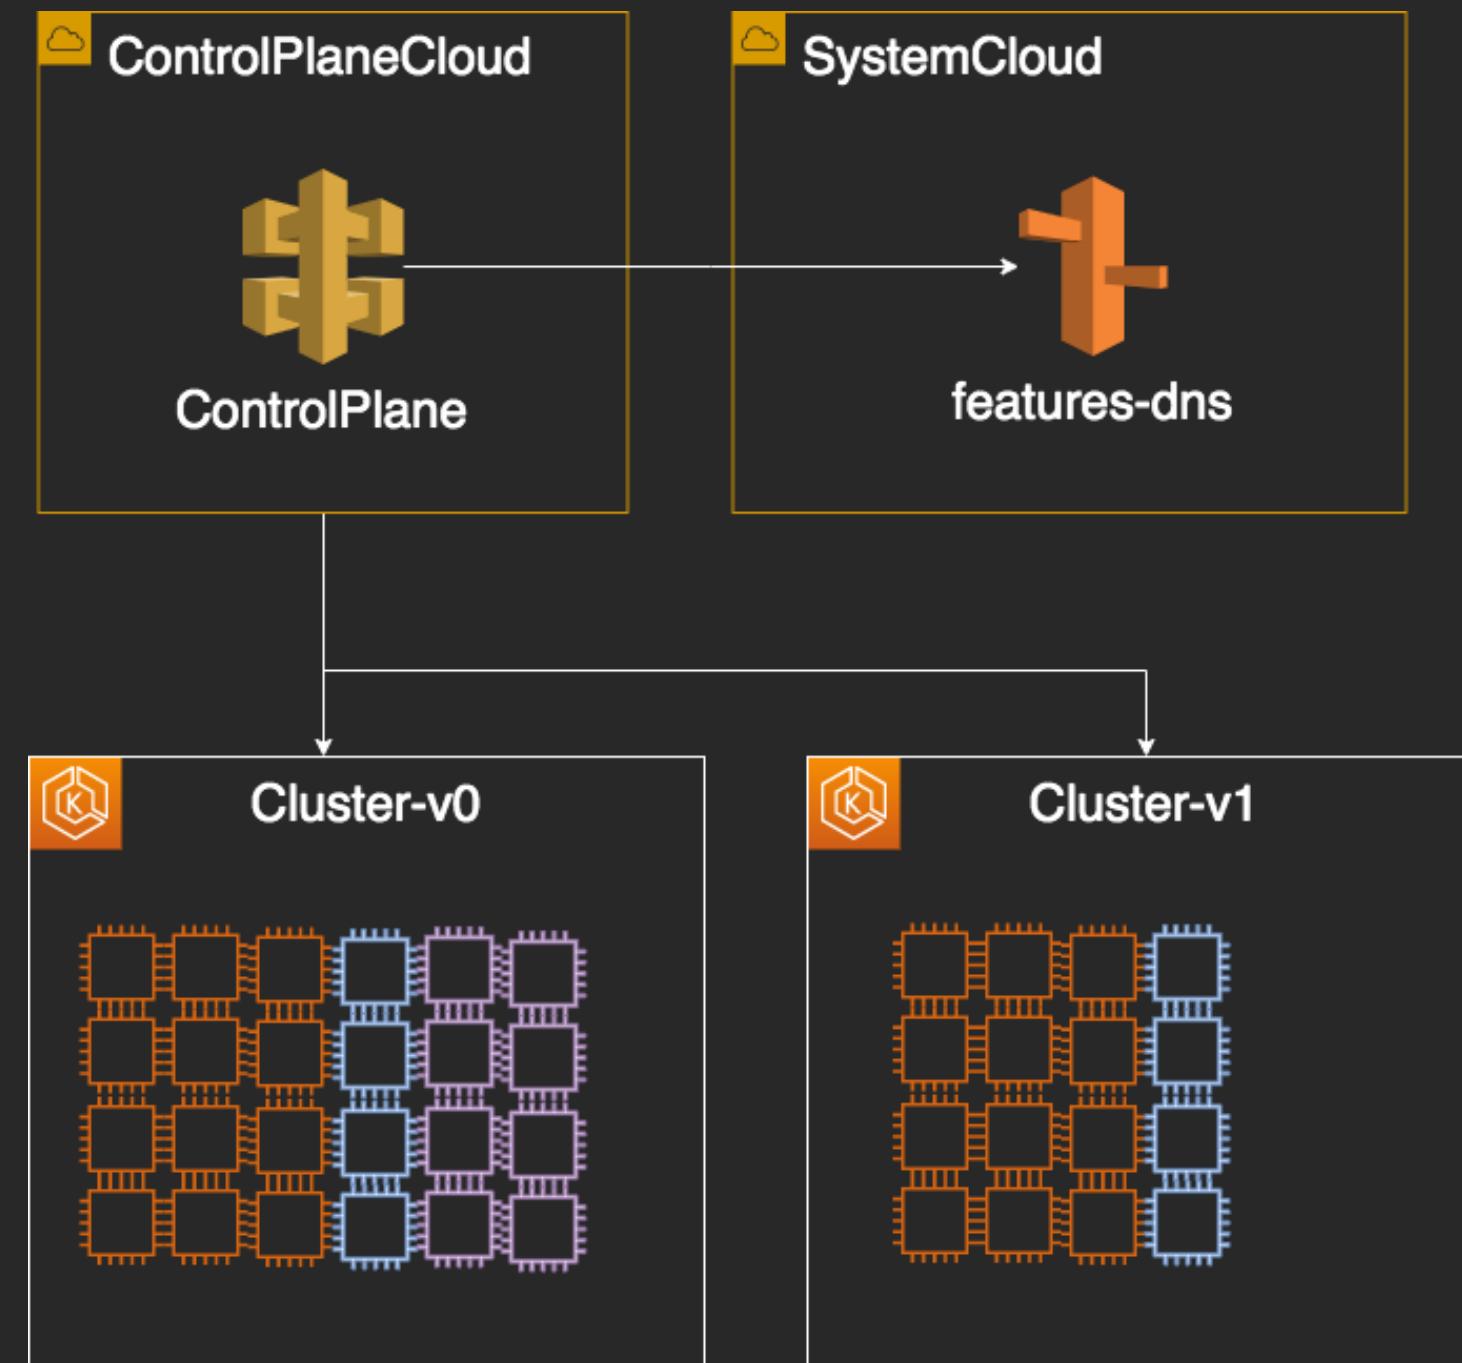
Cluster rollouts



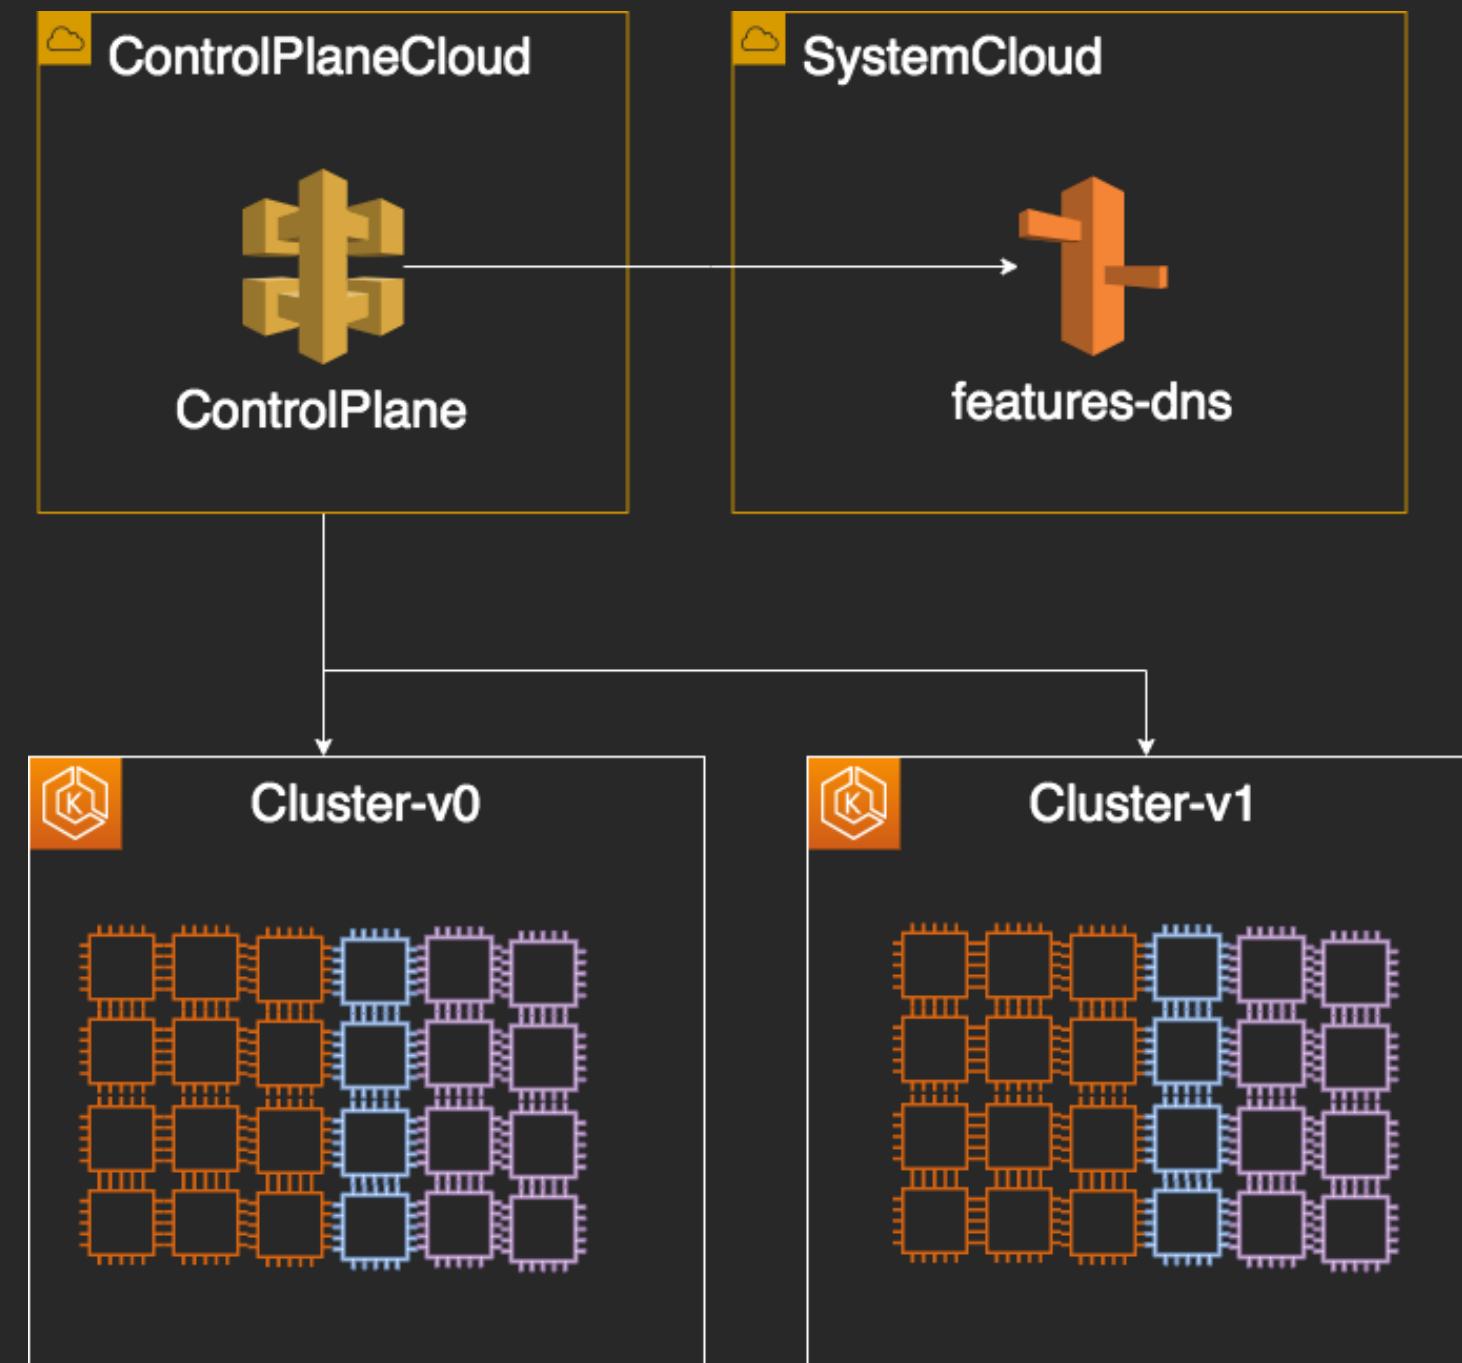
Cluster rollouts



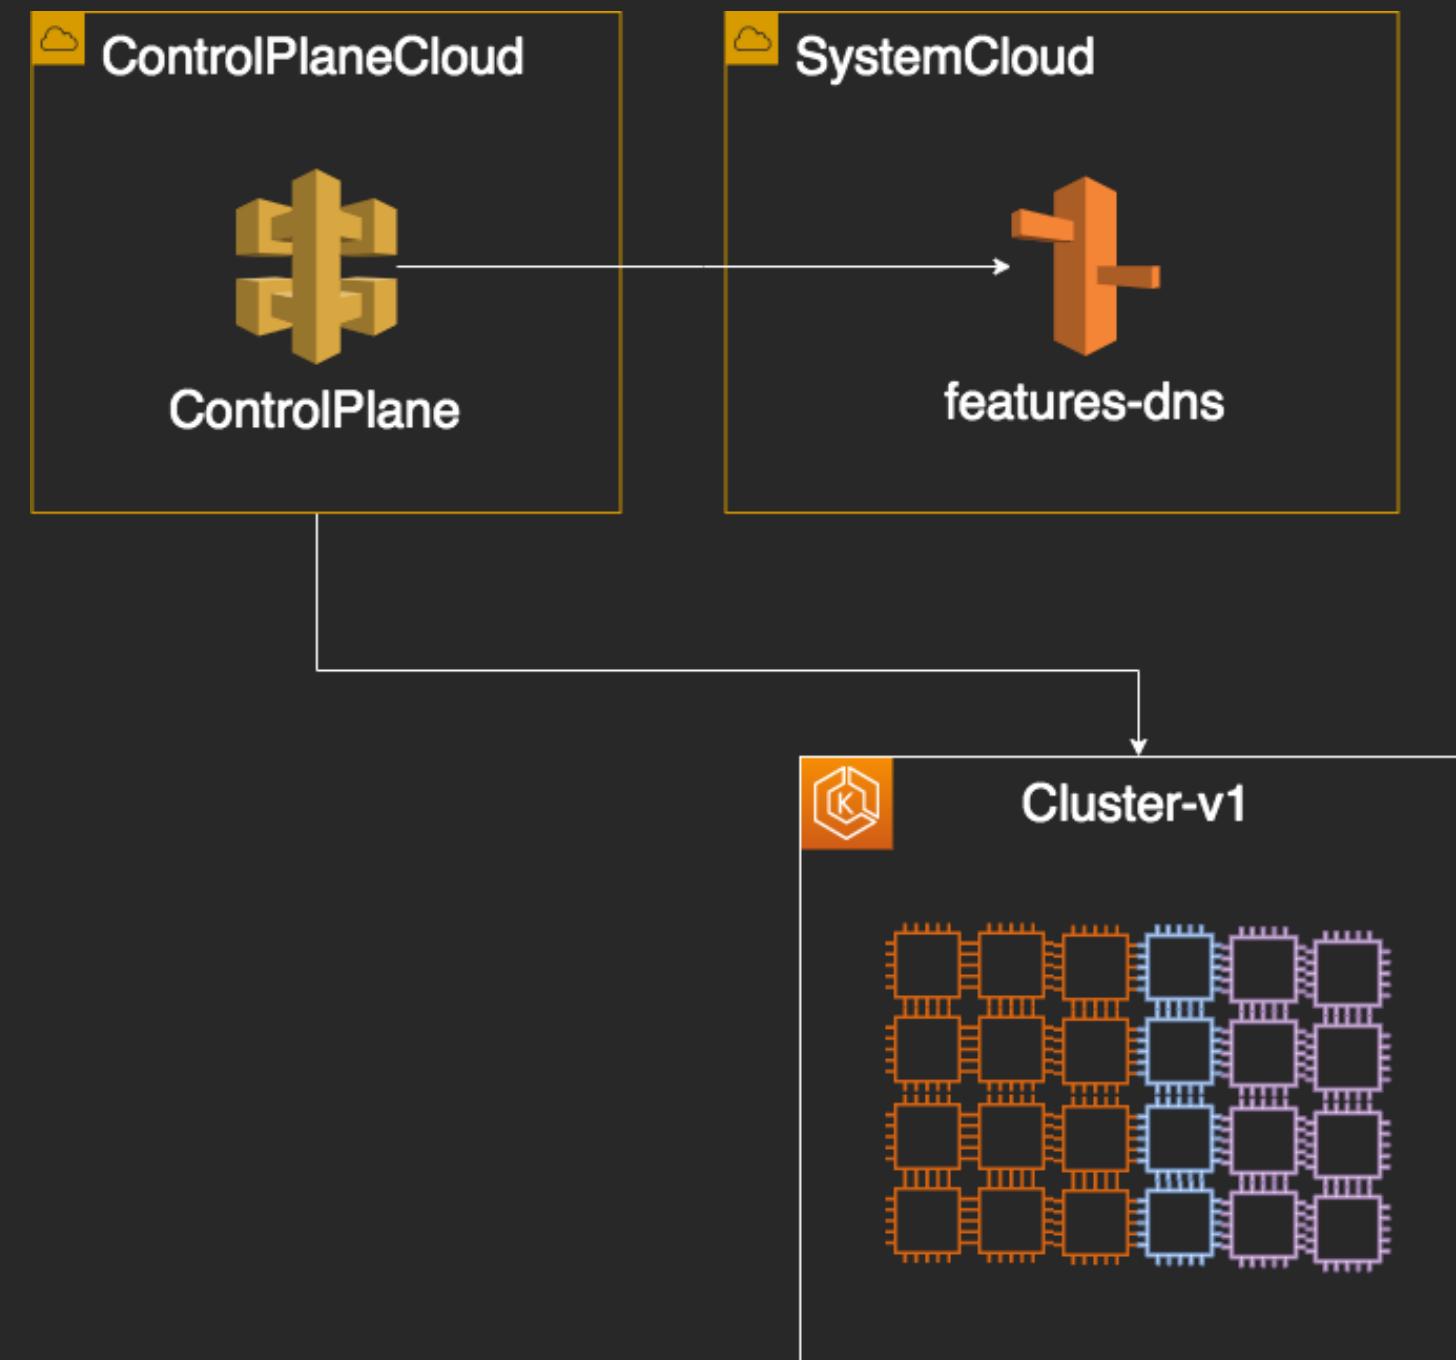
Cluster rollouts



Cluster rollouts



Cluster rollouts



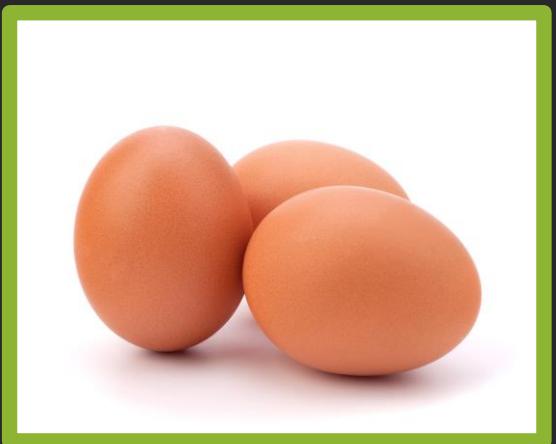
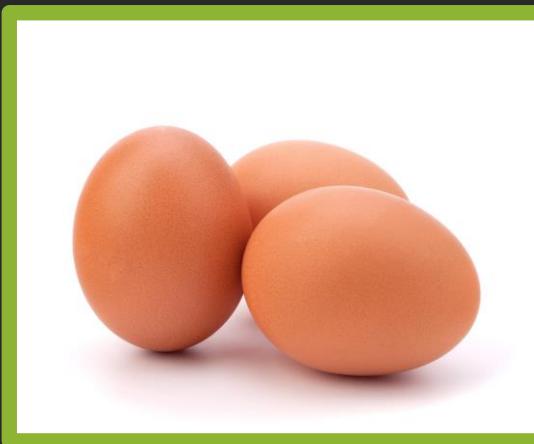
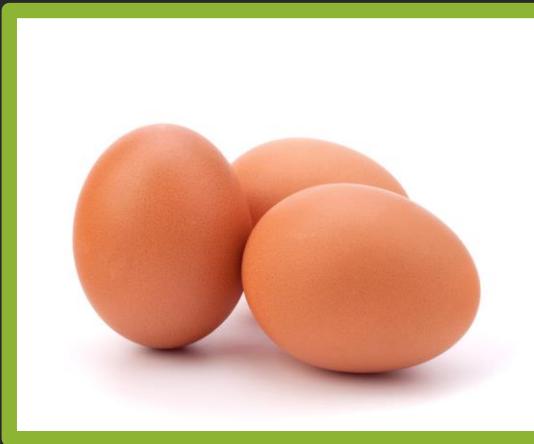
Cell-based architecture

- Don't put your eggs in one bucket

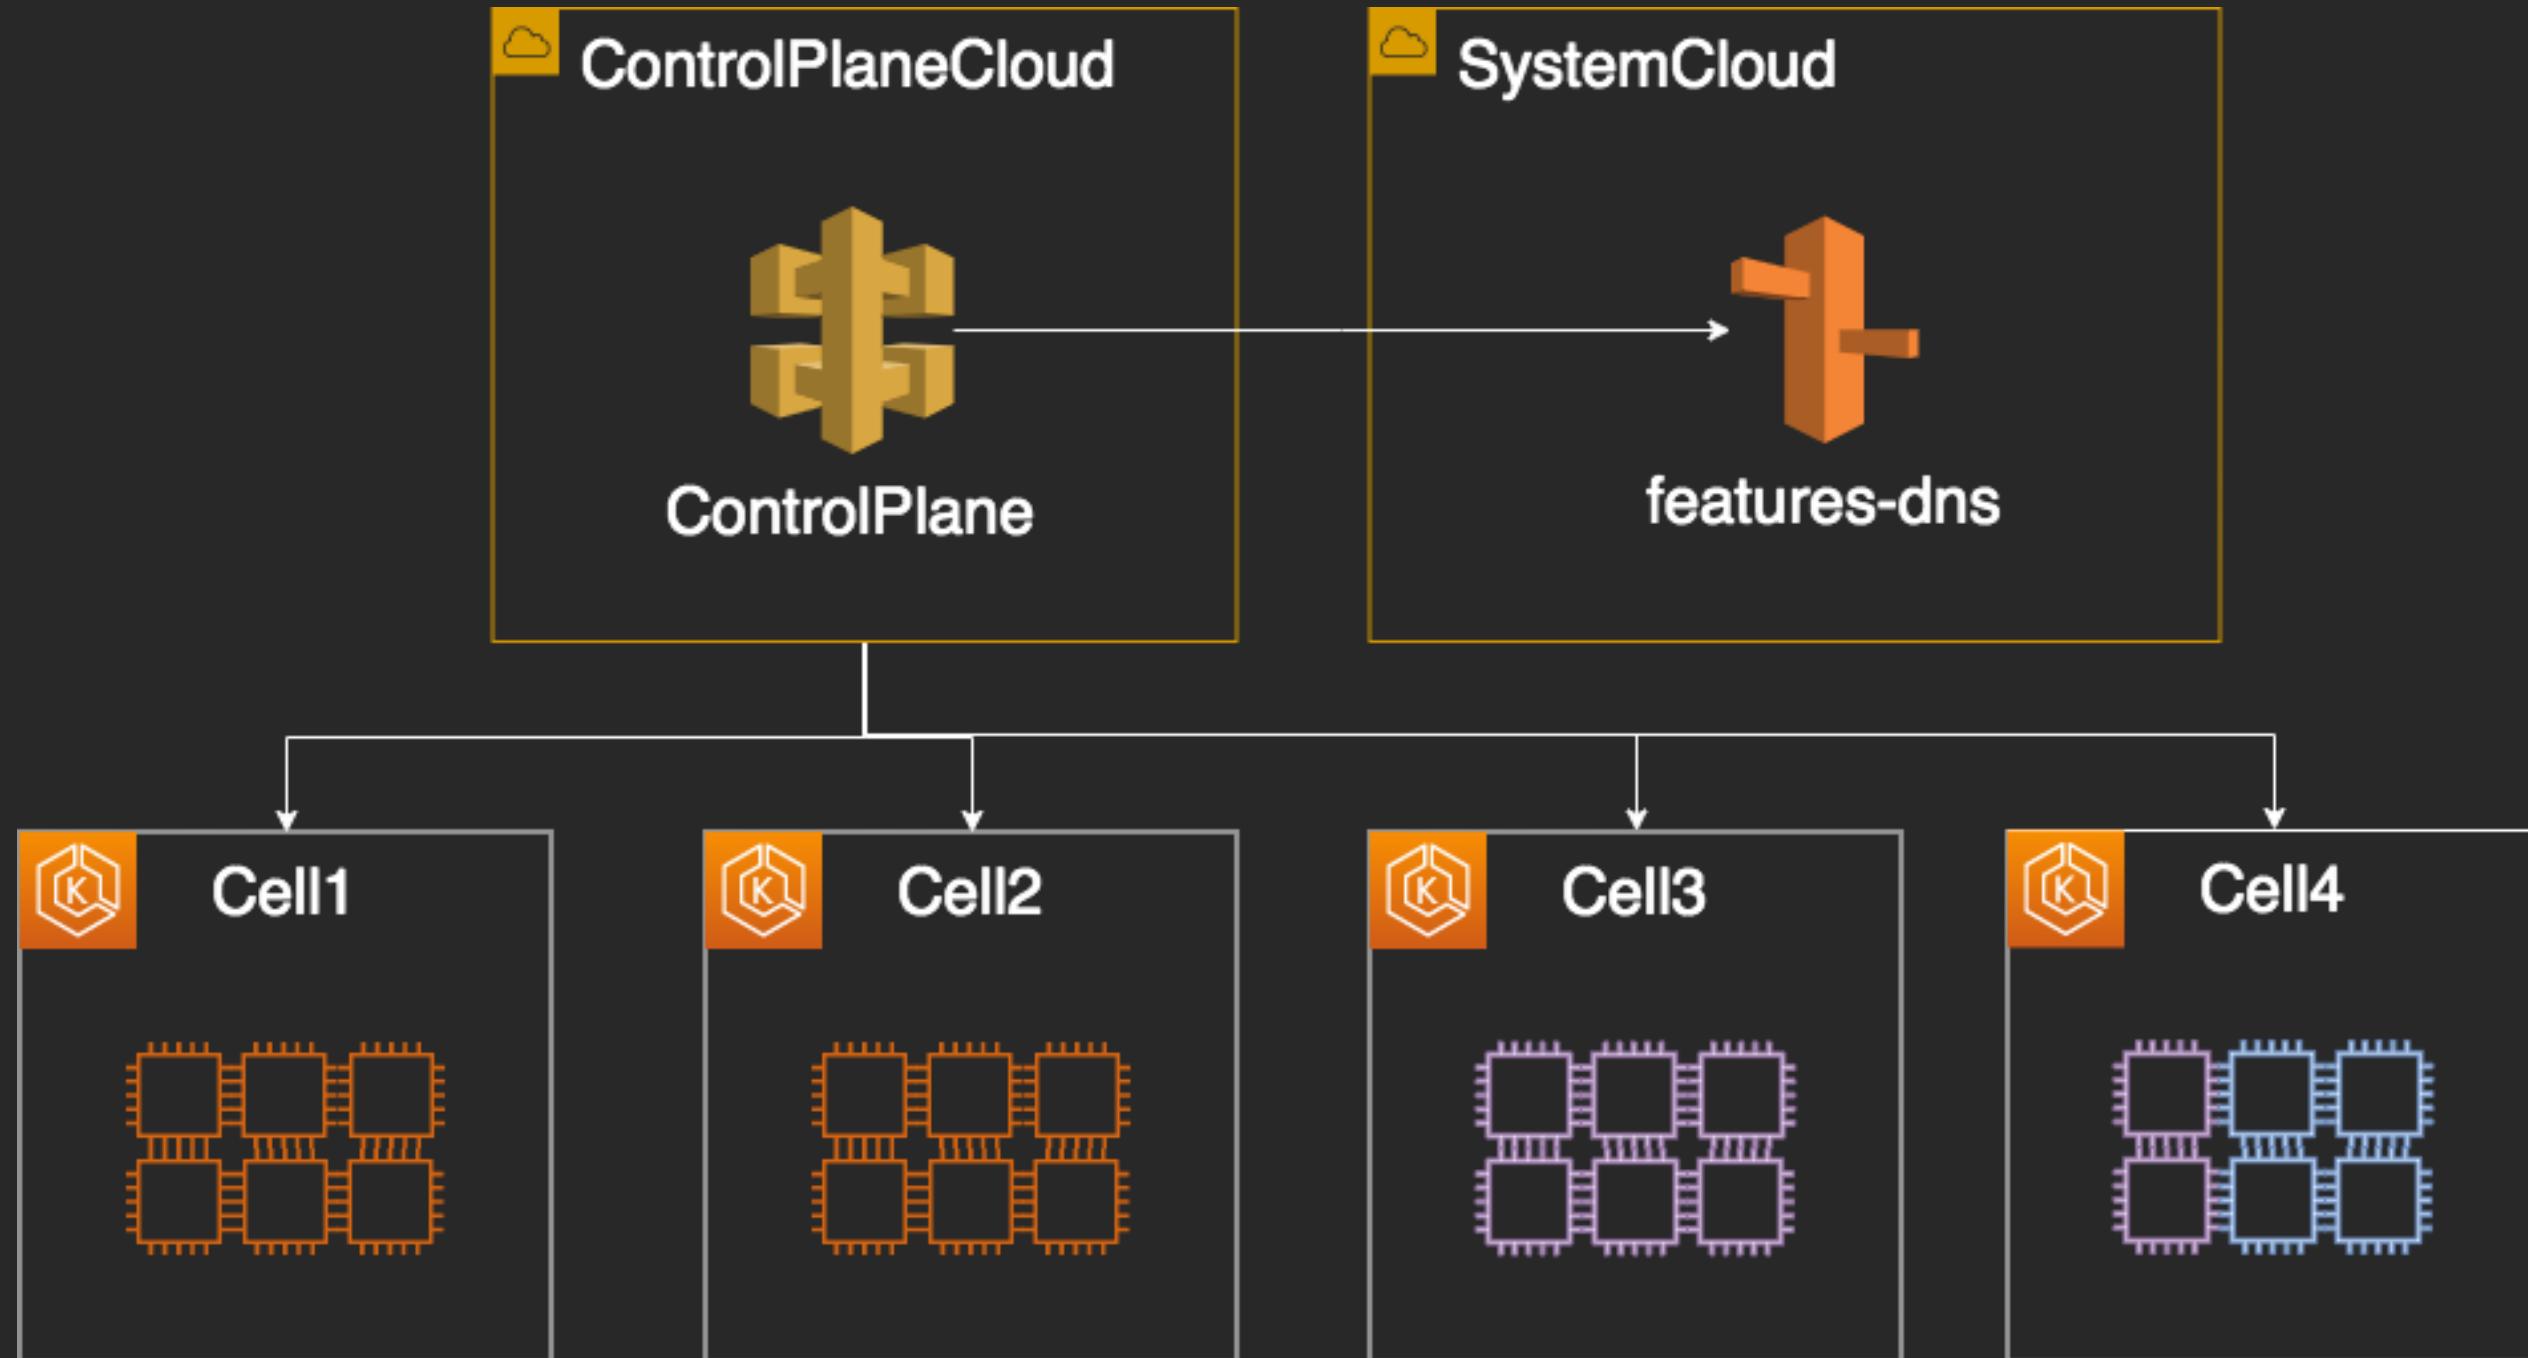


Cell-based architecture

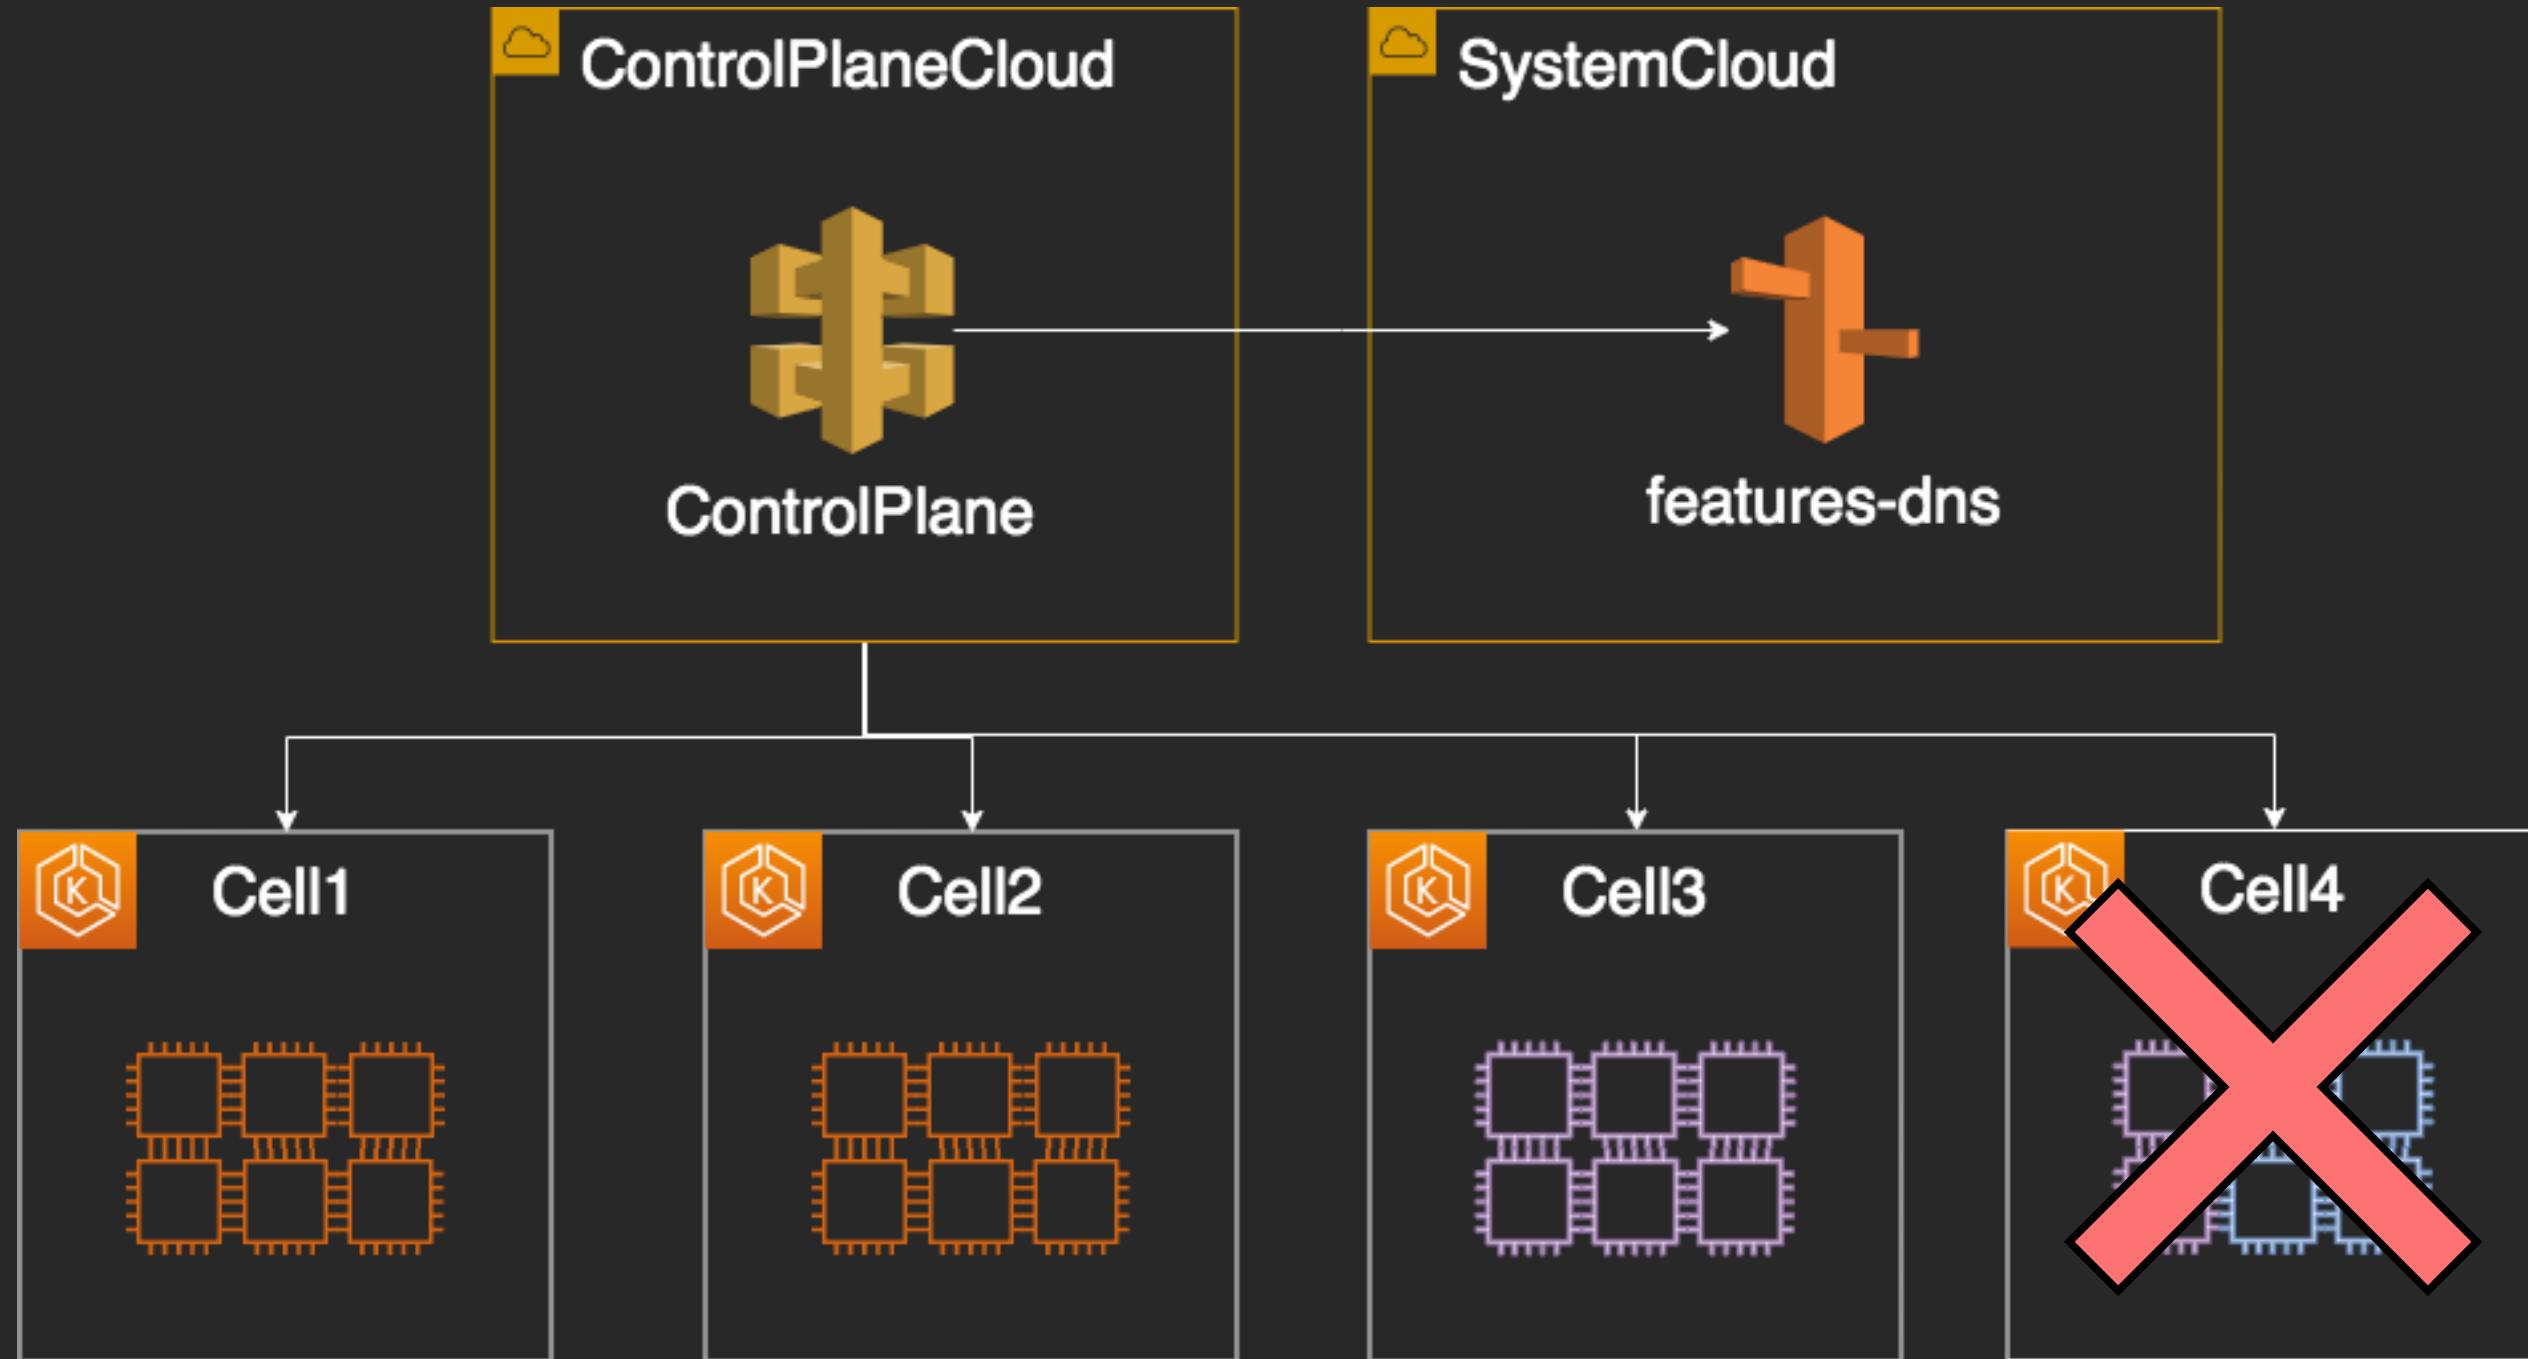
- Don't put your eggs in one basket
- Higher testability
- Higher scalability
- Reduces blast radius



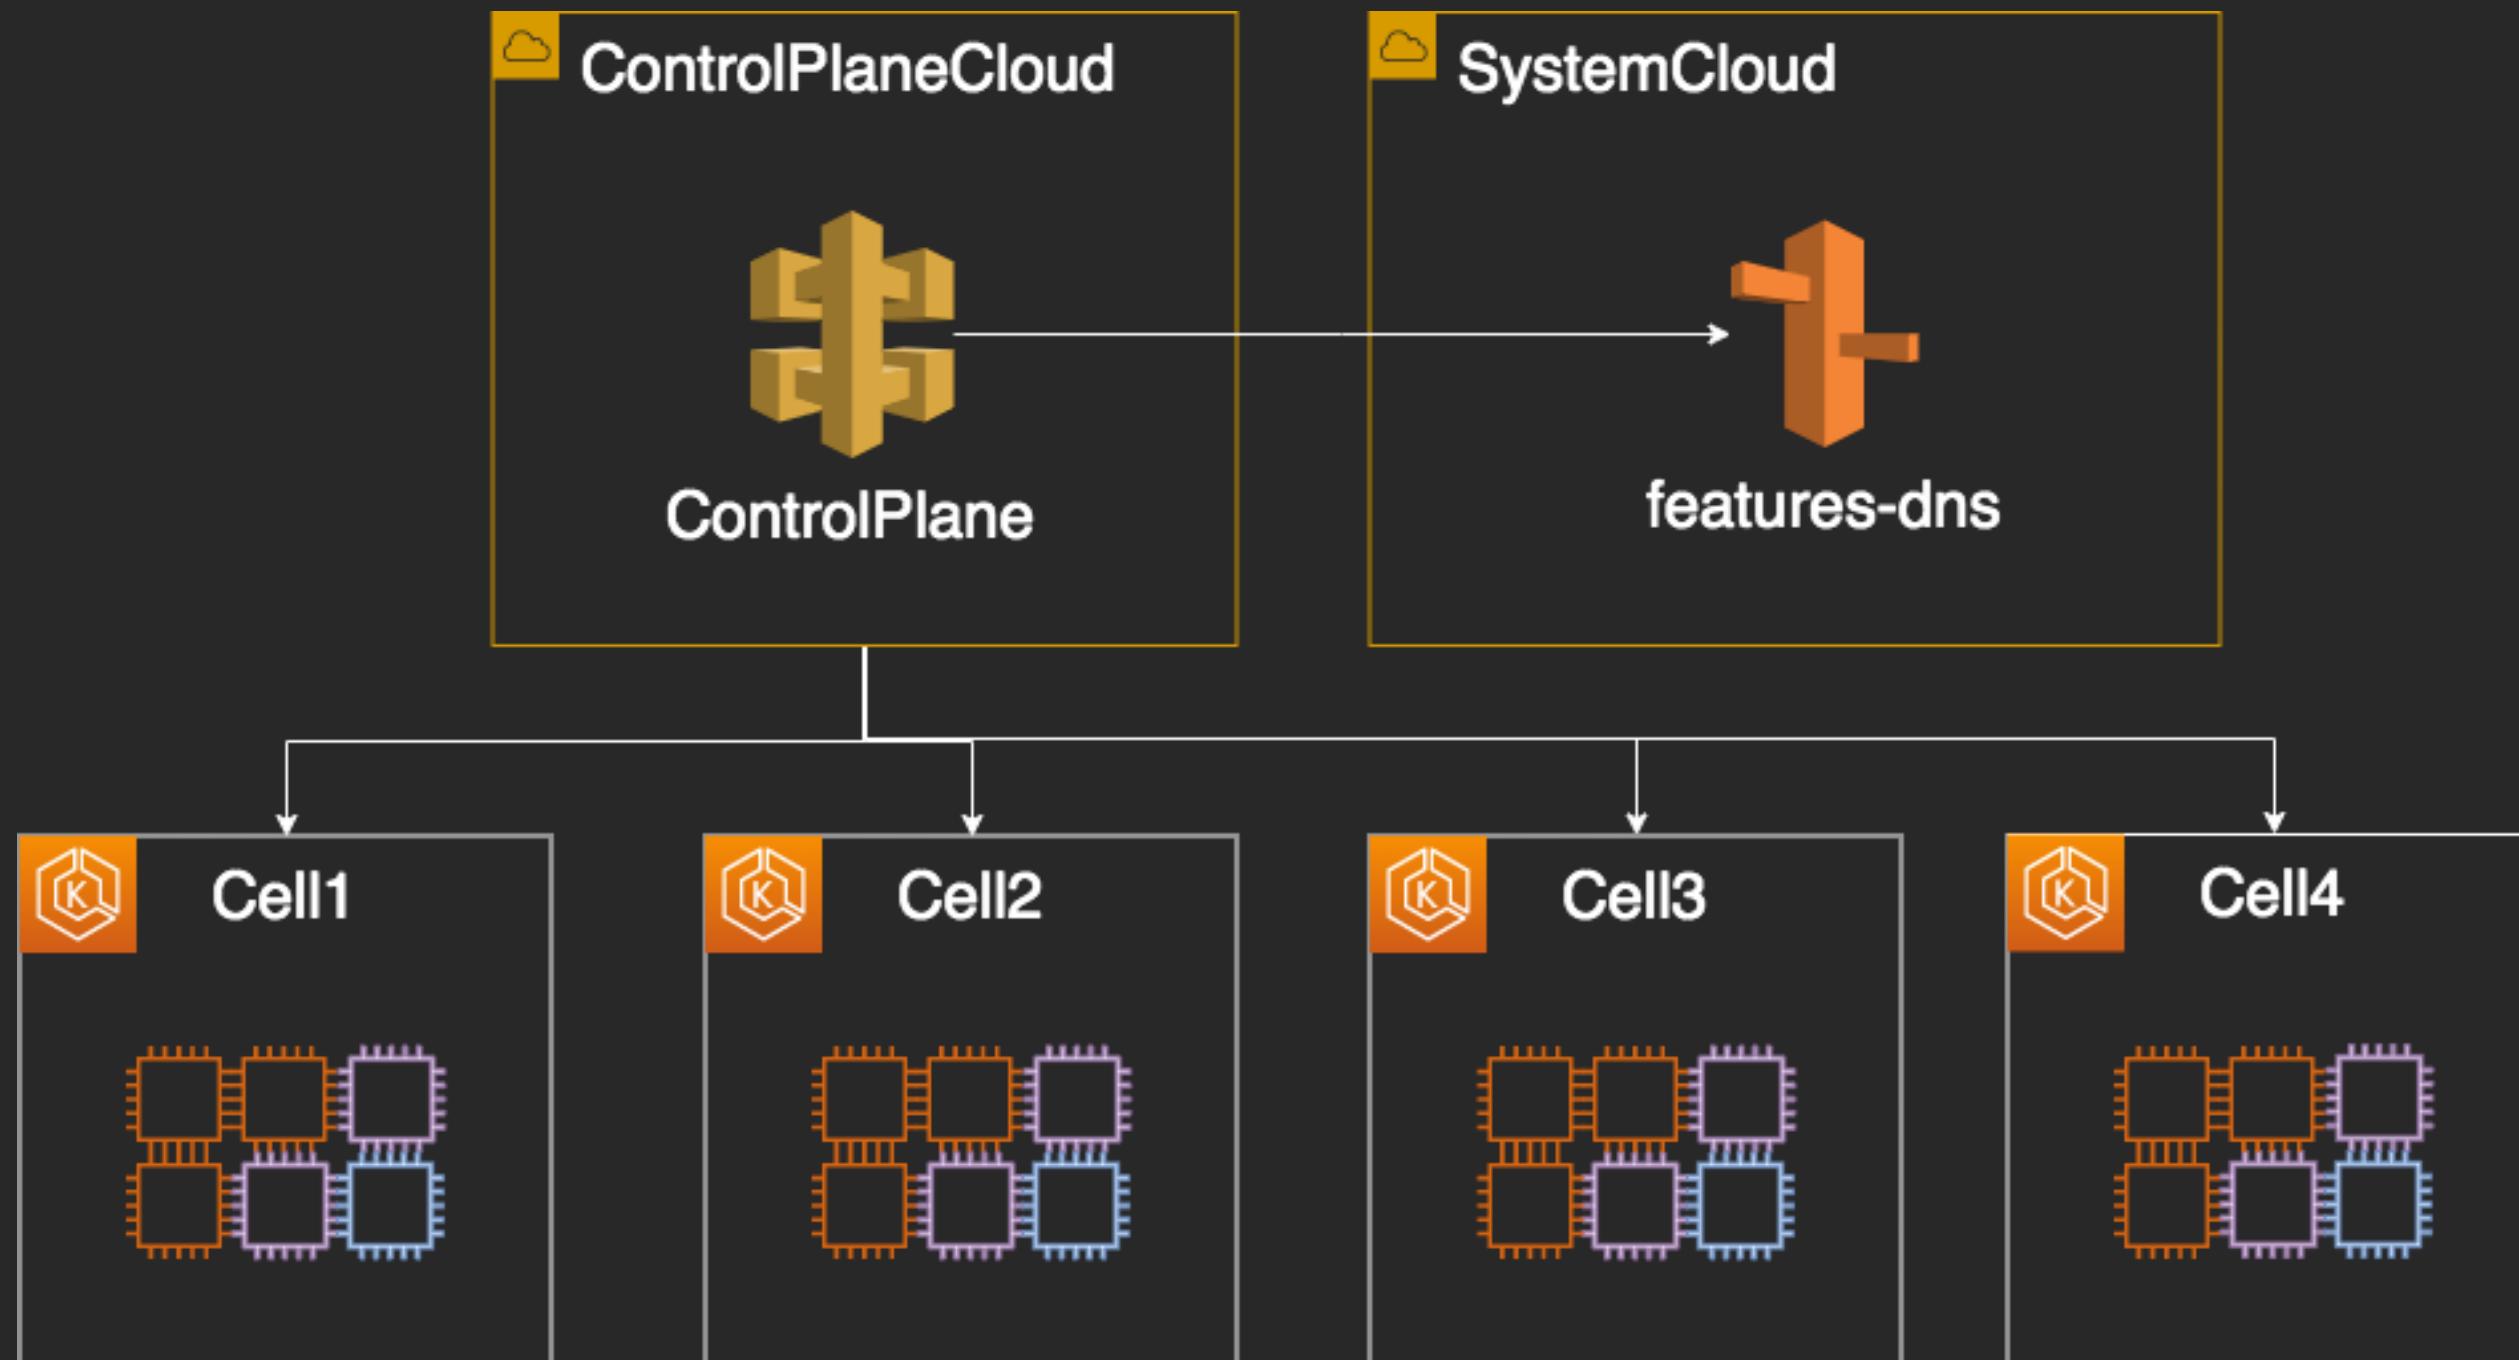
Cell-based architecture



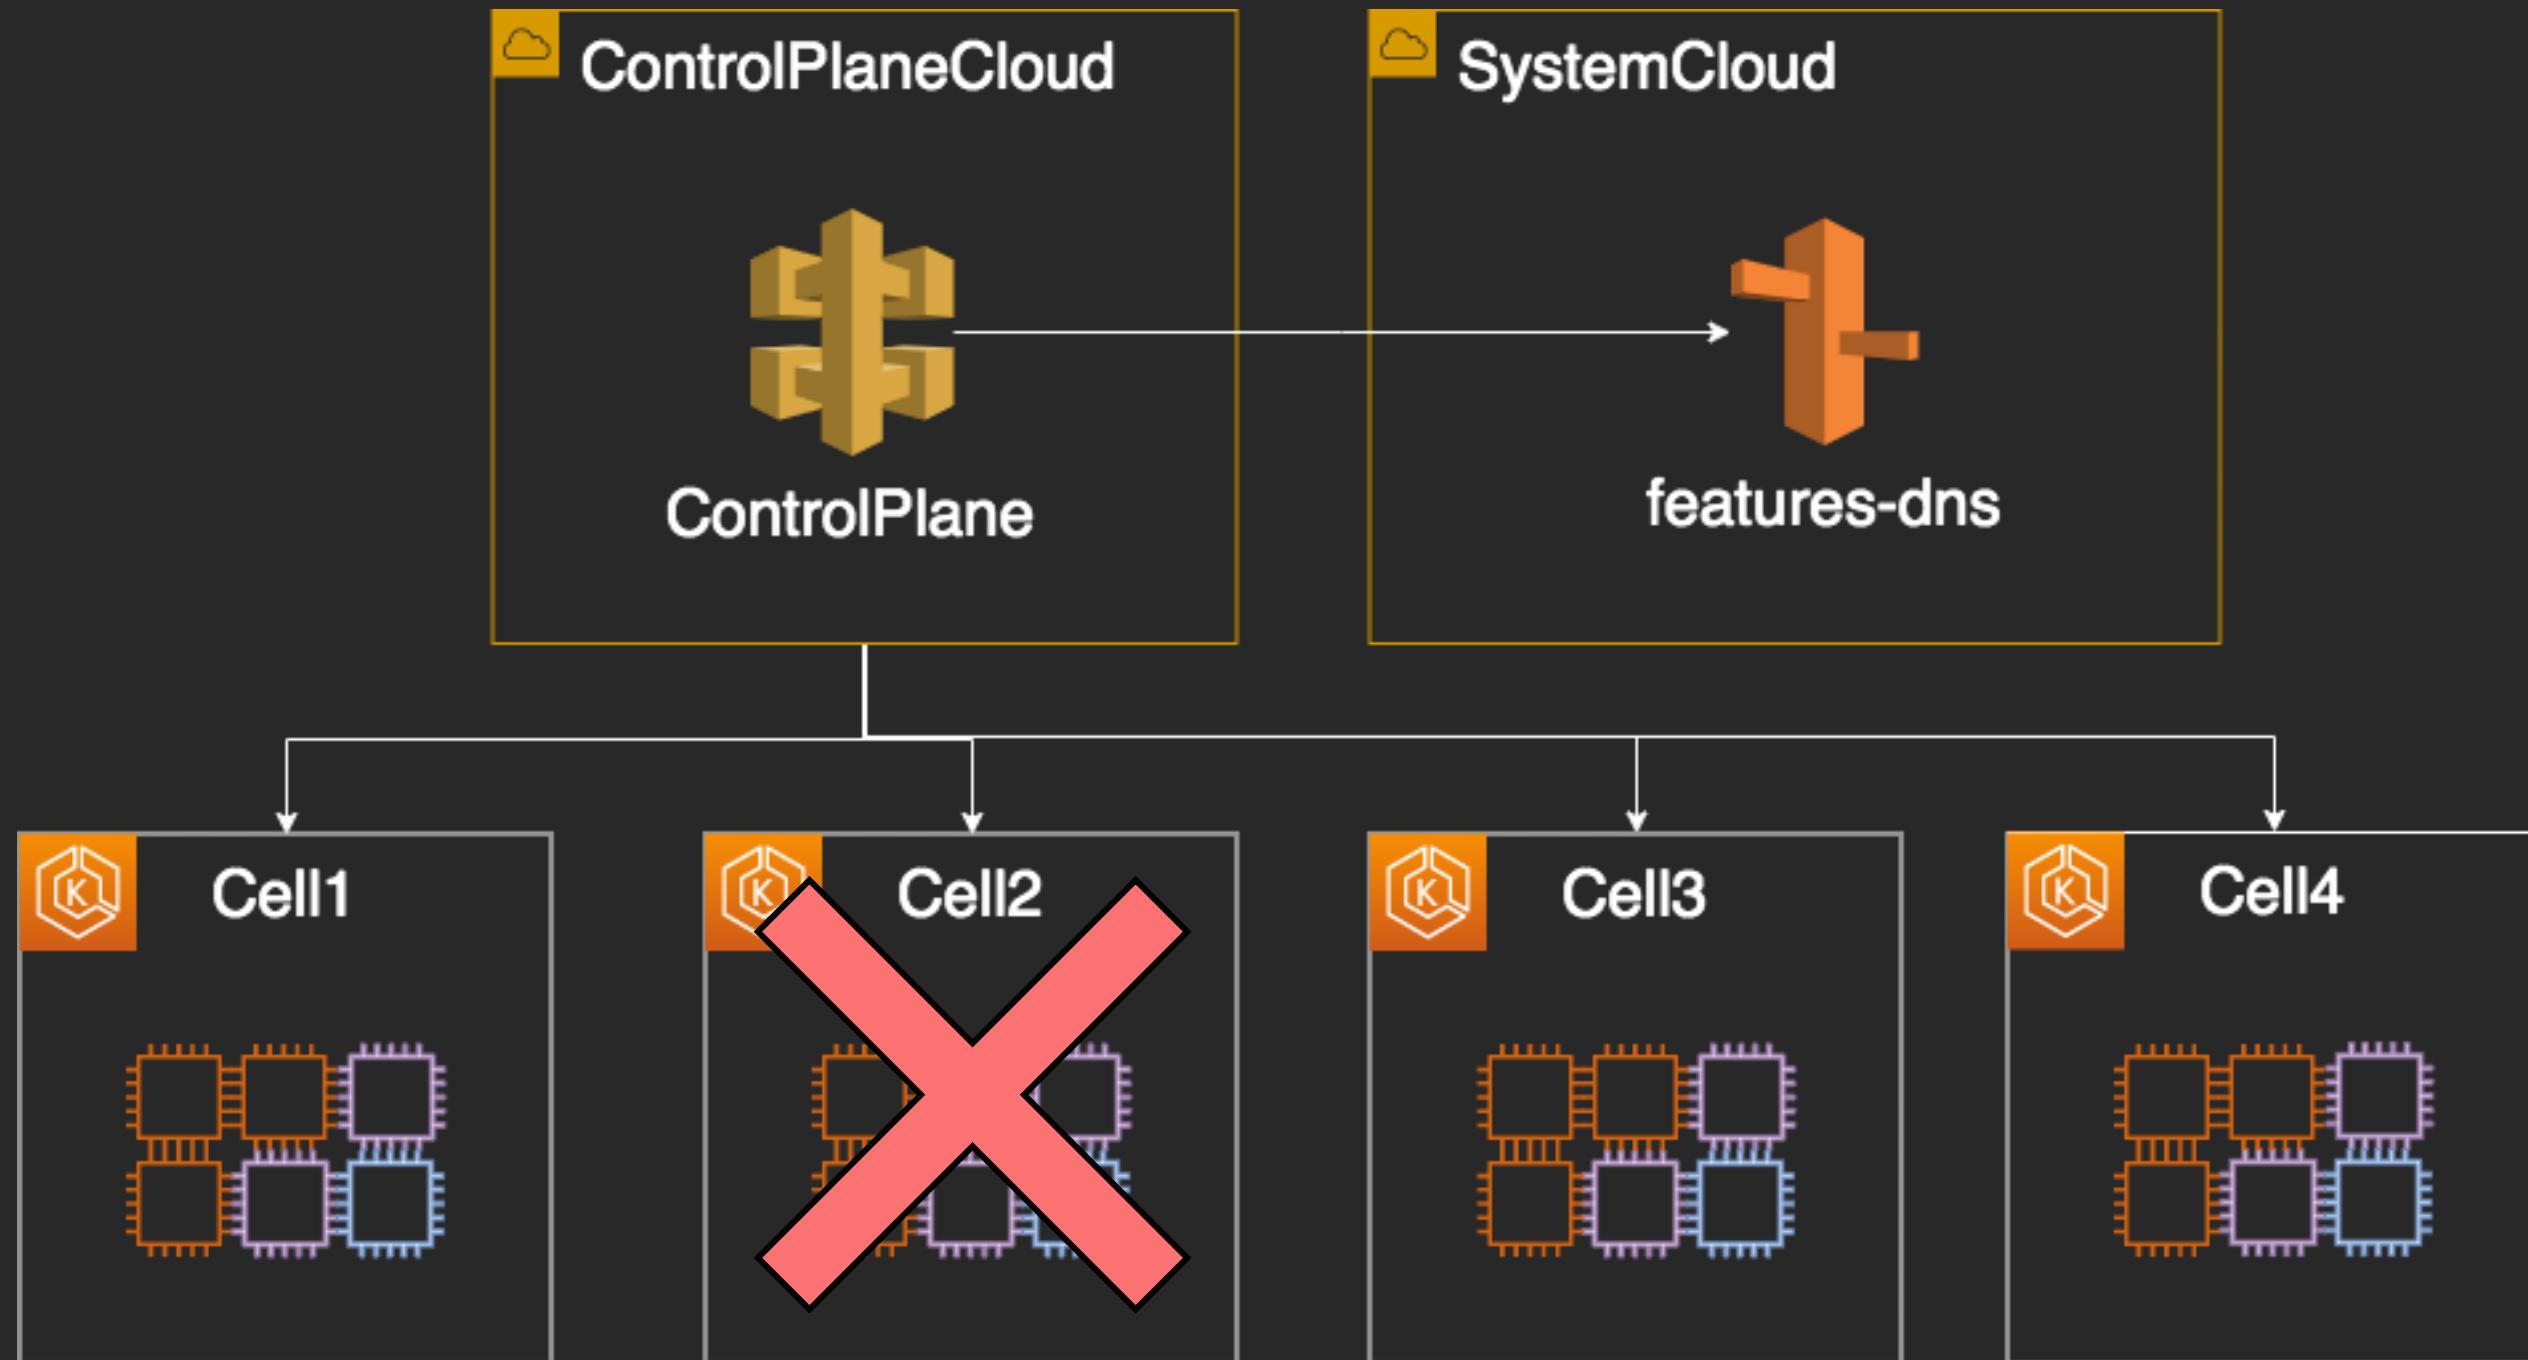
Cell-based architecture

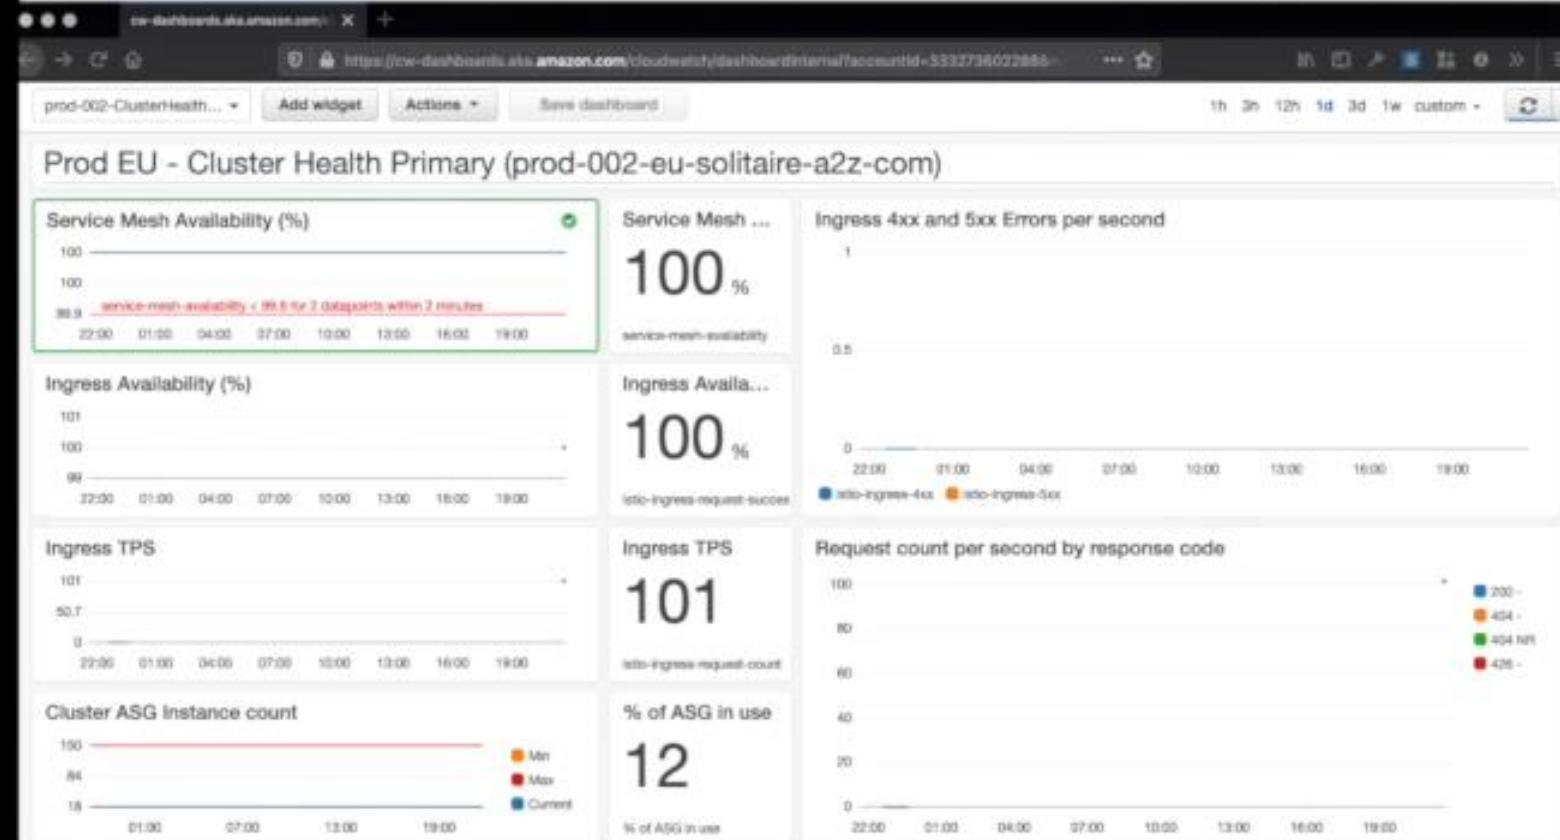
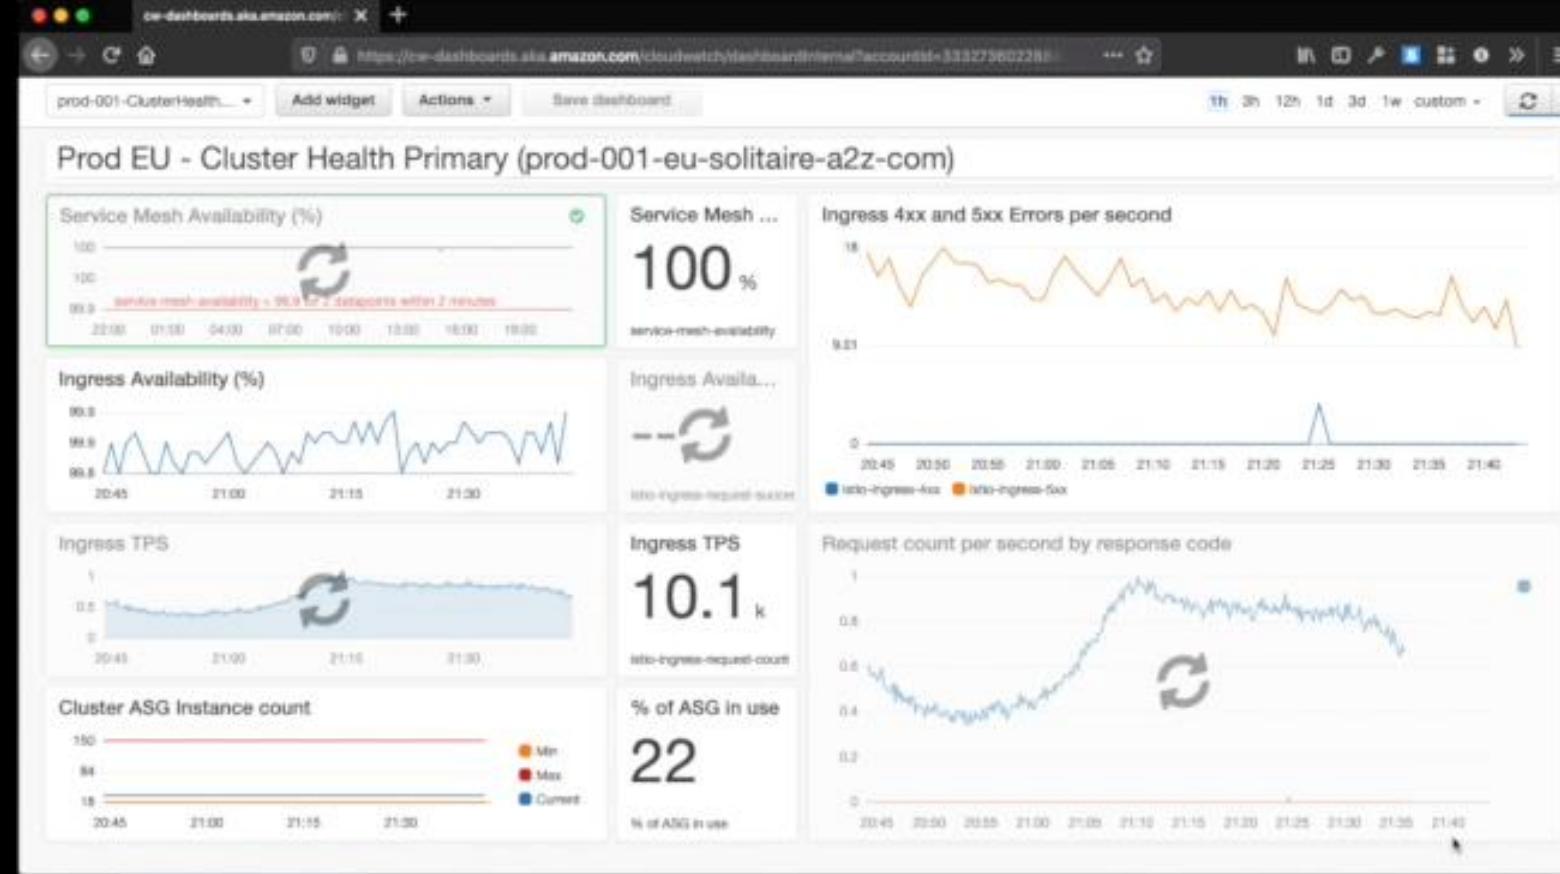


Cell-based architecture



Cell-based architecture





Key takeaways

Ownership is hard, managed services make it easier

Monitor used and reserved resources to optimize over time

Understand your cold start behavior and SLAs to configure elasticity

Design for ephemeral infrastructure

Thank you!



Please complete the session
survey in the mobile app.