

The background of the image features a vibrant, multi-colored gradient. It starts with a dark blue in the top left corner, transitioning through purple, magenta, and red towards the center. From the bottom right corner, there is a large, diagonal band of color that shifts from orange to yellow to light blue. The overall effect is dynamic and modern.

AWS
re:Invent

CON307-S

Top 5 container and Kubernetes best practices

Grace Andrews

Solutions Consultant
New Relic

Safe Harbor

This presentation and the information herein (including any information that may be incorporated by reference) is provided for informational purposes only and should not be construed as an offer, commitment, promise or obligation on behalf of New Relic, Inc. (“New Relic”) to sell securities or deliver any product, material, code, functionality, or other feature. Any information provided hereby is proprietary to New Relic and may not be replicated or disclosed without New Relic’s express written permission.

Such information may contain forward-looking statements within the meaning of federal securities laws. Any statement that is not a historical fact or refers to expectations, projections, future plans, objectives, estimates, goals, or other characterizations of future events is a forward-looking statement. These forward-looking statements can often be identified as such because the context of the statement will include words such as “believes,” “anticipates,” “expects” or words of similar import.

Actual results may differ materially from those expressed in these forward-looking statements, which speak only as of the date hereof, and are subject to change at any time without notice. Existing and prospective investors, customers and other third parties transacting business with New Relic are cautioned not to place undue reliance on this forward-looking information. The achievement or success of the matters covered by such forward-looking statements are based on New Relic’s current assumptions, expectations, and beliefs and are subject to substantial risks, uncertainties, assumptions, and changes in circumstances that may cause the actual results, performance, or achievements to differ materially from those expressed or implied in any forward-looking statement. Further information on factors that could affect such forward-looking statements is included in the filings New Relic makes with the SEC from time to time. Copies of these documents may be obtained by visiting New Relic’s Investor Relations website at ir.newrelic.com or the SEC’s website at www.sec.gov.

New Relic assumes no obligation and does not intend to update these forward-looking statements, except as required by law. New Relic makes no warranties, expressed or implied, in this presentation or otherwise, with respect to the information provided.

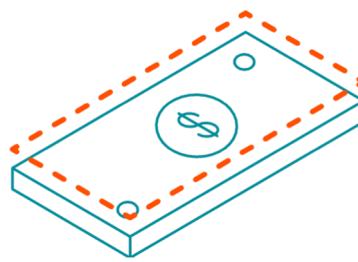


Grace Andrews
Solutions Consultant
Living in Portland, OR

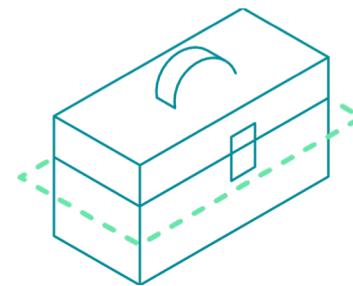




Kubernetes helps you
move faster



Kubernetes is
cost-efficient



Kubernetes is
portable

Kubernetes adoption journey



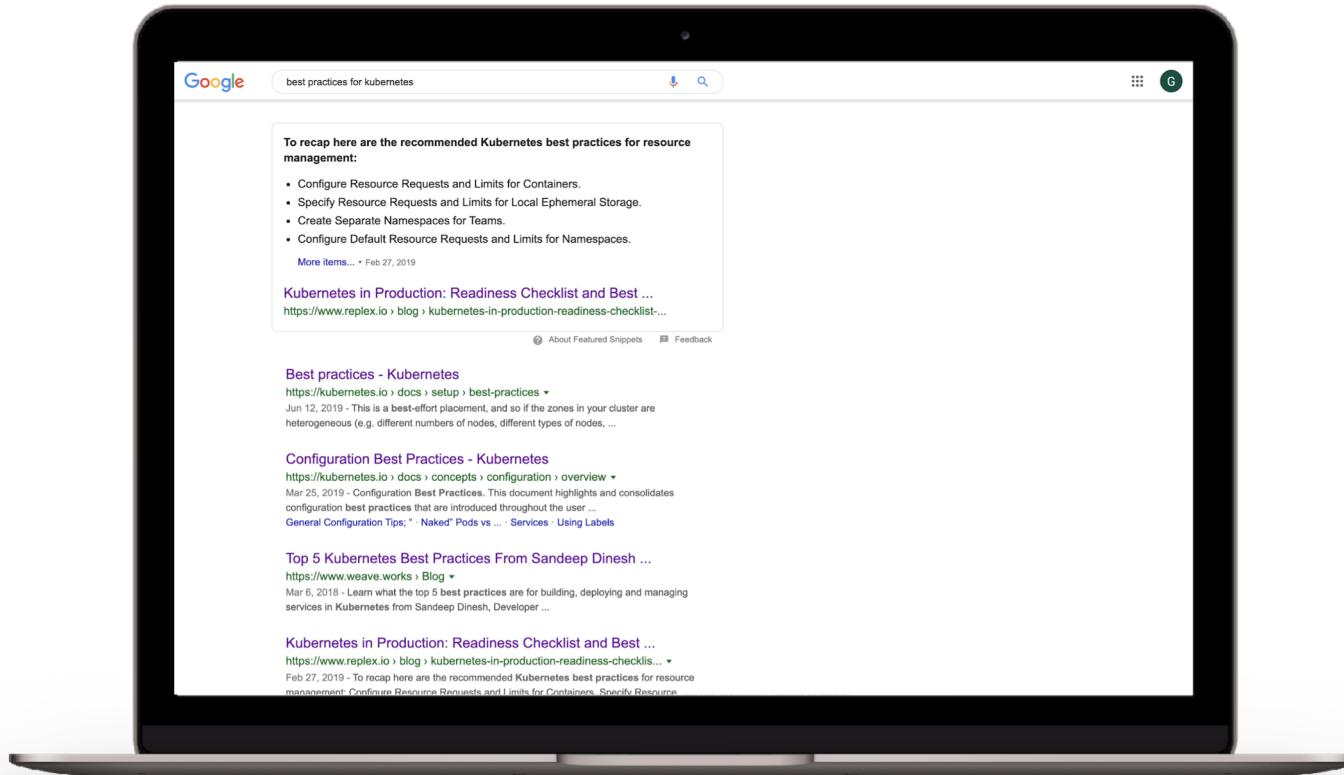
"Help, I need to
Kubernetes"



**Stable
Kubernetes
Cluster**

Kubernetes
Promised
Land

In the beginning... you know the rest

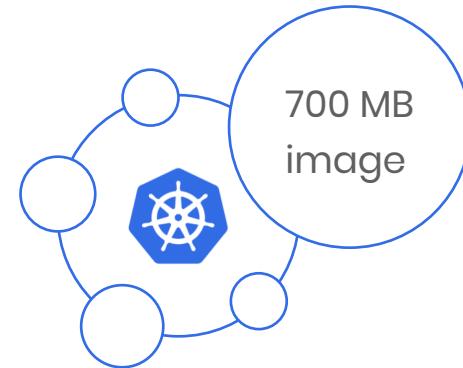


Best practices

-  Build small container images
-  Leverage namespaces
-  Set up health checks
-  Use resource requests & limits; multi-functional alerts
-  Understand the context of problems

Build small container images

Why?



- ✓ Improve time-to-build & time-to-pull
- ✓ Fewer security vulnerabilities

Build small container images

1

Kubernetes builds

Why?

- Quicker builds
- Smaller attack surface
- Less storage
- Faster image pulls/cold starts

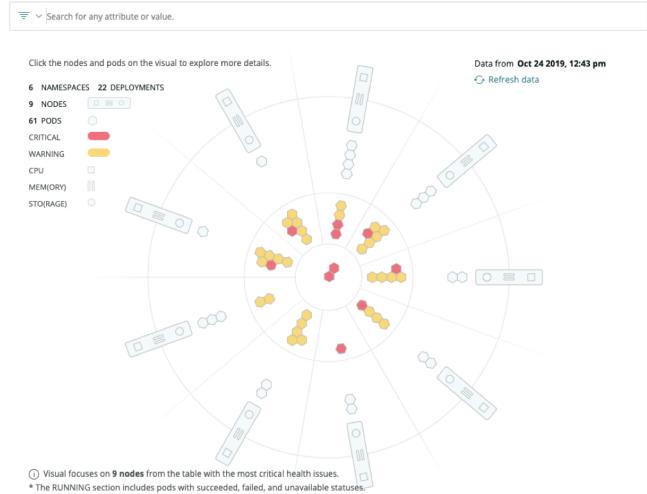
How?

- Identify your app(s) need
- Choose the appropriate base image
- Avoid unnecessary resources (packages + files)

Base node.js image	Alpine
FROM node:latest WORKDIR /app COPY .. RUN npm install -- prod EXPOSE 8080 CMD npm start	FROM node:alpine WORKDIR /app COPY .. RUN npm install -- prod EXPOSE 8080 CMD npm start
700 MB	65 MB

Kubernetes cluster explorer

Reduce time in PENDING state



Leverage the Docker build cache

Layer 5 | *CMD npm start*

Layer 4 | *Expose 8080*

Layer 3 | *RUN npm install -- prod*

Layer 2 | *COPY..*

Layer 1 | *WORKDIR/app*

Base image | *FROM node:alpine*

Layer 6 | *CMD npm start*

Layer 5 | *Expose 8080*

Layer 4 | *COPY SRC_CODE SRC_CODE*

Layer 3 | *RUN npm install --prod*

Layer 2 | *COPY package.json .*

Layer 1 | *WORKDIR /app*

Base image | *FROM node:alpine*

Add your source code as late as possible
Base image and dependencies can be cached

Using multi-stage builds

Layer X | *# Build your final image*

Layer 1 | *COPY --from=dependencies /app/prod_node_modules /node_modules*

Layer 0 | *FROM base AS release*

Layer X | *# Build your dependencies*

Layer 0 | *FROM base AS dependencies*

Layer x | *# Build the base image*

Layer 0 | *FROM node:alpine AS base*

Multi-stage builds! Break down your Dockerfile in multiple stages

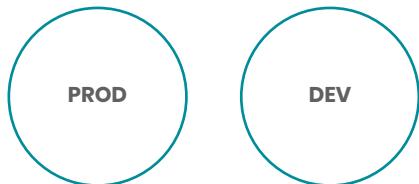
The final image will be based on the latest FROM

Back to the cluster

Example

What's in a name?

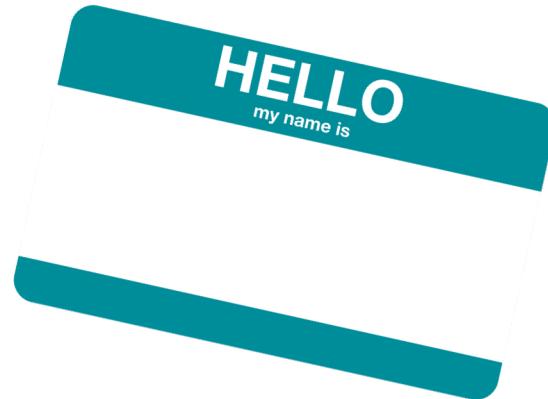
Small team



Multiple teams



Enterprise



Kubernetes observability – Take advantage of namespaces

2

Kubernetes namespace

Error profiles, system health



Why?

Improve manageability & security

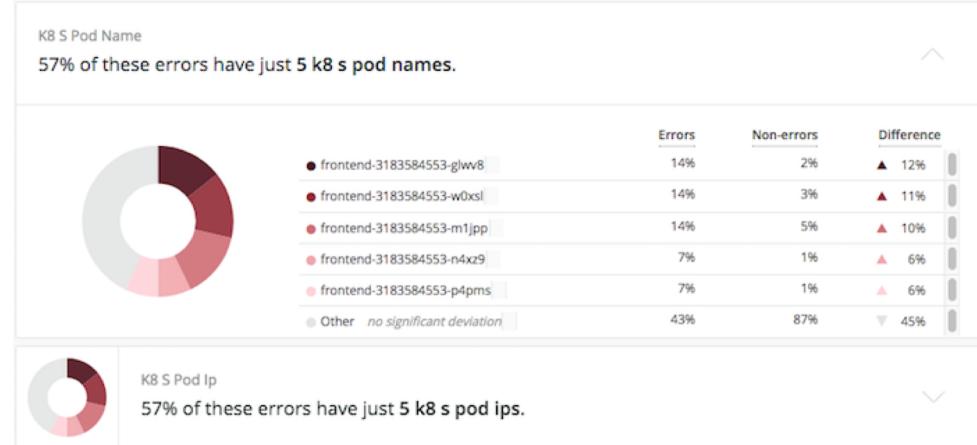
- By isolating Kubernetes resource for different teams

Improve performance

- Kubernetes API working on smaller set of objects

How?

- Modify config file (YAML)
- Update YOUR_CLUSTER_NAME
- Kubectl apply to initiate the changes



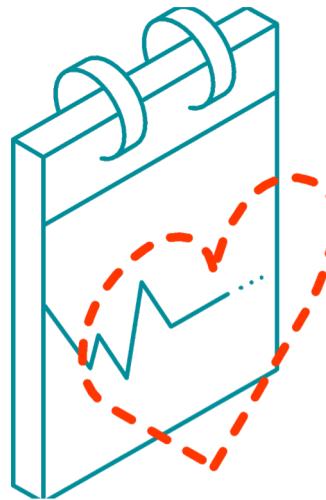
Back to the cluster

Example

Set up health checks

Readiness probes

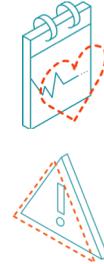
Ready, set, GO!



Liveness probes

Dead or alive?

Kubernetes observability – Health checks



3

Kubernetes health checks

Why?

- Faster troubleshooting and issue resolution
- Watch all events happening in your cluster
- Track specific events such as scaling events

How?

- Download YAML file
- Modify the YAML file with the proper probes
- Remember to set initialDelaySeconds
- Use kubectl apply to update changes

Track vital cluster events

The screenshot displays the New Relic Kubernetes cluster explorer interface. On the left, a sidebar shows a tree view of the cluster structure, including namespaces like 'k8s.default.svc' and pods like 'minecraft-dev-minecraft-5fbfb59b55-thv9'. A red box highlights the 'Events' button in this sidebar. The main area shows two panels: one for 'Events' (last 24 hrs) and another for 'Pods'. The 'Events' panel lists log entries with columns for LAST TIMESTAMP, FIRST TIMESTAMP, REASON, MESSAGE, and COUNT. A red box highlights the 'Events' tab at the top of this panel. The 'Pods' panel shows a list of pods with their status, reason, message, and CPU/Memory usage over time. A red box highlights the 'See logs' button in this panel.

Health checks

```
readinessProbe:  
  exec:  
    command:  
      - cat  
      - /tmp/healthy  
  initialDelaySeconds: 5  
  periodSeconds: 5
```

Readiness probes

Is the application ready to accept traffic or not?

Tip: A synthetics check can travel the full path from the load balancer through the traffic proxy to the container and verify that the network path is working appropriately.

Health checks

Liveness probes

Is our application even alive, or is it dead?

Why is our application dead?

Correlate your k8s data with application data!

```
apiVersion: v1
kind: Pod
metadata:
  labels:
    test: liveness
  name: liveness-exec
spec:
  containers:
  - name: liveness
    image: k8s.gcr.io/busybox
    args:
    - /bin/sh
    - -c
    - touch /tmp/healthy; sleep 30; rm -rf /tmp/healthy; sleep 600
    livenessProbe:
      exec:
        command:
        - cat
        - /tmp/healthy
      initialDelaySeconds: 5
      periodSeconds: 5
```

pods/probe/exec-liveness.yaml 

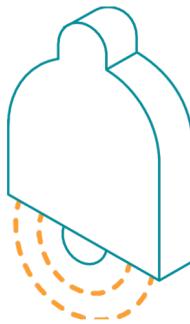
Remark:

Set **initialDelaySeconds** for liveness probe or pod will be restarted before it is ready

Back to the cluster

Example

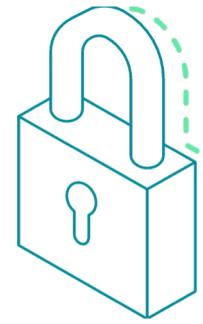
Set up requests & limits



Requests

What the container
is guaranteed to get

Example: When a container requests a resource,
Kubernetes will only schedule it on a node
that can give it that resource



Limits

Make sure a container never
goes above a certain value

The container can only go up to the limit,
and then it's restricted

Kubernetes observability – Requests & limits

4

Kubernetes resource requests & limits

Warning + alerts



Why?

Lack of resources limits creates pod instability and possible termination

- Pod over CPU limit: Throttled
- Pod over memory limit: Terminated
- Pod in PENDING state if resources are not available

Inappropriately sized requests can cause scheduling issues

- If a value larger than the core count of your biggest node is set, your pod will never be scheduled

How?

1. Specify limits and requests within your config file ([Kubernetes.io](#))
2. Use kubectl apply to push changes
3. Use kubectl describe nodes to see a list of requests and limits

```
apiVersion: v1
kind: Pod
metadata:
  name: frontend
spec:
  containers:
    - name: db
      image: mysql
      env:
        - name: MYSQL_ROOT_PASSWORD
          value: "password"
      resources:
        requests:
          memory: "64Mi"
          cpu: "250m"
        limits:
          memory: "128Mi"
          cpu: "500m"
    - name: wp
      image: wordpress
      resources:
        requests:
          memory: "64Mi"
          cpu: "250m"
        limits:
          memory: "128Mi"
          cpu: "500m"
```

Example

Leveraging logs

Gain visibility

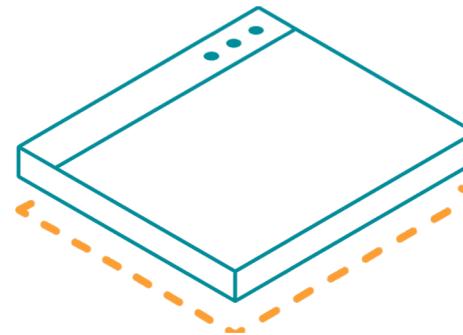
What's happening in the cluster?

Gain visibility

Are the pods behaving properly?

Gain stability

Has my host been compromised?



Kubernetes observability – Understanding problems

5

Kubernetes logging

No more context switching
Logs in context
Fast(!) log search



Why?

- Provide more context for environmental behaviors
- Faster troubleshooting when problems arise
- Can be used at the node, pod, and cluster levels

How?

1. Install log output plugin for Kubernetes
2. Update with proper licensing information
3. Use kubectl apply to begin ingest

The screenshot shows the New Relic Kubernetes log viewer interface. At the top, it displays the account name 'Demotron Rotate' and a search bar with the query 'Find logs where pod_name="order-confirmation-7ddb699b6c-p46q6"'. Below the search bar is a 'Query logs' button. The main area is titled 'Find logs where pod_name="order-confirmation-7ddb699b6c-p46q6"' and shows a timeline from 05:30 AM to 06:00. The log entries are listed in a table with columns for TIMESTAMP, HOSTNAME, SERVICE_NAME, and MESSAGE. The first few log entries are:

TIMESTAMP	HOSTNAME	SERVICE_NAME	MESSAGE
11:36:21.328	ip-172-20-34-116.us....	order-confirmation	stress: FAIL: [1] (415) failed run completed in 0s
11:36:21.700	ip-172-20-34-116.us....	order-confirmation	stress: Info: [1] dispatching hogs: 0
11:36:21.700	ip-172-20-34-116.us....	order-confirmation	cpu: 0 io: 1 vm: 0 hdd

At the bottom left, there is a note: 'Visualize The Run'.

Example

What we learned

- ✓ Think + build small
- ✓ Names matter
- ✓ Health is wealth
- ✓ Ring the alarm: Requests & limits
- ✓ Context is gold

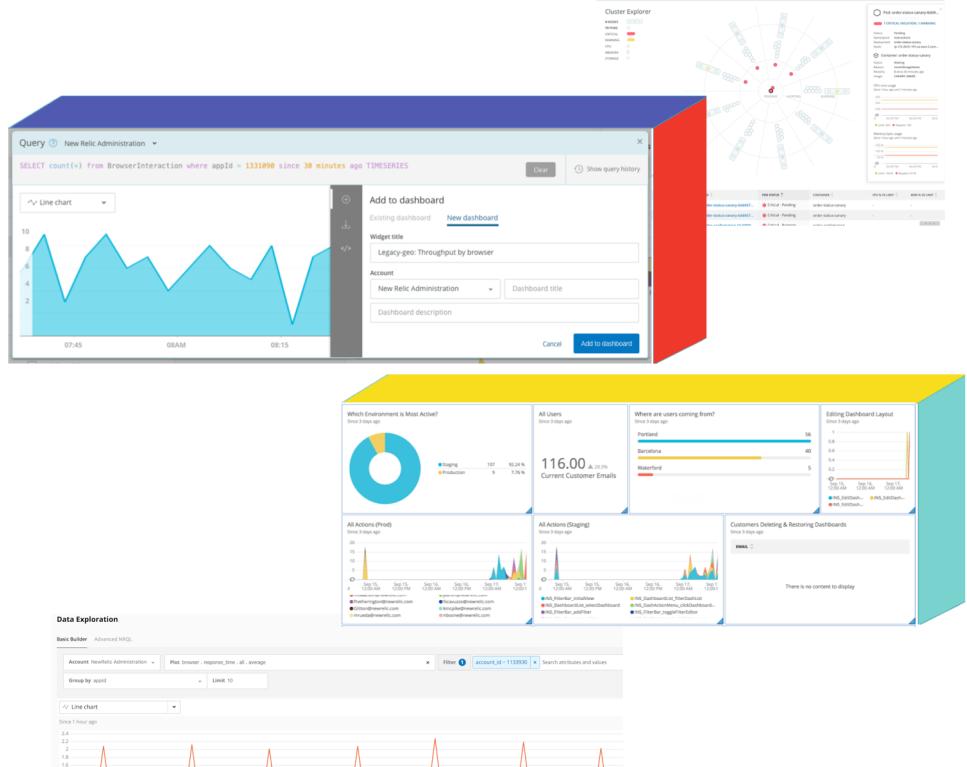




Thank You

Grace Andrews

gandrews@newrelic.com





Please complete the session
survey in the mobile app.