

The background of the image is a vibrant, multi-colored gradient. It features a diagonal split between a blue/purple area on the left and an orange/yellow area on the right. The text 'AWS re:Invent' is positioned on the left side, overlapping the blue/purple gradient. The 'AWS' part is in a small, white, sans-serif font, while 're:Invent' is in a larger, white, sans-serif font. The 're:' is smaller and positioned to the left of 'Invent'.

AWS  
re:Invent

**SVS343-R1**

# Building microservices with AWS Lambda

**Chris Munns**

Senior Manager/Principal Developer Advocate - Serverless  
Amazon Web Services

# About me

Chris Munns - [munns@amazon.com](mailto:munns@amazon.com), [@chrismunns](https://twitter.com/chrismunns)

- **Sr Manager/Principal Developer Advocate – Serverless**
- New Yorker (ehhh...ish.. kids/burbs/ya know?)

## Previously:

- AWS Business Development Manager – DevOps, July 2015 – Feb 2017
- AWS Solutions Architect Nov 2011 – Dec 2014
- Formerly on operations teams @Etsy and @Meetup
- Little time at a hedge fund, Xerox, and a few other startups
- Rochester Institute of Technology: Applied Networking and Systems Administration '05
- Internet infrastructure geek



# Why are we here today?


**containers != microservices**

**microservices != containers**

Today's focus:




AWS Lambda



microservices

# Microservices @ Amazon/AWS



Amazon S3 at launch:



8 separate  
microservices





Amazon S3 @ re:Invent 2018:

More than 235  
distributed  
microservices




# View of a single service

APIs in the front

---

Async in the back



# Common microservices patterns at AWS

## API in the front


- Synchronous communication
- HTTPS clients
- Relying on API gateways or load balancers
- Flexible client interface
- Securing against client

## Async in the back

- Asynchronous communication
- Simple clients
- Queues, topics, buses, streams
- Opinionated event/message production and consumption
- Back off and retry consumption from event source

# Today let's focus on a simple application

One public  
interface →



← One private  
internal backend  
service

# MyService architecture



Frontend

The diagram consists of a single light green rounded rectangle with the word 'Frontend' centered inside it in white text. This rectangle is positioned on the left side of a dark gray background.

# MyService architecture



The diagram illustrates the architecture of 'MyService'. It consists of two main components: a 'Frontend' represented by a green rounded rectangle on the left, and a 'Backend' represented by a blue rounded rectangle on the right. The components are positioned side-by-side on a dark gray background.

Frontend

Backend

# MyService architecture



```
graph TD; Frontend[Frontend] --- Backend[Backend]; Frontend --- Shared[Shared capabilities]; Backend --- Shared;
```

Frontend

Backend

Shared capabilities



AWS Lambda



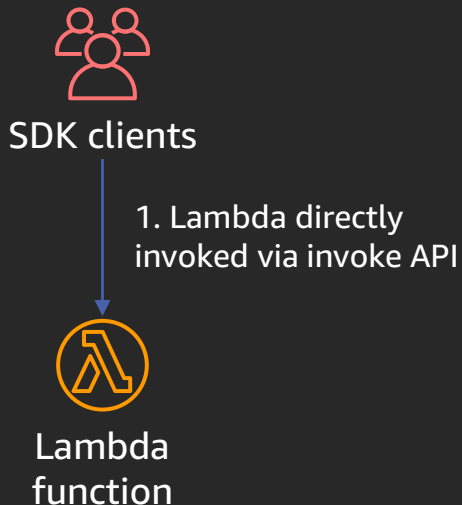
# AWS Lambda

- Serverless compute service
- Code deployed as functions aligned to an event source
- Unique characteristics:
  - 3gb RAM max, 15 minute duration max, max 250mb application artifact, max 512mb /tmp
- Security, monitoring, logging built in
  - Scoped at the function level
- Ephemeral worker environments
  - Transient in nature, could be "reaped" due to idle or other lifecycle events
  - No ability to "sticky session" customer data to a worker environment
  - Store all data that matters in a database or data store off function
- Pay per invocation and for execution duration



AWS SAM

# Lambda API



API provided by the Lambda service

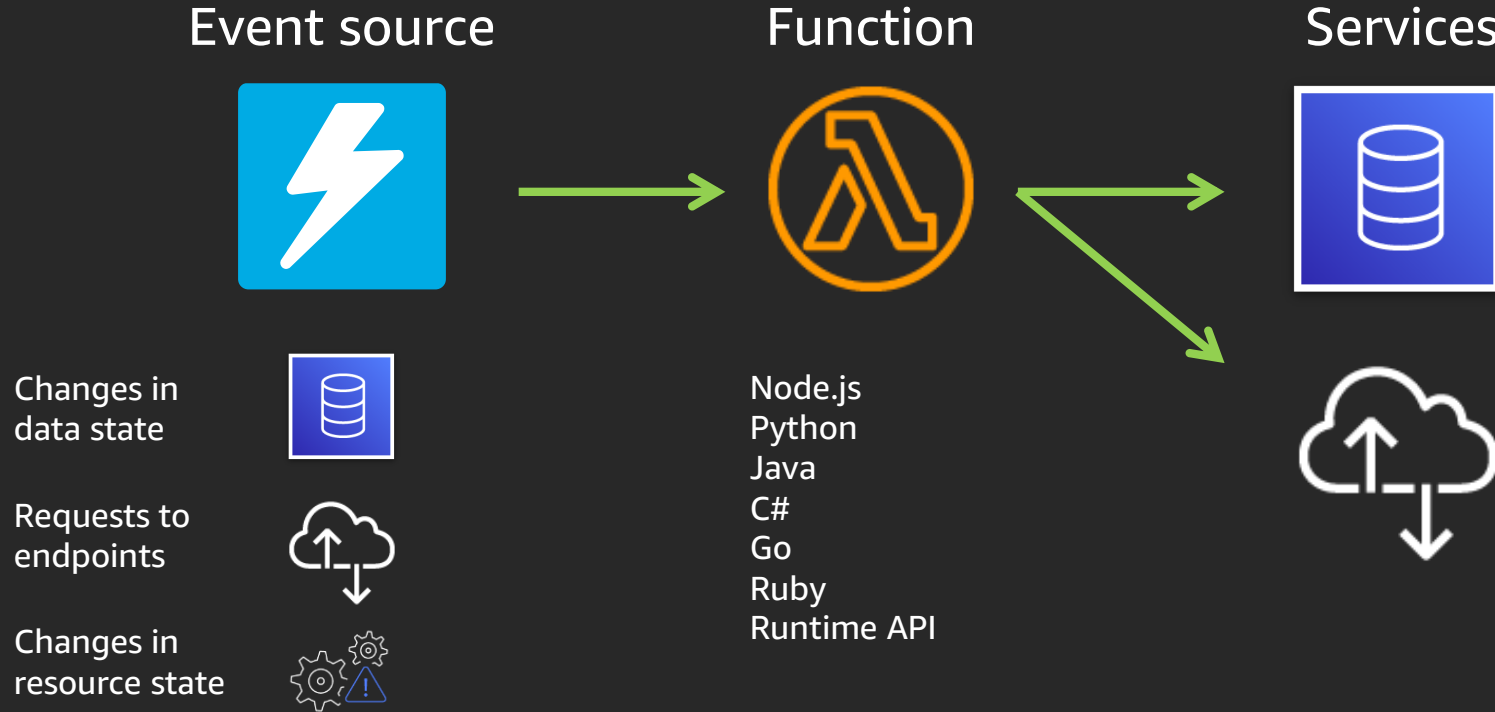
Used by all other services that invoke Lambda across all models

Supports sync and async

Can pass any event payload structure you want

Client included in every SDK

# Serverless applications



# MyService architecture

Frontend



AWS Lambda

Backend



AWS Lambda

Shared capabilities

# MyService architecture

Frontend



AWS Lambda

Backend




AWS Lambda


Shared capabilities

# Accessing the API

Three options for exposing an API:



Amazon API Gateway




Application Load Balancer  
(ALB)



AWS AppSync

# Accessing the API

## Three options for exposing an API:




Amazon API Gateway

- REST & WebSocket support
- Flexible auth options
- Throttling/usage tiers
- Caching
- Client SDK generation
- Edge, regional, private endpoint types
- OpenAPI/Swagger support
- Pay per request & data transferred

See SVS212 & SVS402 for more!

# Accessing the API

Three options for exposing an API:



Application Load Balancer


- HTTP/S support
- Path-based routing
- URL redirects
- Custom HTTP responses
- Container support
- Pay per hour & LCU consumed

See NET413 for more!



# Accessing the API

## Three options for exposing an API:



AWS AppSync

- GraphQL support
  - Can combine data from several sources in a single response
- Polyglot to backend data sources
- Deep integration with Amplify Framework
- Subscriptions & offline sync
- Pay per query and data transfer

See MOB307 & MOB402 for more!

# Picking an API fronting service cheat sheet

Complex API with multiple data sources or very unique queries against data?

- AWS AppSync

WebSockets?

- Amazon API Gateway

Need transforms, throttling, usage tiers, flexible auth?

- Amazon API Gateway

Normal API, potentially high requests per month, no need for added transform capabilities?

- Application Load Balancer

Else: Normal API, with < tens of millions of requests per month?

- Amazon API Gateway

**Warning:**  
May not be  
accurate by  
end of week.

# MyService architecture

## Frontend



Amazon API Gateway



AWS Lambda

## Backend



AWS Lambda

Shared capabilities


# Authorization

- **Open**
  - No authentication or authorization
- **AWS Identity and Access Management (IAM) permissions**
  - Use IAM policies and AWS credentials to grant access
- **Amazon Cognito authorizer**
  - Amazon Cognito is a managed user directory for authentication
  - Connect to Amazon Cognito User Pool and together with OAuth scopes to enable authorization
- **Lambda authorizers**
  - Use Lambda to validate a bearer token (OAuth or SAML as examples) or request parameters and grant access

# A quick bit on where to handle request routing

Who should handle the routing to backend logic?

api.myservice.aws/  
    /users  
    /products  
    /support  
    /search  
    /burgerandfries



Amazon API Gateway

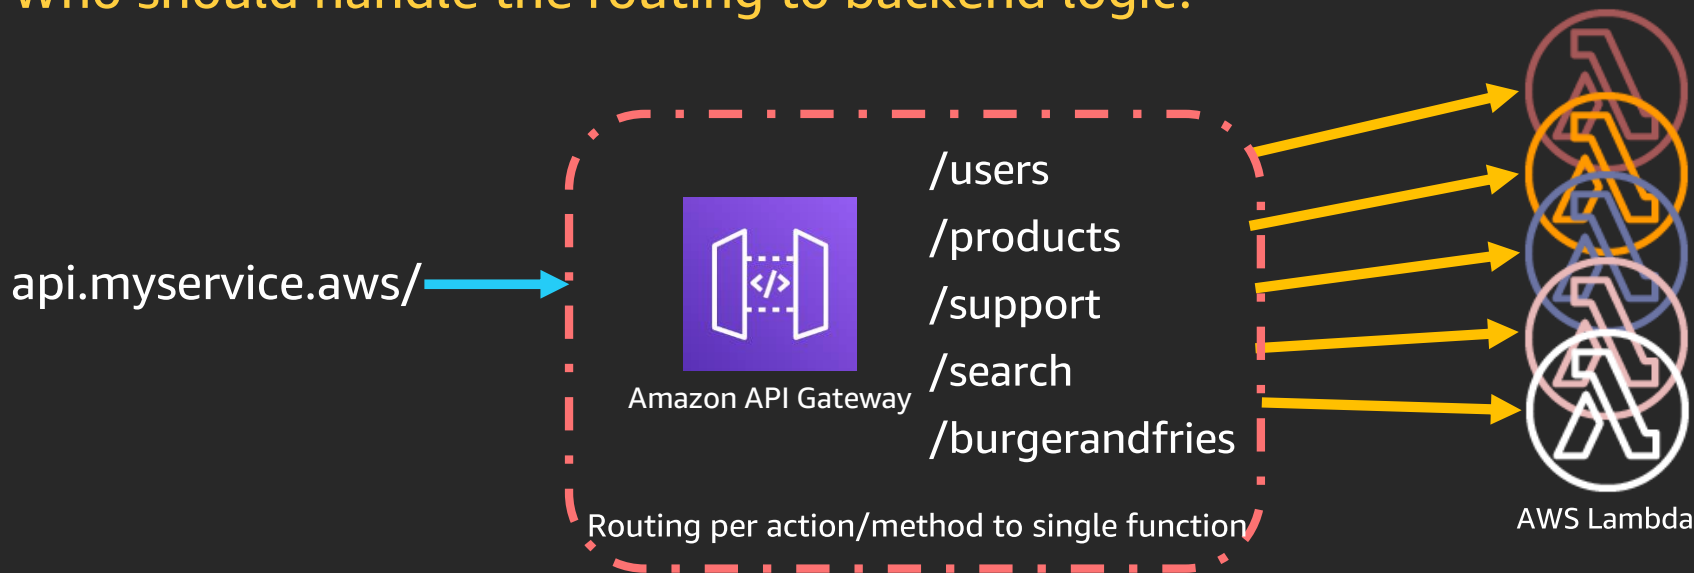


AWS Lambda

Certain frameworks will put that logic in either place. Does it matter?

# A quick bit on where to handle request routing


Who should handle the routing to backend logic?



Certain frameworks will put that logic in either place. Does it matter?

# A quick bit on where to handle request routing

Who should handle the routing to backend logic?



Certain frameworks will put that logic in either place. Does it matter?

# Lambda + in function routing model

The characteristics we mentioned earlier are what put customers into the “danger zone” with the function routing model:

- Security constructs applied to the whole
- Performance settings applied to the whole
- Limited amount of application size
- Limited duration
  - Might be better to chain logic out to other functions with AWS Step Functions or via the async model
- The “Lambda-lith” can grow too complex, leaving you stuck rethinking the whole logic routing model




AWS Lambda



# Where to handle request routing

## Team API Gateway

- Using the benefits of API Gateway
- Better compatibility with AWS's tools
- Better security granularity
- Better performance granularity



Amazon API Gateway

## Team in code

- Native code/framework capabilities
- Potentially better portability of code
- Fewer security constructs (IAM roles, policies, etc.)
- Fewer cold starts? (questionable based on Yan Cui's blog\*)



AWS Lambda

\*<https://hackernoon.com/aws-lambda-should-you-have-few-monolithic-functions-or-many-single-purposed-functions-8c3872d4338f>

# Where to handle request routing

Team API Gateway

Team in code

- Using the benefits of API Gateway

- Native code/framework capabilities

- Better compatibility with AWS's tools

- Potentially better portability of code

- Better security granularity

- Fewer security constructions (IAM Roles, Policies, etc.)

- Better performance granularity

Honestly, it's a toss up. Customer feedback seems very split.

If you love your framework, stick with it and its capabilities.

(on Yan Qui's blog\*)

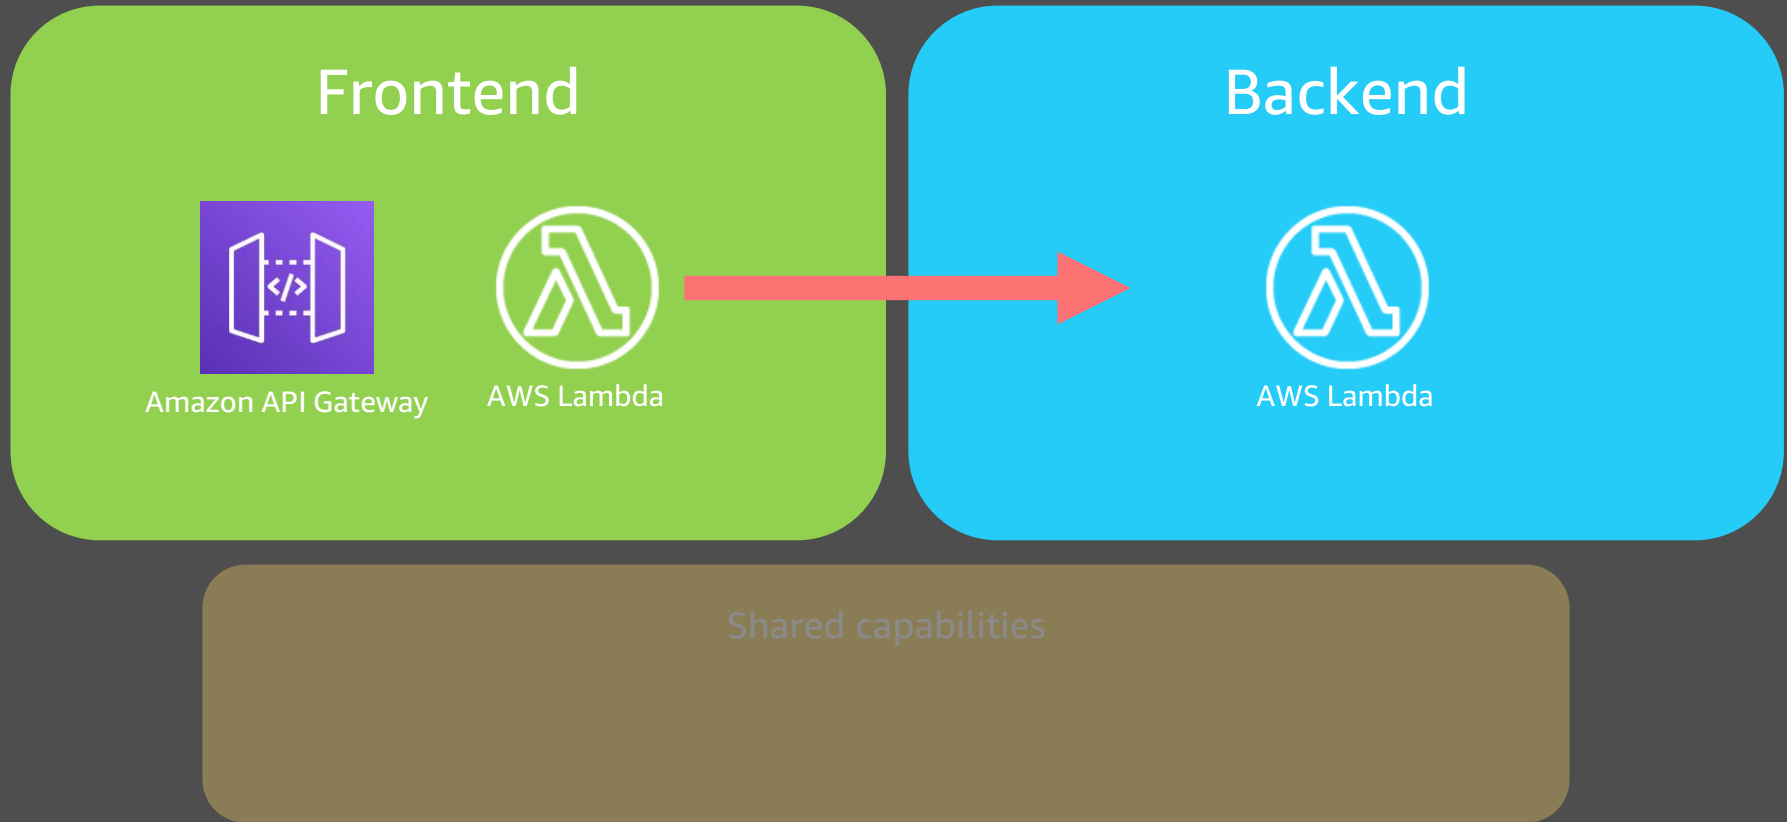
If you want to do less and write less code, go with the benefits of the managed services.

Amazon API Gateway


AWS Lambda

# Topics, streams, queues, and buses! Oh my!


# MyService architecture





# MyService architecture




# Synchronous APIs



# Synchronous APIs




# Synchronous APIs





# Synchronous APIs



# Synchronous APIs




The diagram illustrates a synchronous API interaction. On the left, a 'Client' icon is shown. A blue arrow labeled 'POST /order' points from the Client to a green box labeled 'Order'. From the 'Order' box, a blue arrow labeled 'POST /invoice' points to a blue box labeled 'Invoice'. Below the 'Order' box, a green arrow points back to the Client, and below the 'Invoice' box, a green arrow points back to the 'Order' box. Three blue starburst shapes, representing errors or failures, are placed on the return paths: one on the arrow from the Client to the Order box, one on the arrow from the Invoice box to the Order box, and one on the arrow from the Order box to the Client.


Who own's the retry? For how long?  
Does the client ever know? Etc..

This effectively creates a “tight coupling”  
where failures become harder to recover from

# Thinking asynchronously




Synchronous  
commands



Asynchronous  
events

# Asynchronous APIs




"The time spent to try making a process async will pay for itself in you gaining deeper understanding of what is really happening with your data."


**-Me**

Right Now, 2019


# Topics, streams, queues, and buses! Oh my!



Amazon Simple  
Notification Service  
(Amazon SNS)



Amazon Simple Queue  
Service (Amazon SQS)



Amazon EventBridge




Amazon Kinesis  
Data Streams

# Ways to compare




Scale/concurrency  
controls



Durability



Persistence



Consumption  
models



Retries



Pricing

# Recent announcements for async event sources

- SQS FIFO as invoke source for Lambda
- SNS Dead Letter Queues (DLQs)
- Lambda Destinations:
  - Capture success/failure from functions
- For streamed events:
  - MaximumRetryAttempts, MaximumRecordAgeInSeconds, BisectBatchOnFunctionError, On-failure destination
  - Batch Window
  - Parallelization Factor
- For async events:
  - MaximumRetryAttempts
  - MaximumEventAgeInSeconds



ICYMI: Serverless  
pre:Invent 2019



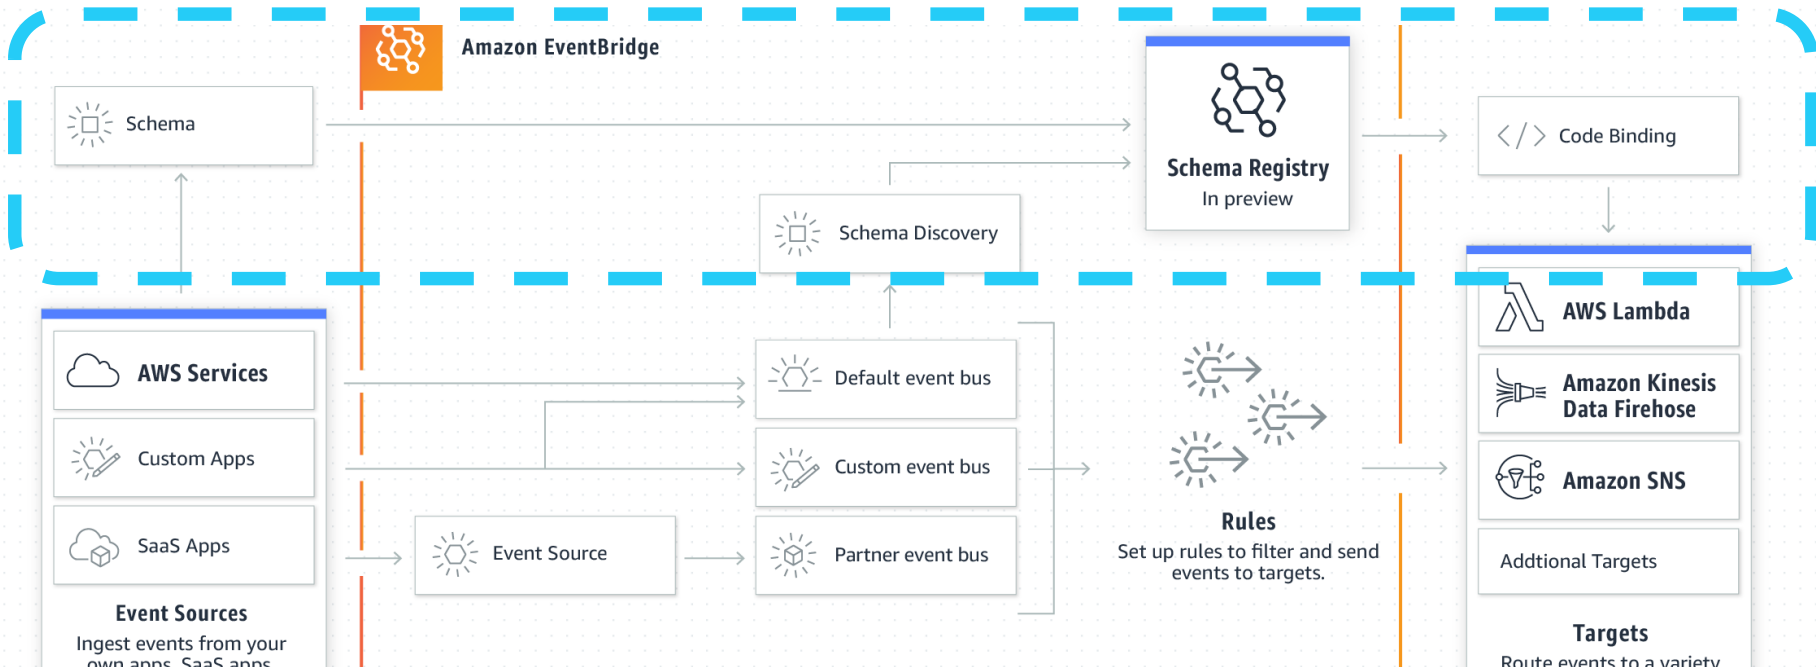
Last night!

# Announcements for async event sources

AWS Compute Blog

## Introducing Amazon EventBridge schema registry and discovery – In preview

by Julian Wood | on 01 DEC 2019 | in [Amazon EventBridge](#), [Serverless](#) | [Permalink](#) | [Comments](#) | [Share](#)



“Woah”

**-Me**

Right Now, 2019

# Async service decider cheat sheet

Massive throughput/ordering/multiple consumers/replay?

- Kinesis Data Streams

One to mostly one or minimal fanout, direct to Lambda/HTTP target?

- Amazon SNS

Buffer requests until they can be consumed, whether ordered or not?

- Amazon SQS

One to many fanout, lots of different consumer targets, schema matching, granular target rules?

- EventBridge

# Async service decider cheat sheet

Massive throughput/ordering/multiple consumers/replay?

- Kinesis Data Streams

This is a big topic itself with a lot of things to consider.  
Evaluate along the 6 criteria mentioned previously.

- SNS

EventBridge changes so much of this because it can target  
to all of these as well. ::mind blown gif::


- SQS

One to many fanout, lots of different consumer targets, schema

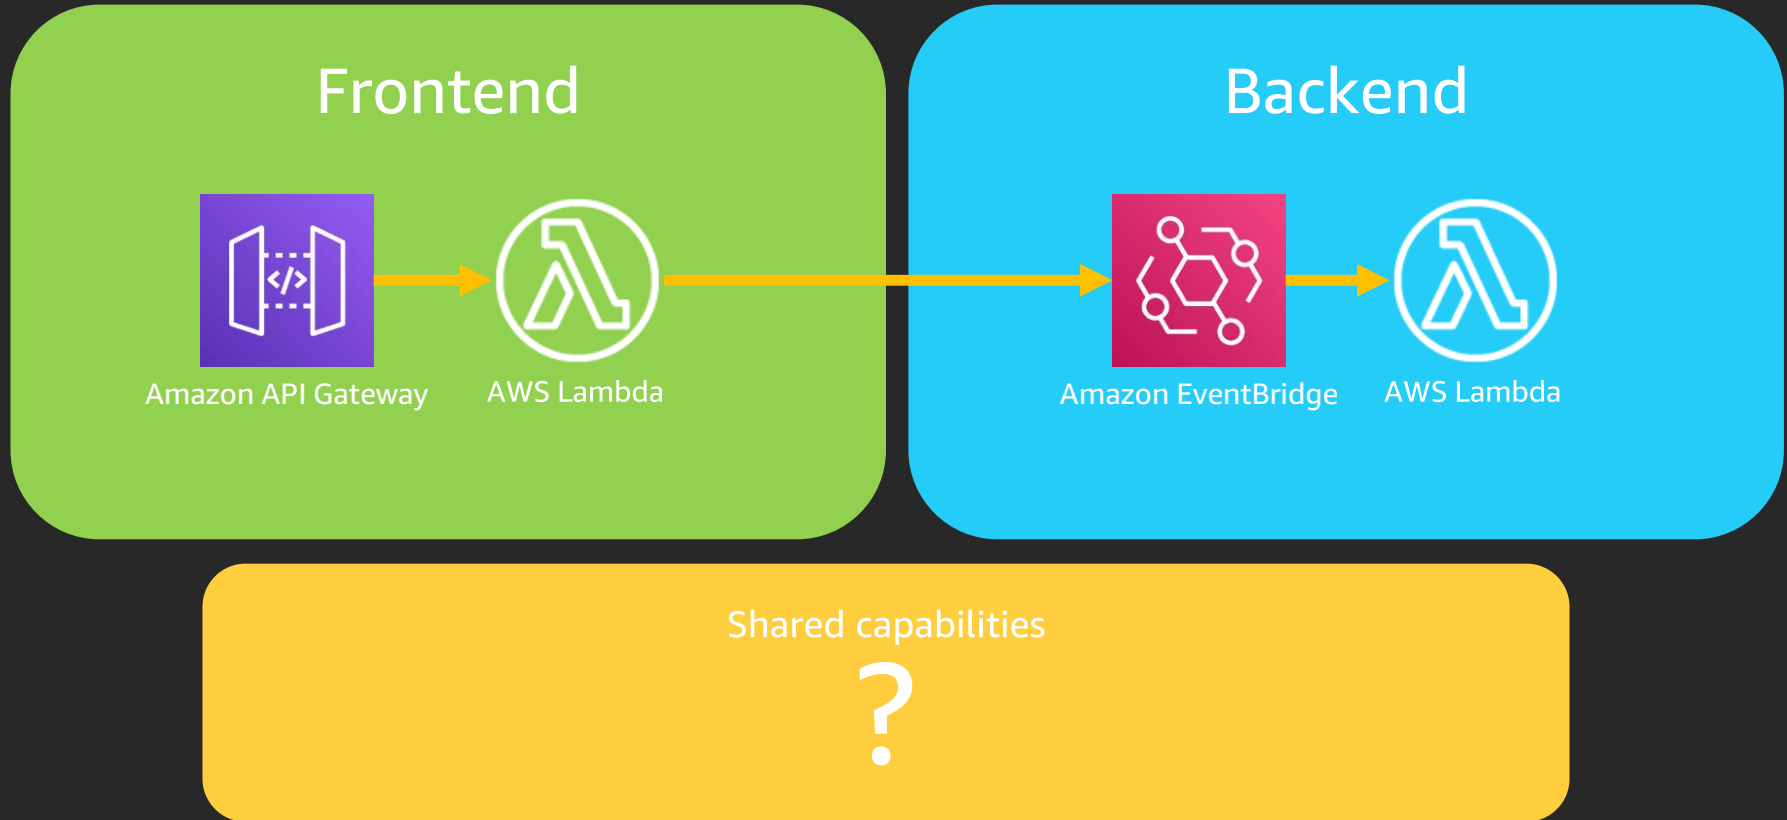
See API304, API315, API320, SVS308, SVS317 for more!

- EventBridge

# MyService architecture



# MyService architecture



# Shared capabilities

# There are still a lot of things to think about

- Secrets/configuration management
- Simplifying code management
- Debugging/troubleshooting
- Performance controls
- Security



AWS Serverless  
Application Model  
(AWS SAM)



# AWS Lambda environment variables

- Key-value pairs that you can dynamically pass to your function
- Available via standard environment variable APIs such as `process.env` for Node.js or `os.environ` for Python
- Can optionally be encrypted via AWS Key Management Service (AWS KMS)
  - Allows you to specify in IAM what roles have access to the keys to decrypt the information
- Useful for creating environments per stage (i.e., dev, testing, production)



# AWS Systems Manager – Parameter Store

## Centralized store to manage your configuration data

- Supports hierarchies
- Plaintext or encrypted with AWS KMS
- Can send notifications of changes to Amazon SNS/AWS Lambda
- Can be secured with IAM
- Calls recorded in AWS CloudTrail
- Can be tagged
- Integrated with AWS Secrets Manager
- Available via API/SDK

Useful for centralized environment variables, secrets control, feature flags

```
from __future__ import print_function
import json
import boto3
ssm = boto3.client('ssm', 'us-east-1')

def get_parameters():
    response = ssm.get_parameters(
        Names=['LambdaSecureString'], withDecryption=True
    )
    for parameter in response['Parameters']:
        return parameter['Value']

def lambda_handler(event, context):
    value = get_parameters()
    print("value1 = " + value)
    return value  # Echo back the first value
```

```
Import sdk
Import http-lib
Import subFunctionA
```

## Dependencies, configuration information, common helper functions

```
Pre-handler-secret-getter()
Pre-handler-db-connect()
```

```
Function myhandler(event, context) {
  <Event handling logic> {
    result = SubfunctionA()
  }else {
    result = SubfunctionB()

  }

  return result;
}
```

## Your handler

```
Function Pre-handler-secret-getter() {
}
```

## Common helper functions

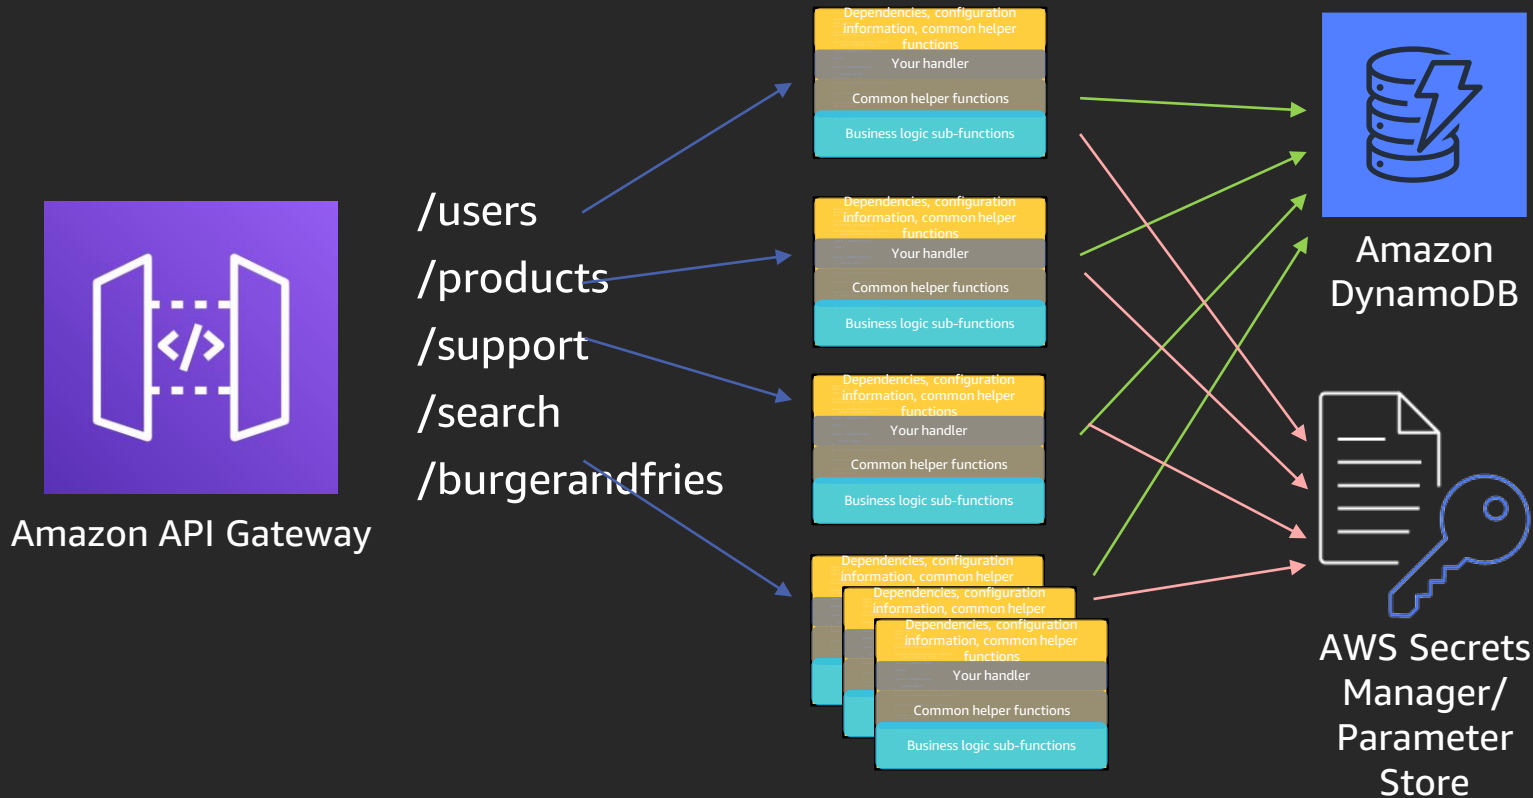
```
Function Pre-handler-db-connect(){
}
```

```
Function subFunctionA(thing){
  ## logic here
}
```


## Business logic sub-functions

```
Function subFunctionB(thing){
  ## logic here
}
```

# Anatomy of a serverless application



# Anatomy of a serverless application



Amazon API Gateway

/users





/products

/support


/search

/burgerandfries

There could be a lot  
of duplicated code  
here!



Amazon  
DynamoDB



AWS Secrets  
Manager/  
Parameter  
Store

# Lambda Layers



Lets functions easily share code: Upload layer once, reference within any function

Layer can be anything: dependencies, training data, configuration files, etc

Promote separation of responsibilities, lets developers iterate faster on writing business logic

Built-in support for secure sharing by ecosystem

# Using Lambda Layers

- Put common components in a ZIP file and upload it as a Lambda Layer
- Layers are immutable and can be versioned to manage updates
- When a version is deleted or permissions to use it are revoked, functions that used it previously will continue to work, but you won't be able to create new ones
- You can reference up to five layers, one of which can optionally be a custom runtime



Lambda  
Layers

arn:aws:lambda:region:accountId:layer:shared-lib:1



Lambda  
Layers

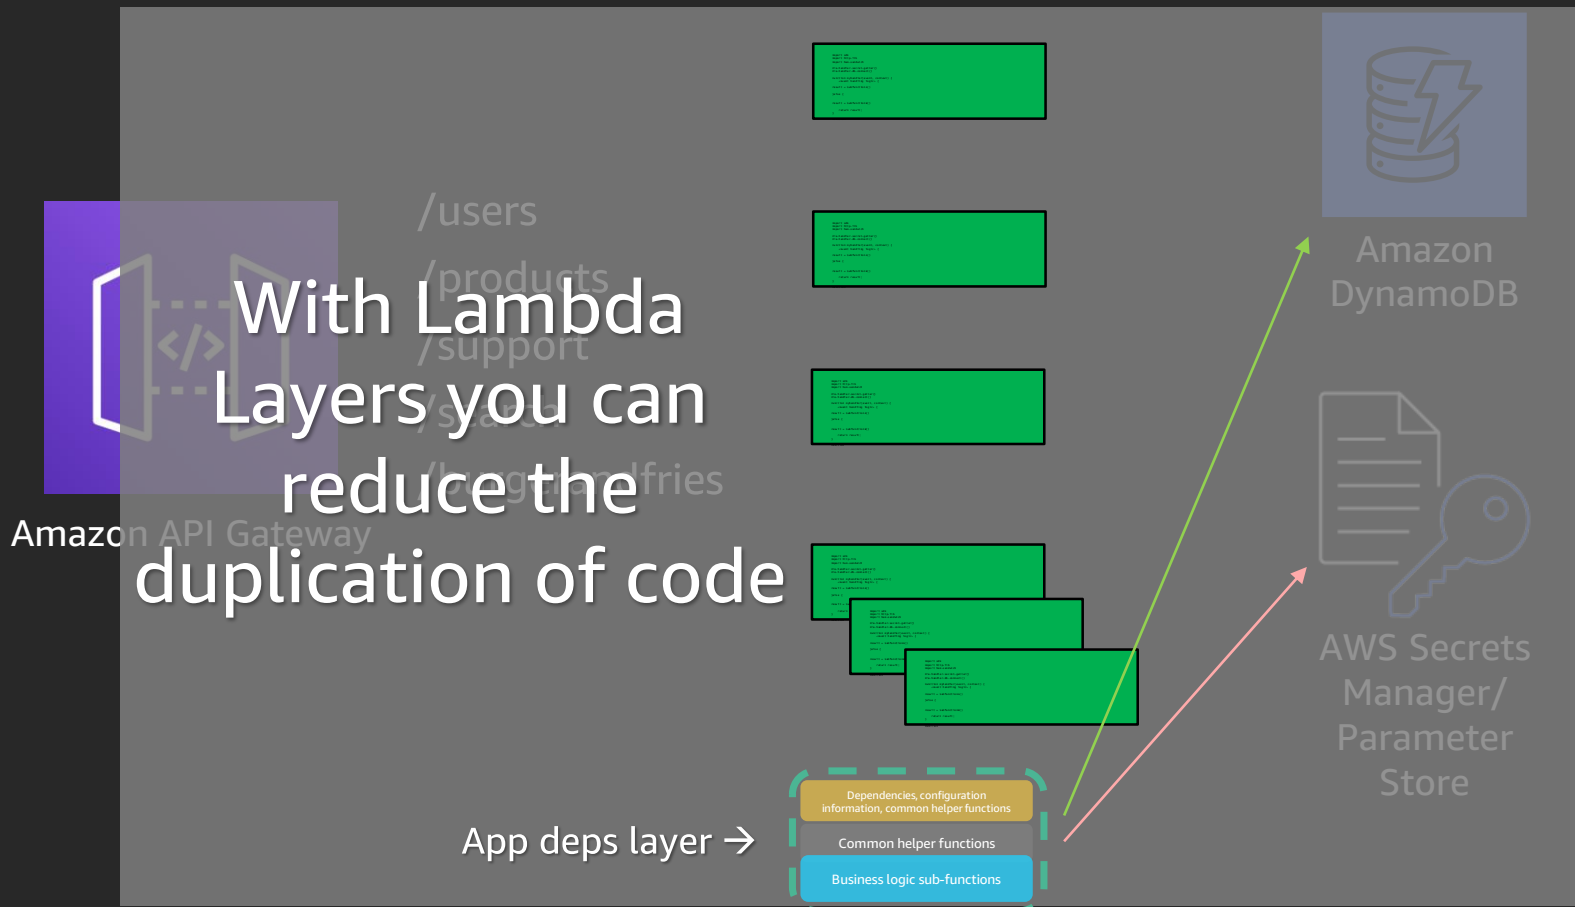
arn:aws:lambda:region:accountId:layer:shared-lib:2



Lambda  
Layers

arn:aws:lambda:region:accountId:layer:shared-lib:3

# Anatomy of a serverless application






# AWS X-Ray

Profile and troubleshoot serverless applications:

- Lambda instruments incoming requests for all supported languages and can capture calls made in code
- API Gateway inserts a tracing header into HTTP calls as well as reports data back to X-Ray itself

```
var AWSXRay = require('aws-xray-sdk-core');  
var AWS = AWSXRay.captureAWS(require('aws-sdk'));  
S3Client = AWS.S3();
```

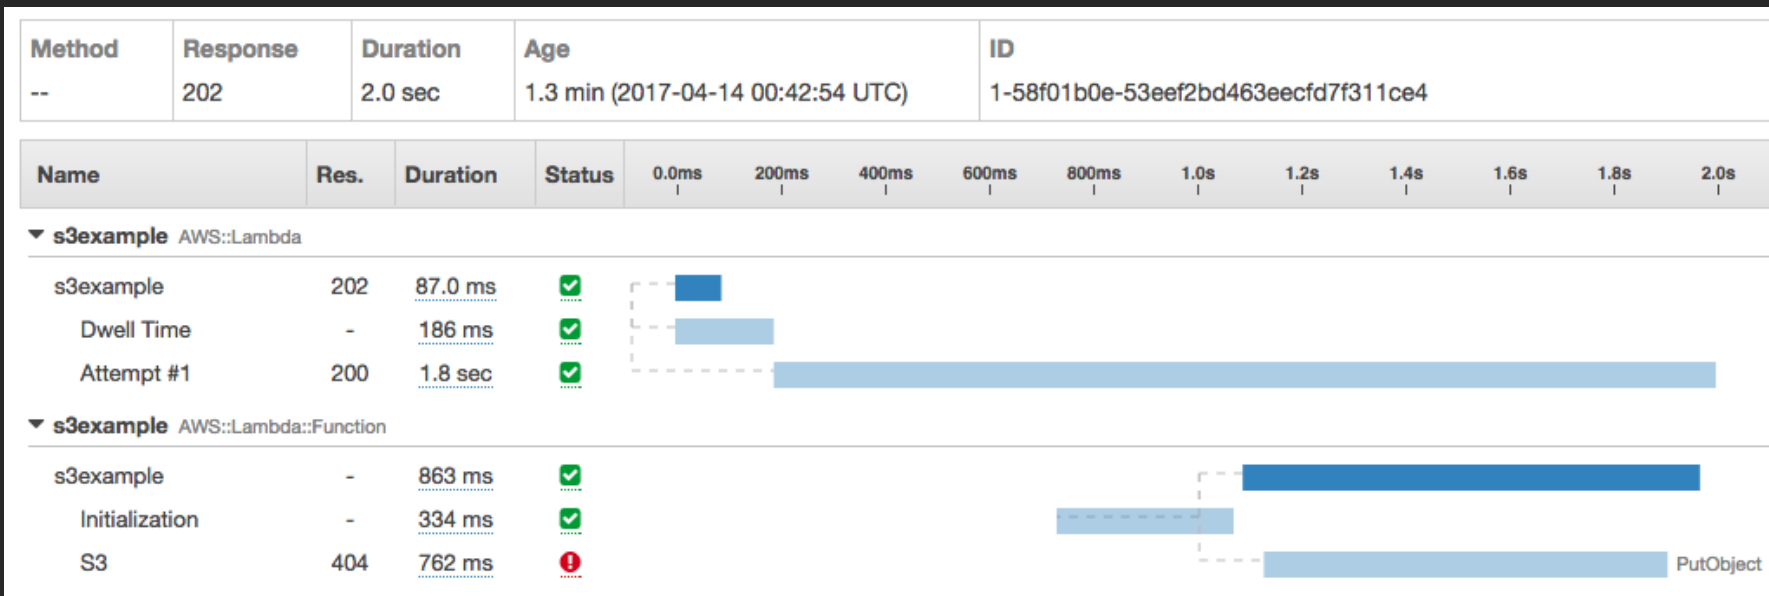


Enable X-Ray Tracing ☒

Enable active tracing [Info](#)



# X-Ray trace example



# Recent announcements for X-Ray & Amazon CloudWatch


- **CloudWatch ServiceLens**
  - A “single pane of glass” monitoring tool
- **CloudWatch Synthetics**
  - Generate “canary” tests for your APIs and other services
- **Embedded Metric Format for CloudWatch Logs**
- **CloudWatch Contributor Insights**
  - Identify top talkers, common patterns in logs
- **X-Ray Trace Maps**
  - Map a single event through your distributed application
- **X-Ray integration with CloudWatch Synthetics**



ICYMI: Serverless  
pre:Invent 2019

# Lambda per function concurrency controls

- Concurrency is a shared pool by default
- Separate using per function concurrency settings
  - Acts as reservation
- Also acts as max concurrency per function
  - Especially critical for downstream resources like databases
- “Kill switch” – set per function concurrency to zero



**Friends don't let friends  
"Action": "s3:\*"**

# Lambda permissions model

- **Function policies:**

- Example: "Actions on bucket X can invoke Lambda function Z"
- Resource policies allow for cross account access
- Used for sync and async invocations

- **Execution role:**

- Example: "Lambda function A can read from DynamoDB table users"
- Define what AWS resources/API calls this function can access via IAM
- Used in streaming invocations



# AWS Serverless Application Model (AWS SAM)



AWS CloudFormation extension optimized for serverless

Special serverless resource types: functions, APIs, tables, layers, and applications

Supports anything AWS CloudFormation supports

Open specification (Apache 2.0)

<https://aws.amazon.com/serverless/sam>

# AWS SAM template

AWSTemplateFormatVersion: '2010-09-09'

Transform: AWS::Serverless-2016-10-31

Resources:

GetProductsFunction:

Type: AWS::Serverless::Function

Properties:

Handler: index.getProducts

Runtime: nodejs8.10

CodeUri: src/

Policies:

- DynamoDBReadPolicy:

TableName: !Ref ProductTable

Events:

GetResource:

Type: Api

Properties:

Path: /products/{productId}


Method: get

ProductTable:

Type: AWS::Serverless::SimpleTable

Just 20 lines to create:

- Lambda function
- IAM role
- API Gateway
- DynamoDB table





# AWS SAM policy templates

GetProductsFunction:

Type: `AWS::Serverless::Function`

Properties:

...

Policies:

- DynamoDBReadPolicy:

TableName: `!Ref ProductTable`

...

ProductTable:

Type: `AWS::Serverless::SimpleTable`



```
3  "Templates": {
4    "SQSPollerPolicy": {
5      "Description": "Gives permissions to poll an SQS Queue",
6      "Parameters": {
7        "QueueName": {
8          "Description": "Name of the SQS Queue"
9        }
10     },
11     "Definition": {
12       "Statement": [
13         {
14           "Effect": "Allow",
15           "Action": [
16             "sqs:ChangeMessageVisibility",
17             "sqs:ChangeMessageVisibilityBatch",
18             "sqs:DeleteMessage",
19             "sqs:DeleteMessageBatch",
20             "sqs:GetQueueAttributes",
21             "sqs:ReceiveMessage"
22           ],
23           "Resource": {
24             "Fn::Sub": [
25               "arn:${AWS::Partition}:sqs:${AWS::Region}:${AWS::AccountId}:${queueName}",
26               {
27                 "queueName": {
28                   "Ref": "QueueName"
29                 }
30             ]
31           }
32         }
33       ]
34     }
35   }
36 }
```

50+ predefined  
policies  
all found here:  
<https://bit.ly/2xWycnj>

# AWS SAM Command Line Interface (AWS SAM CLI)



CLI tool for local development, debugging, testing, deploying, and monitoring of serverless applications

Supports API Gateway “proxy-style” and Lambda service API testing

Response object and function logs available on your local machine

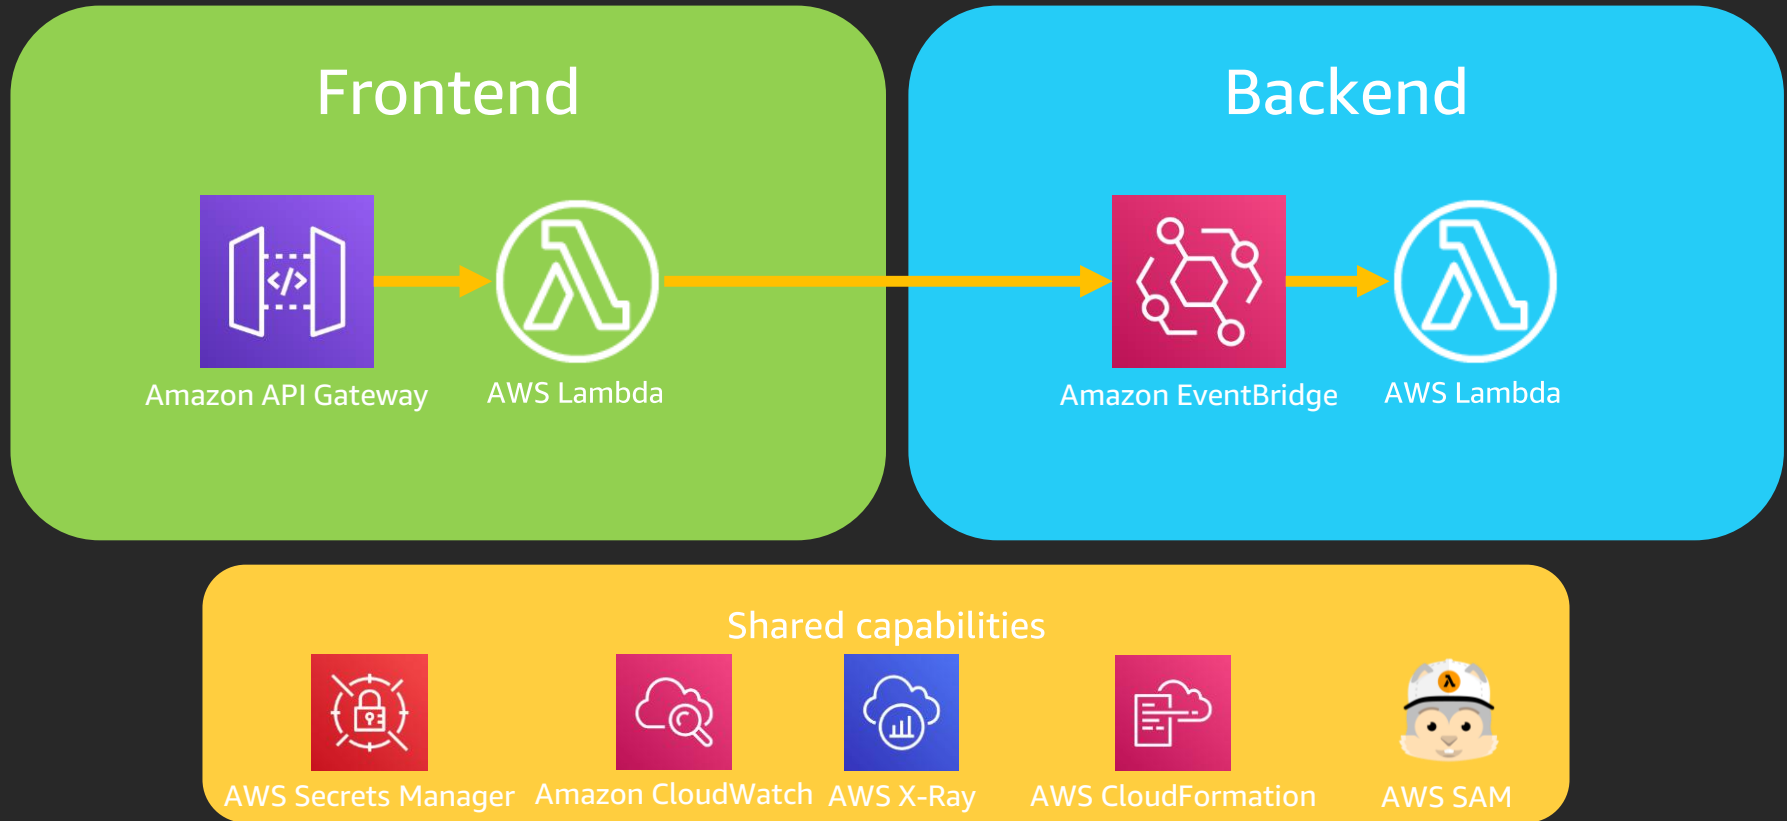
Uses open source docker-lambda images to mimic Lambda’s execution environment such as timeout, memory limits, runtimes

Can tail production logs from CloudWatch Logs

Can help you build in native dependencies

<https://aws.amazon.com/serverless/sam>

# MyService architecture



# FIN/ACK – Building microservices with AWS Lambda

The characteristics of Lambda make you think differently about how to build microservices:

- **Lambda's service limits impact your architectural decisions**
  - Store state elsewhere
- **3 ways to do APIs: Amazon API Gateway, ALB, AWS AppSync**
- **4 stream/async methods: Amazon SNS, Amazon SQS, EventBridge, Kinesis Data Streams**
- **Ecosystem of tools simplify a lot of things:**
  - X-Ray
  - CloudWatch metrics & logs
  - AWS SAM + AWS CloudFormation
  - Secrets Manager + Parameter Store


# Learn Serverless with AWS Training and Certification

Resources created by the experts at AWS to help you learn modern application development



## Free, on-demand courses on serverless, including:

- Introduction to Serverless Development
- Getting in the Serverless Mindset
- AWS Lambda Foundations
- Amazon API Gateway for Serverless Applications
- Amazon DynamoDB for Serverless Architectures



Additional digital and classroom trainings cover modern application development and computing

Visit the Learning Library at <https://aws.training>

# Thank you!

**Chris Munns**

[munns@amazon.com](mailto:munns@amazon.com)

[@chrismunns](#)



Please complete the session  
survey in the mobile app.