



AWS
re:Invent

STG 202

What's new in AWS file storage

Mark Bixler

Sr. Cloud Engineer
MINDBODY



Amazon FSx w/ DFS & IaC

...a presentation on acronyms!

FI+ME+RIX
by  MINDBODY®

Life Before Amazon FSx – Existing Environment



Fitmetrix was an acquisition by MINDBODY



Software that provides real-time statistical data regarding your workouts

Heart Rate

Cadence

Zones

etc...



- Windows 2012 R2
- IIS 8.5
- Static Servers
- DFS on Each Box



- MSFT Web Deploy
- n x Storage Costs
- n x Backups to Manage!
- n x Web Configs!





- Limited Ability to Scale
- Errors with code pushes
- DFS Not in Sync
- Developer Credentials



- 24–72 hours to rebuild machine
- Deployments 8–10 Weeks

Life Before Amazon FSx Volumes – Rapid Aging



Life After Amazon FSx New Environment w/ IaC!

- Flexibility to Implement IaC
- Single Volume to deploy code
- One backup to manage
- Shared IIS Web Config
- Auto Scaling groups and SSM
- Updated Images w/ Patching & IIS
- New Server Spin up Time – 7 Minutes!
- Scale up AND Out – Saves Money
- Deployments every 2 weeks!



pulumi

Azure DevOps

WAIT!

What about HA?



We experienced a service outage during a FSx Maintenance Window



FSx at the time our implementation was NOT HA



Quickly put on our Engineering Hats



Designed IaC with DFS Servers

2 FSx Volumes in different AZs
2 DFS Hosts across different AZs



No More Downtime!



Rumor mill...FSx HA coming soon?



 MINDBODY®
Storage Migration

CODENAME: ASPEN

Moving to the Cloud!

...where the beer flows like wine

Existing Environment



100+ Million Files!

- EMC Isilon (48 TB)
- Replicating to Different DC with AWS Backup Isilon
- Shutting down DC
 - What do we do now?
- Not Your Typical 48 TB
 - OVER ~100 MILLION FILES
- Many traditional approaches didn't work
 - Stub Files bigger than actual files!

Migration & Solution



AWS Snowball Edge



W/ Help of AWS Support



Batched up Files ~70,000 Subdirectories



File Partition lists of 10k / SubDir using fpart



Tar all partition lists



Looped through tar files doing 10 parallel process using GNU Parallel



Sequential transfer using straight `<aws s3 cp>` 3+ years!!



28 Days from submitting the Snowball Job in AWS until the Files were in the S3 Bucket

Post Migration: AWS DataSync



Setup AWS DataSync to keep files in Sync between on-premises array and Amazon S3 bucket



Using Various combinations of Include & Exclude Filters to separate folders for processing




Broke our ~70K sub directories down alphabetically



Run 28 Agents to handle the workload



Used IaC to build AWS Lambda & Amazon CloudWatch Events to invoke the various 28 agents on staggered scheduling



Amazon Storage Gateway endpoint to do any restoring from Amazon S3

Perfectly staged inoffensive group hug



Thank you!

Mark Bixler

mark.bixler@mindbodyonline.com
@markbixler.com



Please complete the session
survey in the mobile app.