
AWS
re:Invent
NOV. 28 – DEC. 2, 2022 | LAS VEGAS, NV

ANT205-R

Achieving your modern data architecture

Santosh Chandrachood

General Manager
AWS Glue

Agenda

Modern data architecture on AWS

End-to-end data life cycle on the modern data architecture

Data governance and data mesh in action

Journey towards modern data architecture

Deriving insights from data is hard

AWS CAN HELP

Data silos

People silos

Business silos

Challenges in data silos

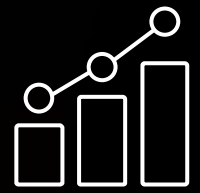

Data is growing exponentially

From new sources

Increasingly diverse

More users needing secure access

Cost and performance

Challenges in people silos

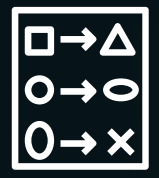

UNIQUE USER SKILLSETS

TOOL PREFERENCES

REQUIRED PROCESSES

Challenges in business silos

COST

LEGACY INFRASTRUCTURE

AGILITY

The five pillars of a modern data architecture

Unified analytics

Highest performance at the lowest cost

Machine learning integration

Unified data access, security and governance

Insights for everyone

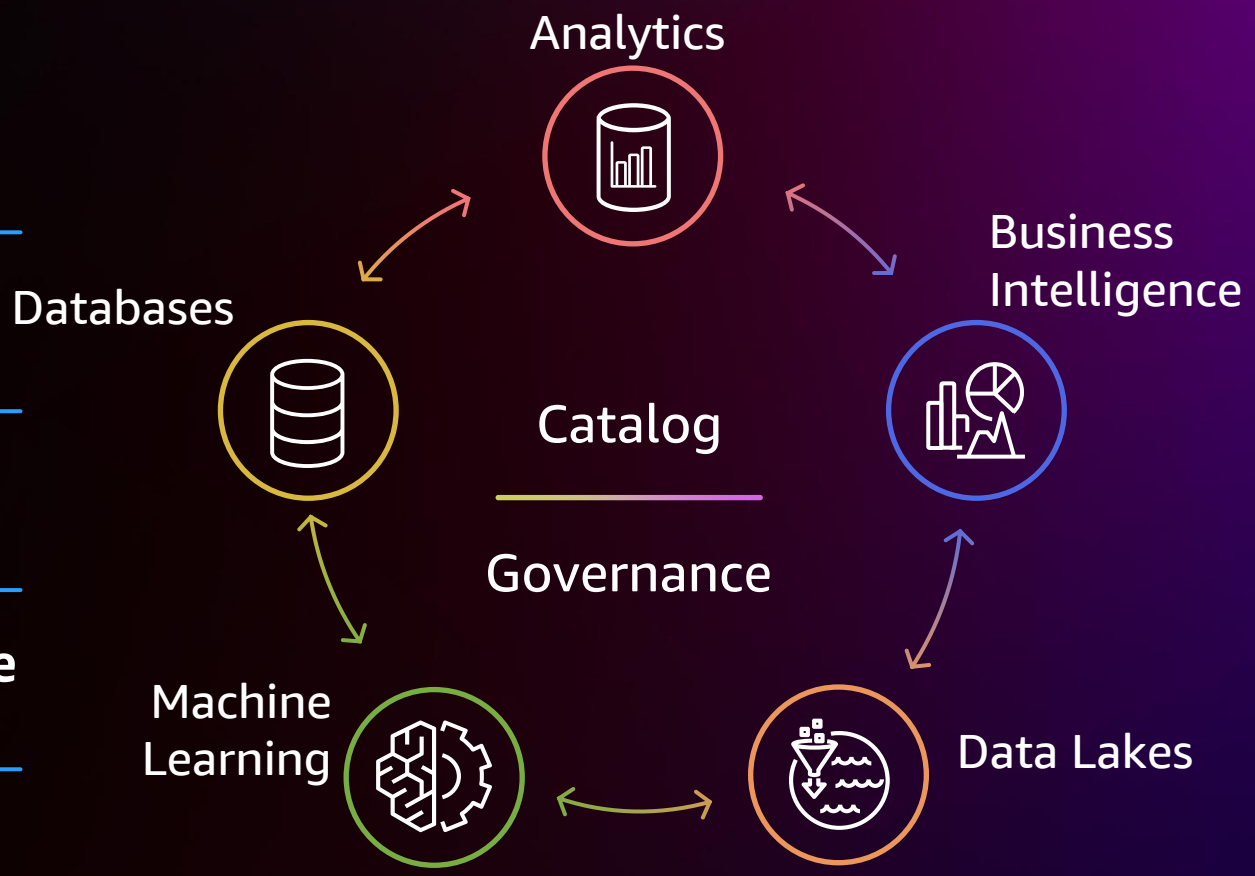

Modern data architecture

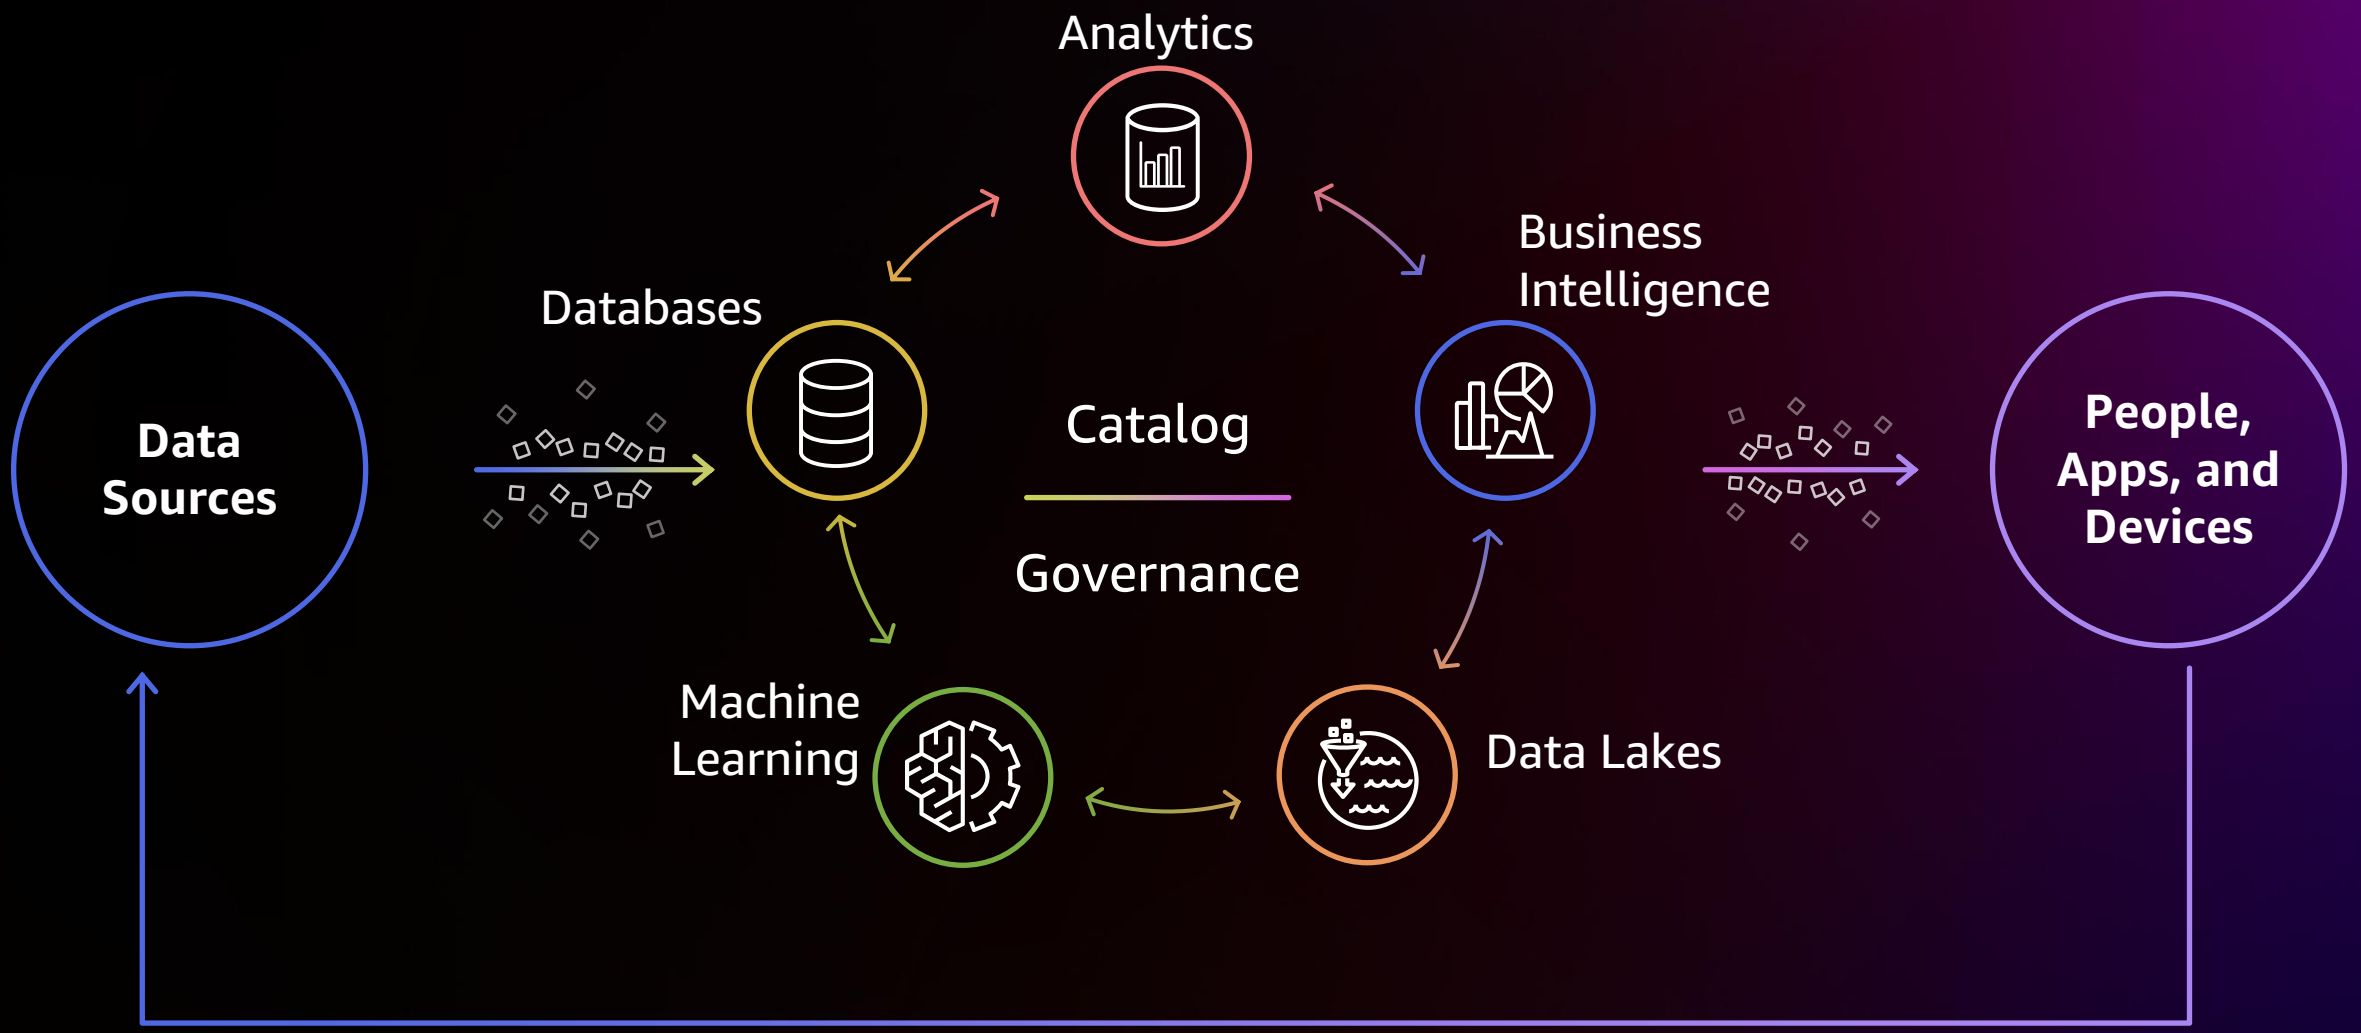

Warner Bros. Games scales data analytics

Challenge:

Measure the business and provide meaningful feedback without getting in the way of the creative process. The business isn't static. They have data velocity, volume, variety, and voracity. They need a scalable infrastructure, data federation and democracy, and a way to act on the data.

Solution:

A modern data architecture with Amazon Redshift, Amazon EMR, and AWS Glue.

Result:

One version of the data for all stakeholders to access with increased scalability, lowered overhead costs, and more compute and memory
for the same cost.

Amazon Redshift AWS Glue Amazon EMR

aws

Modern data architecture's five benefits

Unified analytics

Highest performance at the lowest cost

Machine learning integration

Unified data access, security and governance

Insights for everyone

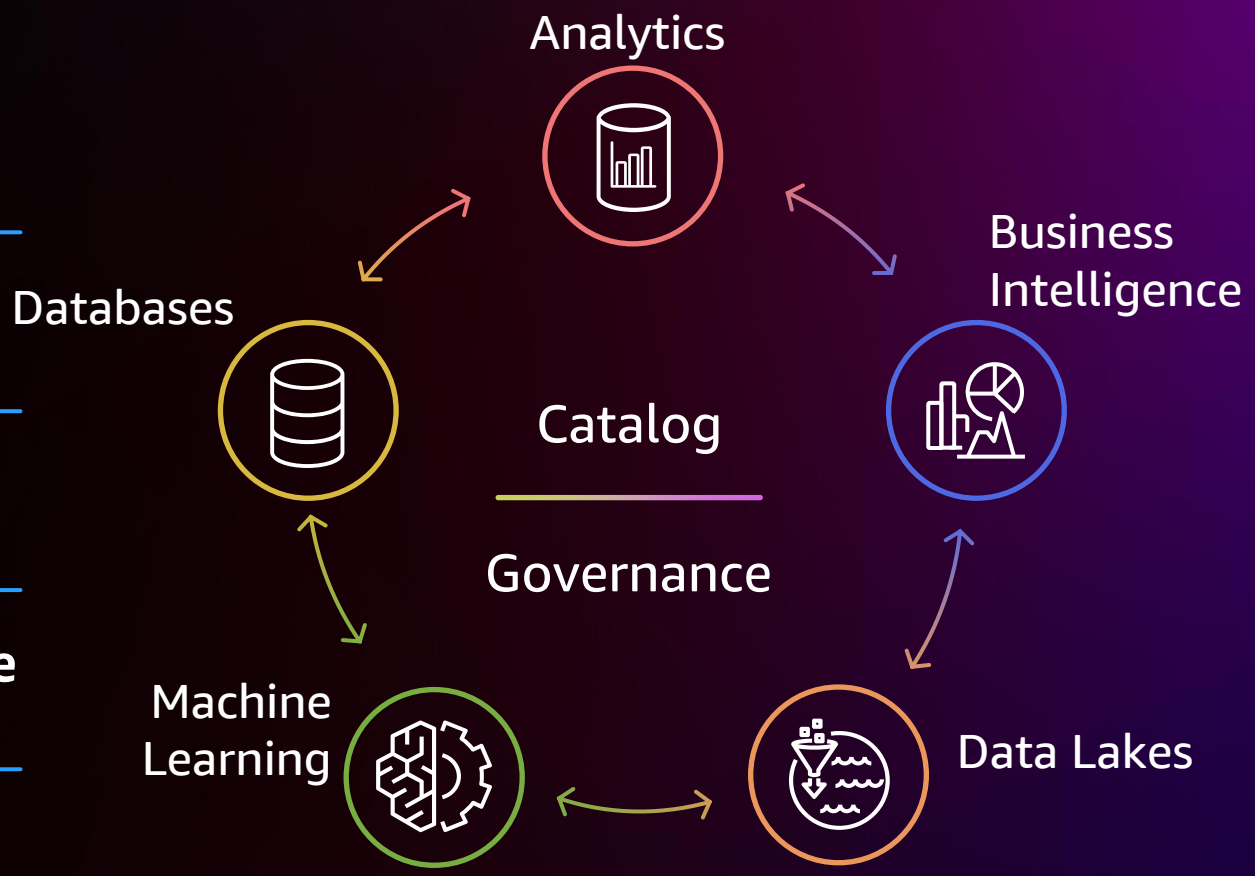

Unified analytics

Data access anywhere

Bringing your data in

Right tools for analytics

Federated query in Amazon Redshift and Amazon Athena

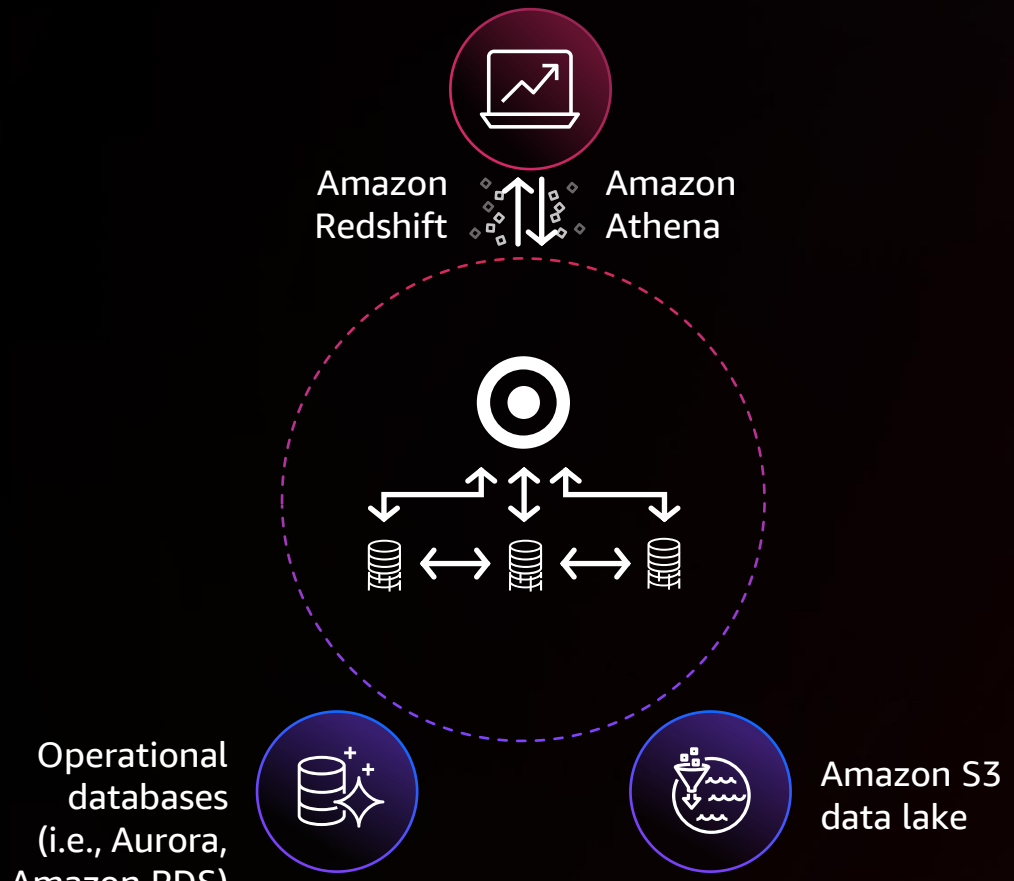

Integrate varied data stores with data warehouse and Amazon S3 data lake

Analytics on combined data **without data movement** and ETL delays

Flexible and easy way to ingest data, avoiding complex ETL pipelines

NEW [PREVIEW]

Amazon Aurora Zero-ETL to Amazon Redshift

Eliminates the need to build and maintain complex ETL pipelines

Run near-real-time analytics and machine learning using Amazon Redshift on petabytes of transactional data from Amazon Aurora

Derive insights using advanced analytics in Amazon Redshift from data consolidated from multiple Amazon Aurora database clusters

aws

NEW [GA]

Amazon Redshift integration for Apache Spark

SIMPLIFY AND SPEED UP APACHE SPARK APPLICATIONS ACCESSING AMAZON REDSHIFT DATA FROM AWS ANALYTICS SERVICES

Author Apache Spark applications using Java, Python, Scala, with access to rich, curated data in your data warehouse

No manual setup and maintenance of uncertified versions of Spark-Amazon Redshift open-source connectors

Advanced pushdown optimizations in the Apache Spark-Amazon Redshift connector accelerate 3 TB out-of-the-box TPC-DS queries by 10x

Improved security with IAM-based credentials

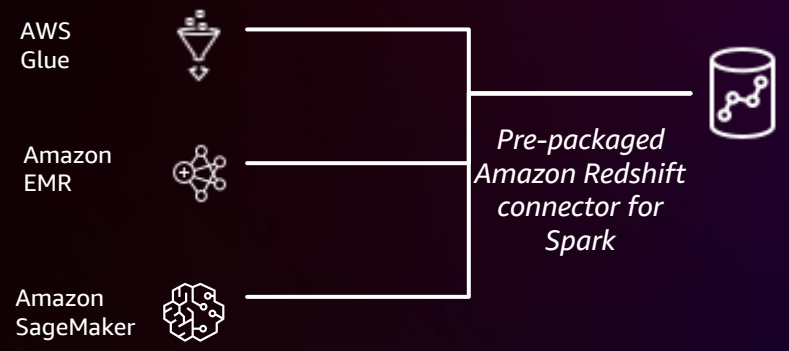

Broadest and most cost-effective set of analytics services

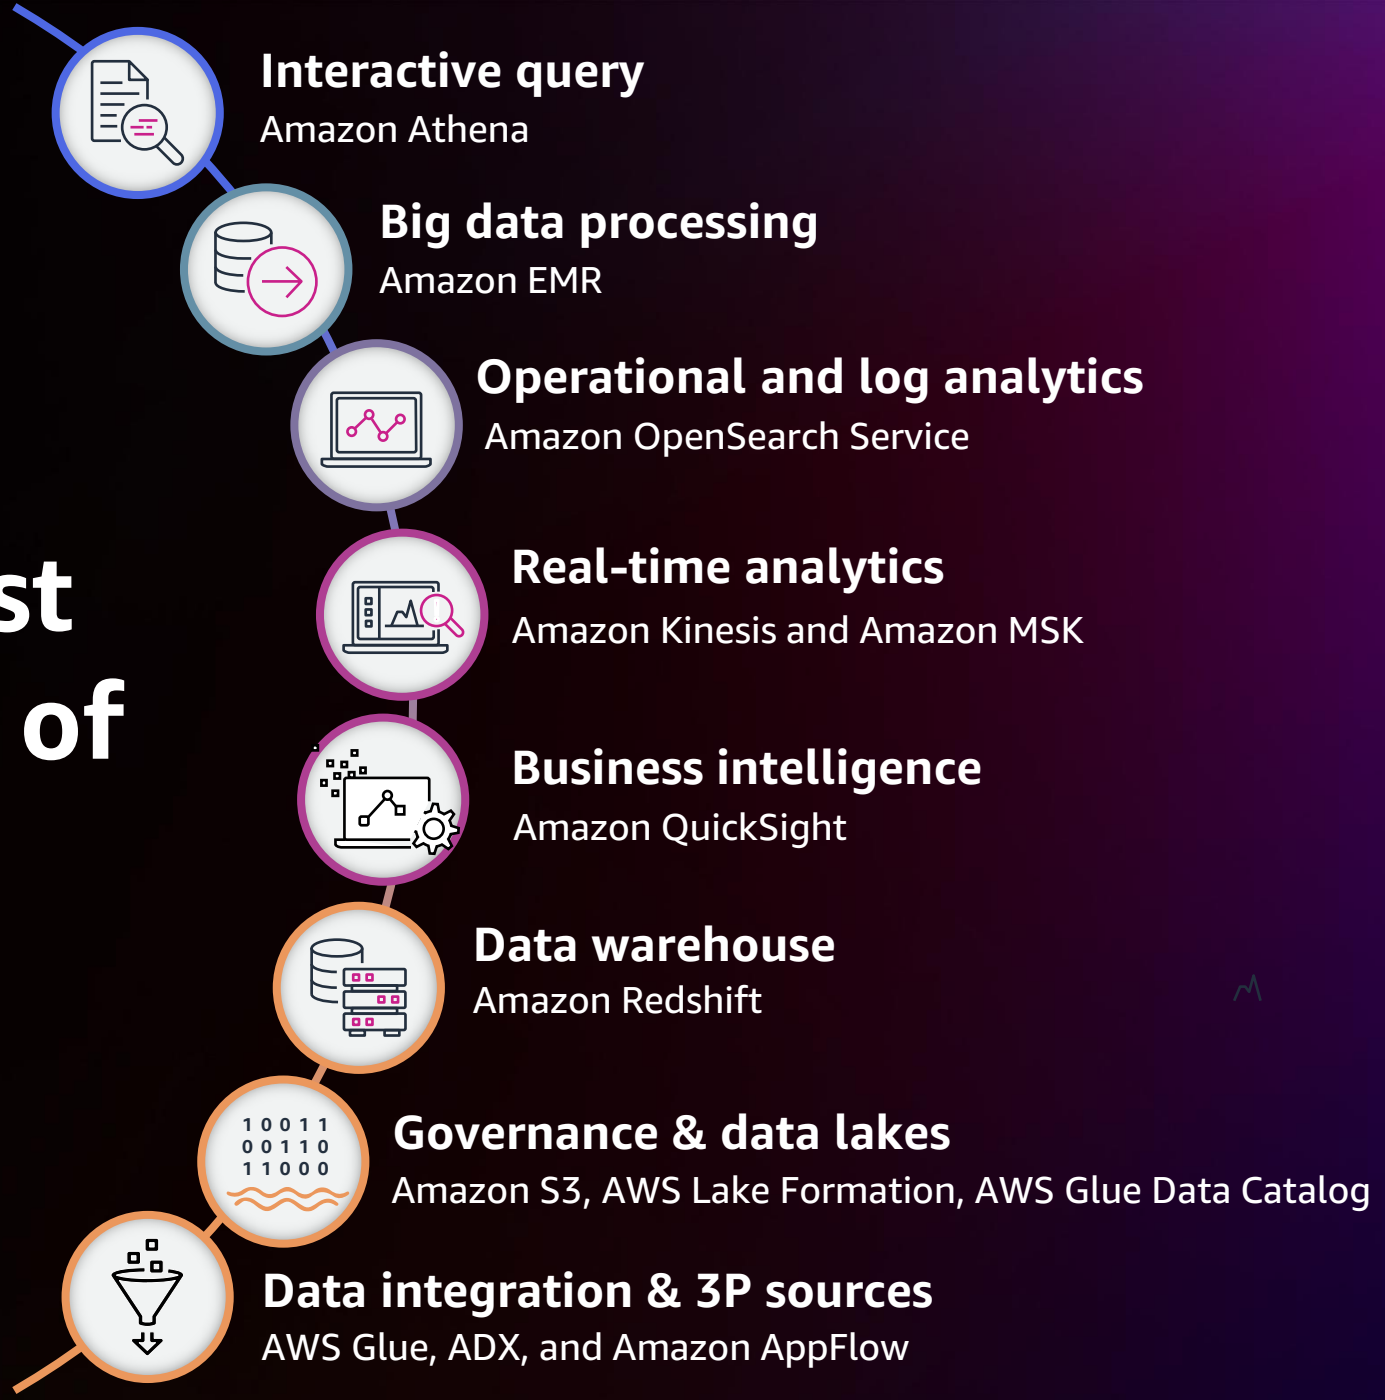

Modern data architecture's five benefits

Unified analytics

Highest performance at the lowest cost

Machine learning integration

Unified data access, security and governance

Insights for everyone

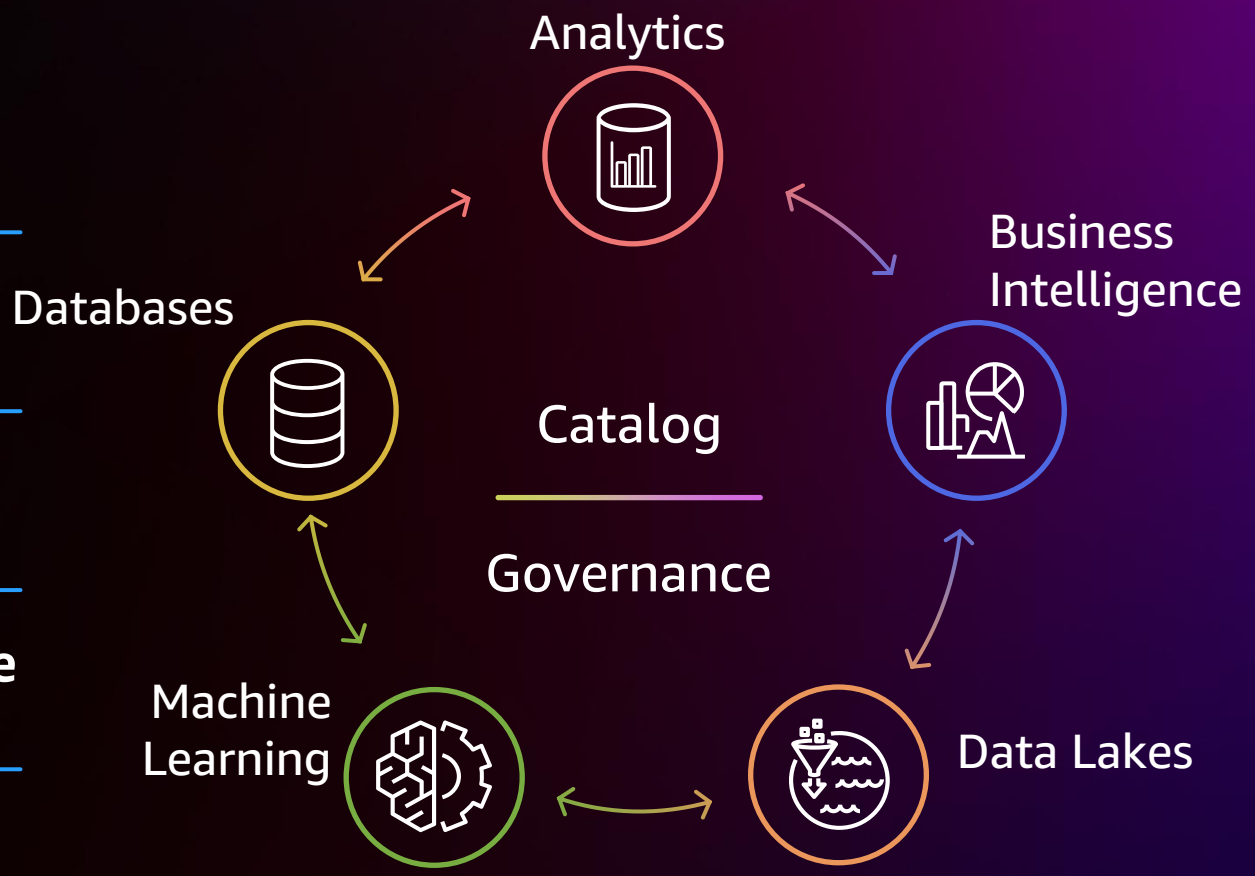

Highest performance at the lowest cost

Scale linearly with predictable high performance

Maximize your cost savings

Self-learning, self-tuning system to enhance performance

Price performance innovations in 2022

Amazon EMR

EMR Serverless GA

11-16% performance improvement with Graviton2 at 20%+ reduced cost

68% performance increase for Spark

19% performance boost on Graviton3

Amazon Athena

Applications start up in under a second

AWS Glue

Auto Scaling

Flex Jobs

2.4x performance increase for Spark

Amazon OpenSearch Service

40% query boost on Graviton2

38% indexing boost on Graviton2

Amazon Redshift

Up to 5x better price performance vs. other cloud data warehouses

Up to 7x better price performance vs. other cloud data warehouses on high concurrency

Low latency workloads like dashboarding applications

Amazon Redshift performance improvements

CONTINUOUS OUT-OF-THE-BOX IMPROVEMENTS FOR BETTER PRICE PERFORMANCE AT ANY SCALE

Compute	System	Autonomics
Vectorized scans for Amazon Redshift tables	String-encoding for in-memory perf	Performance mode
Write/Commit performance	CaaS cache pre-warming	Auto WLM enhancements
Snapshot isolation	CaaS region expansion	ATO enhancements
Concurrency scaling writes (GA)	Incremental updates of MVs on datashares	Advisor enhancements

Differentiated performance

ON SPARK, PRESTO, AND HIVE

3.9x

faster than standard Apache Spark 3.0 in TPC-DS 3TB benchmark

4.2x

faster than standard OSS Trino 388 in TPC-DS 3TB benchmarks

11-16%

performance improvement with Graviton2 at **20%+ reduced cost**

100%

open-source API compliant

30%

Better price-performance with Graviton2

AWS Glue cost optimization

AWS GLUE JOB RUNS AT LOWER COST

Auto scaling

Automatically resize compute for lower cost

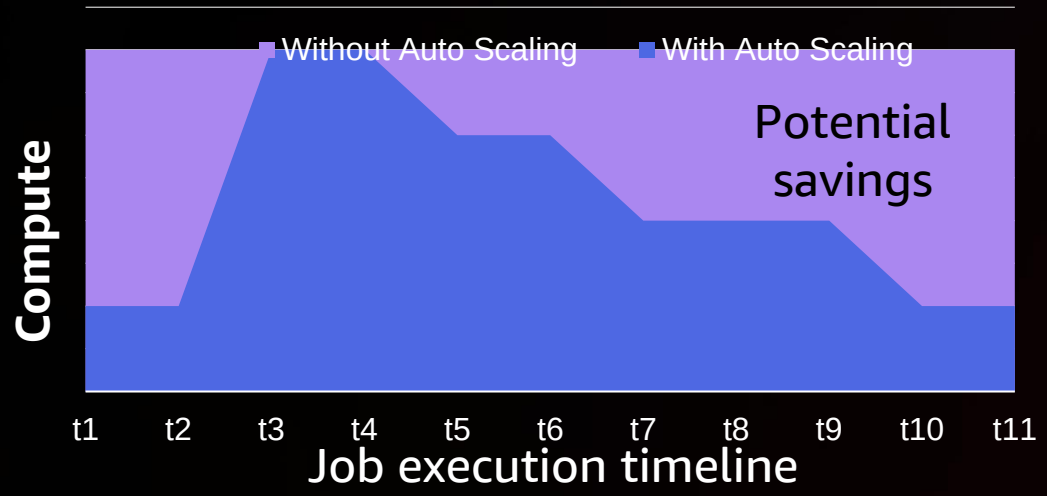

Reduce cost by 20-40%

Simplify capacity and performance planning

Flex

spare capacity execution

Up to **34% cost savings**

Cost effective for **one-time data-load** workloads

Modern data architecture's five benefits

Unified analytics

Highest performance at the lowest cost

Machine learning integration

Unified data access, security and governance

Insights for everyone

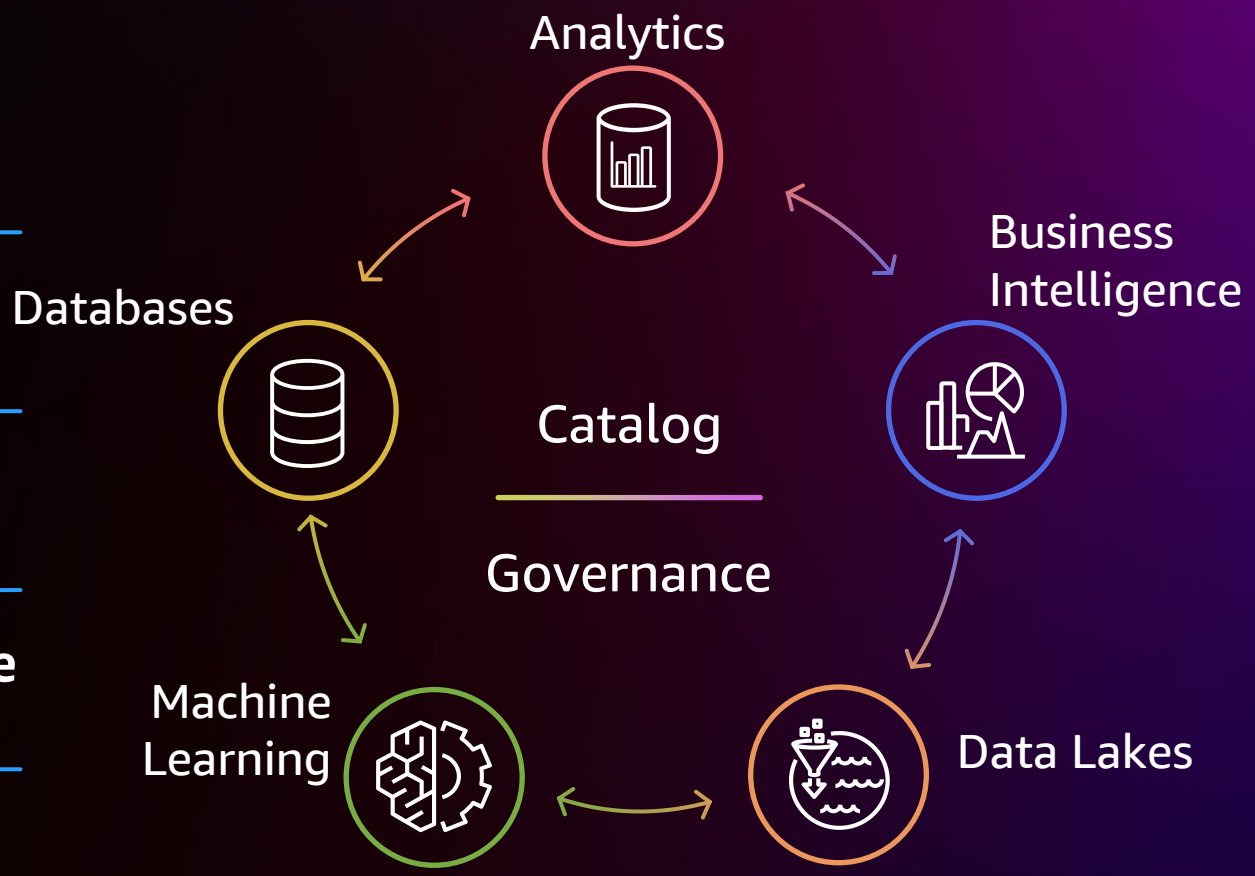

Connecting data services and ML to drive more value

Databases

Data warehouses and data lakes

Business intelligence tools

AMAZON AURORA ML

AMAZON NEPTUNE ML

AMAZON REDSHIFT ML

AMAZON ATHENA ML

AMAZON QUICKSIGHT Q

Amazon Redshift ML

EASILY CREATE AND TRAIN ML MODELS USING SQL QUERIES WITH AMAZON SAGEMAKER

80+ billion

predictions per week

Train and create ML models using SQL

Automatic pre-processing, creation, training, deployment, and inferencing of models

SageMaker models for in-database or remote inference

Supervised and unsupervised trainings

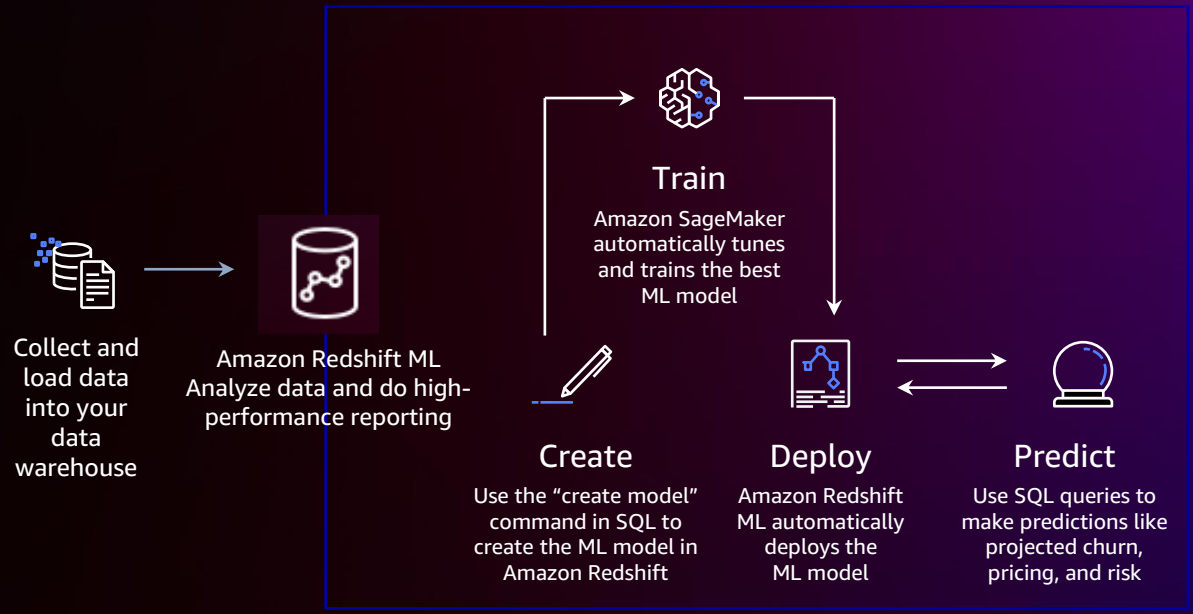

jobcase

Amazon SageMaker Studio Universal Notebooks

PERFORM DATA ENGINEERING, ANALYTICS AND ML IN ONE NOTEBOOK

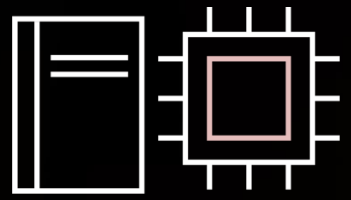

Built-in data and analytics integration

Connect with Amazon EMR, AWS Glue and data lakes on Amazon S3

Interactive data preparation

Interactively query, analyze and transform wide range of data

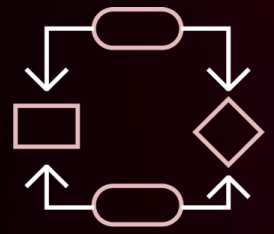

Inline debugging and monitoring

Visually debug and monitor Spark jobs inline in same notebook

Build ML workflows

Build end to end ML workflows without leaving the notebook

Modern data architecture's five benefits

Unified analytics

Highest performance at the lowest cost

Machine learning integration

Unified data access, security and governance

Insights for everyone

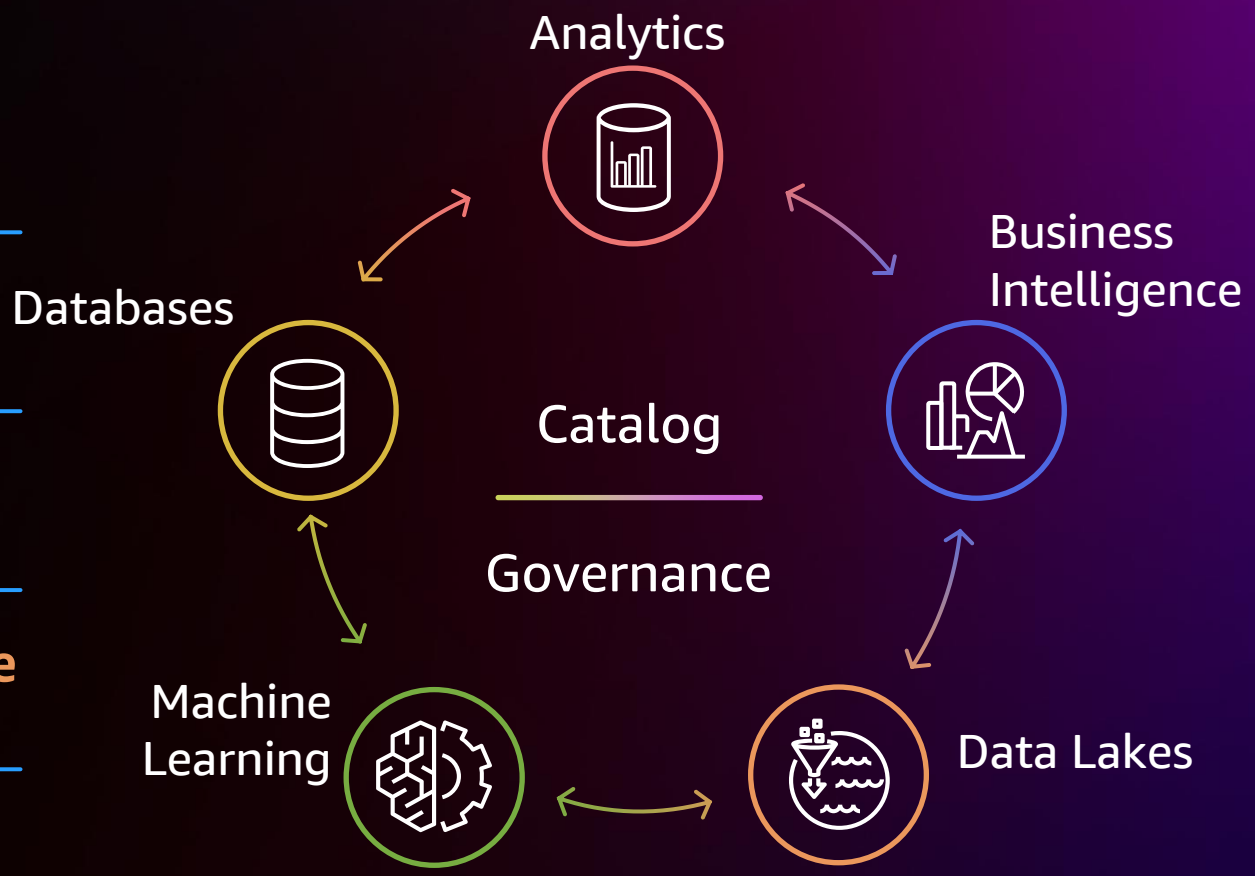

Data Governance starts with business

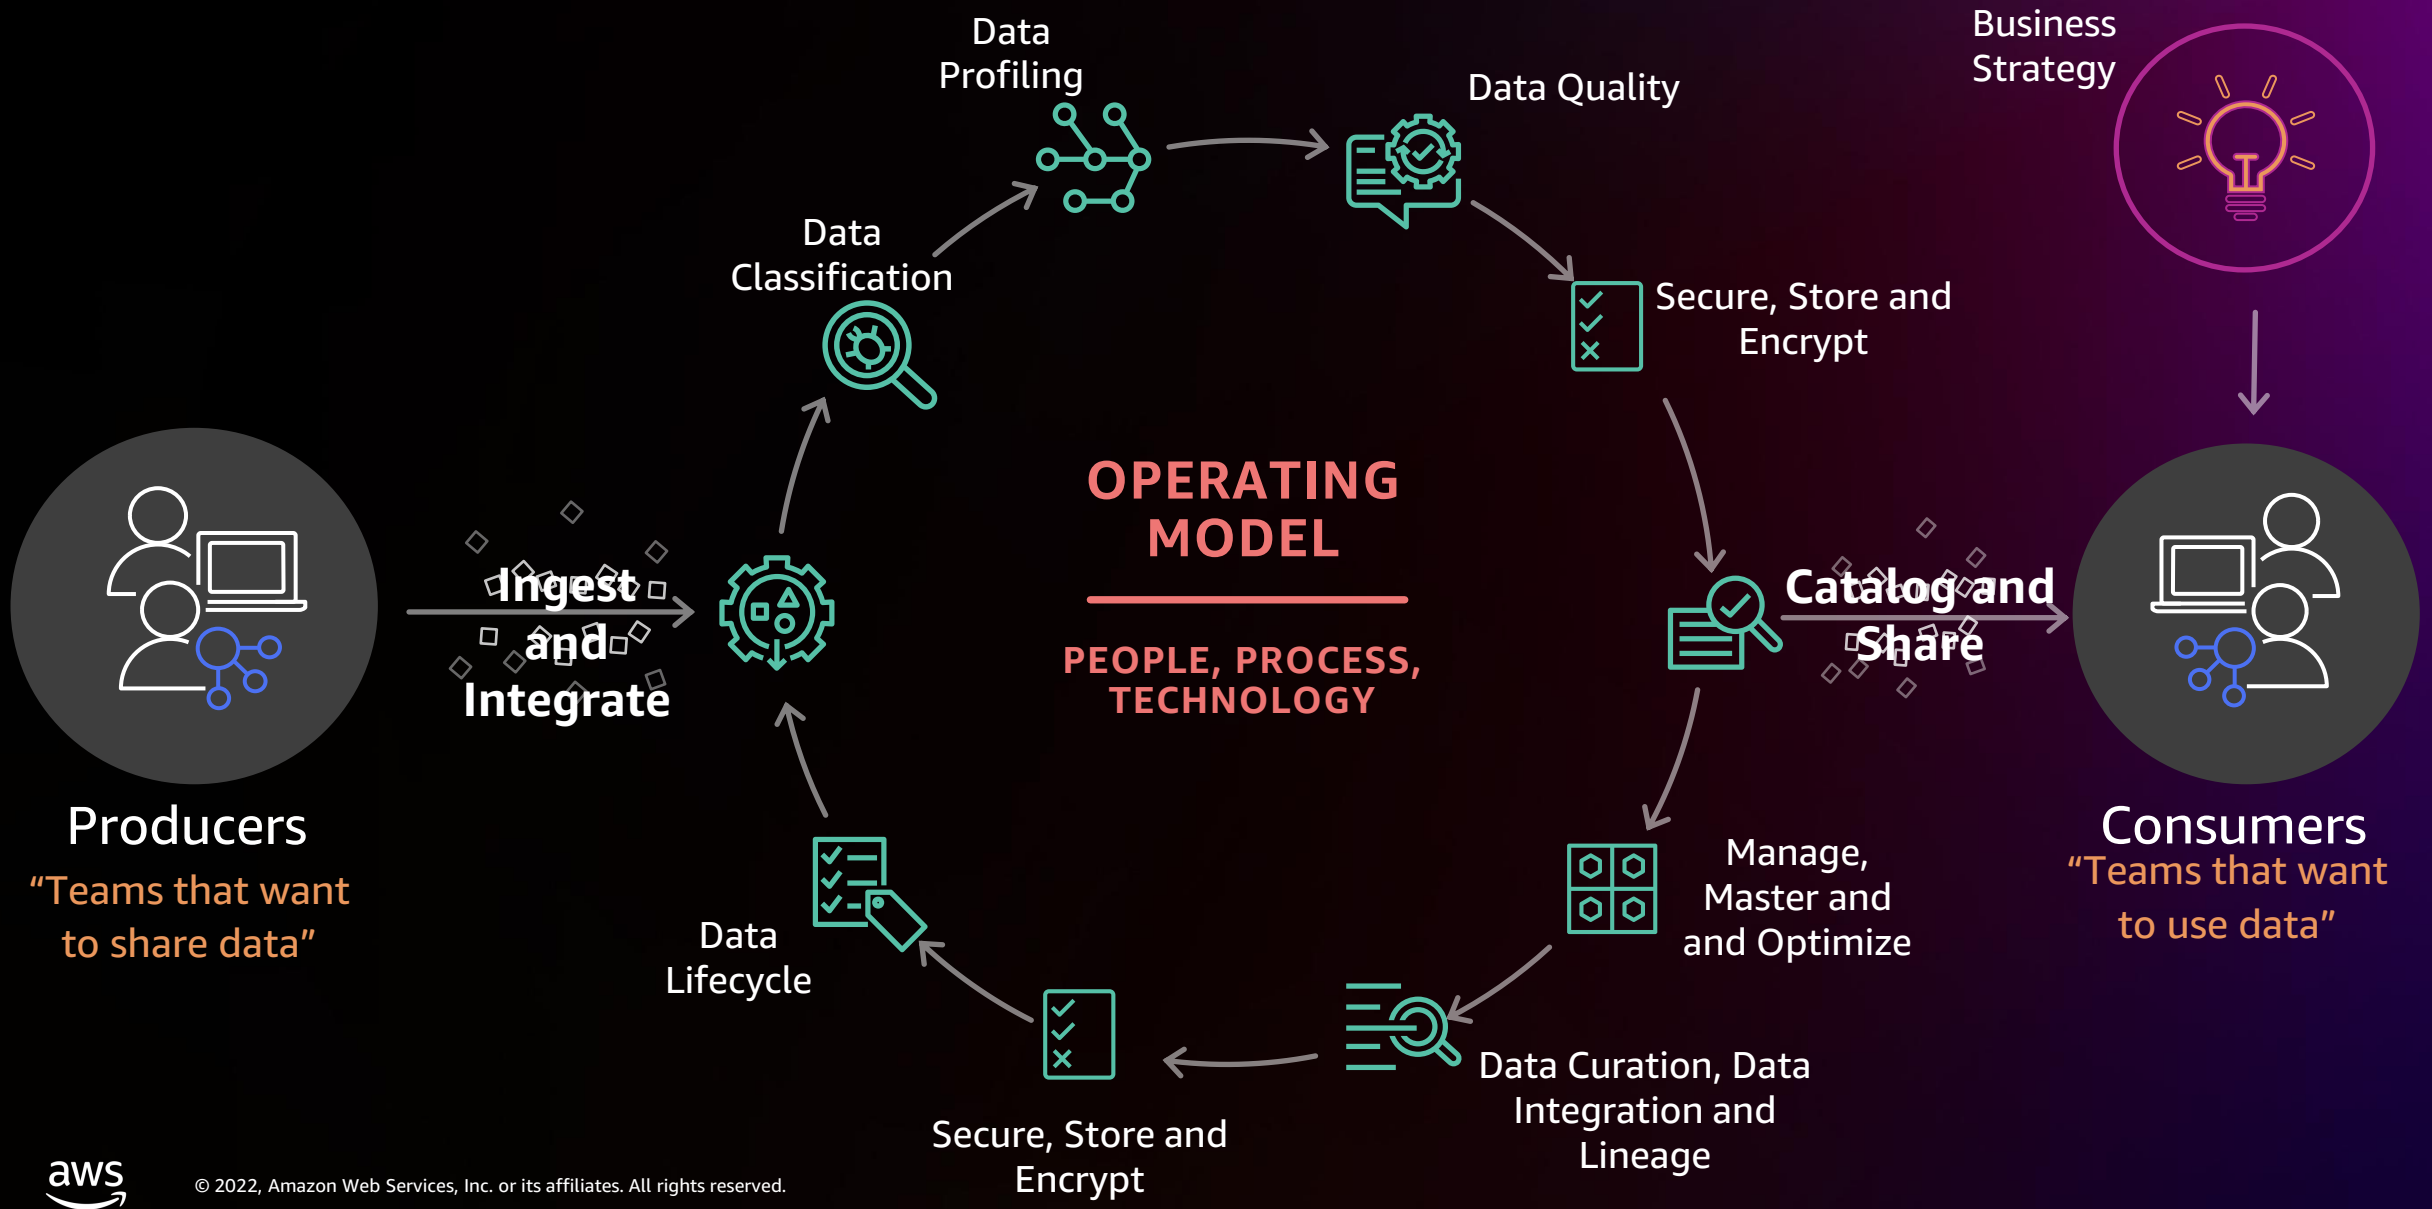

AWS Lake Formation unifies data governance

Simplify security management with Lake Formation

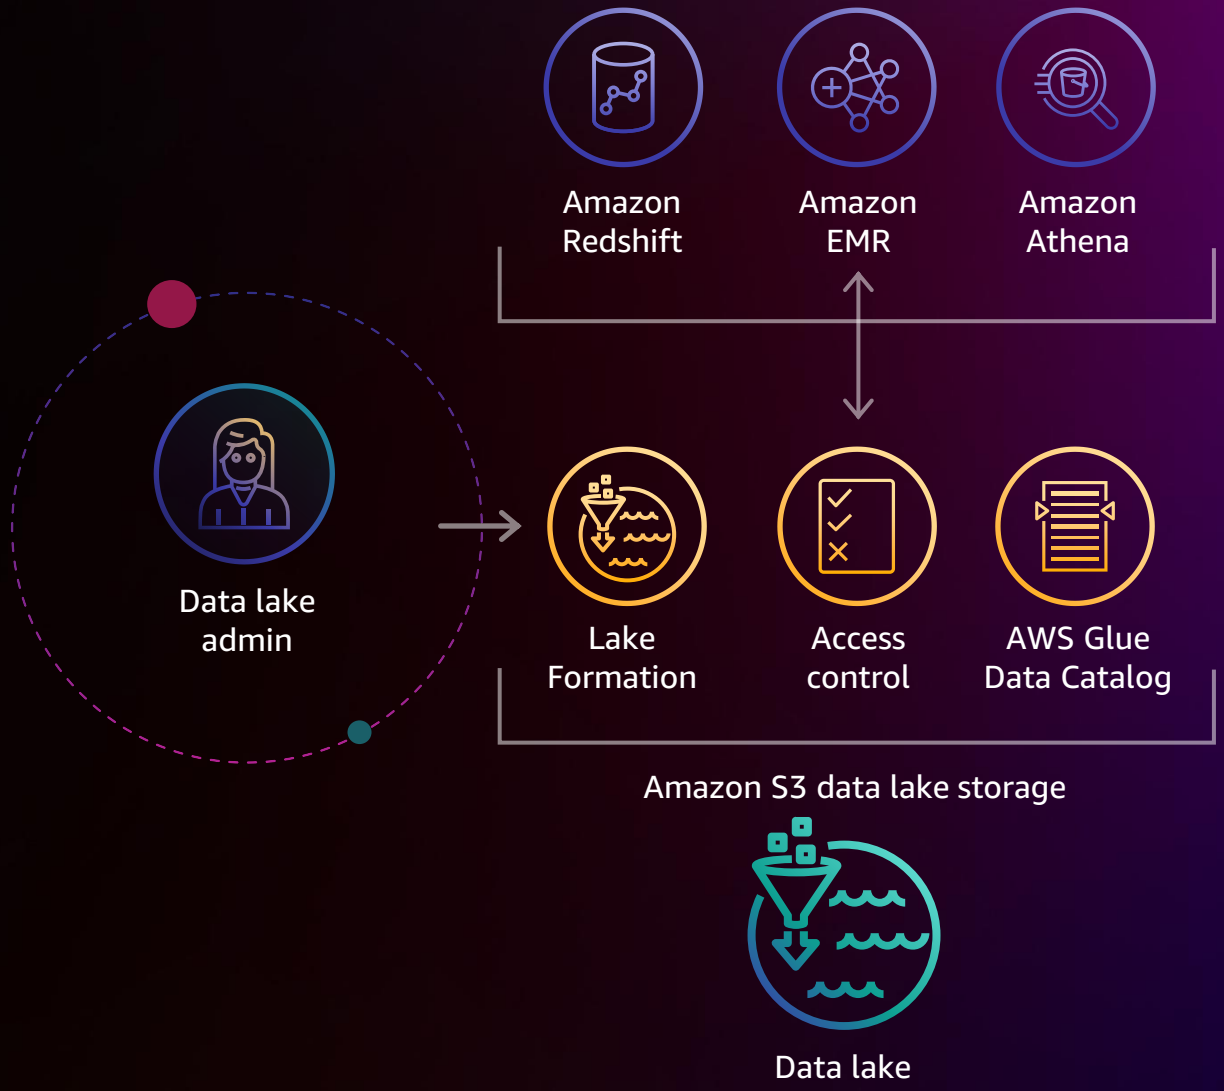

NEW [PREVIEW]

Amazon DataZone

UNLOCK DATA ACCESS FOR ALL USERS WITH TRUSTED AUTONOMY

Catalog your data with business context

Manage organization-wide governance in one place

Simplify access to analytics for everyone in your organization

Work with data to solve specific business use cases

Core components of Amazon DataZone

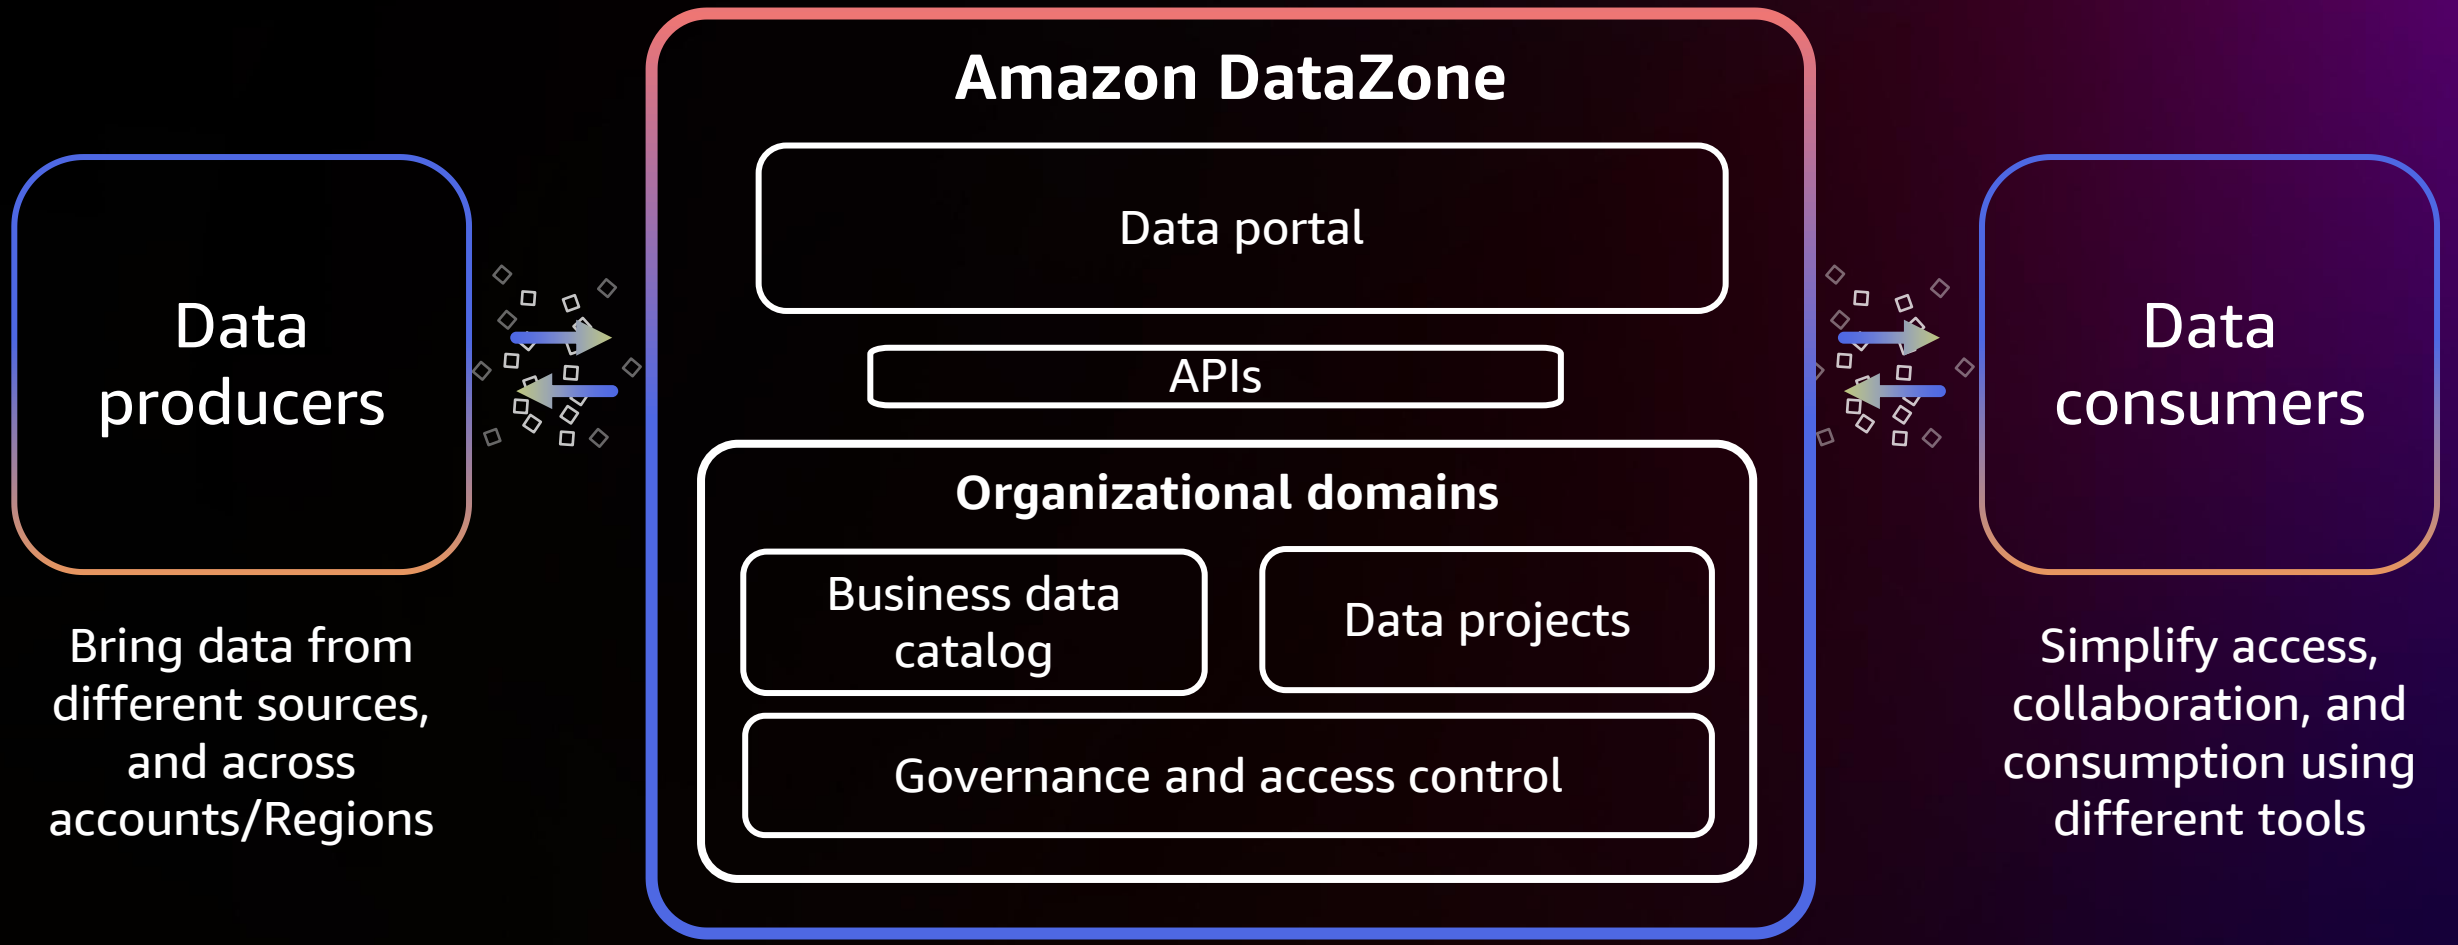

Bring data from different sources, and across accounts/Regions

Simplify access, collaboration, and consumption using different tools

Modern data architecture's five benefits

Unified analytics

Highest performance at the lowest cost

Machine learning integration

Unified data access, security and governance

Insights for everyone

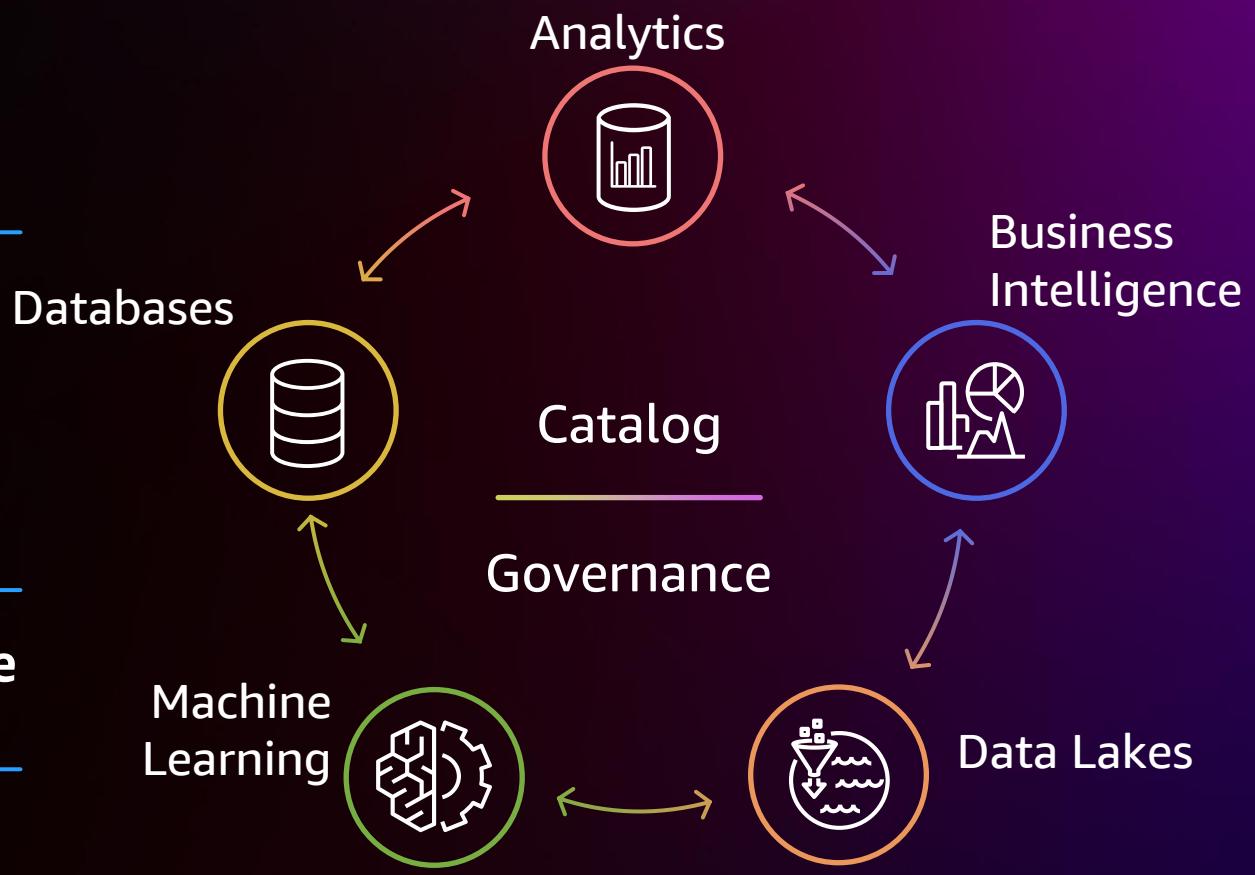

Insights for everyone

Focus on data without managing infrastructure

Choose your tools based on your skillset

Serverless is a key for your data infrastructure

The benefits of serverless

A full stack of serverless options for data analytics in the cloud

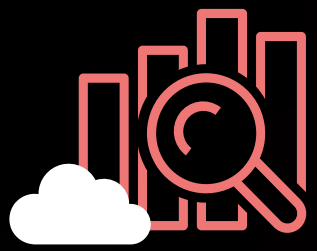

Amazon OpenSearch Serverless

Amazon OpenSearch Service securely unlocks real-time search, monitoring, and analysis of operational data

Easy to administer

No sizing, scaling, and tuning of clusters, and no shard and index lifecycle management

Fast

Automatically scale resources to maintain consistently fast data ingestion rates and query response times

Ecosystem

Get started in seconds using the same OpenSearch Service clients, pipelines, and APIs

Cost-effective

Pay only for the resources consumed

Amazon SageMaker Canvas

Build ML models and generate accurate predictions — no code required

Quickly access and prepare data for Machine Learning

Built-in AutoML to build models and generate accurate predictions

Share ML models and collaborate with data science teams

Usage-based pricing to avoid licensing fees and reduce TCO

Insights through Amazon QuickSight

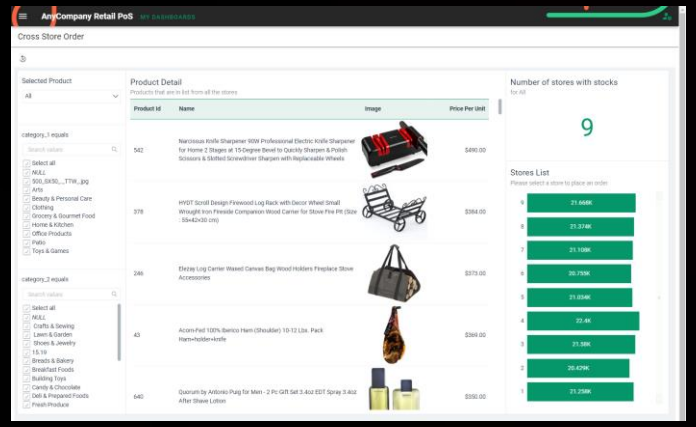

Embedded insights

Enhance customer-facing products and monetize data assets

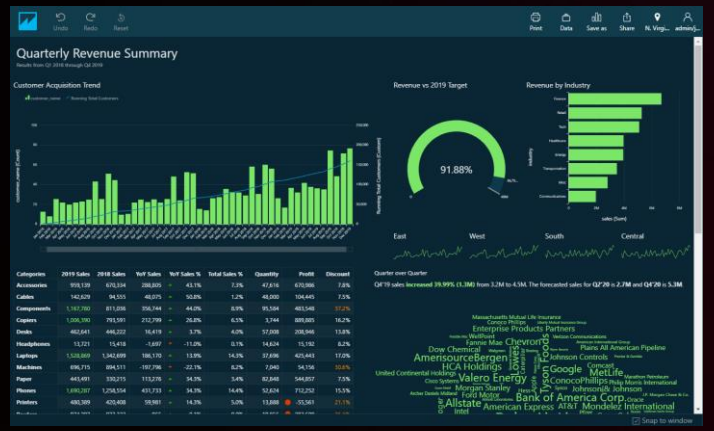

Interactive dashboards, meaningful insights

Dashboards, visualizations, and ad-hoc analysis primarily for internal audiences

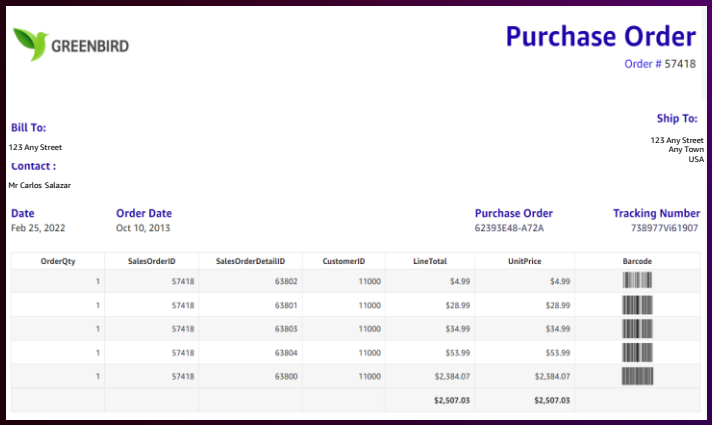

Enterprise reporting

Static, highly formatted, email-based reporting distributed to large internal or external audiences

Amazon QuickSight Q

Ask natural language questions about your data
and get answers in seconds

Modern Data Architecture is making it easier to unlock the value of data across the end-to-end data journey

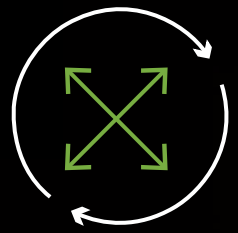

Scalable

Performance at scale

Unified

Connect to all your data

Comprehensive

Tools for all your workloads

Governed

End-to-end governance

Agenda

Modern data architecture on AWS

End-to-end data life cycle on the modern data architecture

Data governance and data mesh in action

Journey towards modern data architecture

End to end data life cycle

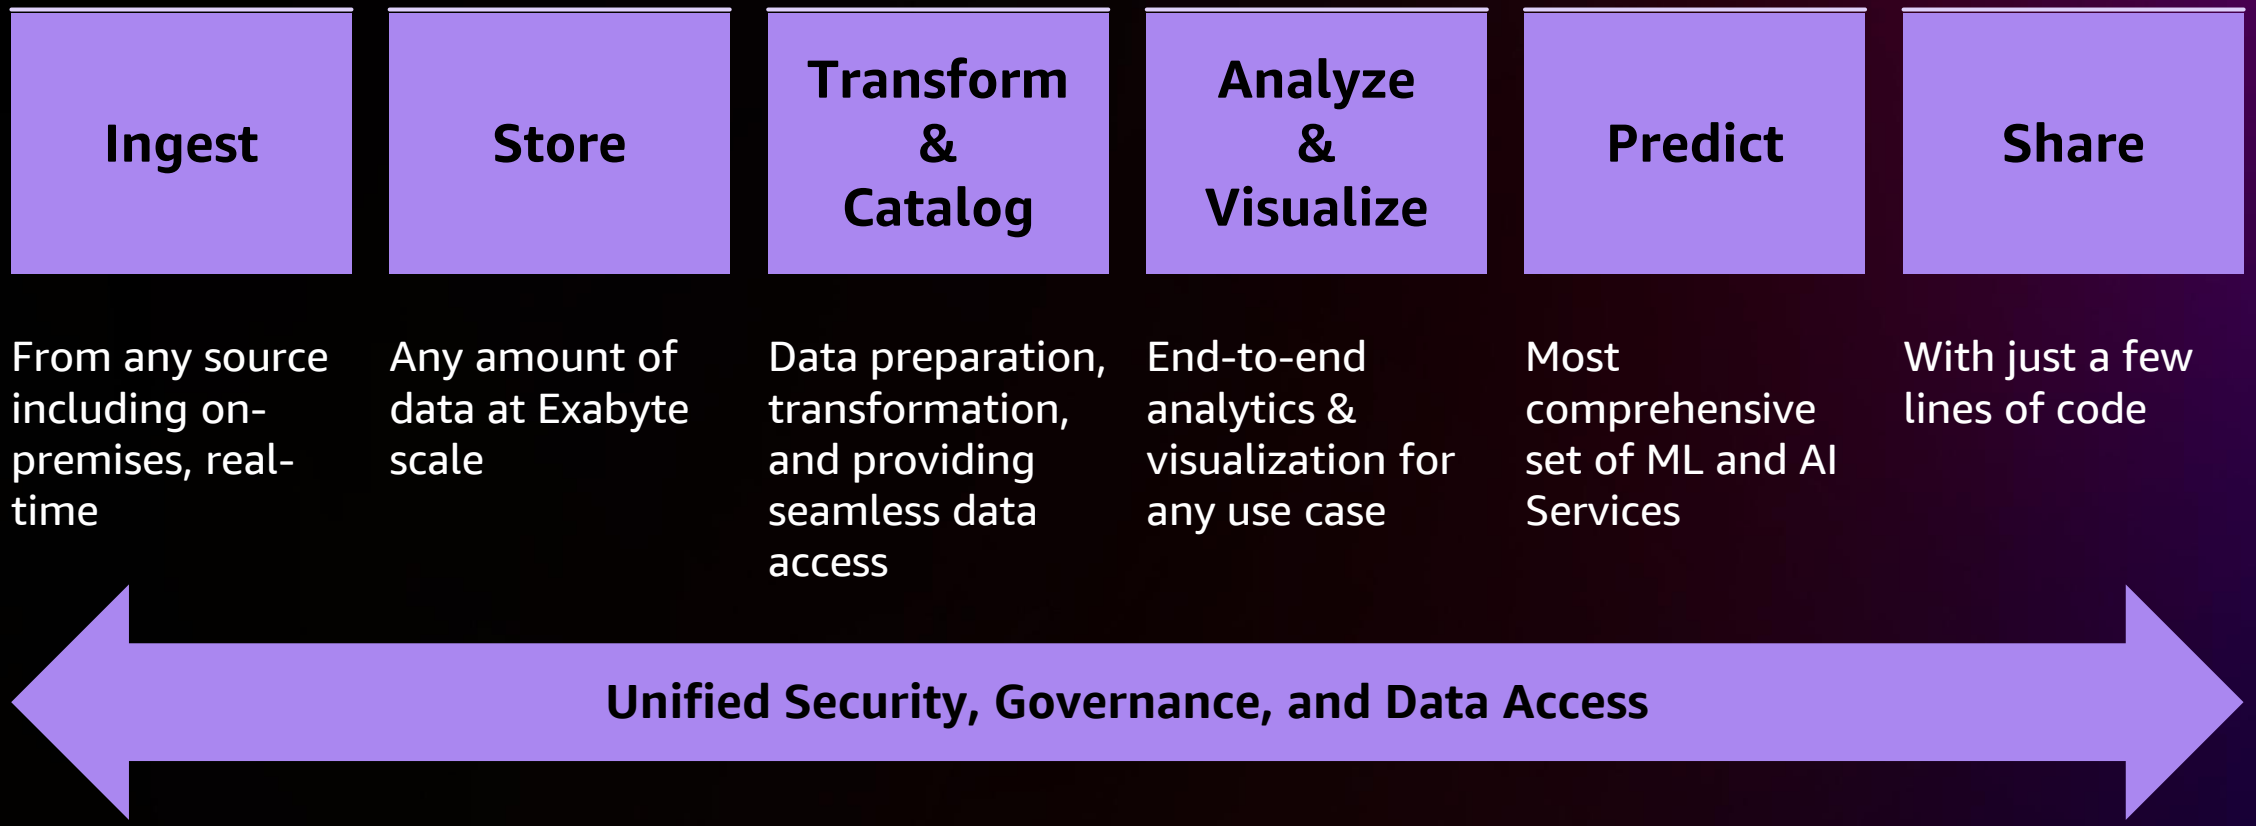

Typical architecture

#1: Ingest

Ingesting data from any source including on-premise sources and data that is generated in real-time.

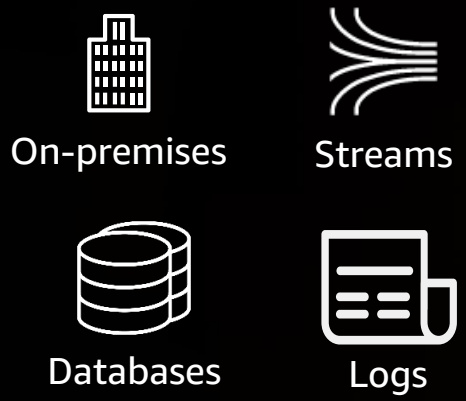

#2: Store

Storing both transactional data in databases and analytical data in data warehouses and data lakes at any scale.

#3: Transform & Catalog

Making it easy to access their data and keeping their data in sync regardless of where it lives

#4: Analyze & Visualize

Analyzing data using any of ad hoc queries, distributed frameworks and search engines, and visualize the data on dashboards

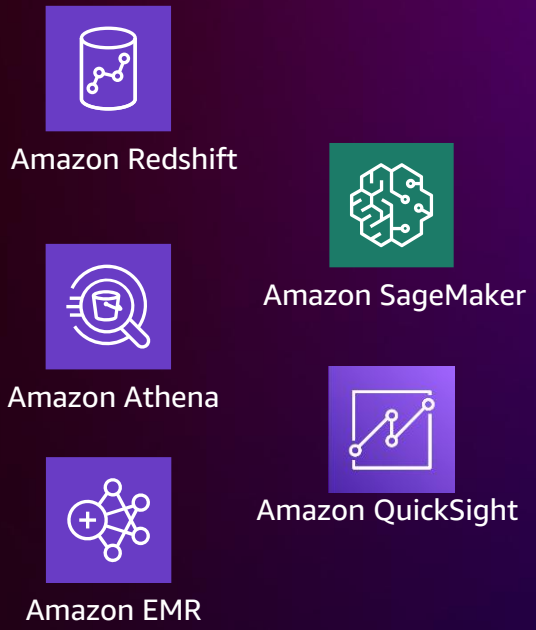

#5: Predict

Adding ML-based intelligence to applications without needing ML skills

#6: Share

Sharing new insights to take intelligent, data-driven actions

Ingest

#1: Ingest

Ingesting data from any source including on-premise sources and data that is generated in real-time.

AWS Glue ETL

Amazon S3

#2: Store

Storing both transactional data in databases and analytical data in data warehouses and data lakes at any scale.

#3: Trans atalog

Making it easy to access their data and keeping their data in sync regardless of where it lives

AWS Transfer Family

AWS DataSync

Amazon Kinesis Data Streams

Amazon Kinesis Data Firehose

AWS Glue Data Catalog

Amazon AppFlow

AWS Glue

#4: Analyze & Visualize

Analyzing data using any of ad hoc queries ed frameworks and search nd visualize the data on dash

AWS Database Migration Service (AWS DMS)

Amazon Managed Streaming for Apache Kafka (Amazon MSK)

AWS Data Exchange

Amazon SageMaker

Amazon QuickSight

#5: Predict

Adding ML-based intelligence to applications without needing ML skills

#6: Share

Sharing new insights to take intelligent, data-driven actions

Store

#1: Ingest

Ingesting data from any source including on-premise sources and data that is generated in real-time.

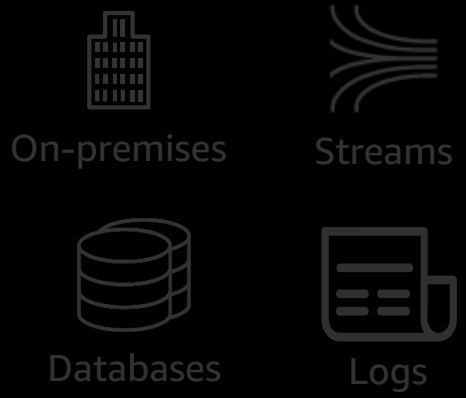

#3: Transform & Catalog

Making it easy to access their data and keeping their data in sync regardless of where it lives

#4: Analyze & Visualize

Analyzing data using any of ad hoc queries, distributed frameworks and search engines, and visualize the data on dashboards

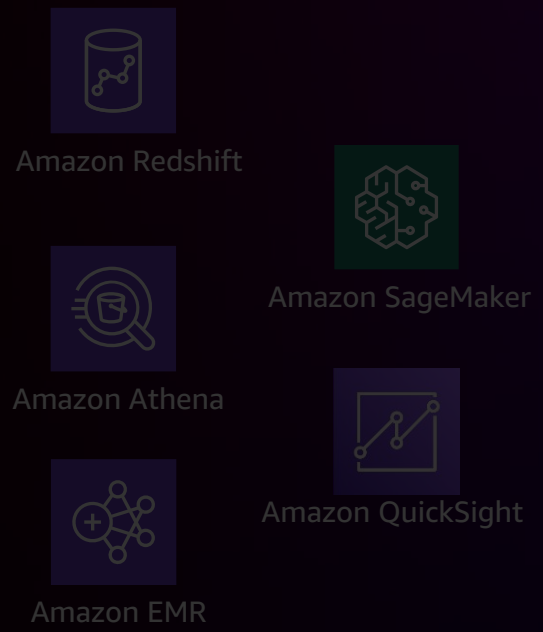

#2: Store

Storing both transactional data in databases and analytical data in data warehouses and data lakes at any scale.

#5: Predict

Adding ML-based intelligence to applications without needing ML skills

#6: Share

Sharing new insights to take intelligent, data-driven actions

Amazon S3 for data lakes

AN OBJECT STORAGE SERVICE OFFERING INDUSTRY-LEADING SCALABILITY, DATA AVAILABILITY, SECURITY, AND PERFORMANCE

Durability, availability, and **scalability**

Easy to use with **cost optimization:** Intelligent tiering

Most ways to **get data in**

Most **object-level controls**

Broad portfolio of analytics tools

Security, compliance, and **audit** capabilities

Cold storage and archive capabilities

Transform & Catalog

#1: Ingest

Ingesting data from any source including on-premise sources and data that is generated in real-time.

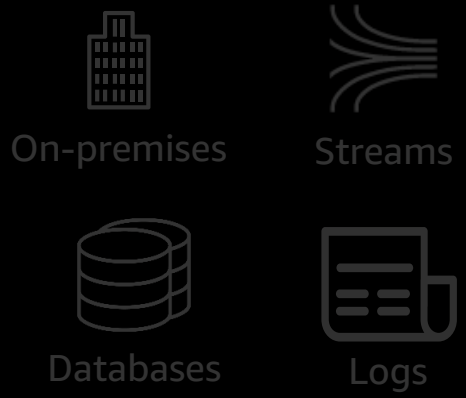

#2: Store

Storing both transactional data in databases and analytical data in data warehouses and data lakes at any scale.

#3: Transform & Catalog

Making it easy to access their data and keeping their data in sync regardless of where it lives

#4: Analyze & Visualize

Analyzing data using any of ad hoc queries, distributed frameworks and search engines, and visualize the data on dashboards

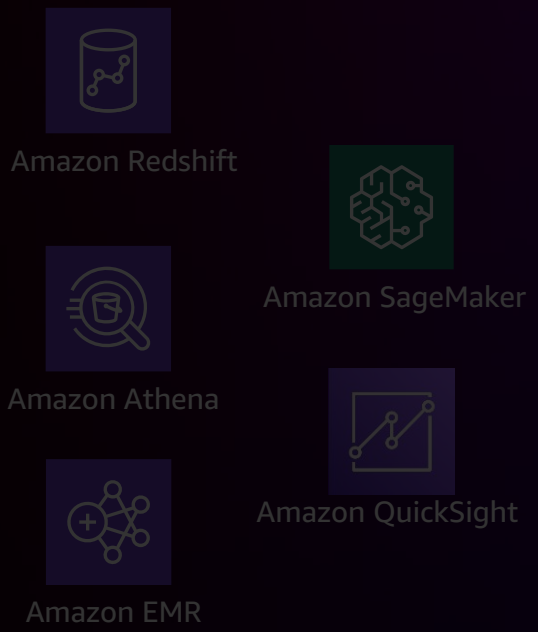

#5: Predict

Adding ML-based intelligence to applications without needing ML skills

#6: Share

Sharing new insights to take intelligent, data-driven actions

Analyze & Visualize

#1: Ingest

Ingesting data from any source including on-premise sources and data that is generated in real-time.

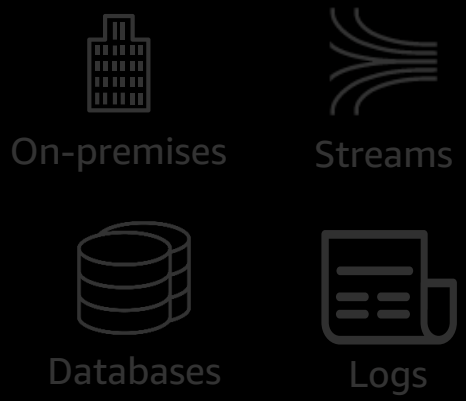

#2: Store

Storing both transactional data in databases and analytical data in data warehouses and data lakes at any scale.

#3: Transform & Catalog

Making it easy to access their data and keeping their data in sync regardless of where it lives

#4: Analyze & Visualize

Analyzing data using any of ad hoc queries, distributed frameworks and search engines, and visualize the data on dashboards

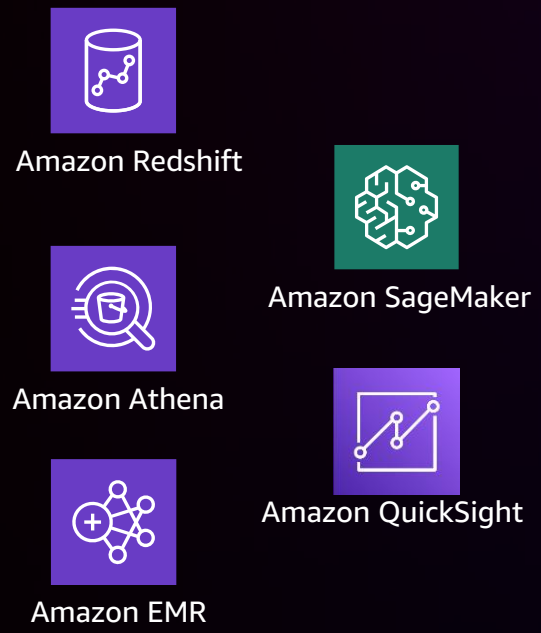

#5: Predict

Adding ML-based intelligence to applications without needing ML skills

#6: Share

Sharing new insights to take intelligent, data-driven actions

Predict and share

#1: Ingest

Ingesting data from any source including on-premise sources and data that is generated in real-time.

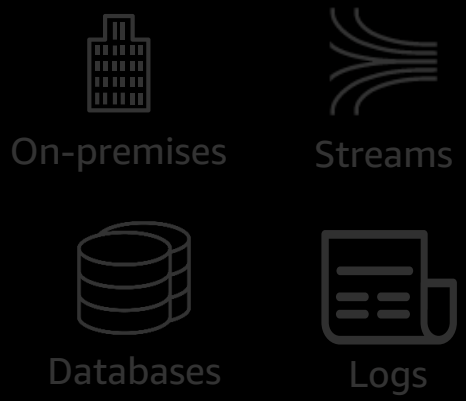

#2: Store

Storing both transactional data in databases and analytical data in data warehouses and data lakes at any scale.

#3: Transform & Catalog

Making it easy to access their data and keeping their data in sync regardless of where it lives

#4: Analyze & Visualize

Analyzing data using any of ad hoc queries, distributed frameworks and search engines, and visualize the data on dashboards

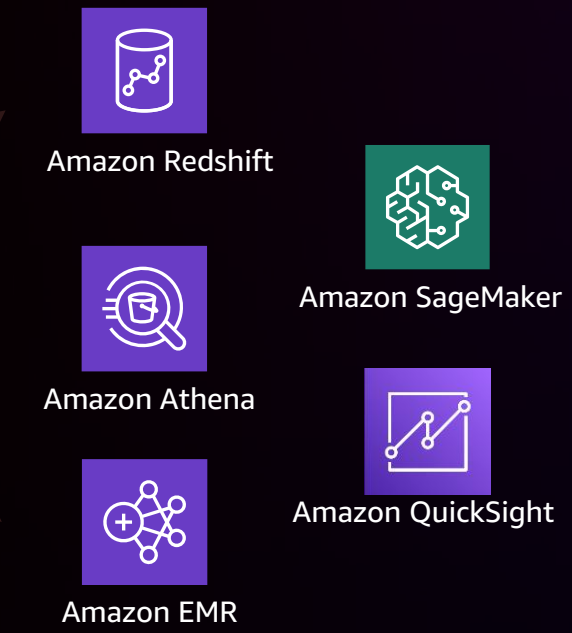

#5: Predict

Adding ML-based intelligence to applications without needing ML skills

#6: Share

Sharing new insights to take intelligent, data-driven actions

Unified Security, Governance, and Data Access

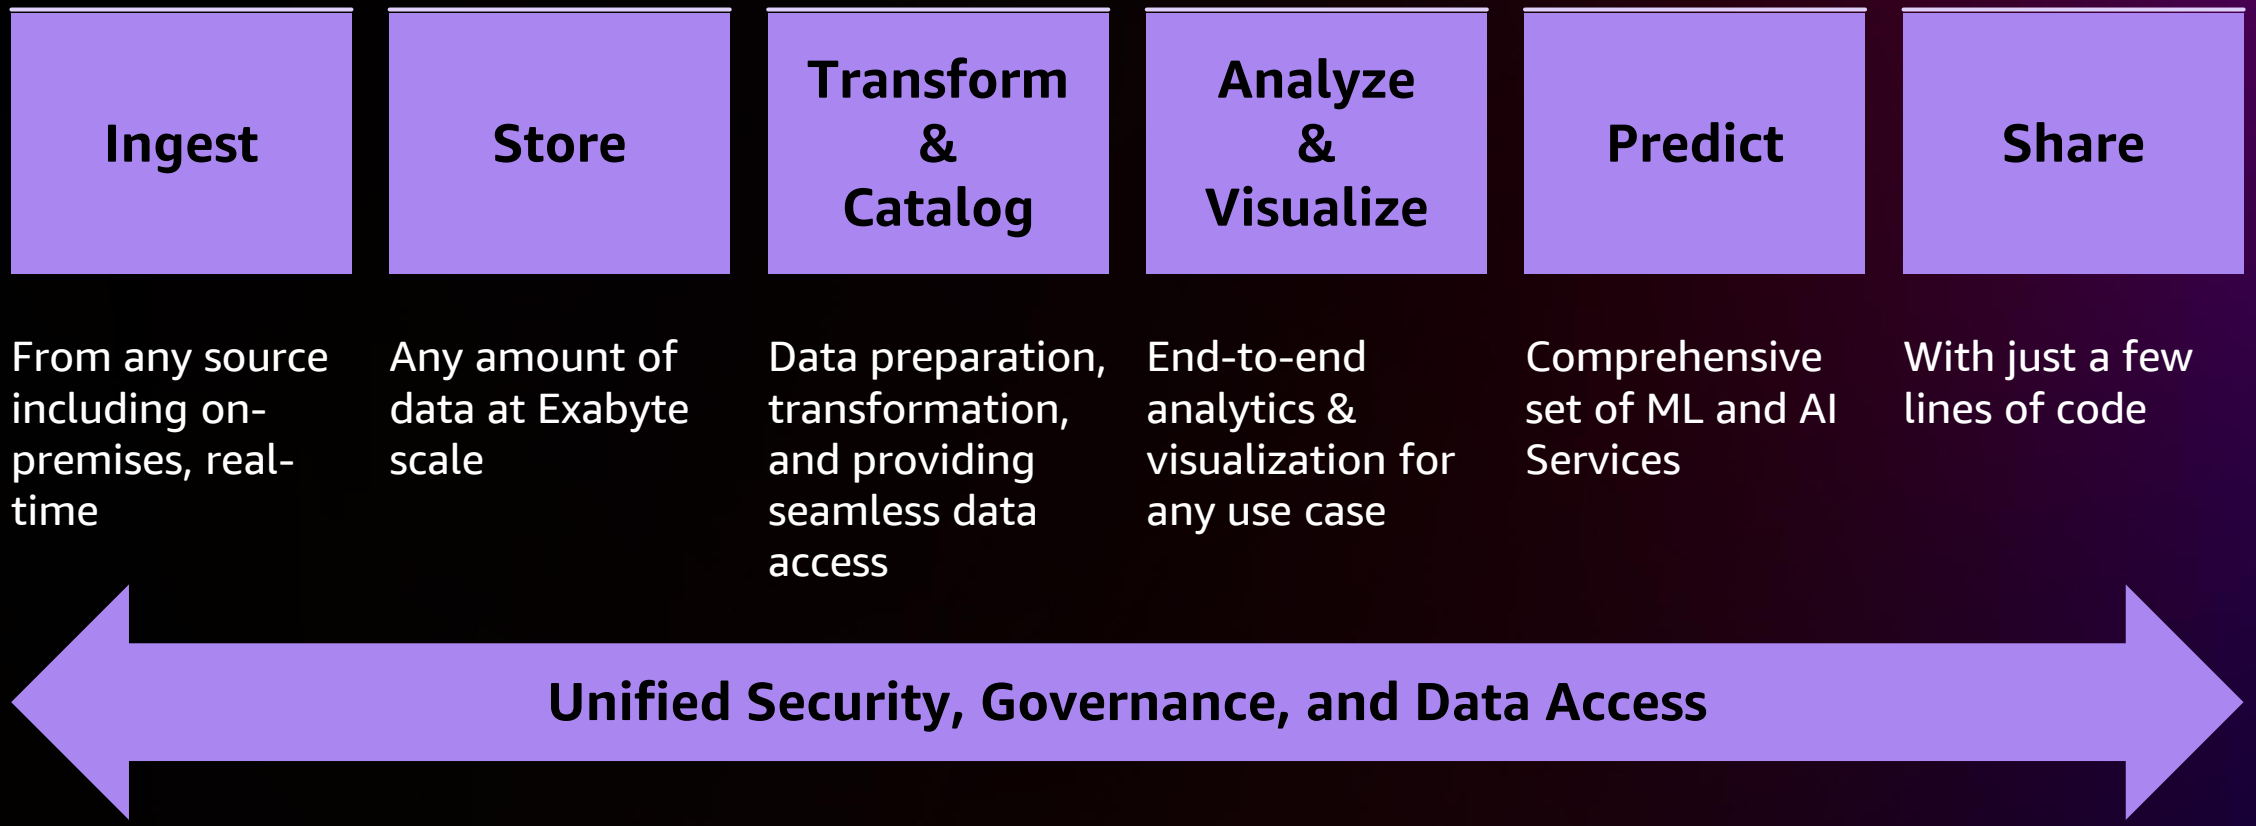

Agenda

Modern data architecture on AWS

End-to-end data life cycle on the modern data architecture

Data governance and data mesh in action

Journey towards modern data architecture

Challenges securing and sharing data

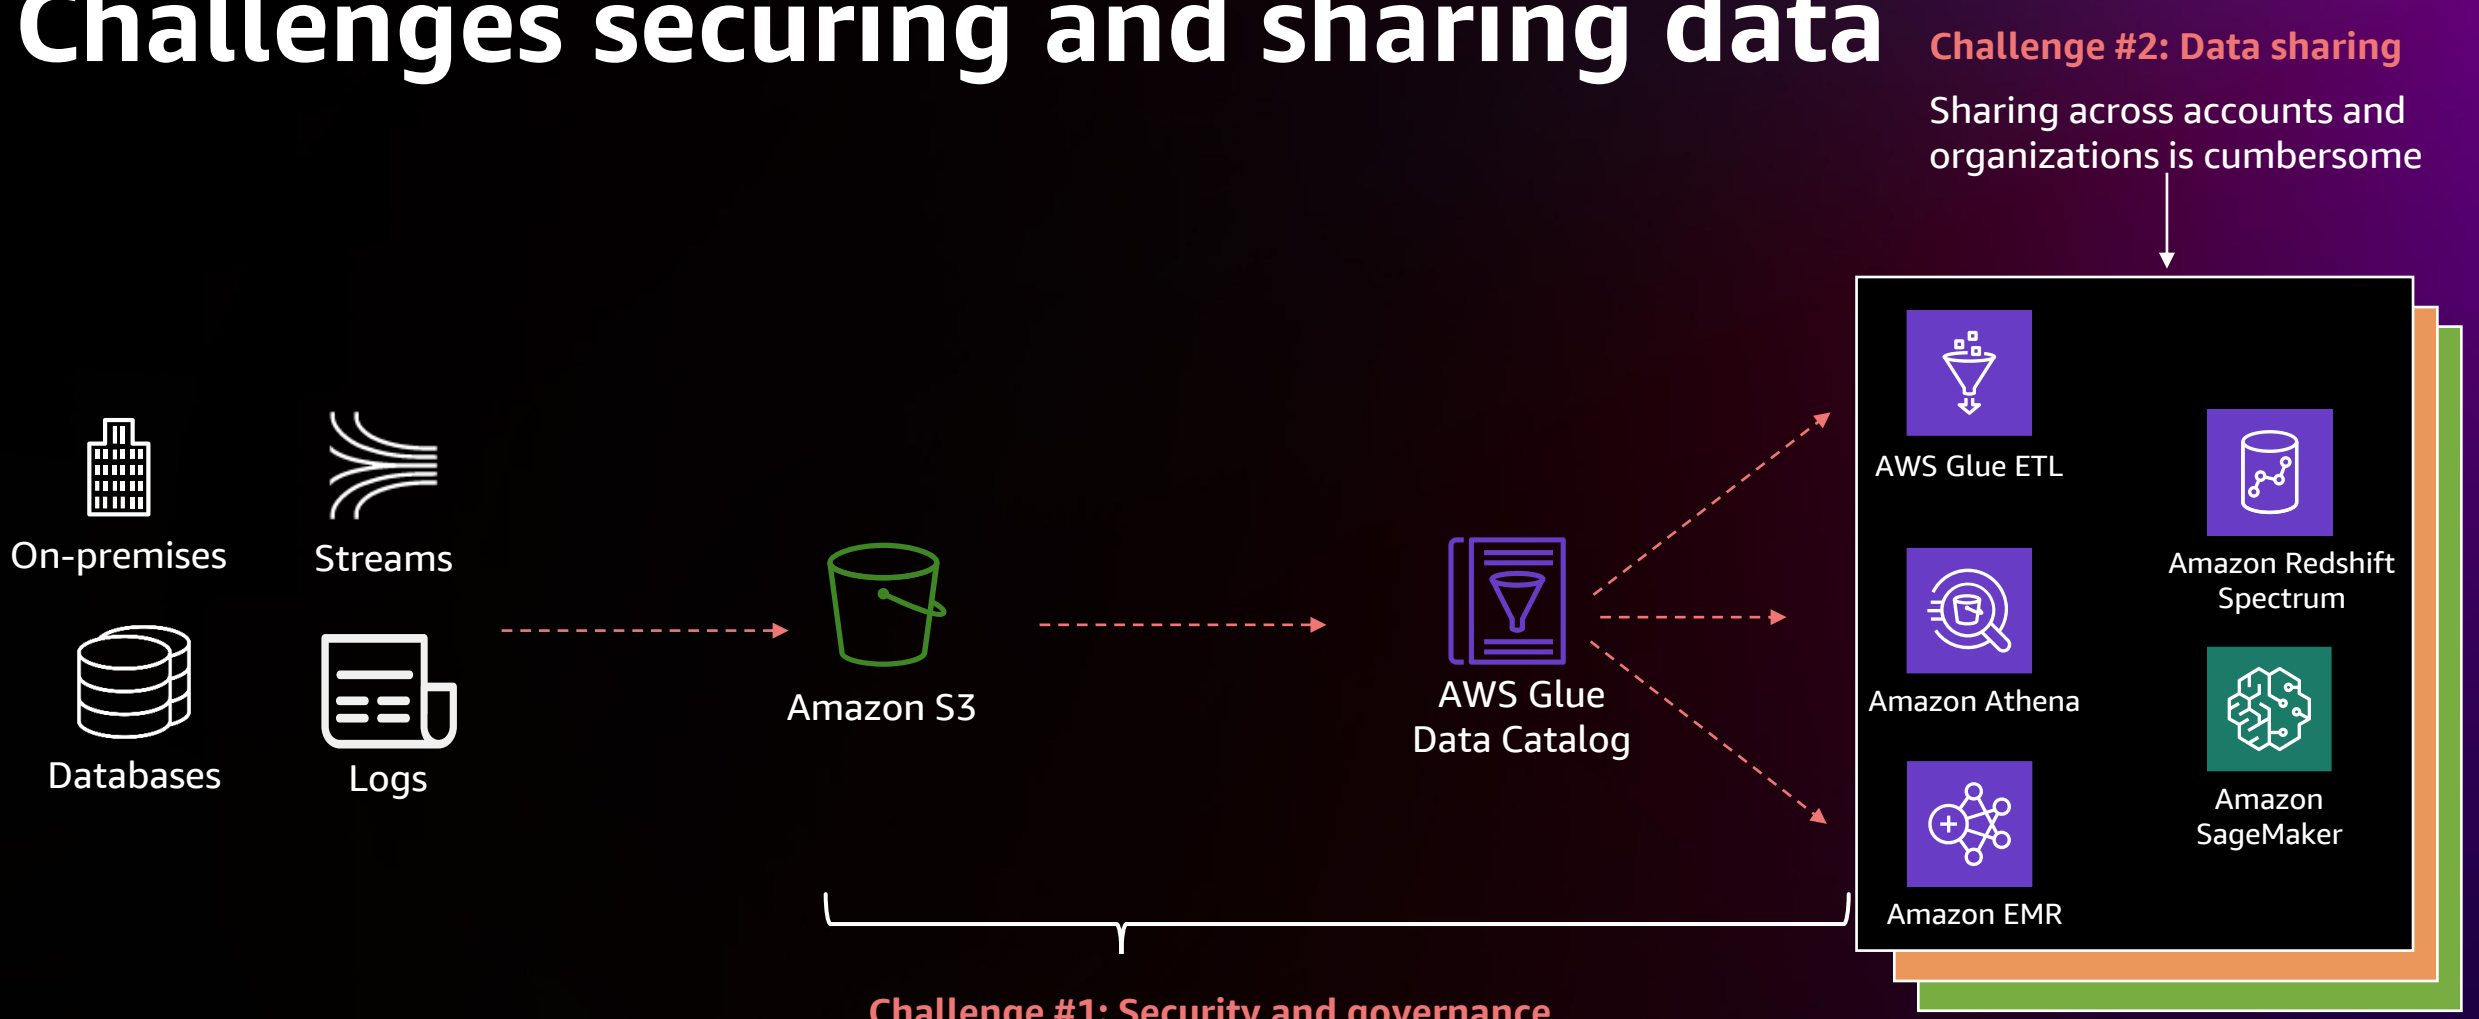

Why is managing data lake permissions hard?

Unifying permissions across the data lake stack

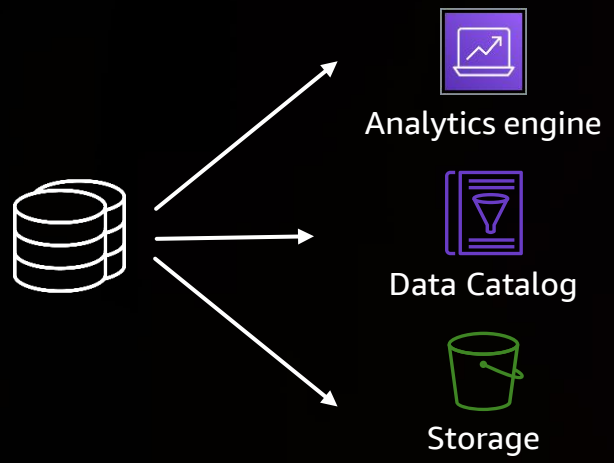

Split storage, metadata, and compute

Each system has different permissions

Syncing permissions is error-prone

Enforcing fine-grained permissions to restrict access

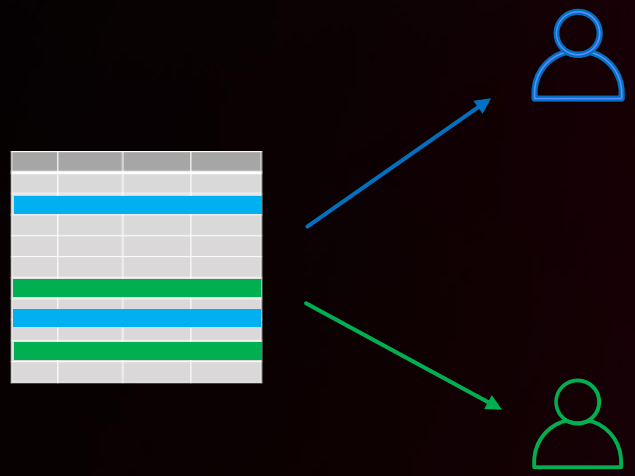

Data lakes contain a lot of data

Users should only access portions

Thousands of resources and tens of thousands of users

Ensuring that data access complies with regulations

Democratize data access

Regulations and governance

Monitor and audit data access

Lake Formation permissions model

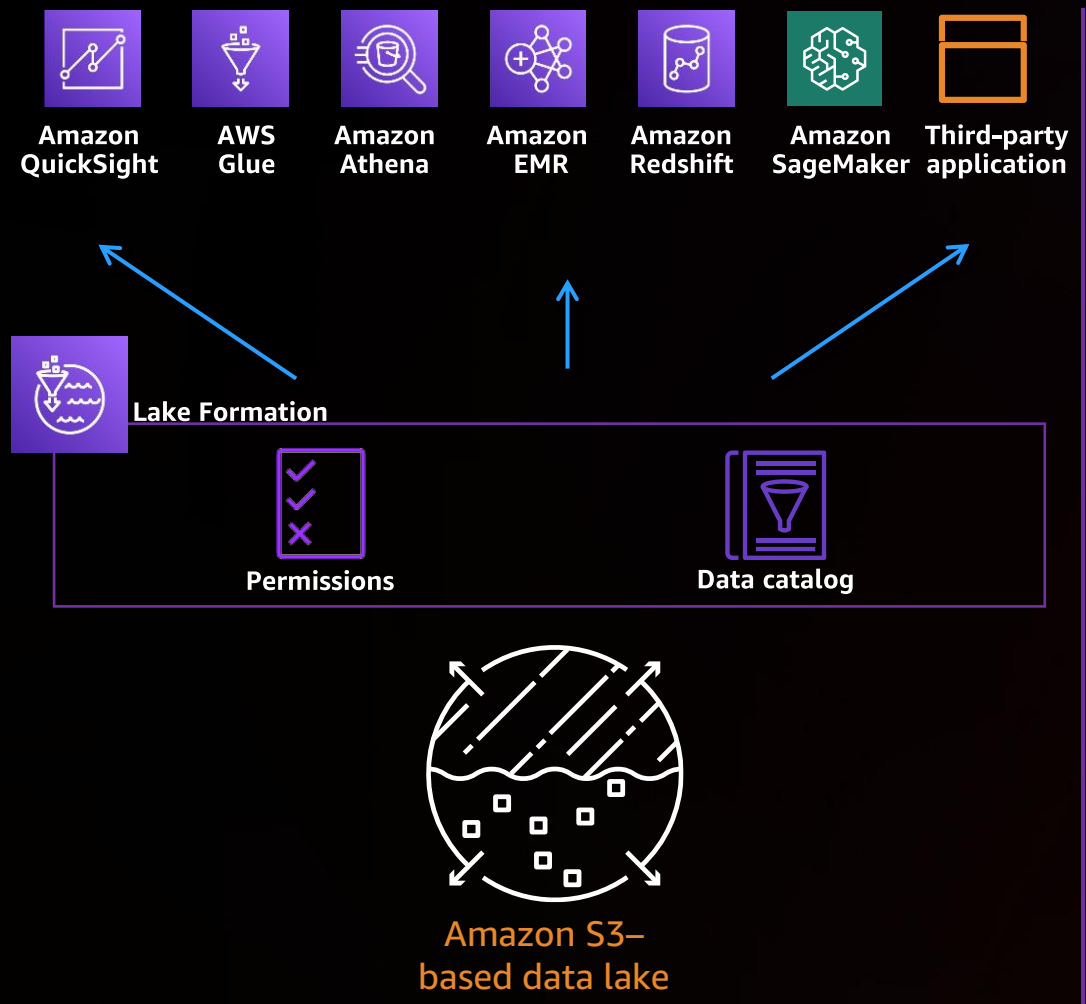

DB-style fine-grained permissions on resources

Scale permissions management Lake Formation Tag-Based Access Control (LF-TBAC)

Unified Amazon S3 permissions

Integrated with services and tools

Easy to audit permissions and access

Challenge: Data sharing

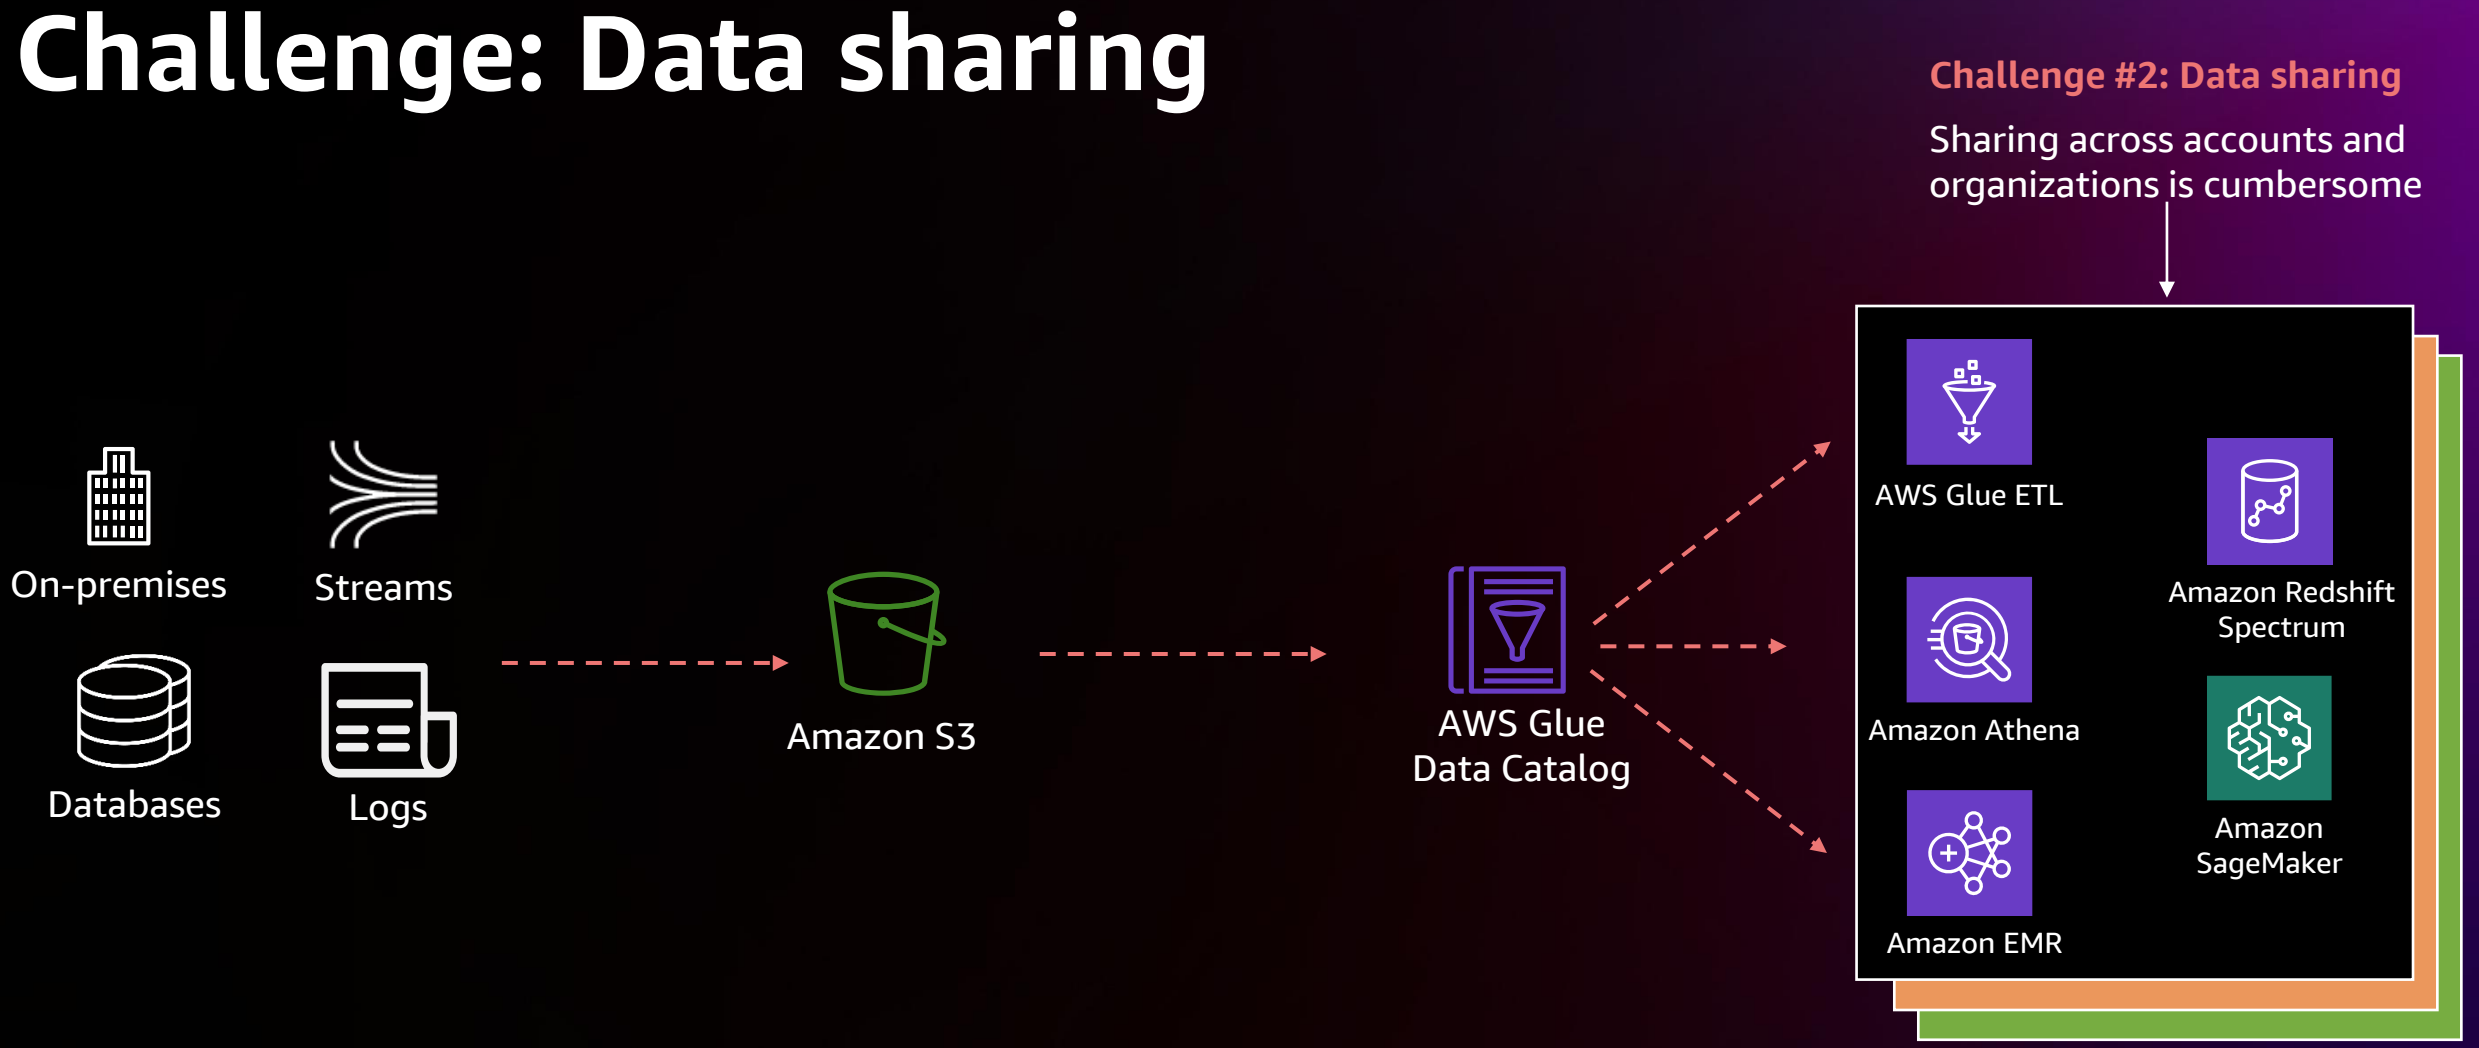

Why is sharing data across accounts hard?

To share data

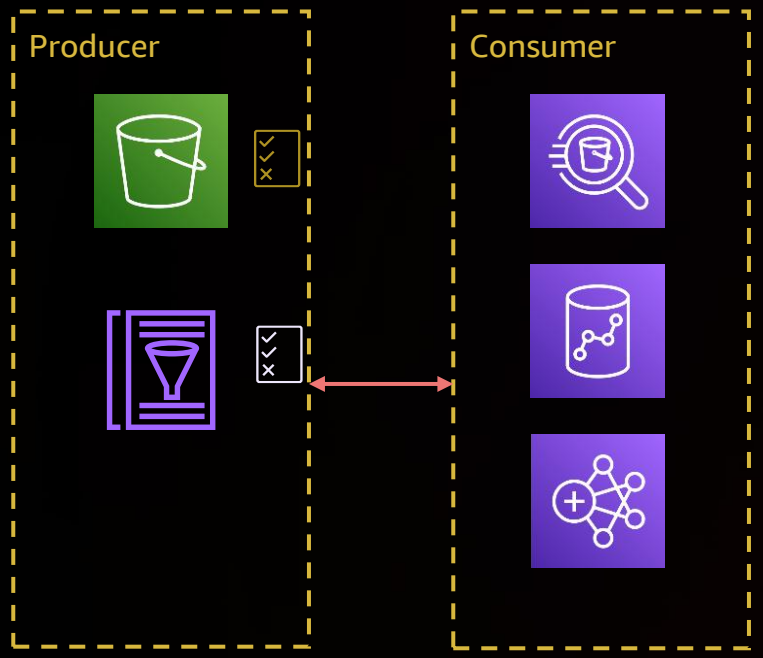

Manges multiple Amazon S3 and IAM policies

Lacks discoverability

Policy size limits (coarse grained)

Duplicating data

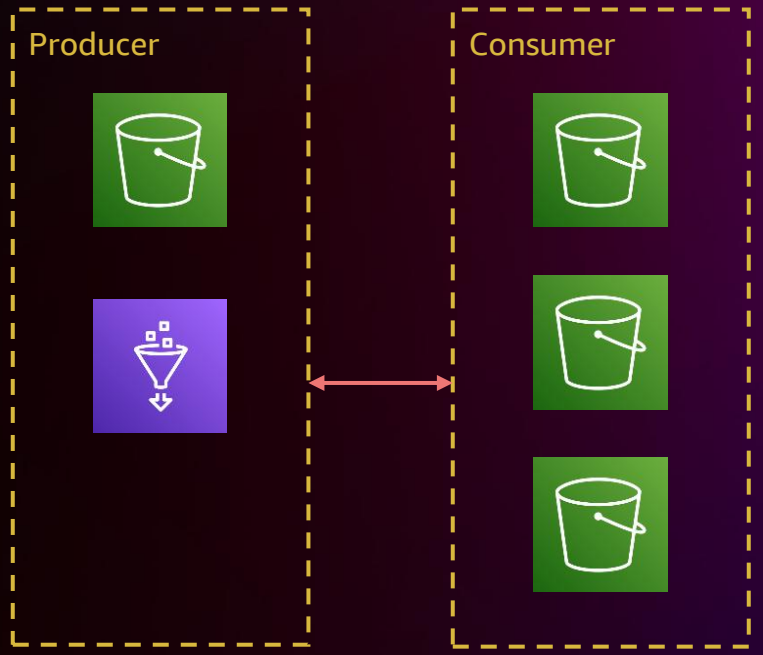

ETL pipelines

Multiple redacted copies

Expensive, brittle, and error-prone

Common data sharing patterns

Single account

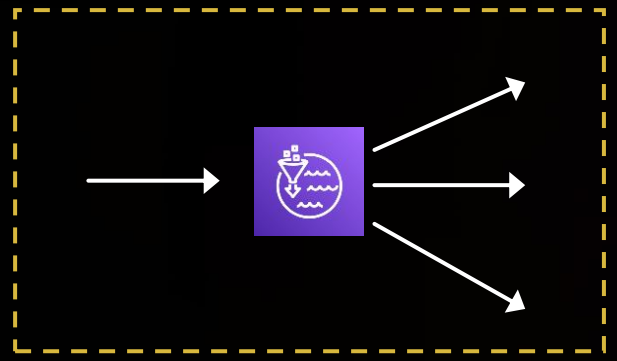

Centralized single account

Simple to get started

Centralized solution (Hub and spoke)

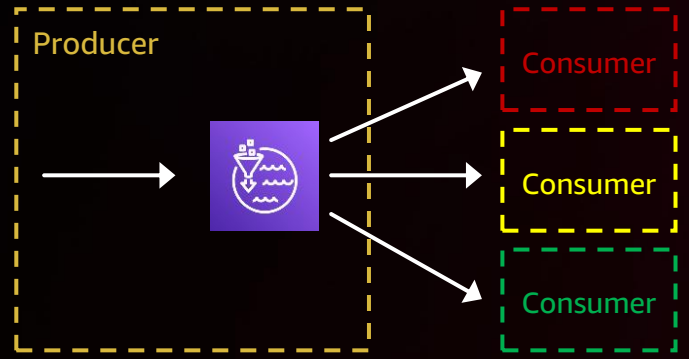

Hub and spoke multi-account

Cross-organization

Data mesh

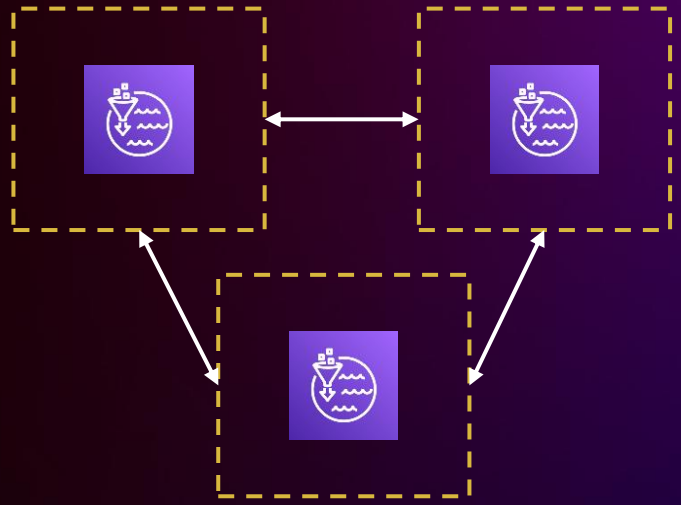

Data mesh central governance

Organizational autonomy

Centralized solution challenges

Central data team becomes the bottleneck

Fail to scale data consumers

Unable to discover and consume the data

No central Data Governance

Lack of data auditability

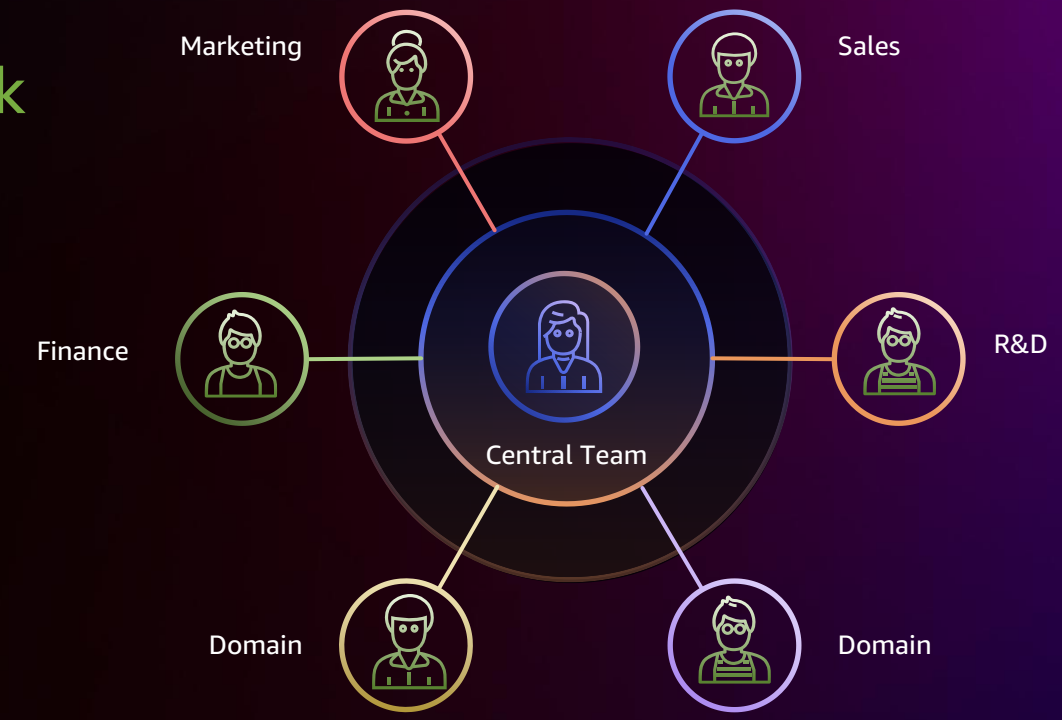

Why data mesh?

Treats existing data platforms as **independent** domains

Improves **data governance** by pushing access policy to data domains

Establishes a **central** mechanism for **data discovery**

Provides **self-service data sharing** features

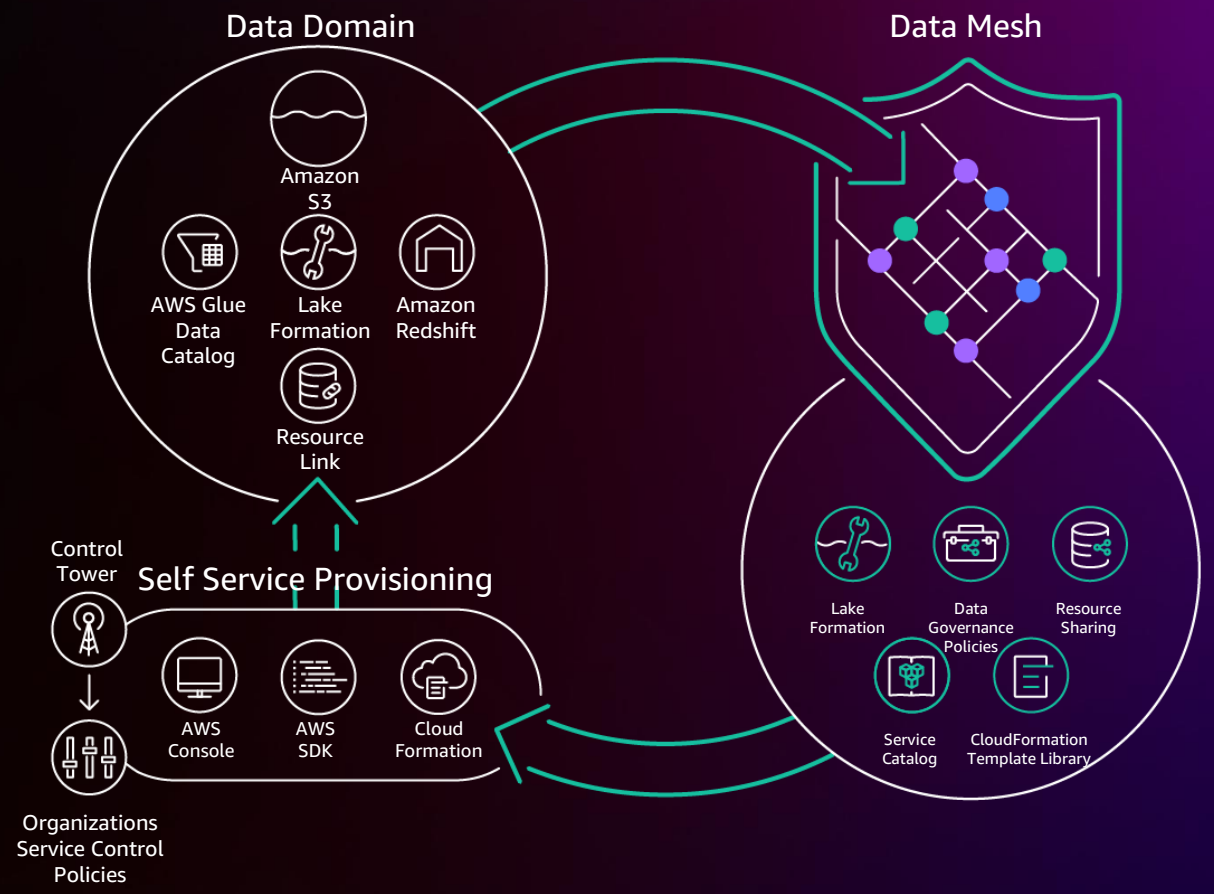

Data mesh – Four core principles

Data Owner

Data Domain Ownership

Data Steward

Federated Computational Governance

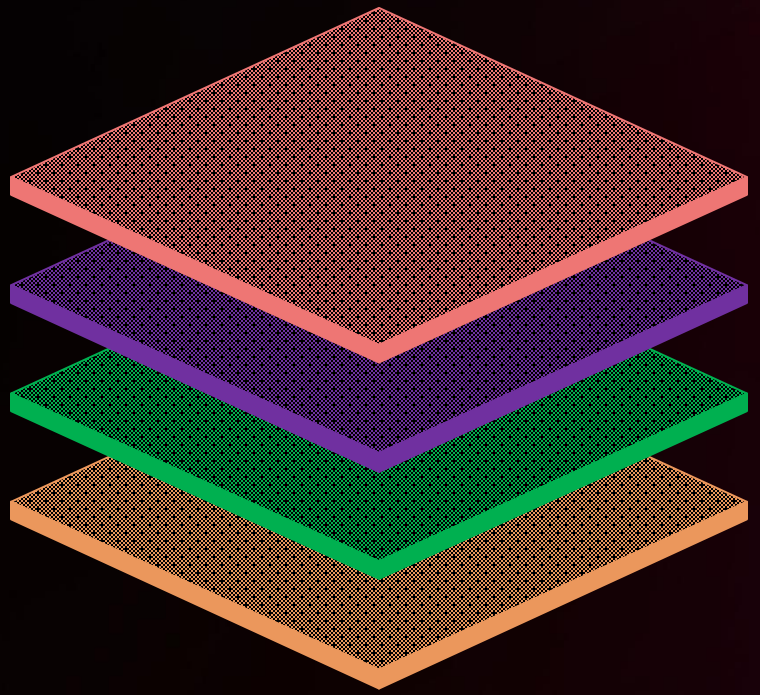

Data Engineer

Data as a Product

Data Consumer

Self-Serve Sharing

Data mesh architecture

DECENTRALIZED, LIGHTWEIGHT FEDERATED GOVERNANCE ACROSS DOMAIN-ORIENTED DATA SYSTEMS TO DRIVE GOVERNED SHARING

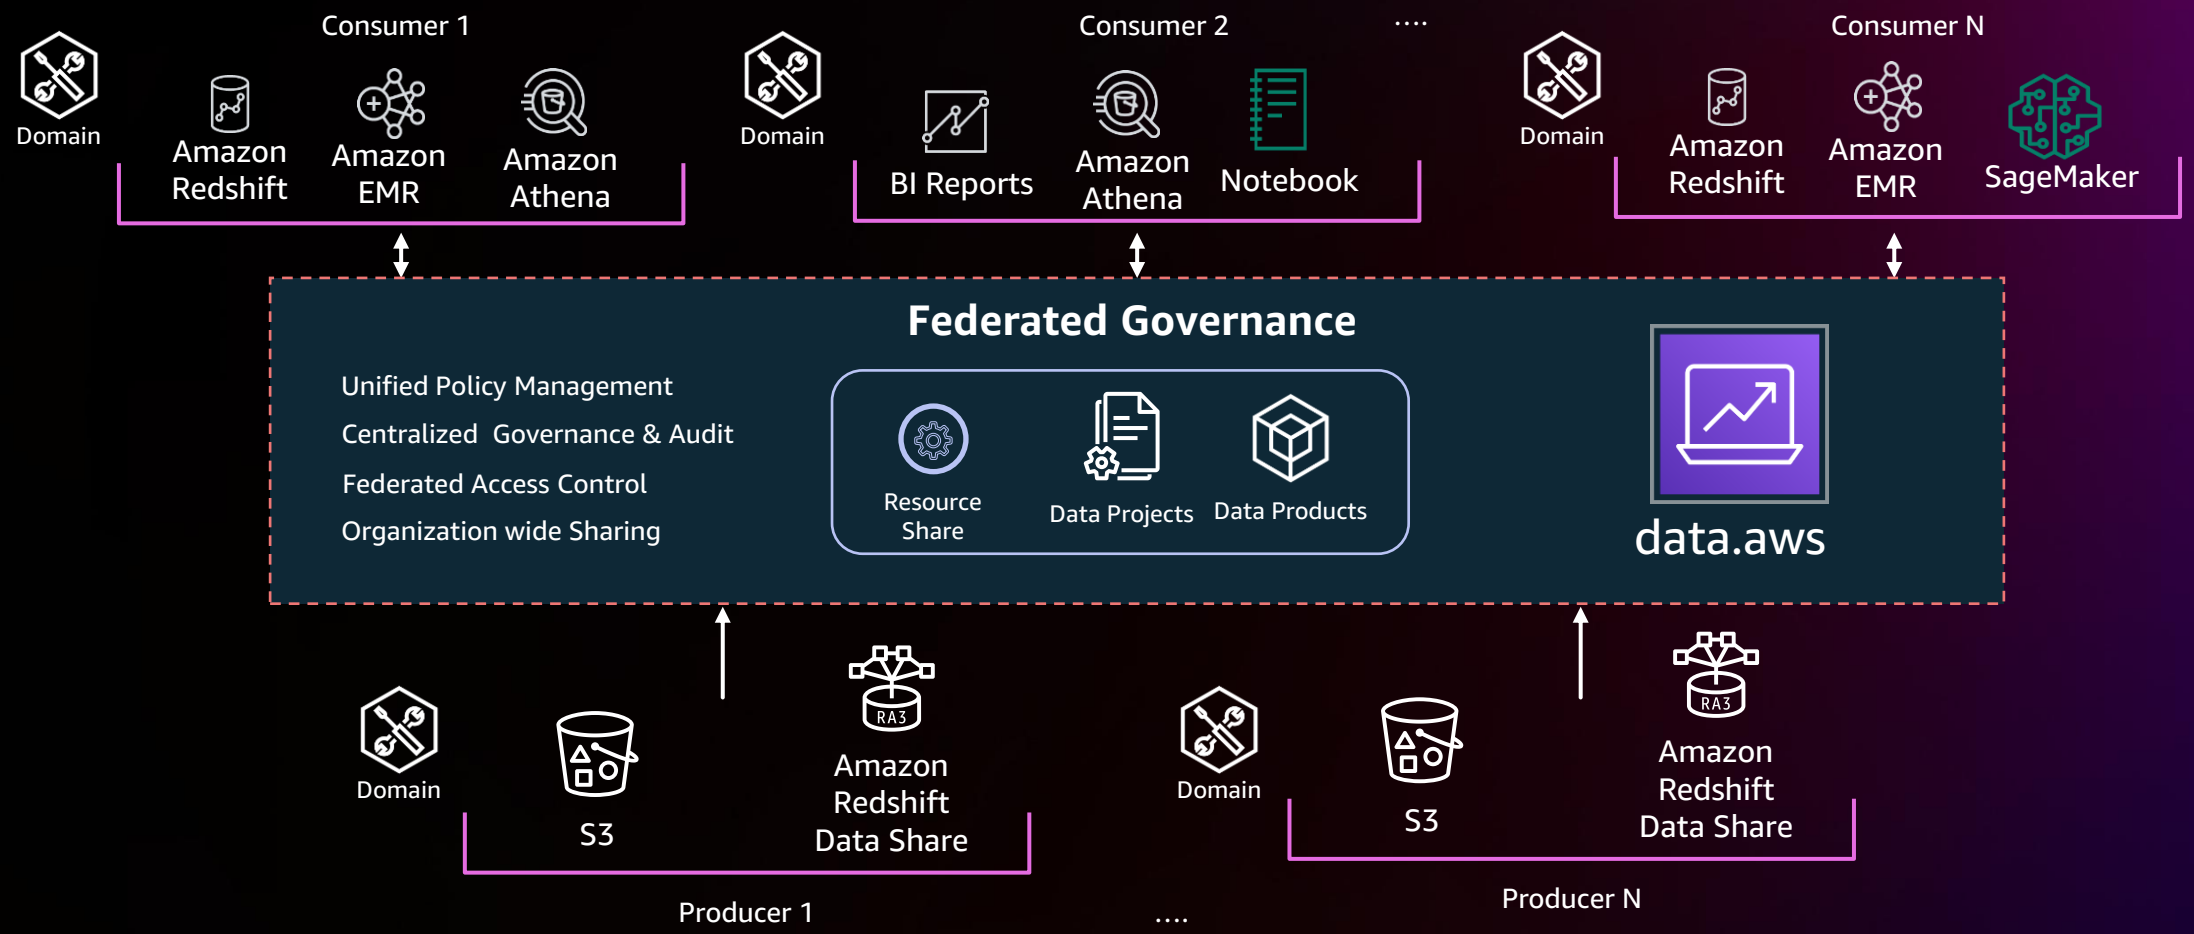

Data mesh architecture pattern: data lake assets collections sharing

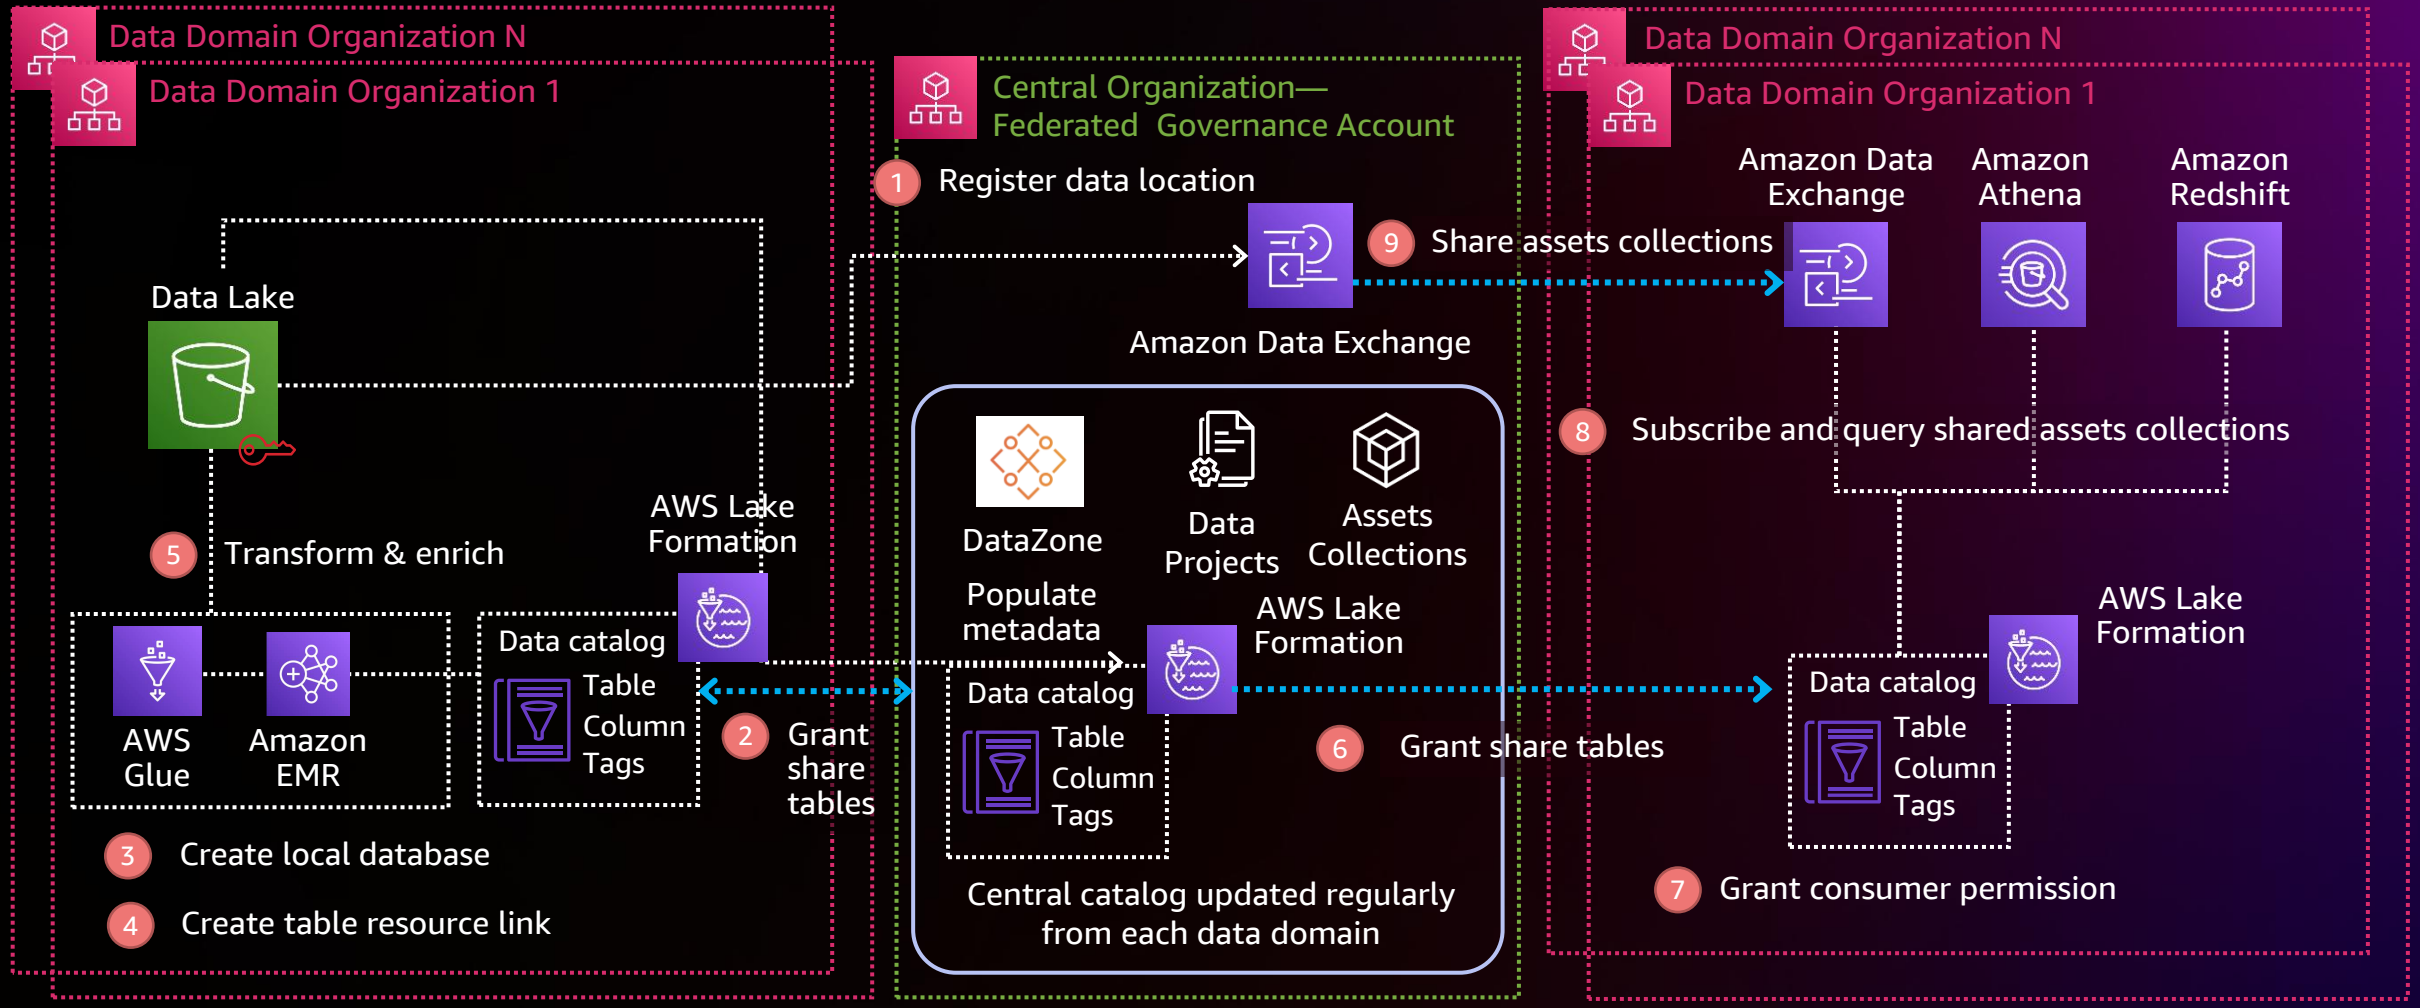

Security, Governance and Data Access with Data Mesh

Data mesh **unifies** Security, Governance, and Data Access of modern data architecture

Agenda

Modern data architecture on AWS

End-to-end data life cycle on the modern data architecture

Data governance and data mesh in action

Journey towards modern data architecture

Journey towards modern data architecture

Start small

Scale fast

Think big

Journey towards modern data architecture

Start small

Start from **small subset**
in your data platform

Journey towards modern data architecture

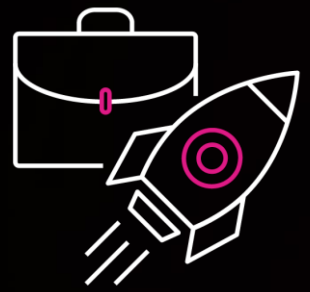

Scale fast

Scale **fast** to achieve your business goal based on data

Journey towards modern data architecture

Think big

Expand your data platform
to support **more workload**
and **advanced goals**

Conclusion

Unified analytics

Highest performance at the lowest cost

Machine learning integration

Unified data access, security and governance

Insights for everyone

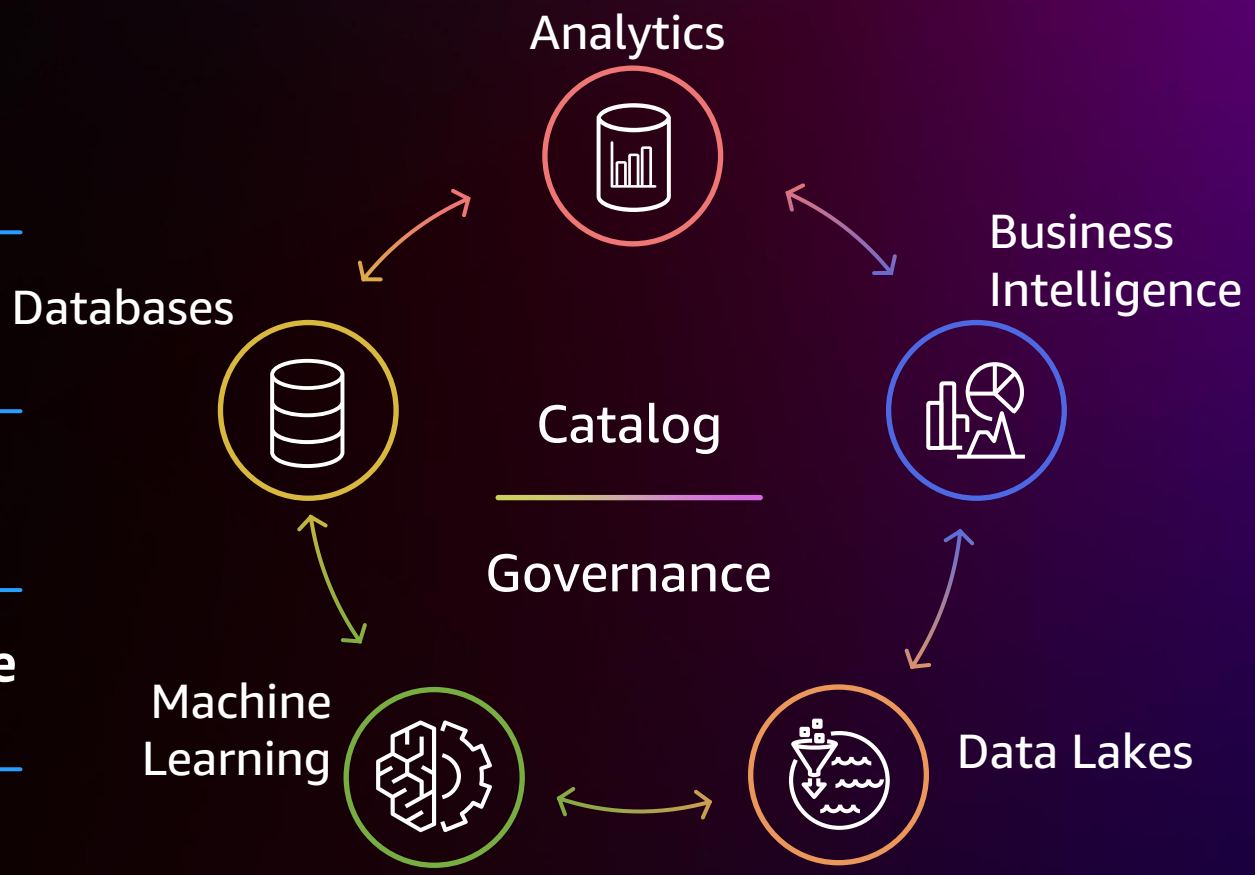

Learn more at re:Invent 2022

SESSIONS RELATED TO MODERN DATA ARCHITECTURE

Swami Sivasubramanian, Vice President of AWS Data and Machine Learning – Keynote
Wednesday November 30 | 8:30 AM – 10:30 AM PST | The Venetian

ANT203-L (LVL 200) Unlock the value of your data with AWS analytics
Wednesday November 30 | 2:30 PM – 3:30 PM PST | The Venetian

ANT223 (LVL 200) Simplify and accelerate data integration & ETL modernization with AWS Glue
Wednesday November 30 | 12:15 PM – 1:15 PM PST | MGM Chairmans 368

ANT310 (LVL 300) Build a data mesh with AWS Lake Formation and AWS Glue
Wednesday November 30 | 05:30 PM – 07:30 PM PST | MGM Grand

ANT344 (LVL 300) Democratize data with governance – Connect people, data, and tools with Amazon DataZone
Wednesday November 30 | 02:30 PM – 03:30 PM PST | MGM Grand

Thank you!

Santosh Chandrachood
sanchas@amazon.com

Please complete the session survey in the **mobile app**