
AWS
re:Invent
NOV. 28 – DEC. 2, 2022 | LAS VEGAS, NV

You like...

Vacation at...

Movie: action or horror?

Bike or hike?

Early bird or night owl?

Pineapple!?


AWS
re:Invent
NOV. 28 – DEC. 2, 2022 | LAS VEGAS, NV

AMZ303

Diving deep with Amazon Ads: Analytics at scale

Tom Skinner

Director,
Amazon Ads

Joshua Anghel

Principal Engineer,
Amazon Ads

Varun Kamalakaran

Principal, Customer Solutions
AWS

Agenda

Introduction

- Amazon advertising + Ads measurement team

Analytics at the core

- Analytics lifecycle
- Challenges, success measures, and key indications
- AWS collaboration

Technical solution – deep dive

- Big picture
- Architecture
  - Orchestration
  - Data applications
  - Reporting engines and APIs
- Lessons learned

amazon ads

amazon

alexa

freevee

music

fire tv

NFL

prime video

twitch

WHOLE FOODS MARKET

Amazon Ads Products

An array of advertising offerings to engage audiences and share the brand story

Sponsored Products

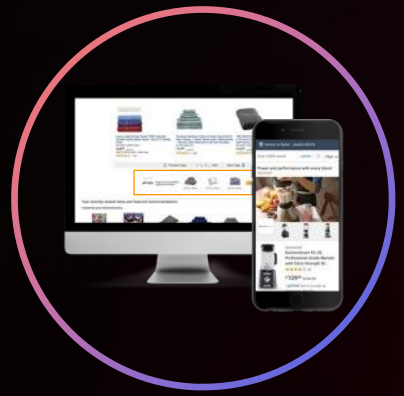

Sponsored Brands

Sponsored Display

Display Ads

OOH Ads

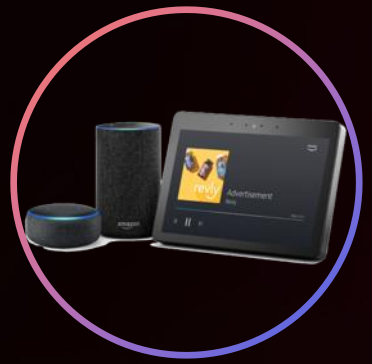

Audio Ads

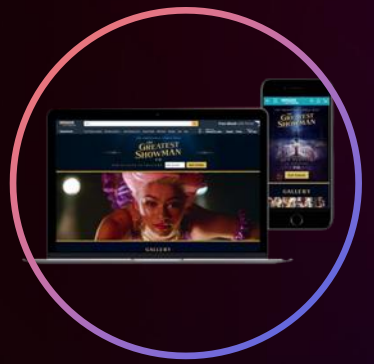

Custom Ads

Video Ads

Ads measurement

100PB+
data

40B+
reports/year

400M
reports/day

1T+
event/week

3M+
TPS

Analytics at the core

Tom Skinner
Director
Amazon Ads

Lifecycle—advertising analytics

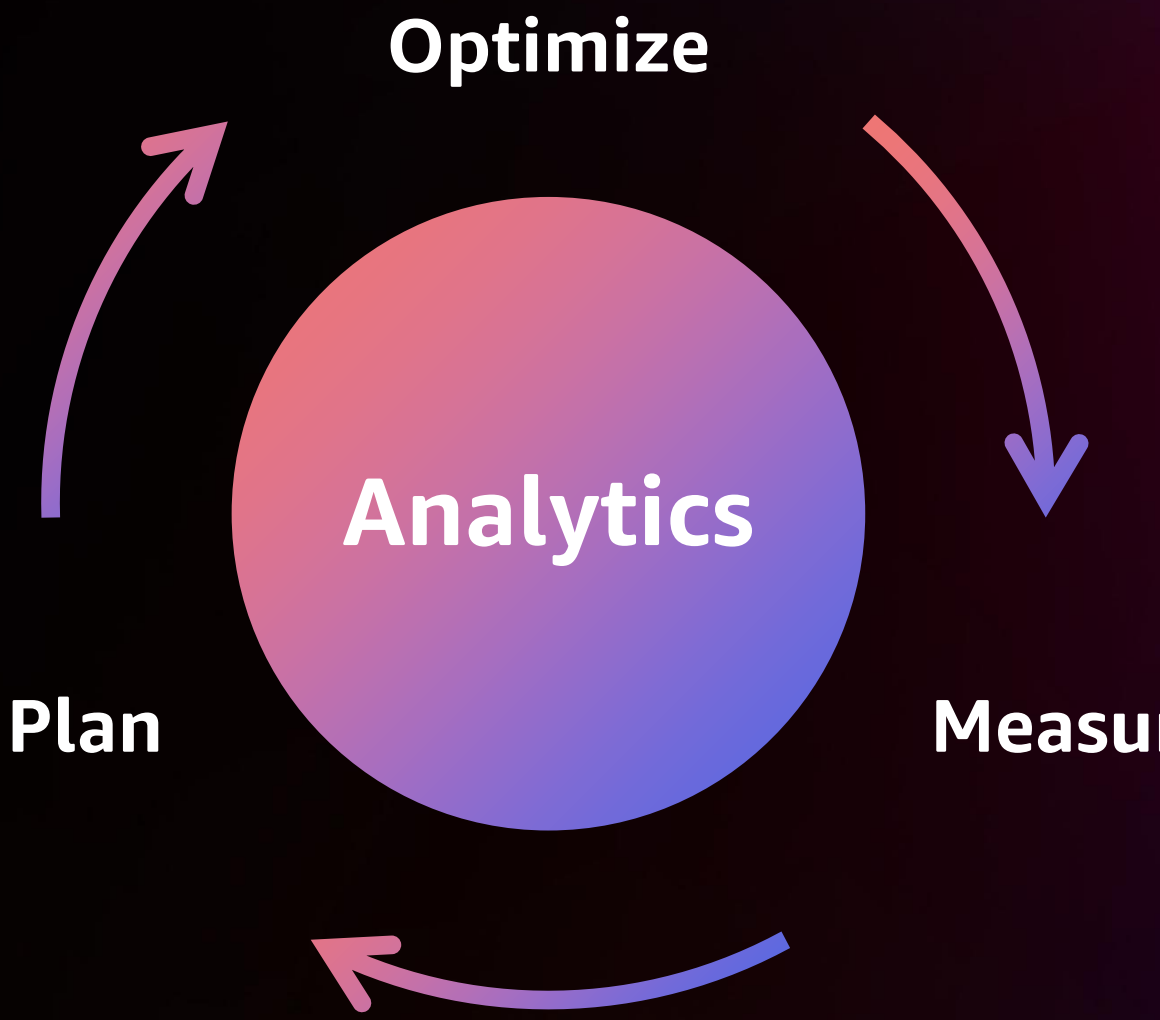

Data flows – Advertising analytics

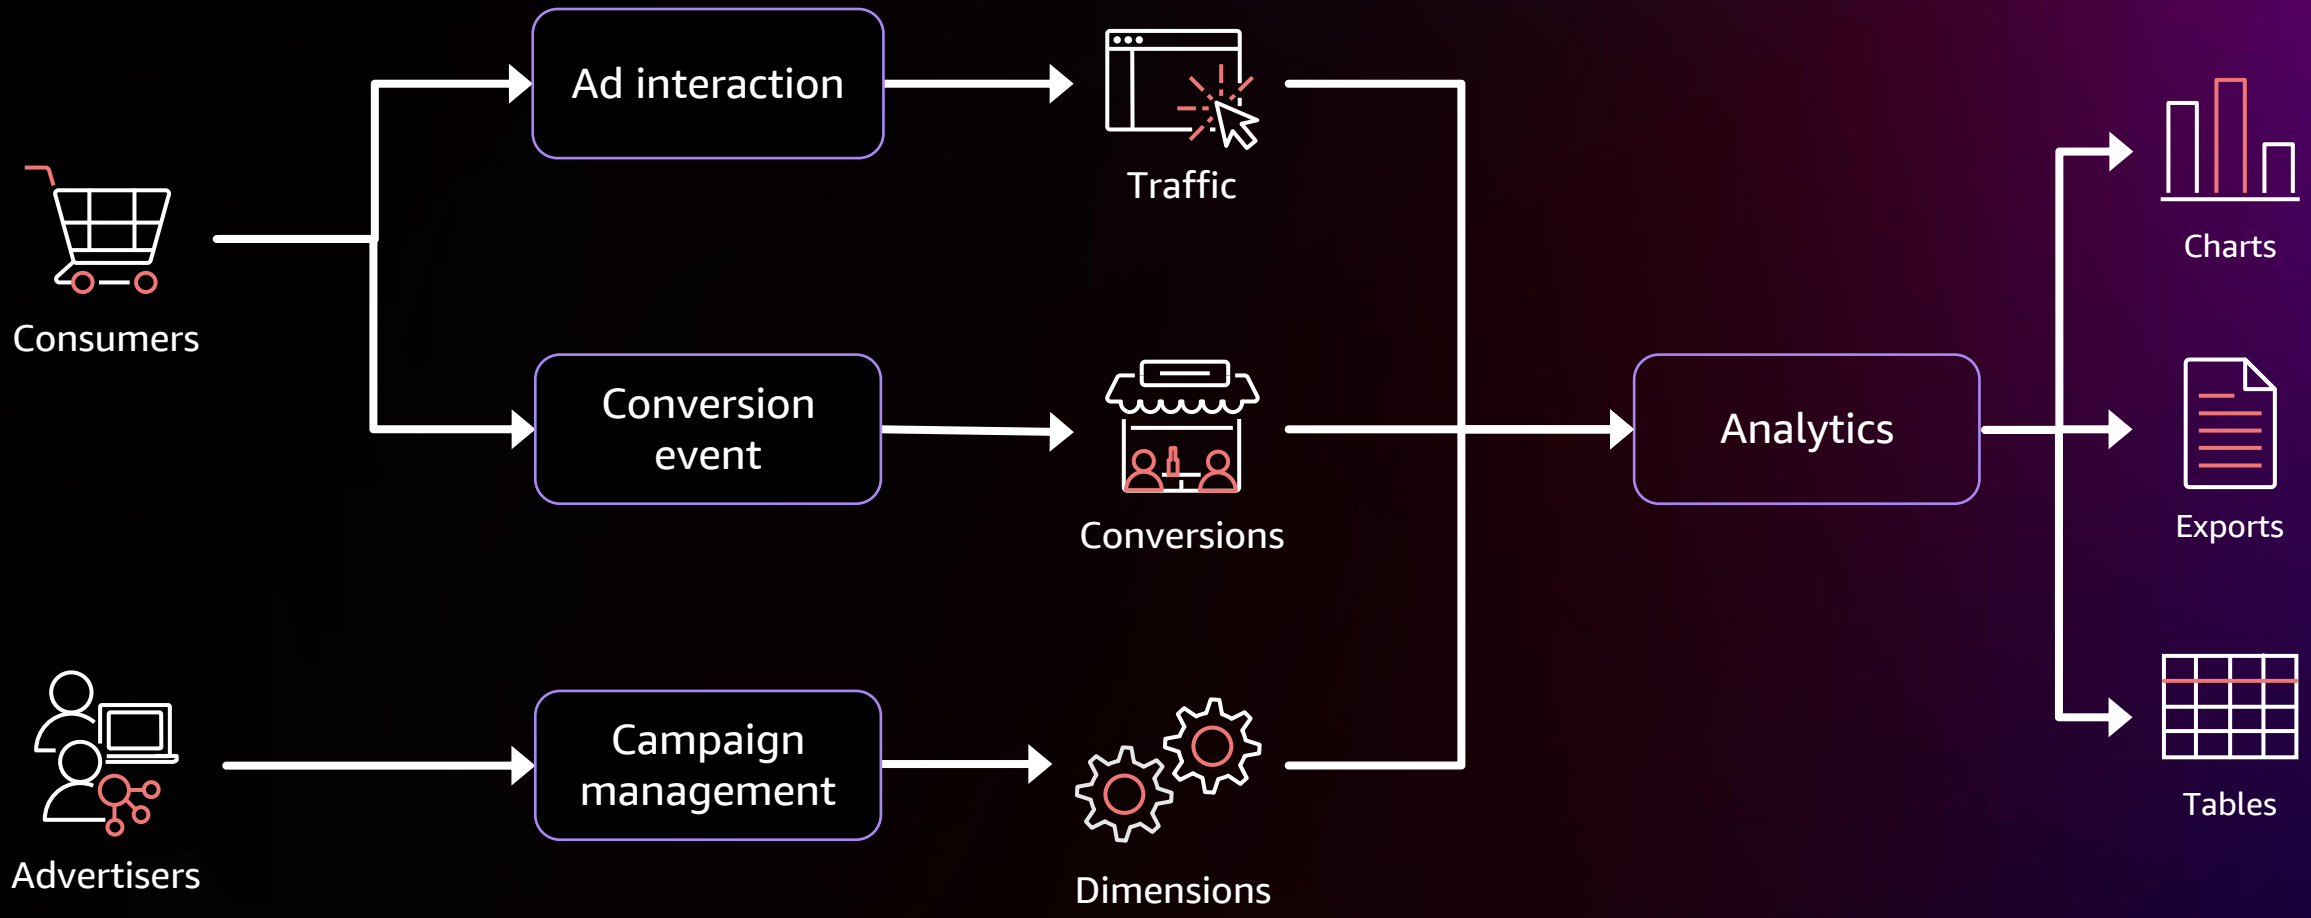

Challenges

Critical to customer experience

Exactly once processing

Late arriving data

Long lookback calculations

Focus on the inputs, measure outputs

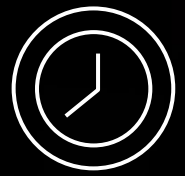

Availability

Data quality

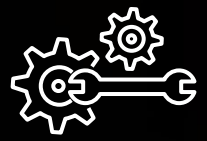

Operational load

Developer velocity

Cost efficiency

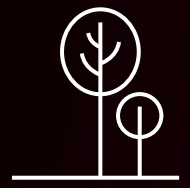

Sustainability

Focus on the inputs, measure outputs

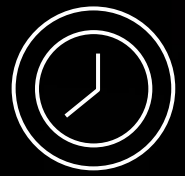

Availability

Data quality

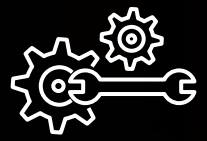

Operational load

Developer velocity

Cost efficiency

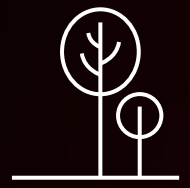

Sustainability

Availability KPIs

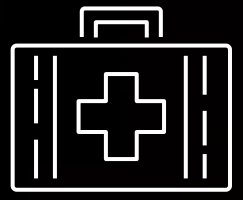

Health check

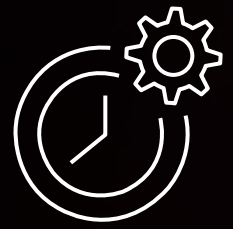

Service availability

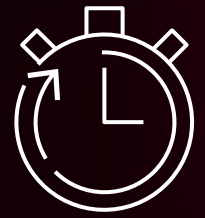

Query latency

Data age

Focus on the inputs, measure outputs

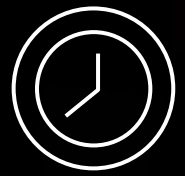

Availability

Data quality

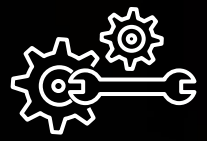

Operational load

Developer velocity

Cost efficiency

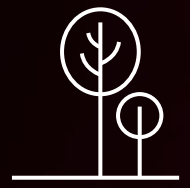

Sustainability

Data quality KPIs

Loss rate

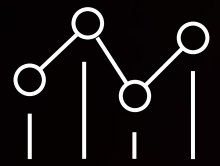

Measurement benchmarks

Anomalies

Customer tickets

Focus on the inputs, measure outputs

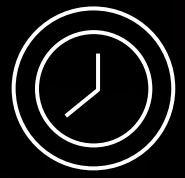

Availability

Data quality

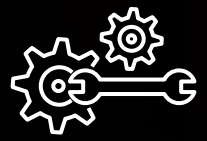

Operational load

Developer velocity

Cost efficiency

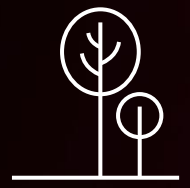

Sustainability

Operational load KPIs

Operational stories/engineer

CI/CD pipeline metrics

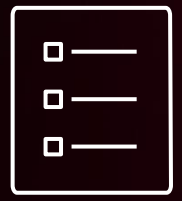

Ticket queue depth

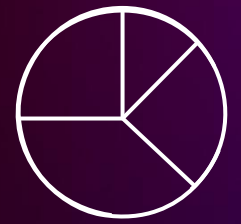

Ticket severity counts

Focus on the inputs, measure outputs

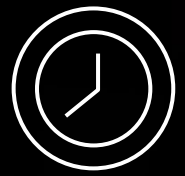

Availability

Data quality

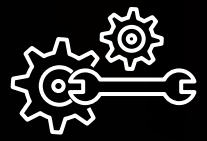

Operational load

Developer velocity

Cost efficiency

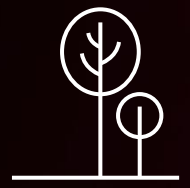

Sustainability

Developer velocity KPIs

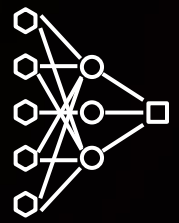

Complexity

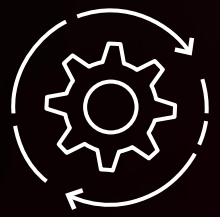

Automation percentage

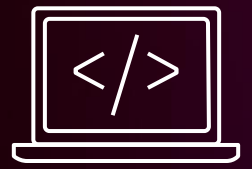

Developer days

Focus on the inputs, measure outputs

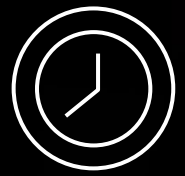

Availability

Data quality

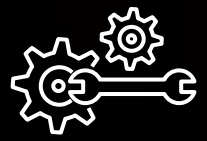

Operational load

Developer velocity

Cost efficiency

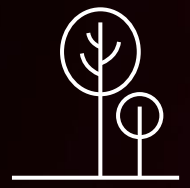

Sustainability

Cost efficiency KPIs

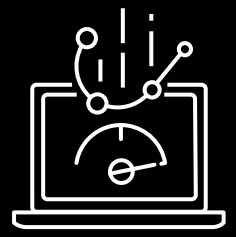

System utilization

Year-over-year trends

Data processing cost

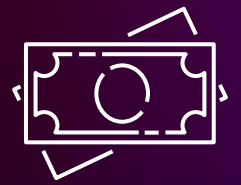

Cost/user action

Focus on the inputs, measure outputs

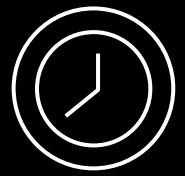

Availability

Data quality

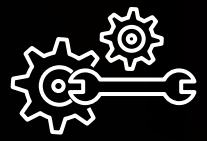

Operational load

Developer velocity

Cost efficiency

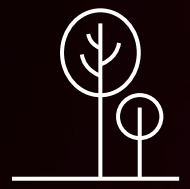

Sustainability

Sustainability KPIs

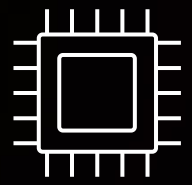

Instance type selection

Region selection

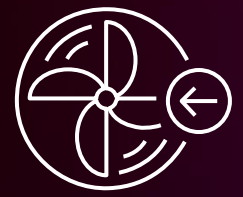

Carbon footprint

Advertising – AWS services


Amazon EBS


Amazon S3

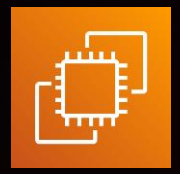
Amazon EC2


AWS Lambda


AWS Glue


Amazon Kinesis


Amazon ECR


Amazon EKS


Amazon SNS


Amazon SQS


Amazon EMR


Amazon Athena


Amazon ECS


AWS Fargate

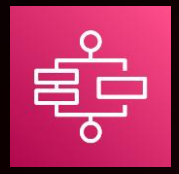
AWS Step Functions


Amazon DynamoDB


Amazon CloudWatch

Technical solution

Josh Anghel
Principal Engineer
Amazon Ads

Big picture – Analytics

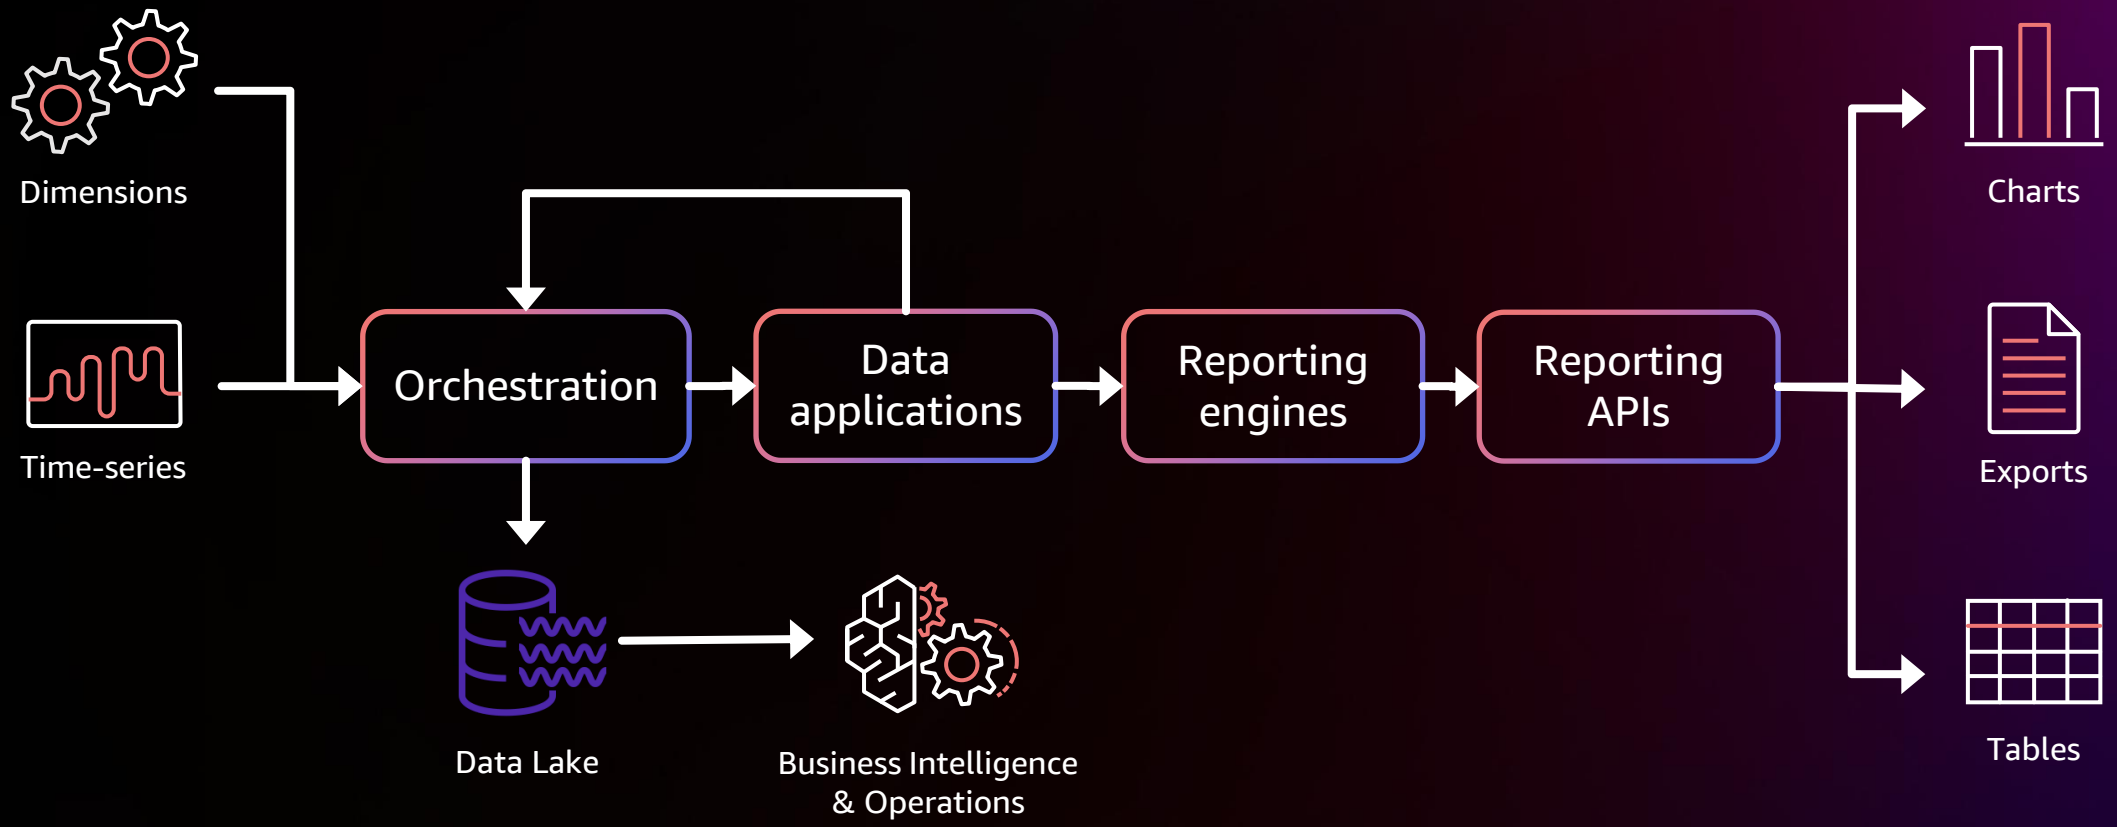

Orchestration

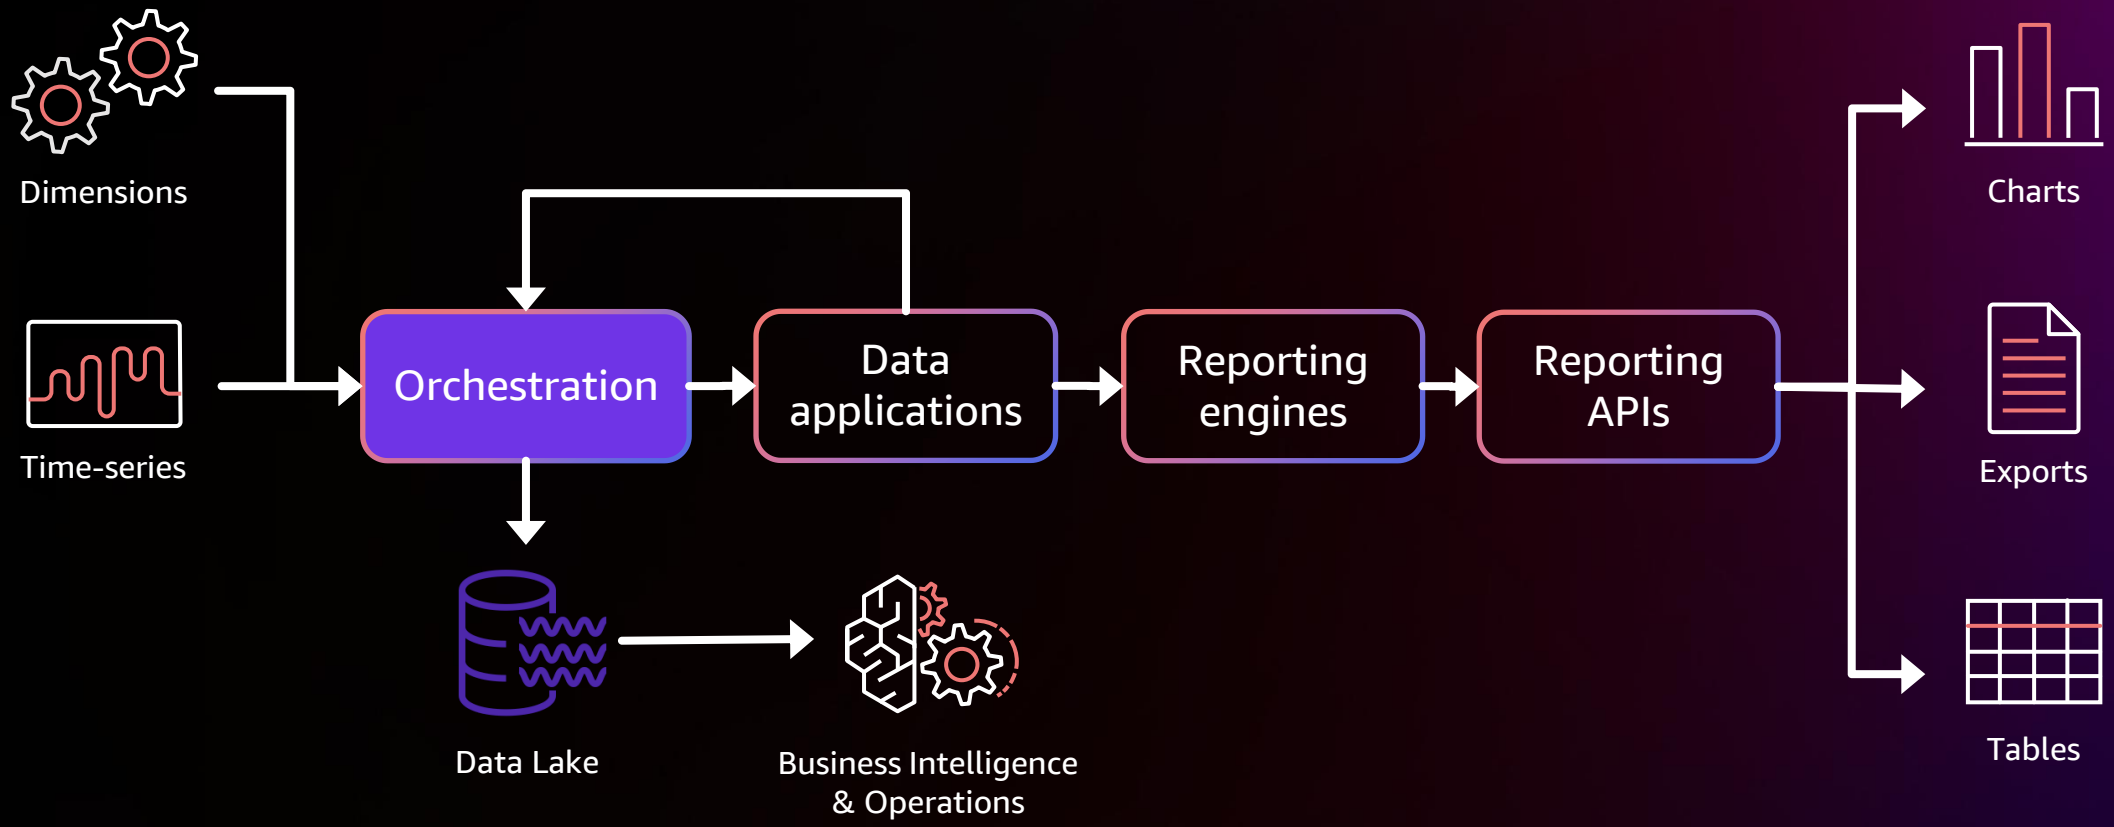

Orchestration

200
clusters

150k
jobs/day

1T+
events/week

Exactly once

Orchestration V1 – Data warehouse

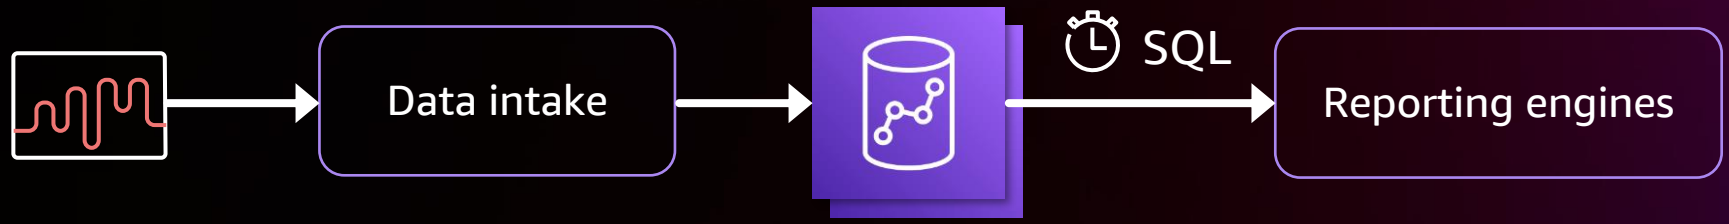

Orchestration V1 – Data warehouse

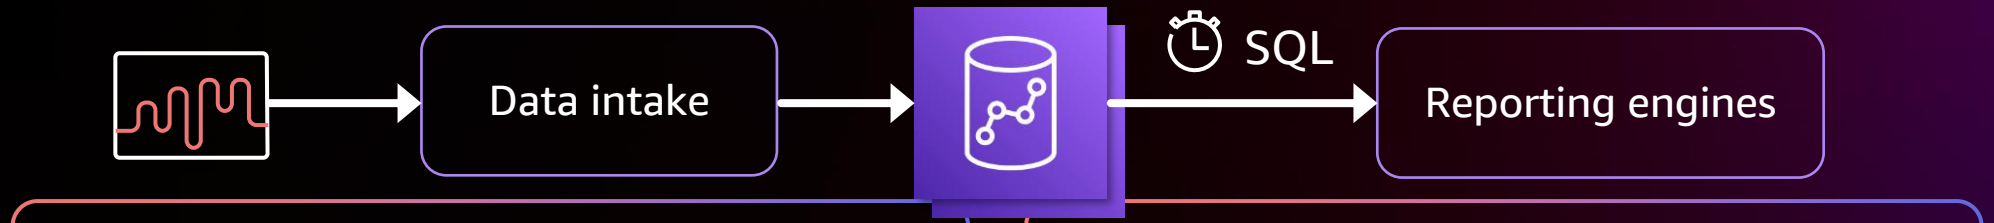

Data is stored on disk

Storage and compute are coupled

Source of record was unclear

Large monolithic clusters

Orchestration V2 – Data lake

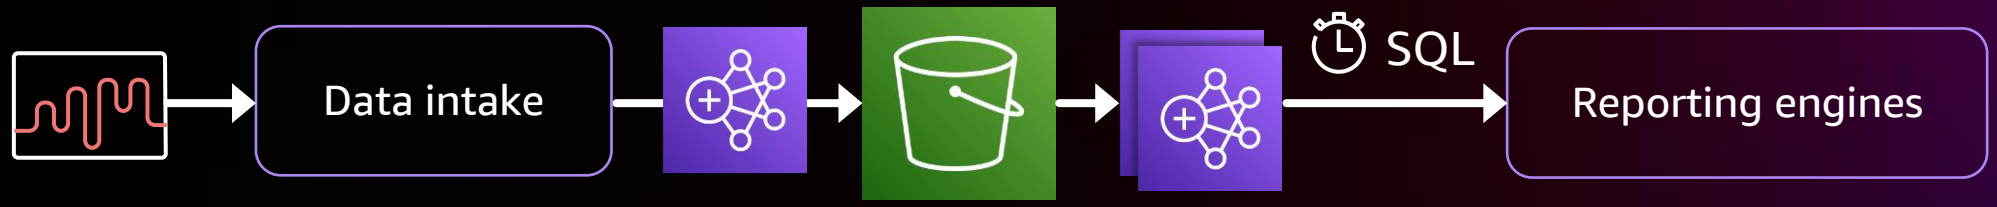

Orchestration V2 – Data lake

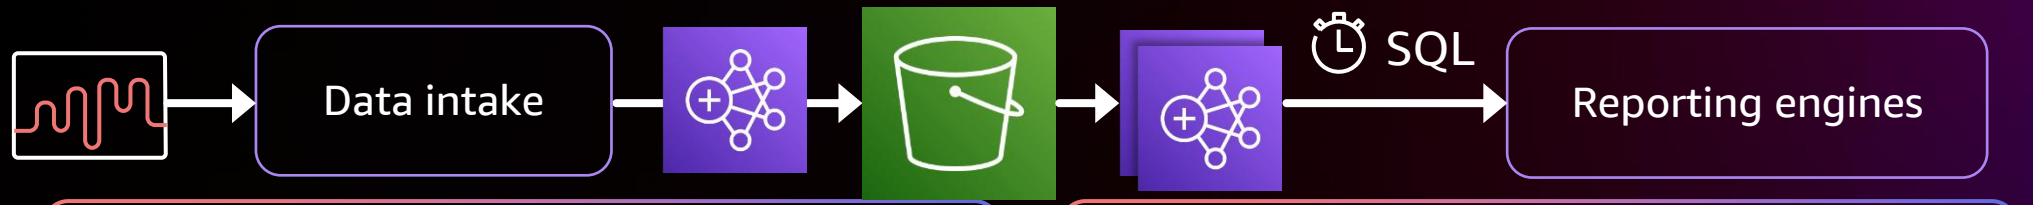

Data is stored in Amazon S3

Separate storage & compute

Single source of record

Optimal cluster sizes

Data lakes

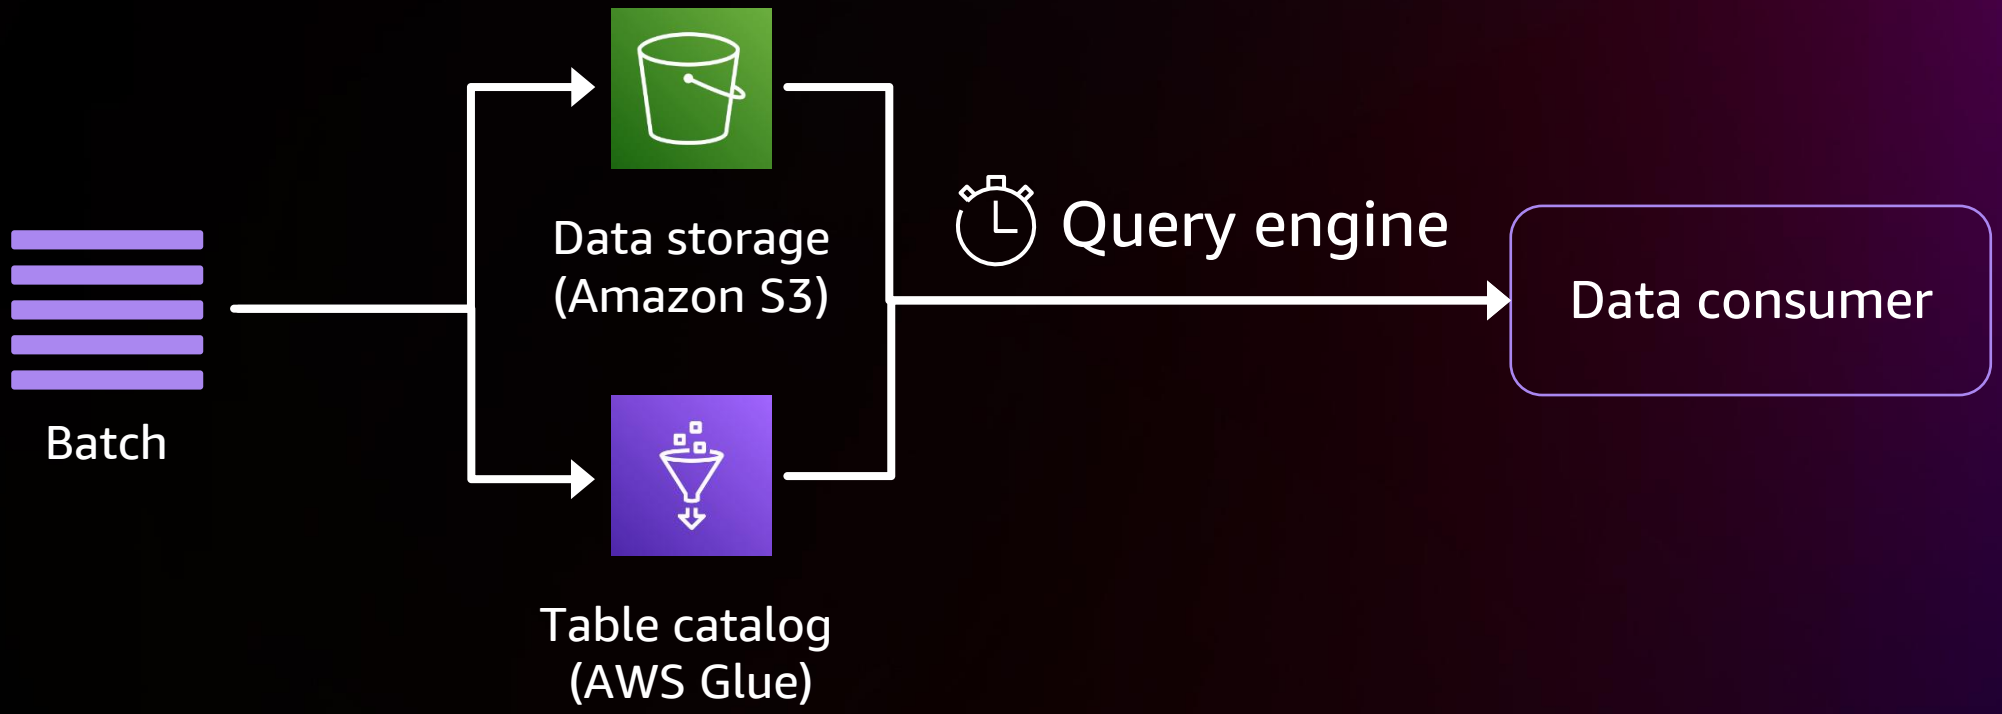

Data streams

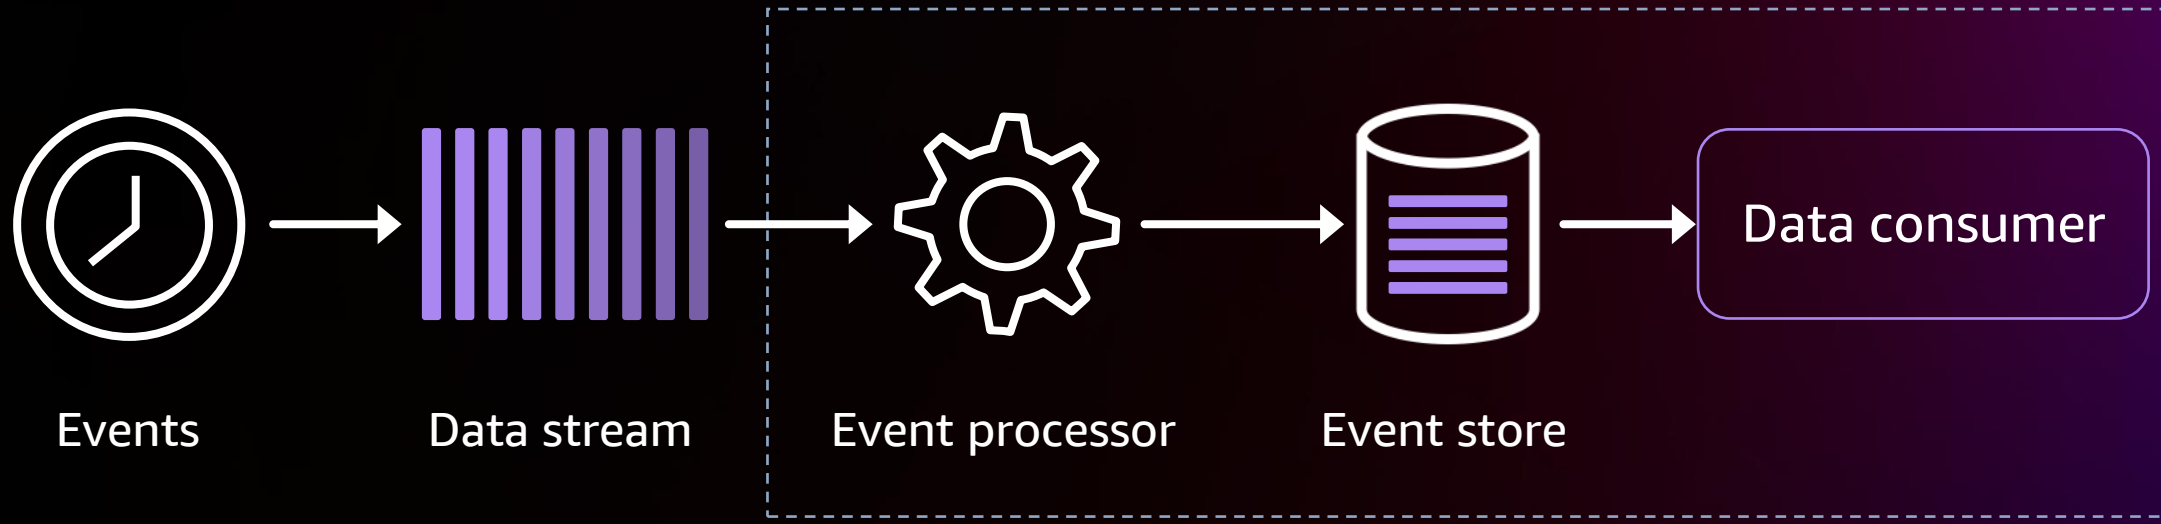

Data streams – Limitations

Event state storage

Developer velocity

Fault tolerance

Lookbacks & aggregation

Cold-starts and backfills

On-demand scaling

Data streams – Limitations

Event state storage

Developer velocity

Fault tolerance

Lookbacks & aggregation

Cold-starts and backfills

On-demand scaling

Data streams – Limitations

Event state storage

Developer velocity

Fault tolerance

Lookbacks & aggregation

Cold-starts and backfills

On-demand scaling

Data streams – Limitations

Event state storage

Developer velocity

Fault tolerance

Lookbacks & aggregation

Cold-starts and backfills

On-demand scaling

Data streams – Limitations

Event state storage

Developer velocity

Fault tolerance

Lookbacks & aggregation

Cold-starts and backfills

On-demand scaling

Data streams – Limitations

Event state storage

Developer velocity

Fault tolerance

Lookbacks & aggregation

Cold-starts and backfills

On-demand scaling

Orchestration V3

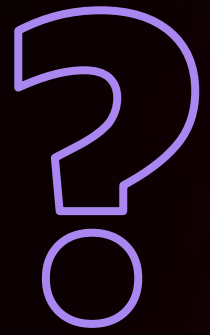

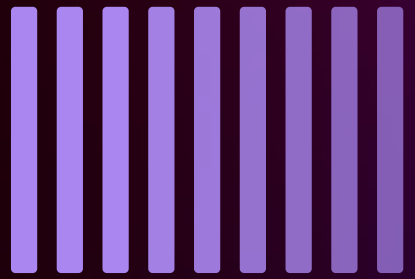

Orchestration V3 – Data rivers

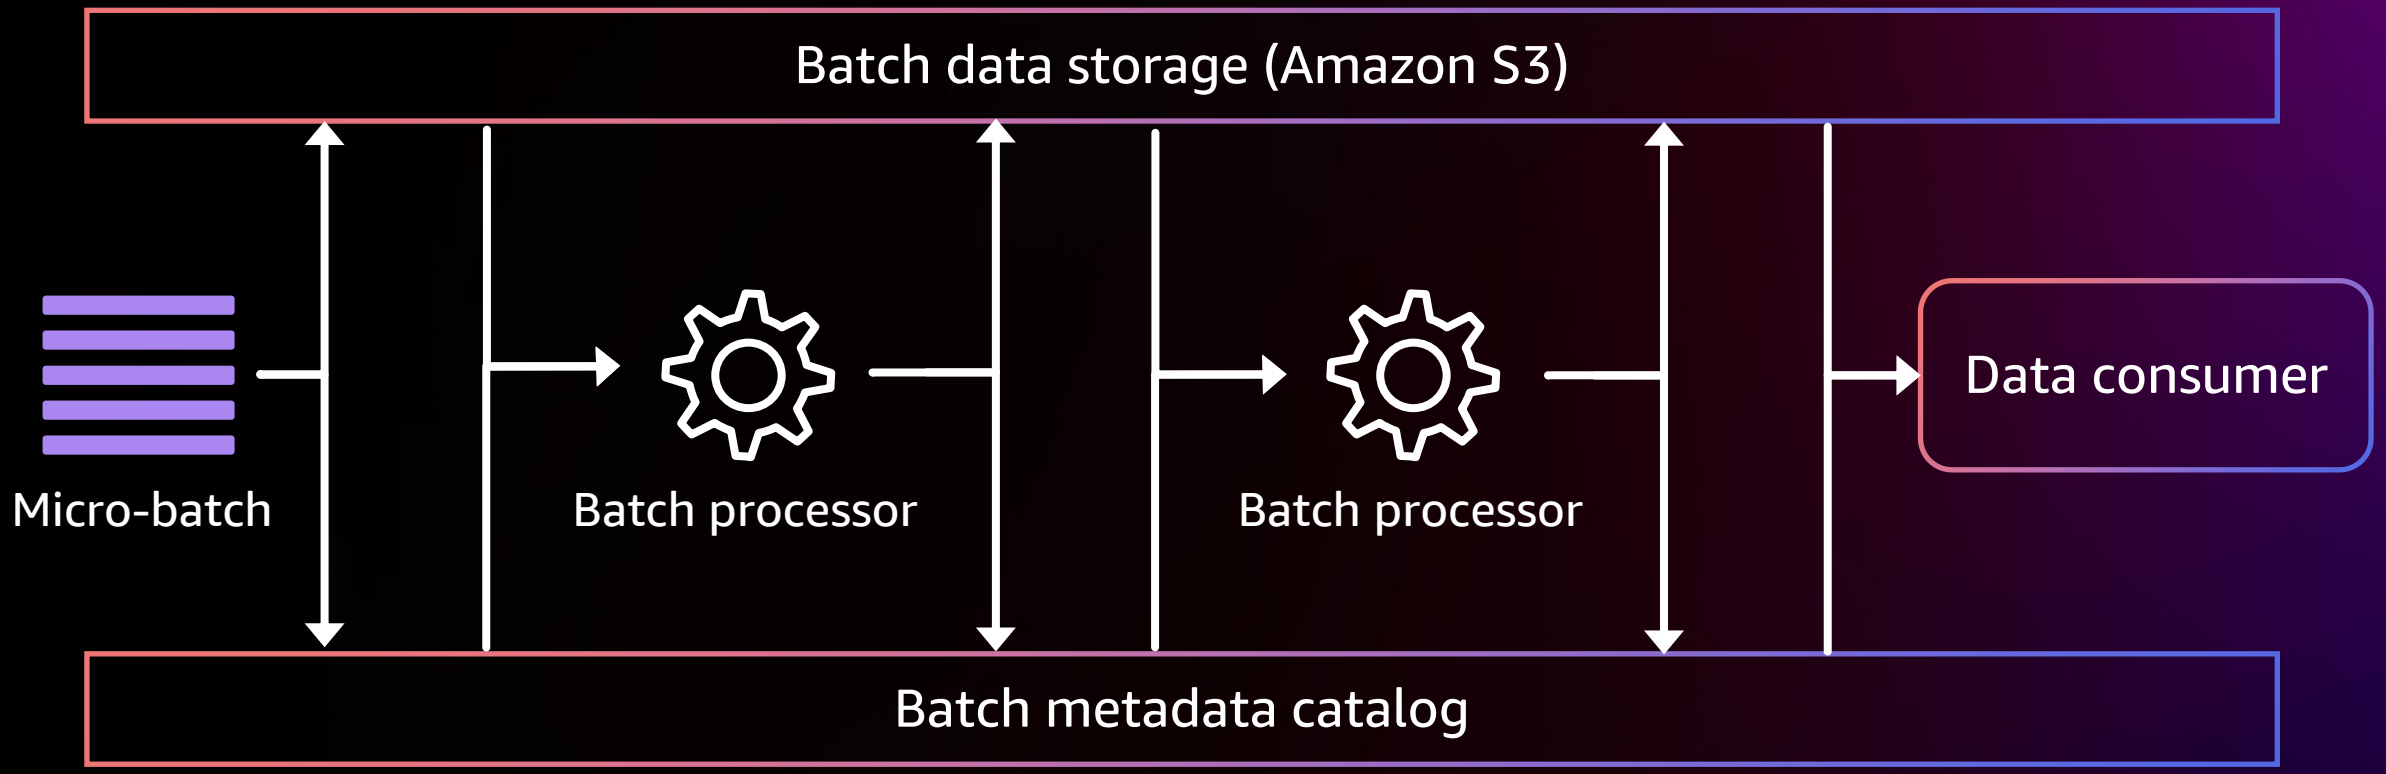

Orchestration V3 – Data rivers

	Data lake	Data river	Data stream
Data processing cost	★★★★★	★★★★	★★
Operational overhead	★★★★	★★★	★★
Developer velocity	★★★★	★★★★	★★
Data age	★	★★★★	★★★★★

Legend: more ★ indicates better results

Orchestration V3 – Data rivers

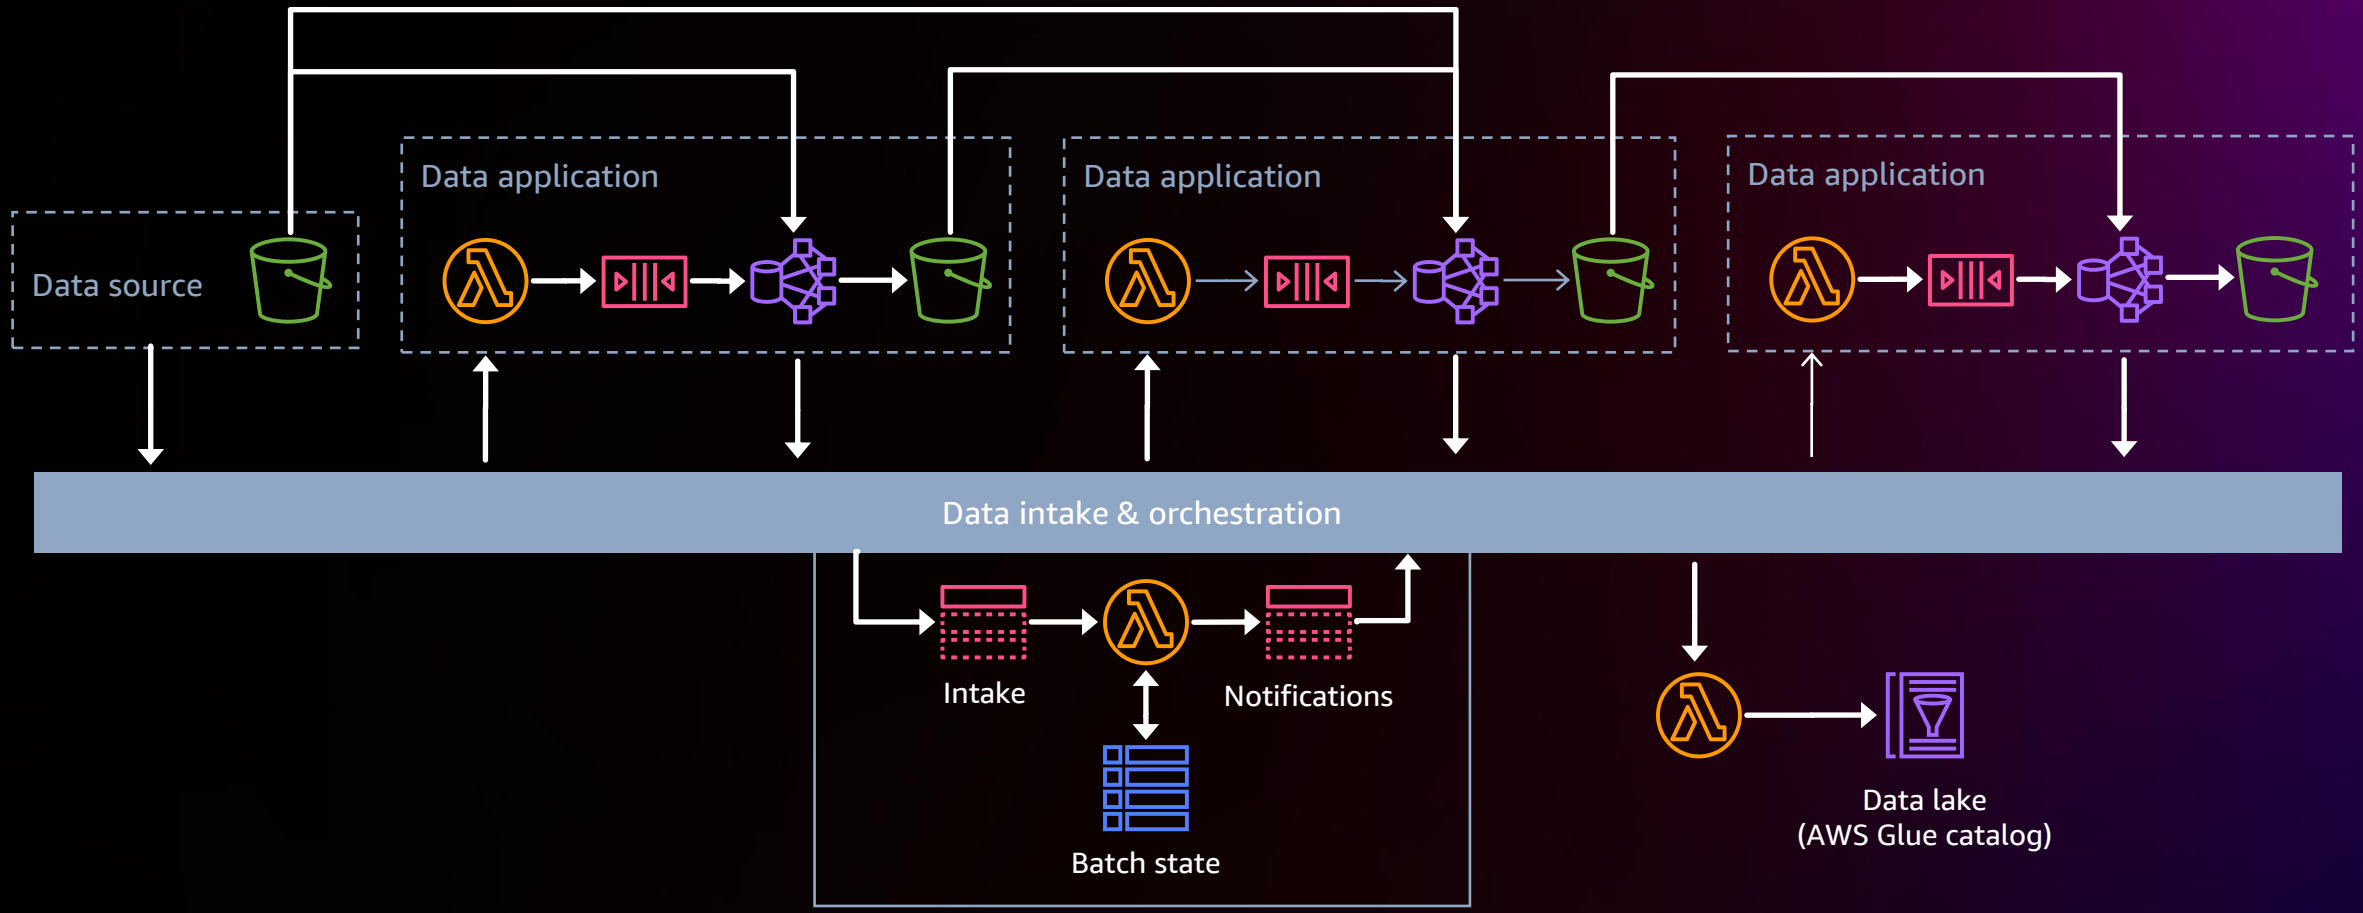

Data rivers – Outcome

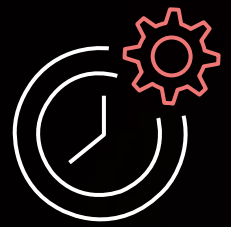

Data availability

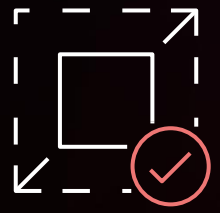

On-demand scaling

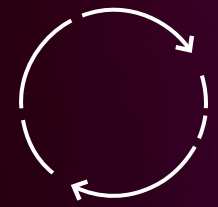

Continuous

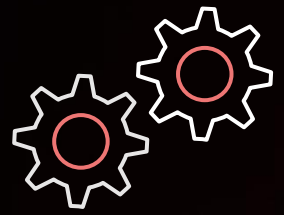

Operational savings

Data processing cost savings

Orchestration – Lessons learned

- Separate storage and compute
- Avoid monolithic clusters
- Choose the right technology
- Process data in small batches
- Data rivers offer a compromise between lake and stream

Data applications

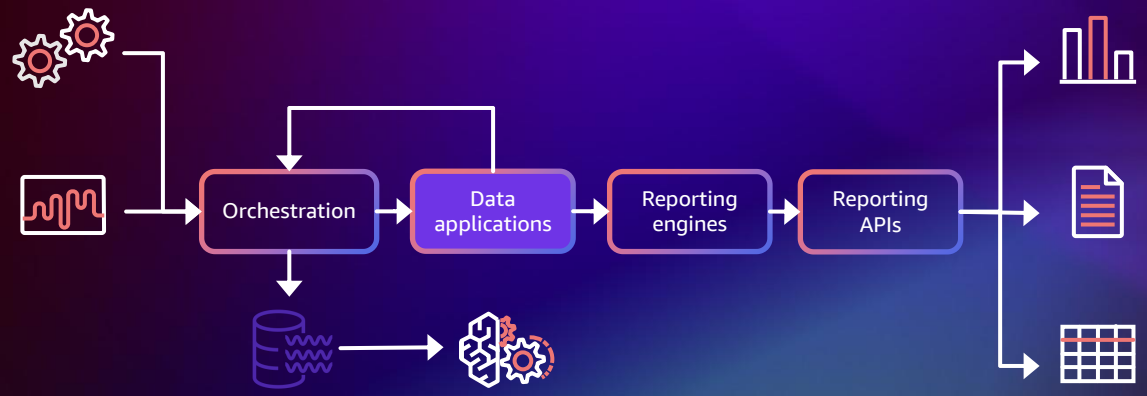

Data application

Partitioning

Aggregation

Batches

Restatements

Data applications

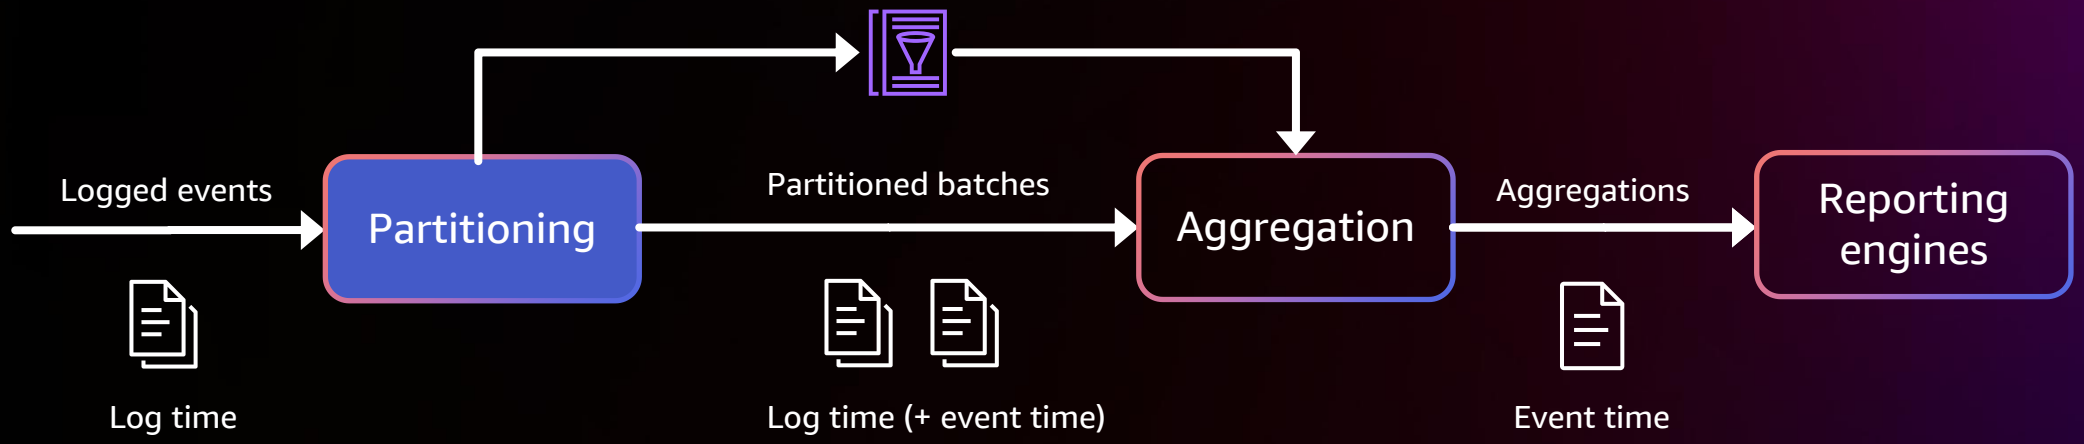

Data applications

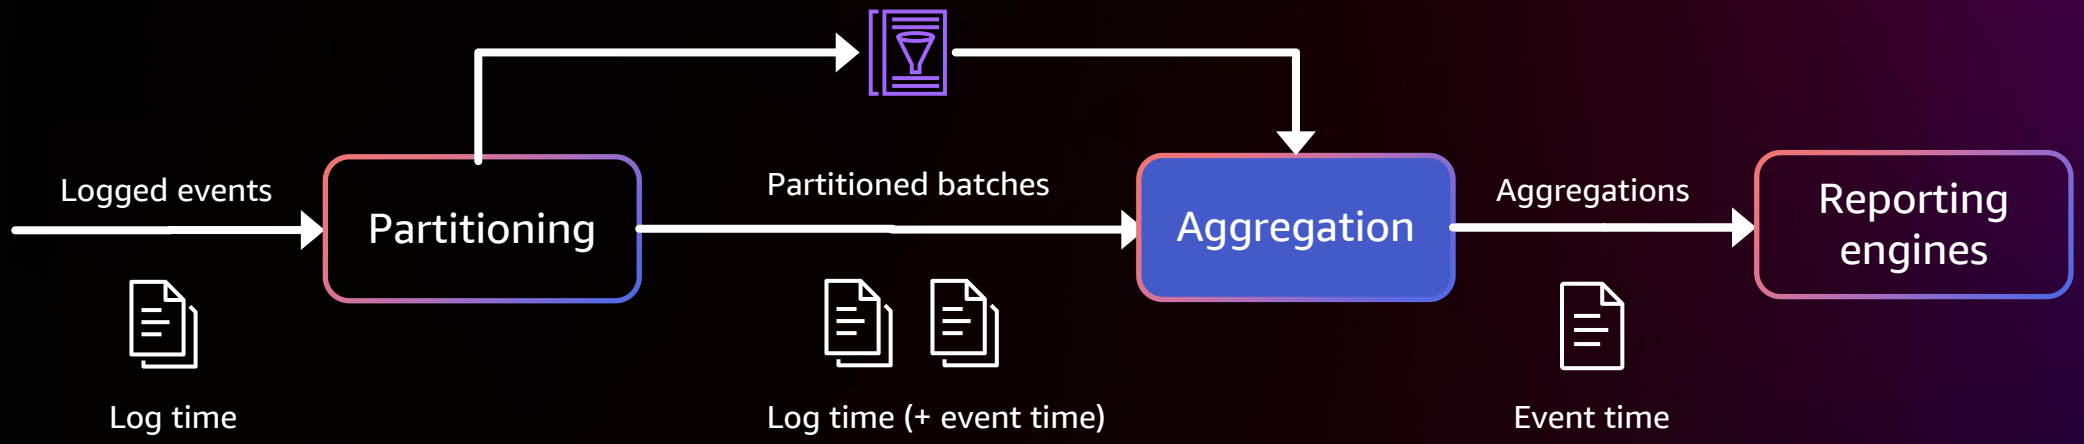

Reporting engines

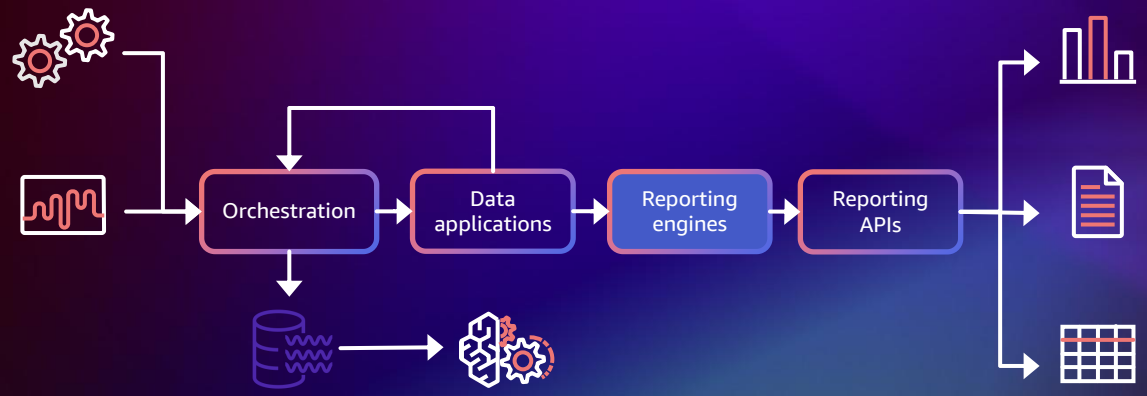

What is a reporting?

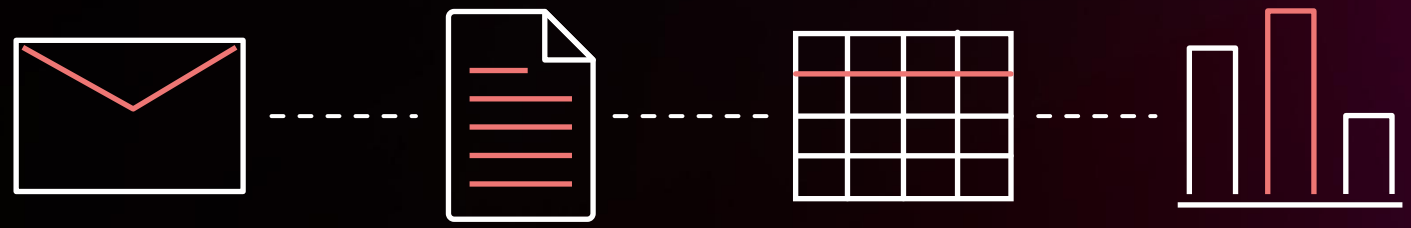

1.5M indexed per second

300M scanned per second

28M aggregated per second

8k requests per second

<2s request latency

Reporting engines

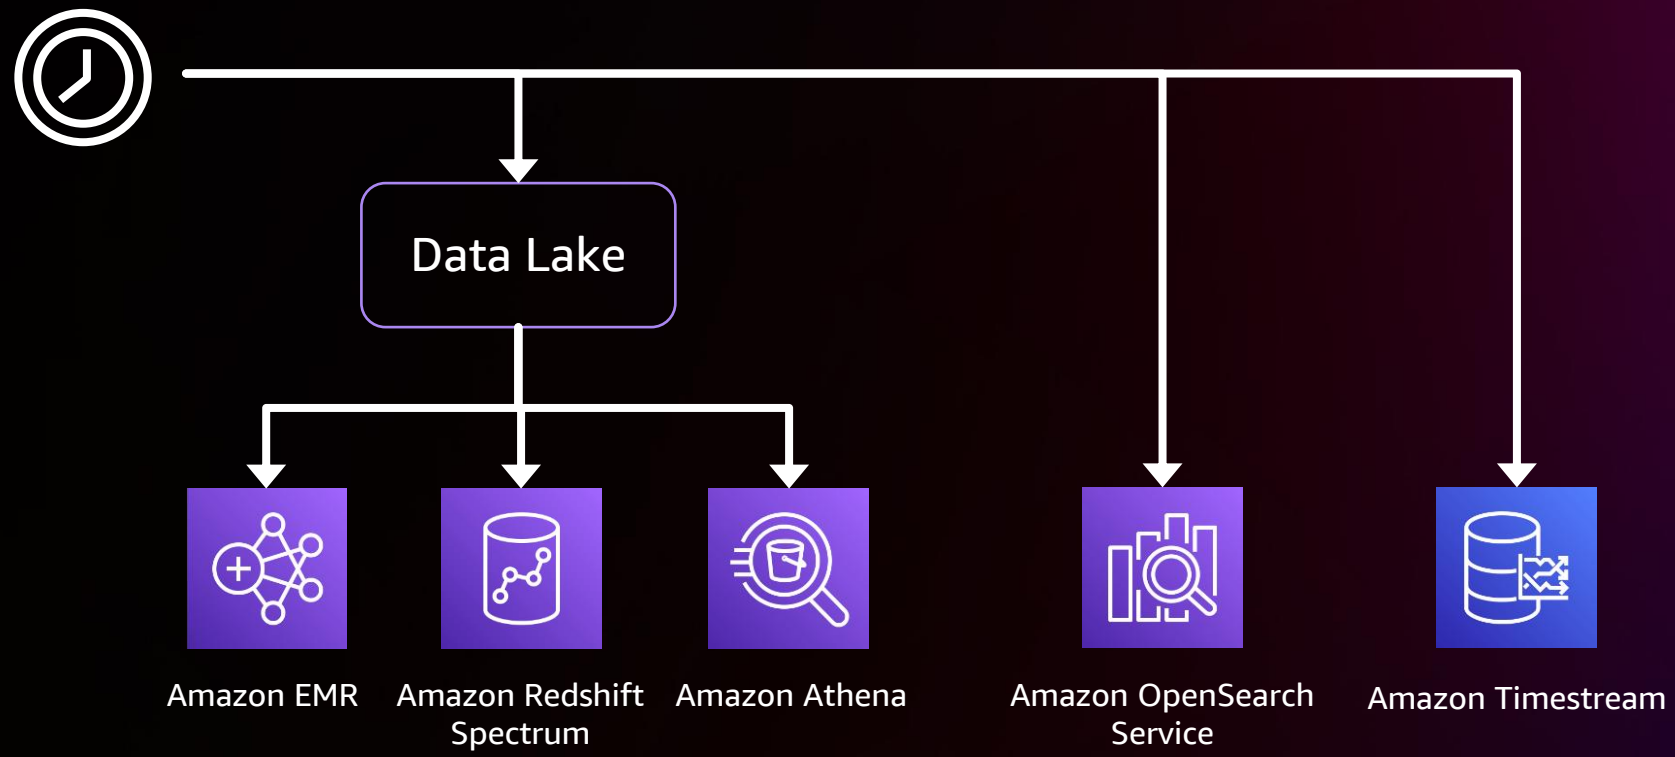

Reporting engines V1

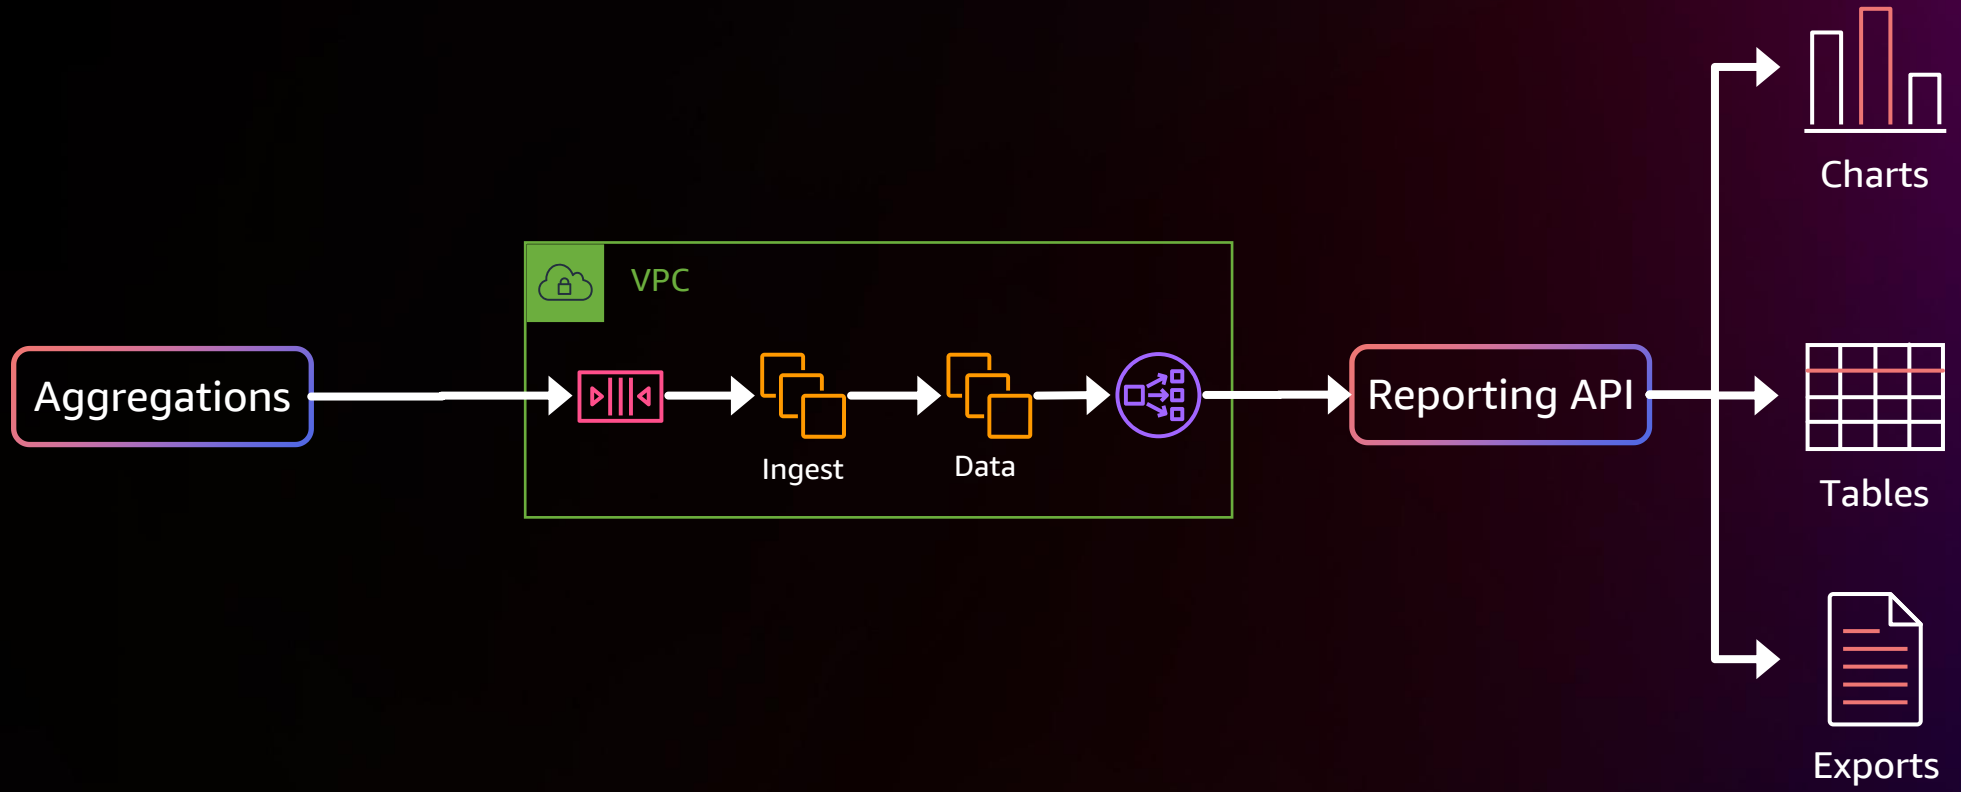

Reporting engines V1

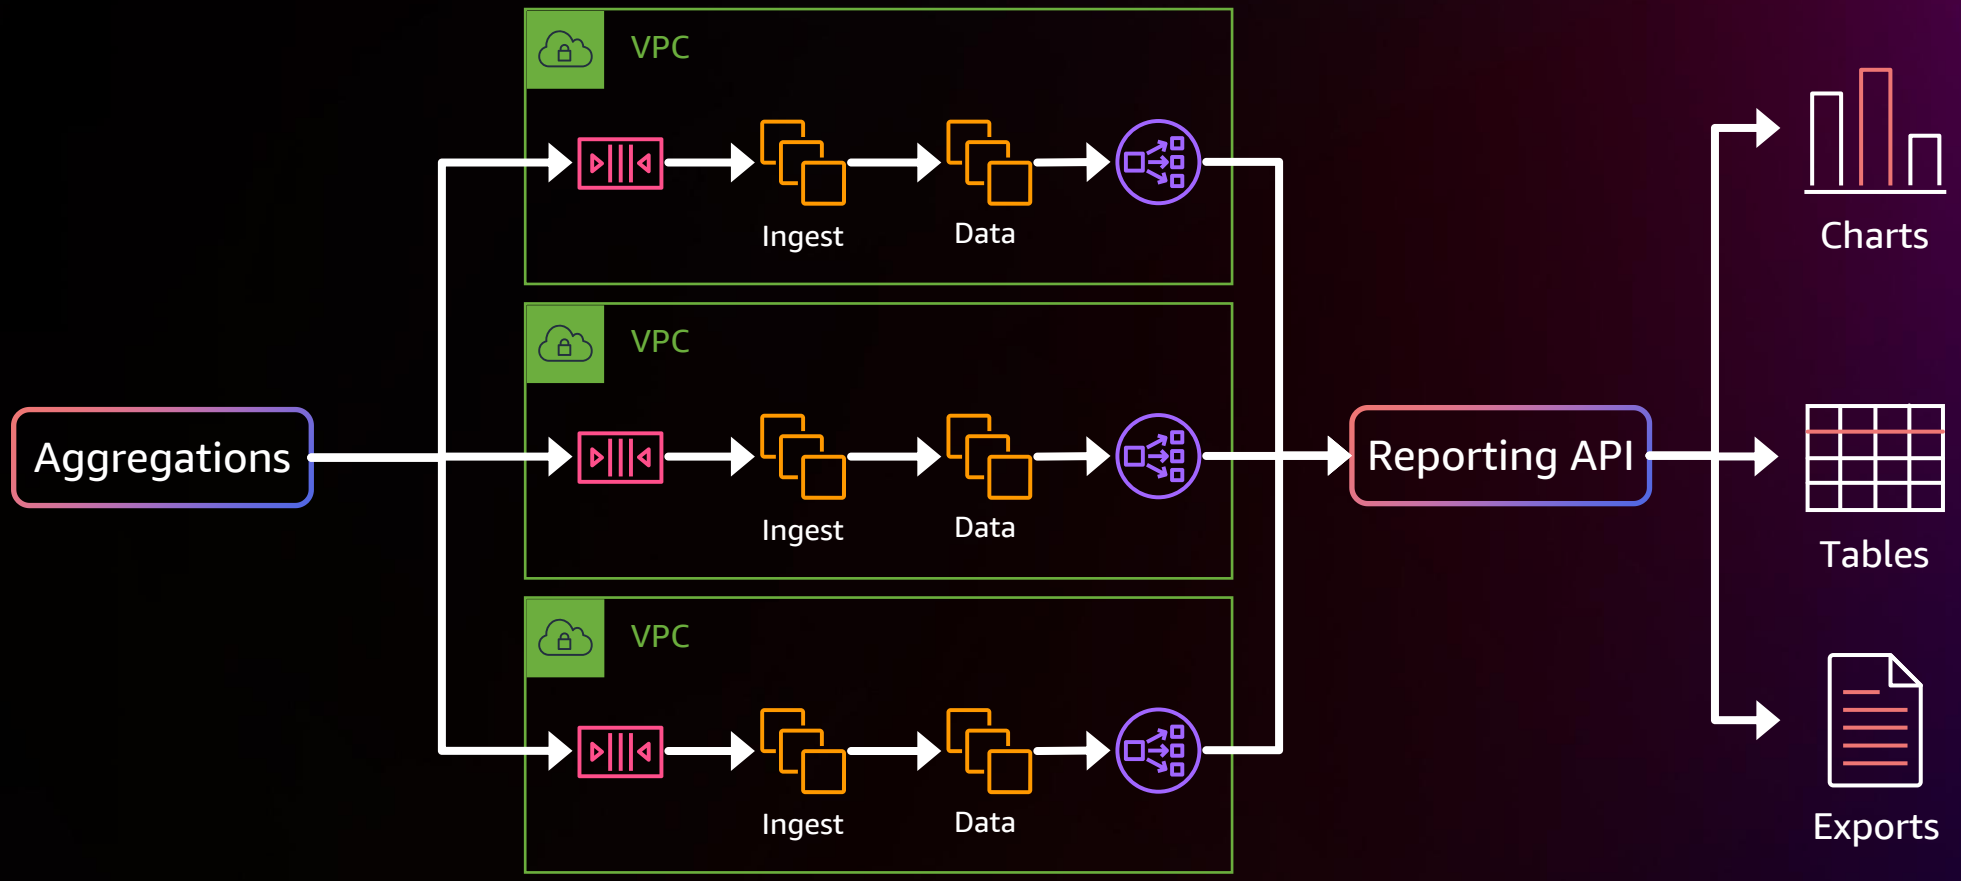

Reporting engines V2

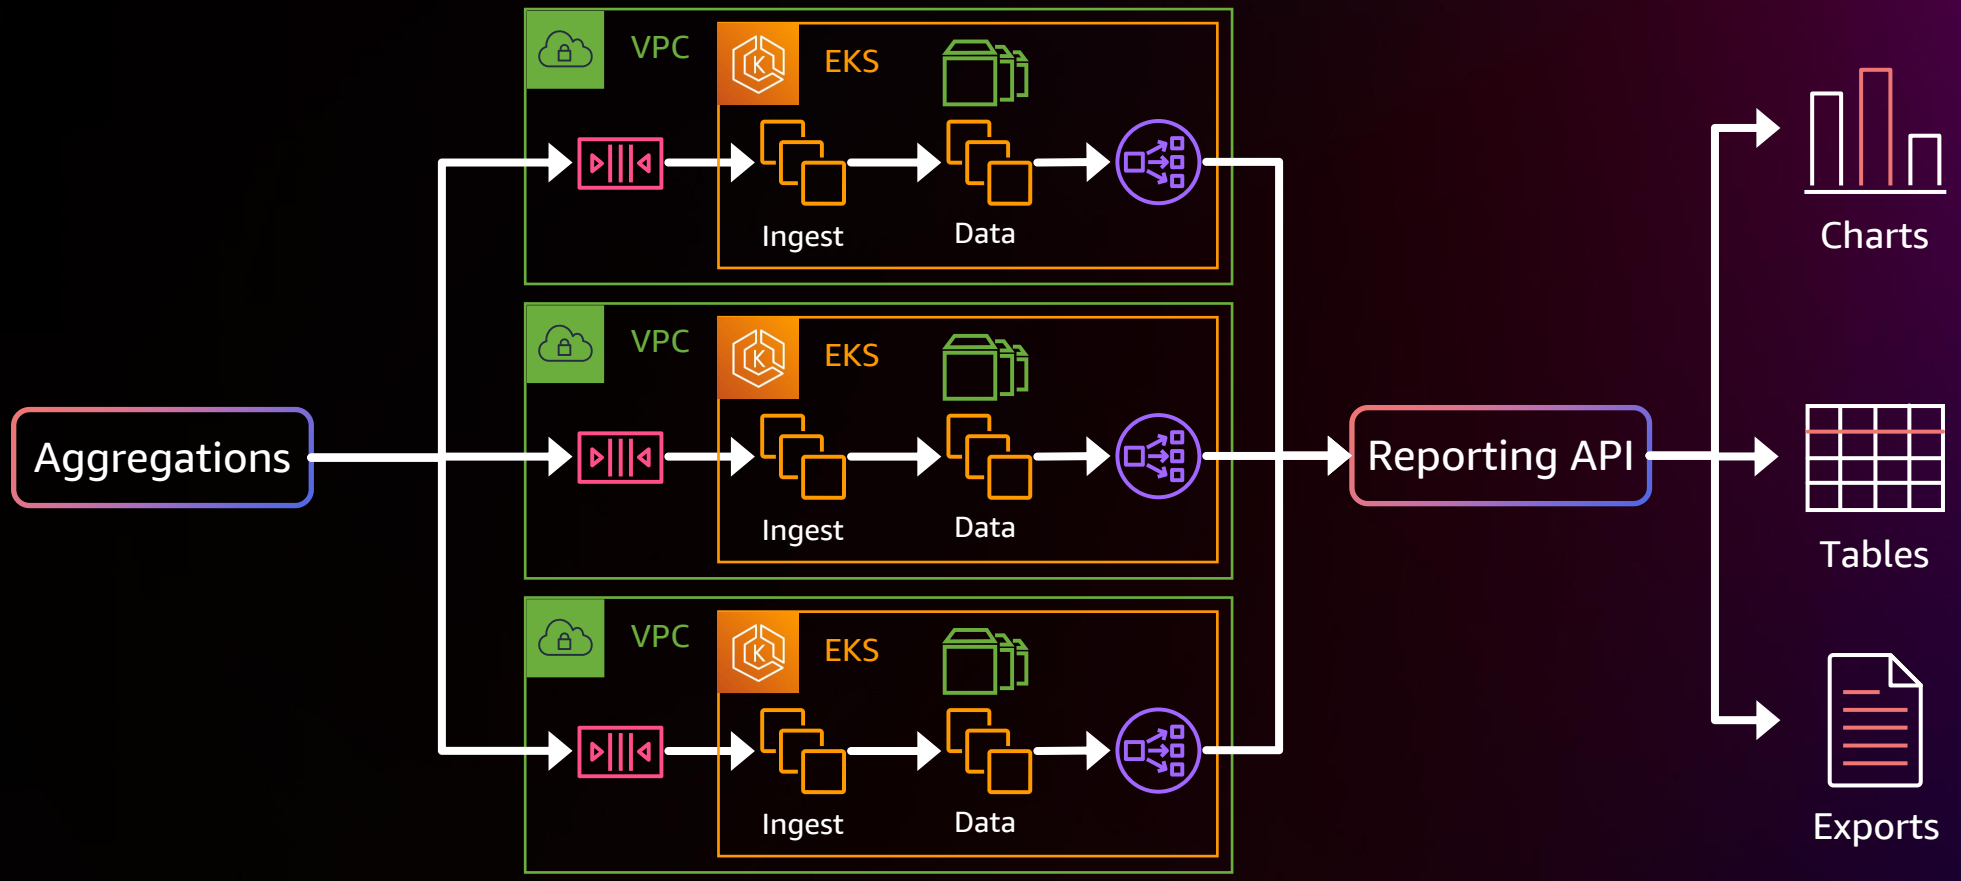

Reporting engines – Lessons learned

- Separate storage and compute
- Avoid monolithic clusters
- Choose the right technology
- Containers enable deploying multi-component services quickly

Reporting API

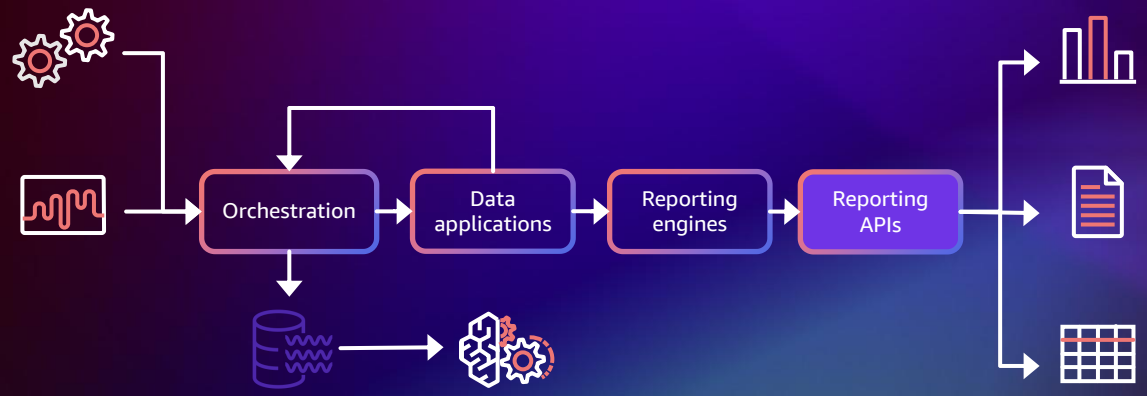

Reporting API – Exports

Reporting API – Exports

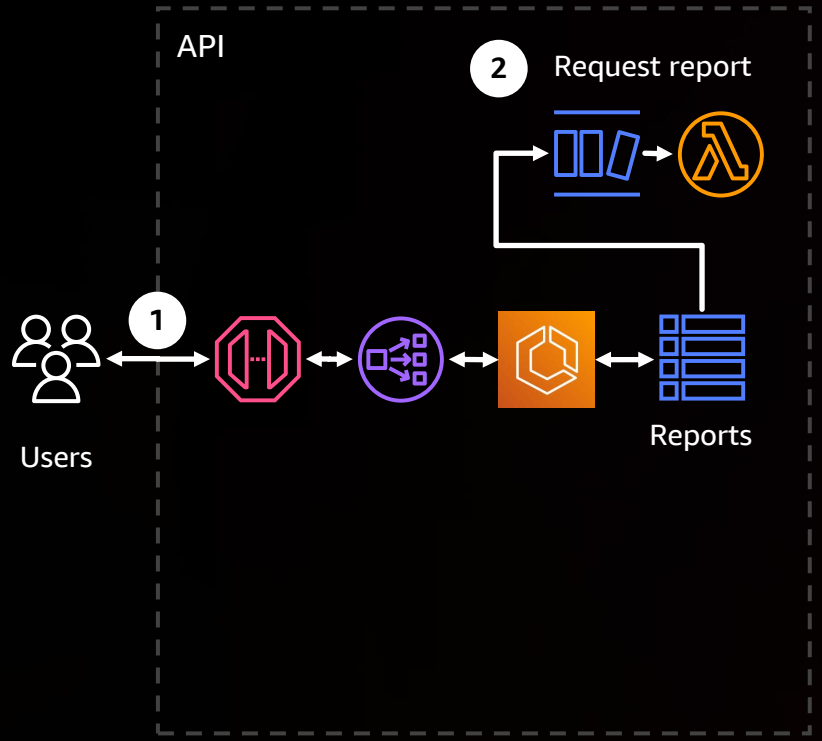

Reporting API – Exports

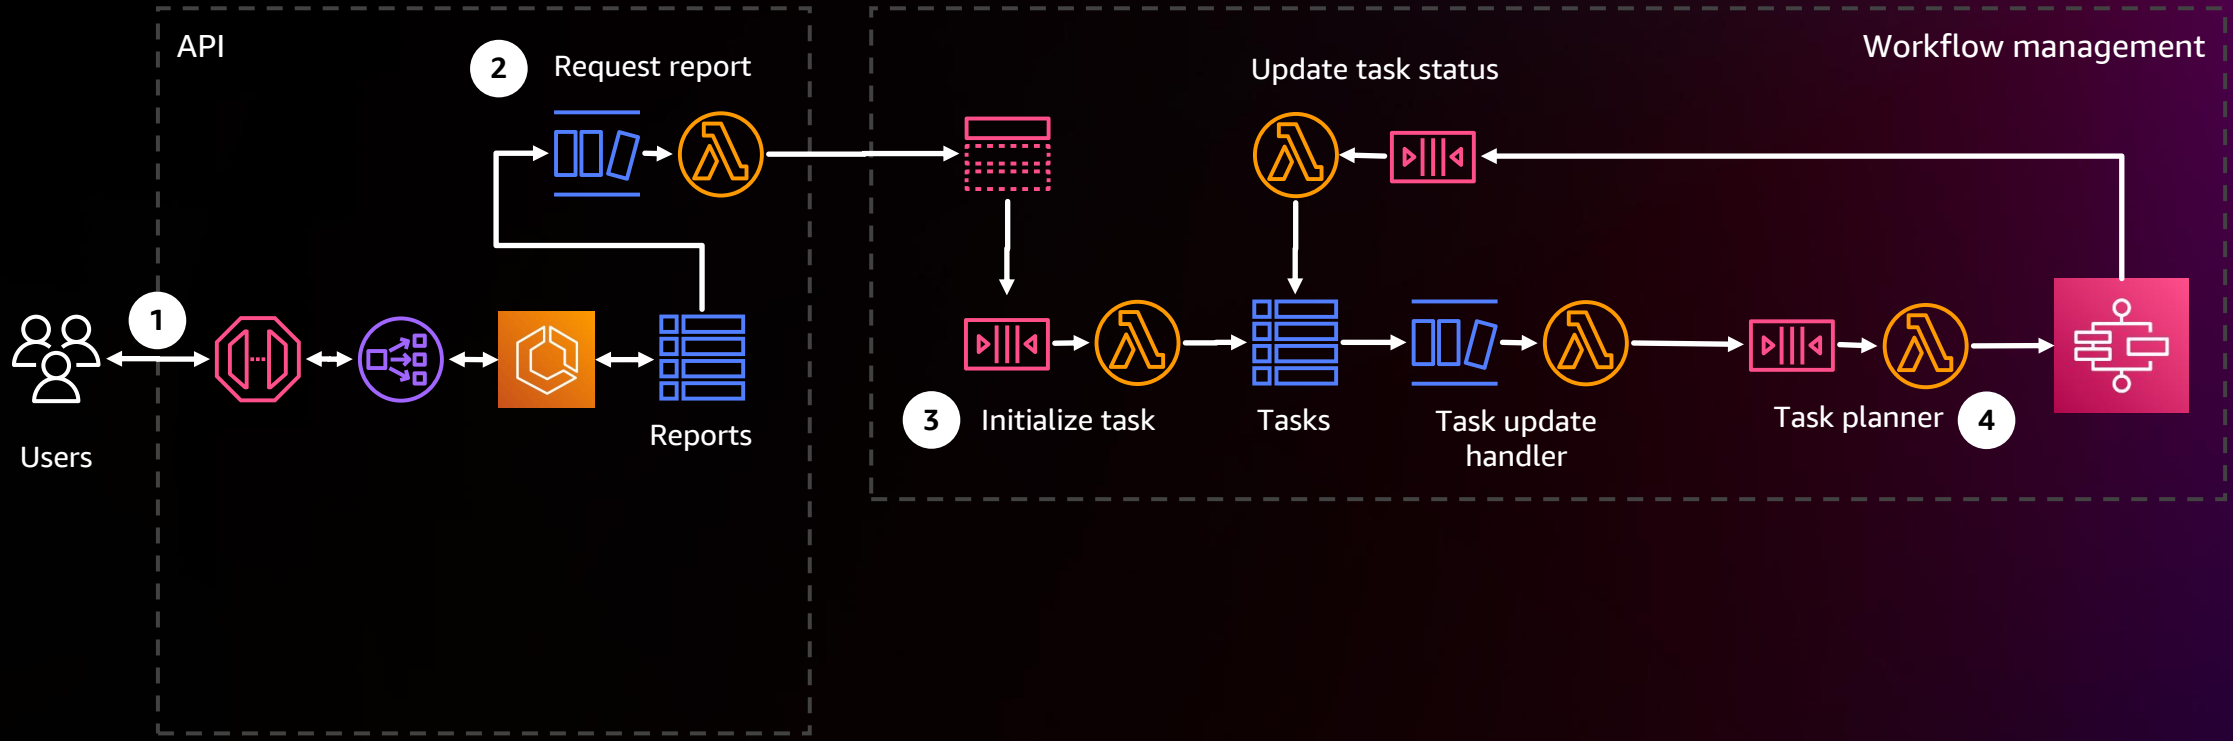

Reporting API – Exports

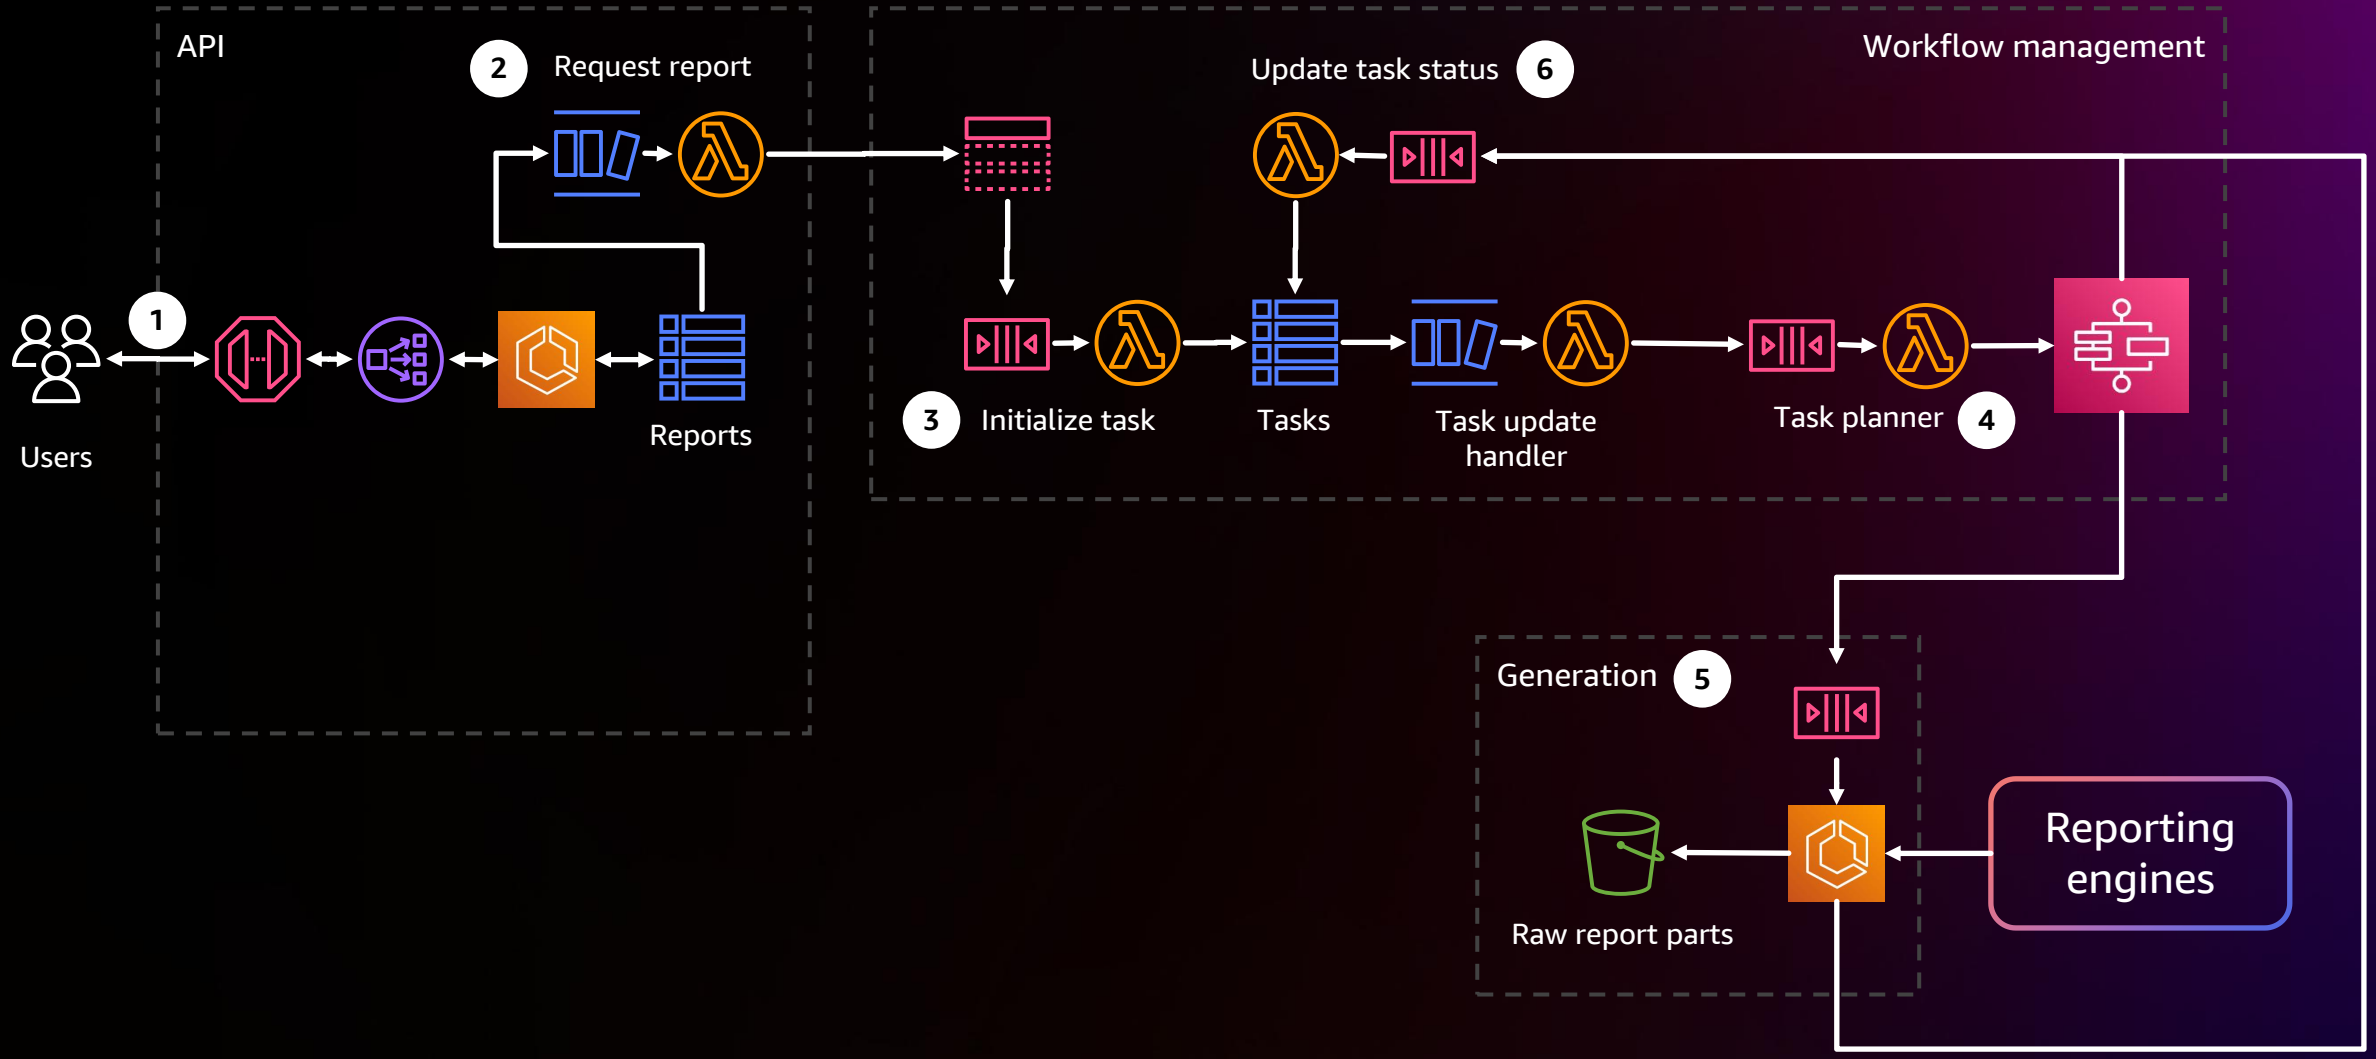

Reporting API – Exports

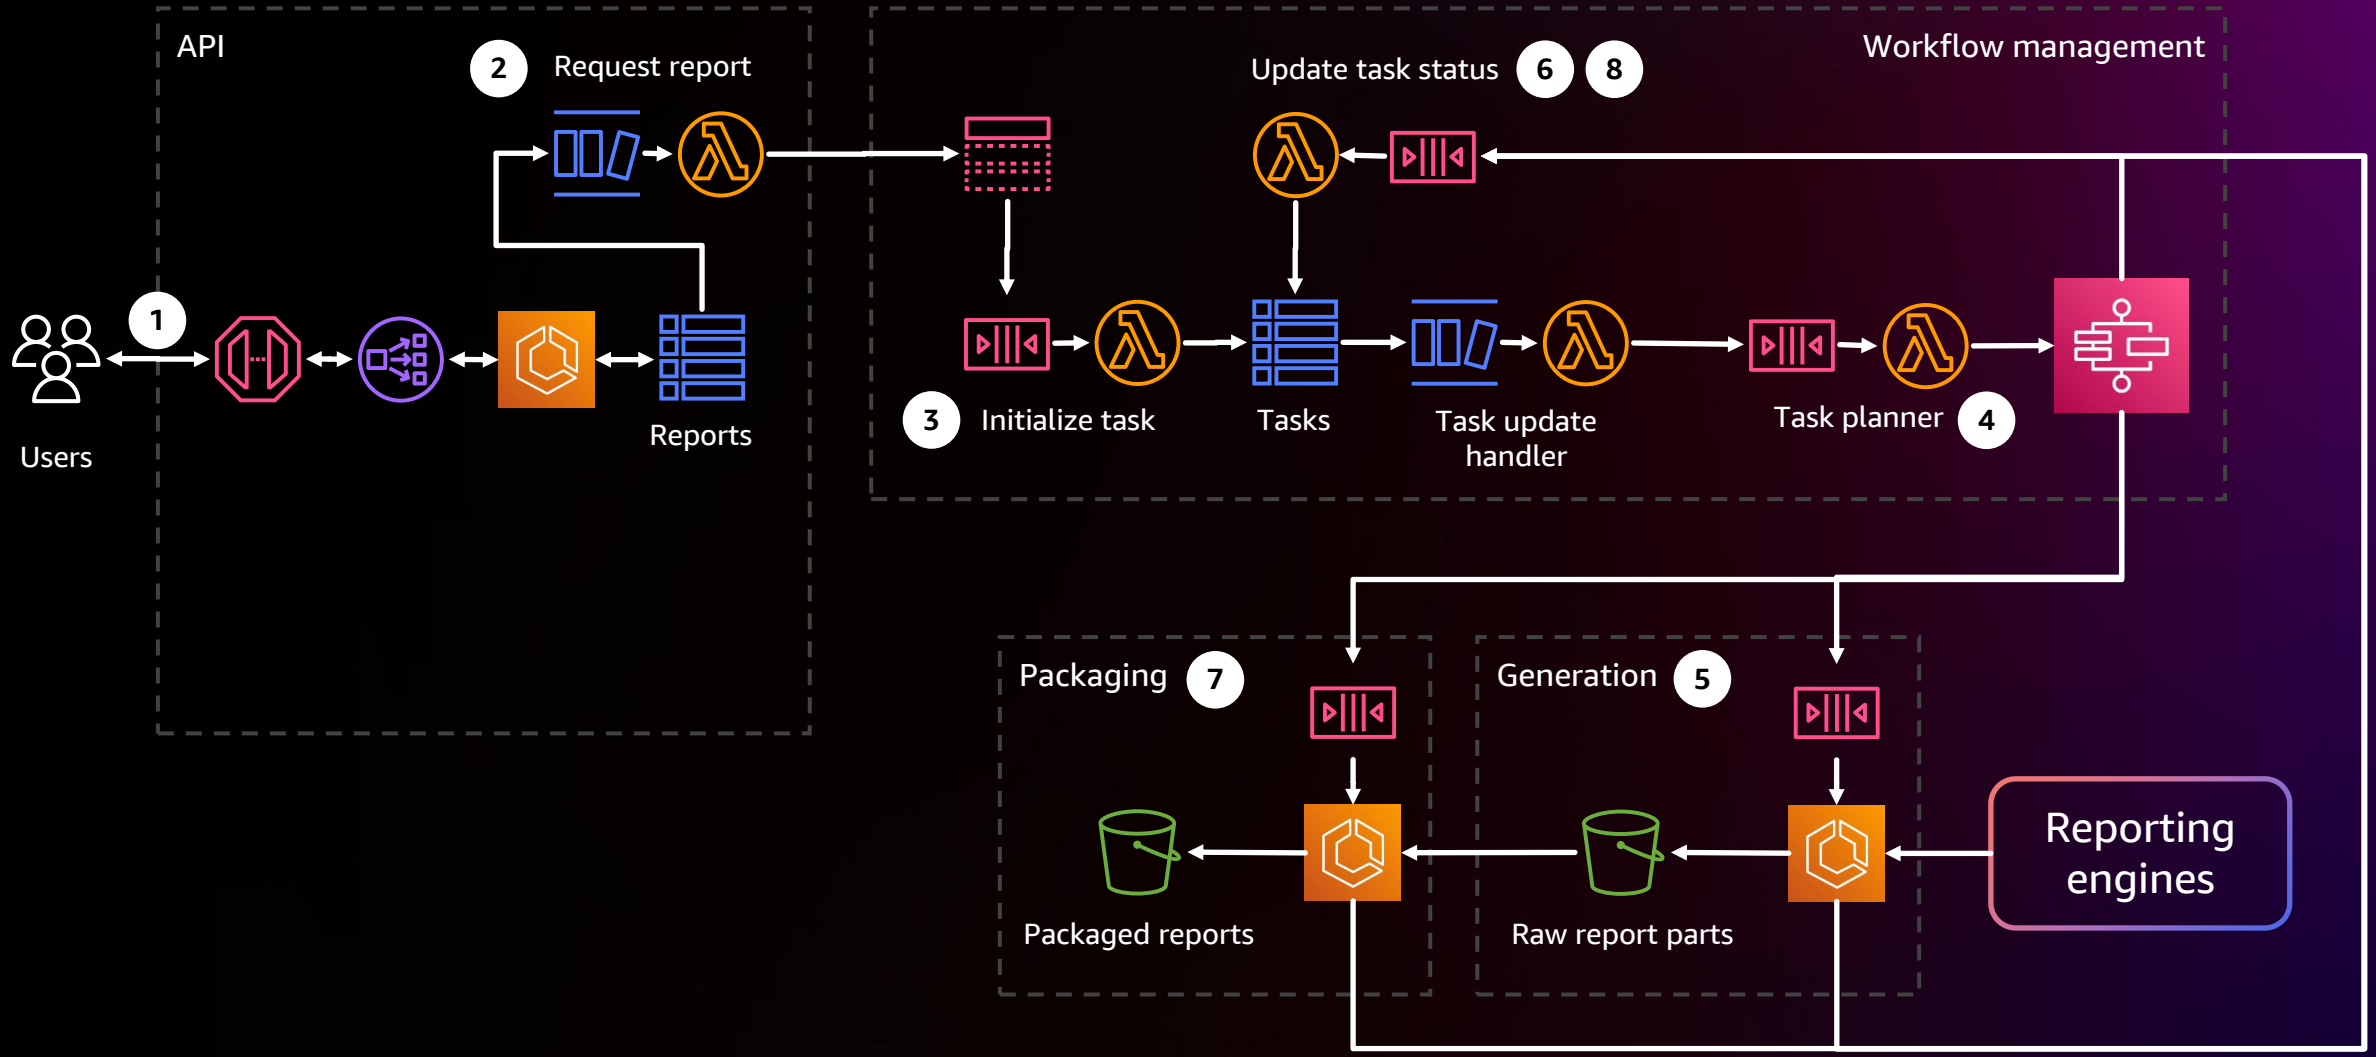

Reporting API – Exports

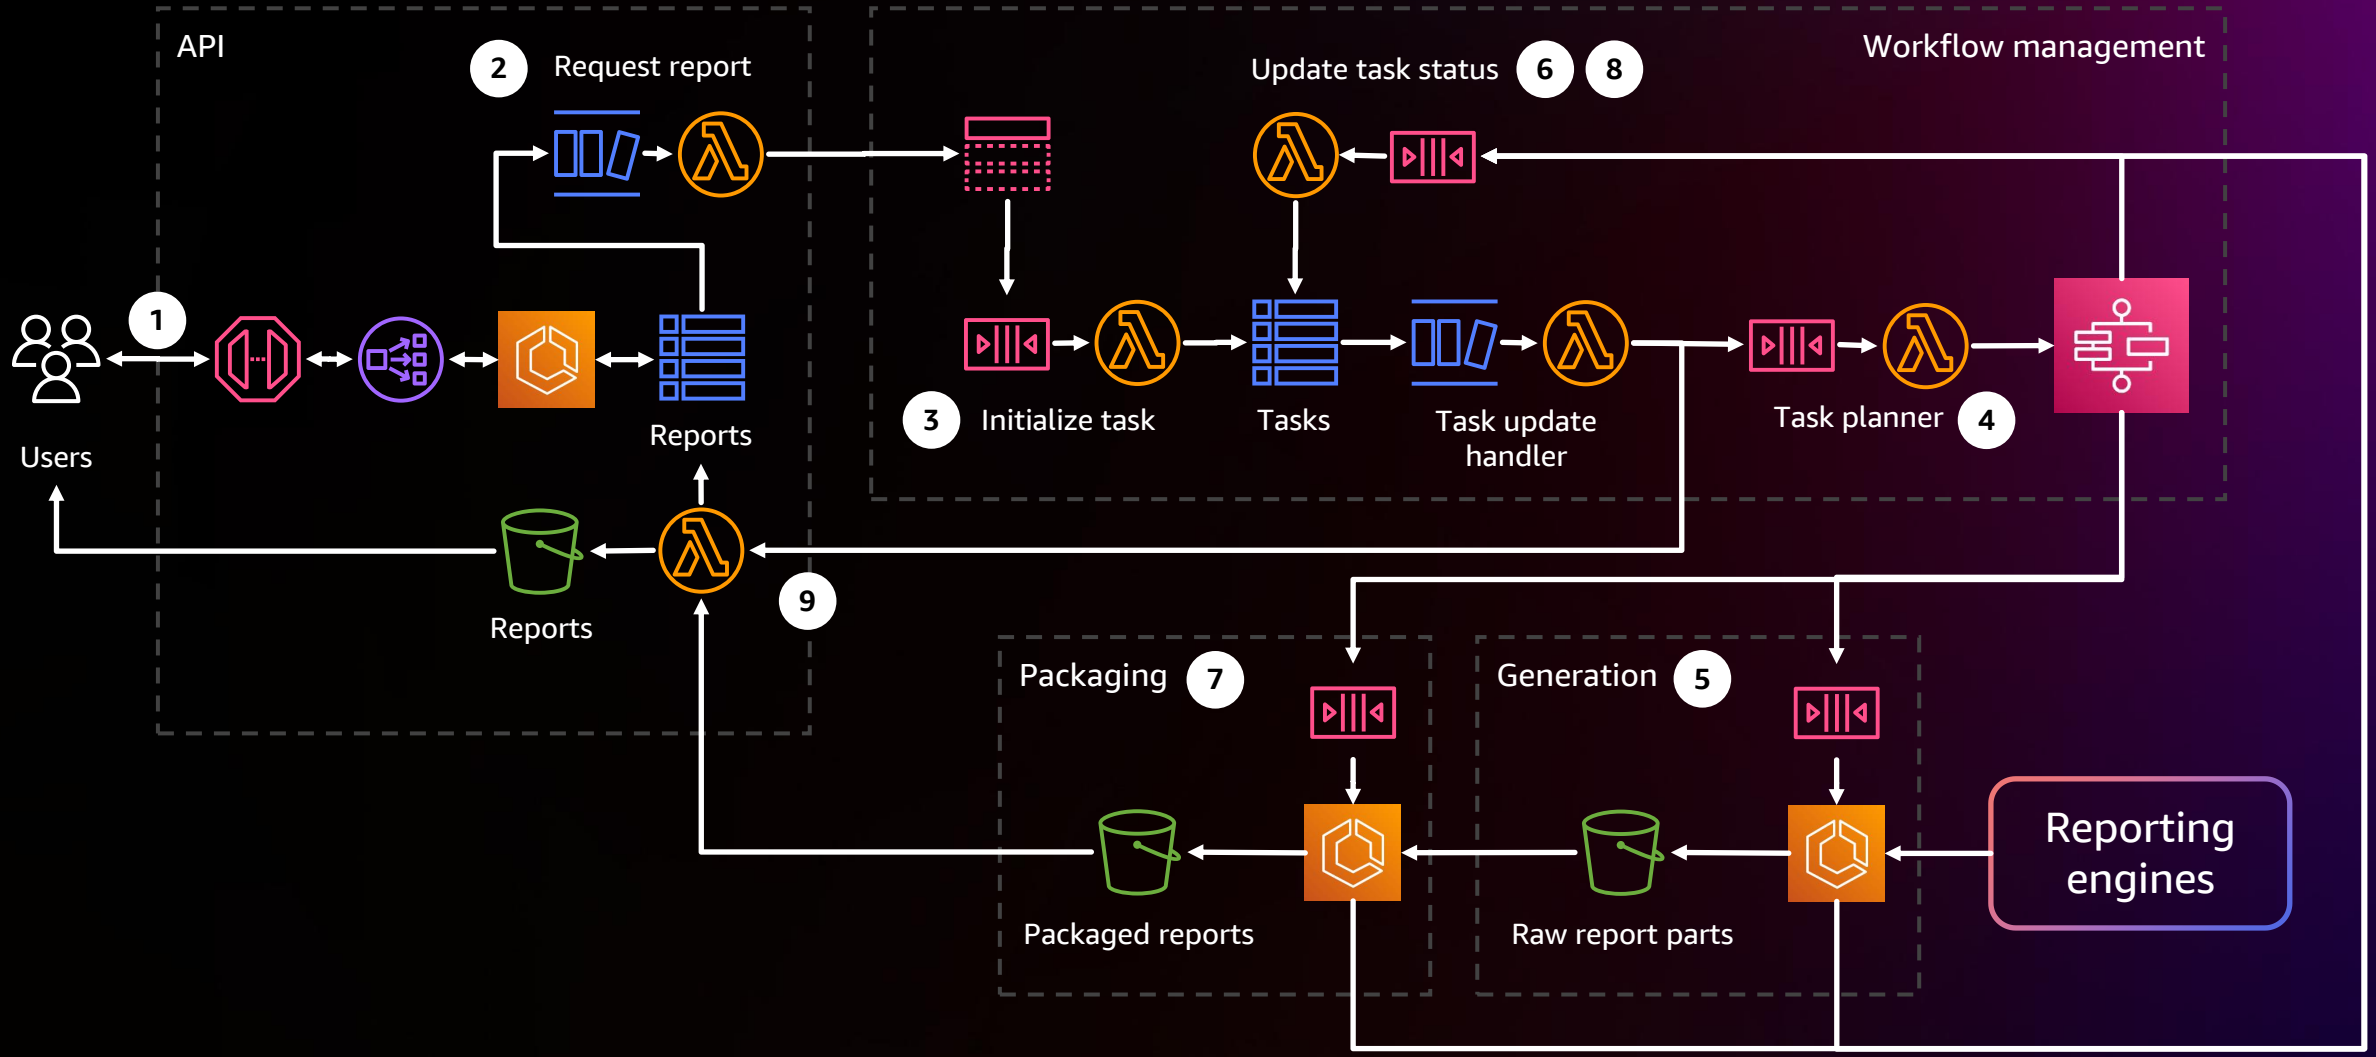

Reporting API – Lessons learned

Choose the right technology

Process data in small batches

Recap

Orchestration with data rivers

Partitioning and aggregating in a river

Scaling high-TPS, low-latency reporting with Amazon EKS and OpenSearch

Using step functions to build reporting exports incrementally

Thank you!

Amazon Ads booth: 2040 (behind AWS village)

Please complete the session survey in the **mobile app**