
AWS
re:Invent
NOV. 28 – DEC. 2, 2022 | LAS VEGAS, NV

ANT306-R

Beyond monitoring: Observability with operational analytics data

Muhammad Ali (he/him)

Sr. Analytics Specialist Solutions Architect
AWS

Rafael Gumiero (he/him)

Sr. Analytics Specialist SA
AWS

Improving customer experience

Reproducing problems

Performance

Creating innovative features

Debugging failures

Scale challenges

Service outages

Networking connectivity

Hardware failure

Availability

Slow response analysis

Resolving downtime

Cost optimization

Long mean time to detect failure

Component interdependency

Long mean time to repair failure

Rolling out to new geography

New releases

Improve developer experience

Building automations

Monitoring recovery

Resolving insufficient capacity

Improving customer experience

Reproducing problems

Performance

Creating innovative features

Debugging failures

Scale challenges

Service outages

Networking connectivity

Hardware failure

Availability

Resolving downtime

Slow response analysis

Cost optimization

Long mean time to detect failure

Component interdependency

Long mean time to repair failure

Rolling out to new geography

New releases

Improve developer experience

Building automations

Monitoring recovery

Resolving insufficient capacity

Agenda

Observability overview

Observability with Amazon OpenSearch Service

Observability ingestion tools with OpenSearch Service

Lab: Microservice Observability with Amazon OpenSearch Service Workshop

Lab setup

Observability overview

What is observability?

Observability is the ability of a system to be observed by gathering telemetry data to EFFICIENTLY detect, investigate, and remediate issues faster

What are the observability pillars?

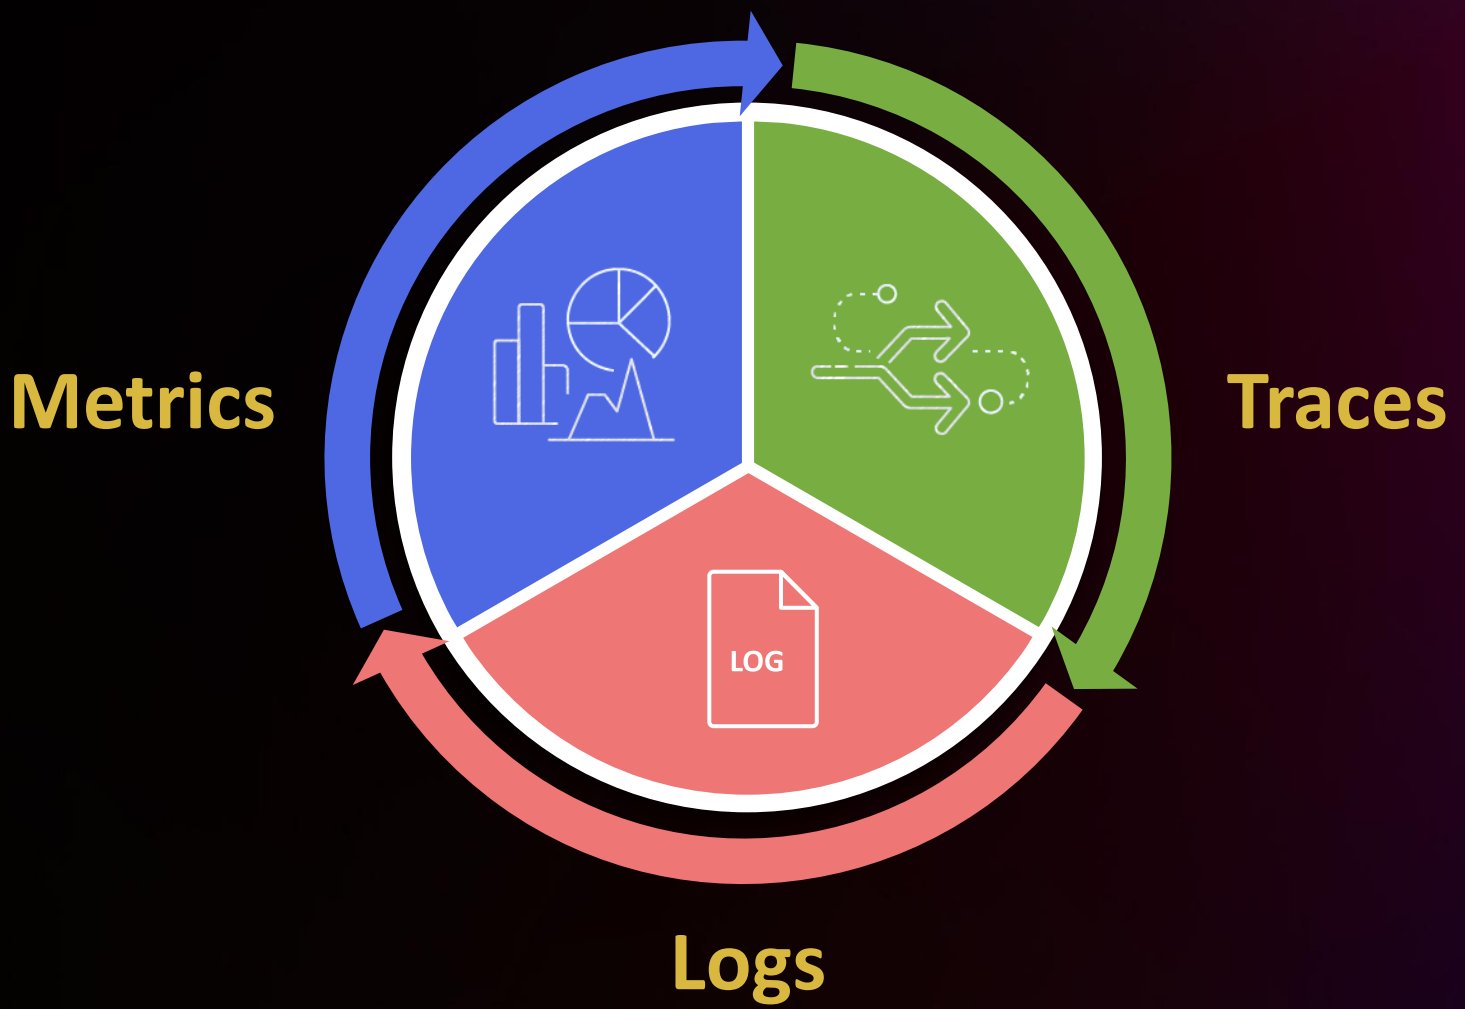

What are the main observability use cases?

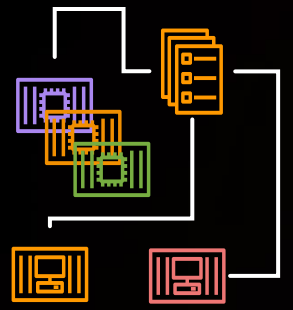

Microservices and containers

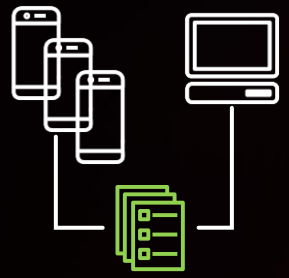

Digital experience monitoring (DEM)

AIOps and DevOps

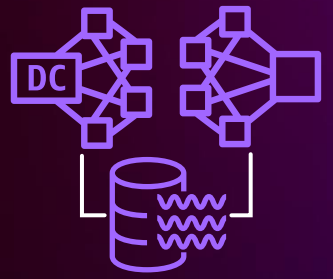

Data lakes

Observability with OpenSearch Service

Amazon OpenSearch Service

OpenSearch Service securely unlocks real-time search, monitoring, and analysis of operational data

Managed: Increase operational excellence by using a popular open-source solution

Secure: Audit and secure your data with a data center and network architecture and built-in certifications

Observability: Systematically detect potential threats and react to a system's state through an open-source solution for machine learning, alerting, and visualization

Cost-conscious: Optimize time and resources for strategic work

Observability with OpenSearch Service

APM – Application analytics to correlate logs, traces, and metrics

Distributed tracing with trace analytics

Log monitoring, event analytics

- Expanded log analytics capabilities
- PPL parse commands support

Out-of-the-box operational analytics
(coming soon – metric extraction, alerting, anomaly detection, Prometheus integration)

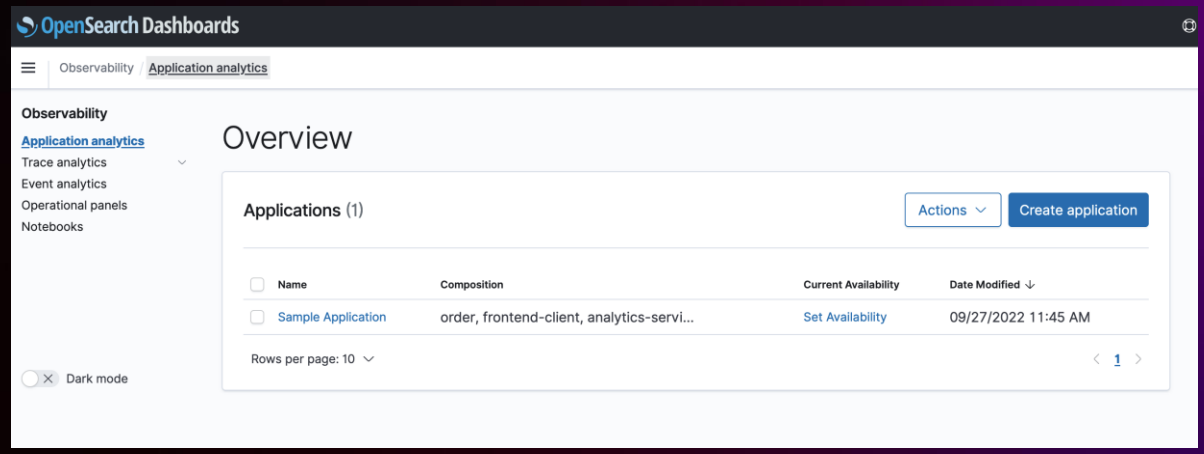

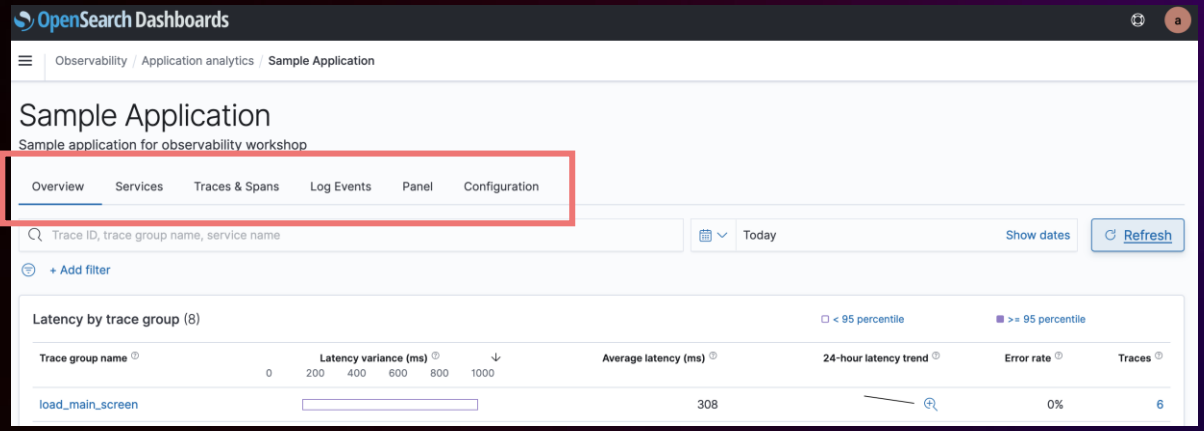

What is distributed tracing?

Distributed tracing refers to methods of observing requests as they propagate through distributed systems

A trace tells the story of a transaction or workflow as it propagates through a system

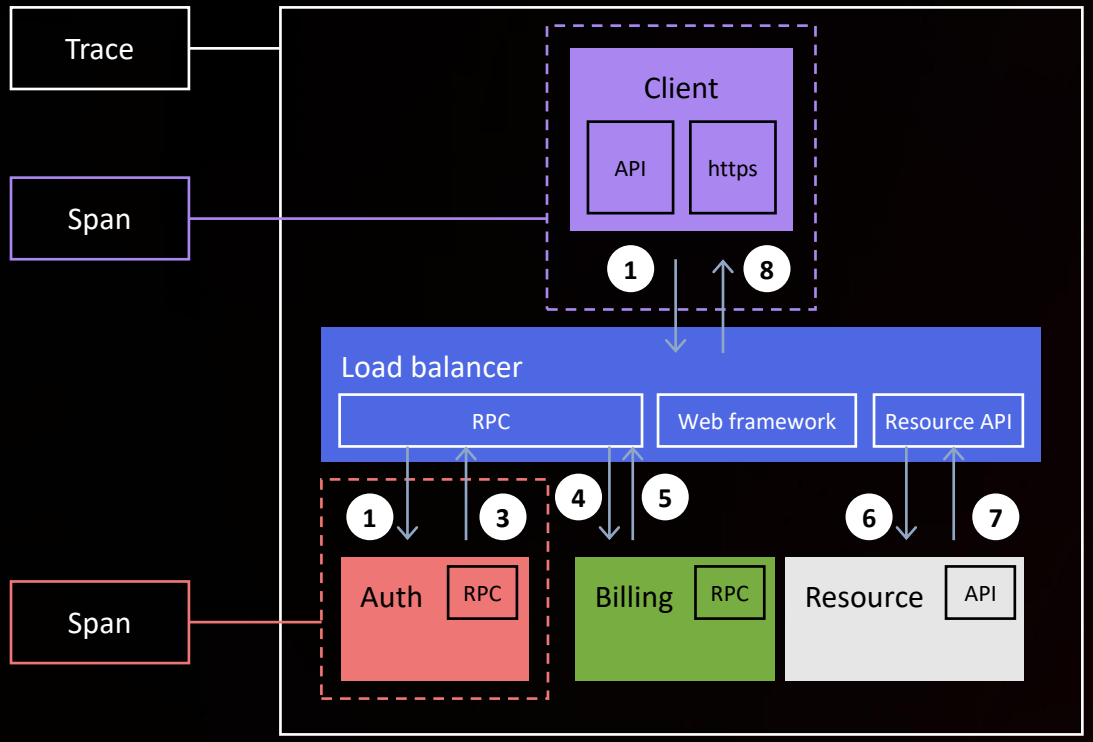

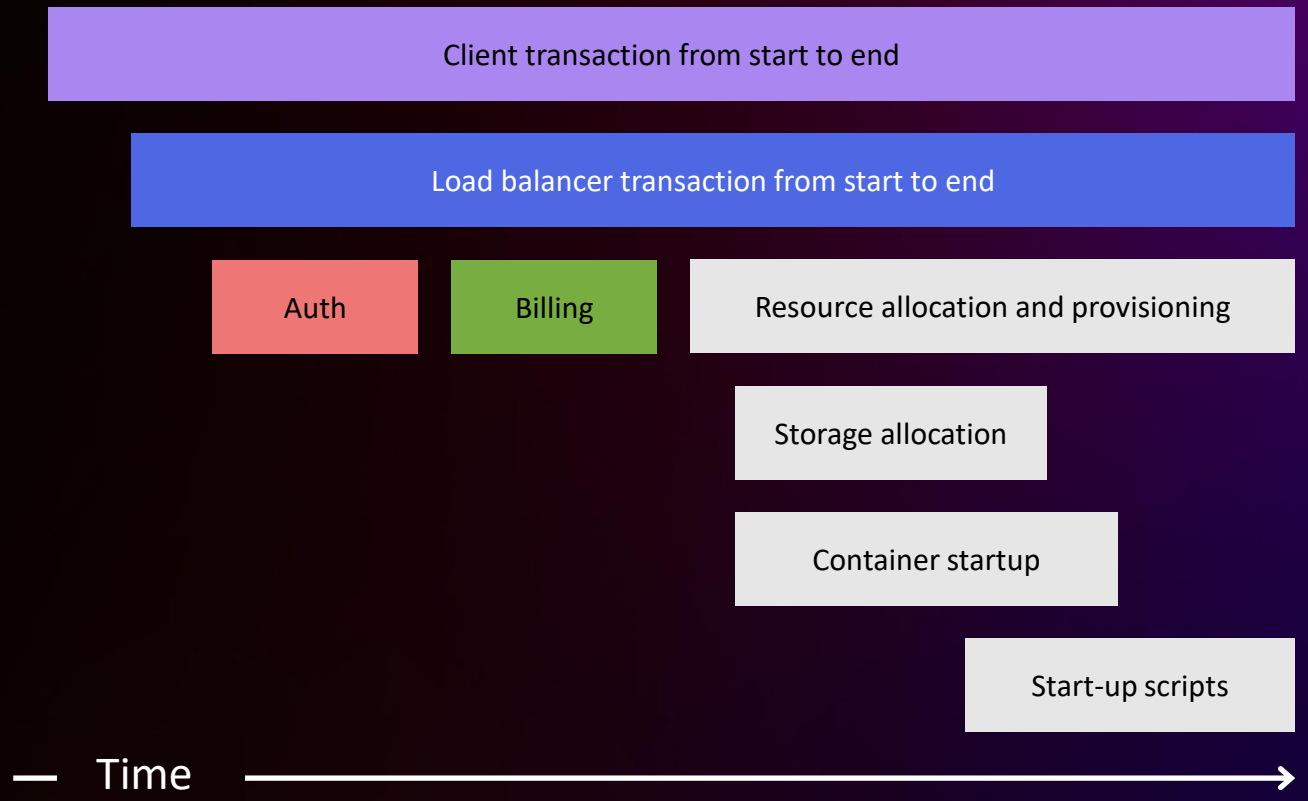

What is trace analytics?

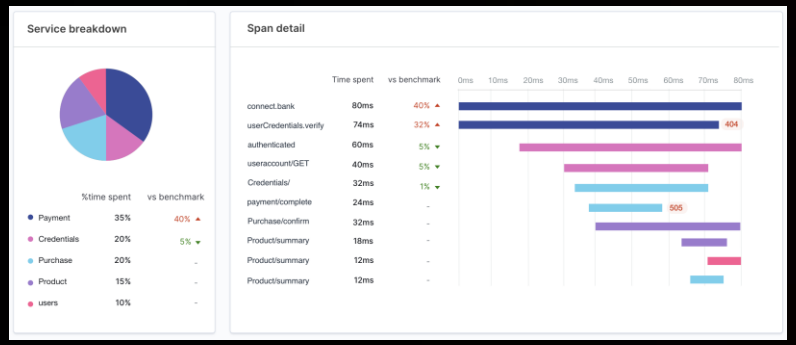

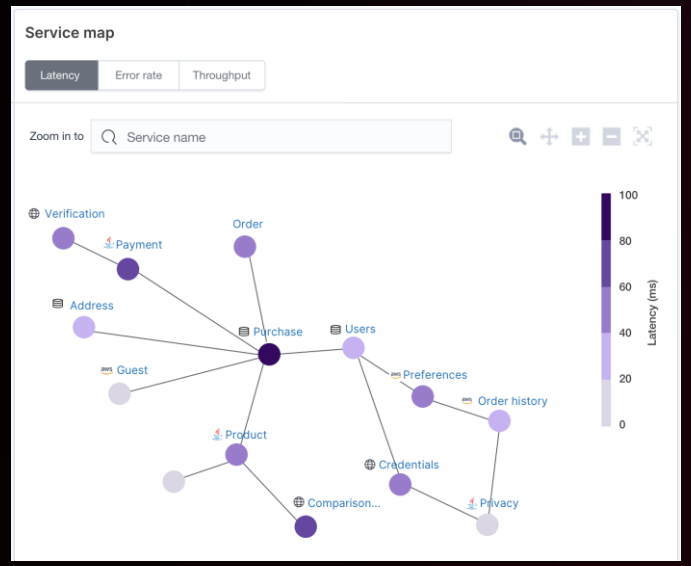

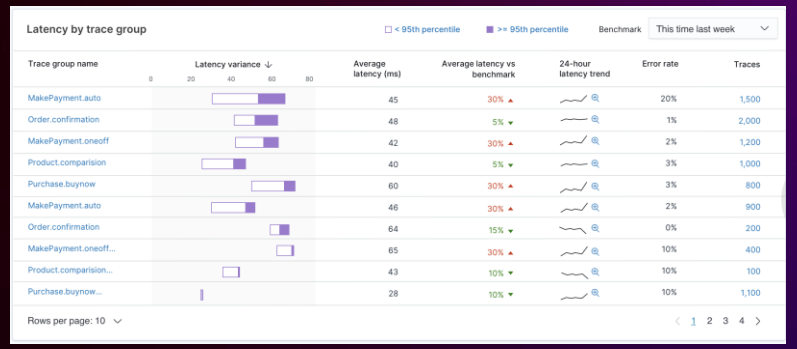

Trace-span details

- Single request performance
- Diagnose root cause

Service maps

- End-to-end view
- Isolate issues to services

Trace groups

- Monitor performance
- Identify issues early

Observability ingestion tools with OpenSearch Service

Open source instrumentation and transformation

- OpenTelemetry is a community-driven, open-source project designed for the creation and management of **telemetry data** such as traces, metrics, and logs
- Supports many popular open source wire formats including Jaeger, Zipkin, and Prometheus
- Currently supports traces (GA), metrics (GA), and logs (experimental)

- Fluent Bit is a community-driven, open source, super fast, lightweight, and highly scalable logging and metrics processor and forwarder
- Plugin-based ecosystem that collects, filters, transforms and augments data for logs, metrics, and traces
- Extensible and integrates with Prometheus, OpenSearch Service, Amazon CloudWatch, AWS X-Ray, Amazon S3, Amazon Kinesis Data Streams, and other AWS managed open-source offerings

- Data Prepper is a community-driven, open source data collector for processing observability data
- Provides features to filter, enrich, transform, normalize, and aggregate data for downstream analytics and visualization
- Currently supports processing of distributed trace data and log ingestion with plans to support metric data in the future
- Integrations with Jeager, Zipkin, OpenTelemetry, and Fluent Bit

Observability ingestion flow

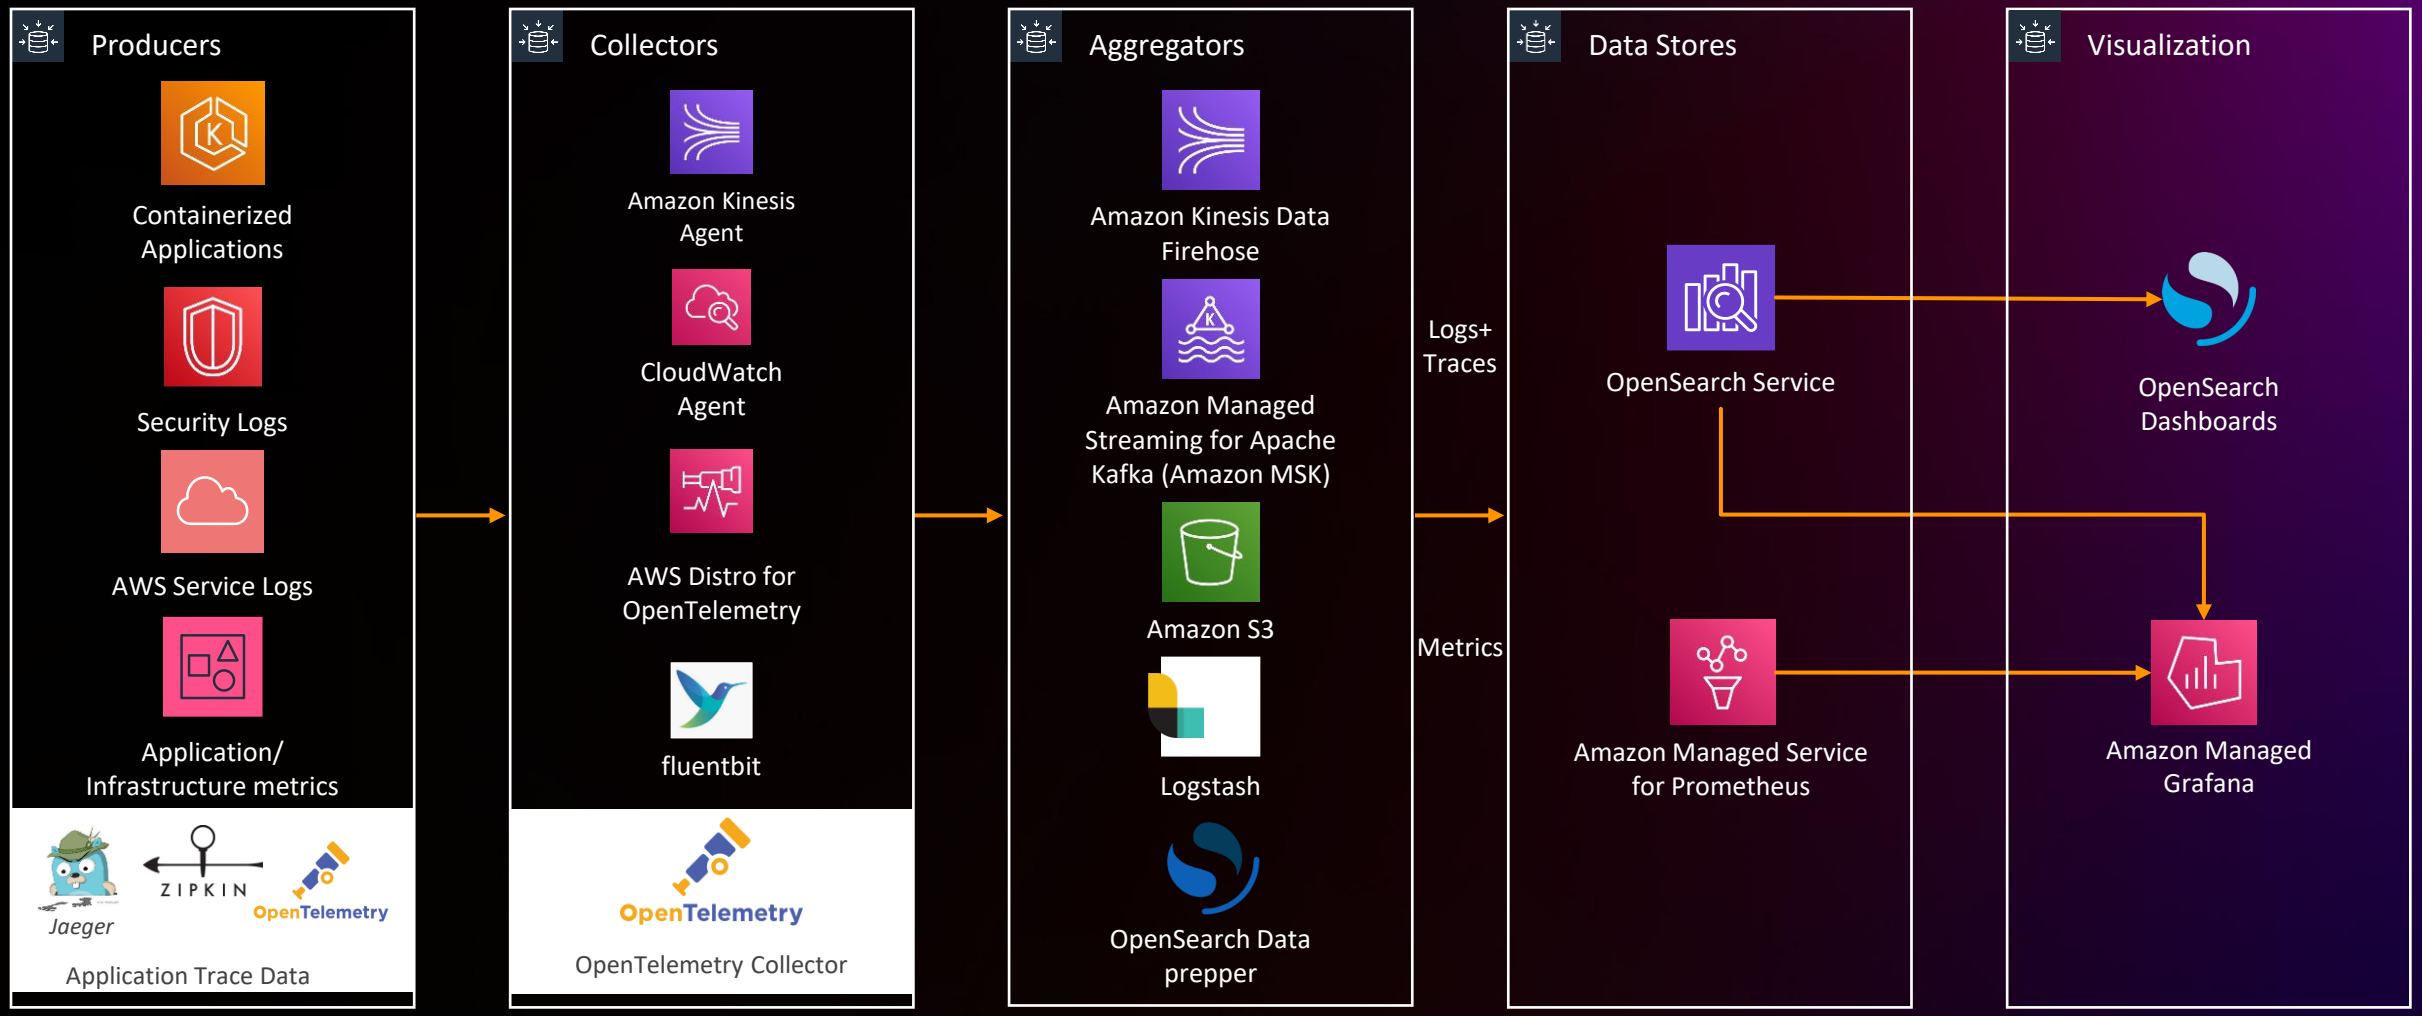

Lab: Microservice observability with Amazon OpenSearch Service workshop

Workshop objectives

1. Discover new observability capabilities in OpenSearch Service
   - Trace analytics and application analytics
   - Log-trace correlation (metrics coming soon)
   - Interactive log visualization
2. Use data collection and ingestion technologies
   - AWS Distro for OpenTelemetry to collect traces
   - Fluent Bit to collect logs
   - OpenSearch Data Prepper for transformation and ingestion
3. Perform a root cause analysis using application analytics

Use case statement

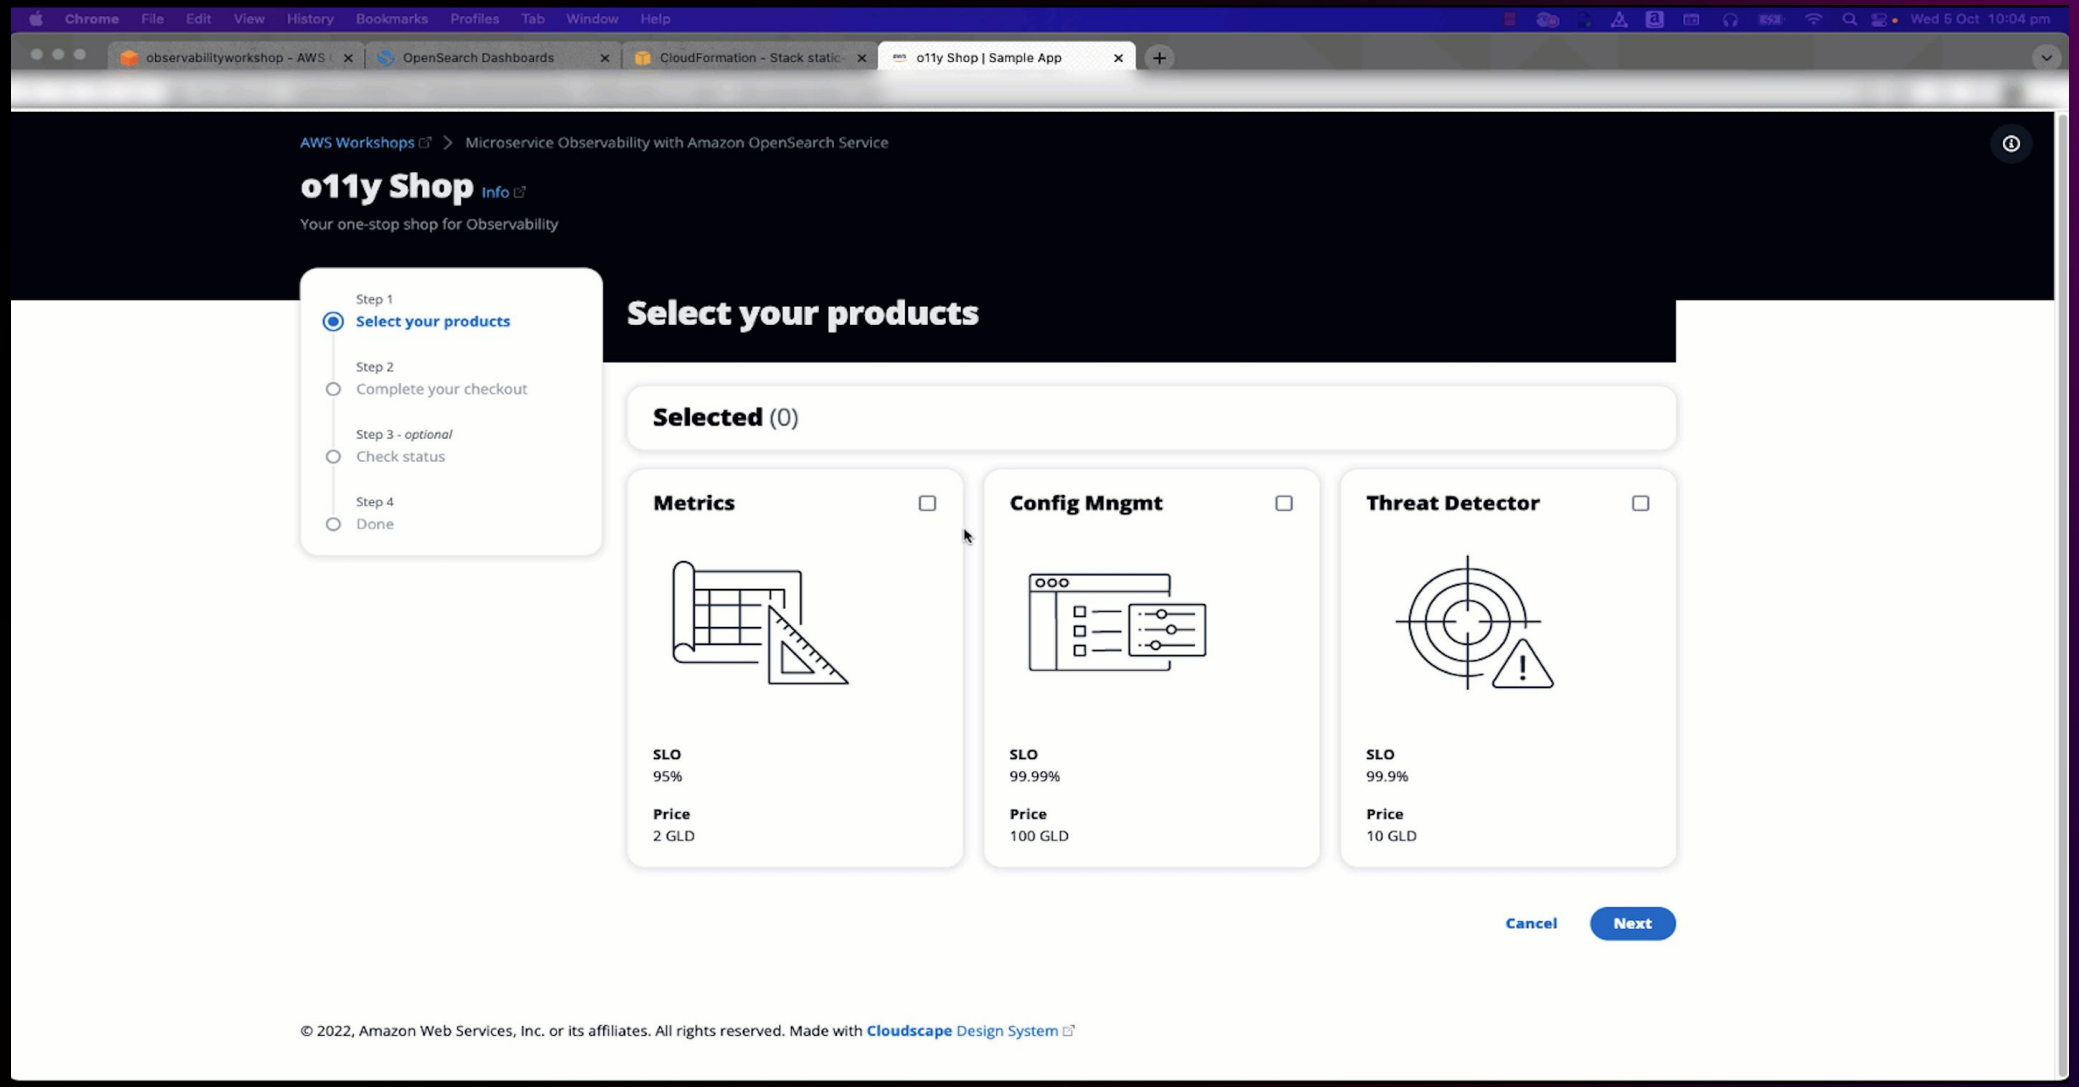

Workshop architecture

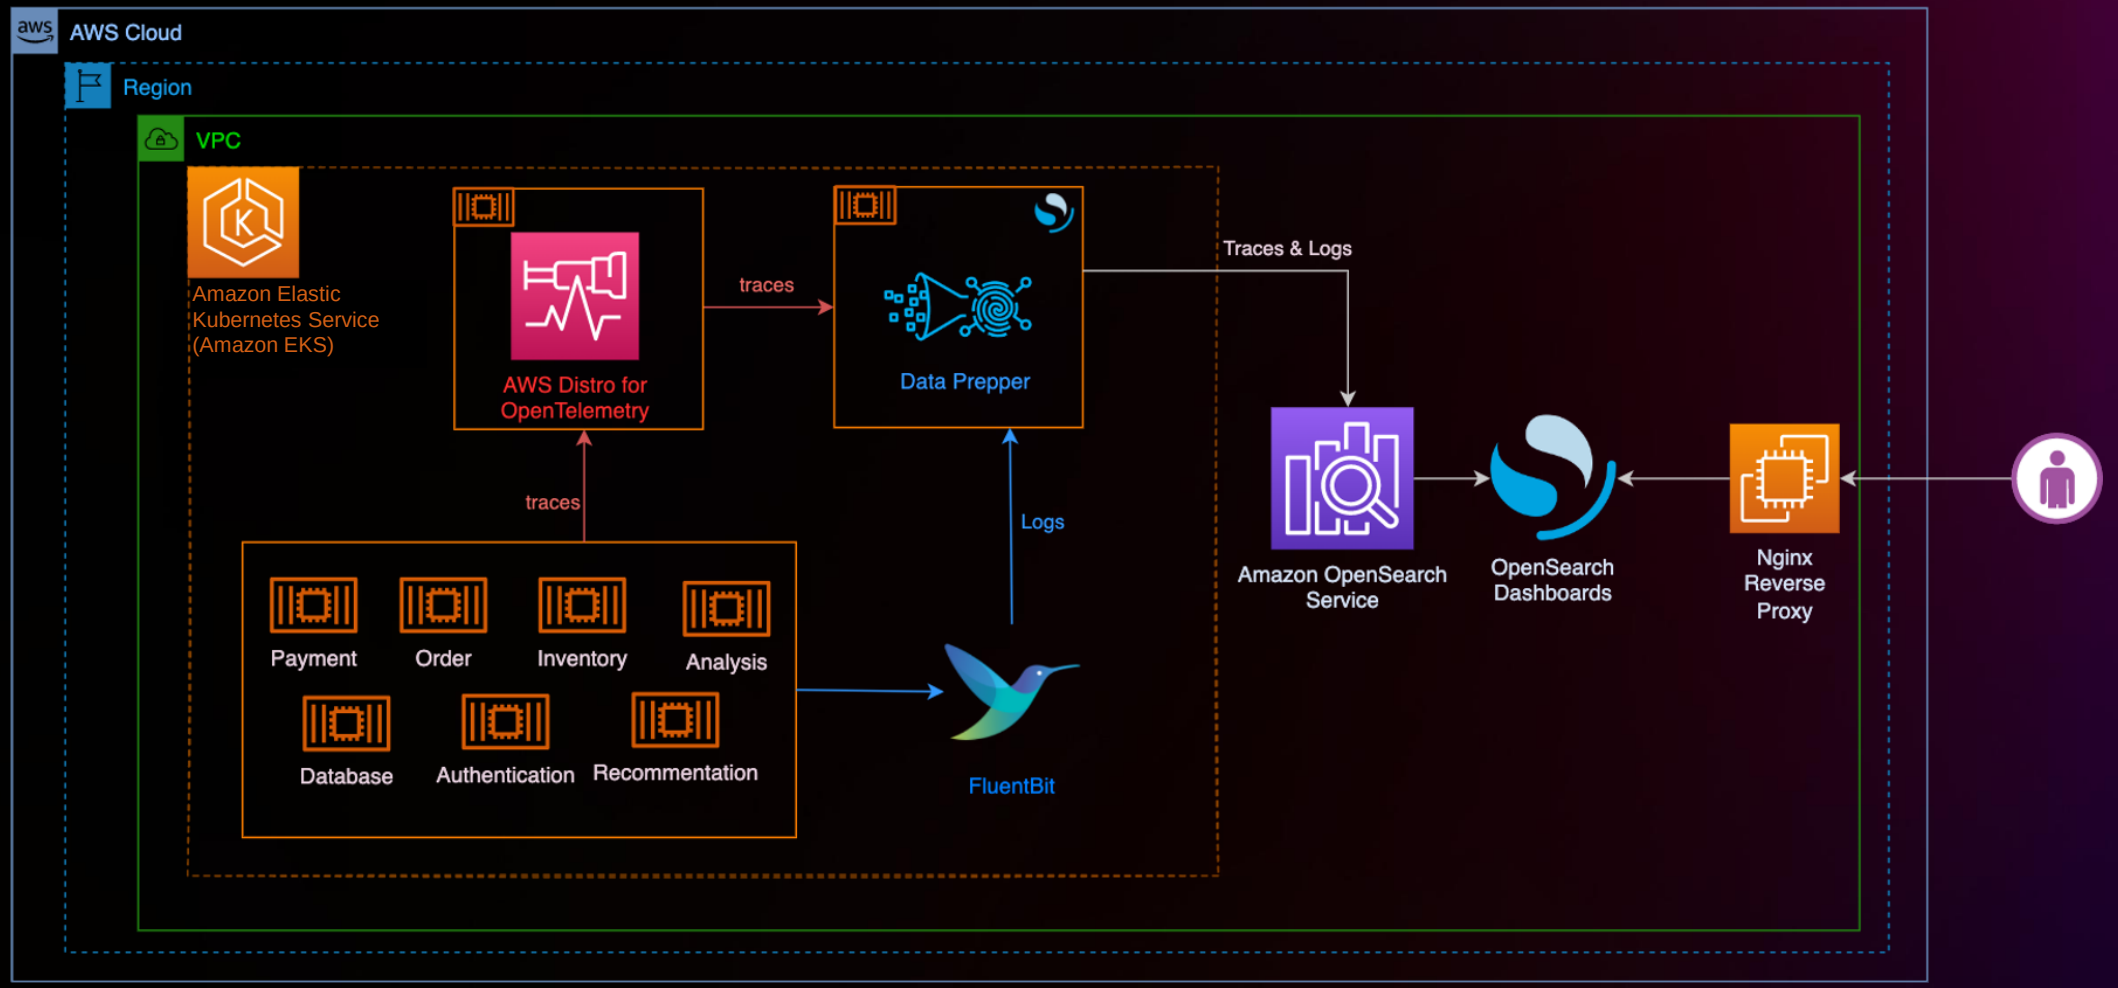

Root cause analysis workflow

1. Define your observability application composition: log sources, services, and entities
2. Observe the latency, error rate, and trends associated with each trace group and service
3. Filter the traces and spans involved in the failing service
4. Correlate the log events and the trace data to identify the issue
5. Fix the code line identified in error log message
6. Observe the system after changes

Lab setup

How to access the workshop lab

Lab instructions

https://bit.ly/3kDk9vp

Lab access

http://bit.ly/3giC9Nb

Access Code: 2596-03af0c-97

Let’s build!

Checkpoint

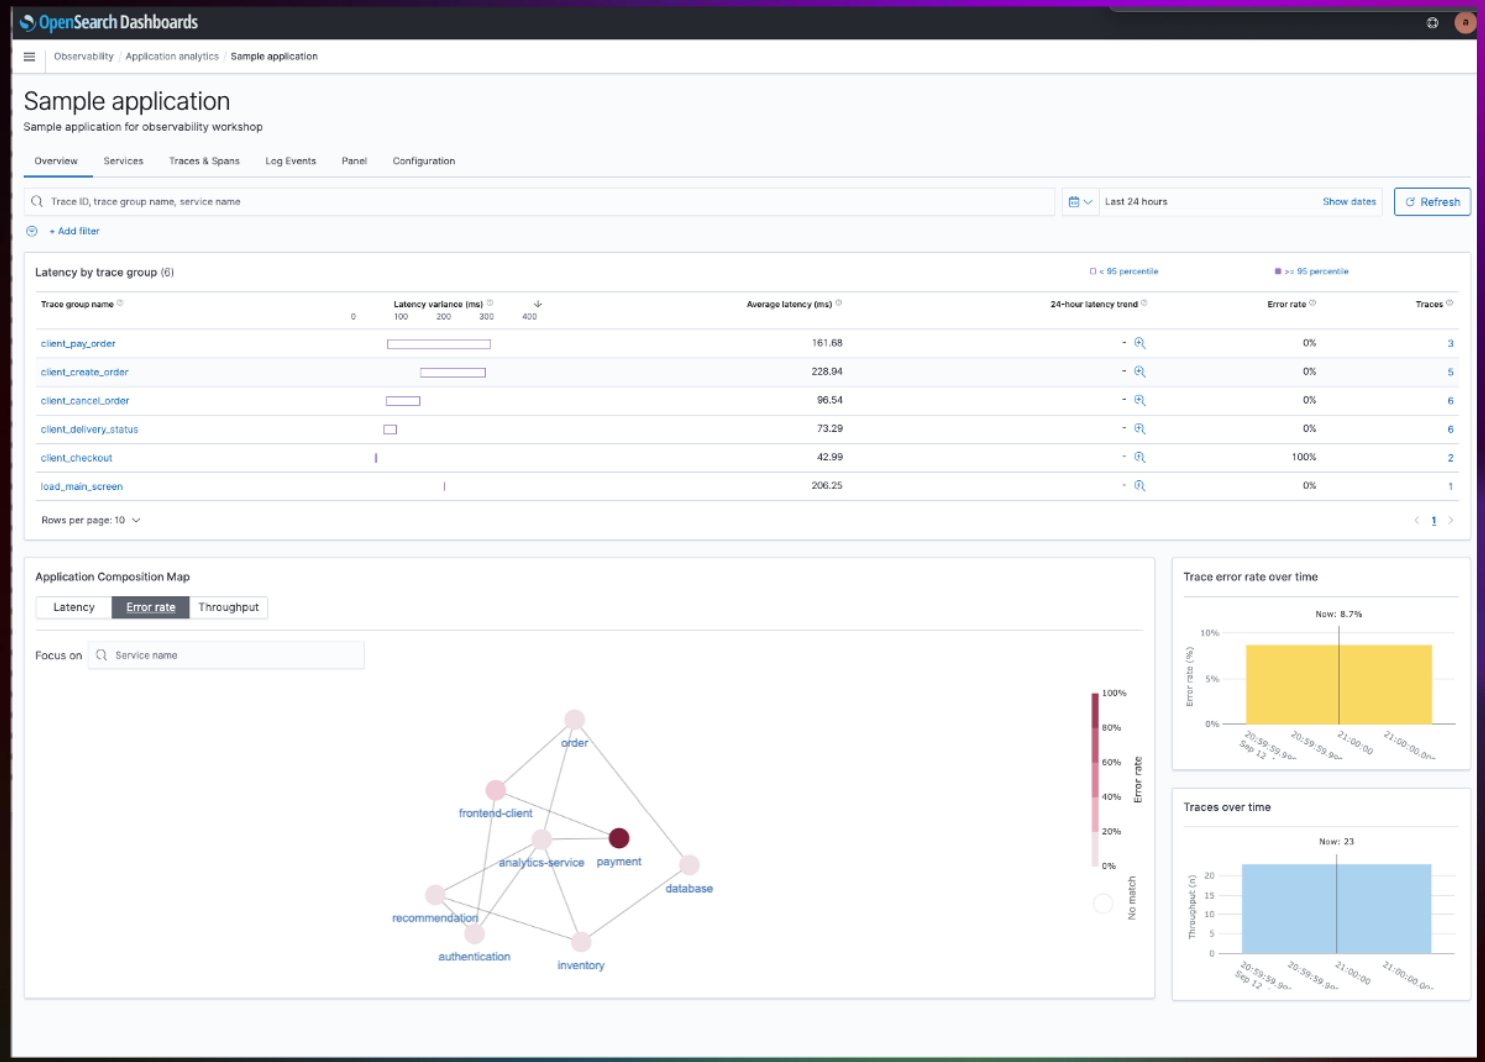

Thank you!

Please complete the session survey in the **mobile app**