

AWS re:Invent

NOV. 28 – DEC. 2, 2022 | LAS VEGAS, NV

ARC305

Reducing your area of impact and surviving difficult days

Bruno Emer

Principal Solutions Architect
Critical Capabilities
AWS

Byron Arnao

Principal Solutions Architect
Critical Capabilities
AWS



© 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.

A black and white portrait of Werner Vogels, a middle-aged man with a beard and mustache, wearing a dark jacket over a dark t-shirt. He is looking directly at the camera with a slight smile. The background is a blurred outdoor scene with trees and a building.

**“Everything fails,
all the time.”**

Werner Vogels

VP and CTO, Amazon.com




ed the Turing tes


Categories of failure




Code deployments
and configuration




Data and state



Dependencies

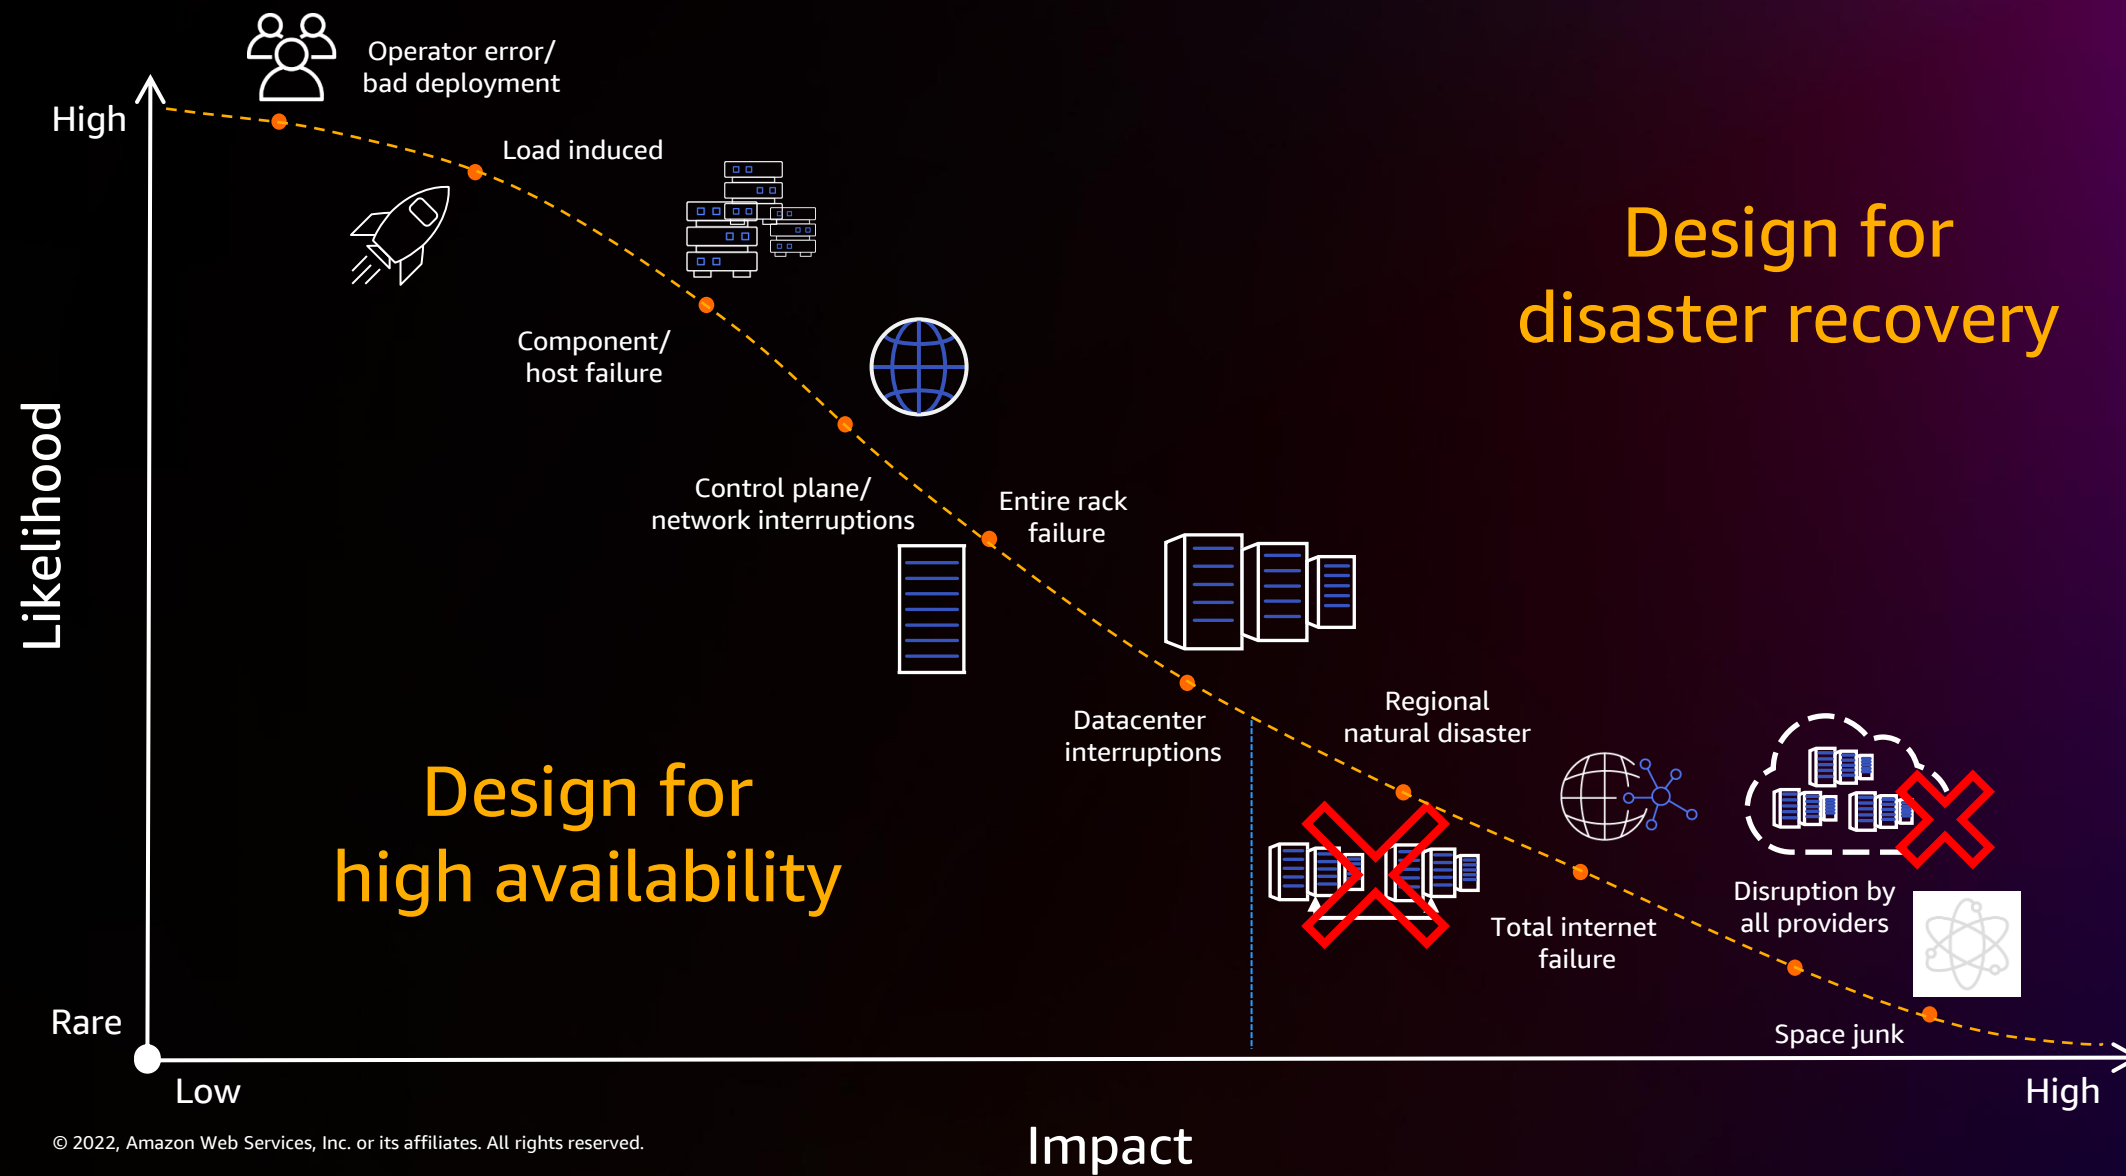


Core
infrastructure









Highly unlikely
scenarios

Various causes of failure



How do we usually protect against failures?

Failure mode			Means of protection	
Rack down		←	Autoscaling groups	✓
AZ impairment		←	Multi-AZ deployments	✓
Regional impairment		←	Region failover	✓
Bad deployment		←	Canary deployments	!
Human failure		←	Processes, runbooks	!
Poison pill		←	Extensive testing	!

Typically only for compute

Can't be fully mitigated

What happens when a workload becomes too big to fail?


**Let's talk about a
real-world scenario . . .**

How do we do this today?


Designing resilient architectures

High Availability


Multi-zonal active-active



Multi-region active-active




Advanced resilience




Traditional enterprise recovery



Multi-zonal failover

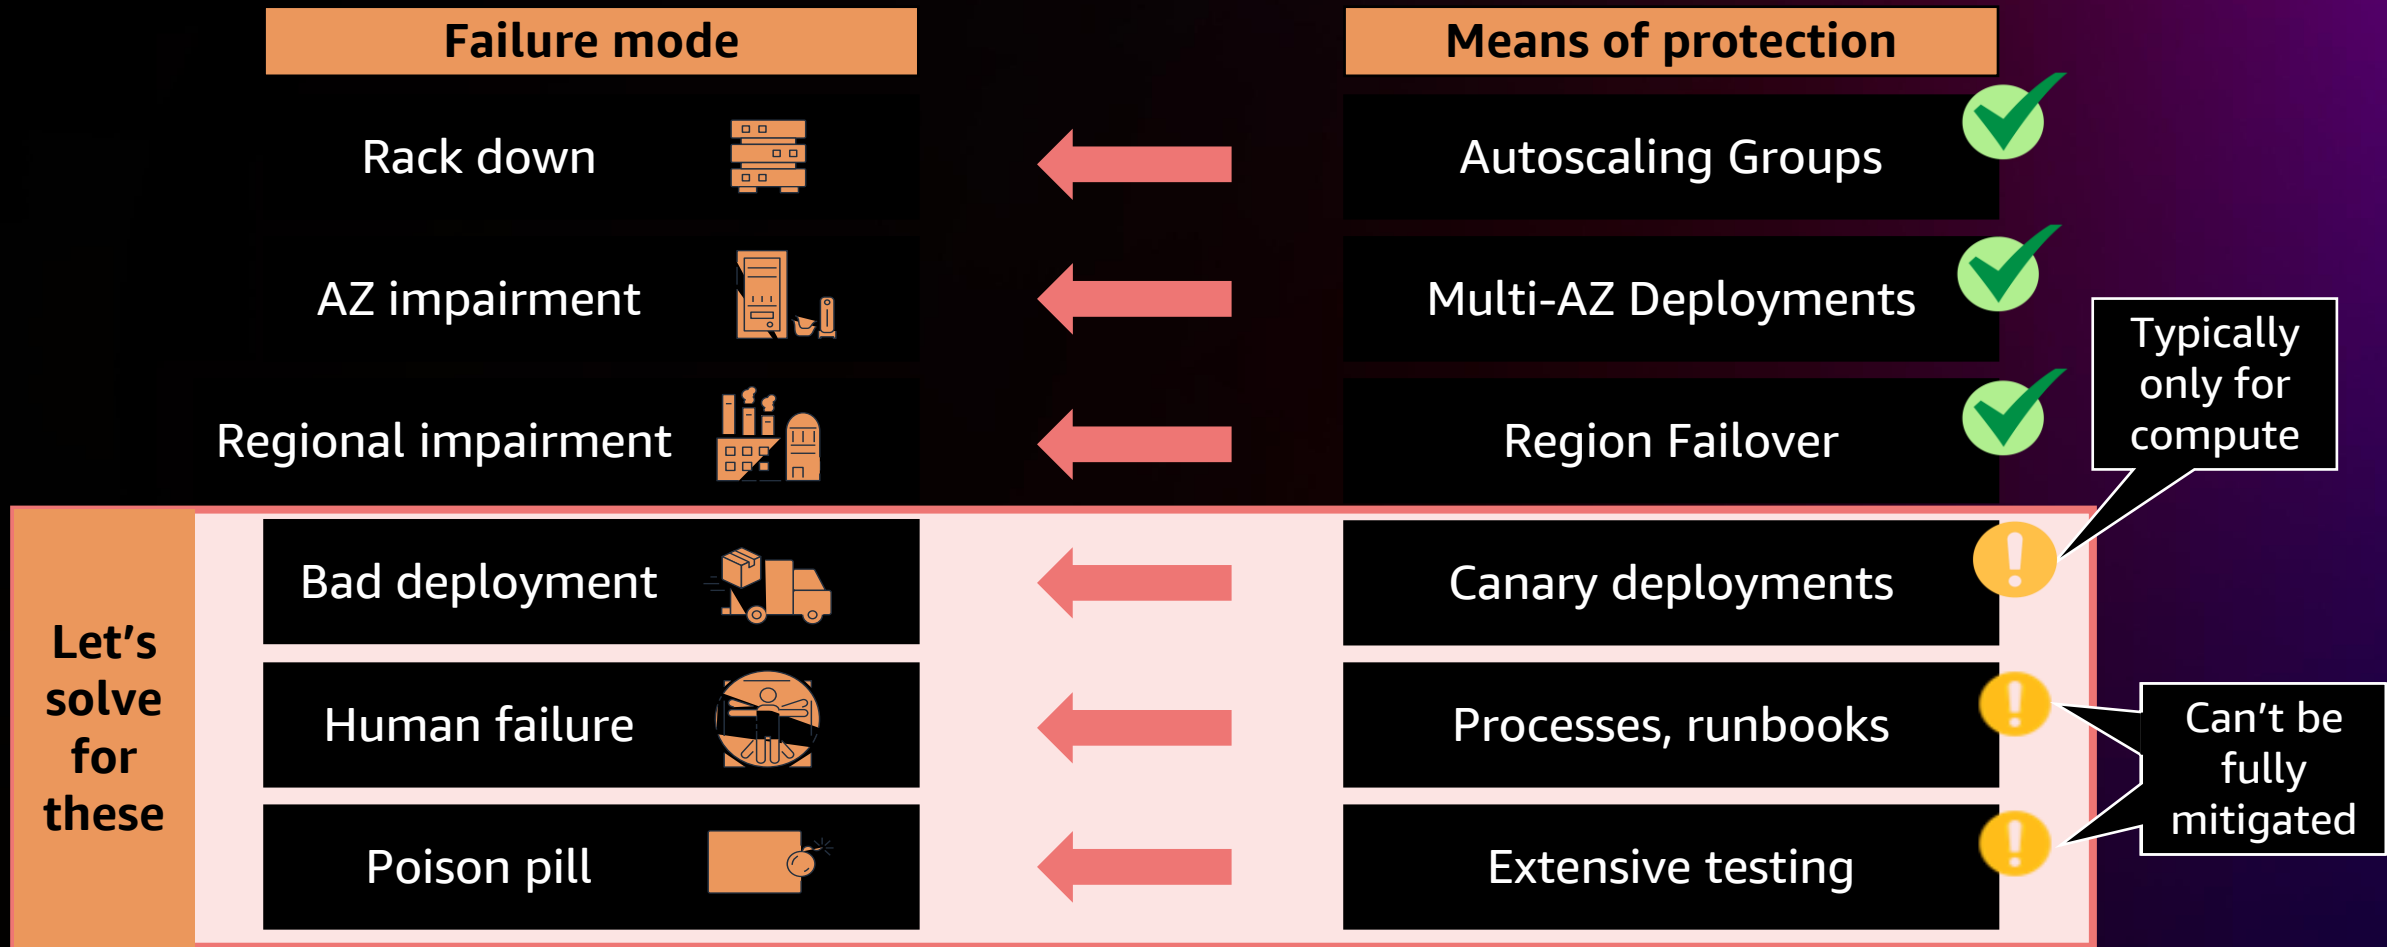


Multi-region failover




Disaster Recovery

How do we usually protect against failures?




Limiting the impact of failures



- Bulkhead pattern
- Concept from **naval engineering**
- Internal chambers **isolate** the hull
- This prevents water from filling the **entire ship**

How do we build a bulkhead with software?


Cell-based architectures




- Multiple copies of the **entire application** deployed in each cell
- Data is **partitioned** – there is no replication between cells
- **Complete isolation** between cells limits and contains failure
- Cells provide a **predictable scale** unit

Cell-based architectures properties

- Workload **isolation**
- Failure **containment**
- **Scale-out** vs. scale-up
- Cells have **maximum size**
- **Testability**
- **Manageability**



Cells reduce the area of impact for failures




+

Failure



"The application was down for everyone!"



+


Failure



"Some users report problems accessing the application"


How to achieve the benefits of isolation with access patterns?

Traditional architecture

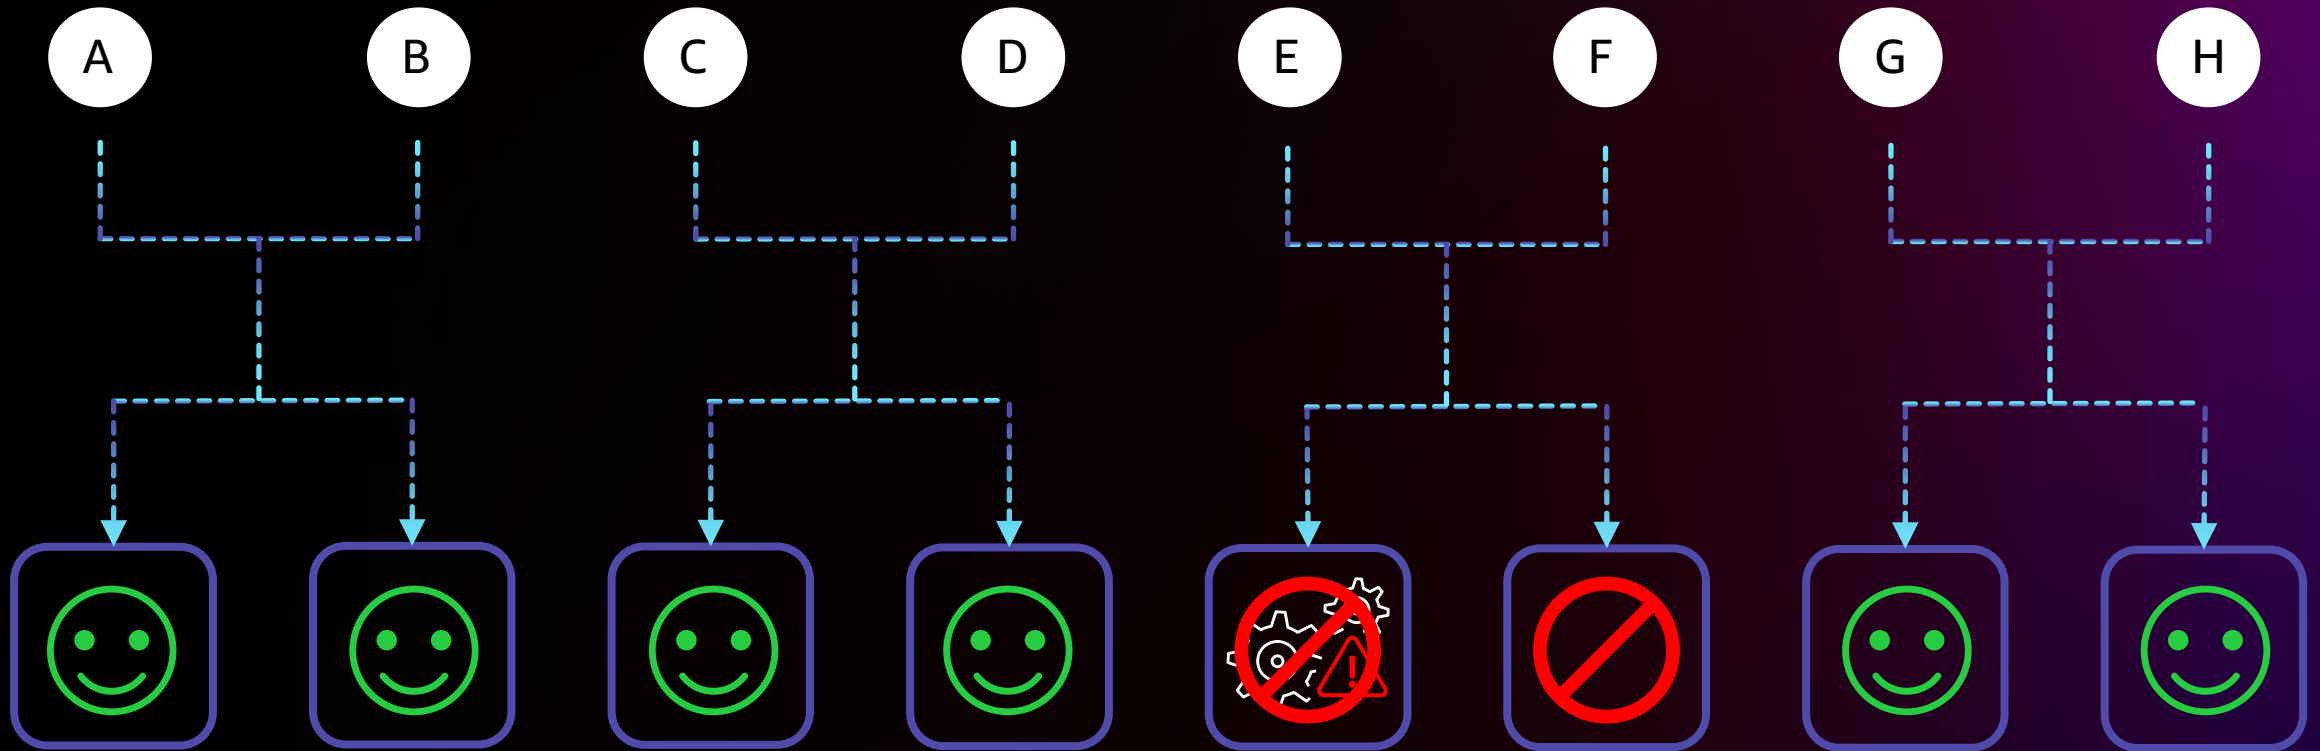


Area of impact = All clients

Sharding




Sharding



$$\text{Area of impact} = \frac{\text{Clients}}{\text{Shards}}$$

How will we route the requests to the right cells?


Routing mechanisms



- Add a **thin** routing layer
- Needs to be **resilient**
- Keep it **simple**
- **Statically stable**
- **Ephemeral**


Routing mechanisms

Routing through DNS



- ✓ Can leverage Amazon Route 53 HA
- Clients need to map users to DNS names

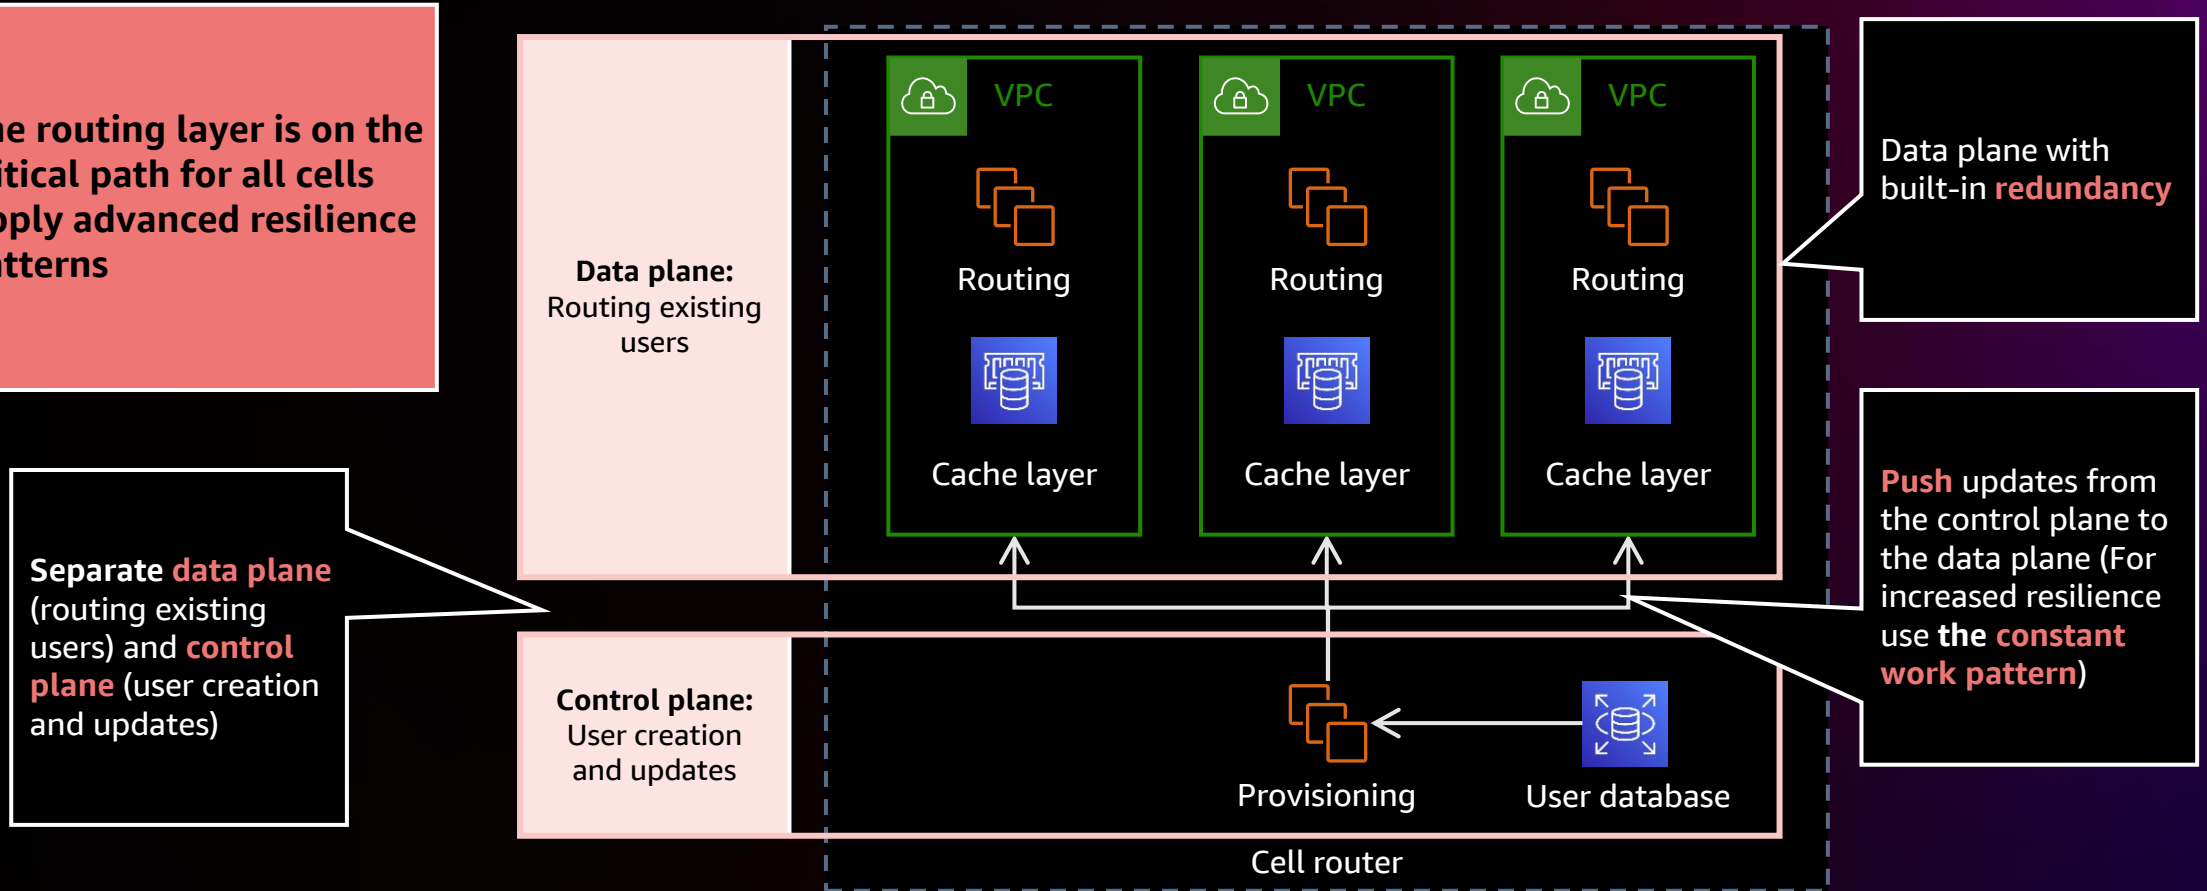
Cell router as load balancer



- ✓ Transparent to clients
- Cell router is critical to ongoing transactions



Hardening the routing layer

- The routing layer is on the critical path for all cells
- Apply advanced resilience patterns



Cells are not Availability Zones or Regions

Cells are a logical isolation mechanism

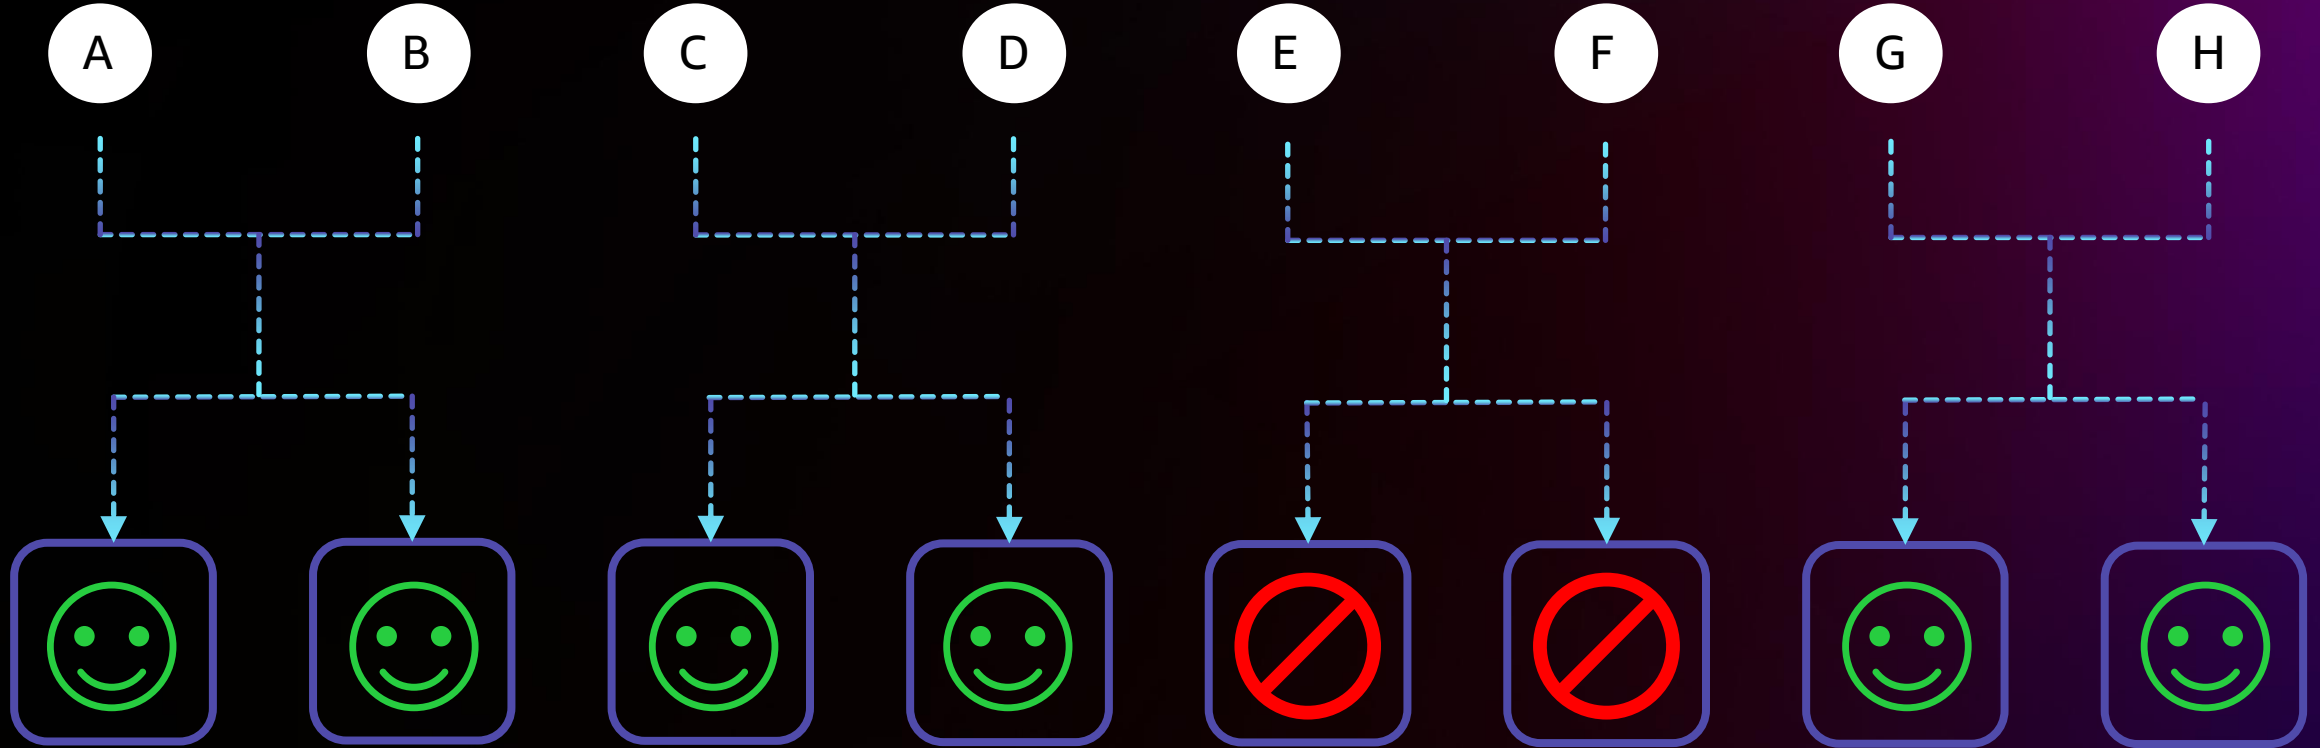


- Availability Zones and Regions provide **redundancy** and **physical isolation**
- Cells provide **compartmentalization**
- Design for the **combination**; cells span AZs or Regions (if necessary)


Cells are not a scaling mechanism

Can we make it even more resilient?


Sharding




Shuffle sharding




Shuffle sharding



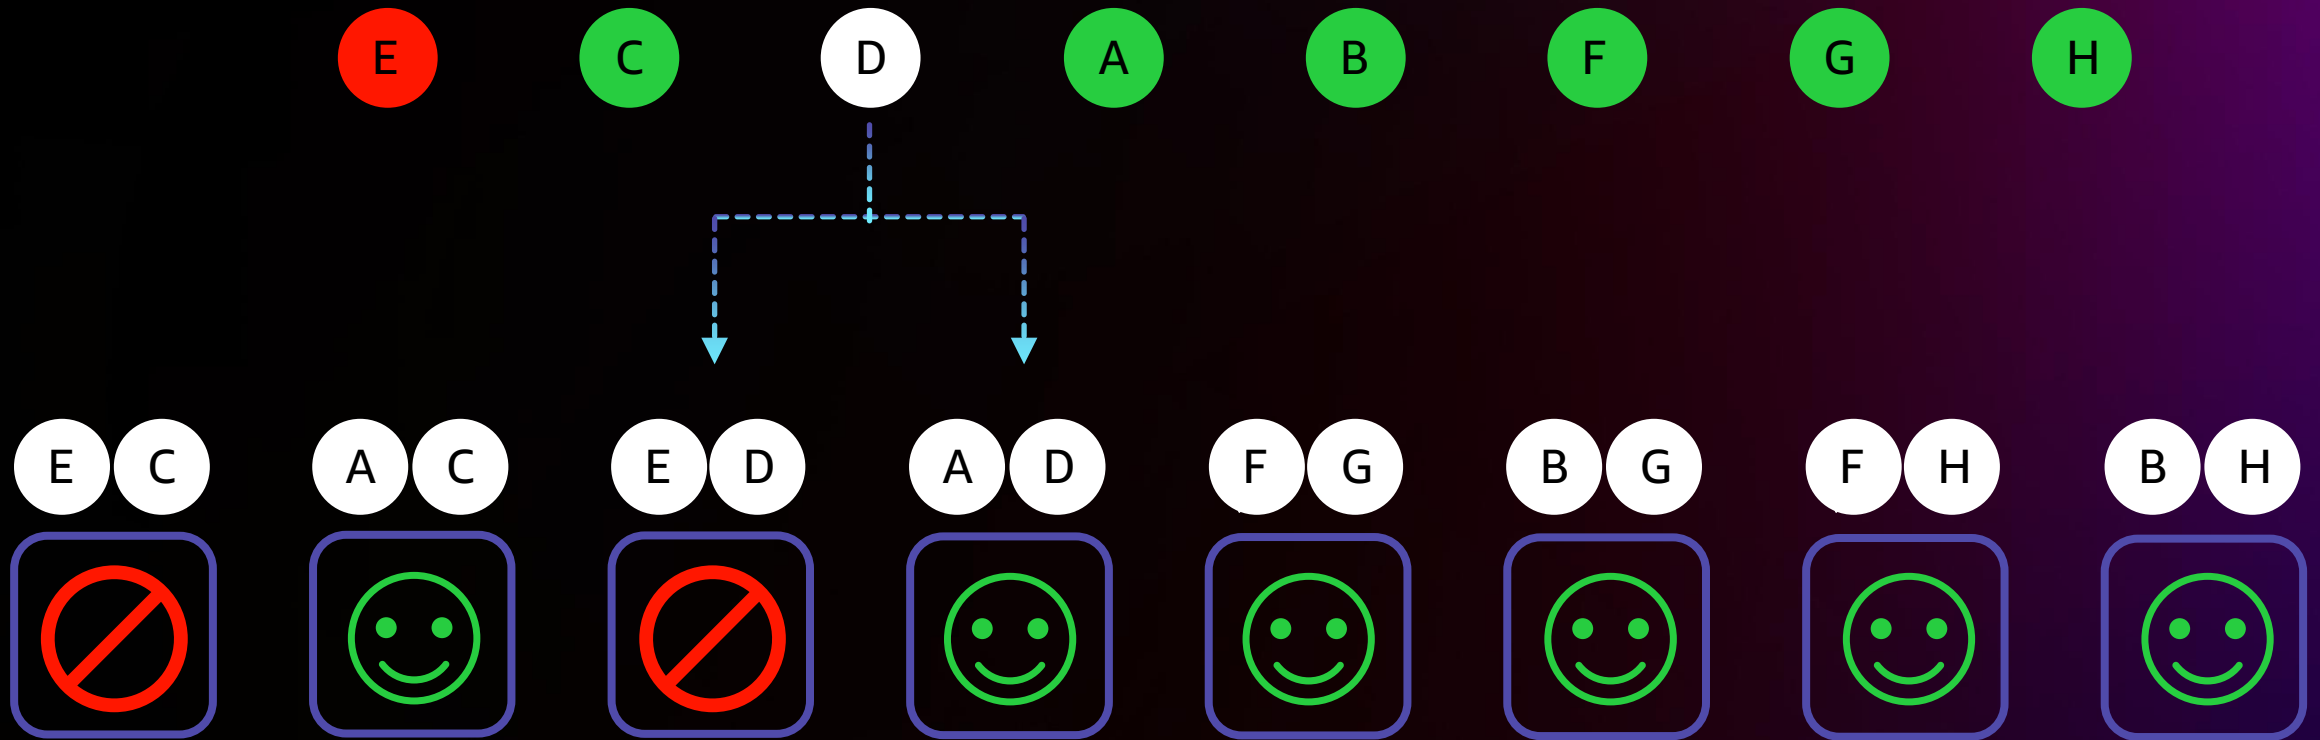
Shuffle sharding



Shuffle sharding



Shuffle sharding



$$\text{Area of impact} = \frac{\text{Clients}}{\text{Combinations}}$$

Observations

- Traditional architecture – complete outage
- Sharding
 - Impact localized to customers on the same shard
 - 25% of customers affected
- Shuffle sharding
 - Impact localized to customers having the same combination of nodes
 - 12.5% of customers affected

Observations

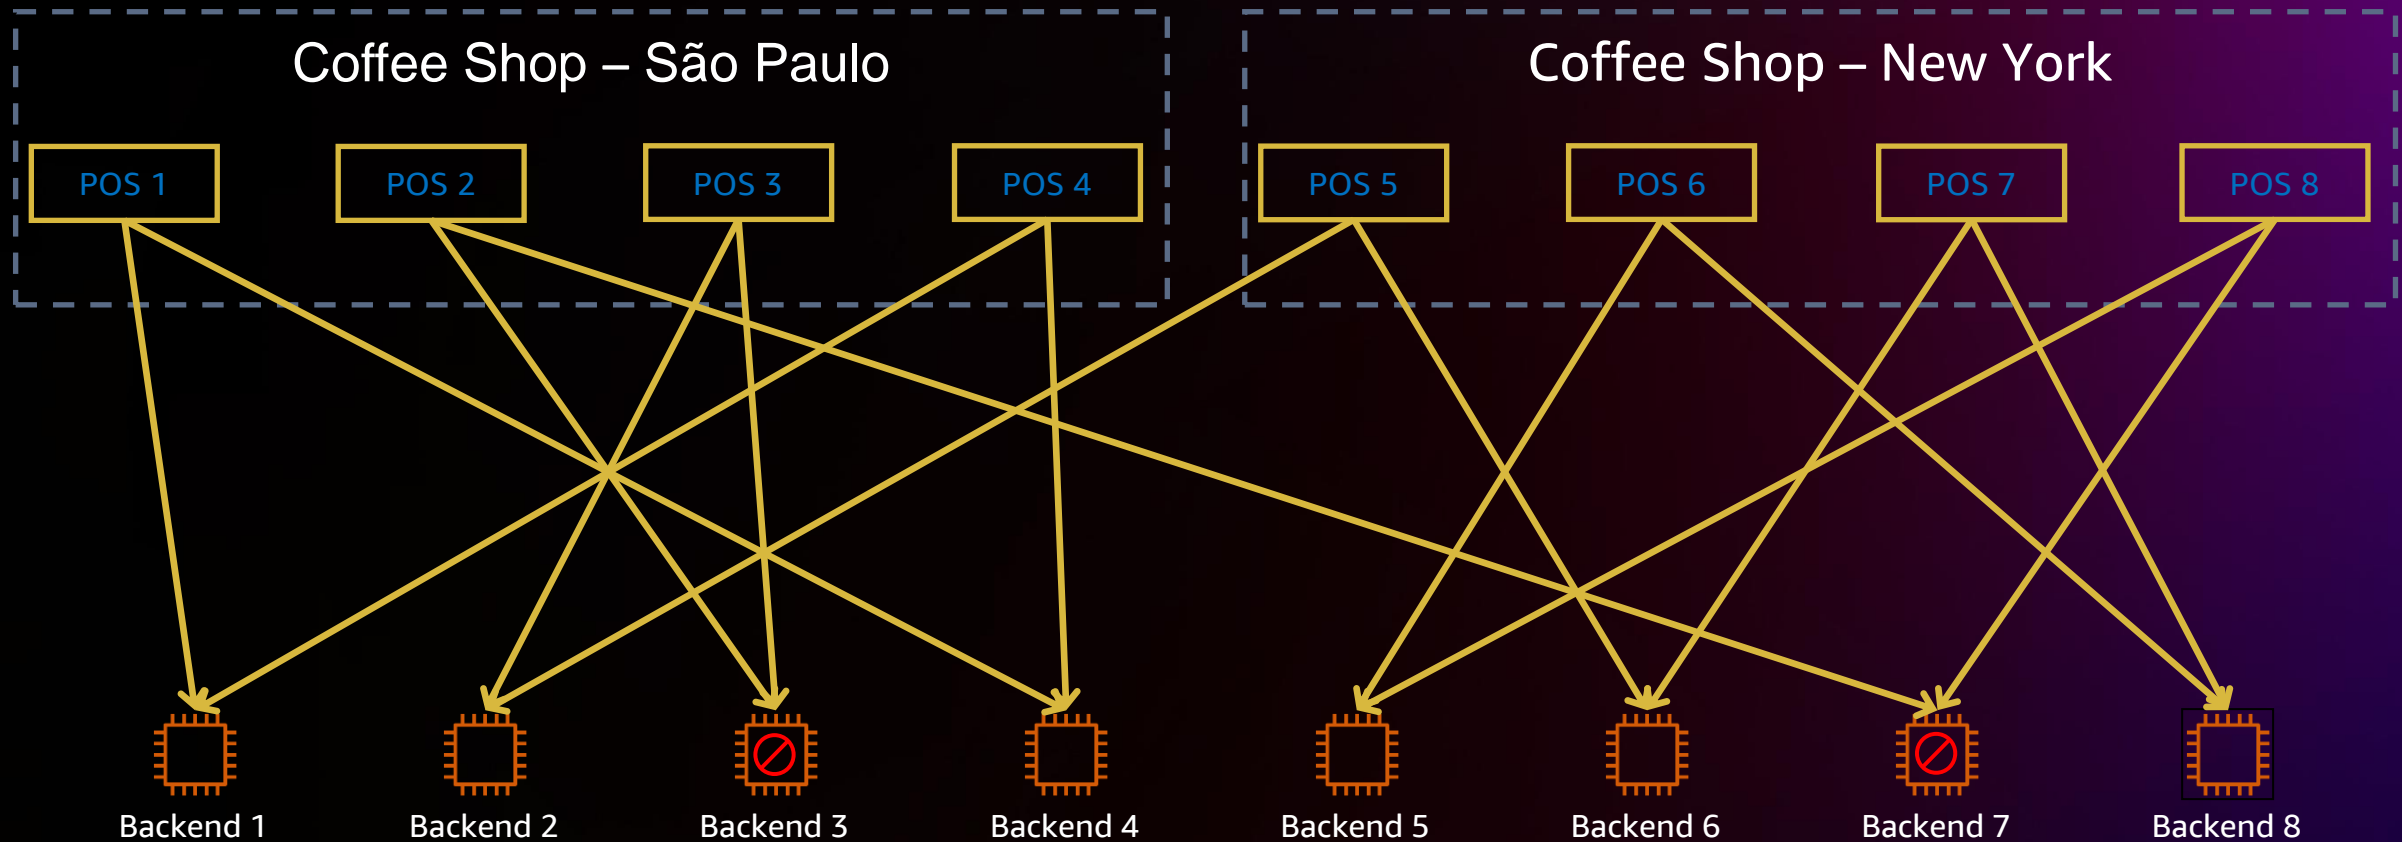
100 nodes, 5 nodes per shard

Overlap	% customers
0	77%
1	21%
2	1.8%
3	0.06%
4	0.0006%
5	0.0000013%



Amazon Route 53




Architecture – Shuffle sharding



How do we operate these cells?

Cell management

- Always build a **rebalancing/migration tool**
- Think about **load distribution** issues
- Cells can **fail** at any time
- Tech stack will **influence** management



*Sharding based on customers

Monitoring in cell-based architectures

Monitor each cell individually

Aggregate metrics into
healthy/unhealthy cells per workload


Make sure messages and errors can
easily be correlated with cells

Is it better to have small or big cells?

Small versus large cells

Smaller cells


- Reduced area of impact
- Easier to test
- Cells easier to operate



VS.

Larger cells

- Cost efficiency
- Reduced splits
- System easier to operate



It depends!

Other topics to consider

Authentication?

Where are credentials stored?

How do we **network**?

Does each cell have its own
SSL certificate?

What do we have on a cell basis?
VPCs? **AWS accounts**?

Where are **team boundaries**?
Is each type of cell maintained by
one or multiple teams?

How do we handle
infrastructure failures?
How do we do **disaster recovery**?

**So, this solution
will solve . . .**

Recap . . .

- Cells provide logical isolation
 - There is a need to focus on data modeling
 - Build a thin, statically stable routing layer
- Monitoring changes at a cell and service level
 - Think about deployments
 - Cell orchestration needs to be implemented
 - Extreme resilience can be achieved with shuffle sharding
- Cells are not AZs
 - Cells are not Regions
 - Cells are not a scaling mechanism
 - Data is not shared between cells
 - Cells have no inter-dependent logic

Thank you!



Please complete the session survey in the **mobile app**

