

AWS re:Invent

NOV. 28 – DEC. 2, 2022 | LAS VEGAS, NV

ARC306-R

Multi-Region design patterns and best practices

Neeraj Kumar

Principal Solutions Architect
AWS

John Formento, Jr.

Principal Solutions Architect
AWS

Barry Sheward

Chief Cloud Architect
The Vanguard Group



Agenda

Multi-Region fundamentals

The Vanguard Group's global multi-Region strategy

Multi-Region architecture patterns

Multi-Region fundamentals



Fundamental #1

Understand the requirements

**“We need to create a
multi-Region architecture.”**

Maybe, but first...

- What are the business requirements?
 - Do we need to meet an extreme resilience requirement?
 - Do we need to be able to prove the workload can run from another Region?
- Weigh the requirements against the cost and complexity

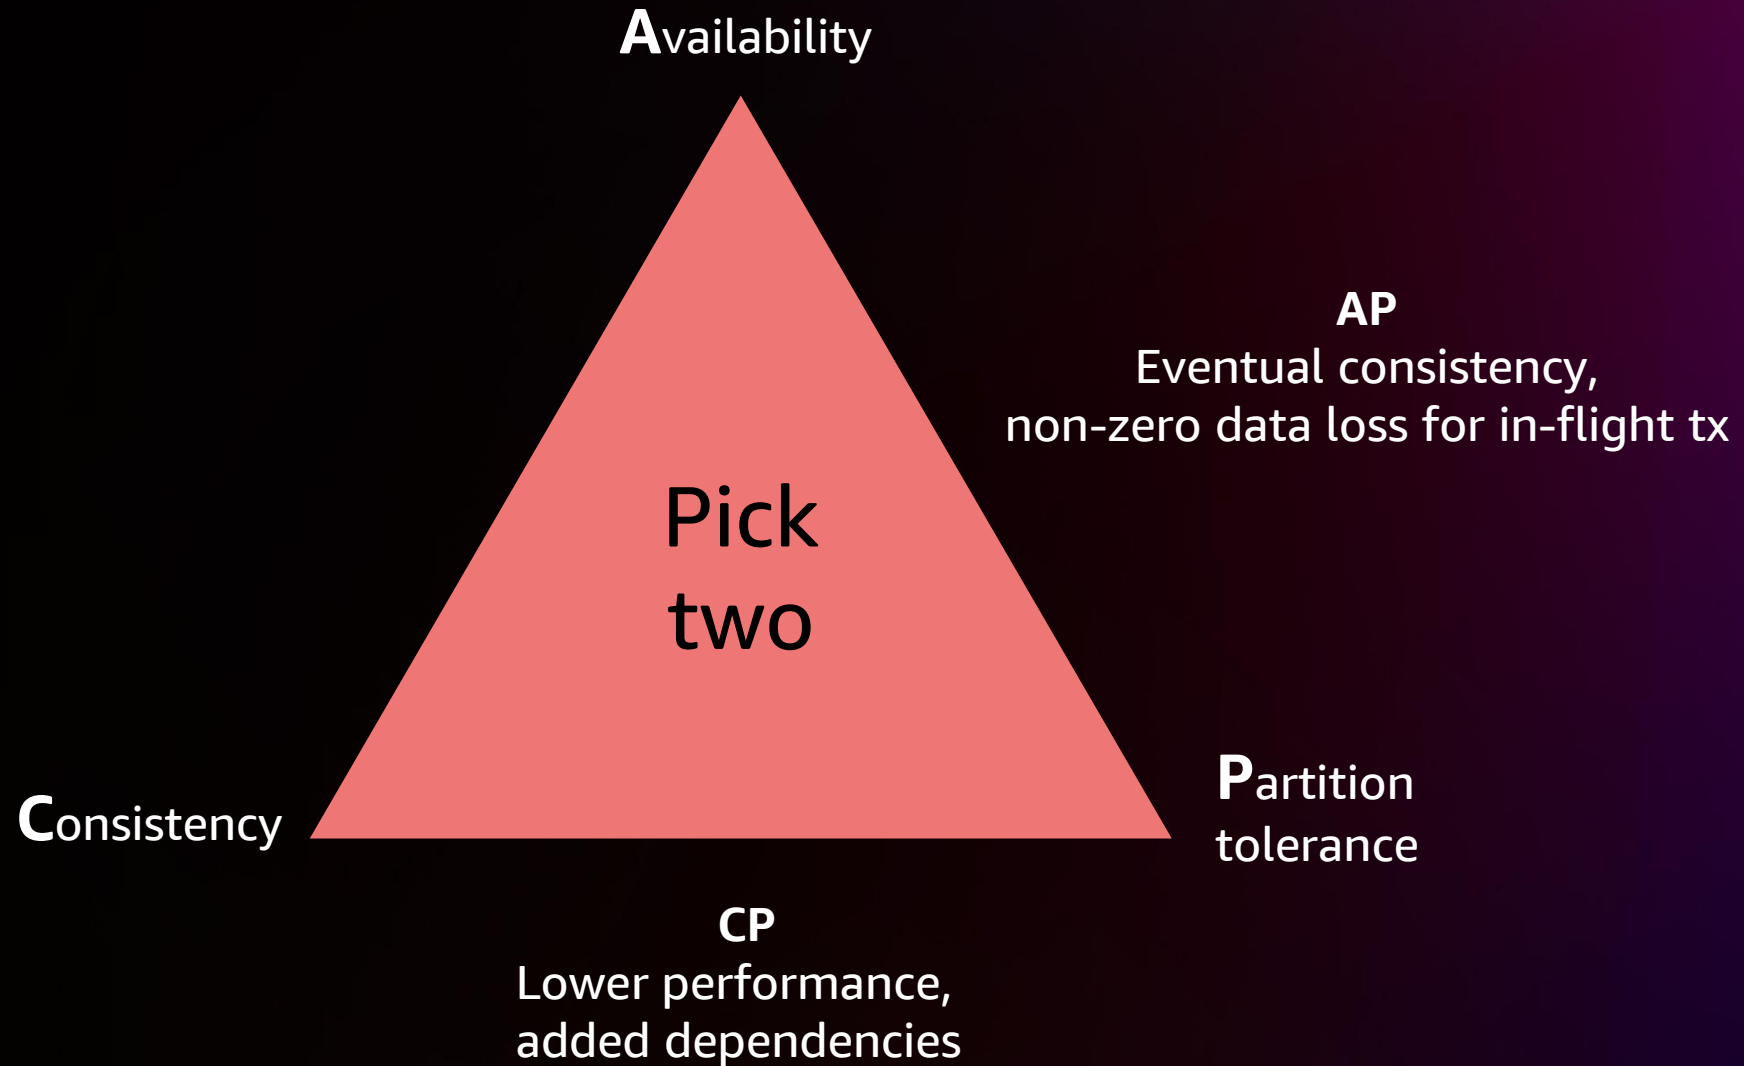
Tier	Max RTO	Max RPO	Criteria	Cost
Platinum	15 min	5 min	Mission-critical workloads	\$\$\$
Gold	15 min–6 hrs	2 hrs	Important, but not mission-critical workloads	\$\$
Silver	6 hrs–few days	24 hrs	Noncritical workloads	\$

Fundamental #2

Understand the data

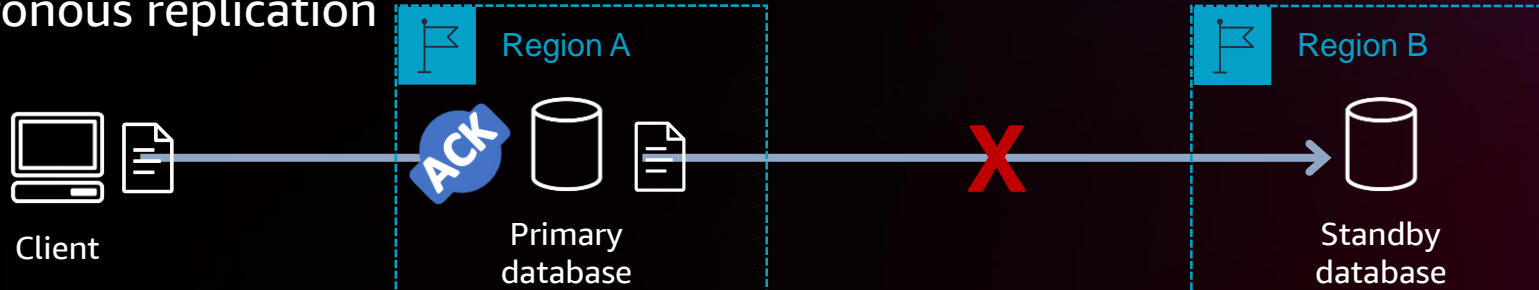


Data consistency requirements



Data consistency requirements

Asynchronous replication




Synchronous replication



Data access patterns

Understand the access characteristics




Read
Write

Fundamental #3

Understand dependencies

AWS services

- Determine what services are available in the desired Regions
- Identify what features will aid in your multi-Region journey



Compute




Storage



Database



Network
and Content
Delivery



Analytics



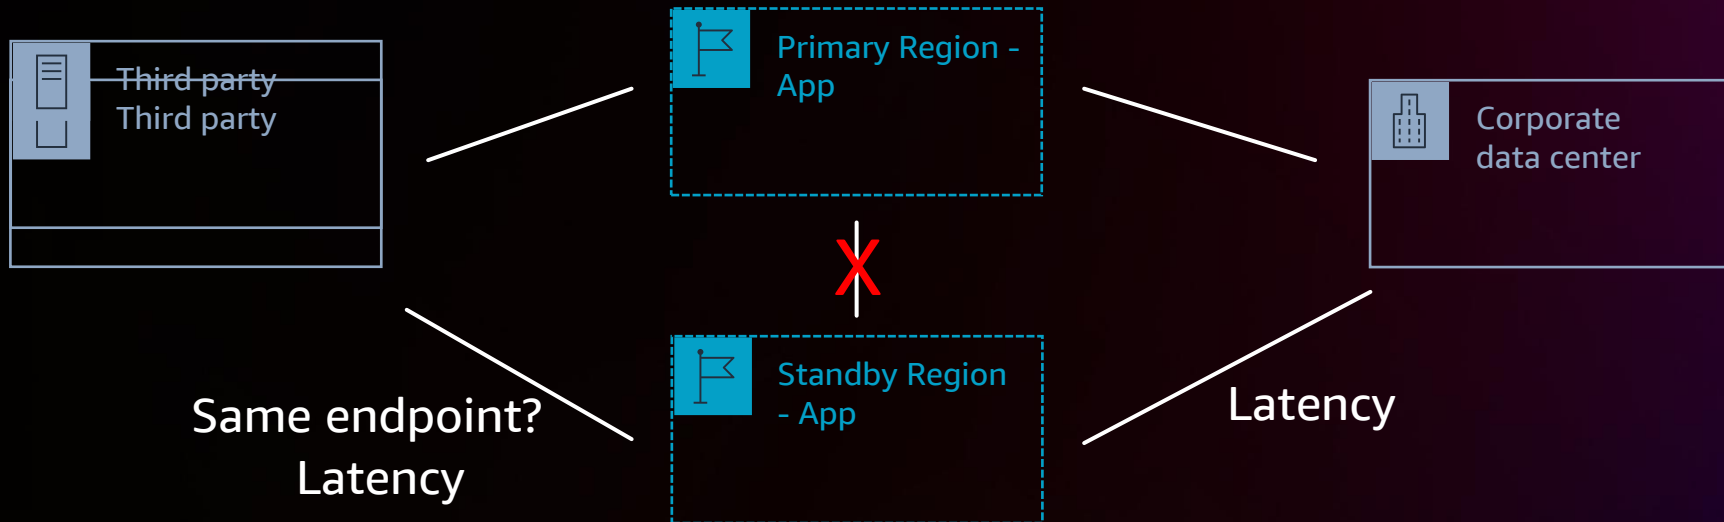
Management
and Governance



Security


Internal and third-party dependencies

- Identify any on-premises dependencies
- Identify any third-party solutions in the workload



Failover mechanisms

- Never take a dependency on your primary Region for failover
- Scrutinize all dependencies in your failover mechanism

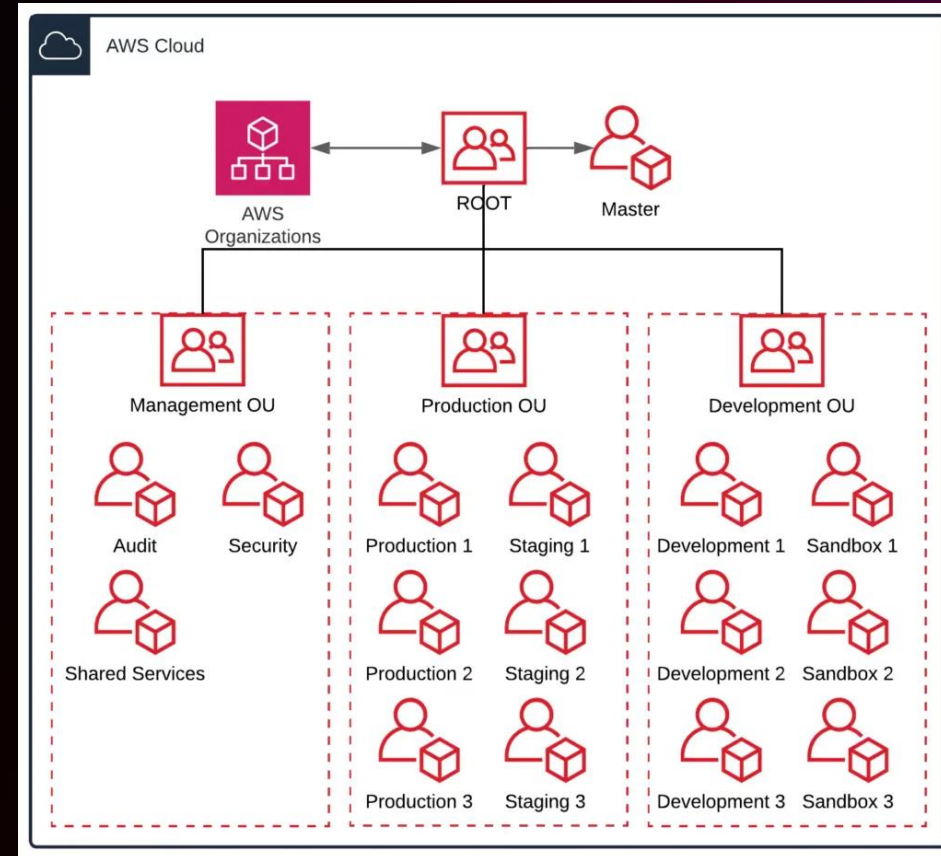


Fundamental #4

Operational readiness



AWS Account Management

- AWS service quotas
- IAM permissions
- Service control policies




Deploy to a Region at a time

BLUE/GREEN OR CANARY DEPLOYMENTS



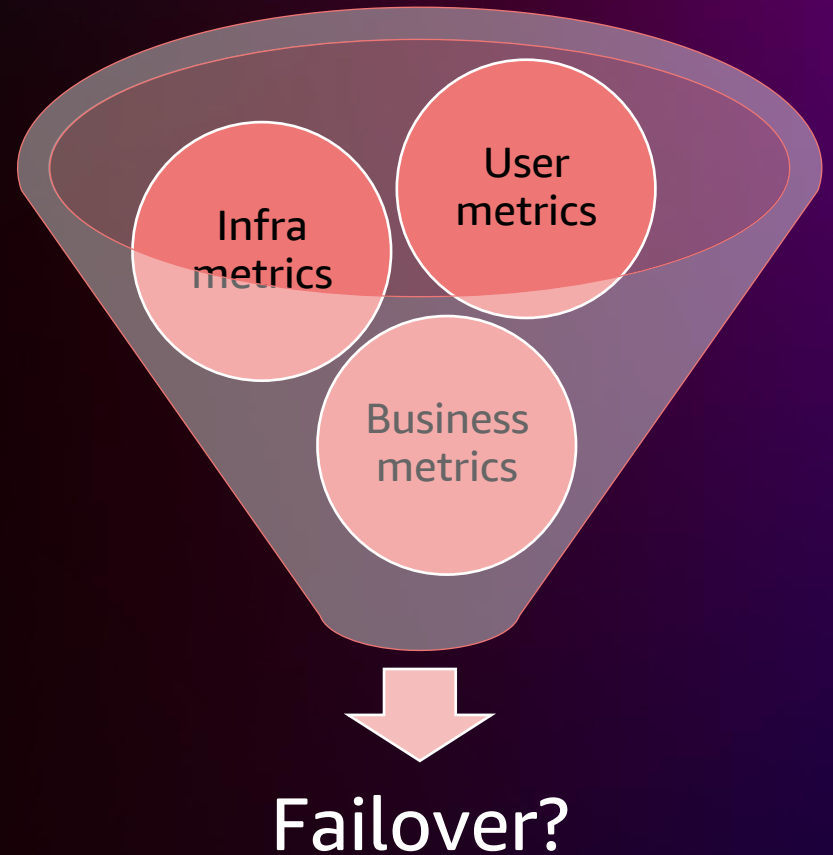
Monitoring: What changes in multi-Region?



- Replicate only high-level metrics to the other Region
 - User experience
 - User count
 - Transaction count
 - Replication status
- Monitor the health of the application in Region A from tooling in Region B
- Use telemetry from Region B for determining when to failover from application in Region A

Processes and people

- Untested DR strategy = no DR strategy
- Scope of failover
 - Individual microservices?
 - Business capability?
- Not only does the failover mechanism need to be tested regularly
 - Process for when to failover needs to be defined



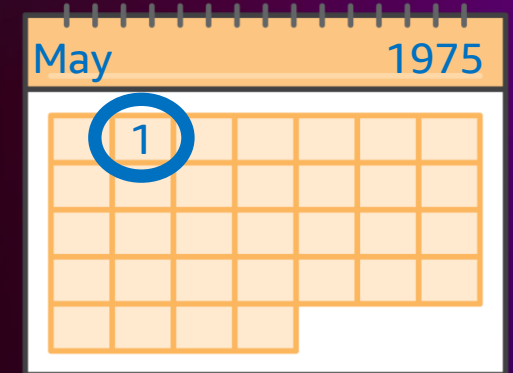
Cost and complexity

- Engineering effort
 - Is the application(s) ready for multi-Region?
- Operational overhead
 - People and process
 - Testing
 - Enhanced deployment processes
- Cost
 - Isolation of application stacks between Regions can result in 2x cost
 - Increased costs from the operational efforts and engineer efforts are on top of infrastructure costs

The Vanguard Group's global multi-Region strategy

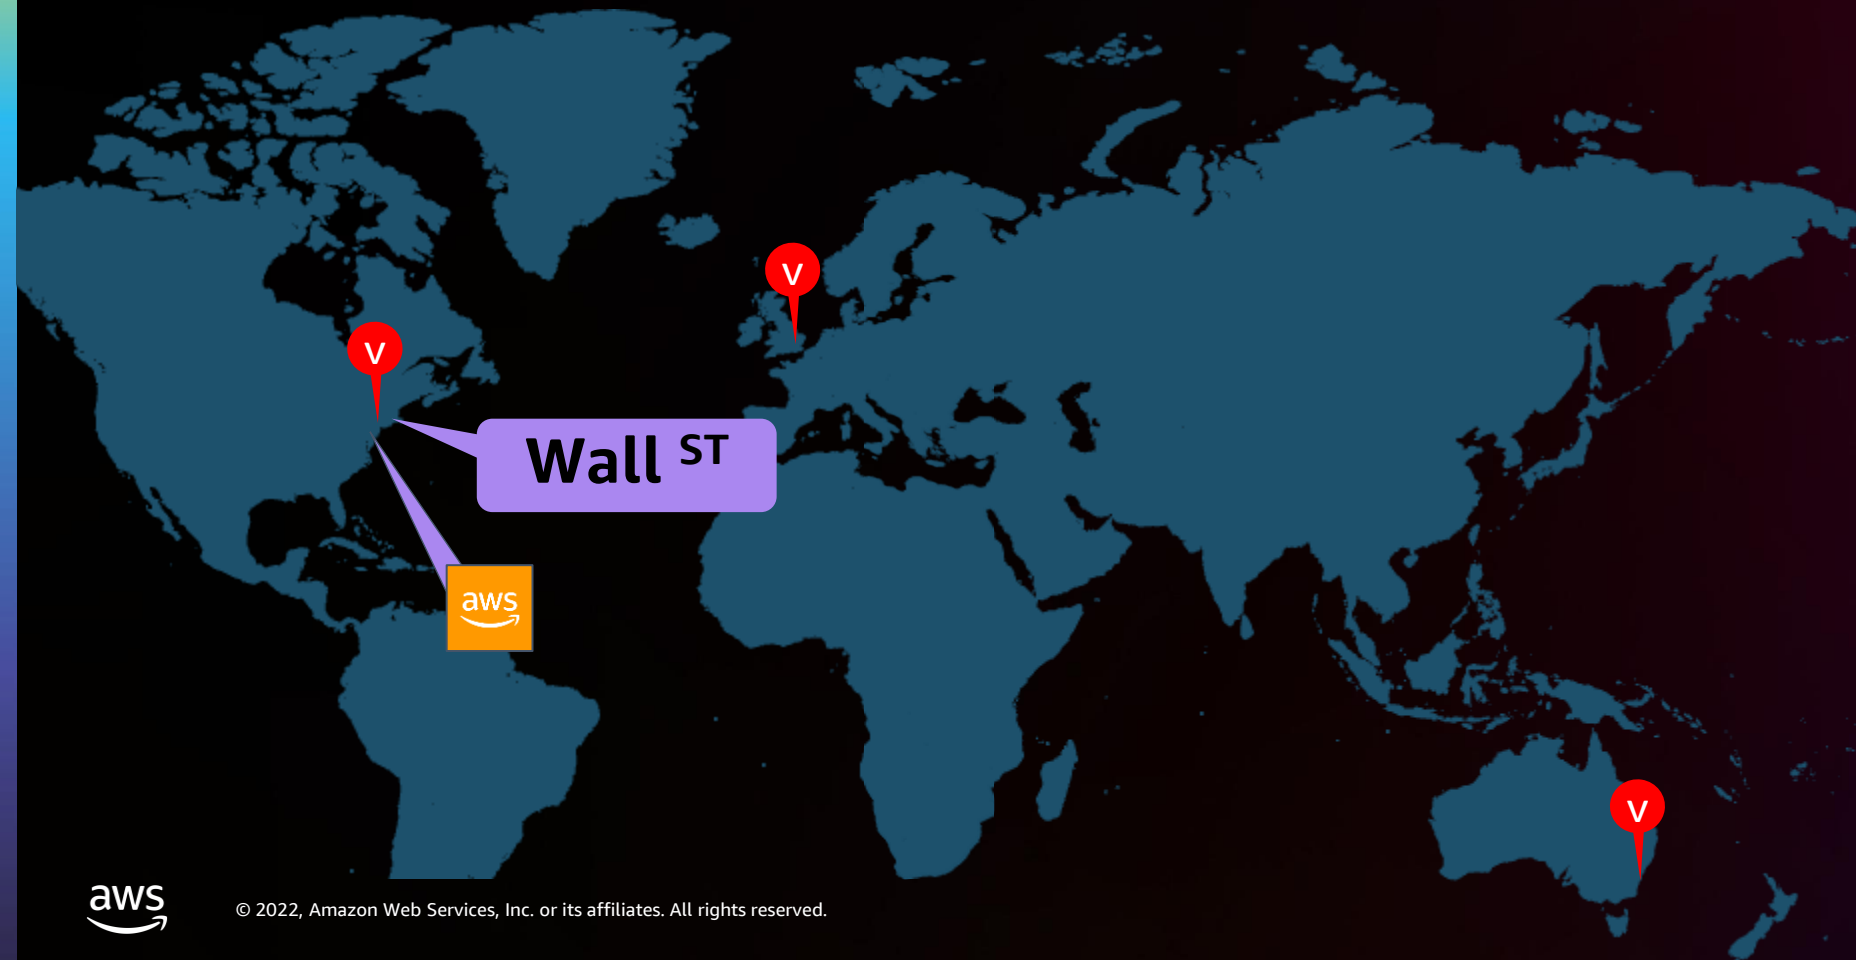
Vanguard: Background

One of the world's largest investment companies, offering a large selection of low-cost mutual funds, ETFs, advice, and related services




Began operations on May 1, 1975, in Valley Forge, PA

Headquartered in Malvern, PA, with operations in Arizona; London, UK; and Sydney, Australia



Understanding the requirements




User experience



High availability

Understanding the data

Centralized – *Hub and spoke*



Some data doesn't change between *start of day* and *end of day*.
Data change is workday specific.

Pass the book/trade – *Follow-the-sun*



Works with both



Amazon DynamoDB

Multi-primary data stores
(e.g., DynamoDB)



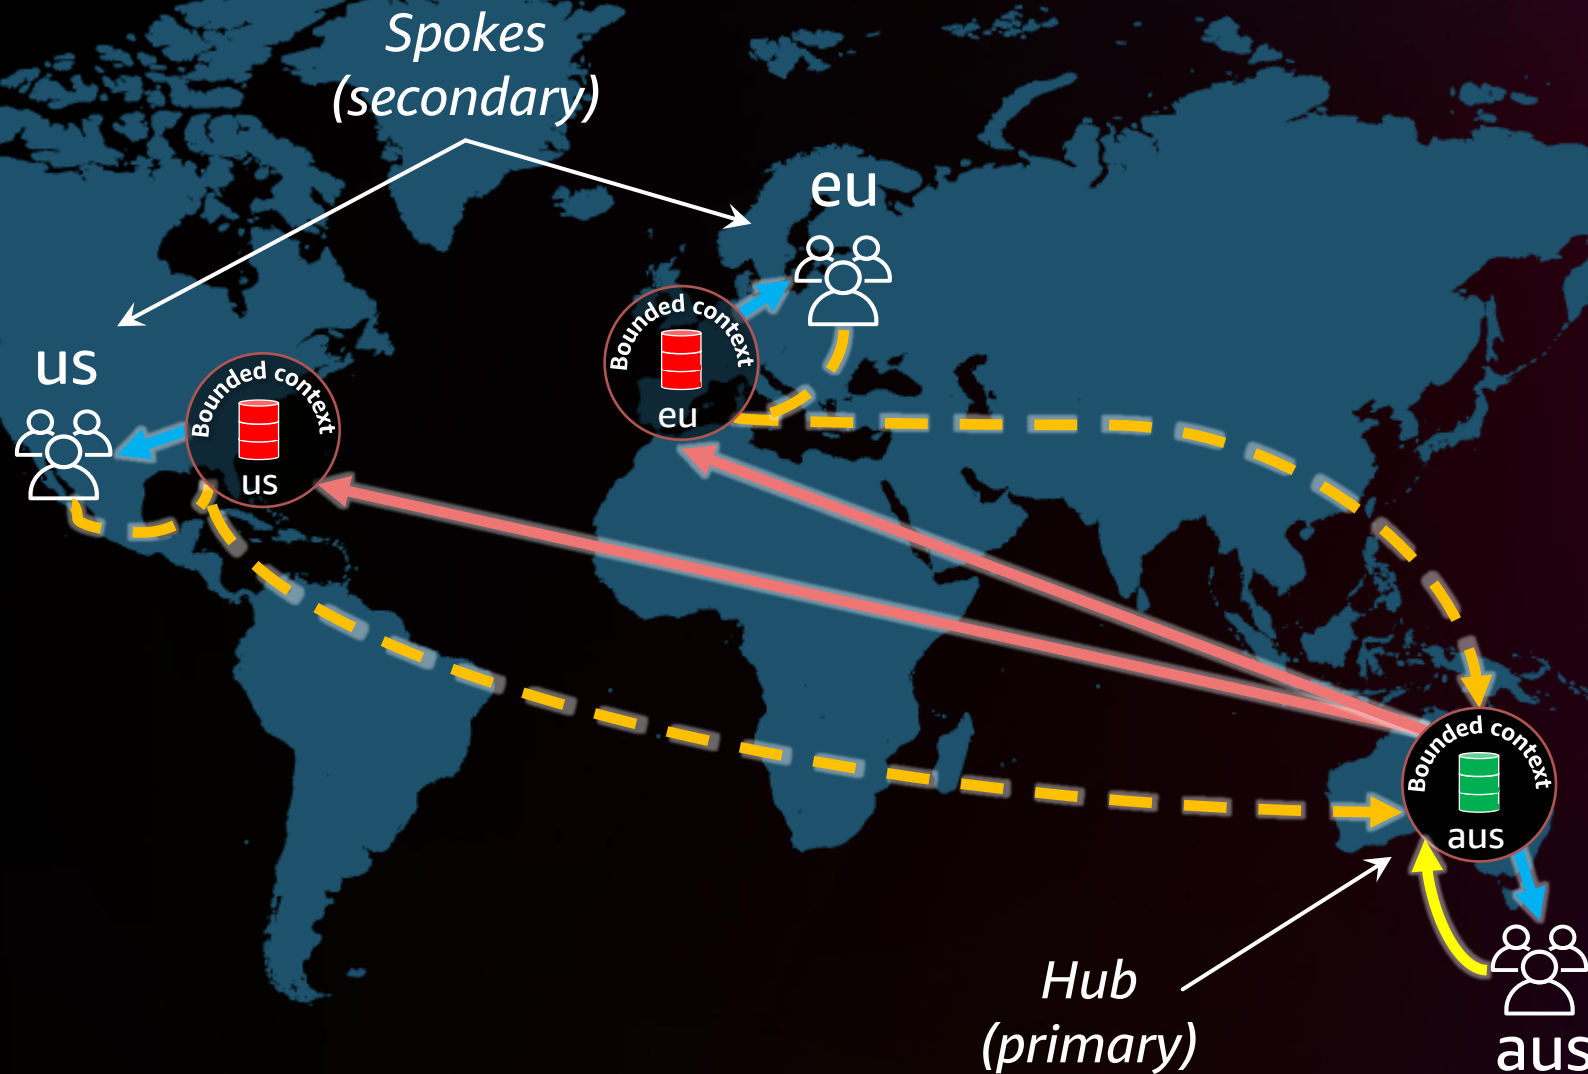
Amazon Aurora

Primary/secondary data
stores (e.g., Aurora
Global Database)

Data access patterns – Hub and spoke



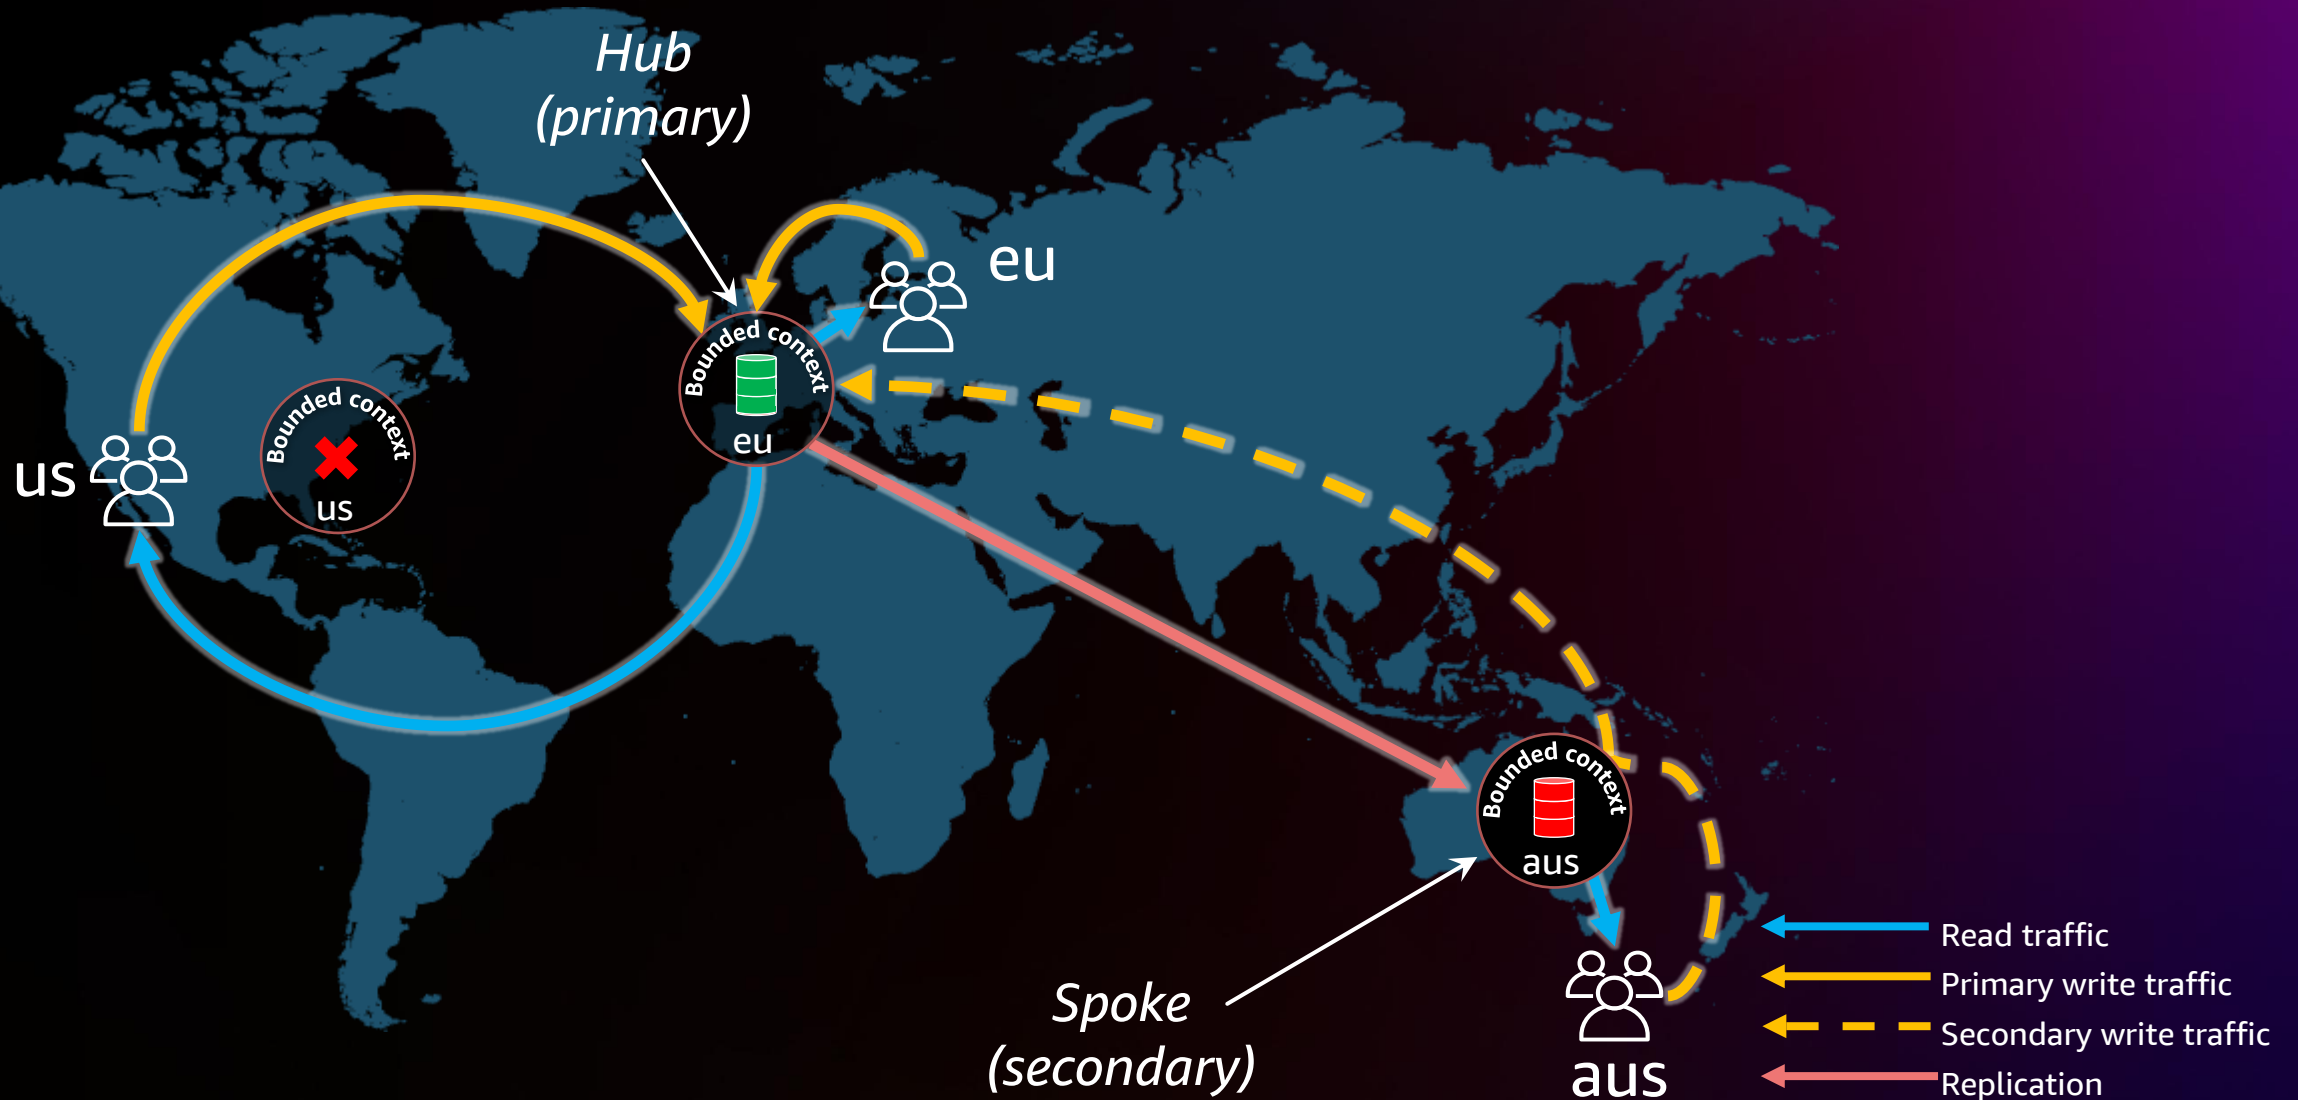
Data access patterns – Switchable hubs



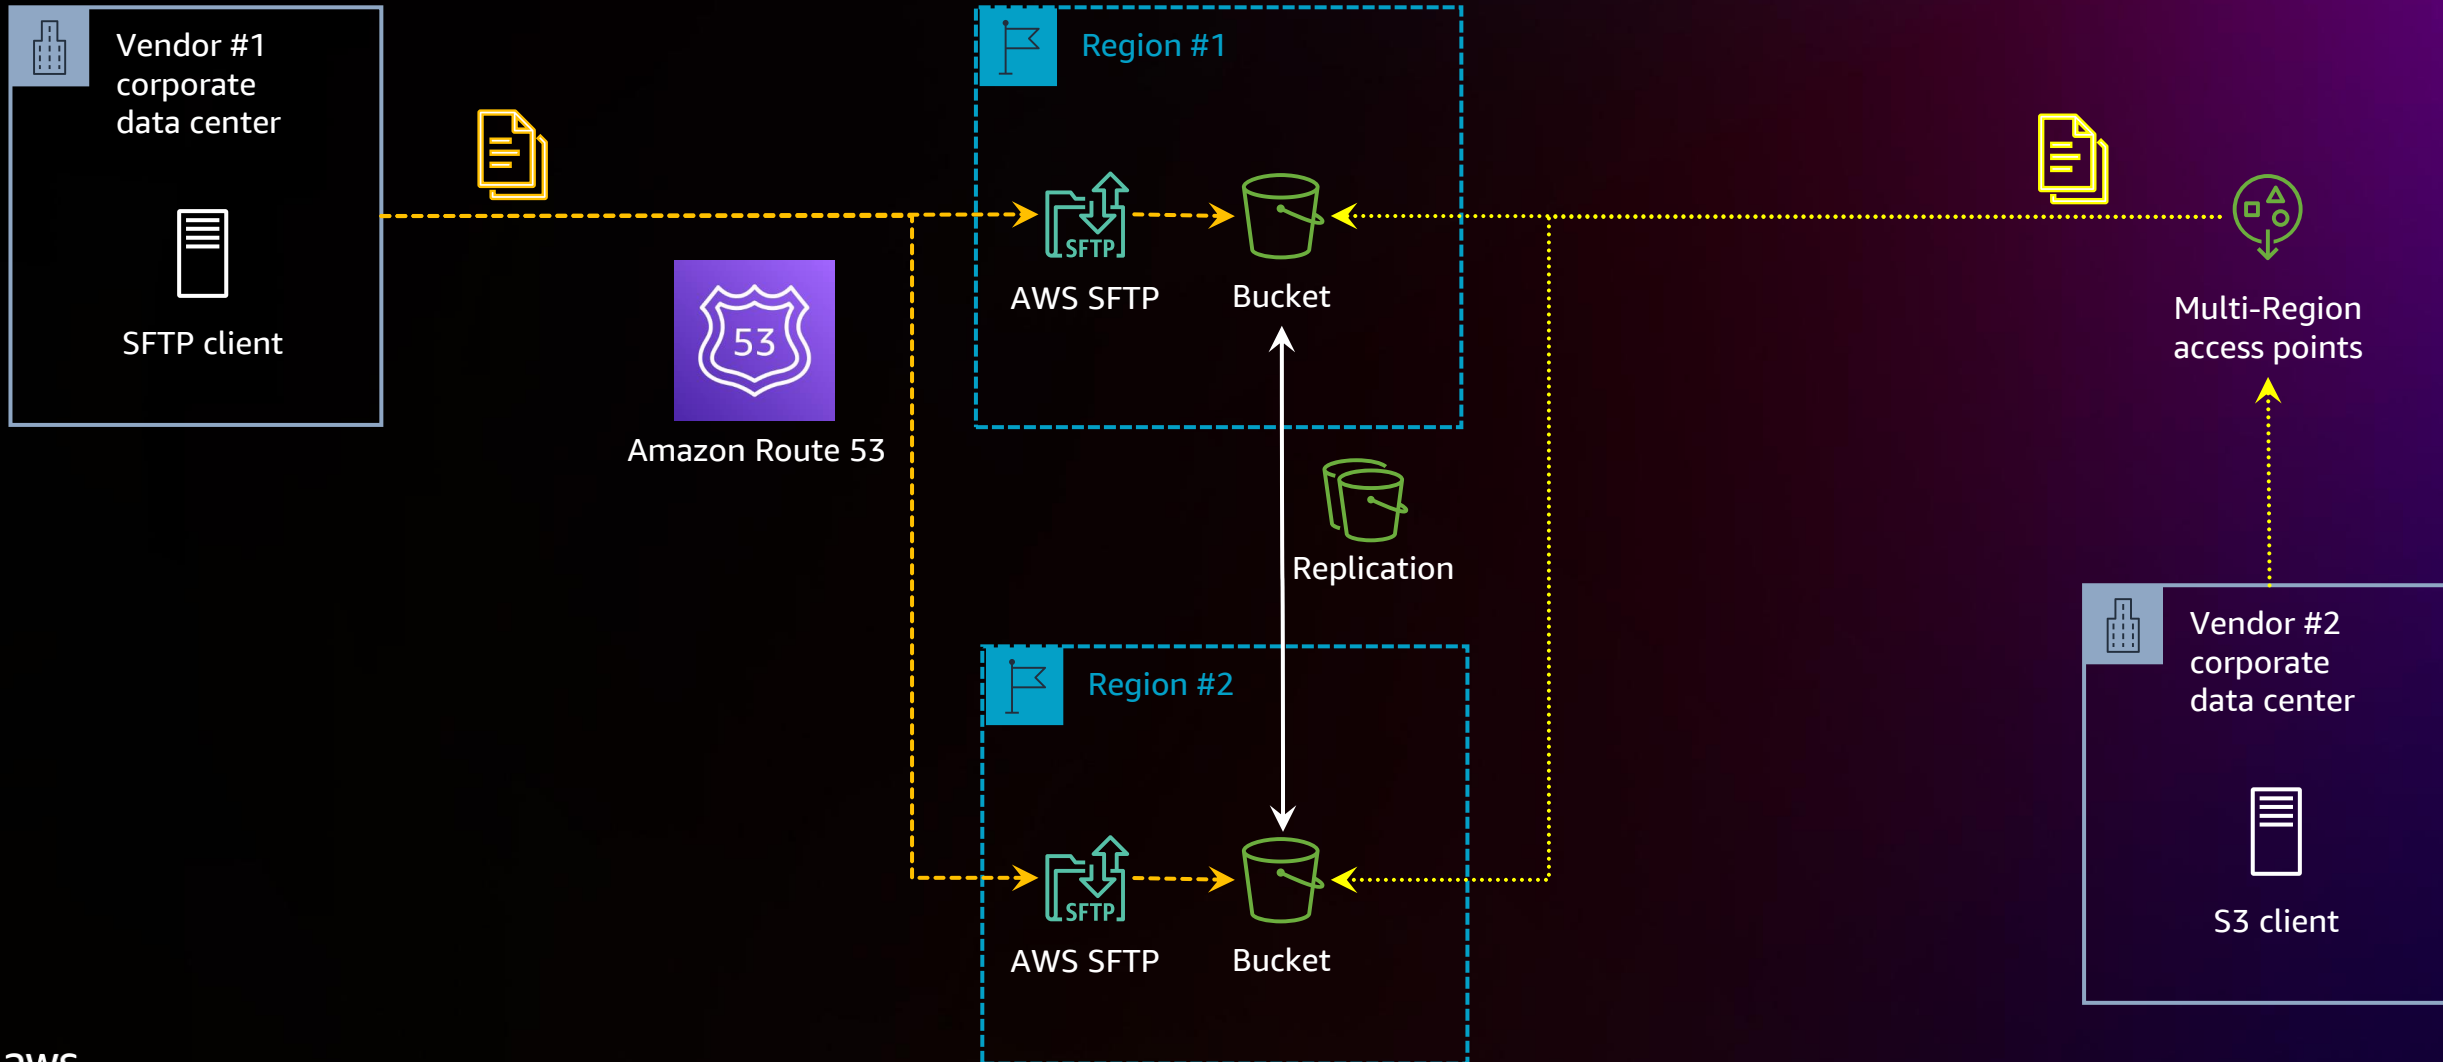
Using Aurora Global Database planned failover to move the hub (primary) during the day, we get *follow-the-sun*

DynamoDB doesn't need planned failover

Data access patterns – Failures



Understanding dependencies – Third party



Understanding dependencies – Failover

GOaST


(Global Orchestration and Status Tool)



Global state



App state



Amazon Aurora failover

I'm alive!
So is...

```
{  
  Primary: us-east-1,  
  Secondary: ap-southeast-2  
}
```

GMRLib – Global Multi-Region library

named op + data

buffer

op + data op + data op + data

retry

write forward

```
defineNamedOp  
  .name("updateIndexData")  
  .doIfPrimary( xxx )  
  .else( yyy );
```

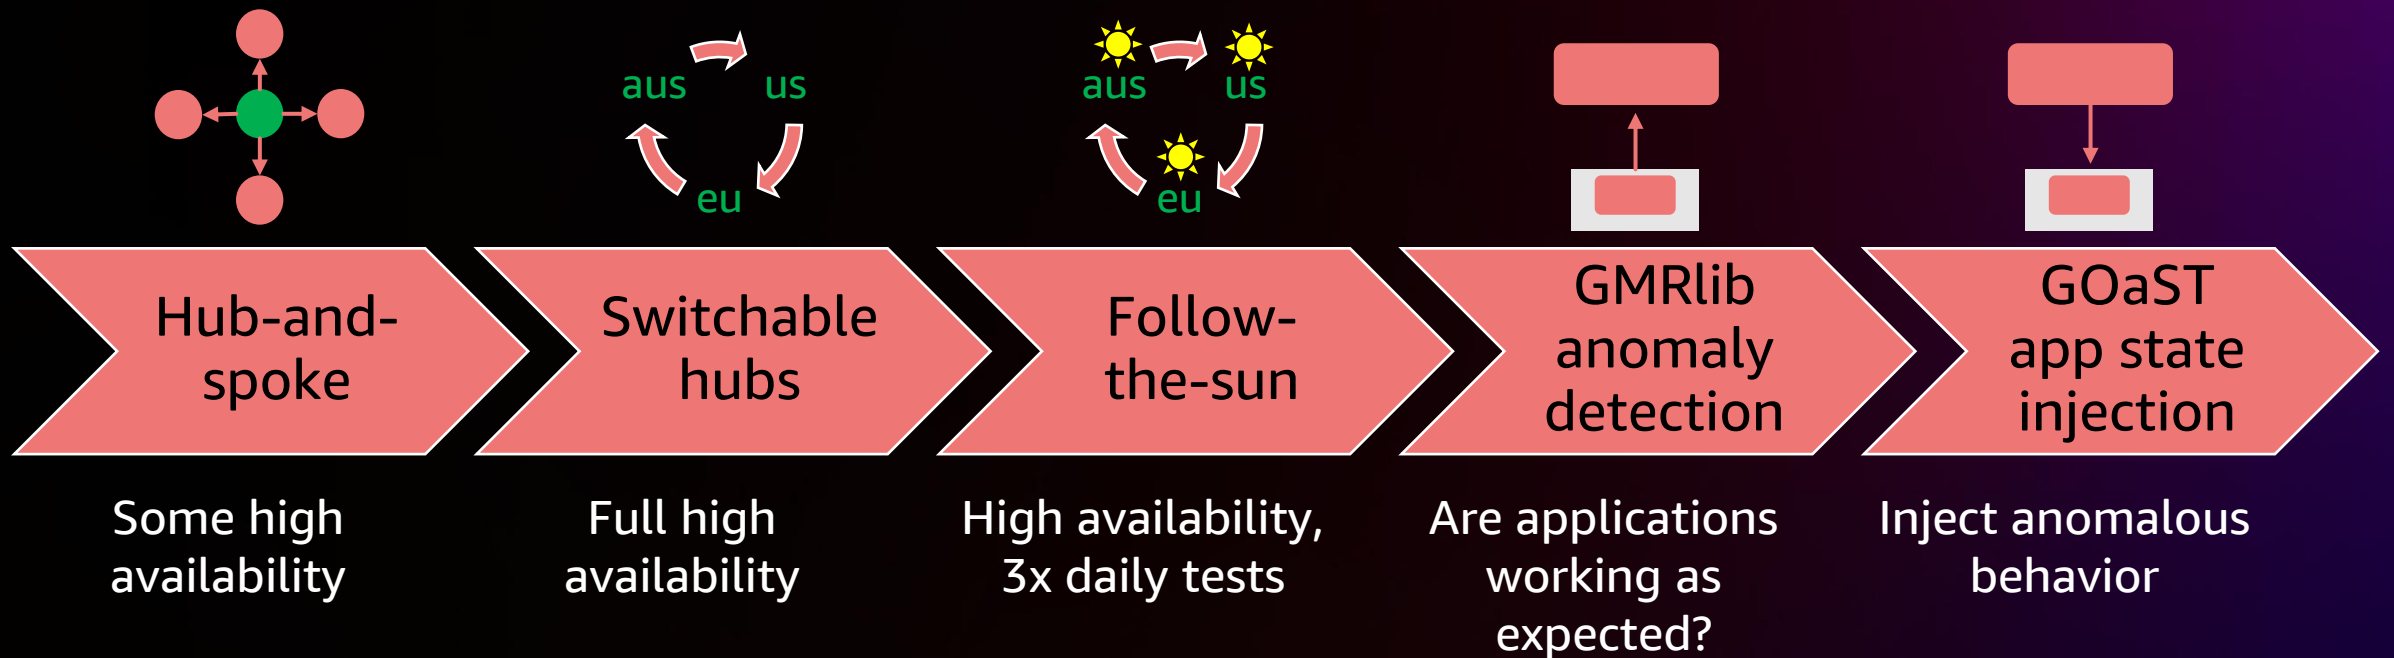
```
runNamedOp(  
  "updateIndexData",  
  data);
```

Microservice


Operational readiness

Remember “an untested DR strategy...is no DR strategy”?

With follow-the-sun, we made failover part of the standard operating procedure



Outcomes



Resiliency



User
experience



Single points
of failure

Costs **\$2x**

but for one workload, a single incident avoided could cover the lifetime cost!

Lessons learned

In summary: global multi-Region is difficult!


- Understand the requirements
- Understand business processes
- Understand the data
- Understand technology limitations
- Understand dependencies
- Set user expectations
- Prepare and test

Expanded multi-Region architecture patterns

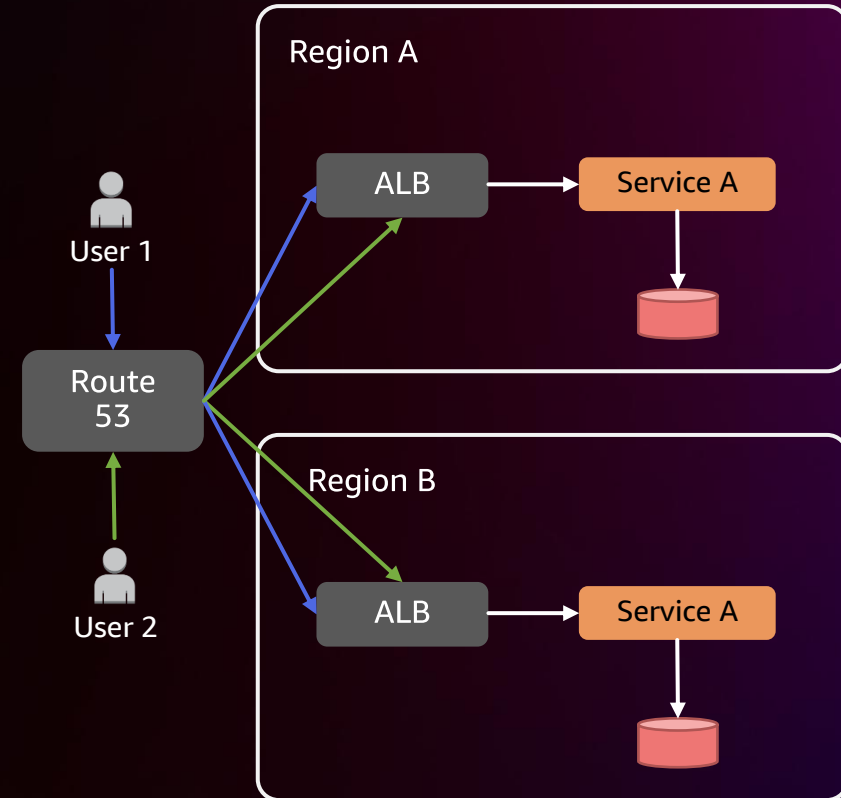


Two macro patterns

Active/passive



Active/active



Pattern #1: Active/passive

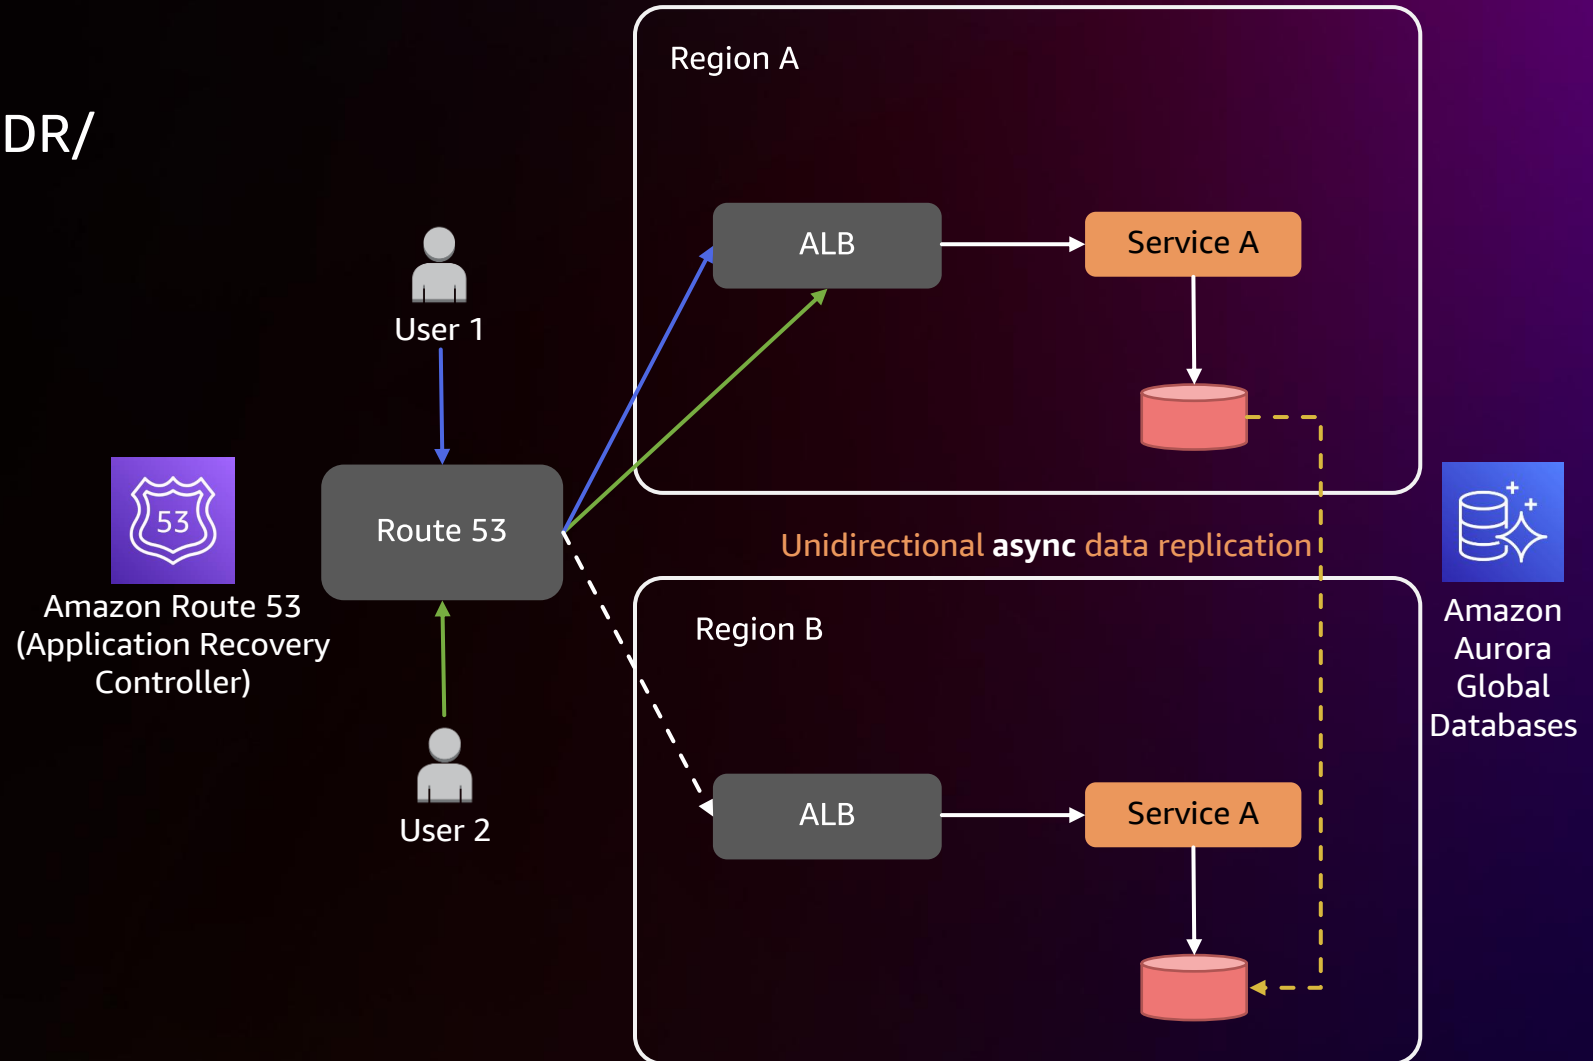
PILOT LIGHT, WARM STANDBY, ACTIVE/HOT STANDBY

- Use cases

- Regional failover/failback for DR/operational continuity
- Regulatory requirements

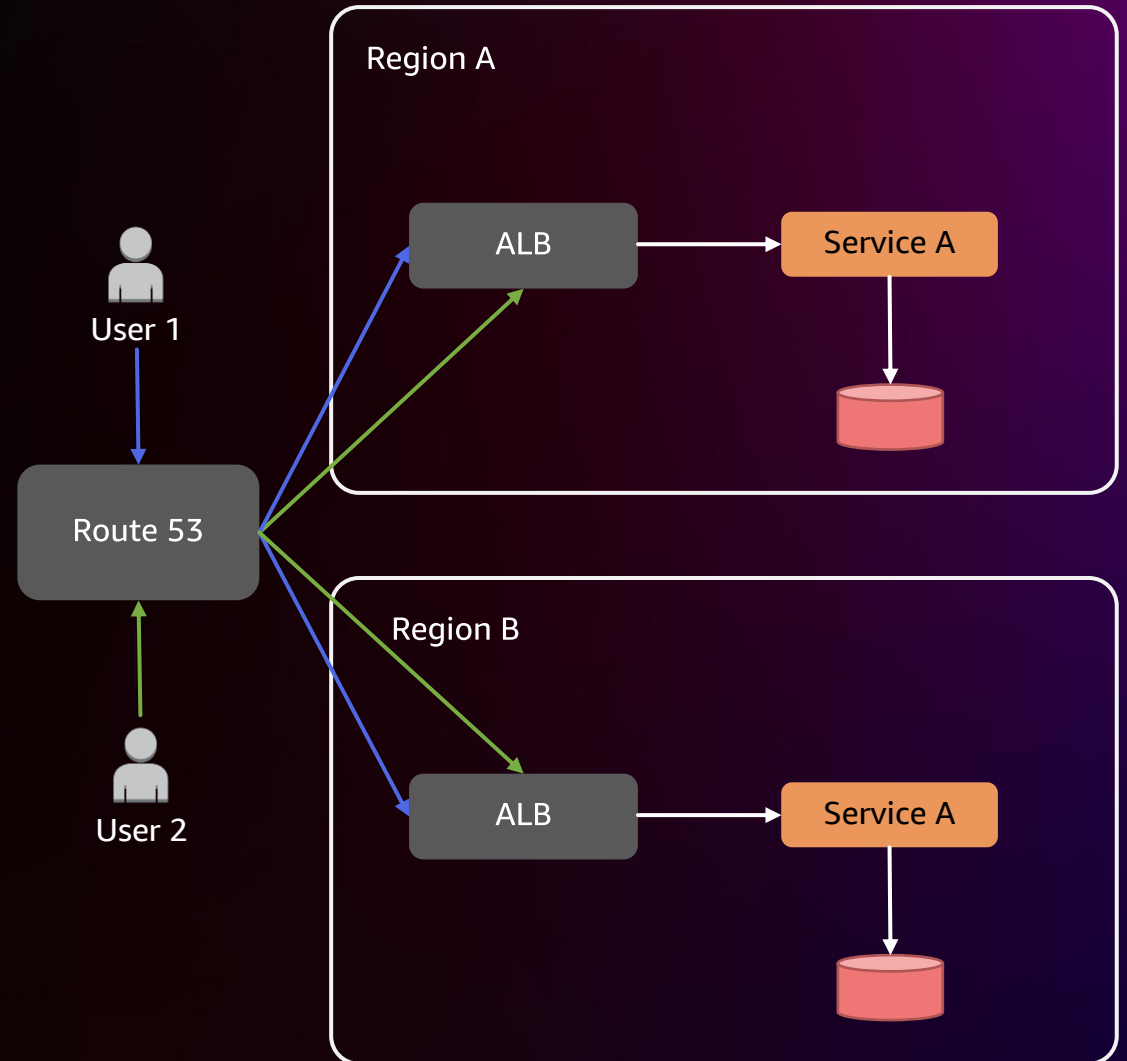
- Key design considerations

- Observability (health checks)
- CI/CD, code and configuration drifts
- Routing strategy
- Failover/failback procedures



Pattern #2: Active/active

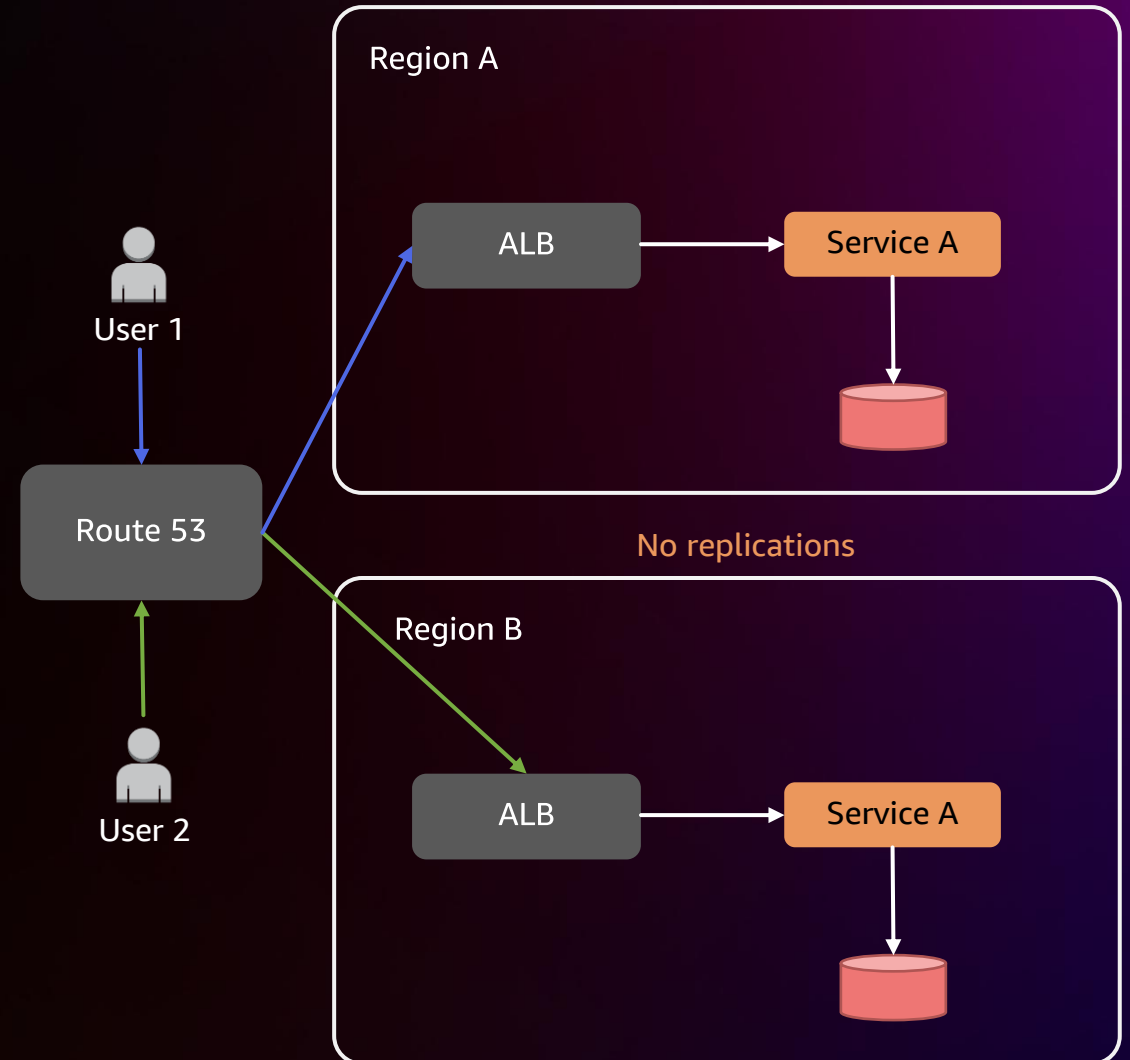
- Use cases
 - Extreme high availability
 - Performance – latency sensitivity
 - Low RTO/RPO



Pattern #2.1: Active/active

REGIONAL SHARDING

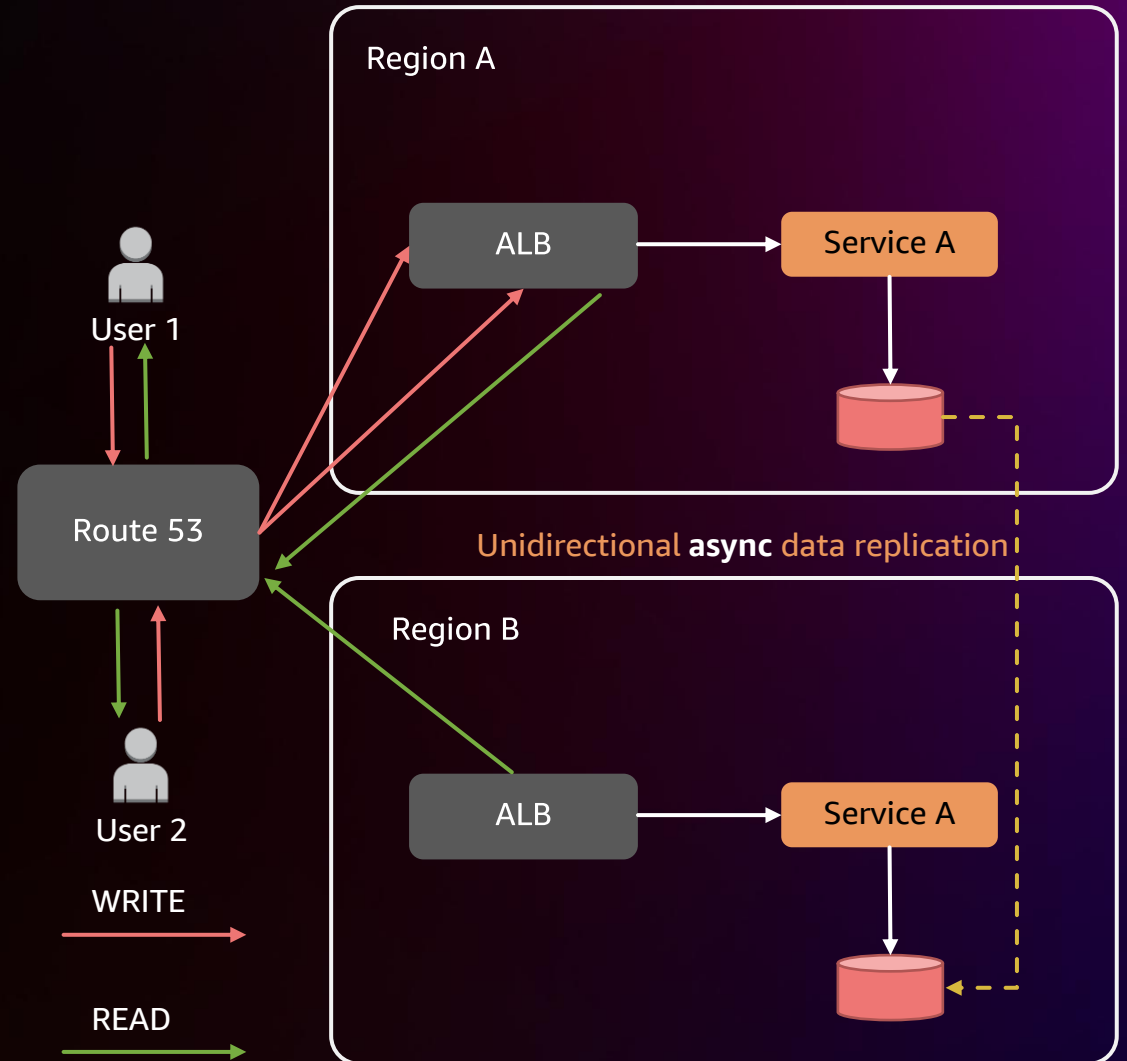
- Use cases
 - Global expansion
 - Data locality, regulatory
 - Performance – latency-sensitive
- Key design considerations
 - Observability (health checks)
 - CI/CD, code and configuration drifts
 - Routing strategy



Pattern #2.2: Active/active

SINGLE WRITER; MULTIPLE READERS

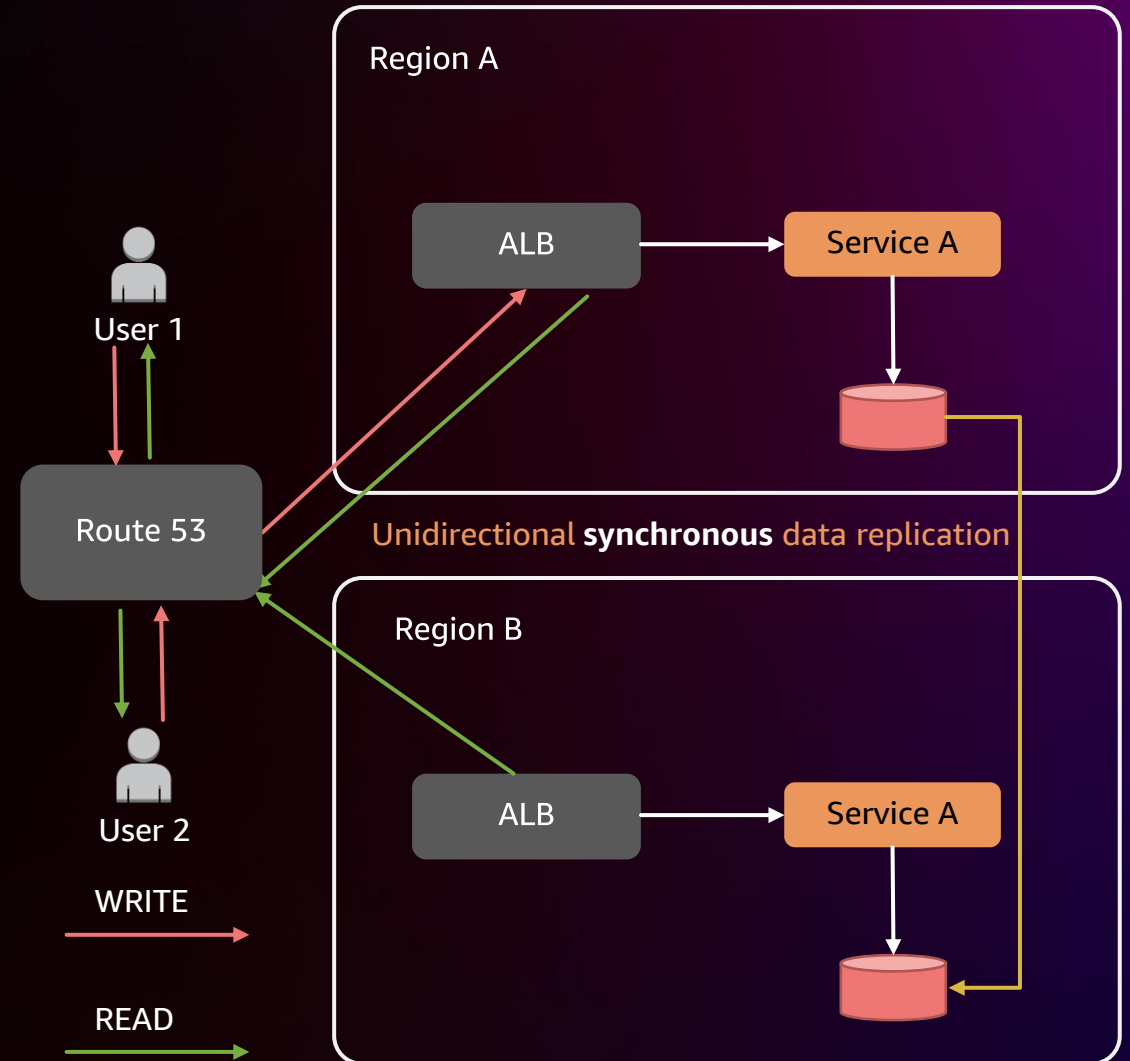
- Use cases
 - High availability
 - Latency sensitivity
 - Eventual consistency – async replication
 - Read-heavy workloads
- Key design considerations
 - Observability (health checks)
 - CI/CD, code and configuration drifts
 - Routing strategy
 - **Failover/failback procedures (writer)**
 - **Read after write consistency**



Pattern #2.3: Active/active

SINGLE WRITER; MULTIPLE READERS (STRONG CONSISTENCY)

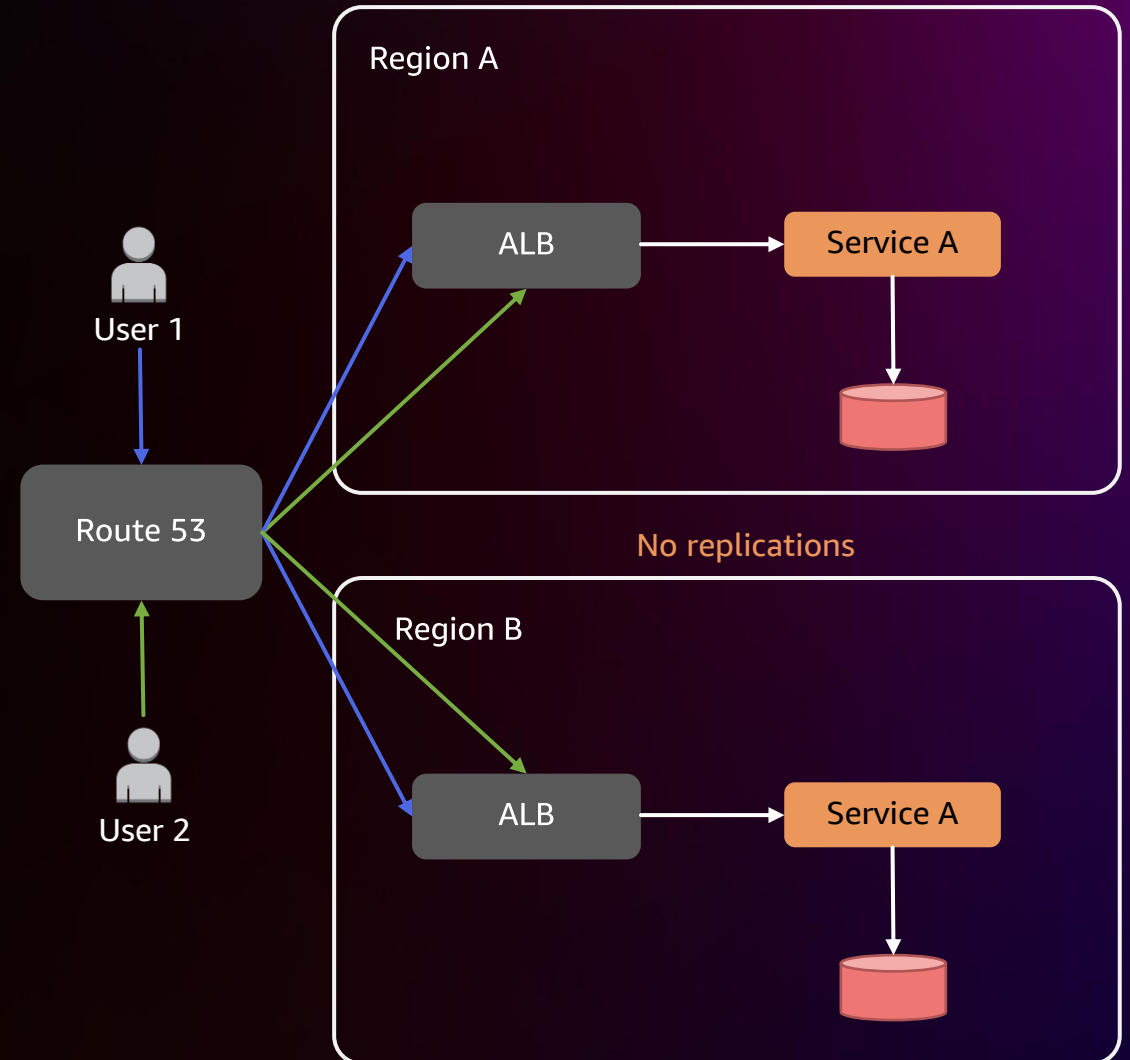
- Use cases
 - Read-heavy workloads
 - Strong read after write consistency
 - Regional failover/failback for DR/operational continuity
 - RPO near 0; regulatory requirements
- Key design considerations
 - Observability (health checks)
 - CI/CD, code and configuration drifts
 - Routing strategy
 - Failover/failback procedures (writer)
 - Increased latency – performance



Pattern #2.4: Active/active

DUAL WRITES

- Use cases
 - Read-heavy workloads
 - Strong read after write consistency
 - Regional failover/failback for DR/operational continuity
 - RPO near 0; regulatory requirements
- Key design considerations
 - Observability (health checks)
 - CI/CD, code and configuration drifts
 - Routing strategy
 - **Idempotency**
 - **Managing data atomicity at the client side (highly complex)**



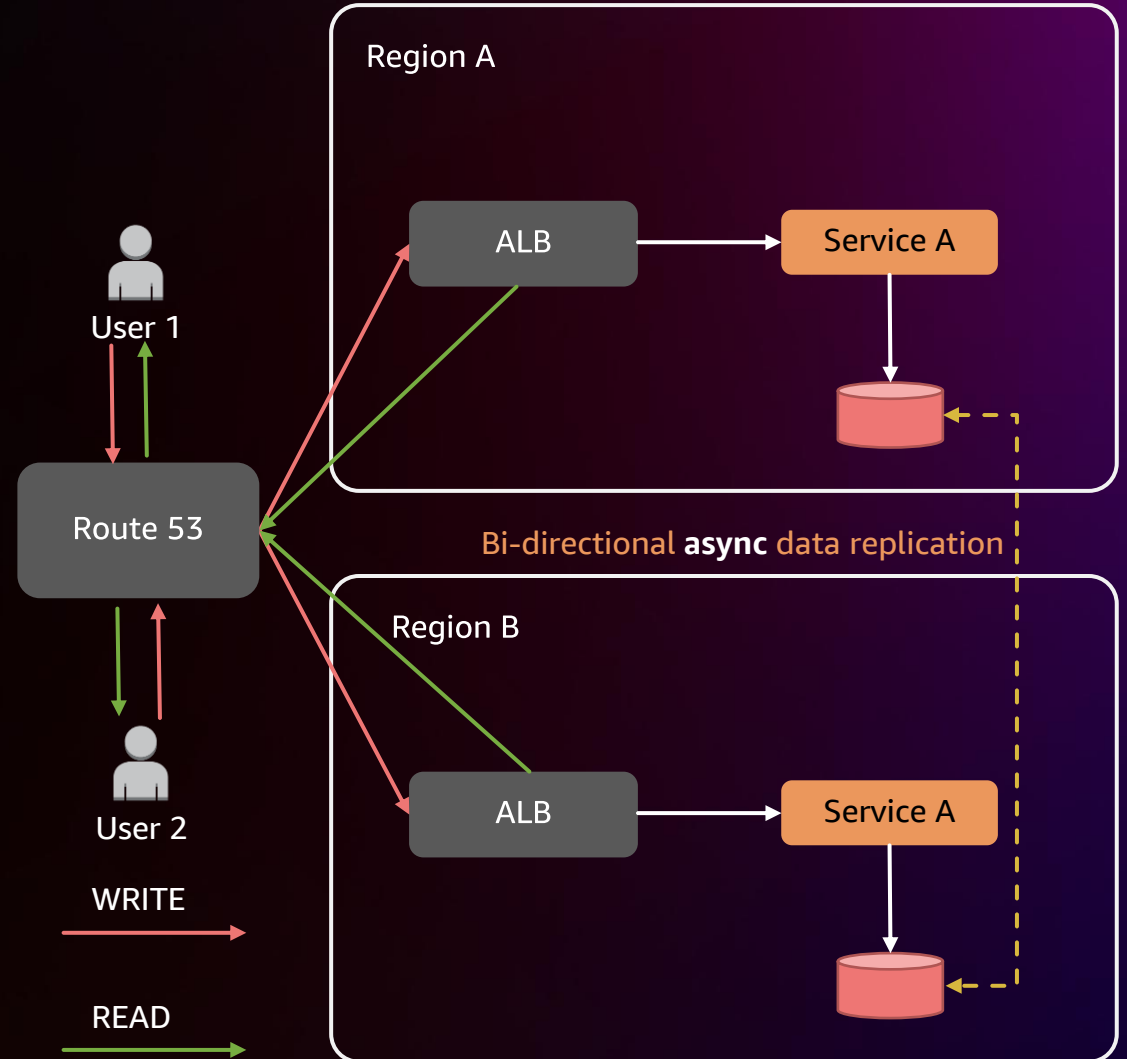
Pattern #2.5: Active/active

MULTIPLE WRITERS; MULTIPLE READERS

- Use cases
 - High availability
 - Latency sensitivity
 - Eventual consistency – async replication
 - Read/write-heavy workloads – scalability
- Key design considerations
 - Observability (health checks)
 - CI/CD, code and configuration drifts
 - Routing strategy
 - **Conflict resolution**
 - **Read after write consistency**



Amazon DynamoDB




Routing patterns



Amazon Route 53
(Application Recovery Controller)



Amazon CloudFront
(Lambda@Edge)



AWS Global Accelerator

Key takeaways

- Do not start with multi-Region as a starting point for resilience
- Multi-Region is not just a technology decision – it's a business decision
- Build a business case first
- Don't forget about dependencies – internal or external
- Use Regions for what they are designed for – independent failure domains
- Build deep observability – do you know when to press the big red button?
- Multi-Region can be complex – choose the right patterns for right reasons

Thank you!

John Formento, Jr.



<https://www.linkedin.com/in/johnformentojr/>

Neeraj Kumar



<https://www.linkedin.com/in/neerajkumar78/>

Barry P. Sheward



<https://www.linkedin.com/in/barrysheward/>



Please complete the session survey in the **mobile app**

