

AWS re:Invent

NOV. 28 – DEC. 2, 2022 | LAS VEGAS, NV

AIM205

Fidelity elevates developer operations with ML-powered insights

Vijay Ranganathan

Vice President, SRE
Fidelity Investments

Ashok Kanjamala

Principal Product Manager, Amazon DevOps Guru
AWS



© 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Agenda

About Fidelity

Observability Maturity Model, DoD (Definition of Done)

What is Amazon DevOps Guru?

Fidelity's DevOps Guru engagement approach

Issues and remediations, integrations, self-healing

Key takeaways



1946



Fidelity Management & Research Company founded



1965

Fidelity purchases first computer



1995

Fidelity becomes an internet pioneer with first mutual fund to create a home page



2016

Fidelity's first product application deployed to the cloud



2019

Fidelity launches multi-cloud hybrid strategy



Today

5,700 + applications on public cloud




Massive scale



Shared responsibility

Each individual team to look after their own application's health




Applications

Hundreds of workloads moving to the cloud across the enterprise
Applications are tiered



Accounts

Hundreds of accounts spread across multiple regions managed by multiple distinct internal business units



Organizational units

Variety of internal business units are migrating to the cloud, all at different levels of maturity in their cloud journey



Developers

17,000 technologists spread across the globe

Application availability (focus areas)

Senior tech executive

I want a consistent way to capture anomalous behavior so my engineering teams can manage and correlate metrics to quickly resolve issues.

I need to determine where we are doing a good job and areas where we can improve.

BU owners

I want my availability expectations defined and documented so that I can measure and report. I want to relay expectations to other dependencies.

I want to continuously measure the performance of my business service and user experience so that I can detect before customers and alert on our error budget burn-down.

I want to enable my teams to focus on the roadmap items.

Operations leader

I need visibility of the service adoption and the impact on the production support data (incidents, MTTR).

I want to enable teams to significantly reduce operational downtime.

I want to enable teams to track and meet their availability targets.

Software engineering team and architect

I want to leverage the ML-powered insights to help identify, correlate, and remediate operational issues.

I want to utilize the integration with ServiceNow to quickly end efficiently deliver issue/recommendations to the right people at the right time.

I want to use the proactive insights to identify problems around resource exhaustion in advance.

SRE and platform engineering

I want a common set of anomalous metrics and recommendations for each AWS service, so I can guide application teams appropriately.

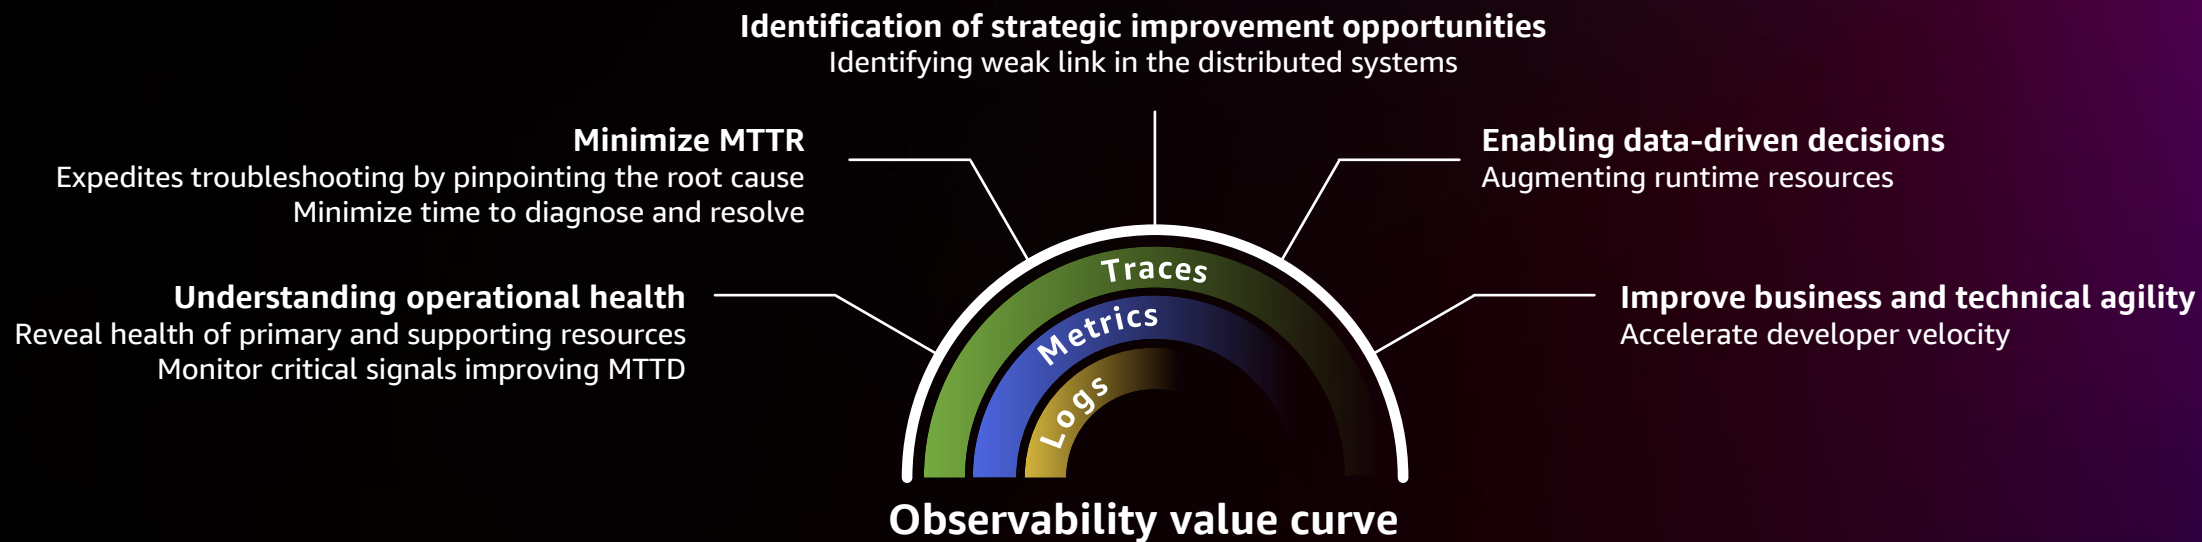
I want to have well-documented self-service onboarding steps, so I can help teams with their adoption.

I want to measure the value-add with respect to MTTR.

Enterprise-wide data transparency (same data with different filter/lens)



Why our observability strategy matters



The three I's (eyes) of observability

Instrumentation

- Common patterns to capture the right signal in the right way
- Emphasize metrics-based SLIs and SLOs
- Leverage auto-instrumental tracing
- Enforce standards for log messages

Inspection

- Use each health indicator (metrics, traces, logs)
- Provide common platform for interacting and exploring observability signals

Insights

- Deliver actionable alerting capabilities
- Provide trend detection to warn of issues before they impact customers
- Consistently applied metadata to afford quickly connecting various signal categories

Observability Maturity Model + DoD

Transactional tracing across apps/bus
Predictive AIOps

Level 4
Self-healing and optimize

Cross-app observability and tracing
Self-optimizing apps and infrastructure

RUM
Region health checks
Customer performance

Level 3
Customer availability

Real customer experience
Error budgets prioritizing backlog

Synthetic monitoring
Ops or SLI/SLO dashboard
Distributed tracing for cloud platforms

Level 2
Common and predictive

Common SLI/SLO
Application dependency metrics

Alerting
Correlated logging

Level 1
Basic (responsive)

Runtime measurements
Faster application recovery

Disk I/O, CPU, memory, ping, network
I/O, logging, process
Logging, process

Level 0
Monitoring

Are my systems running?
Have I exceeded system capacity?

Amazon DevOps Guru



Amazon DevOps Guru

ML-POWERED CLOUD OPERATIONS SERVICE TO IMPROVE APPLICATION AVAILABILITY



DevOps Guru is an ML-powered service that makes it easy for developers and operators to automatically detect issues to improve application availability and reduce expensive downtime – no machine learning experience required

Get started with Amazon DevOps Guru

Select your application
by tags, CFN stacks, or
entire account



Tags

OR



AWS CloudFormation
stack

OR



Account

Select
coverage

DevOps Guru will start
analyzing metrics/logs
to generate insights



Generate
insights

Integrate with
OpsCenter, Amazon SNS,
Amazon EventBridge, and
third-party platforms



Amazon SNS



Amazon
EventBridge



Integrate with
your workflow

Engagement approach

Service type engagement

Targeted 4 AWS services
(Lambda, ELB, Amazon API Gateway, Amazon SQS)

Direct engagement with 3 business units

Prove the value-add

Production incident engagement

Target high-priority production incidents

Direct engagement with 2 business units

Prioritized 2 AWS services (Amazon EKS, DynamoDB)

Prove the value-add

Issues and recommendations (AWS services)

Resource	Issue	Auto healing job
Amazon SQS	Dead letter queue for SQS queue not configured	Set up an SQS dead letter queue. Without an SQS dead letter queue, failure to process the message will result in the message being continually retried by Lambda until the message expires from the queue and is lost. The retries use up your resources and add to your cost.
Amazon Kinesis	RetryLimit for Lambda Kinesis consumer has not been configured	Set MaximumRetryAttempts. Without limits on retries, Lambda will continually retry a failed batch until the messages expire from the stream. When this happens, the Lambda function is effectively blocked.
Amazon EKS	Amazon EKS deployment is having trouble scaling	Set up Auto Scaling for your pods to prevent this situation.
DynamoDB	Failure destination for Lambda DynamoDB has not been configured	Set up dead letter queue to prevent data loss in certain situations.
DynamoDB	DynamoDB table point-in-time recovery not enabled	Enable point-in-time recovery to prevent data loss.
Lambda	Lambda concurrent executions reaching account limit	Increase the Lambda function concurrency to handle your current load.
Lambda	Lambda provisioned concurrency usage lower than expected	Reduce the Lambda function concurrency to save cost.

Integrations with ServiceNow, self-healing



DevOps
Guru



Amazon
SNS



AWS
Lambda



SNOW
API



servicenow



DevOps
Guru



Amazon
SNS



AWS
Lambda



Retrieve the current RCU for the table
Increase the RCU by a certain number
Enable Auto Scaling
Update scaling policy

Self-healing automation platform

Key takeaways

- Self-healing is the key value-add using proactive insights
- Started with 6 AWS services, plans to add 3 more (Amazon RDS, Amazon ECS, AWS Step Functions)
- Share learnings (corrections), share techniques, common issues, common patterns across business units
- Focused approach to resolving the critical signals consistently (working with AWS to improve)

Thank you!



Please complete the session survey in the **mobile app**

