

AWS re:Invent

NOV. 28 – DEC. 2, 2022 | LAS VEGAS, NV

CON207

Your first container in the cloud

Omar Lari

Principal Product Manager
Amazon Web Service

Miller Fanning

Principal Software Engineer
Cox Automotive



© 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Agenda

One, two, go! – Two easy ways to get started with AWS App Runner

What is App Runner and why did we build it?

Operating App Runner in production






Bryan Cantrill

@bcantrill

 Follow

Question I've been posing at [#KubeCon](#): are we near Peak Confusion in the container space?
Consensus: no -- confusion still accelerating!

AWS App Runner



**Your web
applications**

MANAGED BY AWS

No EC2 instances, clusters, load balancers that you need to provision, scale, or manage.


ELASTIC

Scale up and down seamlessly. Pay only for what you use.

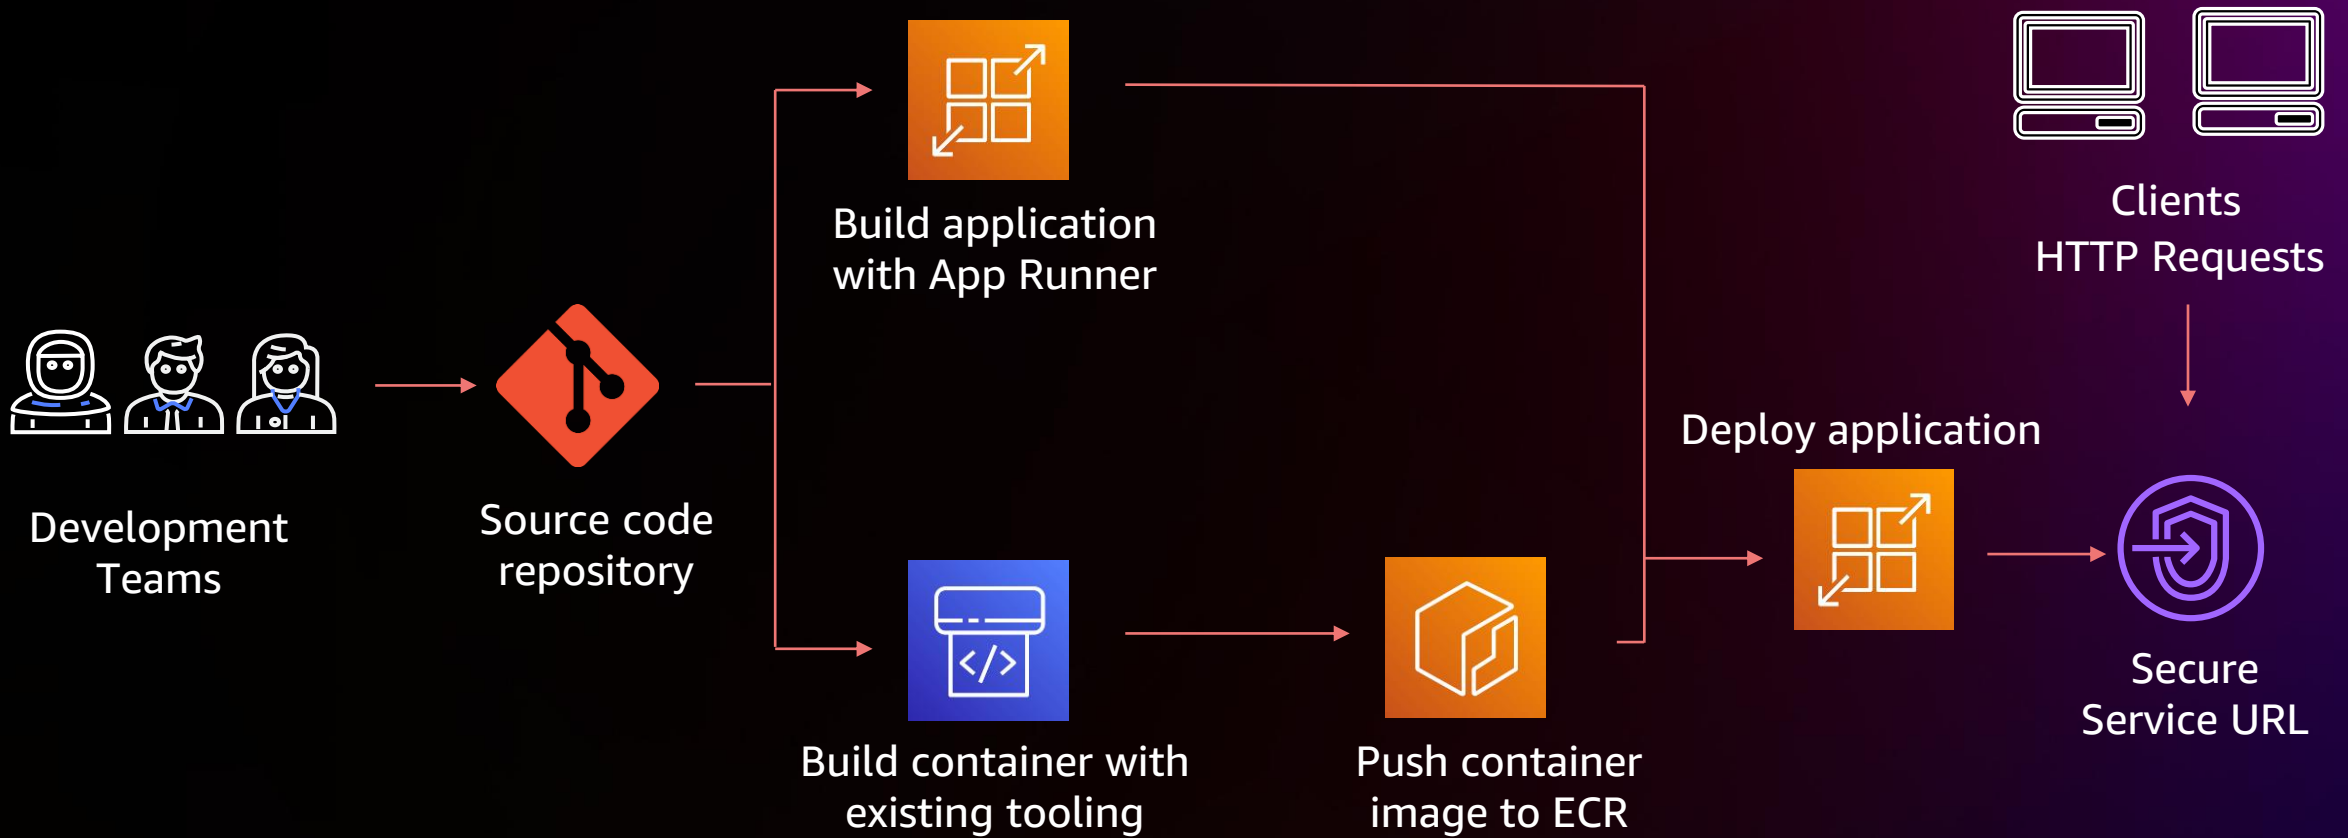
INTEGRATED

With the AWS portfolio of services: VPC Networking, Amazon CloudWatch, AWS Certificate Manager (ACM), and more.

Idea, Code, Launch



Deploy from source or container image



Pay for compute that scales with requests

▼ **Auto scaling** [Info](#)
Configure automatic scaling behavior.

Auto scaling configuration


☒ **Default configuration**
Use the App Runner default configuration.

☐ **Custom configuration**
Use your own configuration.


Concurrency
80 requests

Minimum size
1 instance(s)

Maximum size
2 instances



Pricing dimensions



Compute pricing

\$0.064/vCPU-hour

\$0.007/GB-hour



Automatic deployments

\$1/application, per month



Build fee

\$0.005/build minute

Pricing scenarios

Development or test app

Built by individual developers, these apps serve test or development traffic and use manual deployments. Developers often pause these services to manage costs.

\$5/month

Lightweight, latency-sensitive API

Lightweight commercial API's or web services see low request volume and do not undergo any peak hour activity.


\$25/month

High-volume production app

Production app powers your most latency-sensitive business logic. They see several peak traffic hours daily.


\$102/month

Cox AUTOMOTIVE



COX AUTOMOTIVE
MOBILITY

DEALER.COM

Dealertrack 

 F&I Express

FYUSION



Kelley Blue Book



Manheim

NEXTGEAR
CAPITAL



**SPIERS
NEW
TECHNOLOGIES**
COX AUTOMOTIVE / MOBILITY

vAuto
LIVE MARKET VIEW


VinSolutions

 xtime




© 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.


Model delivery At Cox Automotive



Our goal is to do for data science what DevOps did for software development




One of the problems that we had to solve was how to run models as APIs at scale



Can we build a managed service-driven platform to deliver a wide range of models?


Containers in the cloud



AWS App Runner



Amazon ECS

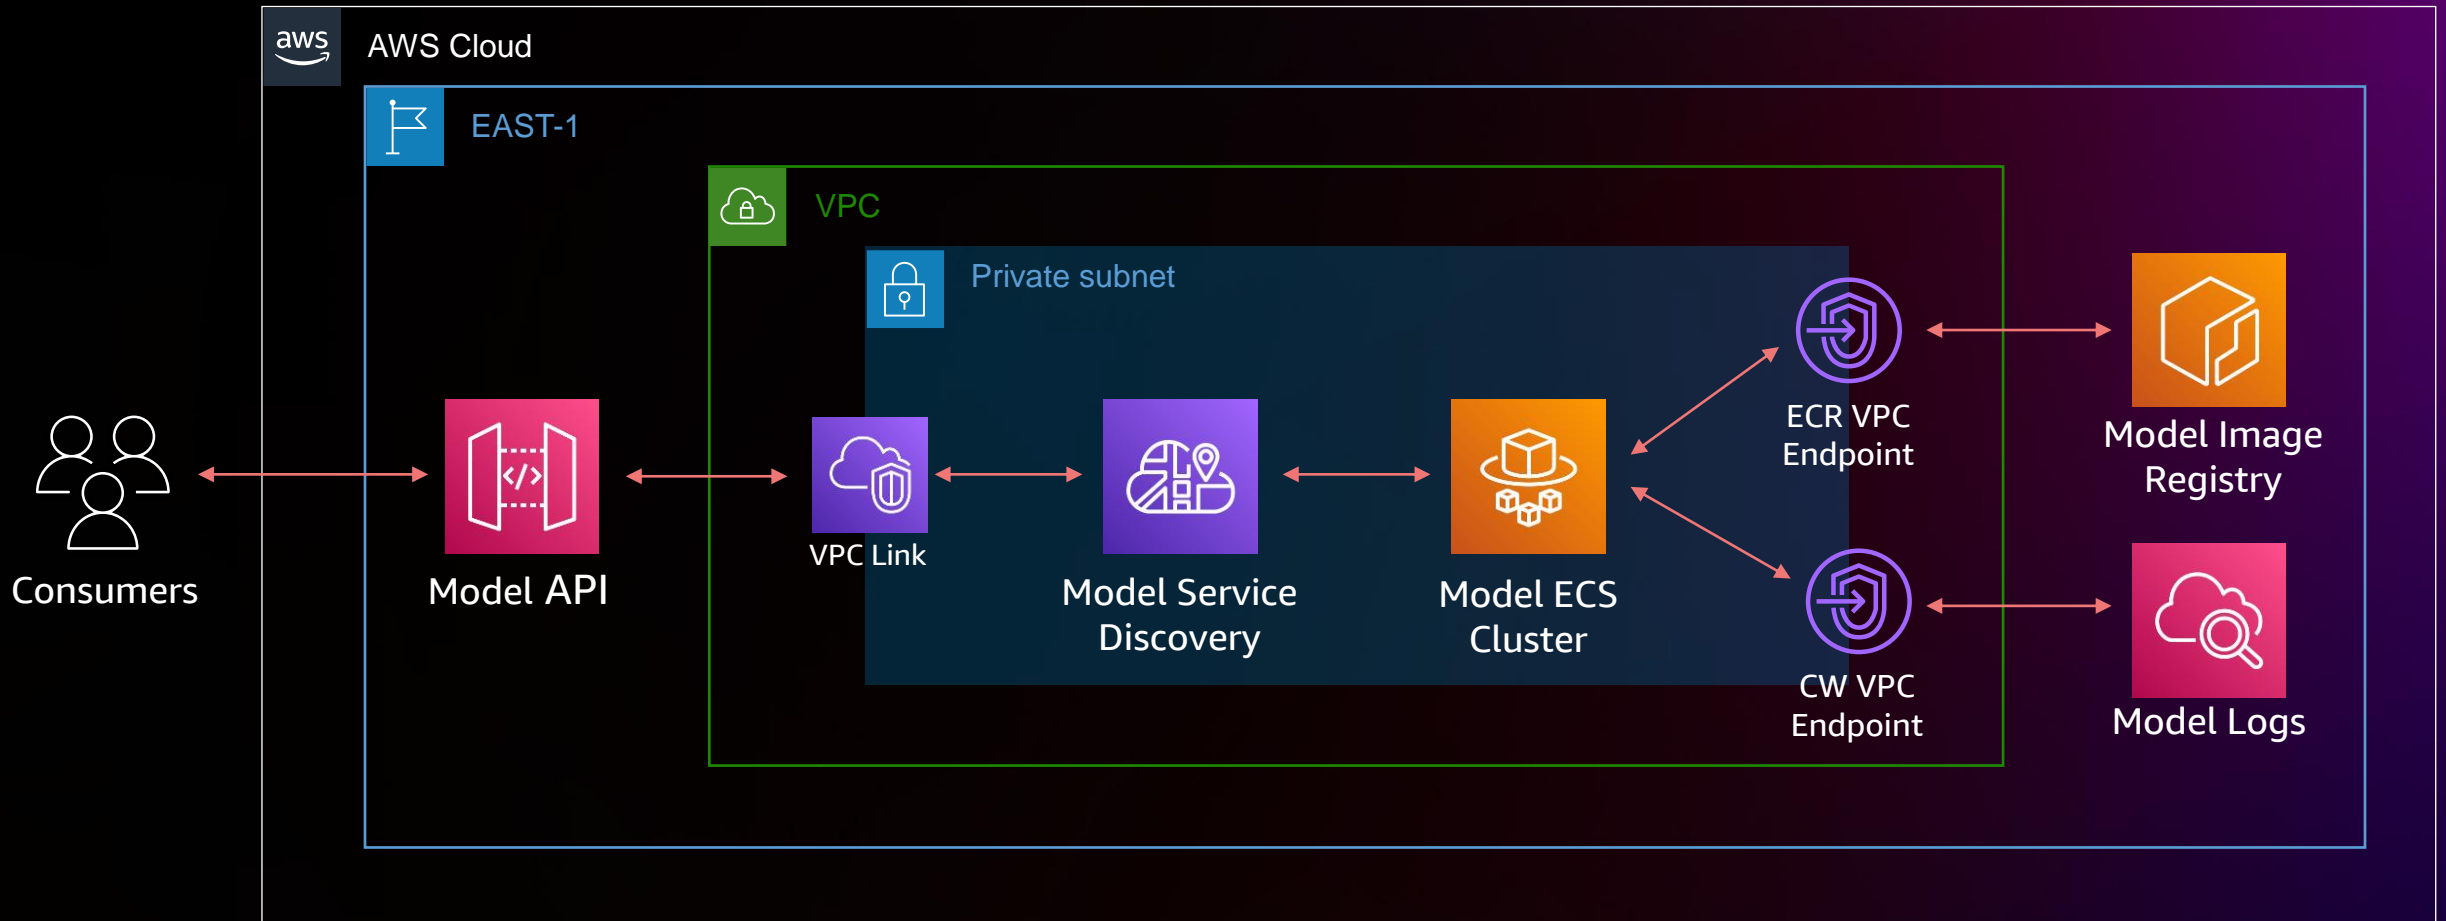


AWS Fargate

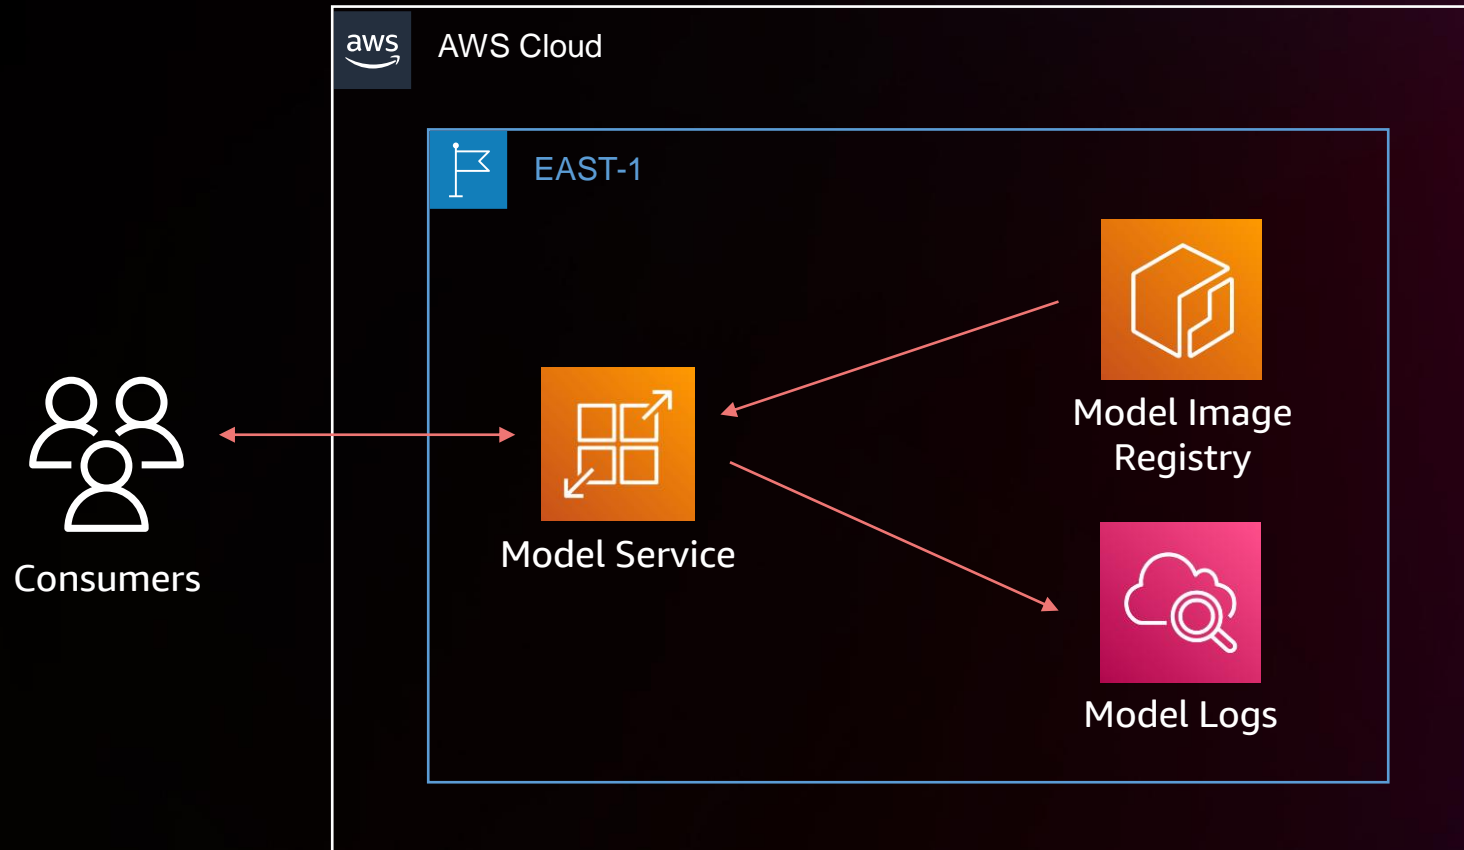


Amazon EKS

Known ECS pattern



Fitting into the App Runner pattern



Simplified operations

Over-provisioning

- Solve for burst traffic
- Previous model versions scaled
- Lower environments mirror production

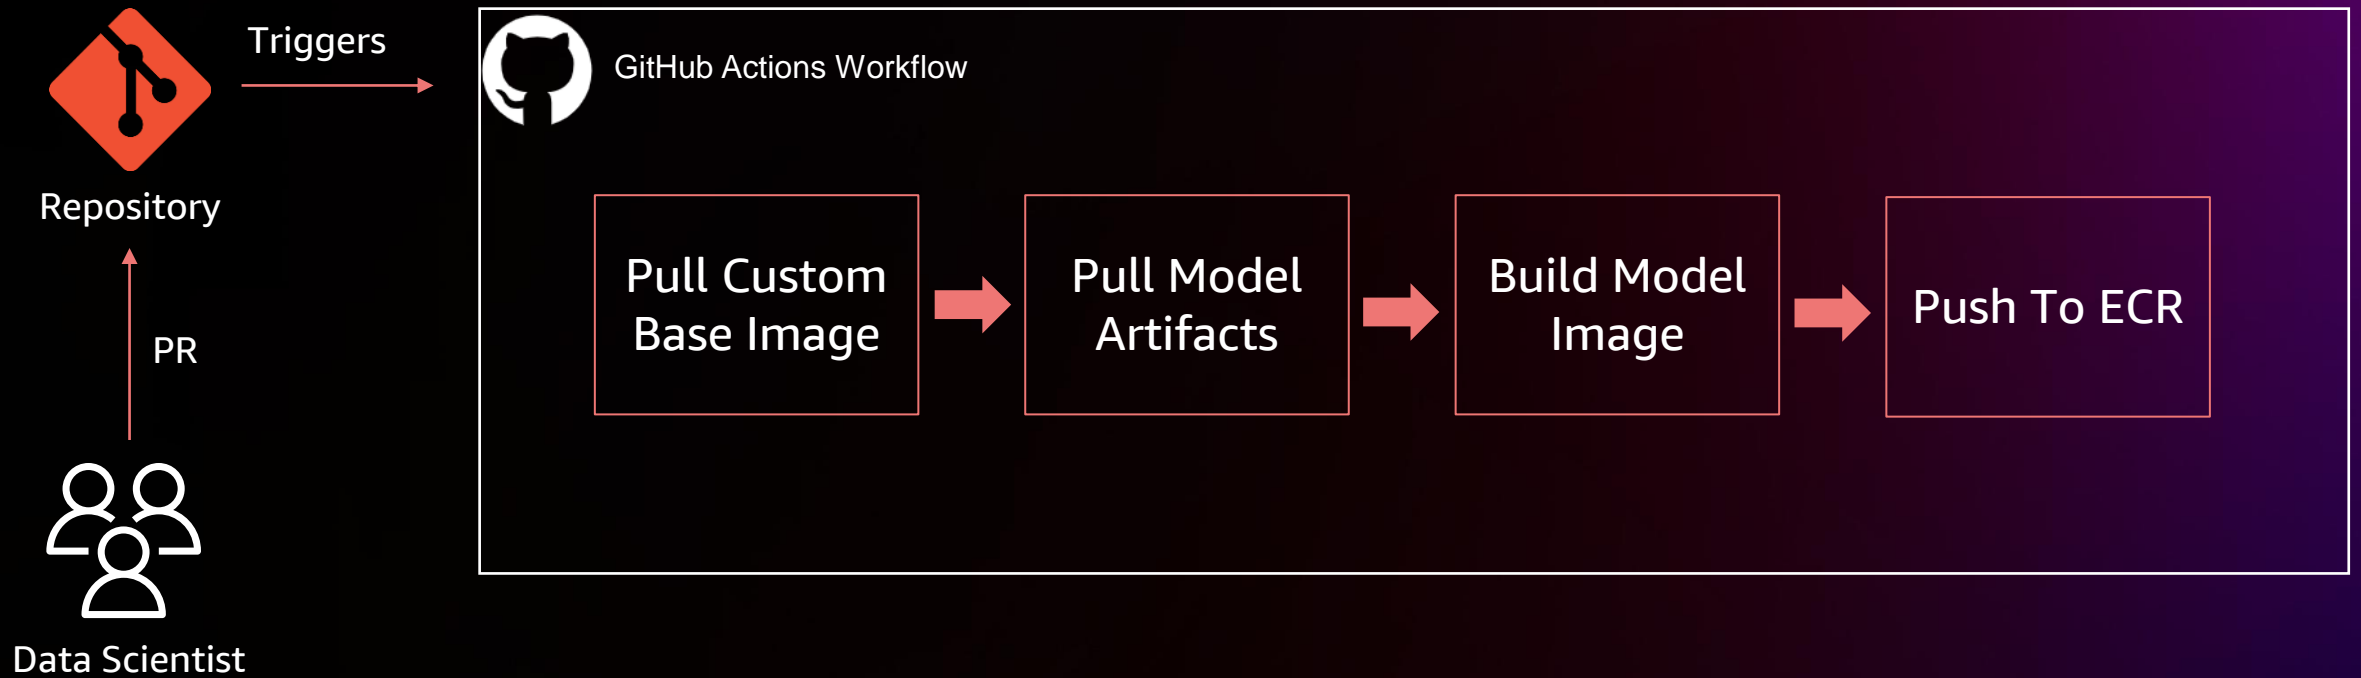
Avoiding scale to 0

- Keep lower environments warm for rapid testing
- Minimal cost for infrequently used models

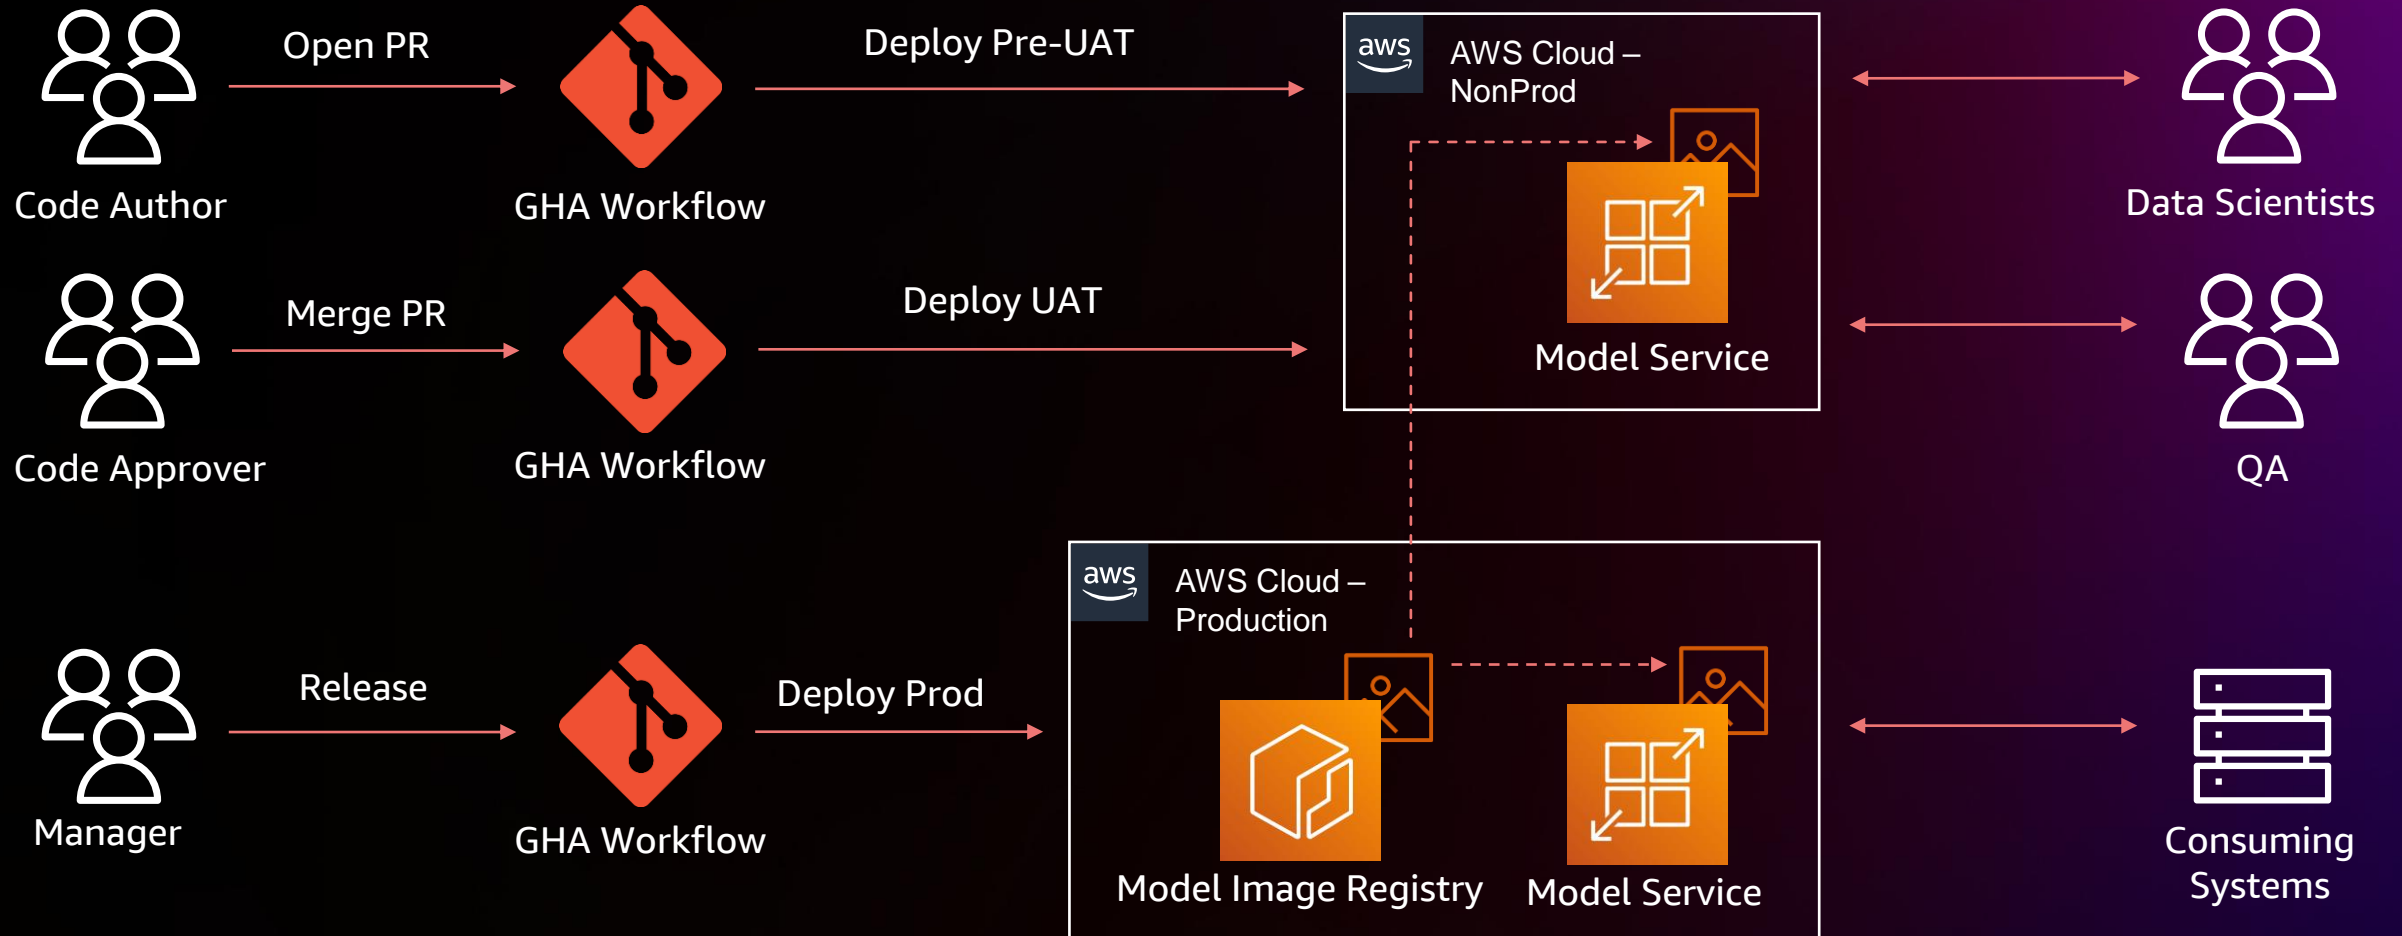
Taking YOUR first container to the cloud



Building the container image

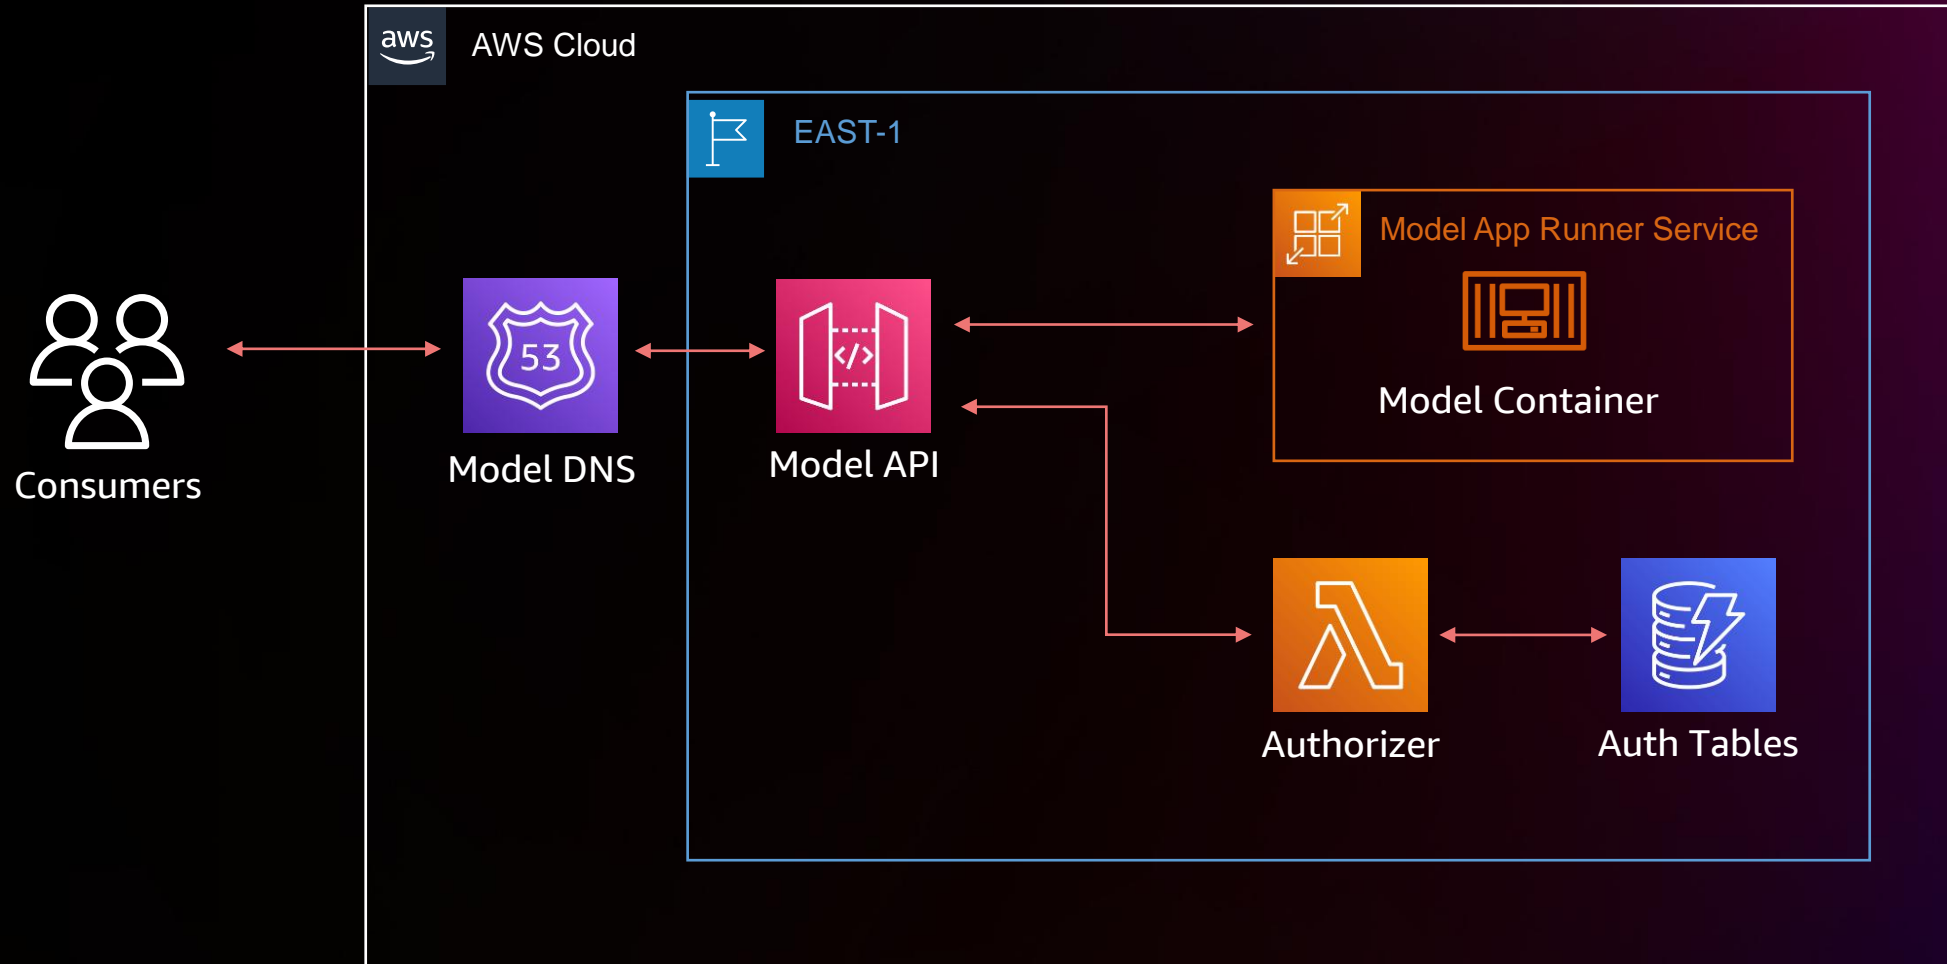


Path to production

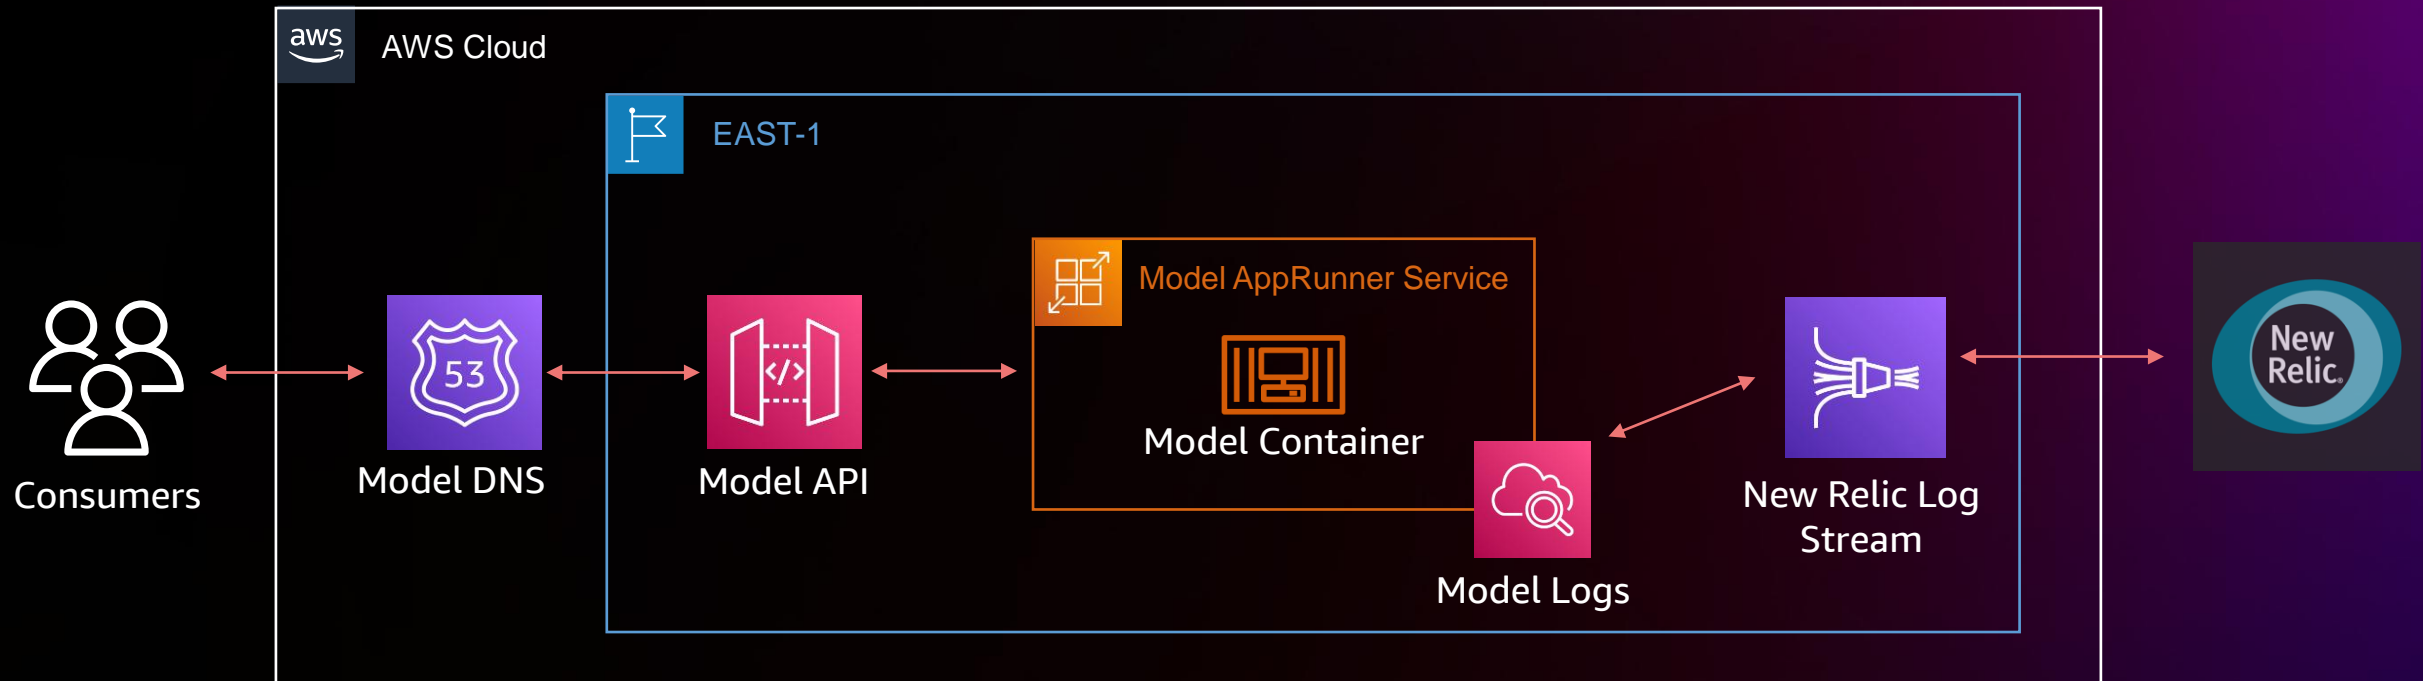


Extending YOUR first container to the cloud

Amazon API Gateway integration




Integrating App Runner logging capabilities



Opportunities for improvement

- AWS CloudFormation support
- Compute SKUs
- Steady-state optimization

App Runner – Launch timeline



AWS App Runner launches in 2022

February – VPC Egress

February – Java run time and Node.js 14

March – FIPS endpoint support

April – AWS X-Ray integration

August – Support for alias records

September – Cross-Region ECR support

September – Node.js 16 runtime

October – Support for .NET core, PHP, Ruby, and Go

October – VPC Ingress (private services)

Please visit our GitHub roadmap to upvote:

<https://github.com/aws/apprunner-roadmap/issues>

Want to learn more?

- CON406 – A close look at AWS Fargate and AWS App Runner
- CON305 – Deploy Node.js using App Runner for applications at scale
- CON408 – Deploying containerized applications in minutes with AWS App Runner

Thank you!

Omar Lari

Twitter: @omarlari

Miller Fanning

LinkedIn: miller-fanning



Please complete the session survey in the **mobile app**



© 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.