

AWS re:Invent

NOV. 28 – DEC. 2, 2022 | LAS VEGAS, NV

CON311-R

Best practices for deploying microservices on Amazon ECS

Uttara Sridhar

Software Development Manager
Amazon Web Services

Vibhav Agarwal

Senior Product Manager
Amazon Web Services



© 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Agenda

Microservices: What and why

Building microservices with Amazon ECS


Best practices for Amazon ECS deployments:

- Continuous deployment

- Safety

- Speed

Amazon Elastic Container Service (Amazon ECS) overview



Amazon EC2

AWS Regions, AWS Wavelength,
AWS Local Zones, AWS Outposts



AWS Fargate

Serverless container
compute engine



Amazon ECS Anywhere

Customer owned
container hosts

Amazon ECS is a serverless container orchestrator



Managed by AWS, no cluster versions, no data store to maintain or scale, free to use with Amazon EC2 or AWS Fargate

Works with different “compute” types, including AWS Fargate

Fargate is a serverless compute engine for containers



Managed by AWS, no AMIs to maintain, no Amazon EC2 instances to provision, scale or manage, pay for what you need

Fargate is where your containers will run; you don't interact with it directly

Amazon ECS customers innovate quickly and efficiently



COVID-19 vaccines



Address exponential demand
for delivering essentials



Stream entertainment
globally




Microservices: What and why

What are microservices?

“A **microservice architectural style** is an approach to developing a single application as a **suite of small, isolated services**, each running in its **own process** and communicating with lightweight mechanisms, often an HTTP resource API.”

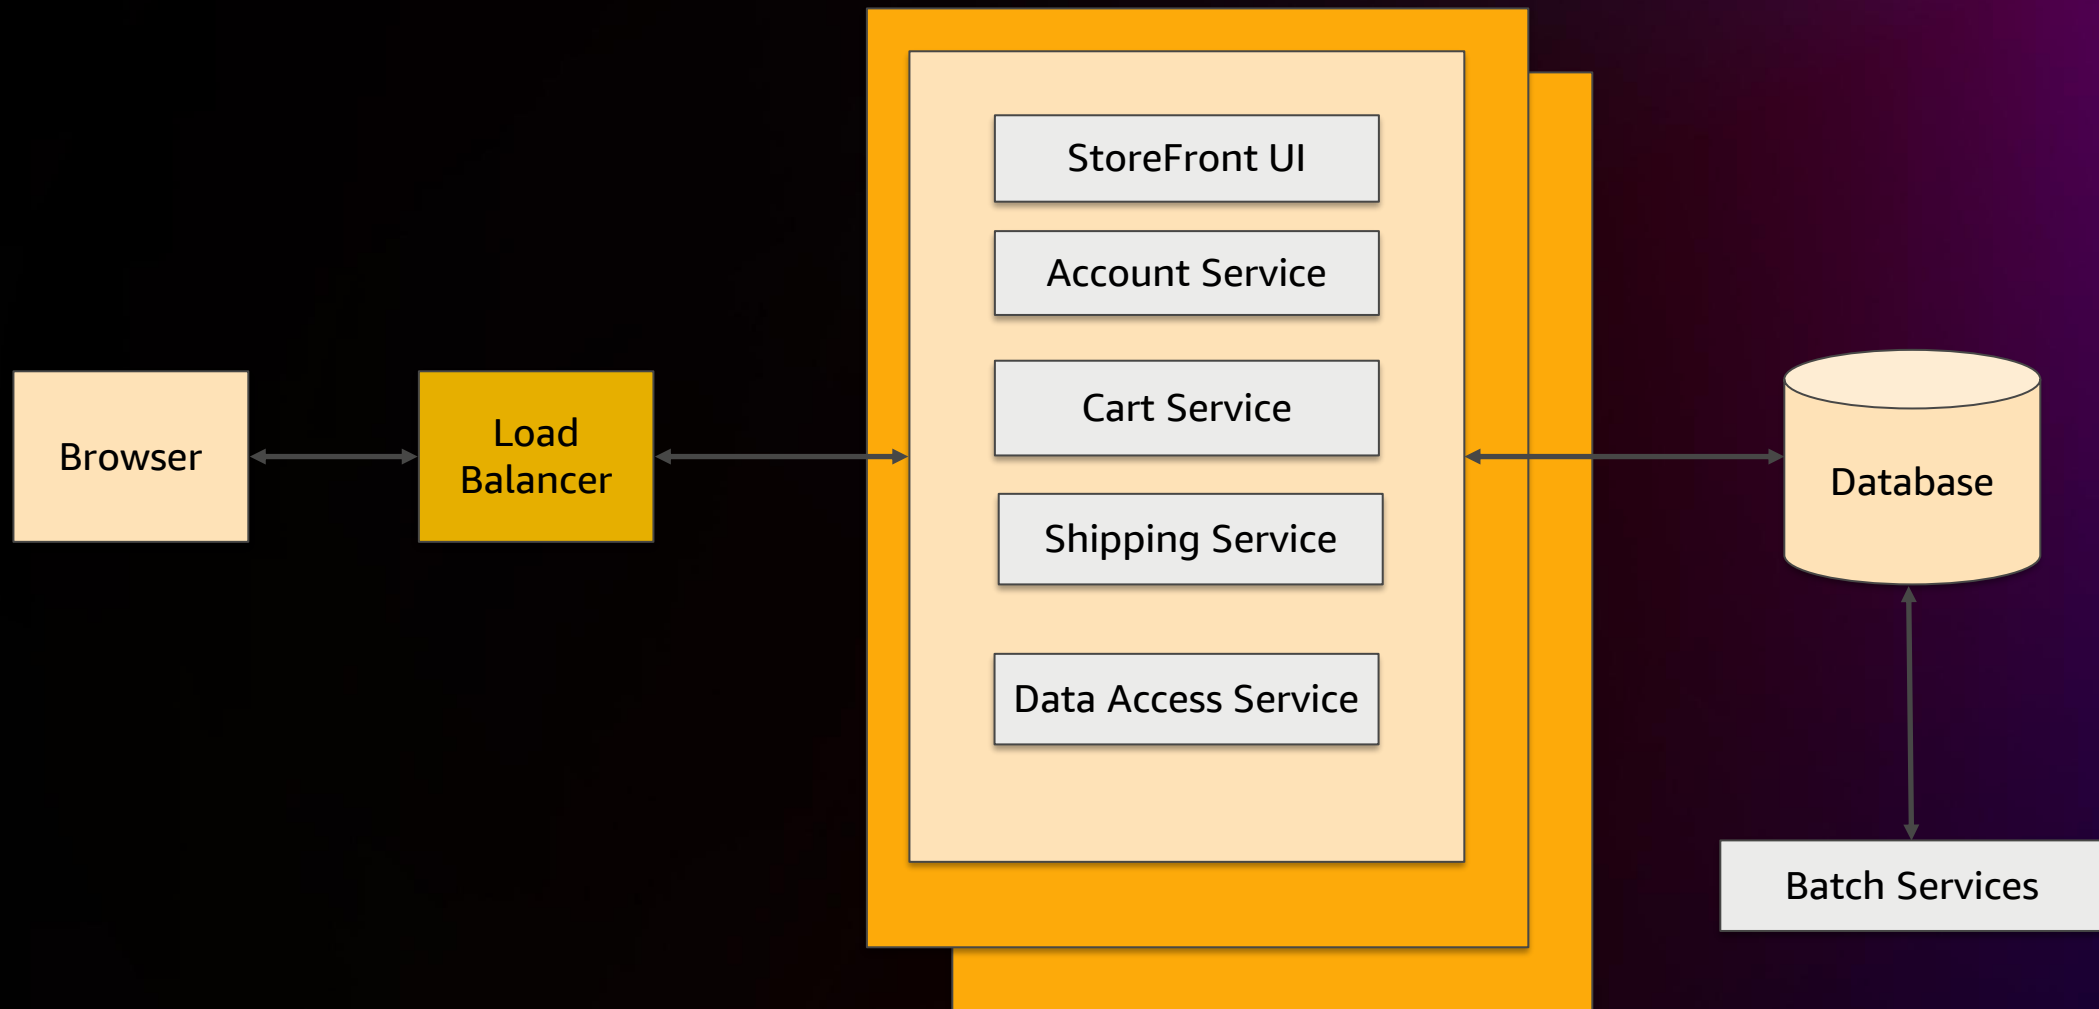
James Lewis and Martin Fowler

Microservice tenets

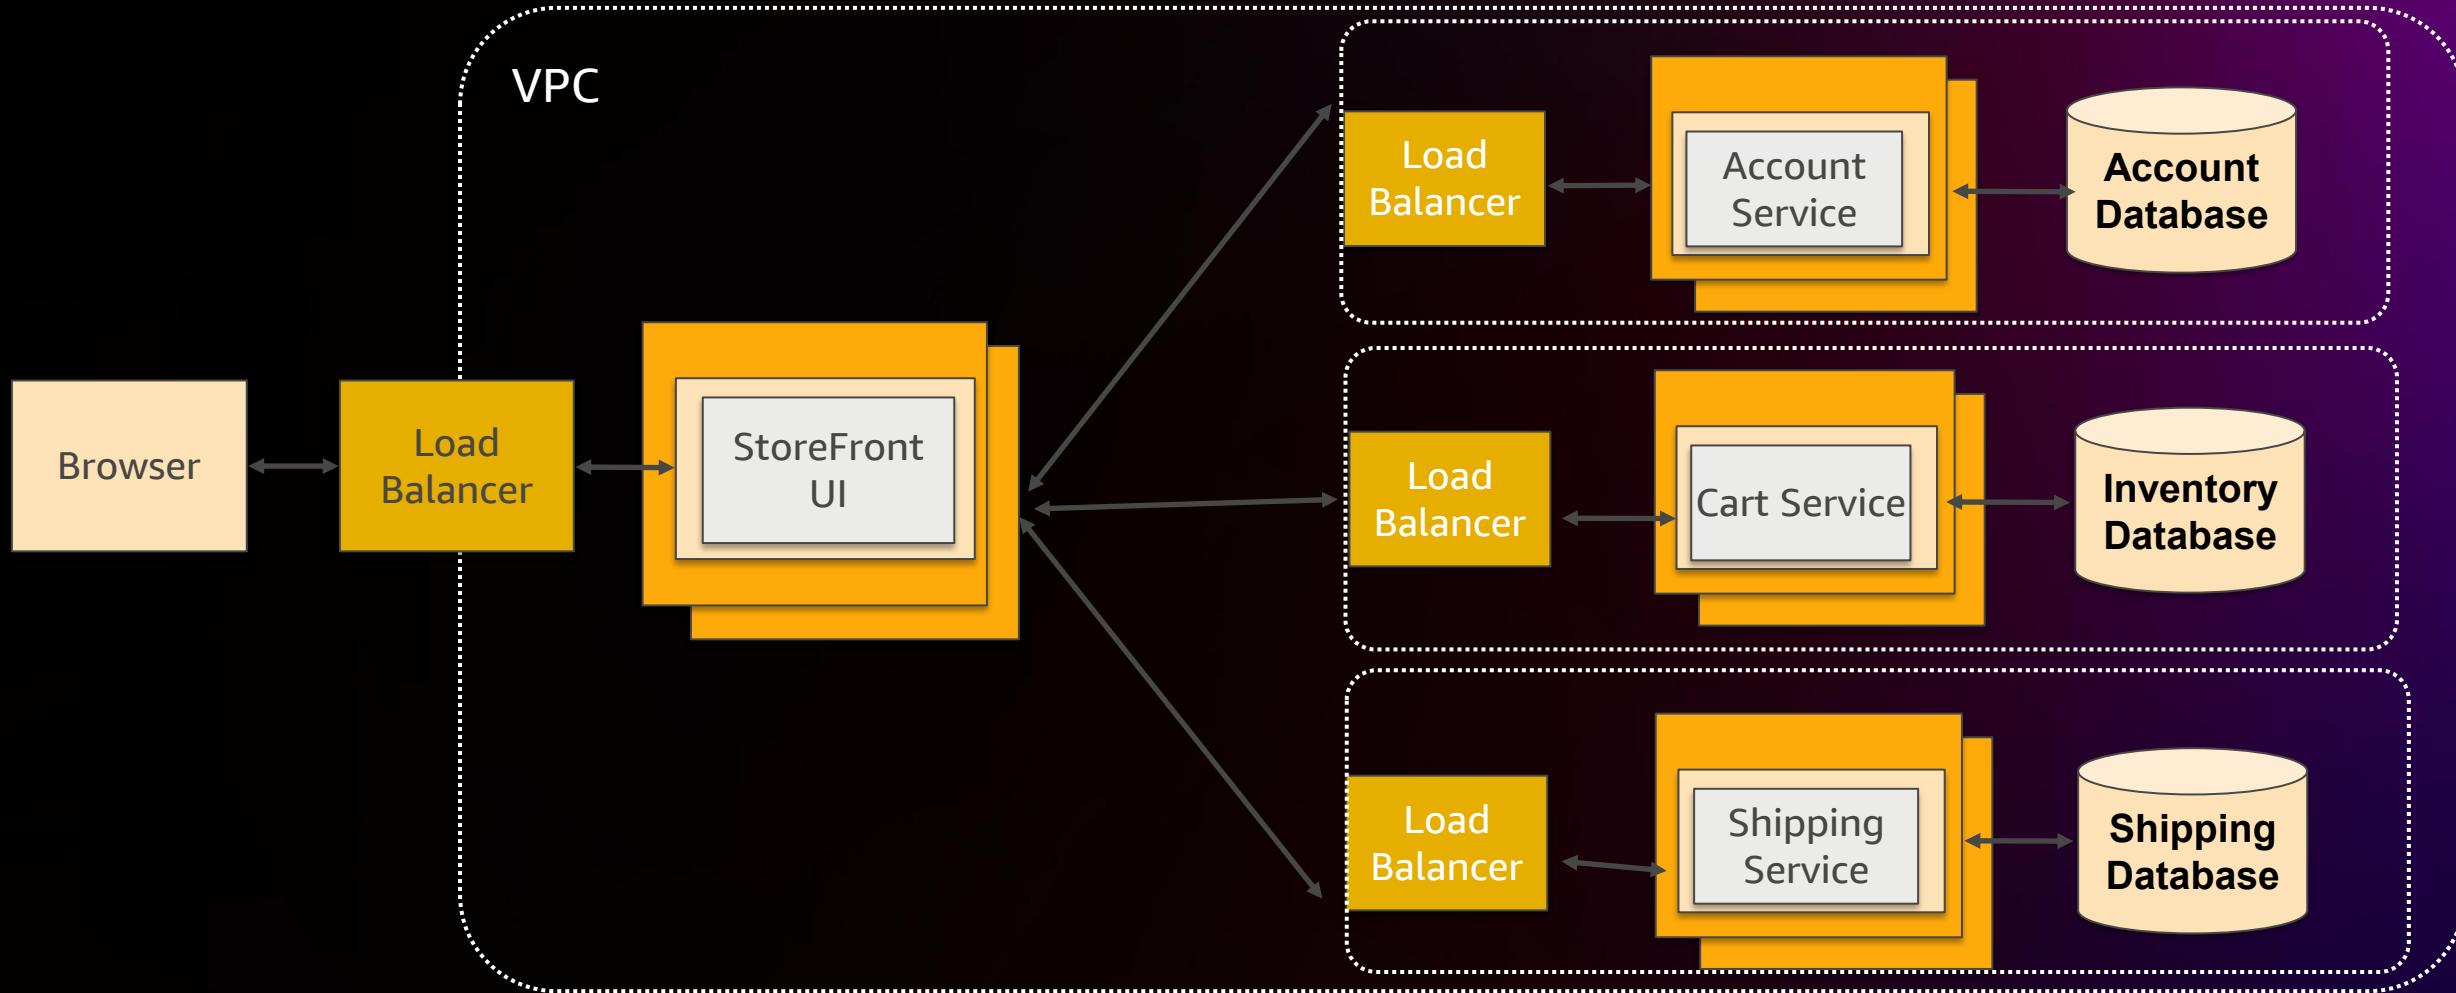


- Well-defined contract
- Autonomous
- Deployed separately
- Scaled independently
- Single responsibility
- Owns, manages, and encapsulates its data
- Decentralized ownership

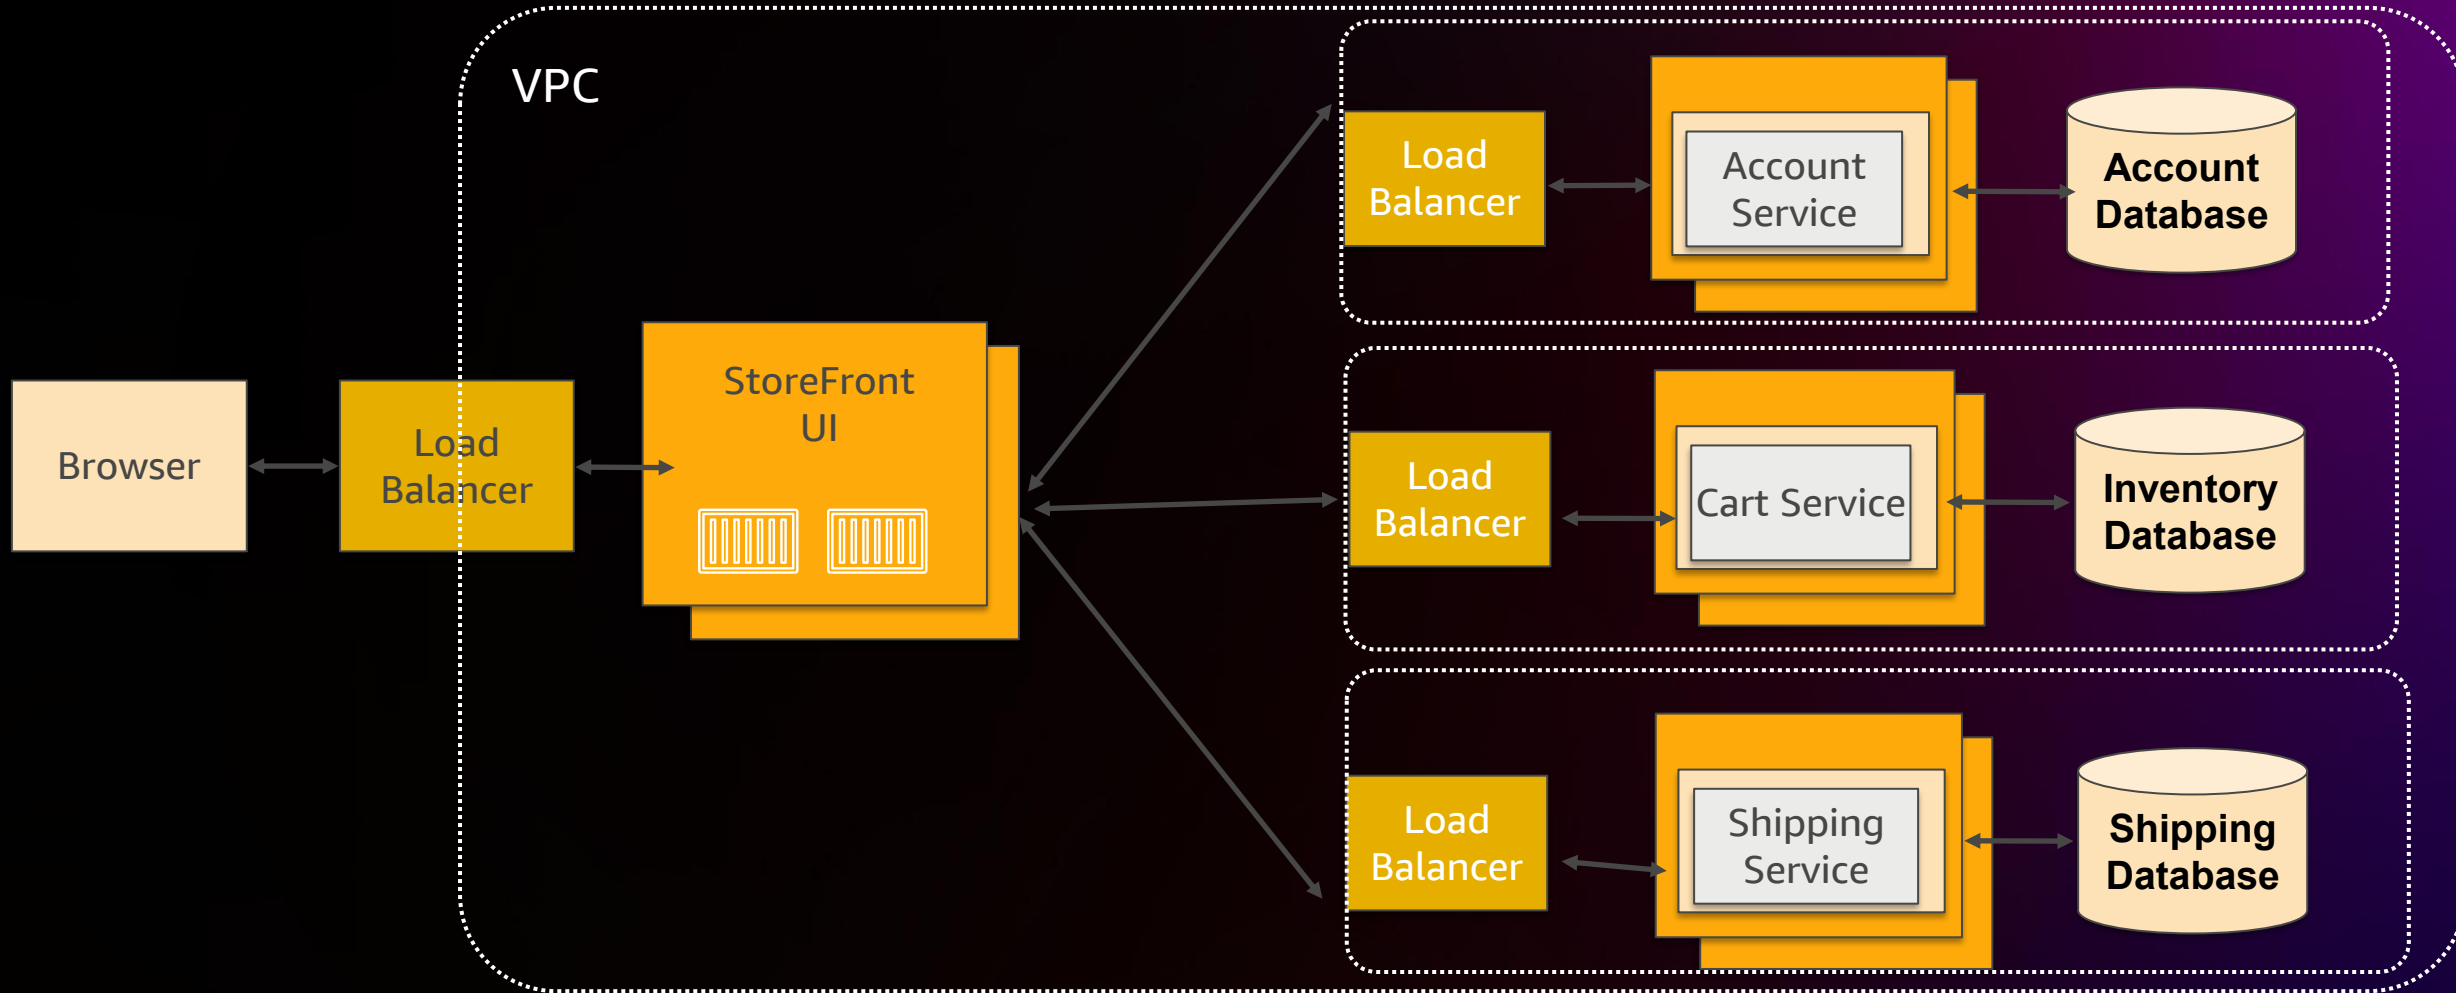
Monolithic architectures are tightly coupled



Microservices architecture



How to run microservices? Containers!

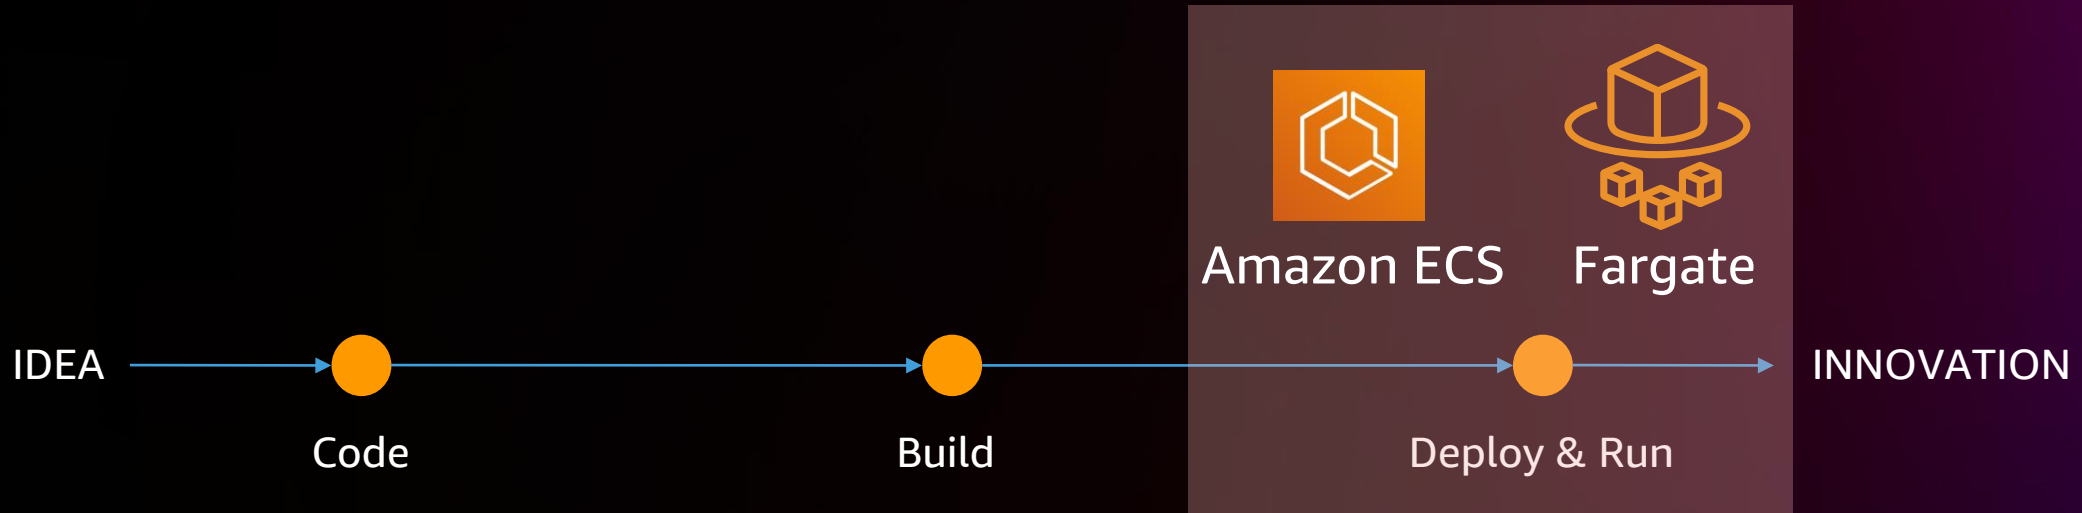


Containers + microservices: Modern digital assembly line for software delivery

Deploying applications as containerized microservices **accelerates innovation** while maintaining **reliability** and **cost efficiency**




Amazon ECS + Fargate: Simple and powerful microservice deployments



Microservice deployments in Amazon ECS

Amazon ECS constructs

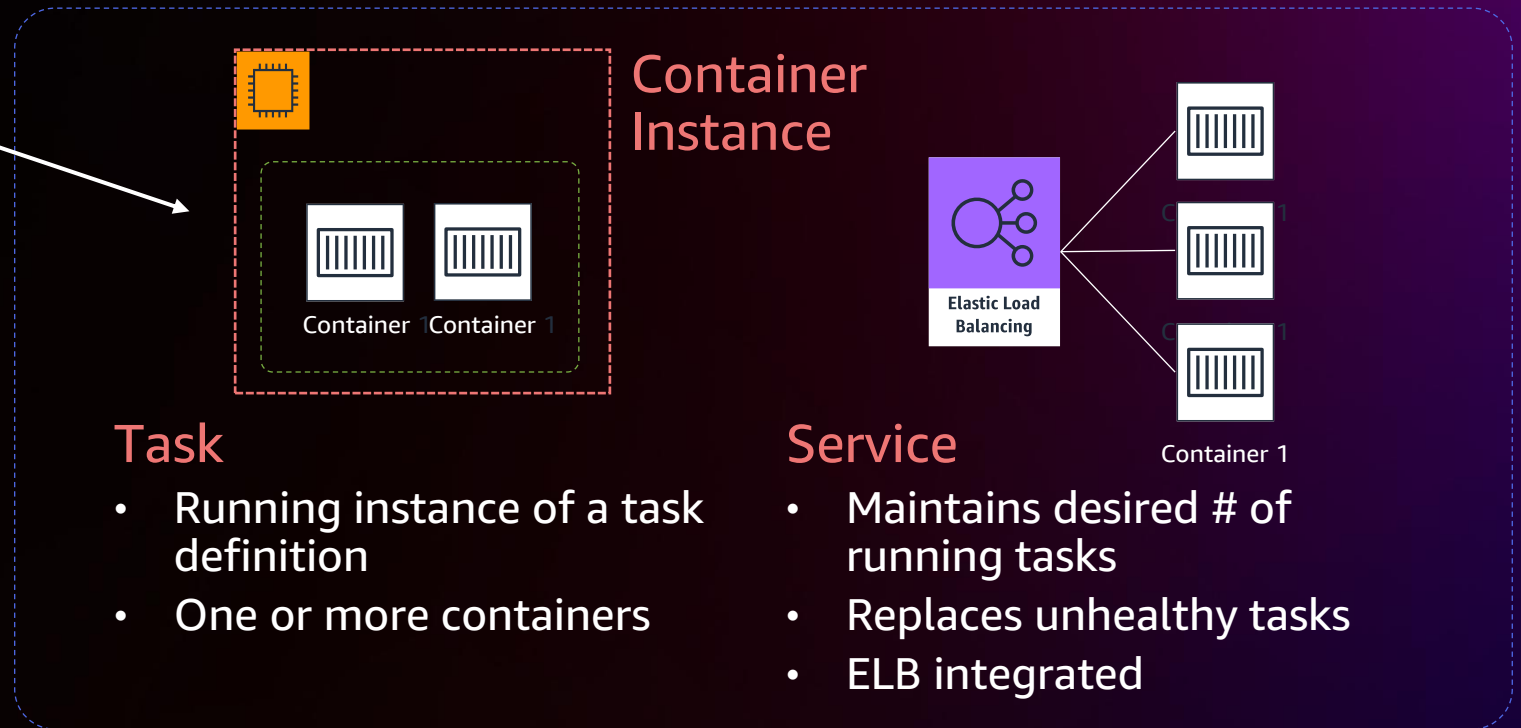


Task definition

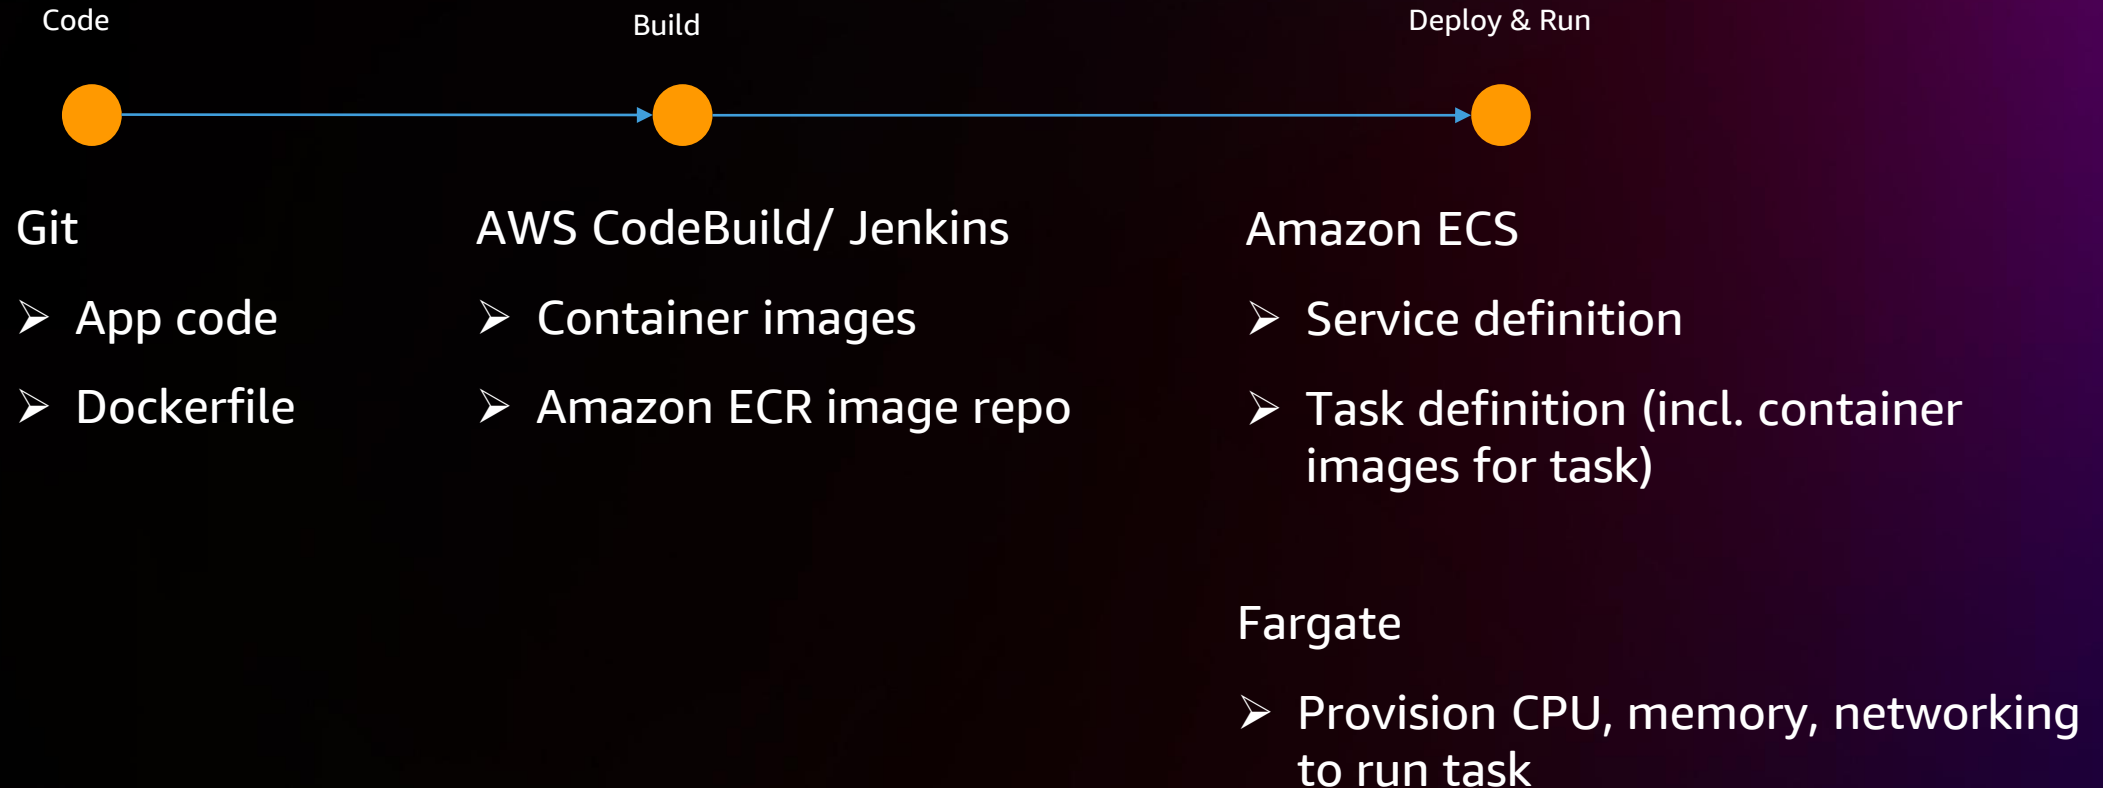
- Template used by Amazon ECS to launch tasks
- Parallels to docker run parameters
- Defines requirements:
 - CPU/memory
 - Container image(s)
 - Logging
 - IAM role
 - Etc.

Cluster

- Resource grouping and isolation
- IAM permissions boundary



Amazon ECS services are commonly used to deploy microservices



Anatomy of an Amazon ECS service deployment

1. Task definition registration
2. Create or update a service
 - i. Amazon ECS service scheduler
 - Provision compute and ENI for each task, Pull container images for task, Task runs
 - For Load Balancer service: Amazon ECS registers task to target group, Amazon ECS checks LB health check status for tasks
3. Service is deployed (desired count of tasks are running and healthy)

What is NOT an Amazon ECS deployment

Scaling a service

Autoscaling

Updating the image tag



A deployment means rolling out a new application version and replacing the old version (if any)

Two key ways to deploy to Amazon ECS

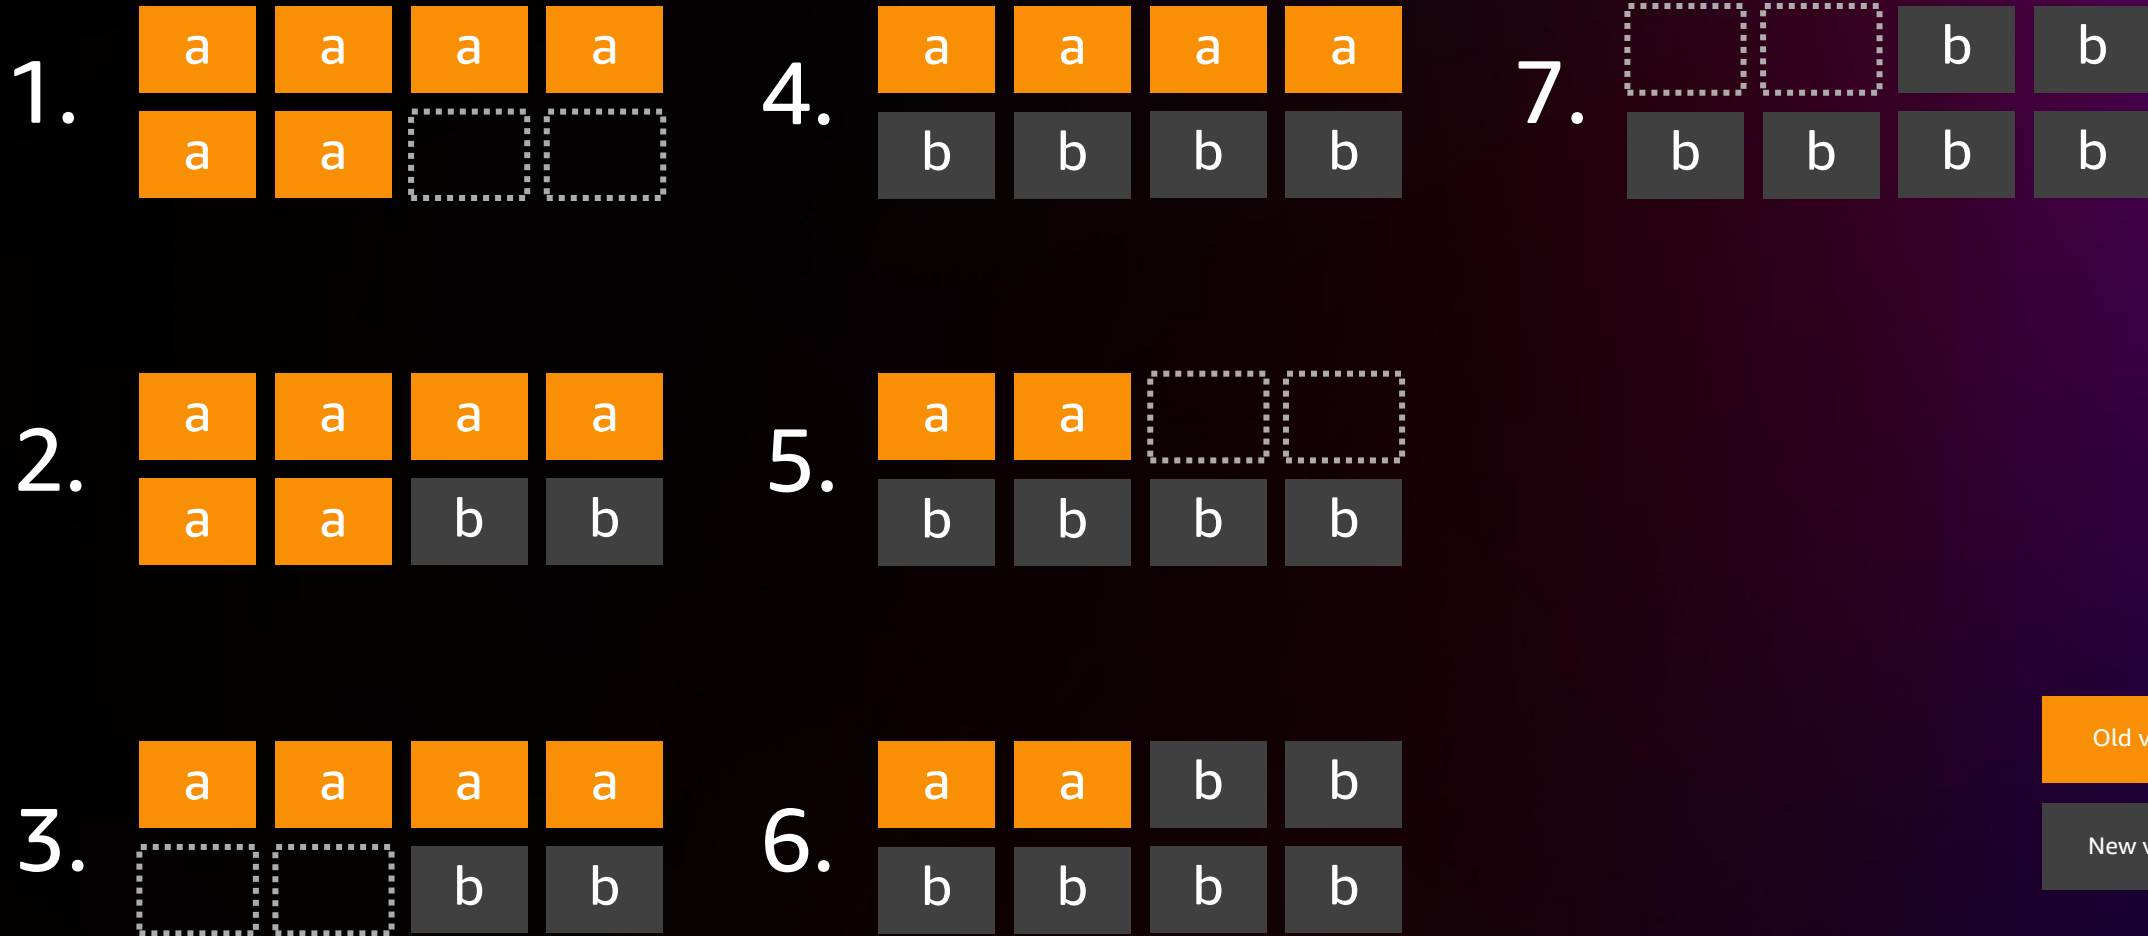
Amazon ECS orchestrated rolling deployment

Incrementally increasing percentage of requests is serviced by the newly deployed application version until 100% is reached

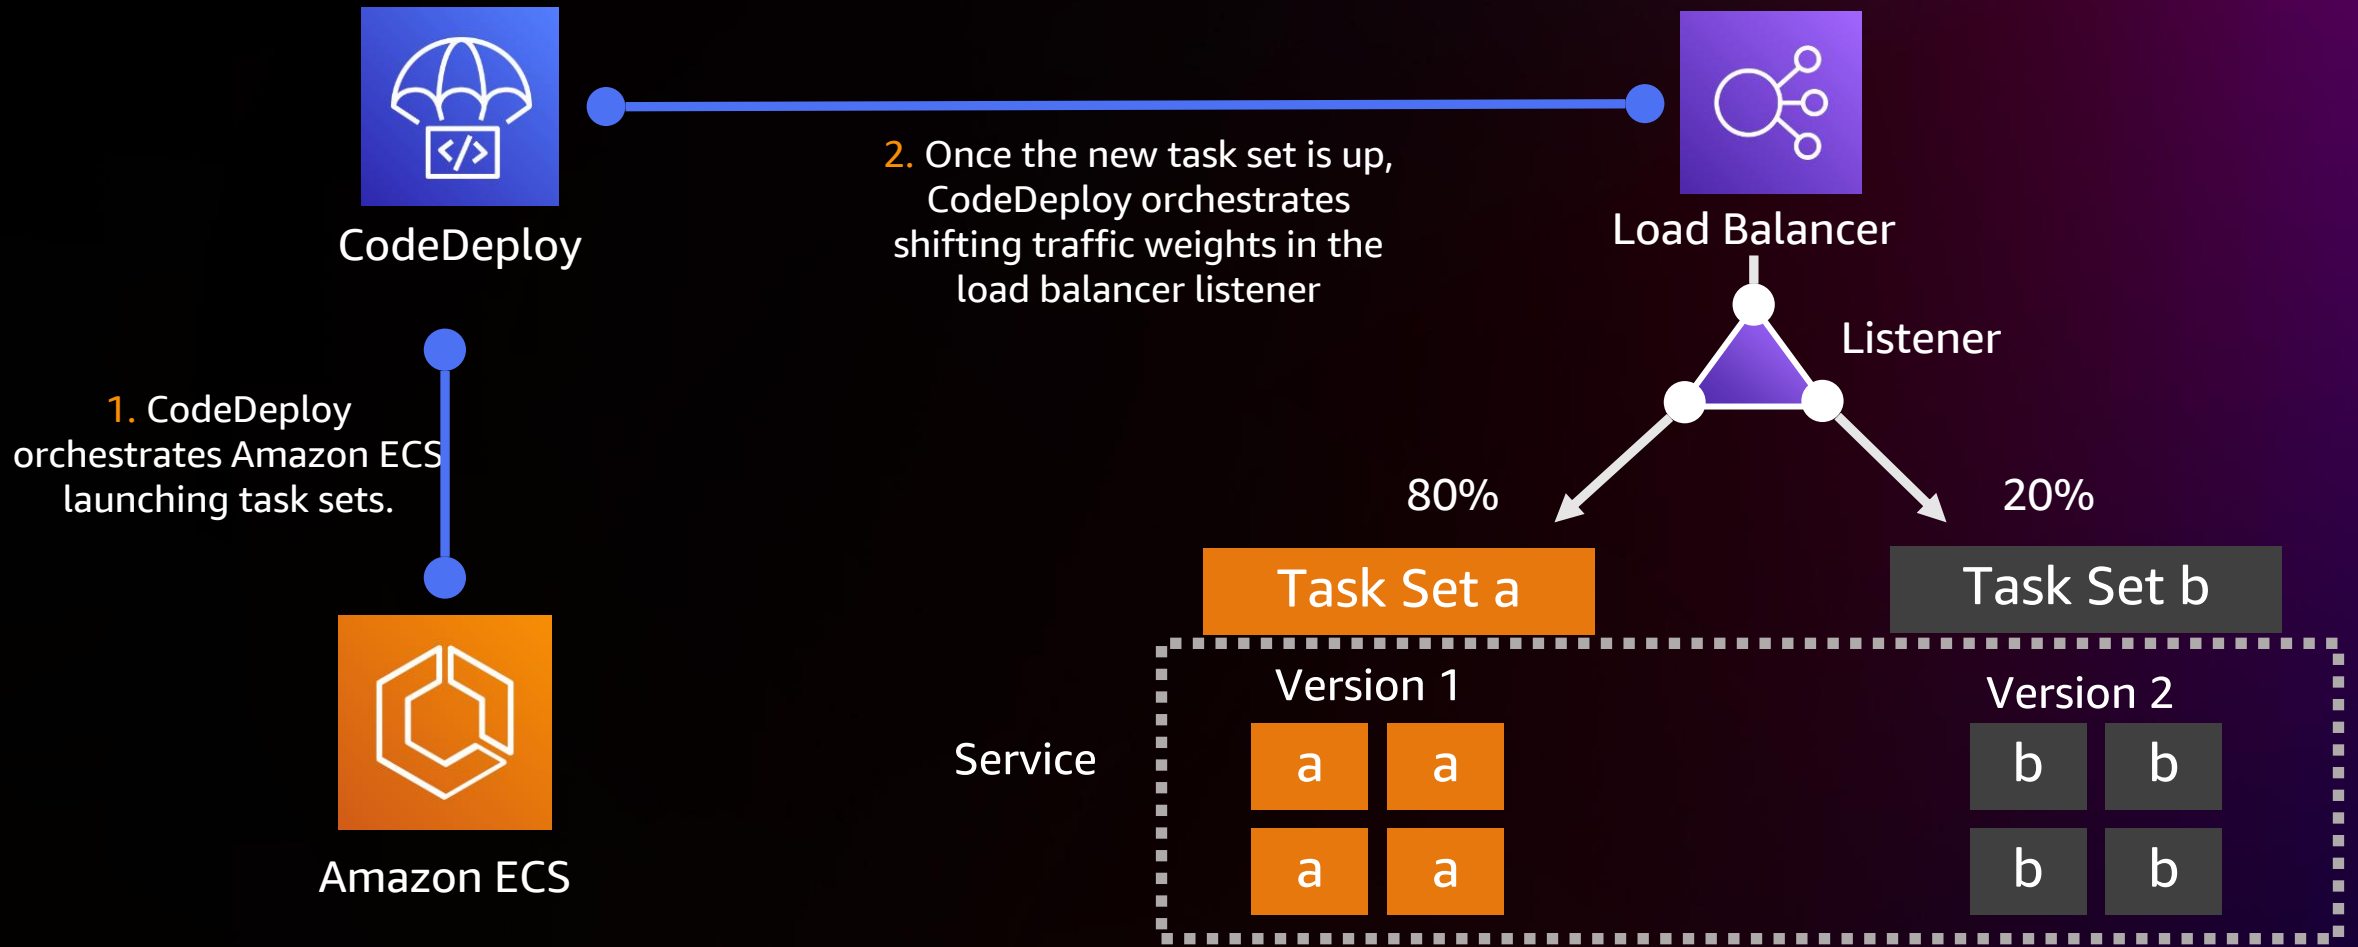
Blue-green and canary deployment with AWS CodeDeploy

Deploy new application version, shift requests to the new infrastructure incrementally, but retain the old infrastructure for rollback purposes

Amazon ECS rolling deployment iteratively creates new tasks and removes old



CodeDeploy orchestrates blue/green deployment with Amazon ECS and ELB



There are many ways to deploy to Amazon ECS

Console

API

Copilot

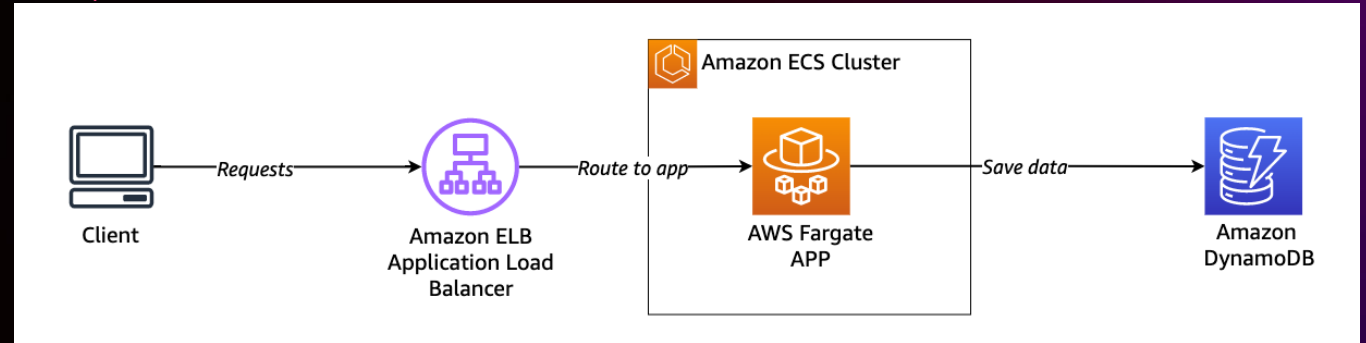
CloudFormation

CDK


Terraform

Github actions

And many more . . .



Application-first interfaces to deploy with Amazon ECS



1. Amazon ECS Blueprints

Terraform templates to deploy apps in minutes

2. Amazon ECS CDK extensions

Program architecture patterns in your choice language

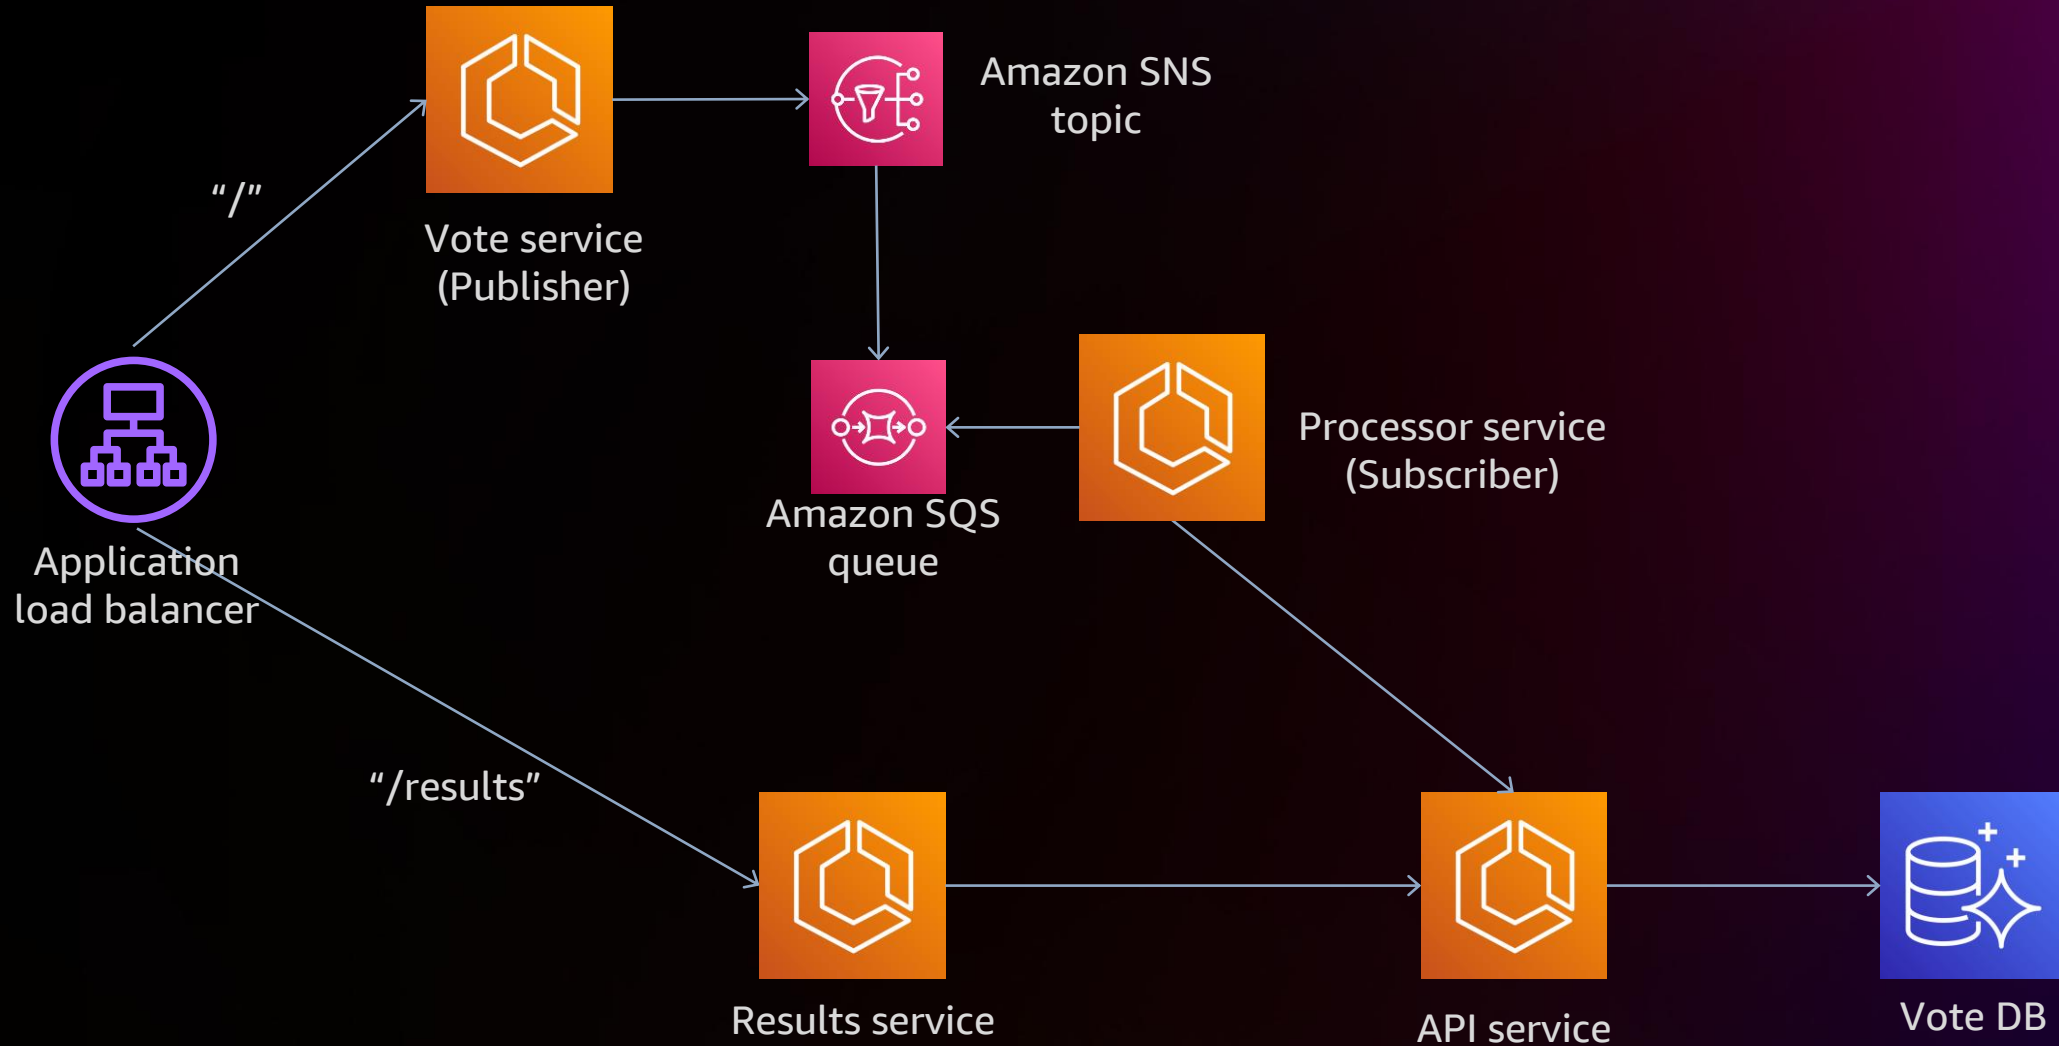
3. AWS Copilot CLI

Configuration-ready experience for common cloud architectures. Preconfigured continuous delivery pipeline with auto scaling, metrics, alarms, and logs.

**Let's see an Amazon ECS
deployment in action!**



First, let's look at the application architecture



Deployment best practices



3 key dimensions for deployments



**Continuous
deployment**



Safety



Speed


Best practices: Continuous deployment

When the impact of change is small, release velocity can increase



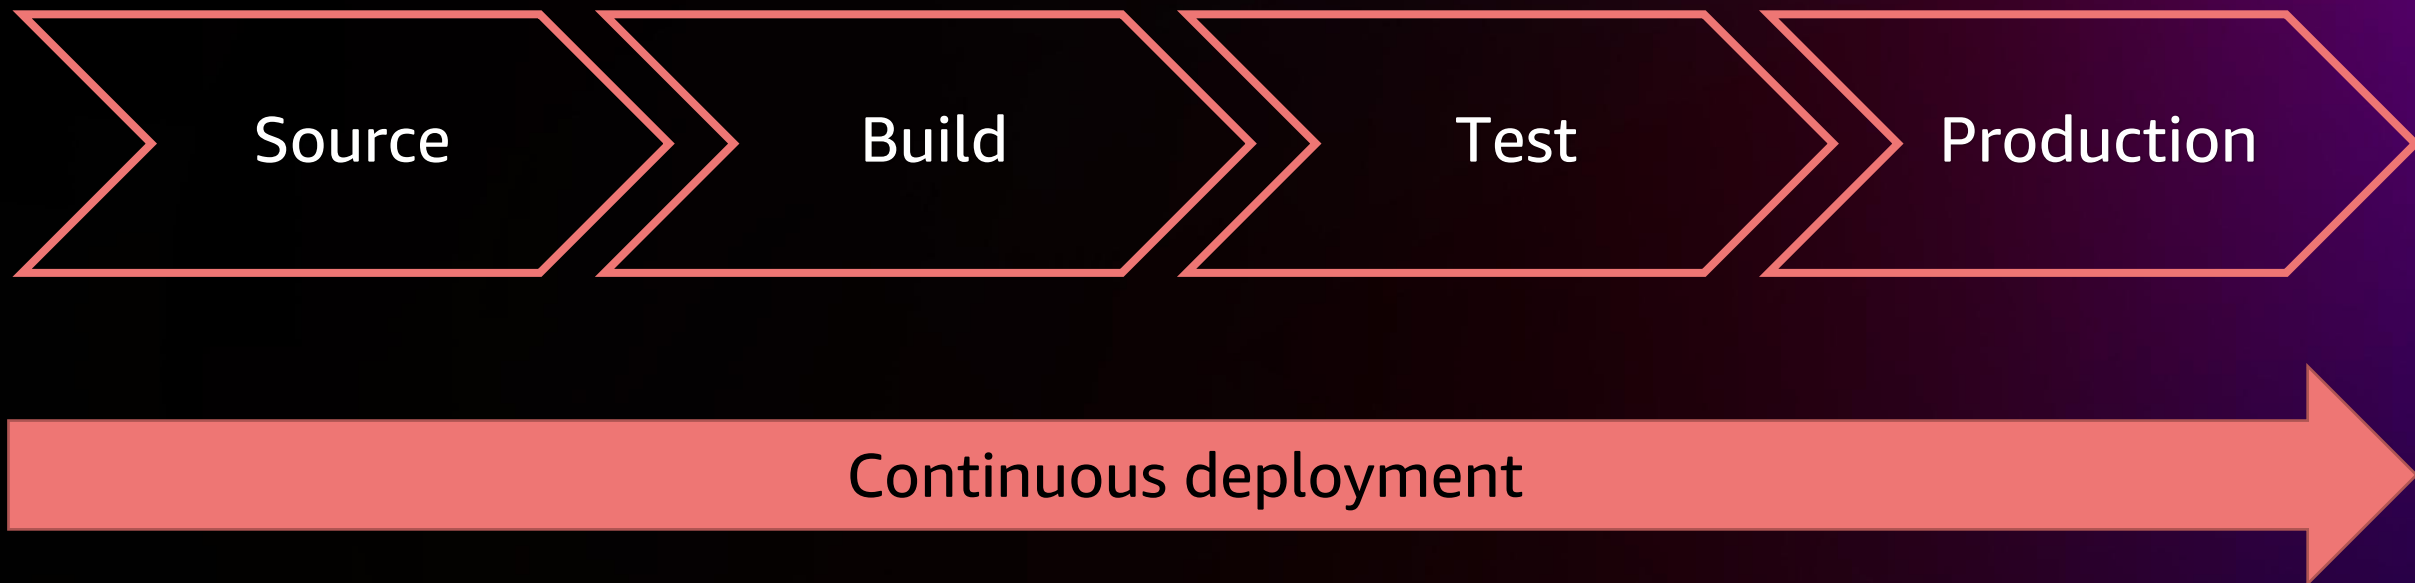
Monolith

Does everything

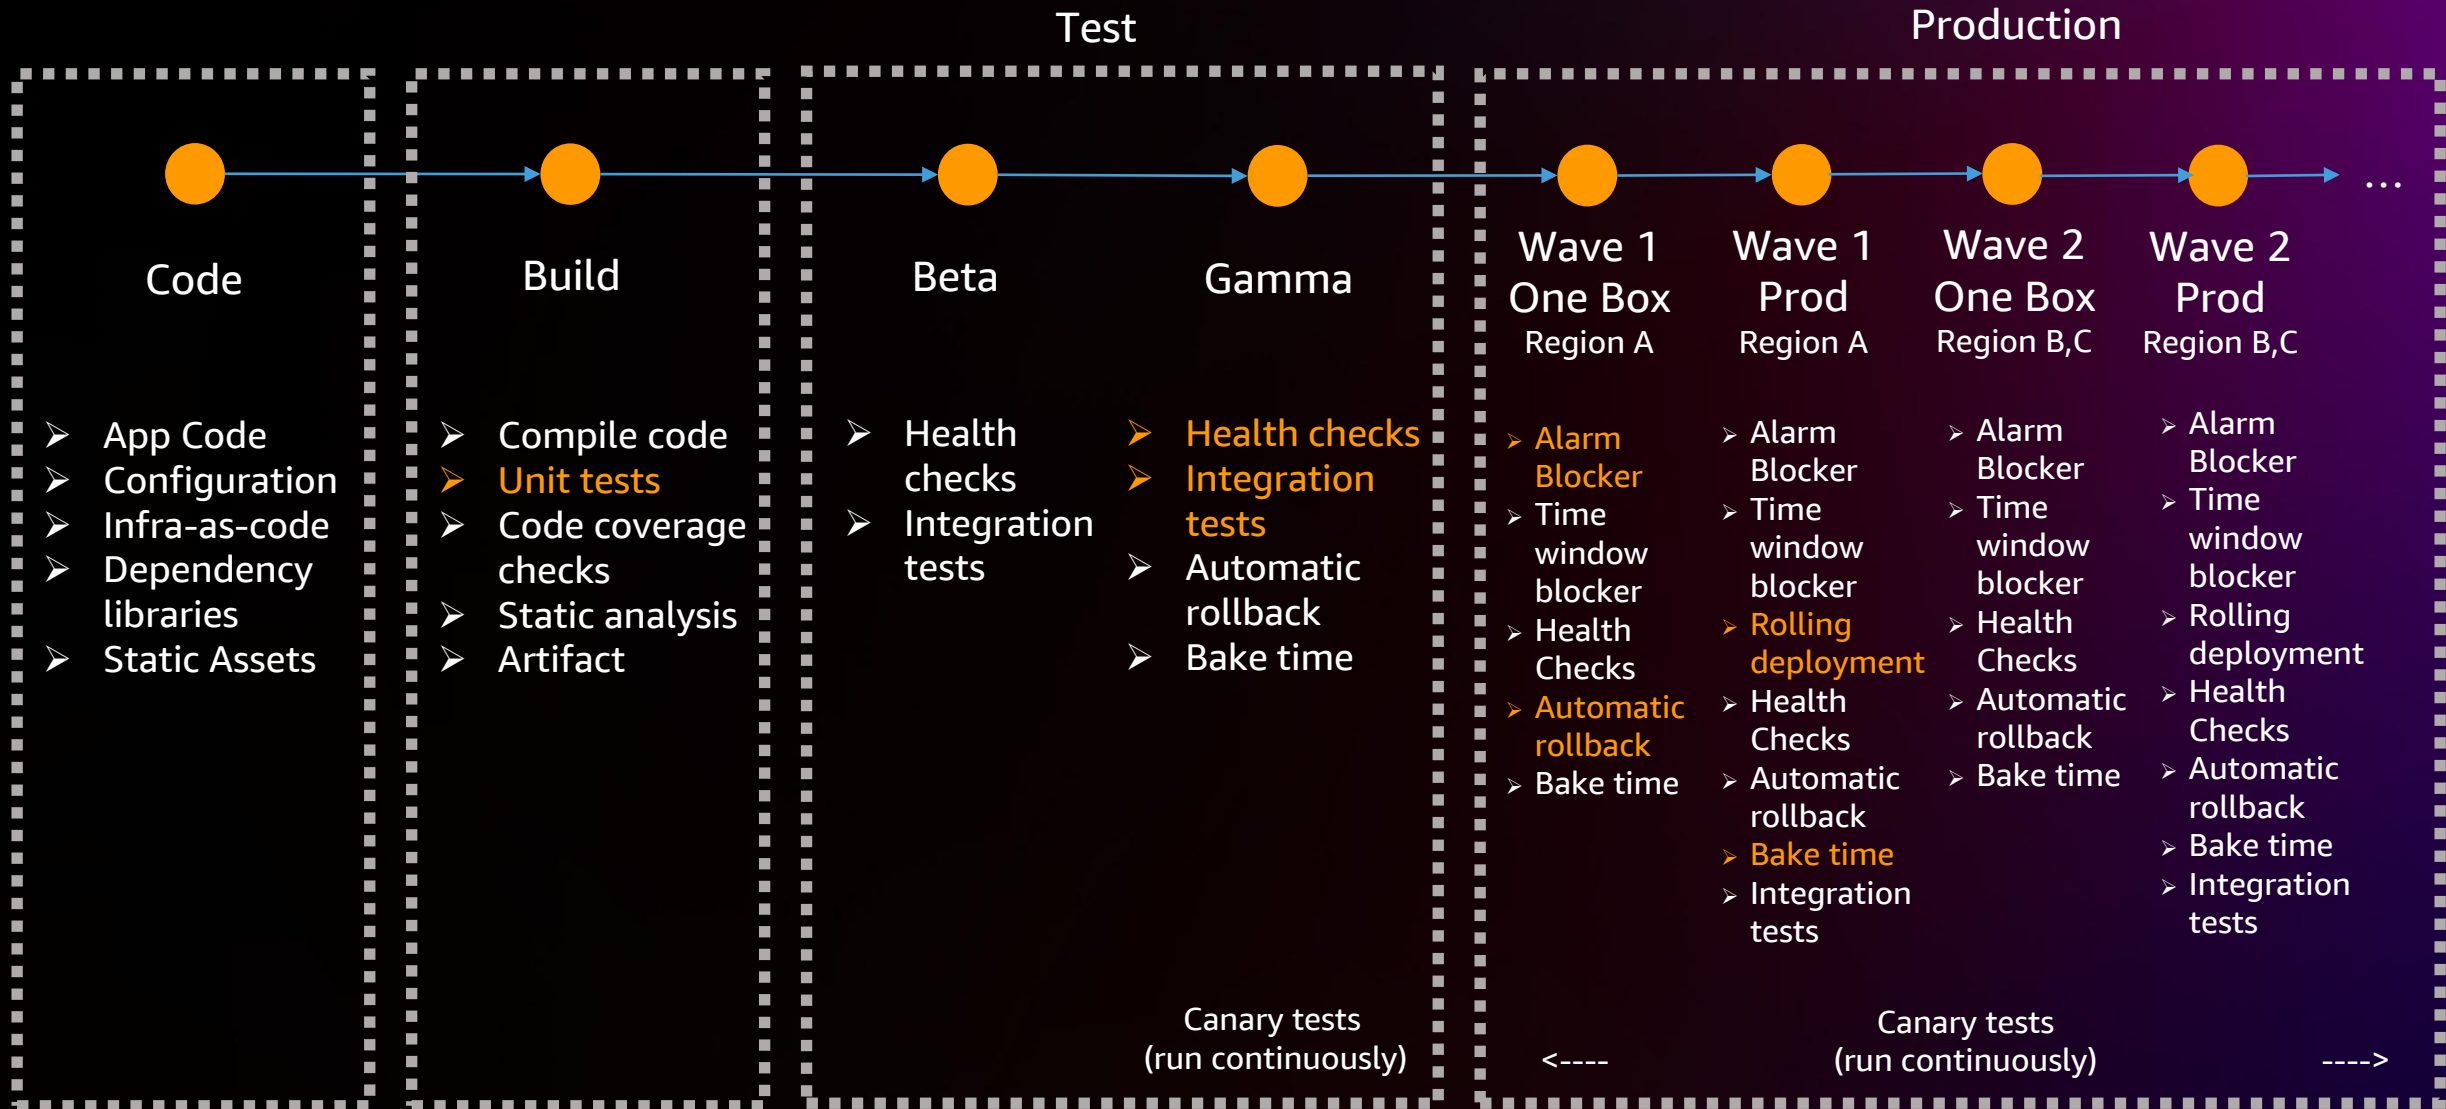


Microservices

Does one thing



Safe continuous deployments at Amazon



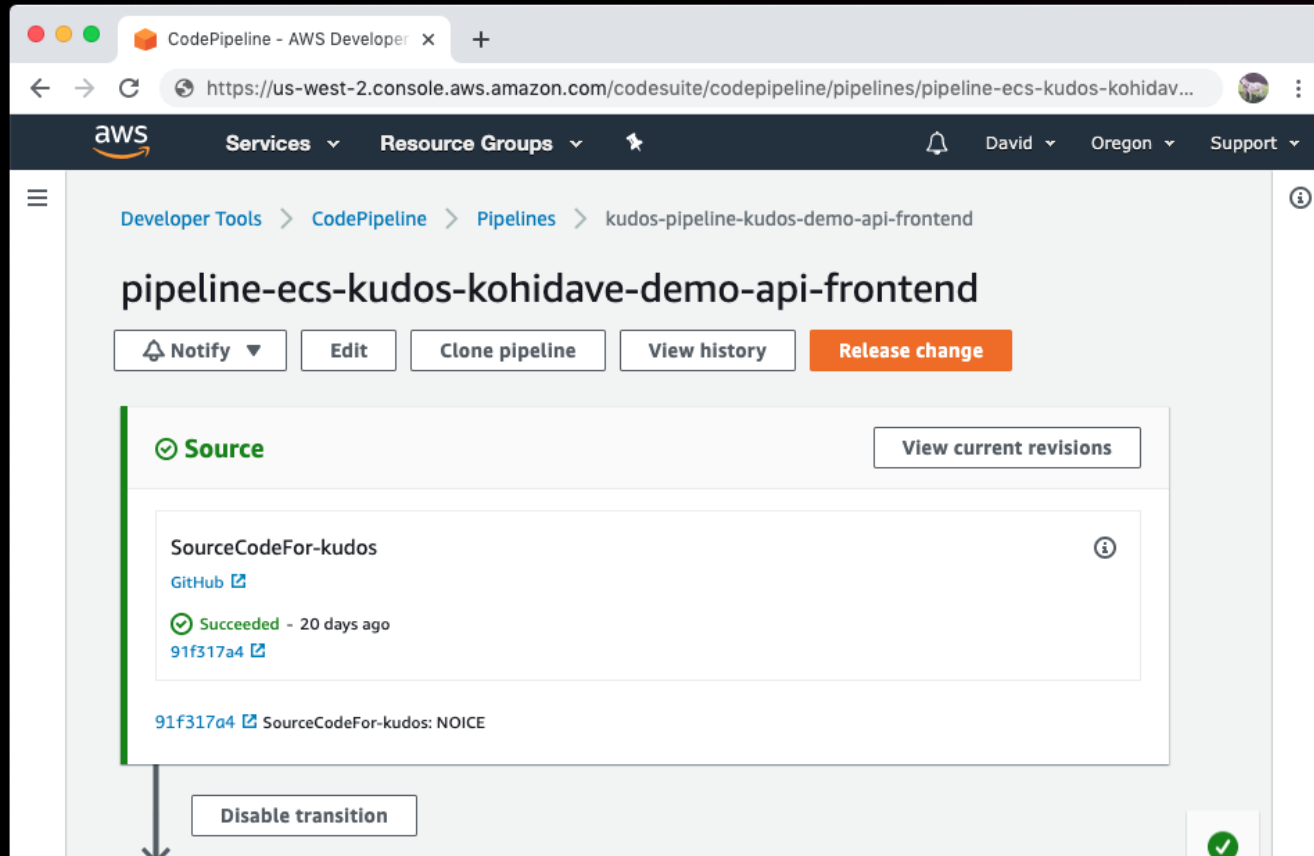
Continuous deployment goals



Continuous deployment

1. Automatically deploy new changes to staging environments for testing
2. Deploy to production safely without affecting customers
3. Deliver to customers faster: Increase deployment frequency, and reduce change lead time and change failure rate

AWS Copilot simplifies continuous deployments



Uses AWS CodePipeline

Deploy to one or many environments

Can configure tests

Fully customize or use own CI/CD

Best practices: Safety

2 key goals for deployment safety

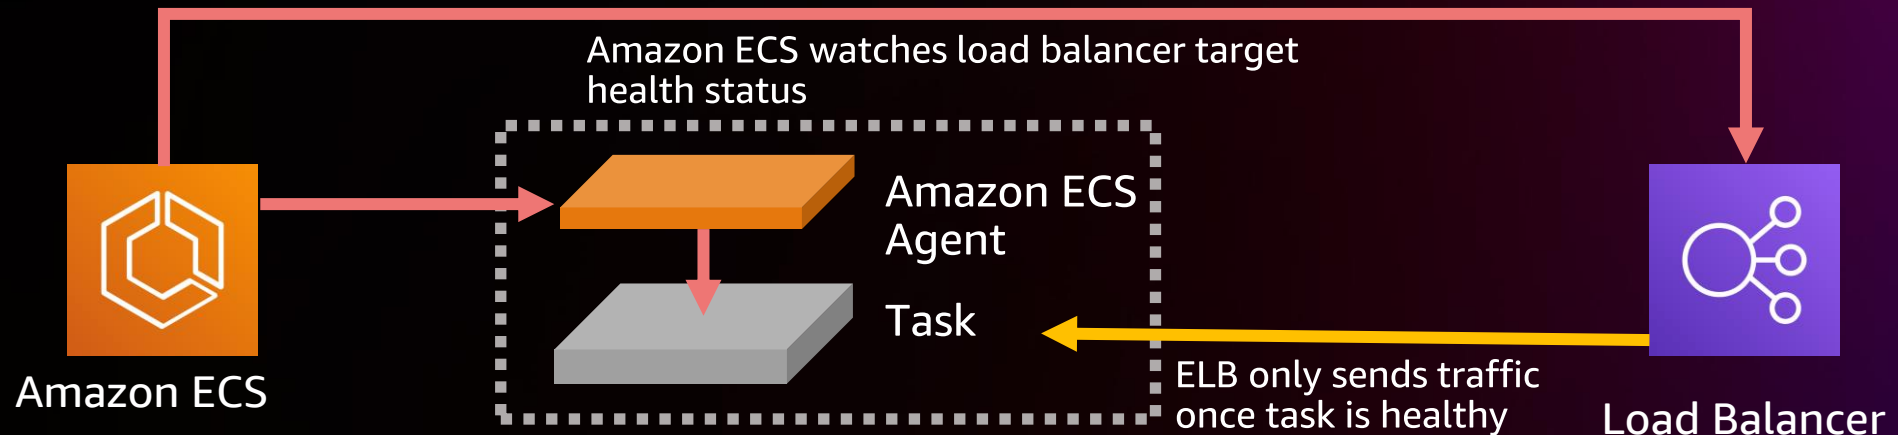
Prevent bad changes from reaching end users

Minimize downtime

- Handle shutdown gracefully
- Detect failures quickly and rollback



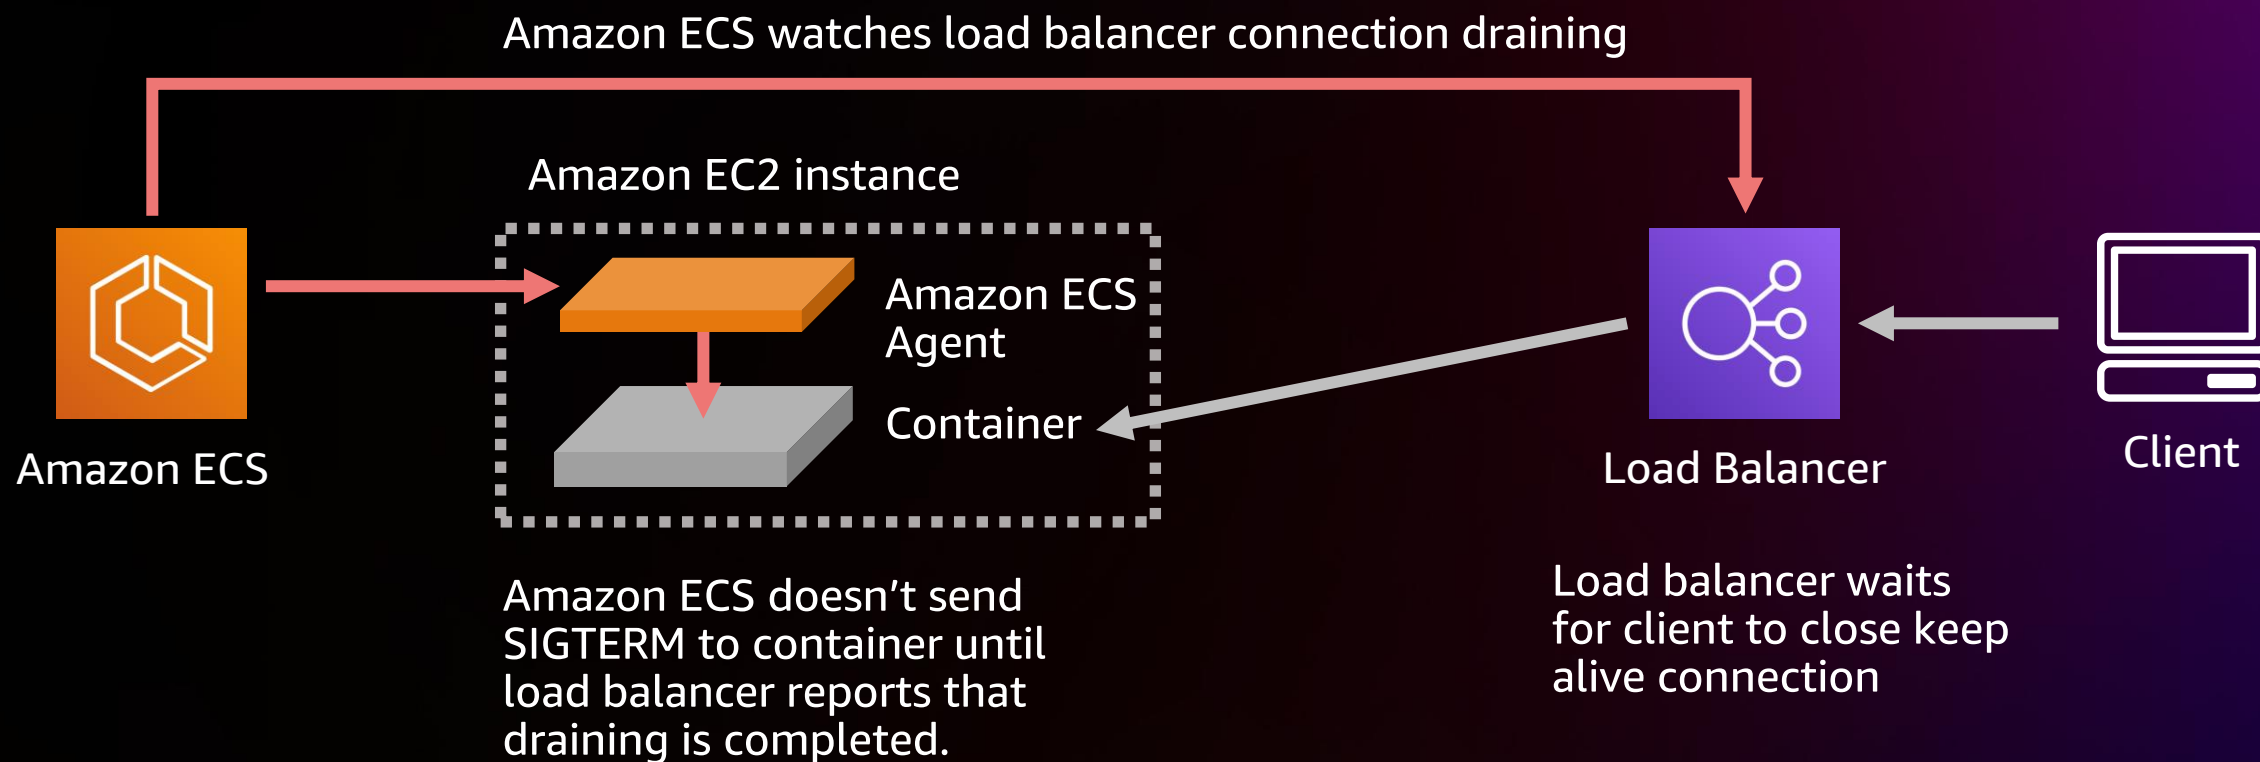
Amazon ECS + ELB health checks prevent bad changes from reaching production



Tip: Add a dedicated health check endpoint for better control

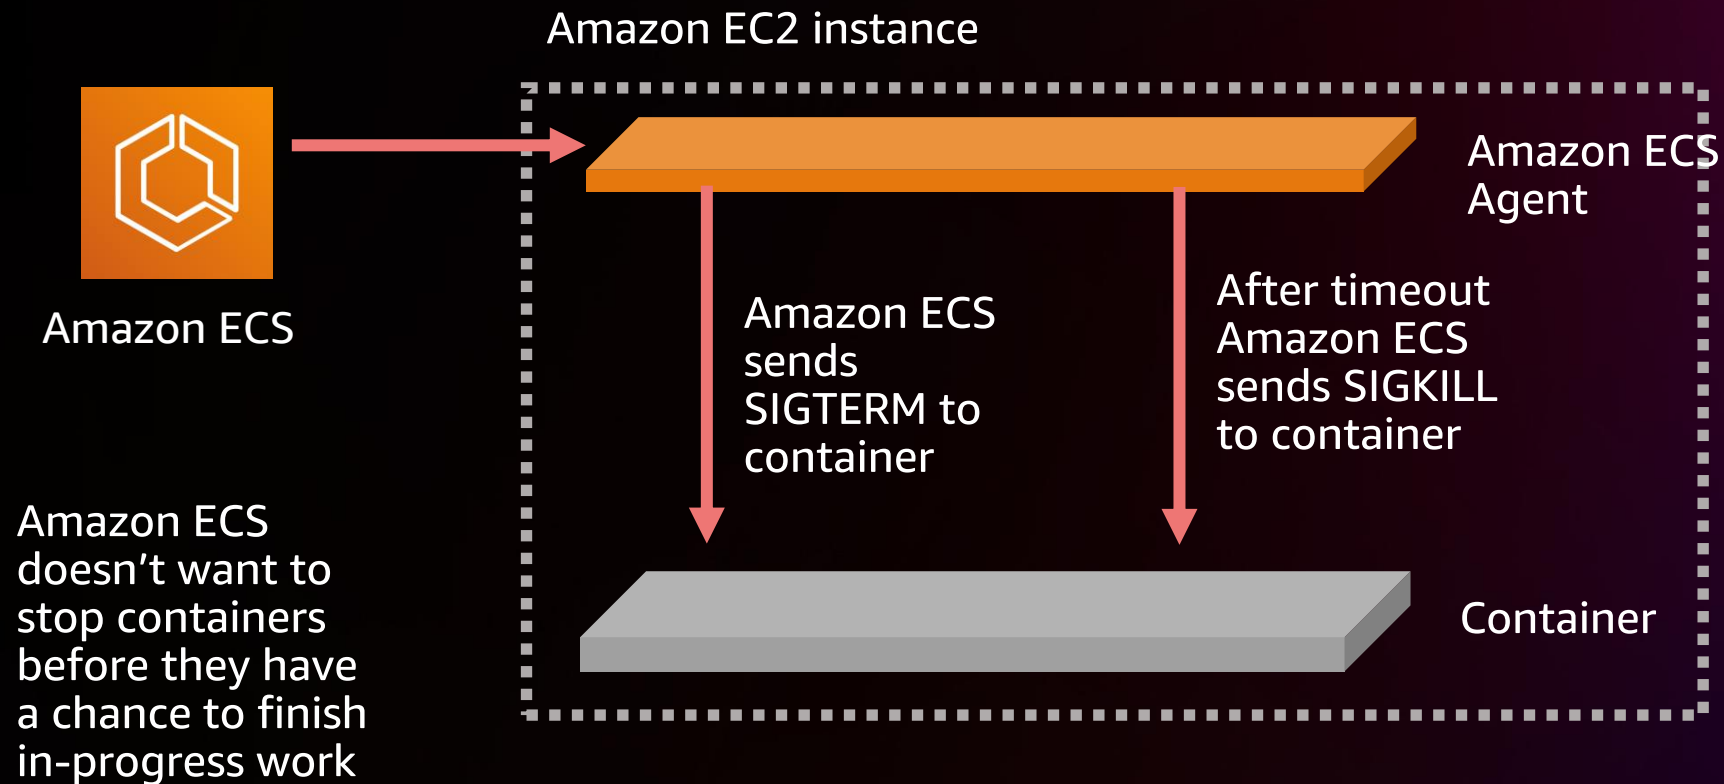
Amazon ECS waits for load balancers to drain connections

`deregistration_delay.timeout_seconds` (default = 300s)



Amazon ECS asks containers to exit nicely, before stopping them

Amazon `ECS_CONTAINER_STOP_TIMEOUT` (default = 30s)



Use container health checks for microservices that don't use a load balancer

Amazon ECS container health check allows you to run user-configured commands in shell

▼ Advanced container configuration

HEALTHCHECK

Command


Interval* second(s)

Timeout* second(s)

Start period* second(s)

Retries*


Other symptoms to detect bad deployments



Task Launch failures



Application errors



Abnormal CPU/
memory utilization

Amazon ECS deployment circuit breaker automatically monitors and rolls back bad deployments

Amazon ECS deployment circuit breaker monitors task launch and health check failures


If failures exceed threshold,
Amazon ECS rolls back or stops deployment

One-click setup



Best practices: Speed

Amazon ECS is continuously making deployments faster to enable faster innovation



Amazon ECS deployment performance

Benchmarks using Amazon ECS services	Size	Duration	Rate
AWS Fargate on-demand capacity, no load balancer	1000 tasks	208 seconds	4.8 tasks/sec
AWS Fargate spot capacity, no load balancer	1000 tasks	353 seconds	2.8 tasks/sec
AWS Fargate 50/50 on-demand and spot capacity	1000 tasks	213 seconds	4.7 tasks/sec
AWS Fargate on-demand, no public IP address	1000 tasks	199 seconds	5 tasks/sec
AWS Fargate on-demand, load balanced service	1000 tasks	252 seconds	4 tasks/sec
EC2 capacity, capacity provider starts empty	1000 tasks	1118 seconds	0.8 tasks/sec
EC2 capacity, with all the EC2 instances up already	1000 tasks	270 seconds	3.7 tasks/sec
EC2 capacity, host networking instead of AWS VPC	1000 tasks	270 seconds	3.7 tasks/sec

Full blog: <https://aws.amazon.com/blogs/containers/under-the-hood-amazon-elastic-container-service-and-aws-fargate-increase-task-launch-rates/>




3 key controls to optimize deployment speed

Load balancer configurations

Amazon ECS service configuration

Task launch time



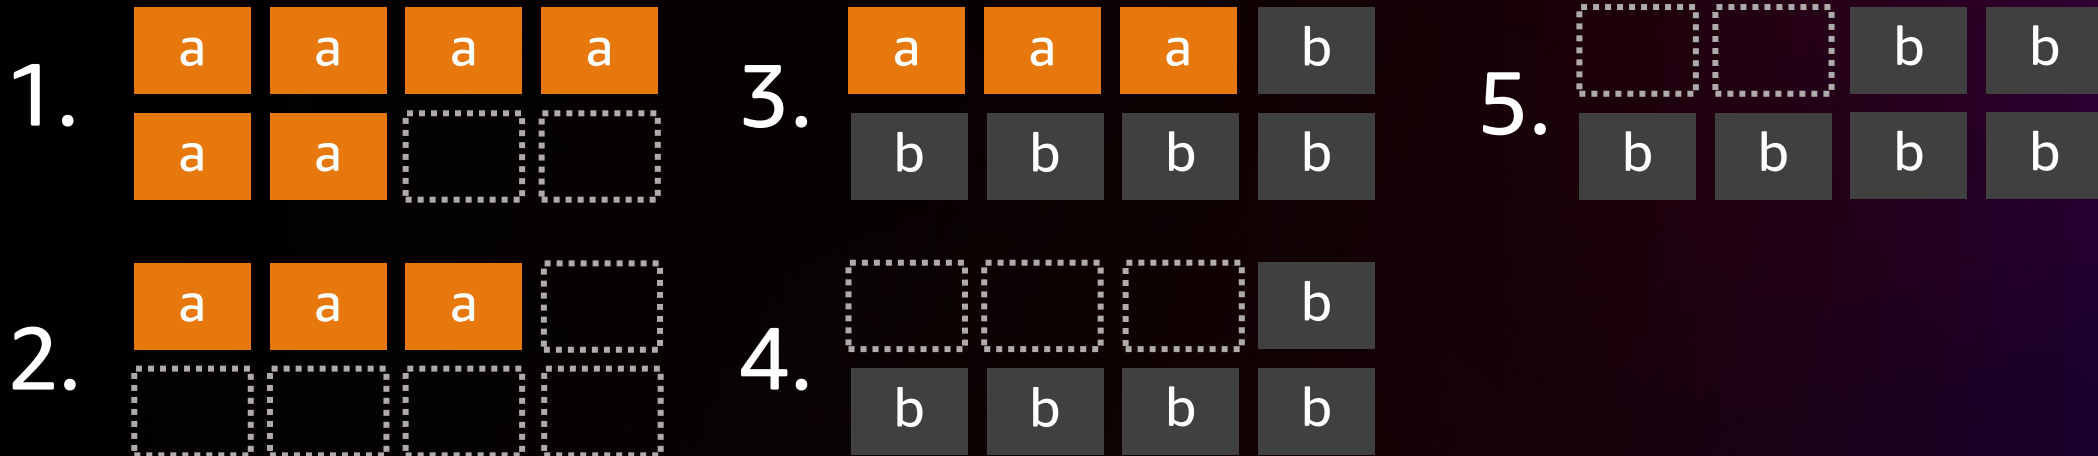
Important controls for Amazon ECS rolling deployment

minimumHealthyPercent – lower limit on the number of tasks in a service that must remain in the RUNNING

maximumPercent – upper limit on the number of tasks in a service that are allowed in the RUNNING or PENDING

These impact resources and time taken to roll out.

E.g. minimumHealthyPercent – 50% ; maximumHealthyPercent – 200%



Availability vs. speed

Availability ← → Speed

For production you want to maintain high availability and handle load

minimumHealthyPercent – 100%
maximumPercent – 200%

For a development environment optimize for speed of rolling a deploy

minimumHealthyPercent – 50%
maximumPercent – 200%

Load Balancer configurations impact deployment speed too

Registering new tasks

HealthCheckIntervalSeconds *(default 30s for ALB with Target as IP)*

HealthyThresholdCount *(default 5 for ALB)*

Terminating old tasks

deregistration_delay.timeout_seconds *(default = 300s)*

You can mutate these configurations to increase deployment speed

Speeding up task launches

Smaller container images

Use cached images on Amazon EC2

Amazon `ECS_IMAGE_PULL_BEHAVIOR` – once or prefer-cached

Compressed container images (zstd) for faster startup on Fargate **(NEW)**

Takeaways



Best practices for deployments

Configure dedicated health check endpoint for your load balanced Amazon ECS services (container health check if service doesn't use ELB)

Use circuit breaker as a default safeguard

Mutate Amazon ECS service and ELB configurations to manage deployment speed

CDK/Copilot/Blueprints for application-focused constructs

Pipelines for automation

Keep images small and/or use caching (EC2)/compression (Fargate)

Resources

Best practices to speed up Amazon ECS deployments: <https://bit.ly/3VrjEoG>

Automating hands off deployments: <https://go.aws/3EFwB7H>

Graceful shutdowns in Amazon ECS: <https://go.aws/3gFK6fz>

Faster startup times on Fargate with zstd compression: <https://go.aws/3F8XRwW>

Sample application: <https://bit.ly/3gIpUd6>



Thank you!

Uttara Sridhar
Uttara@amazon.com

Vibhav Agarwal
agvibhav@amazon.com
@VibhavAgarwal22



Please complete the session
survey in the **mobile app**

